
The InfoEdge Interview Companion

Classical ML, Deep Learning, SQL, Python (Round 3) and Case Studies (Round 4)

Every concept taught from zero, with small hand-worked examples and real Naukri, 99acres and Jeevansathi style use cases.

Interview questions answered the way a strong candidate answers them aloud: one-line answer first, then the full logic, a worked example, the follow-ups and the traps.

What this book covers (the InfoEdge Pre-Placement Talk 2025–26 syllabus, with SQL in Part IX):

Rounds 1 and 2, Classical ML and Deep Learning (66.66%): Bagging, Boosting, Gradient Descent, Decision Trees, Linear and Logistic Regression, Random Forests, Gradient Boosting Machines, XGBoost, LightGBM, SVM and the Kernel Trick, Regularisation, Bias–Variance Tradeoff, Metrics (Precision, Recall, F1, TPR . . .), Confusion Matrix, Missing Data Imputation, Matrix Factorisations, Density Estimation (KDE, K-Means . . .), MAP, MLE, Basic NLP, TF-IDF, Embeddings (Word2Vec, FastText, BERT), Architectures (LSTM, CNN, RNN, Transformers, GANs).

Round 3, Programming: Python, NumPy, Pandas: data reading, column naming, statistical parameters, group-by, basic visualisation, basic text processing; scikit-learn: train–test split, training a model, performance testing.

Round 4, Case Study: real world data science problems: data challenges, modelling and design choices, model assessment, thought process.

SQL (Part IX): four levels from zero (SELECT, NULL logic, joins, GROUP BY, window functions, CTEs) to interview patterns, leak-free feature tables, indexes and schema design; every query run on a small job-portal database.

Plus the “unconventional variants” InfoEdge rewards (GLMs, Poisson regression, SVR, quantile regression, kernel methods beyond SVM . . .) and basic GenAI/LLM knowledge mentioned in the 2026–27 talk.

How to use this book

What this book is, and how it fits with your OA book

You already have the *OA question bank* (675 short questions for the 40-question online test). That book trains speed. This book trains **understanding and speaking**: Rounds 1 and 2 are “drilling” rounds in which an interviewer asks three or four questions per topic and keeps asking “why?” until you either explain the idea from first principles or run out of road. Round 3 is a live Python notebook. Round 4 is an open case study. None of these can be passed by remembering formulas; all of them can be passed by *understanding*.

So every chapter has two halves.

1. **Teach.** The concept is built from nothing. No prerequisite is assumed: if a chapter needs a derivative, a dot product or a probability, it says in plain words what that is at the point it is needed. Every idea is followed by a *small example* you can do with a pen in under a minute, and an *in practice* box that places it in a Naukri, 99acres, Jeevansathi or Shiksha setting. Red boxes list the mistakes that interviewers fish for.
2. **Interview questions.** Each question is answered the way a strong candidate answers aloud:
 - **In one line** – the sentence you say first, so the interviewer knows you know.
 - **The full answer** – the logic, step by step, with the “because” behind every claim.
 - **Worked example** – numbers or a concrete scenario, so the idea is checkable.
 - **If the interviewer pushes** – the next two or three questions they will ask, already answered.
 - **Trap to avoid** – the wrong answer that sounds right.

The EDA Lab and its notebook

EDA carries heavy weight in the interviews, so Part VIII (Chapters 41–44) is a hands-on lab: a 12-step workflow, a reusable toolkit, 25 pandas/NumPy recipes and four complete case studies (MBA placement, HR job change of data scientists, Titanic, Pima diabetes) with real code, real printed output and plots. Every analysis is in the Colab-ready notebook `InfoEdge_Interview_Companion/eda_lab/EDA_Lab.ipynb`; type the cells yourself before the interview.

The SQL part and its notebook

SQL is one of the six topics drilled in Rounds 1 and 2. Part IX (Chapters 45–48) teaches it in four levels, from “what is a table” to gaps-and-islands, retention, funnels, leak-free feature tables and query plans. Every query runs on *MiniNaukri*, a seven-table job-portal database small enough to check by eye, and the output printed under each query is the real result. The Colab notebook `InfoEdge_Interview_Companion/sql_lab/SQL_Lab.ipynb` builds the database in Python’s built-in SQLite and runs every query, with practice exercises.

The tags

-  **Most asked** – (SQL part only) a standard SQL interview question. The InfoEdge candidate confirmed that SQL is asked in Rounds 1–2 but did not recall the exact questions, so these are *not* marked as PYQs.
-  **PYQ** – a question reported by InfoEdge candidates (online test, interview rounds, or the coding round). They come first in each chapter. InfoEdge does not publish papers, so the wording is a faithful reconstruction of what was reported.
-  **Think deeper** – a thought-provoking or “unconventional” question of the kind the InfoEdge hire in the experience video called out: *can an SVM do regression? what do you do with count data? is logistic regression a GLM? can a random forest extrapolate?* You asked for these to be at least 15% of all questions. This edition contains **184 such questions out of 717** in total, and **110 PYQs**.

At the foot of *every page* there is a “Stuck on this page?” link to one specific video for the topic on that page. Where a direct video link could not be verified, the link is a YouTube search for the exact video title, which lands on that video as the first result.

The InfoEdge selection process, from the Pre-Placement Talk

Stage	What is tested	Time
Pre-screening	CV shortlisting + online test: about 40 MCQs in 40 minutes to 1 hour, from ML, DL, programming and statistics. Programming MCQs are “quite easy” (for example converting a string to an int and spotting the right output).	40–60 min
Round 1	Technical: Maths & Statistics (33.33%) and Classical ML & Deep Learning (66.66%).	30 min
Round 2	Same split, deeper. Reported as the round that “tests your ability to explain the architecture of ML and DL models” on pen and paper.	45 min
Round 3	Programming. Part 1, data handling (Python, NumPy, Pandas): reading data, column naming, statistical parameters, group-by, basic visualisation, basic text processing. Part 2, ML basics (scikit-learn): train–test split, training a model, performance testing.	30–45 min
Round 4	Case study on a real-world data science problem: data challenges, modelling and design choices, model assessment, thought process and team fit. “There is generally no single best solution.”	45–60 min
HR	Company knowledge, why InfoEdge, work model.	15–20 min

What the 2026–27 talk added

- **Emphasis:** “a lot more emphasis will be given to ML and DL”; GenAI is evolving, so *basic* LLM knowledge is expected, but “if you don’t know topics like quantisation it will be fine”. The work is “largely ML based, with some knowledge of DL and GenAI”.
- **No DSA**, unless the discussion with a very strong candidate goes there.
- The interviewers want “a young data scientist who actually wants to understand business”.
- **A spam example was used to explain what they look for:** do you look at recall and not just accuracy, can you build the right dataframe, do you frame it properly as a classification or a regression problem. Chapter 11 and Chapter 27 are built around exactly this.
- **Products shown:** the Naukri job posting engine (100 crore+ applies per year), seeker acquisition and reactivation, apply generation, search and recommendation, the job–CV matching engine,

Resdex search (60K recruiters, 16 lakh searches and 30 lakh candidates a day) organised as *query intelligence* → *retrieval* + *ranking* → *personalisation*, PremiumX hybrid retrieval (dense + sparse, LLM ranking features, explainability), the SilverSkAI taxonomy (50 lakh raw skill strings normalised to 70K), AI REX agentic hiring (mandate, sourcing, screening, outreach, calibration agents), the Data Factory (400+ models; salary engine, company engine), Talent Pulse analytics, the content feed, Jeevansathi match recommendation (content, activity and collaborative models), and 99acres ML-based telecalling (user requirement, user-project matching, user ranking, then calls to high-intent leads). Every one of these appears as a case study in Part VI.

- **Tools named:** classifiers, ensembles, kernel methods, dimensionality reduction, feature learning, Bayesian inference; ANNs, RNNs/LSTMs, Transformers, CNNs, label attention networks, Siamese networks; PyTorch, Keras/TensorFlow, Spark; Transformers, GPT, BERT, open LLMs, TTS/ASR, fine-tuning, agentic AI.

What a recent hire reported (IIT BHU, InfoEdge DS)

- **Statistics:** p-value vs alpha. **Probability:** Naive Bayes built up from Bayes' theorem, with intuition rather than formulas.
- **ML:** PCA via SVD and the scree plot; t-SNE and perplexity; Fisher LDA (push class means apart, shrink within-class spread); a little on decision trees.
- **DL (the heaviest):** vanishing and exploding gradients; weight decay; how a network learns (backprop); preventing overfitting; Xavier vs He initialisation; time-series models (RNN, LSTM, GRU and their internals).
- **Linear algebra:** rank, eigen-things, null space, and a Markov chain question.
- **Round 3:** an MBA placement dataset. He spotted *target leakage* first (salary is only present for placed students), checked missing values, checked outliers with `describe()` and *did not* blindly remove genuine extreme marks, did value counts, a correlation matrix, logistic regression, a confusion matrix, precision, recall, F1 and accuracy.
- **Round 4:** features predicting whether a 99acres user will buy (price-band concentration, location of recent searches); linear regression assumptions and what to do under heteroscedasticity (he proposed a GLM; the interviewer expected transformations; both were accepted); "you are sales head of Naukri, give three ways to grow revenue, no clarifying questions".
- **Depth in basic ML:** count data needs Poisson regression (a GLM); SVMs can do regression (SVR).

How to answer aloud: the five-beat structure

Interviewers grade *structure* almost as much as content. For any "explain X" question:

1. **Define in one sentence** what X is and what problem it solves.
2. **Give the intuition** with a picture or an everyday analogy.
3. **Give the mechanism:** the formula or algorithm, and say what each symbol means.
4. **Give a tiny example**, ideally with numbers.
5. **Give the trade-off:** when it fails, and what you would use instead.

If you are not sure, say what you *do* know and reason aloud. "I have not used it, but from the name and from how X works I would expect. . . ." is a strong answer. Silence is the only bad answer.

A four-week plan

Week 1	Part I (Chapters 1–5) and Chapters 6–9. Say every “In one line” aloud.
Week 2	Chapters 10–17. Re-derive every worked example on paper without looking.
Week 3	Part IV (deep learning) and Chapter 23. Draw LSTM, Transformer and backprop diagrams from memory.
Week 4	Part V in a real notebook (type every snippet) and Part VI with a friend acting as interviewer.

Contents

How to use this book	ii
I Foundations of Machine Learning	1
1 What Learning Means: Generalisation, Bias–Variance and Regularisation	2
1.1 The problem machine learning solves	2
1.2 Generalisation: the only thing that matters	3
1.3 Underfitting, overfitting and model capacity	4
1.4 The bias–variance decomposition, built slowly	5
1.5 Regularisation: paying for complexity	7
1.6 No free lunch, and why baselines matter	8
1.7 Interview questions	8
2 Linear Regression, from a Line to the Normal Equations	18
2.1 The model in words and symbols	18
2.2 Choosing the line: least squares	19
2.3 Measuring the fit	20
2.4 Assumptions: what OLS needs, and for what	20
2.5 Regularised linear regression	22
2.6 Outliers, leverage and influence	22
2.7 Interview questions	22
3 Logistic Regression and Multiclass Classification	31
3.1 Why not just use a straight line for yes/no?	31
3.2 The sigmoid: squeezing a score into a probability	31
3.3 Odds and log-odds: how to read the coefficients	31
3.4 The decision boundary is a straight line	32
3.5 Training: maximum likelihood and log loss	32
3.6 Regularisation and separable data	33
3.7 Probabilities and calibration	33
3.8 More than two classes	33
3.9 Multi-label is different	34
3.10 Interview questions	35

4	Generalised Linear Models and the Regression Family	41
4.1	The question GLMs answer	41
4.2	The three ingredients of a GLM	41
4.3	Poisson regression, built from scratch	42
4.4	Gamma regression: positive, skewed, multiplicative noise	43
4.5	Tweedie: when many values are exactly zero	43
4.6	Beyond the mean: other regression losses	44
4.7	Choosing a regression model: a decision guide	45
4.8	Interview questions	46
5	Gradient Descent and Optimisers	53
5.1	Walking downhill in fog	53
5.2	The learning rate	53
5.3	Batch, stochastic and mini-batch	54
5.4	Convexity, local minima and saddle points	54
5.5	Momentum, RMSProp and Adam, built up step by step	55
5.6	Second-order methods	56
5.7	Practical checklist when training is not going well	56
5.8	Interview questions	57
II	Classical Models	64
6	Decision Trees	65
6.1	A tree is a game of twenty questions	65
6.2	How to choose a question: impurity	65
6.3	The algorithm (pseudo-code)	67
6.4	The classic algorithms	67
6.5	Overfitting and how to stop it	68
6.6	Strengths and weaknesses	68
6.7	Feature importance	68
6.8	Interview questions	69
7	Ensembles, Bagging and Random Forests	75
7.1	Wisdom of crowds	75
7.2	Three families of ensembles	75
7.3	Why averaging reduces variance: the formula	75
7.4	Majority voting: the maths of the crowd	76
7.5	Bootstrap sampling	76

7.6	Random forest = bagged trees + random feature subsets	76
7.7	Out-of-bag error: free validation	77
7.8	Feature importance in a forest	77
7.9	Variants and practical notes	77
7.10	Random forest vs gradient boosting	78
7.11	Stacking done right	78
7.12	Interview questions	78
8	Boosting: AdaBoost, Gradient Boosting, XGBoost and LightGBM	85
8.1	The idea: learn from your mistakes, a little at a time	85
8.2	AdaBoost (adaptive boosting)	85
8.3	Gradient boosting: gradient descent on predictions	86
8.4	XGBoost: regularised, second-order, engineered	87
8.5	LightGBM: faster on big data	88
8.6	Learning to rank with boosted trees	88
8.7	Interview questions	89
9	Support Vector Machines and the Kernel Trick	95
9.1	Which line is best?	95
9.2	The geometry you need	95
9.3	Hard margin: the optimisation problem	96
9.4	Soft margin: allowing mistakes, and the C parameter	96
9.5	The dual problem: where kernels come from	97
9.6	The kernel trick	97
9.7	Practicalities	98
9.8	Interview questions	99
10	Three Classic Classifiers: k-NN, Naive Bayes and Fisher's LDA	104
10.1	k -nearest neighbours	104
10.2	Naive Bayes, built from Bayes' theorem	105
10.3	Fisher's Linear Discriminant Analysis	106
10.4	Interview questions	108
11	Evaluation Metrics and the Confusion Matrix	113
11.1	The confusion matrix	113
11.2	The metrics, each answering one question	113
11.3	Thresholds and threshold-free metrics	114
11.4	Multiclass averaging	115

11.5 Calibration metrics	115
11.6 Regression and ranking metrics (summary)	116
11.7 Offline metrics vs business metrics	116
11.8 Interview questions	116
12 Maximum Likelihood (MLE) and Maximum a Posteriori (MAP)	122
12.1 Probability vs likelihood: flipping the question	122
12.2 Maximum likelihood estimation: the recipe	122
12.3 Why ML losses are negative log-likelihoods	123
12.4 MAP: adding prior beliefs	123
12.5 MLE vs MAP vs full Bayesian inference	125
12.6 Properties of MLE worth quoting	125
12.7 When there is no closed form: EM	125
12.8 Interview questions	126
13 Missing Data Imputation, Outliers and Class Imbalance	131
13.1 Why missing values are a modelling problem	131
13.2 Three missingness mechanisms (Rubin)	131
13.3 Option 1: deletion	131
13.4 Option 2: simple (univariate) imputation	132
13.5 Option 3: multivariate (model-based) imputation	132
13.6 Option 4: let the model handle it	133
13.7 The missing-indicator trick	133
13.8 Leakage: fit imputers on training data only	133
13.9 Outliers: detect, understand, then decide	133
13.10 Class imbalance	134
13.11 Interview questions	135
14 Matrix Factorisations: SVD, PCA, NMF, Recommenders and t-SNE	141
14.1 What “factorising a matrix” means	141
14.2 A quick refresher: vectors, matrices, eigenvectors	141
14.3 SVD: the master factorisation	141
14.4 PCA is SVD of the centred data	142
14.5 NMF: non-negative parts	143
14.6 Matrix factorisation for recommendations	143
14.7 t-SNE and UMAP: seeing clusters	144
14.8 Interview questions	145

15 Density Estimation and Clustering: Histograms, KDE, K-Means, GMM and EM	151
15.1 What is density estimation?	151
15.2 Histograms	151
15.3 Kernel Density Estimation (KDE)	151
15.4 K-Means clustering	152
15.5 Gaussian Mixture Models and EM	153
15.6 Other clustering methods worth knowing	154
15.7 Interview questions	155
 III Natural Language Processing	 159
16 Basic NLP: Preprocessing, Bag of Words and TF-IDF	160
16.1 Why text needs special treatment	160
16.2 Tokenisation	160
16.3 Normalisation	161
16.4 Bag of Words (BoW)	161
16.5 TF-IDF: down-weighting words everybody uses	161
16.6 Classic text classifiers	162
16.7 Other basic NLP vocabulary	162
16.8 Interview questions	163
 17 Embeddings: Word2Vec, GloVe, FastText and BERT	 168
17.1 From one-hot to dense vectors	168
17.2 Word2Vec	168
17.3 GloVe	169
17.4 FastText: embeddings from character n-grams	170
17.5 Static vs contextual embeddings	170
17.6 BERT	170
17.7 Sentence and document embeddings	171
17.8 Interview questions	171
 IV Deep Learning	 177
18 Neural Network Foundations: Backpropagation, Activations, Initialisation and Regularisation	178
18.1 From a neuron to a network	178
18.2 Activation functions	179
18.3 Loss functions	179

18.4 Training: forward pass, loss, backward pass	179
18.5 Vanishing and exploding gradients	180
18.6 Weight initialisation: Xavier and He	181
18.7 Residual connections (ResNet)	181
18.8 Normalisation layers	181
18.9 Regularisation in deep learning	182
18.10 Learning rate and getting unstuck	182
18.11 RBMs and Deep Belief Networks (historical)	182
18.12 Interview questions	182
19 Convolutional Neural Networks	188
19.1 Why not a dense network on images?	188
19.2 The convolution operation, by hand	188
19.3 Output size: the formula	189
19.4 Parameters and computation	189
19.5 Pooling	189
19.6 Invariance vs equivariance	189
19.7 Receptive field and the feature hierarchy	190
19.8 Landmark architectures	190
19.9 Transfer learning	190
19.10 1-D CNNs	191
19.11 Interview questions	191
20 Recurrent Networks: RNN, LSTM, GRU and Sequence Models	196
20.1 Why sequences need memory	196
20.2 The vanilla RNN	196
20.3 Backpropagation through time (BPTT)	197
20.4 The vanishing/exploding problem in RNNs	197
20.5 LSTM: a memory with gates	197
20.6 GRU: a lighter LSTM	198
20.7 Variants and training tricks	198
20.8 Time series with deep learning	199
20.9 Why Transformers replaced RNNs	199
20.10 Interview questions	199
21 Transformers and Attention	204
21.1 The idea of attention, in plain words	204
21.2 Scaled dot-product attention	204

21.3 Multi-head attention	205
21.4 The Transformer block	205
21.5 Positional information	205
21.6 Masks	206
21.7 Three families	206
21.8 Cost and efficiency	206
21.9 Training and generation	206
21.10 Transformers beyond text	206
21.11 Interview questions	207
22 Generative Models: GANs, Autoencoders and VAEs	212
22.1 Discriminative vs generative, once more	212
22.2 GANs: the forger and the detective	212
22.3 Why GANs are hard to train	213
22.4 Important variants	214
22.5 Evaluating generative models	214
22.6 Autoencoders	214
22.7 Variational autoencoders (VAEs)	214
22.8 GANs vs VAEs vs diffusion vs autoregressive	215
22.9 Interview questions	215
23 GenAI Basics: LLMs, Fine-tuning, RAG and Agents	219
23.1 What a large language model is	219
23.2 Ways to adapt an LLM to a task	220
23.3 Retrieval-augmented generation (RAG)	220
23.4 Agents	220
23.5 Making LLMs cheaper and faster (basic awareness)	221
23.6 Speech: ASR and TTS	221
23.7 Risks and evaluation	221
23.8 Interview questions	222
V Round 3: Programming in Python	225
24 Python for Data Science: the Core Language	226
24.1 Values, types and variables	226
24.2 Operators	226
24.3 Strings	227
24.4 Lists, tuples, dicts and sets	227

24.5 Mutability, references and copying	228
24.6 Control flow and comprehensions	228
24.7 Functions: *args, **kwargs, lambda, scope	228
24.8 Iterators and generators	229
24.9 Decorators	229
24.10 Exceptions	229
24.11 Classes in one page	230
24.12 Performance notes	230
24.13 Interview and OA questions	230
25 NumPy: Arrays, Vectorisation and Broadcasting	235
25.1 Why NumPy	235
25.2 Shape, axes and creation	235
25.3 Indexing	235
25.4 Broadcasting	236
25.5 Useful functions	236
25.6 Linear algebra	237
25.7 Missing values and randomness	237
25.8 Interview and OA questions	237
26 Pandas and Visualisation: Data Handling for Round 3	241
26.1 The two objects: Series and DataFrame	241
26.2 Reading and first look	241
26.3 Column naming	242
26.4 Selecting and filtering	242
26.5 Statistical parameters	242
26.6 Creating and transforming columns	242
26.7 Group-by: split, apply, combine	243
26.8 Missing values	243
26.9 Combining tables	244
26.10 Basic text processing with the .str accessor	244
26.11 Dates	244
26.12 Basic visualisation	244
26.13 Performance	245
26.14 Round 3 style questions	245

27 Scikit-learn and the Complete Round 3 Walkthrough	250
27.1 The scikit-learn API in one idea	250
27.2 Splitting	250
27.3 Preprocessing with Pipeline and ColumnTransformer	250
27.4 Models you should be able to call	251
27.5 Performance testing	251
27.6 The complete Round 3 walkthrough	251
27.7 Regression and other tasks in scikit-learn	254
27.8 Round 3 style questions	254
 VI Round 4: Case Studies and the Final Conversation	 258
28 How to Answer a Data Science Case Study	259
28.1 The ten-step skeleton	259
28.2 Feature engineering: a generator of ideas	260
28.3 Metrics: from business to offline	260
28.4 A/B testing in one page	261
28.5 Product and business questions	261
28.6 Communication and team fit	261
28.7 Framework questions	261
 29 Worked Case Studies from InfoEdge's Own Products	 264
 30 The Final Conversation: Resume Drill, Company Knowledge and HR	 272
30.1 Everything on your resume will be drilled	272
30.2 Company knowledge	272
30.3 "Why InfoEdge?"	272
30.4 Questions to ask them	273
30.5 Questions	273
 VII The InfoEdge Hire's Checklist: EDA, Extra ML/DL, Time Series, Probability, Statistics and Linear Algebra	 274
 Part VII Guide: Checklist Coverage Map	 275
 31 EDA I: Univariate Analysis, Descriptive Statistics and Reading Graphs	 277
31.1 Why we do EDA at all	277
31.2 Types of variables	277

31.3 Measures of central tendency: where is the middle?	277
31.4 Measures of spread: how scattered?	278
31.5 Shape: skewness and kurtosis	279
31.6 Reading graphs: what each plot answers	280
31.7 Interview questions	282
32 EDA II: Bivariate and Multivariate Analysis, Correlations and Multicollinearity	286
32.1 Which measure for which pair of variable types?	286
32.2 Covariance	286
32.3 Pearson correlation	286
32.4 Spearman rank correlation	287
32.5 Point-biserial correlation (numeric vs binary)	287
32.6 Phi coefficient (binary vs binary)	288
32.7 Numeric vs multi-class categorical: ANOVA and η^2	288
32.8 Reading a correlation matrix	288
32.9 Simpson's paradox: when aggregate correlation lies	289
32.10 Multicollinearity	289
32.11 Interview questions	290
33 EDA III: Outliers, Missing Values and Leaky Variables, in Depth	294
33.1 The normal curve and the 68–95–99.7 rule	294
33.2 z-scores	294
33.3 IQR rule and box plots	295
33.4 Multivariate and model-based detection	295
33.5 What to do with an outlier: a decision procedure	295
33.6 Missing values: the full picture	296
33.7 Leaky variables	297
33.8 Interview questions	298
34 EDA IV: Scaling, Transforming Skew, Binning, Encoding and Imbalance	301
34.1 Why scale at all?	301
34.2 The scalers, compared on one column	301
34.3 Curing skew: log, power, Box–Cox, Yeo–Johnson	302
34.4 Numerical to categorical: binning	303
34.5 Categorical to numerical: encodings	303
34.6 Imbalanced datasets: the complete toolkit	305
34.7 Interview questions	306

35 Clustering Variants and the Dimensionality-Reduction Family	309
35.1 The K-Means family	309
35.2 Hard vs soft clustering	310
35.3 The Swiss roll: why “non-linear” matters	310
35.4 SNE vs t-SNE	311
35.5 A taxonomy of dimensionality reduction	312
35.6 Interview questions	313
36 The Perceptron and the Four Application Areas: CV, NLP, RL and Federated Learning	316
36.1 The perceptron	316
36.2 Computer vision basics	317
36.3 NLP basics (a map of Chapters 16, 17 and 21)	317
36.4 Reinforcement learning basics	317
36.5 Federated learning basics	318
36.6 Interview questions	319
37 Time Series: Components, Stationarity, ARIMA, SARIMAX and ANOVA	322
37.1 What makes time series different	322
37.2 Components	322
37.3 Stationarity	322
37.4 Autocorrelation: ACF and PACF	323
37.5 The ARIMA family, built up	323
37.6 Exponential smoothing	324
37.7 Evaluating forecasts	324
37.8 ANOVA (analysis of variance)	324
37.9 Interview questions	325
38 Probability from Zero: Events, Random Variables, Moments and Distributions	328
38.1 Experiments, outcomes and events	328
38.2 Conditional probability, independence and Bayes	328
38.3 Random variables	329
38.4 Expectation, variance and moments	330
38.5 The common distributions	330
38.6 Interview questions	333
39 Statistics for Data Scientists: Sampling, CLT, Confidence Intervals and Hypothesis Testing	336
39.1 Descriptive vs inferential statistics	336
39.2 Population, sample, parameter, statistic	336

39.3	Sampling distributions and the standard error	336
39.4	The central limit theorem	337
39.5	Confidence intervals	337
39.6	Hypothesis testing, step by step	337
39.7	The common tests	338
39.8	Multiple testing	339
39.9	Interview questions	339
40	Linear Algebra for ML: Vectors, Matrices, Rank, Eigenvectors and Markov Chains	343
40.1	Vectors	343
40.2	Matrices as transformations	343
40.3	Systems of equations	344
40.4	Linear independence, span, rank and null space	344
40.5	Eigenvalues and eigenvectors	345
40.6	Markov chains: eigenvectors in action	345
40.7	SVD in one paragraph (recap)	346
40.8	Interview questions	346
VIII	EDA Lab: Hands-on Exploratory Data Analysis on Real Datasets (with a Colab Notebook)	350
41	EDA Lab I: The EDA Playbook in Code	351
41.1	The 12-step EDA workflow	351
41.2	The reusable toolkit	351
41.3	The cookbook: 25 recipes, run on real data	353
41.4	Plot recipes (Matplotlib/Seaborn)	356
41.5	How to narrate EDA in the interview	357
41.6	Interview questions	357
42	EDA Lab II: The MBA Placement Dataset, End to End	360
42.1	The data dictionary	360
42.2	Step 1: load and look	360
42.3	Step 2: data quality	361
42.4	Step 3: the target	361
42.5	Step 4: salary, leakage and “if the package is given, we know the outcome”	361
42.6	Step 5: univariate analysis of the scores	363
42.7	Step 6: which scores separate placed from not placed?	363
42.8	Step 7: thresholds a placement cell can act on	364

42.9	Step 8: categorical features	365
42.10	Step 9: correlations and multicollinearity	366
42.11	Step 10: multivariate view and a sanity-check model	366
42.12	Step 11: the summary you say at the end	367
42.13	Interview questions	367
43	EDA Lab III: Will a Data Scientist Change Jobs? (Real HR Data)	370
43.1	Step 1: load, overview, target	370
43.2	Step 2: types and data quality	370
43.3	Step 3: missing values, with the target	371
43.4	Step 4: numeric features	372
43.5	Step 5: categorical features	373
43.6	Step 6: the high-cardinality city	373
43.7	Step 7: experience and an interaction	374
43.8	Step 8: findings and the modelling plan	375
43.9	Interview questions	375
44	EDA Lab IV: Titanic (Interactions and Text Features) and Pima Diabetes (Disguised Missing Values)	377
44.1	Titanic: who survived?	377
44.2	Pima Indians diabetes: missing values in disguise	379
44.3	Interview questions	381
IX	SQL for Data Scientists: From Zero to Interview Level (with a Colab Notebook)	383
45	SQL Level 1: Tables, Rows and Reading Data from One Table	384
45.1	What a database is, from zero	384
45.2	The practice database: MiniNaukri	385
45.3	SELECT and FROM: choosing columns	386
45.4	WHERE: choosing rows	387
45.5	NULL and three-valued logic	389
45.6	ORDER BY, LIMIT and OFFSET: sorting and taking the top	390
45.7	DISTINCT: unique values	391
45.8	Computed columns, aliases and integer division	392
45.9	CASE WHEN: if-else inside a query	393
45.10	String and date functions	393
45.11	The logical order of a query	394
45.12	Interview questions: Level 1	395

46 SQL Level 2: Aggregation, GROUP BY and Joins	401
46.1 Aggregate functions: many rows in, one number out	401
46.2 GROUP BY: one result row per group	402
46.3 HAVING: filtering groups	403
46.4 Conditional aggregation: the pattern behind every rate	403
46.5 Joins: combining tables	404
46.6 Set operations: UNION, UNION ALL, INTERSECT, EXCEPT	410
46.7 Interview questions: Level 2	411
47 SQL Level 3: Subqueries, CTEs, Window Functions and Changing Data	418
47.1 Subqueries: a query inside a query	418
47.2 CTEs: naming the steps of a query	420
47.3 Window functions: aggregates that keep every row	422
47.4 Changing data: INSERT, UPDATE, DELETE and DDL	427
47.5 Interview questions: Level 3	428
48 SQL Level 4: Interview Patterns, SQL for Machine Learning, Performance and Design	434
48.1 Nth highest value	434
48.2 Top N per group	435
48.3 Duplicates: find, inspect, delete	436
48.4 Self-join classics: employee vs manager	437
48.5 Consecutive days: the gaps-and-islands trick	437
48.6 Daily, rolling and cohort metrics	438
48.7 Funnels and A/B tests	440
48.8 Other patterns worth knowing	441
48.9 SQL for machine learning: feature tables without leakage	444
48.10 Performance: indexes and query plans	446
48.11 Schema design: normalisation, OLTP vs OLAP, star schema	447
48.12 SQL ↔ pandas translation	448
48.13 Interview questions: Level 4	448
Last-Week Checklist	455

Part I

Foundations of Machine Learning

CHAPTER 1

What Learning Means: Generalisation, Bias–Variance and Regularisation

🎥 **Companion video for this chapter:** [StatQuest: Machine Learning Fundamentals, Bias and Variance](#)

Almost every interview question in Rounds 1 and 2 is, underneath, a question about one idea: *a model must do well on data it has never seen*. This chapter builds that idea from nothing. Every later chapter (trees, boosting, SVMs, neural networks) is a different answer to the same problem, so time spent here pays off everywhere.

1.1 The problem machine learning solves

Suppose Naukri wants to show each job seeker an estimate: “jobs like this usually pay about Rs. X lakh”. Nobody can write that rule by hand; it depends on title, city, experience, company, skills and a hundred other things. But Naukri has *examples*: millions of past jobs with their salaries. Machine learning is the craft of turning examples into a rule.

We use a few words for this throughout the book.

- An **example** (also: row, sample, observation) is one item, such as one job posting.
- A **feature** (also: input, variable, column, x) is one measured property of it: years of experience required, city, number of skills listed.
- The **target** (also: label, output, y) is what we want to predict: the salary.
- A **model** is a rule f that takes features and returns a prediction $\hat{y} = f(x)$. The hat on $\hat{y}$ always means “predicted”.
- **Parameters** are the numbers inside the model that training adjusts (for example the slope of a line). **Hyperparameters** are the settings *we* choose before training (for example how deep a tree may grow).
- A **loss** measures how wrong one prediction is. Training means adjusting the parameters to make the average loss small on the examples we have.

📊 Small example: the smallest possible model

Three past jobs: experience $x = 1, 2, 3$ years, salary $y = 4, 6, 8$ lakh. Try the rule “salary = $2x + 2$ ”. Predictions: 4, 6, 8. Every error is zero, so the average loss is zero. A new job asking for $x = 5$ years gets the prediction $2 \cdot 5 + 2 = 12$ lakh. That last step, predicting for an x we have not seen, is the whole point. Fitting the three known points was only a means to it.

Two kinds of problem cover most of the syllabus.

- **Regression**: the target is a number (salary, property price, days to fill a vacancy).
- **Classification**: the target is a category (will the seeker apply: yes/no; which of 500 job categories does this title belong to; is this message spam).

There is also **unsupervised** learning, where there is no target at all and we look for structure (clusters

of similar candidates, a compressed representation of resumes). And **ranking**, where we order items (which 20 of 30 lakh candidates should a recruiter see first); ranking is central at InfoEdge because search and recommendation are its core products.

1.2 Generalisation: the only thing that matters

1.2.1 Training error versus test error

The **training error** is the average loss on the examples used to fit the model. The **test error** (or generalisation error) is the average loss on *new* examples from the same source. We care only about the second; the first is easy to make tiny and tells us little.

💡 Intuition: the student who memorised the answer key

A student who memorises last year's answer key scores 100% on last year's paper (training error zero) and fails this year's paper (test error high). A student who understood the chapter scores 85% on both. We want the second student. Memorising is called **overfitting**; understanding is **generalising**.

1.2.2 Why we split data

Because we cannot see the future, we *simulate* it. We hide part of our data, fit on the rest, and measure error on the hidden part.

- **Training set** (typically 60–80%): used to fit parameters.
- **Validation set** (10–20%): used to choose hyperparameters and compare models.
- **Test set** (10–20%): touched *once*, at the end, to report an honest number.

Why three and not two? Every time you look at a set and make a choice because of it, that set stops being “unseen”. If you try 50 settings and keep the one with the best validation score, that score is optimistic: you picked the luckiest of 50. The untouched test set corrects for that.

📊 Small example: why the best of many is optimistic

Imagine 50 models that are all *exactly* 80% accurate in truth. On a validation set of 200 rows, each one's measured accuracy wobbles by chance, roughly $\pm 2.8\%$ (the standard error $\sqrt{0.8 \cdot 0.2 / 200} \approx 0.028$). The best of 50 such wobbles is usually about 2 standard errors above the truth, so you would report about 85–86% for a model that is really 80%. A fresh test set would show 80%.

1.2.3 Cross-validation

When data are scarce, a single validation split is noisy. **k-fold cross-validation** splits the training data into k equal parts (folds). Train on $k - 1$ folds, validate on the remaining one, repeat k times so every fold is the validation fold once, and average the k scores.

Small example: 5-fold CV on 1,000 rows

Each fold has 200 rows. Round 1: train on rows 201–1000, validate on 1–200. Round 2: train on 1–200 and 401–1000, validate on 201–400, and so on. You get 5 scores, say 0.81, 0.79, 0.83, 0.80, 0.82. Report the mean 0.81 and the spread (standard deviation about 0.016). Every row was used for validation exactly once and for training four times.

Variants you should be able to name:

- **Stratified k -fold**: each fold keeps the class proportions (essential when only 2% are positive; otherwise a fold may contain no positives at all).
- **Group k -fold**: all rows of one group (one user, one company) stay in the same fold. If one recruiter posted 300 jobs, splitting their jobs across train and validation lets the model “recognise the recruiter” and look better than it is.
- **Time-series split**: always train on the past and validate on the future. Random splits on time data let the model peek at tomorrow.
- **Leave-one-out**: $k = n$. Nearly unbiased but expensive and high-variance.

1.2.4 Data leakage

Leakage is when information that would not be available at prediction time sneaks into training. It produces models that look brilliant offline and fail in production. It is the single most common reason real projects fail, which is why the InfoEdge coding round rewards spotting it first.

In practice: the reported Round 3 leakage


An MBA placement dataset had a salary column that was filled only for placed students. If you keep it, the model learns “salary present $\Rightarrow$ placed” and scores near 100%. At prediction time (before placement) salary does not exist. The candidate who said “drop salary, it leaks the target” before writing any model code impressed the panel.

Common leakage patterns:

- **Target leakage**: a feature that is a consequence of the target (salary after placement; “number of interview calls” when predicting “got the job”).
- **Preprocessing leakage**: computing the mean for imputation, the scaler, TF-IDF vocabulary or PCA on the *whole* dataset before splitting. The validation rows then influenced training.
- **Temporal leakage**: using features computed with future data (“total applications this job will receive”).
- **Duplicate leakage**: the same resume uploaded twice landing in both train and test.

1.3 Underfitting, overfitting and model capacity

Capacity (or complexity) is how flexible a model is: how many different shapes it can bend into. A straight line has low capacity; a degree-15 polynomial or a deep tree has high capacity.



- **Underfitting:** training error is high and test error is about as high. The model is too simple to capture the pattern (a straight line through data that curve).
- **Overfitting:** training error is low but test error is much higher. The model learned the noise in the particular sample.
- Training error falls steadily as complexity rises. Test error first falls, then rises: a U shape. The bottom of the U is what we tune for.

Small example: reading the gap

Model	Train acc.	Val. acc.	Diagnosis
A	62%	61%	underfit (both low, small gap)
B	99%	72%	overfit (big gap)
C	88%	86%	good (high and small gap)

1.4 The bias–variance decomposition, built slowly

1.4.1 An experiment in your head

Imagine you could collect *many* different training sets of the same size from the same world, and train the same kind of model on each. For one fixed test input x (say a job with 5 years' experience in Pune), you would get many different predictions, one per training set. Look at that cloud of predictions.

- **Bias** is how far the *centre* of the cloud is from the truth. A model that is systematically wrong (a straight line for a curved truth) has high bias, however much data it gets.
- **Variance** is how *spread out* the cloud is. A model that changes a lot when a few training rows change (a deep tree) has high variance.
- **Noise** is the randomness in y itself that no model can predict: two identical jobs can pay differently because of negotiation.

Intuition: the dartboard

Four archers. Low bias, low variance: tight cluster on the bullseye. High bias, low variance: tight cluster, but off-centre (consistently wrong). Low bias, high variance: centred on average but scattered. High bias, high variance: scattered and off-centre. Training many models on many datasets is throwing many darts.

1.4.2 The formula, and what each piece means

For squared error, the expected test error at a point x splits exactly into three parts:

$$\mathbb{E}[(y - \hat{f}(x))^2] = \underbrace{(\mathbb{E}[\hat{f}(x)] - f(x))^2}_{\text{bias}^2} + \underbrace{\mathbb{E}[(\hat{f}(x) - \mathbb{E}[\hat{f}(x)])^2]}_{\text{variance}} + \underbrace{\sigma^2}_{\text{noise}}.$$

Here $f(x)$ is the true average target at x , $\hat{f}(x)$ is our model's prediction (random, because it depends on the random training set), $\mathbb{E}[\cdot]$ means "average over many training sets", and σ^2 is the noise variance.

Small example: computing the three parts

At x = "5 years, Pune" the true average salary is $f(x) = 10$ lakh and the noise variance is $\sigma^2 = 1$. Over many training sets our model predicts 9, 11, 13, 11 (four sets, for simplicity). Mean prediction = 11, so bias = $11 - 10 = 1$ and $\text{bias}^2 = 1$. Variance = $\frac{(9-11)^2 + 0 + (13-11)^2 + 0}{4} = \frac{8}{4} = 2$. Expected error = $1 + 2 + 1 = 4$, i.e. a typical miss of $\sqrt{4} = 2$ lakh.

1.4.3 Why there is a trade-off

Making a model more flexible lets its average move closer to the truth (bias falls) but also lets it chase the quirks of each sample (variance rises). Making it more rigid does the opposite. Total error is minimised somewhere in between. More *data* is the one lever that lowers variance without raising bias, which is why "get more data" is the right answer to an overfitting model and the wrong answer to an underfitting one.

	High bias (underfit)	High variance (overfit)
Symptom	train and val error both high, close together	train error low, val error much higher
Typical models	linear model on curved data; very shallow tree; k -NN with huge k	deep tree; k -NN with $k = 1$; huge network on little data
Fixes	more features, more flexible model, less regularisation, train longer	more data, regularisation, simpler model, bagging, dropout, early stopping, data augmentation

1.4.4 Learning curves: the diagnostic picture

Plot training and validation error against the *number of training examples*.

- High bias: both curves flatten quickly at a high error and nearly meet. More data will not help.
- High variance: a large gap that narrows slowly as data grow. More data will help.

A second curve InfoEdge reportedly asks about is error against *training iterations* for a neural network: training error keeps falling; validation error falls, bottoms out, then rises. Stopping at the bottom is **early stopping**.

1.4.5 A modern footnote: double descent

Very large models (modern deep nets) can show a second descent: test error rises near the point where the model can just fit the training data exactly, then *falls again* as the model grows far beyond it. The classical U-curve is still the right first answer in an interview; mention double descent as a nuance if asked about deep learning.

1.5 Regularisation: paying for complexity

1.5.1 The idea

Regularisation is anything that discourages a model from becoming more complex than the data justify. The most common form adds a penalty to the loss:

$$\text{minimise} \quad \underbrace{\text{average loss on training data}}_{\text{fit the data}} + \lambda \cdot \underbrace{\text{complexity(parameters)}}_{\text{stay simple}}.$$

$\lambda \geq 0$ (lambda) is a hyperparameter: $\lambda = 0$ means no penalty; large λ means simplicity matters a lot.

💡 Intuition: the budget

Think of every weight as costing money. A feature must “earn” its weight by reducing the loss more than it costs. Weak, noisy features cannot afford a large weight, so their influence shrinks. That is exactly how regularisation cuts variance.

1.5.2 L2 (Ridge) and L1 (Lasso)

For a model with weights $w_1, \dots, w_d$:


- **L2 / Ridge / weight decay:** penalty $\lambda \sum_j w_j^2$. Shrinks all weights smoothly toward zero but rarely makes any exactly zero.
- **L1 / Lasso:** penalty $\lambda \sum_j |w_j|$. Pushes many weights to *exactly* zero, so it also selects features.
- **Elastic net:** both. Useful when features come in correlated groups (Lasso alone picks one of a group arbitrarily).

📊 Small example: one weight, by hand

Suppose without a penalty the best weight would be $w = 3$, and the loss near it is $\frac{1}{2}(w - 3)^2$.

- Ridge with penalty $\frac{\lambda}{2}w^2$, $\lambda = 1$: minimise $\frac{1}{2}(w - 3)^2 + \frac{1}{2}w^2$. Derivative $(w - 3) + w = 0 \Rightarrow w = 1.5$. Shrunk by half, but not zero.
- Lasso with penalty $\lambda|w|$, $\lambda = 1$: for $w > 0$ the derivative is $(w - 3) + 1 = 0 \Rightarrow w = 2$. If instead the unpenalised optimum were $w = 0.8$ (a weak feature), the same calculation gives $w = -0.2$, which contradicts $w > 0$; the minimum is at exactly $w = 0$. Lasso subtracts a fixed amount λ and stops at zero (“soft thresholding”): $w = \text{sign}(w_0) \max(|w_0| - \lambda, 0)$.

Why L1 gives exact zeros: the picture. Minimising loss + $\lambda \cdot$ penalty is equivalent to minimising the loss subject to a budget “penalty $\leq t$ ”. In two dimensions the L2 budget $w_1^2 + w_2^2 \leq t$ is a circle; the L1 budget $|w_1| + |w_2| \leq t$ is a diamond with corners on the axes. The loss contours are ellipses centred at the unpenalised solution; we expand them until they first touch the budget region. A circle is smooth, so the touching point is almost never on an axis. A diamond has sharp corners sticking out, and an expanding ellipse very often hits a corner first. A corner lies on an axis, so one weight is exactly zero.



The Bayesian view (one sentence, often asked). L2 regularisation is the same as assuming a Gaussian prior on the weights (most weights small, centred at zero); L1 is the same as a Laplace prior (sharply peaked at zero). Chapter 12 derives this.

1.5.3 Other regularisers you must be able to name

- **Early stopping:** stop training when validation error starts rising.
- **Dropout** (neural nets): randomly switch off units during training.
- **Data augmentation:** create plausible new training examples (flipped images, paraphrased text).
- **Tree limits:** maximum depth, minimum samples per leaf, pruning.
- **Ensembling** (bagging) reduces variance by averaging.
- **Label smoothing, weight sharing** (CNNs), **batch norm** (a mild side effect).

⚠ Common mistake

Always **standardise features** before Ridge or Lasso. The penalty treats every weight equally, but a weight's size depends on the feature's units: "salary in rupees" needs a tiny weight, "salary in crores" a huge one. Without scaling, the penalty unfairly punishes features measured in large-number units less and small-unit features more.

1.6 No free lunch, and why baselines matter

The **no free lunch theorem** says that averaged over *all* possible problems, no learning algorithm beats any other. Every algorithm wins by making assumptions (smoothness, linearity, locality) that suit the problem. Practically: always start with a simple **baseline** (predict the mean; logistic regression; "most popular jobs") and demand that a complex model beat it by a margin worth its cost.

1.7 Interview questions

Q1.1

★ **PYQ (OA and INT: bias-variance graphs)** Draw the error versus model complexity graph and explain it. Where is the best model?

🕒 **In one line.** Training error falls monotonically with complexity; test error is U-shaped; the best model sits at the bottom of the test-error U, where the sum of bias² and variance is smallest.

⚙️ **The full answer.** On the x-axis is model complexity (depth of a tree, degree of a polynomial, number of parameters, or $1/\lambda$). On the y-axis, error.

1. The **training error** curve starts high (a very simple model cannot fit even the data it sees) and keeps falling. A sufficiently complex model can memorise every training point, so it can reach zero.
2. The **test error** curve first falls together with training error, because extra flexibility is capturing real structure (bias falling). Beyond some point, extra flexibility only captures noise specific to the sample (variance rising), so test error turns upward.
3. The left region is **underfitting** (high bias), the right region **overfitting** (high variance). The vertical gap between the curves is a measure of overfitting.

The best model is at the minimum of the *validation* curve (we estimate the test curve with a validation set or cross-validation).

📊 **Worked example.** A polynomial fitted to 20 noisy points from a sine curve: degree 1 has train MSE 0.30, val 0.32 (underfit); degree 4 has 0.05 and 0.07 (good); degree 15 has 0.001 and 2.4 (overfit). Choose degree 4.

💡 **If the interviewer pushes.** “What if I plot against number of training examples instead?” That is a learning curve: training error rises (harder to fit more points), validation error falls; they converge. A persistent big gap signals variance; both converging at a high value signals bias. “And against training iterations of a neural net?” Training error keeps falling; validation error falls then rises: stop at its minimum (early stopping).

⚠️ **Trap to avoid.** Saying “test error decreases forever with complexity” or putting the optimum at zero training error. Also do not confuse the x-axis: more *data* is not more *complexity*.

Q1.2

★ **PYQ (OA, IIT Roorkee ML section)** A model has 99% training accuracy and 72% validation accuracy. What is happening, and what would you try, in order?

🕒 **In one line.** It is overfitting (high variance); I would first check for leakage or a split problem, then add regularisation or reduce complexity, and get more or augmented data if possible.

⚙️ **The full answer.** A 27-point gap means the model’s behaviour depends heavily on which rows it saw. Before changing the model I verify the measurement:

1. **Is the split sound?** Are there duplicates across train and validation? Is the validation set from a different period or distribution (a time shift)? Is the validation set tiny and noisy?
2. **Is there leakage in training only?** For example a feature that exists in the training extract but is computed differently for validation.

If the measurement is fine, reduce variance:

1. Simplify: shallower trees, fewer features, smaller network.
2. Regularise: L1/L2, dropout, min-samples-per-leaf, early stopping.
3. Average: bagging or a random forest instead of one tree.

4. More data or augmentation: the only fix that reduces variance without adding bias.
5. Feature selection: remove noisy features (high-cardinality IDs are classic culprits).

🏠 Worked example. A decision tree on 5,000 Shiksha leads to predict “will apply to a college”: unlimited depth gives 99%/72%. Limiting `max_depth=6` and `min_samples_leaf=50` gives 84%/81%. A random forest gives 90%/85%.

💡 If the interviewer pushes. “Why not just collect more data first?” It is often slow and expensive; regularisation is free. But if a learning curve shows validation error still dropping steeply with more data, data collection is worth it.

⚠️ Trap to avoid. Jumping straight to “use a bigger model” (that worsens variance), or to SMOTE, which addresses class imbalance, not overfitting.

Q1.3

★ **PYQ (INT: weight decay)** What is weight decay? Is it the same as L2 regularisation?

🎯 In one line. Weight decay shrinks every weight by a small fraction at each update; for plain gradient descent it is exactly equivalent to adding an L2 penalty, but for adaptive optimisers like Adam the two differ, which is why AdamW exists.

⚙️ The full answer. With L2, the loss becomes $L(w) + \frac{\lambda}{2} \|w\|^2$, whose gradient is $\nabla L + \lambda w$. A gradient-descent step is

$$w \leftarrow w - \eta(\nabla L + \lambda w) = (1 - \eta\lambda)w - \eta\nabla L.$$

The factor $(1 - \eta\lambda)$ multiplies (“decays”) the weight every step before the usual update. That is weight decay. So for SGD, “weight decay with coefficient $\eta\lambda$ ” and “L2 with coefficient λ ” are the same thing.

In Adam, the gradient is divided by a running estimate of its size, per weight. If the L2 term is put inside the gradient, it also gets divided, so weights with large gradients are decayed *less*. That is not what we want. **AdamW** applies the decay directly to the weights, outside the adaptive scaling (“decoupled weight decay”), and generalises better; it is the default for training Transformers.

🏠 Worked example. $w = 1$, $\eta = 0.1$, $\lambda = 0.01$, and the data gradient happens to be 0. One SGD step: $w \leftarrow (1 - 0.001) \cdot 1 = 0.999$. With no data pushing it, the weight decays geometrically toward 0.

💡 If the interviewer pushes. “Why does shrinking weights reduce overfitting?” Small weights make the function smoother: a small change in input causes a small change in output, so the model cannot carve sharp, data-specific wiggles. In a sigmoid/tanh network, small weights also keep units in their near-linear range.

⚠️ Trap to avoid. Saying weight decay “reduces the learning rate”. It shrinks the weights, not the step size.

Q1.4

★ **PYQ (INT and OA DL: overfitting)** List the ways to reduce overfitting in a neural network and explain why each works.

🎯 **In one line.** More data or augmentation, smaller network, L2/weight decay, dropout, early stopping, batch norm, label smoothing, and ensembling; each either limits the effective capacity or makes the training signal less specific to the sample.

⚙️ **The full answer.**

- **More data / augmentation:** noise patterns stop being repeatable, so memorising them stops paying. Augmentation (crops, flips, synonym replacement) teaches invariances we know are true.
- **Smaller network:** fewer parameters, fewer ways to memorise.
- **Weight decay:** smooth functions (previous question).
- **Dropout:** randomly zeroes units each step with probability p , so no unit can rely on a specific partner; it behaves like averaging many thinned networks (an implicit ensemble).
- **Early stopping:** weights start small and grow during training; stopping early limits how far they move, acting like an L2 penalty.
- **Batch normalisation:** the batch statistics add a little noise per example, a mild regulariser.
- **Label smoothing:** replaces hard 0/1 targets with 0.05/0.95-style targets, so the network is not rewarded for infinitely confident logits.
- **Ensembling / checkpoint averaging:** averaging reduces variance.

📊 **Worked example.** A CNN classifying 2,000 property photos into 6 room types: 99% train, 70% val. Adding flips/crops and dropout 0.3 gives 93%/82%; switching to a pretrained ResNet fine-tuned with weight decay gives 95%/91%. Transfer learning is often the biggest single fix when data are small.

⚠️ **Trap to avoid.** Listing techniques without the “why”. The panel reported asking exactly this to test intuition.

Q1.5

Explain the bias–variance trade-off to a non-technical product manager in under a minute.

🎯 **In one line.** A model can be wrong because it is too rigid to see the pattern (bias) or because it is so eager that it mistakes coincidences for patterns (variance); we tune its flexibility to balance the two.

⚙️ **The full answer.** “Imagine predicting the price of flats in Noida. If I only use ‘price = average price’, I am always off for big or small flats: that is *bias*, a systematic blind spot. If I memorise every flat sold last month, including the one that sold cheap because the owner was in a hurry, my predictions jump around depending on which flats happened to sell: that is *variance*. The sweet spot is a model that uses size, location and floor but ignores one-off stories. We find it by testing on flats the model has not seen.”

🗣️ **If the interviewer pushes.** “So which one is worse?” Neither in general; the business question decides. For a salary estimate shown to users, high variance (estimates that jump week to week) erodes trust, so we would accept slightly more bias for stability.

Q1.6

Compute the expected squared error at a point where the true value is 10, predictions over four training sets are 8, 12, 14, 10, and the noise variance is 2.

🎯 **In one line.** $\text{Bias}^2 = 1$, variance = 5, noise = 2, so the expected error is 8.

⚙️ **The full answer.** Mean prediction = $(8 + 12 + 14 + 10)/4 = 11$. Bias = $11 - 10 = 1$, $\text{bias}^2 = 1$. Variance = $\frac{(8-11)^2 + (12-11)^2 + (14-11)^2 + (10-11)^2}{4} = \frac{9+1+9+1}{4} = 5$. Total = $1 + 5 + 2 = 8$.

🗣️ **If the interviewer pushes.** “Which part can no model remove?” The noise, 2. “How would bagging change this?” It averages models, cutting the variance term (towards $\rho\sigma^2$ if models are correlated) while leaving bias roughly unchanged.

Q1.7

Why does L1 regularisation produce sparse weights but L2 does not? Answer both geometrically and with derivatives.

🎯 **In one line.** The L1 penalty has a constant-size pull toward zero that does not fade as the weight shrinks, and its constraint region has corners on the axes; the L2 pull is proportional to the weight, so it fades to nothing near zero.

⚙️ **The full answer. Derivatives.** The L2 penalty λw^2 has derivative $2\lambda w$: when w is already small, the push toward zero is tiny, so w settles at a small non-zero value wherever the data’s pull balances it. The L1 penalty $\lambda|w|$ has derivative $\lambda \cdot \text{sign}(w)$: a constant push of size λ however small w is. If the data’s pull on a weight is weaker than λ , nothing can hold the weight away from zero, so it lands exactly at zero.

Geometry. As drawn in Section 1.5, the L1 constraint region is a diamond whose corners lie on the axes; expanding loss ellipses usually touch a corner first, and corners have some coordinates equal to zero. The L2 region is a circle with no corners.

📊 **Worked example.** With an unpenalised optimum $w_0 = 0.8$ and $\lambda = 1$: Lasso gives $\max(0.8 - 1, 0) = 0$ (feature removed). Ridge with the same loss gives $0.8/(1 + 1) = 0.4$ (shrunk, kept).

🗣️ **If the interviewer pushes.** “When would you prefer Ridge anyway?” When many features each carry a little signal (text features, embeddings) or features are highly correlated: Lasso arbitrarily keeps one of a correlated group and can be unstable; Ridge spreads weight across the group.

Q1.8

What is data leakage? Give three concrete examples from a job-portal setting and how you would detect each.

🎯 **In one line.** Leakage is using, during training, information that will not be available when the model makes real predictions; it inflates offline scores and collapses in production.

⚙️ **The full answer.**

1. **Target leakage:** predicting “will this candidate be hired” using “number of offer letters generated”. Detection: a single feature with almost perfect importance or AUC; ask “when is this value created, relative to the event?”
2. **Preprocessing leakage:** fitting a scaler or TF-IDF on all resumes before splitting. Detection: code review; enforce a Pipeline so every transform is fit inside each training fold.
3. **Temporal leakage:** random train/test split on applications from 2023–2025 when the model will predict 2026 behaviour; also features like “total applies this job will get”. Detection: compare a random split score with a time-based split score; a large drop signals leakage.
4. **Group leakage:** the same candidate’s multiple applications in both train and test. Detection: group k -fold by candidate ID.

💡 **If the interviewer pushes.** “Your offline AUC is 0.99 on a hard problem. What do you think first?” Leakage, before celebrating. Then I look at the top features and the split.

Q1.9

Why do we need a separate test set if we already use cross-validation?

🎯 **In one line.** Because choosing the best configuration using CV scores makes those scores optimistic; an untouched test set gives an unbiased final estimate.

⚙️ **The full answer.** CV is used to *select* among hyperparameters and models. Selection on noisy scores favours lucky configurations (the maximum of noisy numbers is biased upward). The test set has not influenced any decision, so its score is an honest estimate of performance on new data. When data are very small, **nested cross-validation** replaces the single test set: an outer CV loop estimates performance, an inner CV loop tunes hyperparameters inside each outer training fold.

Q1.10

When is a random train–test split wrong? Give two cases and the right split.

🎯 **In one line.** When rows are not independent: time-ordered data (use a time-based split) and grouped data such as many rows per user (use a group split).

⚙️ **The full answer.** A random split assumes rows are exchangeable, i.e. the test rows look like unseen future rows. For time series (daily applications to a job), a random split puts tomorrow in training and yesterday in testing: the model interpolates rather than forecasts. Use a forward-chaining split. For grouped data (many sessions per seeker), a random split lets the model memorise each seeker’s habits and get them “right” in test. Use `GroupKFold`. A third case is **distribution shift**: if the model will be deployed in a new city, validate on a held-out city.

Q1.11

 **Think deeper** Can adding a completely useless random feature ever *improve* validation accuracy? Can it hurt?

🎯 **In one line.** It can appear to help by chance on a small validation set, but in expectation it hurts slightly, because it adds variance without adding signal; how much it hurts depends on the model's regularisation.

⚙️ **The full answer.** A random feature has no relationship with the target in the population. A flexible model will still find spurious patterns in it on the training sample (pure noise always correlates a little with anything in a finite sample). Those patterns do not repeat on new data, so predictions get noisier: variance goes up. With strong regularisation (Lasso, a random forest that rarely picks it) the damage is small; with an unregularised model on few rows and many such features, it can be large. On a single small validation set, random chance can make the noisy model score higher, which is exactly why we use CV and look at the spread.

📊 **Worked example.** Linear regression with $n = 30$ rows: adding 20 random columns raises training R^2 from 0.60 to 0.85 (it can only go up) while CV R^2 drops from 0.55 to 0.20.

💬 **If the interviewer pushes.** “How is this used in practice?” A classic feature-selection trick: add a few random “shadow” features (Boruta) and drop any real feature whose importance is not higher than the best random one.

Q1.12

 **Think deeper** A model has *lower* validation loss than training loss. Is something wrong?

🎯 **In one line.** Not necessarily; regularisation that is active only during training (dropout, augmentation), the timing of when training loss is averaged, or an easier validation set can all cause it, but it is worth checking for a leaky or unrepresentative validation set.

⚙️ **The full answer.**

1. **Dropout and augmentation** make training harder: training loss is computed on a crippled network and distorted inputs; validation uses the full network on clean inputs.
2. **Averaging timing:** training loss is often the average over the epoch while the weights were still improving; validation loss is computed at the end of the epoch with the better weights.
3. **Easier validation data:** by chance, or because of a bad split (validation drawn from a cleaner source), or because of duplicates of training rows in validation (leakage).

I would check the third explanation first, since the first two are benign.

Q1.13

What is the difference between parameters and hyperparameters? Give three examples of each for three different models.

🎯 **In one line.** Parameters are learned from data by the training algorithm; hyperparameters are chosen by us before training and control how learning happens.

⚙️ **The full answer.**

Model	Parameters (learned)	Hyperparameters (chosen)
Linear/logistic regression	weights, bias	λ (or C), penalty type
Decision tree	split features, thresholds, leaf values	max depth, min samples per leaf, criterion
Neural network	weights and biases of every layer	learning rate, layers, width, dropout, batch size
k -means	centroid positions	k , initialisation

Hyperparameters are tuned on validation data (grid, random or Bayesian search), never on the test set.

Q1.14

Your model has 62% training and 61% validation accuracy on a problem where humans get 95%. What do you do?

🎯 **In one line.** This is high bias (underfitting): increase capacity or improve features, and reduce regularisation; more data alone will not help.

⚙️ **The full answer.** The gap is tiny, so the model is consistent but systematically limited. Options in order: (1) richer features (interactions, text embeddings, domain features); (2) a more flexible model (gradient boosting instead of a linear model; a deeper network); (3) less regularisation (larger C , smaller λ , less dropout); (4) train longer if it is a neural network that has not converged; (5) check label quality: if labels are noisy, “human 95%” may not be achievable with these labels.

⚠️ **Trap to avoid.** Recommending more data. A learning curve that is flat at 61% tells you more rows will not move it.

Q1.15

🧠 **Think deeper** Is a model with 0% training error always overfitting?

🎯 **In one line.** No. Zero training error means the model can fit the sample; it is overfitting only if validation error is much worse than it could be. Many good models (boosted trees, large neural nets, 1-NN on a clean problem) reach zero training error and still generalise well.

⚙️ **The full answer.** Overfitting is defined by the gap and by comparison with alternatives, not by training error alone. If the classes are perfectly separable and labels are clean (say, digit images), fitting every training point can be correct. Modern deep networks routinely interpolate the training set yet generalise, which is part of the double-descent story. Conversely, if labels are noisy (10% mislabelled), zero training error means the model learned the mislabelled points too, which usually hurts.

Q1.16

What is k -fold cross-validation, how do you choose k , and what does it cost?

🎯 **In one line.** Split data into k folds, train k times each time holding one fold out, and average the scores; $k = 5$ or 10 is the usual compromise between bias of the estimate, its variance, and compute.

⚙️ **The full answer.** Small k (like 2) trains on only half the data each time, so each model is weaker than the final one (pessimistic, biased estimate). Large k (like n , leave-one-out) trains on almost all data (low bias) but the k models are nearly identical, so their errors are correlated and the estimate is noisy; it also costs k full trainings. With 1 million rows, a single holdout is already precise, so CV is often unnecessary; with 2,000 rows, 5- or 10-fold stratified CV is standard.

💡 **If the interviewer pushes.** “After CV, which model do you deploy?” Usually retrain on all training data with the chosen hyperparameters (or average the k fold models, common in competitions).

Q1.17

🧠 **Think deeper** If regularisation always increases training error, why does it ever help?

🎯 **In one line.** Because it trades a small increase in bias for a larger decrease in variance, and test error is the sum of both.

⚙️ **The full answer.** Training error measures fit to this sample. Regularisation deliberately stops the model from fitting the sample perfectly, so training error rises. But the parts of the sample it refuses to fit are disproportionately noise. On new data that noise is different, so not having fitted it is an advantage. In bias–variance terms: variance falls faster than bias² rises, until λ becomes so large that the model ignores real signal too, after which test error rises again. That is why λ is tuned on validation data.

📊 **Worked example.** Ridge on 50 rows, 40 features: $\lambda = 0$: train MSE 0.1, val 5.0. $\lambda = 1$: train 0.8, val 1.2. $\lambda = 100$: train 3.0, val 3.1 (now underfitting).

Q1.18

What is the no-free-lunch theorem and what does it mean for a working data scientist?

🎯 **In one line.** No algorithm is best on every possible problem; an algorithm wins only when its built-in assumptions match the data, so we must try sensible candidates and validate empirically.

⚙️ **The full answer.** A linear model assumes additive effects; a CNN assumes local, translation-equivariant patterns; k -NN assumes nearby points share labels; gradient-boosted trees assume axis-aligned, piecewise-constant structure. Each is excellent when its assumption holds and poor otherwise. Practically: (1) start with a baseline, (2) pick candidates whose assumptions fit the data type (trees for tabular, CNNs for images, Transformers for text), (3) compare on a sound validation scheme.

Q1.19

 **Think deeper** Two models have the same validation accuracy, 85%. One is logistic regression, the other a 2,000-tree gradient-boosted model. Which do you ship?

🎯 **In one line.** Usually the simpler one: equal accuracy plus lower latency, easier debugging, better calibration and interpretability, and typically more stable behaviour under data shift.

⚙️ **The full answer.** Accuracy is one of several criteria. Ask: (1) Is the 85% estimate equally certain for both? Look at CV spread. (2) Latency and cost: a model scoring 30 lakh Resdex candidates per search must be fast. (3) Interpretability: recruiters and compliance may need reasons. (4) Robustness: the simpler model has fewer ways to break when inputs drift. (5) Other metrics: compare calibration, recall at the operating point, and performance on important segments (freshers, tier-2 cities). If the complex model wins clearly on a metric that matters (say recall on rare high-value leads), it can be worth it.

Q1.20

What is stratified sampling and when must you use it?

🎯 **In one line.** Splitting so that each subset keeps the same class proportions as the whole; it is essential for imbalanced classification and small datasets.

⚙️ **The full answer.** With 1% positives and 1,000 rows, a random 20% validation split contains 2 positives on average and quite often 0 or 1, making recall meaningless and unstable. Stratification guarantees exactly 2 per fold. In scikit-learn: `train_test_split(..., stratify=y)` and `StratifiedKFold`.

CHAPTER 2

Linear Regression, from a Line to the Normal Equations

🎥 **Companion video for this chapter:** [StatQuest: Linear Regression, Clearly Explained](#)

Linear regression is the first model interviewers drill and the one they come back to in Round 4 (“what are its assumptions? what if homoscedasticity fails?”). It is worth knowing completely: every idea in it (a loss, a gradient, a closed form, assumptions, diagnostics, regularisation) reappears in every other model.

2.1 The model in words and symbols

2.1.1 One feature

We believe the target y is roughly a straight-line function of one feature x :

$$\hat{y} = wx + b.$$

w is the **slope** (how much $\hat{y}$ rises when x rises by one unit) and b the **intercept** (the prediction when $x = 0$).

📊 Small example: reading the numbers

A fitted salary model: $\widehat{\text{salary}} = 1.8 \times \text{experience} + 3.5$ (lakh). A fresher ($x = 0$) is predicted 3.5 lakh. Each extra year adds 1.8 lakh. Someone with 4 years: $1.8 \cdot 4 + 3.5 = 10.7$ lakh.

2.1.2 Many features

With d features $x_1, \dots, x_d$ (experience, a city indicator, number of skills, ...):

$$\hat{y} = w_1x_1 + w_2x_2 + \dots + w_dx_d + b = \mathbf{w}^\top \mathbf{x} + b.$$

The symbol $\mathbf{w}^\top \mathbf{x}$ is a **dot product**: multiply matching entries and add them up. For $\mathbf{w} = (2, -1)$ and $\mathbf{x} = (3, 4)$, $\mathbf{w}^\top \mathbf{x} = 2 \cdot 3 + (-1) \cdot 4 = 2$.

Stack n examples as rows of a matrix X (n rows, d columns, plus a column of 1s to carry the intercept). Then all predictions at once are $\hat{\mathbf{y}} = X\mathbf{w}$.

2.1.3 Why it is called “linear”

“Linear” means linear in the *parameters* $\mathbf{w}$, not in x . The model $\hat{y} = w_1x + w_2x^2 + w_3 \log x + b$ is still linear regression: we simply created new features x^2 and $\log x$. This is how linear regression fits curves.

2.2 Choosing the line: least squares

2.2.1 Residuals and the loss

For each example, the **residual** is the error $e_i = y_i - \hat{y}_i$. We want residuals small. We cannot just add them up (positive and negative errors would cancel), so we square them:

$$\text{SSE} = \sum_{i=1}^n (y_i - \hat{y}_i)^2, \quad \text{MSE} = \frac{1}{n} \text{SSE}.$$

Ordinary least squares (OLS) chooses w, b to minimise this.

Why squares and not absolute values?

1. Squares are smooth, so calculus gives a clean closed-form answer.
2. Squares punish large errors heavily (an error of 4 costs 16, four errors of 1 cost 4). Often appropriate, sometimes a weakness (outliers dominate).
3. If the noise is Gaussian, least squares is exactly maximum likelihood (Chapter 12). That is the deep reason.

2.2.2 Solving it with one feature, step by step

A **derivative** tells you the slope of a function: where a smooth function is at its lowest point, its slope is zero. So we take the derivative of SSE with respect to b and w and set each to zero.

- Setting $\partial \text{SSE} / \partial b = -2 \sum (y_i - wx_i - b) = 0$ gives $b = \bar{y} - w\bar{x}$. ($\bar{x}$ means the average of the x values.) So the line always passes through $(\bar{x}, \bar{y})$.
- Substituting and setting $\partial \text{SSE} / \partial w = 0$ gives

$$w = \frac{\sum (x_i - \bar{x})(y_i - \bar{y})}{\sum (x_i - \bar{x})^2} = \frac{\text{Cov}(x, y)}{\text{Var}(x)} = r \frac{s_y}{s_x},$$

where r is the correlation and s_x, s_y the standard deviations.

Small example: fit a line to four Noida flats

Area x (hundreds of sq ft): 5, 7, 9, 11. Price y (lakh): 30, 40, 44, 54. $\bar{x} = 8$, $\bar{y} = 42$. Deviations: $x - \bar{x} = -3, -1, 1, 3$; $y - \bar{y} = -12, -2, 2, 12$. $\sum (x - \bar{x})(y - \bar{y}) = 36 + 2 + 2 + 36 = 76$. $\sum (x - \bar{x})^2 = 9 + 1 + 1 + 9 = 20$. $w = 76/20 = 3.8$ lakh per 100 sq ft. $b = 42 - 3.8 \cdot 8 = 11.6$. Line: $\hat{y} = 3.8x + 11.6$. A 1,000 sq ft flat ($x = 10$): 49.6 lakh. Predictions 30.6, 38.2, 45.8, 53.4; residuals $-0.6, 1.8, -1.8, 0.6$; they sum to 0, as they must.

2.2.3 Many features: the normal equations

In matrix form, $\text{SSE} = \|\mathbf{y} - \mathbf{X}\mathbf{w}\|^2$. Its gradient (the vector of all partial derivatives) is $-2\mathbf{X}^\top(\mathbf{y} - \mathbf{X}\mathbf{w})$. Setting it to zero:

$$\mathbf{X}^\top \mathbf{X} \mathbf{w} = \mathbf{X}^\top \mathbf{y} \implies \mathbf{w} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}.$$

These are the **normal equations**. They require $\mathbf{X}^\top \mathbf{X}$ to be invertible, which fails when one feature is an exact combination of others (perfect multicollinearity) or when there are more features than rows.

The geometric picture (a favourite interview follow-up). Think of $\mathbf{y}$ as one point in n -dimensional space (one coordinate per example). Every possible prediction $X\mathbf{w}$ lies in the flat subspace spanned by the columns of X (the **column space**). Least squares picks the point in that subspace closest to $\mathbf{y}$: the **orthogonal projection**. So the residual vector $\mathbf{y} - X\hat{\mathbf{w}}$ is perpendicular to every column of X . That is exactly what $X^T(\mathbf{y} - X\hat{\mathbf{w}}) = 0$ says. With a column of ones in X , perpendicular to it means the residuals sum to zero.

Cost and alternatives. Forming $X^T X$ costs about nd^2 operations and inverting it about d^3 . Fine for d in the thousands. For huge d or streaming data, use gradient descent (Chapter 5). In practice libraries use a QR or SVD decomposition rather than an explicit inverse, which is numerically safer.

2.3 Measuring the fit

2.3.1 R^2 : fraction of variance explained

Compare our model with the dumbest possible model, “always predict $\bar{y}$ ”.

$$R^2 = 1 - \frac{SS_{res}}{SS_{tot}} = 1 - \frac{\sum (y_i - \hat{y}_i)^2}{\sum (y_i - \bar{y})^2}.$$

$R^2 = 1$: perfect; $R^2 = 0$: no better than the mean; on test data R^2 can be negative (worse than the mean). For simple regression with an intercept, $R^2 = r^2$.

Small example: R^2 for the flats

$SS_{res} = 0.36 + 3.24 + 3.24 + 0.36 = 7.2$. $SS_{tot} = 144 + 4 + 4 + 144 = 296$. $R^2 = 1 - 7.2/296 = 0.976$.

2.3.2 Adjusted R^2

Training R^2 never decreases when you add a feature (OLS can always give it weight 0). Adjusted R^2 penalises features: $\bar{R}^2 = 1 - (1 - R^2) \frac{n-1}{n-d-1}$. It goes down if a new feature adds less than chance would.

2.3.3 Error metrics

- **MAE** = $\frac{1}{n} \sum |y - \hat{y}|$: in target units, robust to outliers; minimised by the median.
- **RMSE** = $\sqrt{\text{MSE}}$: in target units, punishes large errors; minimised by the mean.
- **MAPE** = $\frac{1}{n} \sum |y - \hat{y}|/|y|$: percentage, scale-free, explodes near $y = 0$ and penalises over-prediction more than under-prediction.

2.4 Assumptions: what OLS needs, and for what

A crucial distinction interviewers love: some assumptions are needed for the *coefficients to be right on average* (unbiased), others only for *p-values and confidence intervals to be valid*.

Assumption	Plain meaning	How to check	If violated
Linearity	the average of y moves linearly with the features (after any transforms)	residuals vs fitted plot shows a curve	add x^2 , $\log x$, interactions, or use a non-linear model
Independence	errors of different rows are unrelated	Durbin–Watson; residuals vs time	time-series models, clustered standard errors, mixed models
Normality	errors are bell-shaped	Q–Q plot, Shapiro–Wilk	matters only for small-sample inference; transform y , bootstrap, or robust methods
Equal variance (homoscedasticity)	spread of errors is the same everywhere	residual plot shows a funnel; Breusch–Pagan test	transform y (log), weighted least squares, robust SEs, or a GLM
No perfect multicollinearity	no feature is an exact combination of others	VIF, correlation matrix	drop or combine features, PCA, Ridge
Exogeneity	errors are uncorrelated with features (no omitted confounder)	domain reasoning	coefficients biased; add the confounder, instrumental variables

Mnemonic: **LINE** plus “no multicollinearity”. For unbiased coefficients you need linearity, exogeneity and no perfect multicollinearity. Independence and homoscedasticity make OLS the *best* (minimum-variance) linear unbiased estimator (the Gauss–Markov theorem) and make the usual standard errors correct. Normality is only for exact small-sample t- and F-tests.

2.4.1 Heteroscedasticity, slowly

“Hetero” means different, “scedastic” means spread. Salaries are the textbook case: freshers’ salaries vary by ± 1 lakh, senior leaders’ by ± 20 lakh. A residual plot fans out like a funnel.

- **What breaks:** coefficients stay unbiased, but OLS standard errors (and hence p-values) are wrong, and OLS gives too much influence to the noisiest points.
- **Fixes:**
 1. *Transform y :* model $\log(\text{salary})$. Multiplicative noise ($\pm 10\%$) becomes additive, roughly constant noise. (This is what the reported interviewer wanted.)
 2. *Weighted least squares:* give each point weight $1/\sigma_i^2$, so noisy points count less.
 3. *Robust (“sandwich”, HC) standard errors:* keep OLS coefficients, fix the inference.
 4. *A GLM with a variance that grows with the mean,* such as Gamma regression with a log link (what the candidate proposed; Chapter 4 explains it fully).

2.4.2 Multicollinearity and VIF

If two features move together (“years of experience” and “age”), the model cannot tell which one deserves credit. Predictions stay fine, but individual coefficients become unstable: they can swing wildly or flip sign between samples. The **variance inflation factor** of feature j comes from regressing

x_j on all other features and taking its R_j^2 :

$$\text{VIF}_j = \frac{1}{1 - R_j^2}.$$

VIF = 1: no collinearity; VIF > 5 or > 10: worrying. The standard error of w_j is inflated by $\sqrt{\text{VIF}_j}$.

2.5 Regularised linear regression

- **Ridge**: minimise $\text{SSE} + \lambda \|\mathbf{w}\|^2$. Closed form $\mathbf{w} = (X^\top X + \lambda I)^{-1} X^\top \mathbf{y}$. Adding λ to the diagonal makes the matrix always invertible, which is why Ridge cures multicollinearity.
- **Lasso**: minimise $\text{SSE} + \lambda \sum |w_j|$. No closed form; solved by coordinate descent. Sets some weights exactly to zero.
- **Elastic net**: both penalties.

The intercept is not penalised, and features must be standardised first (Chapter 1).

2.6 Outliers, leverage and influence

- An **outlier** has an unusual y given its x (large residual).
- A **high-leverage point** has an unusual x (far from the other x values). It has the power to tilt the line.
- An **influential point** actually does tilt it: high leverage *and* a large residual. **Cook's distance** measures how much all fitted values change when that point is deleted.

Remedies: check whether it is a data error; use robust losses (Huber, MAE); or keep it if it is real and important. Deleting genuine extreme values is a mistake the InfoEdge coding round specifically tested.

In practice: where linear regression lives at InfoEdge

- **Salary engine** (Data Factory slide): estimating a fair salary range for a title, city and experience. Log-salary linear models with many one-hot features are a strong, interpretable baseline.
- **99acres price estimates**: price per square foot by locality, floor, age, amenities.
- As a **diagnostic tool**: coefficients answer “how much does one more year of experience add, holding city fixed?”, which product teams ask constantly.

2.7 Interview questions

Q2.1

★ **PYQ (INT, Round 4 follow-up)** What are the assumptions of linear regression?

🕒 **In one line.** Linearity in the parameters, independent errors, constant error variance (homoscedasticity), normally distributed errors, no perfect multicollinearity, and errors uncorrelated

with the features (exogeneity).

⚙️ **The full answer.** I like to group them by what they buy us.

1. **For coefficients to be correct on average (unbiased):**

- *Linearity:* $\mathbb{E}[y | x]$ is a linear combination of the features we included. Note: features can be transformed (x^2 , $\log x$), so this is about the specification.
- *Exogeneity:* the error has mean zero given the features. Violated when an omitted variable affects both x and y (omit “company tier” from a salary model and the “IIT degree” coefficient absorbs part of its effect).
- *No perfect multicollinearity:* otherwise the coefficients are not even uniquely defined.

2. **For OLS to be the most efficient linear unbiased estimator, and for standard errors to be right:** *independence* of errors and *homoscedasticity* (Gauss–Markov).

3. **For exact small-sample t-tests and confidence intervals:** *normality* of errors. With large n the central limit theorem makes this matter little.

📊 **Worked example.** For each I state a check: residual vs fitted plot (linearity, homoscedasticity), Durbin–Watson or residuals vs time (independence), Q–Q plot (normality), VIF (multicollinearity).

💬 **If the interviewer pushes.** “Do the features need to be normally distributed?” No; the assumptions are about the errors. Binary features are perfectly fine. “Does y need to be normal?” Not marginally; only the errors given x .

⚠️ **Trap to avoid.** Saying “the data must be normal”. Also listing assumptions without saying which consequences follow from violating each.

Q2.2

★ PYQ (INT, Round 4) Homoscedasticity is violated in your model. What do you do?

🎯 **In one line.** First confirm it with a residual plot or Breusch–Pagan test; then either transform the target (log or Box–Cox), use weighted least squares, keep OLS with heteroscedasticity-robust standard errors, or switch to a GLM whose variance grows with the mean.

⚙️ **The full answer.** **What goes wrong.** The coefficients are still unbiased, so predictions of the mean are fine. But OLS treats every point as equally reliable, so it is no longer the most efficient estimator, and the usual standard-error formula (which assumes one common σ^2) is wrong: p-values and confidence intervals cannot be trusted. Prediction intervals are also wrong: too wide for small values, too narrow for large ones.

Fix 1, transform y . If the noise is proportional to the level (salaries vary by about $\pm 15\%$), take logs: $\log(y) = \log(\text{level}) + \log(1 \pm 0.15)$, and the second term has roughly constant spread. Box–Cox searches over power transforms y^λ . Back-transforming needs care: $\exp(\mathbb{E}[\log y]) \neq \mathbb{E}[y]$; multiply by a smearing factor such as $\exp(\hat{\sigma}^2/2)$.

Fix 2, weighted least squares. If the variance of point i is σ_i^2 , minimise $\sum (y_i - \hat{y}_i)^2 / \sigma_i^2$. Noisy points get less weight. Usually σ_i is modelled (for instance proportional to $\hat{y}_i$).

Fix 3, robust standard errors. Keep OLS, but compute standard errors with the White/HC formula that does not assume constant variance. Good when the goal is inference and the coefficients themselves are fine.

Fix 4, a GLM. Model y with a distribution whose variance is tied to the mean: Gamma (variance $\propto \text{mean}^2$, constant coefficient of variation) or Poisson for counts (variance = mean), usually with a log link. This models the mean on the original scale, avoiding the back-transformation problem.

 **Worked example.** Salary vs experience: residual spread is about 1 lakh at 1 year and 8 lakh at 15 years. Fit $\log(\text{salary})$; the residual plot becomes a flat band; the coefficient 0.08 means “each year adds about 8%”.

 **If the interviewer pushes.** “Which would you choose?” For prediction with a multiplicative structure, log or Gamma-log GLM; for explaining effects with valid p-values, robust SEs.

 **Trap to avoid.** Claiming heteroscedasticity biases the coefficients. It does not (unless it comes from misspecification).

Q2.3

★ **PYQ (INT: multicollinearity and VIF)** What is multicollinearity, why is it a problem, and how do you detect and fix it?

 **In one line.** It is strong correlation among features; it makes individual coefficients unstable and their standard errors large, though predictions are largely unaffected; detect it with VIF and fix it by dropping, combining, PCA or Ridge.

 **The full answer.** If x_1 and x_2 almost always move together, many combinations of (w_1, w_2) give nearly the same predictions ($w_1 + w_2 \approx \text{constant}$). The data cannot pin down the split, so small changes in the sample swing the coefficients, sometimes flipping signs. Mathematically, $X^T X$ is nearly singular, so its inverse (which sets the coefficient variances) has huge entries.

Detect: pairwise correlations miss three-way relations, so use $\text{VIF}_j = 1/(1 - R_j^2)$. Also the condition number of X . **Fix:** drop one of the pair; combine them (an index); use PCA; use Ridge (which adds λI and stabilises the inverse); or collect data where they vary independently.

 **Worked example.** Regressing “years of experience” on “age” and “graduation year” gives $R^2 = 0.96$, so $\text{VIF} = 25$: the standard error of the experience coefficient is 5× what it would be without collinearity.

 **If the interviewer pushes.** “Does it matter for a random forest?” Prediction is fine, but importance gets split between the correlated features, so each looks less important than it is.

Q2.4

Derive the OLS slope and intercept for a single feature.

 **In one line.** Set the derivatives of the squared-error sum with respect to b and w to zero; you get $b = \bar{y} - w\bar{x}$ and $w = \sum(x - \bar{x})(y - \bar{y}) / \sum(x - \bar{x})^2$.

 **The full answer.** $S(w, b) = \sum(y_i - wx_i - b)^2$.

$$\frac{\partial S}{\partial b} = -2 \sum (y_i - wx_i - b) = 0 \Rightarrow \sum y_i = w \sum x_i + nb \Rightarrow b = \bar{y} - w\bar{x}.$$

Substitute b : residual = $(y_i - \bar{y}) - w(x_i - \bar{x})$.

$$\frac{\partial S}{\partial w} = -2 \sum x_i [(y_i - \bar{y}) - w(x_i - \bar{x})] = 0.$$

Because $\sum \bar{x}[(y_i - \bar{y}) - w(x_i - \bar{x})] = 0$, we can replace x_i by $(x_i - \bar{x})$:

$$\sum (x_i - \bar{x})(y_i - \bar{y}) = w \sum (x_i - \bar{x})^2 \Rightarrow w = \frac{S_{xy}}{S_{xx}}.$$

The second derivative $2 \sum (x_i - \bar{x})^2 > 0$, so this is a minimum.

 **Worked example.** See the Noida flats example: $w = 76/20 = 3.8$, $b = 11.6$.

Q2.5

Derive the normal equations and give their geometric meaning.

 **In one line.** Minimising $\|\mathbf{y} - X\mathbf{w}\|^2$ gives $X^T X \mathbf{w} = X^T \mathbf{y}$; geometrically, the fitted vector is the orthogonal projection of $\mathbf{y}$ onto the column space of X , so residuals are perpendicular to every feature column.

 **The full answer.** Expand $L = (\mathbf{y} - X\mathbf{w})^T (\mathbf{y} - X\mathbf{w}) = \mathbf{y}^T \mathbf{y} - 2\mathbf{w}^T X^T \mathbf{y} + \mathbf{w}^T X^T X \mathbf{w}$. Gradient: $-2X^T \mathbf{y} + 2X^T X \mathbf{w} = 0$. If $X^T X$ is invertible, $\mathbf{w} = (X^T X)^{-1} X^T \mathbf{y}$. The **hat matrix** $H = X(X^T X)^{-1} X^T$ maps $\mathbf{y}$ to $\hat{\mathbf{y}}$; it is a projection ($H^2 = H$), and its diagonal entries are the leverages. Its trace equals the number of parameters.

 **If the interviewer pushes.** “What if $X^T X$ is not invertible?” Infinitely many solutions fit equally well; use the pseudo-inverse (minimum-norm solution, via SVD), drop redundant columns, or add a Ridge penalty.

Q2.6

What does R^2 mean? Can it be negative? Why do we need adjusted R^2 ?

 **In one line.** R^2 is the fraction of the target’s variance the model explains relative to predicting the mean; on new data it can be negative; adjusted R^2 corrects for the fact that training R^2 never falls when features are added.

 **The full answer.** $R^2 = 1 - SS_{res}/SS_{tot}$. On training data with an intercept, OLS can always do at least as well as the mean, so $0 \leq R^2 \leq 1$. On test data a bad model can have $SS_{res} > SS_{tot}$. Adjusted $R^2 = 1 - (1 - R^2) \frac{n-1}{n-d-1}$ rises only if a new feature improves fit more than expected by chance.

 **Worked example.** $n = 50, d = 4, R^2 = 0.80$: adjusted = $1 - 0.2 \cdot 49/45 = 0.782$.

 **Trap to avoid.** “ $R^2 = 0.9$ means the model is good/causal.” A high R^2 can come from leakage or a trend, and a low R^2 can still be very useful (noisy human behaviour).

Q2.7

 **Think deeper** Can linear regression fit a curve? Can it fit $y = ae^{bx}$?

 **In one line.** Yes: “linear” means linear in the parameters, so adding x^2 , x^3 or $\log x$ features fits curves; an exponential with multiplicative noise becomes linear after taking logs.

 **The full answer.** A polynomial model $\hat{y} = w_0 + w_1x + w_2x^2$ is a weighted sum of fixed functions of x , so OLS solves it with the same normal equations. For $y = ae^{bx}\varepsilon$ (multiplicative noise ε), $\log y = \log a + bx + \log \varepsilon$: regress $\log y$ on x ; the slope estimates b and the intercept estimates $\log a$. If the noise is *additive* ($y = ae^{bx} + \varepsilon$), log-transforming distorts it; then use non-linear least squares or a GLM with a log link (Poisson/Gamma), which models $\log \mathbb{E}[y] = \log a + bx$ directly.

 **Worked example.** Number of applications to a job vs days since posting decays exponentially: regressing $\log(\text{applies} + 1)$ on days gives slope -0.12 , i.e. about 11% fewer applies per extra day ($e^{-0.12} = 0.887$).

Q2.8

 **Think deeper** Linear regression on a 0/1 target: what goes wrong? Is it ever acceptable?

 **In one line.** Predictions can fall outside $[0, 1]$, errors are heteroscedastic by construction, and far-away easy points tilt the line; it is acceptable (the “linear probability model”) for quick effect estimates near the middle of the range, but logistic regression is the right default.

 **The full answer.** For a binary y , the error variance is $p(1 - p)$, which depends on x , so homoscedasticity fails automatically. Squared loss also penalises a correctly classified but extreme point (predicted 1.7 for a true 1) and pulls the line toward it, moving the 0.5 threshold. Economists still use the linear probability model because its coefficients read directly as “change in probability” and with robust SEs it is fine near $p \approx 0.5$.

Q2.9

What is the difference between correlation and the regression slope? When are they equal?

 **In one line.** Slope $= r \cdot s_y / s_x$; it carries units and direction of prediction, correlation is unit-free and symmetric; they are equal when both variables are standardised.

 **The full answer.** Regressing y on x gives slope rs_y/s_x ; regressing x on y gives rs_x/s_y ; the two lines differ unless $|r| = 1$. Their product is r^2 . After converting both variables to z-scores, the slope is exactly r .

 **Worked example.** If $r = 0.6$, $s_y = 5$ lakh, $s_x = 2$ years, the slope is $0.6 \cdot 5/2 = 1.5$ lakh per year.

Q2.10

Explain Ridge regression and why it handles multicollinearity.

🕒 **In one line.** Ridge adds $\lambda \|\mathbf{w}\|^2$ to the squared error; its solution $(X^T X + \lambda I)^{-1} X^T \mathbf{y}$ always exists, and the added λ stabilises directions in which the data carry little information.

⚙️ **The full answer.** With collinear features $X^T X$ has a tiny eigenvalue along the direction where features cancel ($w_1 - w_2$). Inverting divides by that tiny number, so coefficients explode along that direction. Ridge adds λ to every eigenvalue, so the smallest becomes at least λ , and the explosion stops. Equivalently it shrinks the coefficients most along low-variance principal directions. The price is a little bias.

📊 **Worked example.** One standardised feature, $\sum x_i^2 = 10$, $\sum x_i y_i = 20$: OLS $w = 2$; Ridge $\lambda = 10$: $w = 20/(10 + 10) = 1$.

Q2.11

Lasso vs Ridge vs Elastic net: when would you use each?

🕒 **In one line.** Lasso when you believe few features matter and want selection; Ridge when many features each matter a little or are correlated; Elastic net when you want sparsity but have correlated groups.

⚙️ **The full answer.** Lasso's L1 penalty zeroes weights (feature selection) but with correlated features it keeps one arbitrarily and can be unstable, and it can select at most n features when $d > n$. Ridge never zeroes but is stable. Elastic net ($\alpha L1 + (1 - \alpha)L2$) groups correlated features and still sparsifies.

📊 **At InfoEdge scale.** For a salary model with 5,000 one-hot skill features, Elastic net gives a compact, stable model whose non-zero skills are explainable to product teams ("Kubernetes adds 6%").

Q2.12

What are leverage, outliers and influential points? How do you detect them?

🕒 **In one line.** Leverage is unusualness in x , an outlier is unusualness in y given x , and an influential point is one whose removal changes the fit a lot; leverage comes from the hat matrix diagonal, influence from Cook's distance.

⚙️ **The full answer.** A high-leverage point that lies on the trend does no harm and even improves precision. A high-leverage point off the trend acts like a long lever and pulls the line. Cook's distance combines residual size and leverage; values above $4/n$ are commonly flagged. Studentised residuals flag outliers.

💡 **If the interviewer pushes.** "Would you delete it?" Only if it is an error (a price of Rs. 1, a negative area). If it is a real luxury villa, keep it, and consider a robust loss or a separate model for that segment.

Q2.13

Think deeper You fit a regression of salary on “number of certifications” and get a negative coefficient. Does that mean certifications reduce salary?

🕒 **In one line.** Not necessarily: it is a correlation conditional on the other included features, and it can be driven by confounding (juniors collect more certificates) or by collinearity with other features.

⚙️ **The full answer.** A coefficient measures the association of a feature with the target *holding the other features fixed*. If experience is not in the model, “many certificates” may simply indicate a junior person (omitted-variable bias). If experience is in the model and highly collinear, the sign can flip between samples. Causal claims need either an experiment or a careful causal design (controls for confounders, instrumental variables, difference in differences).

📊 **Worked example.** Without experience: coefficient -0.4 lakh/cert. Adding experience: $+0.3$ lakh/cert. This is Simpson’s paradox in regression form.

Q2.14

Gradient descent or normal equations for linear regression: which and when?

🕒 **In one line.** Normal equations (via QR/SVD) for small to medium d because they are exact and need no tuning; gradient descent (or SGD) for very large d , very large or streaming n , or when you add penalties without closed forms.

⚙️ **The full answer.** Normal equations cost about $O(nd^2 + d^3)$ and need the whole data in memory. Gradient descent costs $O(nd)$ per pass, can stream mini-batches, but needs a learning rate, feature scaling and many iterations. With $d = 100$ and $n = 10^6$: normal equations are instant. With $d = 10^6$ sparse text features: SGD.

Q2.15

Think deeper Why do we minimise squared error and not absolute error? When would absolute error be better?

🕒 **In one line.** Squared error is smooth, has a closed form, and is the maximum-likelihood loss for Gaussian noise; absolute error is better when there are outliers or heavy-tailed noise, and it targets the median instead of the mean.

⚙️ **The full answer.** The value c that minimises $\sum (y_i - c)^2$ is the mean; the one that minimises $\sum |y_i - c|$ is the median. So OLS estimates the conditional mean and least-absolute-deviation regression estimates the conditional median. With one data-entry error of Rs. 10 crore in a property dataset, the mean (and OLS) moves a lot; the median barely moves. Huber loss is a compromise: squared for small errors, absolute for large ones.

📊 **Worked example.** Values $\{3, 4, 5, 6, 100\}$: mean = 23.6, median = 5.

Q2.16

Interpret the coefficient 0.08 in a regression of $\log(\text{salary})$ on years of experience. What if both sides were logged?

🕒 **In one line.** Each extra year multiplies salary by $e^{0.08} \approx 1.083$, about +8.3%; in a log-log model the coefficient is an elasticity: a 1% increase in x gives about a $\beta\%$ increase in y .

⚙️ **The full answer.** $\log y = \beta_0 + 0.08x$; raising x by 1 adds 0.08 to $\log y$, i.e. multiplies y by $e^{0.08}$. For small β , $e^\beta - 1 \approx \beta$. In $\log y = \beta_0 + \beta \log x$, $\frac{d \log y}{d \log x} = \beta$ (elasticity).

📊 **Worked example.** Price per sq ft vs area in log-log with $\beta = -0.15$: a 10% larger flat has about 1.5% lower price per sq ft (bulk discount).

Q2.17

How would you check the assumptions of a fitted linear model in Python?

🕒 **In one line.** Plot residuals against fitted values and against time, draw a Q-Q plot, compute VIFs, and run Breusch-Pagan and Durbin-Watson tests using statsmodels.

⚙️ **The full answer.**

```
import statsmodels.api as sm
from statsmodels.stats.diagnostic import het_breuschpagan
from statsmodels.stats.outliers_influence import variance_inflation_factor
Xc = sm.add_constant(X)
m = sm.OLS(y, Xc).fit()
print(m.summary())                                # coefficients, p-values, R2, Durbin-Watson
resid, fitted = m.resid, m.fittedvalues
# plt.scatter(fitted, resid) -> look for curves (linearity) and funnels (
# heteroscedasticity)
sm.qqplot(resid, line="s")                        # normality
print(het_breuschpagan(resid, Xc))                # small p-value -> heteroscedastic
vifs = [variance_inflation_factor(Xc.values, i) for i in range(1, Xc.shape[1])]
m_robust = sm.OLS(y, Xc).fit(cov_type="HC3")      # robust standard errors
```

Q2.18

💡 **Think deeper** Your linear model's training R^2 is 0.35. A colleague says linear regression is useless here. How do you respond?

🕒 **In one line.** A low R^2 only says much of the variance is unexplained; the model can still be useful if its coefficients are reliable or its ranking or average predictions are good enough for the decision, and the gap may be closable with better features rather than a different model.

⚙️ **The full answer.** Human behaviour (who applies, who buys) is noisy; an R^2 of 0.35 can be excellent. Ask: (1) what is the business metric? For ranking leads, ordering matters, not exact values; (2) is it underfitting? Compare with a gradient-boosted model: if it also gets 0.37, the ceiling is the noise, not linearity; if it gets 0.60, add non-linear terms or switch; (3) are features missing? Domain features (price-band concentration in 99acres) often beat model changes.

Q2.19

What does the intercept mean, and should you ever remove it?

🎯 **In one line.** The intercept is the prediction when all features are zero; remove it only when theory demands the line pass through the origin, because forcing it biases the slope otherwise.

⚙️ **The full answer.** Without an intercept, residuals no longer need to sum to zero and R^2 can be misleading (some software reports an uncentred R^2). If features are centred, the intercept equals $\bar{y}$, which is easy to interpret. A case for no intercept: cost = price $\times$ quantity, where zero quantity truly means zero cost.

Q2.20

What is the Gauss–Markov theorem, in plain words?

🎯 **In one line.** Among all estimators that are linear in y and unbiased, OLS has the smallest variance, provided errors have mean zero, constant variance and are uncorrelated; it is “BLUE”.

⚙️ **The full answer.** “Best linear unbiased estimator.” It does not require normality. It does not say OLS beats biased estimators: Ridge is biased but can have lower total error, which is the bias–variance trade-off again.

Q2.21

🧠 **Think deeper** Can you do linear regression when you have more features than rows ($d > n$)?

🎯 **In one line.** Not with plain OLS, because $X^T X$ is singular and infinitely many weight vectors fit the data exactly; with Ridge, Lasso, PCA regression or the minimum-norm solution you can.

⚙️ **The full answer.** With 1,000 skill features and 200 salaries, any 200-dimensional pattern can be fitted perfectly, so the training error is zero and meaningless. Ridge makes the problem well-posed; Lasso selects at most n features; the pseudo-inverse picks the smallest-norm exact fit (which, interestingly, is what gradient descent from zero converges to). Always validate, since all of these can still overfit.

CHAPTER 3

Logistic Regression and Multiclass Classification

🎥 **Companion video for this chapter:** [StatQuest: Logistic Regression](#)

Logistic regression was the model the InfoEdge Round 3 candidate trained in the coding round, and it is still the baseline for click, apply and lead-conversion prediction across the industry. Know it well enough to derive its gradient on paper and to interpret every coefficient.


3.1 Why not just use a straight line for yes/no?

Suppose $y = 1$ if a seeker applies to a job we show them and 0 otherwise. We would like a model that outputs the *probability* of applying. A straight line $w^\top x + b$ can output -3 or 7 , which are not probabilities. We need a function that takes any real number and squeezes it into $(0, 1)$.

3.2 The sigmoid: squeezing a score into a probability

$$\sigma(z) = \frac{1}{1 + e^{-z}}.$$

- z very negative: e^{-z} is huge, so $\sigma(z) \approx 0$.
- $z = 0$: $\sigma(0) = 1/2$.
- z very positive: $e^{-z} \approx 0$, so $\sigma(z) \approx 1$.
- Handy values: $\sigma(1) = 0.731$, $\sigma(2) = 0.881$, $\sigma(-2) = 0.119$, $\sigma(\ln 3) = 0.75$.
- Its derivative is beautifully simple: $\sigma'(z) = \sigma(z)(1 - \sigma(z))$, largest (0.25) at $z = 0$.



Logistic regression computes a linear score $z = \mathbf{w}^\top \mathbf{x} + b$ and outputs $p = \sigma(z) = P(y = 1 \mid \mathbf{x})$.

3.3 Odds and log-odds: how to read the coefficients

The **odds** of an event with probability p are $p/(1 - p)$. Probability 0.8 means odds 4 (“4 to 1”). Inverting the sigmoid gives

$$\log \frac{p}{1 - p} = \mathbf{w}^\top \mathbf{x} + b.$$

📺 **Stuck on this page?** Watch: [StatQuest: Logistic Regression](#)

The left side is the **log-odds** or **logit**. So logistic regression is a *linear model for the log-odds*. This gives the standard interpretation: *increasing x_j by one unit adds w_j to the log-odds, i.e. multiplies the odds by e^{w_j} , holding other features fixed.*

Small example: interpreting a coefficient

In an apply-prediction model, the feature “job is in the seeker’s preferred city” (0/1) has $w = 1.1$. $e^{1.1} = 3.0$: the odds of applying triple. If a seeker’s base probability was 0.10 (odds 0.111), the new odds are 0.333 and the new probability is $0.333/1.333 = 0.25$. The probability did not triple; the *odds* did.

3.4 The decision boundary is a straight line

To turn a probability into a class we use a threshold t , often 0.5. Predict 1 when $\sigma(z) \geq 0.5$, i.e. when $z \geq 0$. The set $\mathbf{w}^\top \mathbf{x} + b = 0$ is a line (a hyperplane in higher dimensions). So, despite the curved sigmoid, logistic regression is a **linear classifier**. A different threshold t just moves the line parallel to itself ($z \geq \log \frac{t}{1-t}$). Curved boundaries need engineered features ($x_1^2, x_1 x_2$).

3.5 Training: maximum likelihood and log loss

3.5.1 From probabilities to a loss

For one example, the model assigns probability p to $y = 1$ and $1 - p$ to $y = 0$. The probability it assigns to what actually happened is $p^y(1 - p)^{1-y}$ (this is p if $y = 1$ and $1 - p$ if $y = 0$). Assuming examples are independent, the probability of the whole dataset is the product of these. **Maximum likelihood** picks the weights that make the observed data most probable. Products are awkward, so take logs (which turns products into sums and does not move the maximum) and flip the sign to get something to minimise:

$$L(\mathbf{w}) = -\frac{1}{n} \sum_{i=1}^n \left[y_i \log p_i + (1 - y_i) \log(1 - p_i) \right].$$

This is the **log loss**, also called **binary cross-entropy**.

Small example: how log loss punishes confidence

True label $y = 1$. If $p = 0.9$: loss = $-\ln 0.9 = 0.105$. If $p = 0.5$: 0.693. If $p = 0.1$: 2.303. If $p = 0.01$: 4.605. Being confidently wrong is punished brutally, which forces honest probabilities.

3.5.2 The gradient, derived

With $p = \sigma(z)$ and $z = \mathbf{w}^\top \mathbf{x} + b$, for one example $\ell = -[y \log p + (1 - y) \log(1 - p)]$. By the chain rule:

$$\frac{\partial \ell}{\partial p} = -\frac{y}{p} + \frac{1 - y}{1 - p}, \quad \frac{\partial p}{\partial z} = p(1 - p), \quad \frac{\partial z}{\partial w_j} = x_j.$$

Multiplying, $\frac{\partial \ell}{\partial z} = -y(1 - p) + (1 - y)p = p - y$, so

$$\frac{\partial \ell}{\partial w_j} = (p - y) x_j, \quad \frac{\partial \ell}{\partial b} = p - y.$$

“Prediction minus truth, times the input”: exactly the same form as linear regression’s gradient. The $p(1 - p)$ from the sigmoid cancelled perfectly; this is why log loss and sigmoid are a natural pair.

3.5.3 No closed form, but a convex problem

Setting the gradient to zero gives equations with no algebraic solution, so we iterate: gradient descent, or Newton’s method (called **IRLS**, iteratively reweighted least squares, because each Newton step is a weighted least-squares fit). The log loss is **convex** in $\mathbf{w}$: shaped like a bowl, with no false valleys, so any of these methods reaches the global optimum.

Small example: one gradient step by hand

One feature, $w = 0.5$, $b = 0$, learning rate $\eta = 0.1$, example $x = 2$, $y = 0$. $z = 1$, $p = \sigma(1) = 0.731$. Gradients: $(p - y)x = 1.462$ and $p - y = 0.731$. Update: $w = 0.5 - 0.1462 = 0.354$, $b = -0.073$. Both moved to lower p , correctly, since the true label is 0.

3.5.4 Why not squared error with a sigmoid?

1. It is non-convex in $\mathbf{w}$: there can be flat plateaus and local minima.
2. Its gradient contains $p(1 - p)$, which is near zero when $p \approx 0$ or 1. A confidently wrong model ($p = 0.001$ for a positive) barely learns: the gradient vanishes exactly when the error is largest.
3. Log loss is the maximum-likelihood loss for Bernoulli labels; squared error corresponds to Gaussian noise, which binary labels do not have.

3.6 Regularisation and separable data

If the classes can be perfectly separated by a line, the log loss can be pushed toward zero by making $\|\mathbf{w}\|$ larger and larger (the sigmoid gets steeper; every probability approaches 0 or 1). Unregularised training never converges: weights grow without bound. An L2 penalty fixes this. scikit-learn’s LogisticRegression applies L2 by default with $C=1.0$, where $C = 1/\lambda$: *smaller C means stronger regularisation*.

3.7 Probabilities and calibration

A model is **calibrated** if, among cases where it says 0.3, about 30% are positive. Logistic regression fitted by maximum likelihood tends to be well calibrated when its features are adequate, which is a big reason it is still used for bidding, ranking and expected-value decisions (“expected applies = sum of probabilities”). Two things break calibration: heavy regularisation, and resampling the classes (training on 50/50 data when reality is 1/99 shifts the intercept; correct it with $b \leftarrow b + \log \frac{\pi_{\text{true}}/(1-\pi_{\text{true}})}{\pi_{\text{train}}/(1-\pi_{\text{train}})}$).

3.8 More than two classes

3.8.1 Softmax (multinomial) regression

One score per class, $z_k = \mathbf{w}_k^\top \mathbf{x} + b_k$, turned into probabilities that sum to one:

$$p_k = \frac{e^{z_k}}{\sum_j e^{z_j}}.$$

Loss: $-\log p_{\text{true class}}$ (categorical cross-entropy).

Small example: softmax by hand

Scores (2, 1, 0): $e^2 = 7.389$, $e^1 = 2.718$, $e^0 = 1$, sum 11.107. Probabilities (0.665, 0.245, 0.090). If the true class is the second, loss = $-\ln 0.245 = 1.41$. Subtracting the maximum score before exponentiating gives the same probabilities and avoids overflow.

With two classes softmax reduces exactly to the sigmoid on $z_1 - z_2$.

3.8.2 One-vs-Rest and One-vs-One

For any binary classifier:

- **One-vs-Rest (OvR, OvA):** K classifiers, “class k vs everything else”; predict the class with the highest score. Each model sees all data but is imbalanced.
- **One-vs-One (OvO):** $\binom{K}{2} = K(K-1)/2$ classifiers, one per pair, each trained only on those two classes; predict by majority vote. Used by kernel SVMs because many small problems are cheaper than a few large ones.

For 5 classes: OvR trains 5 models, OvO trains 10.

3.9 Multi-label is different

Multiclass: exactly one label (a job belongs to one function: IT, Sales, ...). **Multi-label:** any number of labels (a resume can have many skills). Multi-label uses an independent sigmoid per label with binary cross-entropy each, not a softmax.

In practice: logistic regression at InfoEdge scale

- **Apply / click prediction** in the Apply Generation engine and job recommendation: billions of rows, sparse one-hot and hashed features, logistic regression trained with SGD. Fast to score, calibrated, and the coefficients explain behaviour.
- **99acres telecalling:** “probability this user is a high-intent buyer” ranks whom to call first. Calibration lets the business compute expected conversions per caller-hour.
- **Spam / fake profile detection** as a binary classifier, where the panel explicitly said to look at recall.

3.10 Interview questions

Q3.1

★ **PYQ (Round 3 coding round)** Walk me through logistic regression: what it models, how it is trained, and how you would evaluate it on the placement dataset.

🕒 **In one line.** It models the log-odds of the positive class as a linear function of the features, is trained by minimising log loss (maximum likelihood) with gradient-based optimisation, and is evaluated with a confusion matrix, precision, recall, F1 and ROC AUC on a held-out set.

⚙️ **The full answer.** **Model.** Score $z = \mathbf{w}^T \mathbf{x} + b$, probability $p = \sigma(z)$. Equivalently $\log \frac{p}{1-p} = \mathbf{w}^T \mathbf{x} + b$. **Training.** Minimise $-\frac{1}{n} \sum [y \log p + (1 - y) \log(1 - p)]$ plus an L2 penalty. The gradient is $(p - y)\mathbf{x}$; the problem is convex, so the optimiser (lbfgs by default in scikit-learn) finds the global minimum. **Workflow on the placement data.** Drop the leaking salary column; one-hot encode board, stream and work experience; scale the percentage columns (the penalty treats all weights alike); stratified train/test split; fit; predict; print confusion matrix and classification report; look at ROC AUC; choose a threshold to suit the cost of errors.

📄 Worked example.

```
X = df.drop(columns=["status", "salary"]) # salary leaks the target
y = (df["status"] == "Placed").astype(int)
num = X.select_dtypes("number").columns; cat = X.columns.difference(num)
pre = ColumnTransformer([("num", StandardScaler(), num),
                        ("cat", OneHotEncoder(handle_unknown="ignore"), cat)])
clf = Pipeline([("pre", pre), ("lr", LogisticRegression(max_iter=1000))])
Xtr, Xte, ytr, yte = train_test_split(X, y, test_size=0.2, stratify=y, random_state=42)
clf.fit(Xtr, ytr); pred = clf.predict(Xte)
print(confusion_matrix(yte, pred)); print(classification_report(yte, pred))
print(roc_auc_score(yte, clf.predict_proba(Xte)[:, 1]))
```

💡 **If the interviewer pushes.** “How do you read the coefficients?” After scaling, e^{w_j} is the odds multiplier per one standard deviation of feature j . “Why scale?” For the penalty’s fairness and for faster convergence.

Q3.2

Derive the gradient of the log loss for logistic regression.

🕒 **In one line.** By the chain rule the derivative with respect to the score is $p - y$, so the gradient with respect to w_j is $(p - y)x_j$.

⚙️ **The full answer.** Written out in the theory section: $\partial \ell / \partial p = -y/p + (1 - y)/(1 - p)$; $\partial p / \partial z = p(1 - p)$; product $= -y(1 - p) + (1 - y)p = p - y$; times $\partial z / \partial w_j = x_j$. For the whole dataset, $\nabla_{\mathbf{w}} L = \frac{1}{n} \mathbf{X}^T (\mathbf{p} - \mathbf{y})$.

💡 **If the interviewer pushes.** “What is the Hessian?” $\frac{1}{n} \mathbf{X}^T \mathbf{S} \mathbf{X}$ with $\mathbf{S} = \text{diag}(p_i(1 - p_i))$. It is positive semi-definite, which proves convexity, and it is what Newton/IRLS uses.

Q3.3

Why is it called *logistic regression* if it is a classifier?

🎯 **In one line.** Because it regresses (fits a linear model for) the log-odds of the outcome, producing a continuous probability; it becomes a classifier only when you add a threshold.

⚙️ **The full answer.** It is a generalised linear model: a linear predictor, a link function (logit) and a Bernoulli distribution (Chapter 4). The output is a probability, a continuous quantity; classification is a decision we layer on top.

Q3.4

A coefficient for “has premium account” is 0.693. Interpret it.

🎯 **In one line.** Premium users have twice the odds of applying ($e^{0.693} = 2$), other features held fixed; this is not a doubling of probability.

⚙️ **The full answer.** Log-odds increase by 0.693, so odds multiply by 2. From $p = 0.5$ (odds 1) the new odds are 2 and $p = 2/3$; from $p = 0.05$ (odds 0.0526) the new odds are 0.105 and $p = 0.095$. Near small probabilities, odds ratios approximate risk ratios.

Q3.5

Why do we not use mean squared error to train logistic regression?

🎯 **In one line.** Because with a sigmoid output MSE is non-convex and its gradient vanishes when the model is confidently wrong, whereas log loss is convex, is the maximum-likelihood loss for binary labels, and gives the clean gradient $(p - y)x$.

⚙️ **The full answer.** The MSE gradient is $(p - y)p(1 - p)x$. For a positive example predicted $p = 0.001$, $p(1 - p) \approx 0.001$, so the update is a thousand times weaker than it should be. Log loss cancels that factor.

📊 **Worked example.** $y = 1$, $p = 0.01$: log-loss gradient w.r.t. z is -0.99 ; MSE gradient is $-0.99 \cdot 0.0099 \approx -0.0098$.

Q3.6

What happens when the training data are perfectly separable?

🎯 **In one line.** The unregularised maximum-likelihood solution does not exist: weights grow without bound as the optimiser keeps steepening the sigmoid; regularisation gives a finite answer.

⚙️ **The full answer.** If a line separates the classes, scaling $\mathbf{w}$ by $c > 1$ keeps the same boundary but makes every probability more extreme and every loss term smaller, so the loss keeps decreasing toward 0. Software prints convergence warnings and coefficients with huge standard errors. L2 penalty, early stopping, or Firth’s correction fixes it.

Q3.7

 Think deeper Is logistic regression a linear model? Can it learn XOR?

🕒 **In one line.** It is linear in its decision boundary (a hyperplane in feature space), so on raw features it cannot learn XOR; with an interaction feature x_1x_2 it can.

⚙️ **The full answer.** For XOR points $(0,0) \rightarrow 0, (1,1) \rightarrow 0, (0,1) \rightarrow 1, (1,0) \rightarrow 1$, no line separates the classes. Adding $x_3 = x_1x_2$: the score $z = 10x_1 + 10x_2 - 20x_3 - 5$ gives $-5, 5, 5, -5$ for the four points: correct. The model is still linear in its parameters; the feature space was enriched. This is the same idea as kernels (Chapter 9) and hidden layers (Chapter 18).

Q3.8

 PYQ (OA: one-vs-one and one-vs-all) You have 6 job functions. How many binary classifiers do OvR and OvO train, and when would you prefer each?

🕒 **In one line.** OvR trains 6, OvO trains 15; OvR is simpler and cheaper for most models, OvO suits algorithms whose cost grows faster than linearly with data size, like kernel SVMs.

⚙️ **The full answer.** OvO: $\binom{6}{2} = 15$, each on only two classes' data (about a third of the rows here). For kernel SVMs, training cost is between $O(n^2)$ and $O(n^3)$, so 15 small problems can be cheaper than 6 big ones. OvR gives each model a heavily imbalanced problem (one class vs five) and its scores may not be comparable across models. For logistic regression, use true multinomial softmax instead of either.

Q3.9

Softmax vs independent sigmoids: when do you use which?

🕒 **In one line.** Softmax when classes are mutually exclusive (exactly one is true); independent sigmoids when labels can co-occur (multi-label).

⚙️ **The full answer.** Softmax probabilities sum to 1, so raising one lowers the others; that is right for "which job function" and wrong for "which skills does this resume contain". For skills, use one sigmoid per skill with binary cross-entropy, and tune a threshold per skill.

Q3.10

How does the decision threshold affect precision and recall? How do you choose it?

🕒 **In one line.** Raising the threshold makes the model predict positive less often: precision usually rises and recall falls; choose it from business costs, a target recall or precision, or by maximising F1 on validation data.

⚙️ **The full answer.** The model's probabilities are fixed; the threshold only decides where to cut. For spam the panel wanted recall considered; for a telecalling queue with 100 calls per day, choose the cut that fills the queue with the highest probabilities. Formally, with cost C_{FP} and C_{FN} and calibrated p , predict positive when $p > C_{FP}/(C_{FP} + C_{FN})$.

 **Worked example.** If missing a fraudulent recruiter costs 10 times a false alarm, threshold $= 1/(1 + 10) = 0.09$.

Q3.11

 **Think deeper** Your training set was rebalanced from 2% positives to 50% by undersampling. What happens to the predicted probabilities and how do you fix them?

 **In one line.** They become inflated, because the model learned a 50% base rate; the ranking is mostly unaffected, and you fix the probabilities by shifting the intercept by the log of the odds ratio between true and training base rates (or by recalibrating on untouched data).

 **The full answer.** Undersampling negatives by a factor β multiplies the odds of every prediction by $1/\beta$, which in a logistic model only changes the intercept. With true rate $\pi = 0.02$ (odds 0.0204) and training rate 0.5 (odds 1), subtract $\log(1/0.0204) = 3.89$ from b . Alternatively fit isotonic or Platt calibration on a validation set with the real class balance.

 **Worked example.** A raw output of 0.6 (log-odds 0.405) corrects to log-odds $0.405 - 3.89 = -3.49$, i.e. $p = 0.030$.

Q3.12

What is C in scikit-learn's LogisticRegression? What are the defaults?

 **In one line.** C is the inverse regularisation strength ($C = 1/\lambda$); the default is L2 with $C = 1.0$, so smaller C means stronger shrinkage.

 **The full answer.** The objective is $\frac{1}{2}\|\mathbf{w}\|^2 + C \sum_i \ell_i$. Large C trusts the data more (less regularisation), small C shrinks more. Default solver is lbfgs (L2 only); liblinear or saga support L1. Tune C on a log grid such as $10^{-3}, \dots, 10^3$ with LogisticRegressionCV.

Q3.13

 **Think deeper** Logistic regression vs a linear SVM: what is really different?

 **In one line.** Both learn a linear boundary with an L2 penalty; they differ only in the loss: log loss (smooth, every point contributes, outputs probabilities) versus hinge loss (zero for points beyond the margin, only support vectors matter, outputs scores).

 **The full answer.** Hinge $\max(0, 1 - yz)$ with $y \in \{-1, 1\}$ is exactly zero once a point is correctly classified with margin, so confidently correct points stop influencing the boundary. Log loss $\log(1 + e^{-yz})$ never reaches zero, so all points keep a small influence. In practice accuracy is often similar; logistic regression is preferred when calibrated probabilities are needed, SVM when the margin-based robustness to far-away points helps or with kernels.

Q3.14

What is the decision boundary of logistic regression with features x_1 and x_2 and weights $w = (2, -1)$, $b = 1$? Classify the point $(1, 4)$.

🕒 **In one line.** The boundary is the line $2x_1 - x_2 + 1 = 0$; the point $(1, 4)$ has score -1 , probability 0.269 , so it is class 0 at threshold 0.5 .

⚙️ **The full answer.** $z = 2 \cdot 1 - 4 + 1 = -1$; $\sigma(-1) = 1/(1 + e) = 0.269$. Points above the line $x_2 = 2x_1 + 1$ are class 0.

Q3.15

How would you handle a categorical feature with 3,000 levels (for example company name) in logistic regression?

🕒 **In one line.** One-hot encode it with regularisation (sparse is fine for linear models), or group rare levels, hash it, or use target encoding computed out-of-fold.

⚙️ **The full answer.** One-hot creates 3,000 sparse columns; with L2 and enough data per level this works well and is standard in click models. Rare companies (a few rows) get weights shrunk toward 0, which is sensible. Alternatives: map to a taxonomy (the SilverSkAI-style normalisation collapses variants like “TCS”, “Tata Consultancy”), an “other” bucket for levels with < 20 rows, feature hashing, or out-of-fold target encoding (never in-fold: that leaks the label).

Q3.16

🧠 **Think deeper** Can logistic regression be used for regression of a proportion (for example the fraction of recruiters who open a profile)?

🕒 **In one line.** Yes: fit a binomial GLM with the logit link, using counts of successes and trials, or a fractional logit; the sigmoid keeps predictions inside $[0, 1]$.

⚙️ **The full answer.** The binomial GLM models k successes out of n trials with $P = \sigma(\mathbf{w}^\top \mathbf{x})$; the log loss weighted by n is its likelihood. scikit-learn's `LogisticRegression` accepts only 0/1 labels but you can expand rows or use `sample_weight`; statsmodels' `GLM(family=Binomial())` accepts proportions with `var_weights`. This avoids the main weaknesses of linear regression on proportions (predictions outside $[0, 1]$ and non-constant variance $p(1 - p)/n$).

Q3.17

What is the log loss of a model that always predicts the base rate π ? Why is it a useful baseline?

🕒 **In one line.** It is the entropy of the labels, $-\pi \log \pi - (1 - \pi) \log(1 - \pi)$; any useful model must beat it.

⚙️ **The full answer.** With $\pi = 0.1$: $-(0.1 \ln 0.1 + 0.9 \ln 0.9) = 0.230 + 0.095 = 0.325$. A model at 0.30 is only modestly informative. Normalised cross-entropy (model log loss divided by this) is a common metric in ad/click teams.

Q3.18

How do you get feature importance from logistic regression, and what can go wrong?

🎯 **In one line.** Use the absolute values of coefficients after standardising features (or odds ratios), and their confidence intervals; collinearity, unscaled features and regularisation can all distort the picture.

⚙️ **The full answer.** Without scaling, a coefficient depends on units. With correlated features, credit is split arbitrarily. With L1, one of a correlated group survives. Permutation importance on validation data is a model-agnostic check.

Q3.19

🧠 **Think deeper** Logistic regression gives probability 0.97 for a profile being fake. Should you auto-ban the user?

🎯 **In one line.** Only if the model is calibrated at that range, the cost of a false ban has been weighed, and the precision at that threshold is verified on recent labelled data; otherwise route to human review.

⚙️ **The full answer.** A probability is only meaningful if calibrated (check a reliability curve on recent data). Even at 97% precision, 3% of bans are wrong; for paying recruiters that may be unacceptable. A common design: auto-action above a very high threshold with measured precision > 99.5%, human review between, no action below. Monitor drift, because fraud adapts.

Q3.20

Compute the probability and log loss for $z = 1.2$ and true label 1.

🎯 **In one line.** $p = \sigma(1.2) = 0.769$, $\log \text{loss} = -\ln 0.769 = 0.263$.

⚙️ **The full answer.** $e^{-1.2} = 0.301$; $p = 1/1.301 = 0.769$; $-\ln(0.769) = 0.263$.

CHAPTER 4

Generalised Linear Models and the Regression Family

🎥 **Companion video for this chapter:** [StatQuest: Generalised Linear Models \(search\)](#)

The InfoEdge hire in the experience video said the single biggest change in his understanding came from learning the “variants most courses skip”: count data needs *Poisson regression*, which comes from GLMs; SVMs can do *regression*. In Round 4 he answered a heteroscedasticity question with a GLM and the interviewer accepted it. This chapter makes you fluent in that family. It is a high-return chapter because most candidates have never seen it.

4.1 The question GLMs answer

Ordinary linear regression quietly assumes three things about the target y :

1. its average is a straight-line function of the features;
2. it can be any real number, negative included;
3. its noise has the same spread everywhere (and is roughly bell-shaped).

Many real targets break all three:

- **Counts:** number of applications a job gets today ($0, 1, 2, \dots$). Never negative; a job that averages 2 applies varies by about ± 1.4 , one that averages 200 varies by about ± 14 . Spread grows with the mean.
- **Positive, skewed amounts:** salary, property price, time-to-hire. Never negative; spread grows in proportion to the level.
- **Yes/no and proportions:** did the user apply; what fraction of mails were opened. Must stay in $[0, 1]$.

A **generalised linear model (GLM)** keeps the convenient “weighted sum of features” but lets us choose (a) what distribution the target follows and (b) how the weighted sum is connected to the target’s average.

4.2 The three ingredients of a GLM

1. **A linear predictor** (the familiar part): $\eta = \mathbf{w}^\top \mathbf{x} + b$. It can be any real number.
2. **A link function** g that connects the average of y , written $\mu = \mathbb{E}[y \mid \mathbf{x}]$, to the linear predictor: $g(\mu) = \eta$, or equivalently $\mu = g^{-1}(\eta)$. The link’s job is to map the allowed range of μ onto the whole real line.
3. **A distribution** for y around its mean, from the *exponential family* (Gaussian, Bernoulli/Binomial, Poisson, Gamma, Inverse Gaussian, Tweedie, $\dots$). The distribution fixes how the variance depends on the mean: $\text{Var}(y) = \phi V(\mu)$, where V is the **variance function** and ϕ a dispersion constant.

Target type	Distribution	Variance $V(\mu)$	Usual link	Name
Any real number	Gaussian	1 (constant)	identity $\mu = \eta$	ordinary linear regression
0/1 or proportion	Bernoulli / Binomial	$\mu(1 - \mu)$	logit $\log \frac{\mu}{1-\mu} = \eta$	logistic regression
Count $0, 1, 2, \dots$	Poisson	μ	$\log \log \mu = \eta$	Poisson regression
Overdispersed count	Negative binomial	$\mu + \mu^2/k$	log	NB regression
Positive, skewed	Gamma	μ^2	log (or inverse)	Gamma regression
Zeros + positive amounts	Tweedie ($1 < p < 2$)	μ^p	log	Tweedie regression

💡 Intuition: the link is a translator, the distribution is a noise model

The linear predictor speaks “any real number”. A count’s mean speaks “positive numbers”. The log link translates: $\mu = e^\eta$ is always positive, whatever η is. The distribution then says how scattered the real counts are around that mean. Ordinary regression is the special case where no translation is needed and the scatter is the same everywhere.

4.3 Poisson regression, built from scratch

4.3.1 The Poisson distribution in one paragraph

If events happen independently at an average rate μ per interval, the count Y in one interval follows $P(Y = k) = e^{-\mu} \mu^k / k!$ for $k = 0, 1, 2, \dots$. Its mean *and* its variance are both μ .

📊 Small example: Poisson probabilities

A job averages $\mu = 2$ applies per hour. $P(0) = e^{-2} = 0.135$; $P(1) = 2e^{-2} = 0.271$; $P(2) = 2e^{-2} = 0.271$; $P(3) = \frac{8}{6}e^{-2} = 0.180$.

4.3.2 The model

$$y \mid \mathbf{x} \sim \text{Poisson}(\mu), \quad \log \mu = \mathbf{w}^\top \mathbf{x} + b \iff \mu = e^{\mathbf{w}^\top \mathbf{x} + b}.$$

Effects are *multiplicative*: raising x_j by one unit multiplies the expected count by e^{w_j} .

📊 Small example: reading a Poisson coefficient

Daily applies to a job: $\log \mu = 1.5 + 0.4 \cdot \text{remote} - 0.05 \cdot \text{days_since_posting}$. A fresh, non-remote job: $\mu = e^{1.5} = 4.48$ applies/day. Remote: $e^{0.4} = 1.49$, so 49% more, 6.7/day. After 10 days: multiply by $e^{-0.5} = 0.61$.

4.3.3 How it is fitted

Maximum likelihood again. For one row, $\log P(y \mid \mu) = y \log \mu - \mu - \log y!$. With $\log \mu = \eta$, the negative log-likelihood to minimise (dropping the constant) is

$$L(\mathbf{w}) = \sum_i (e^{\eta_i} - y_i \eta_i).$$

Its gradient with respect to $\mathbf{w}$ is $\sum_i (\mu_i - y_i) \mathbf{x}_i$: *prediction minus truth, times input*, exactly the same form as linear and logistic regression. This is no coincidence: it is a property of every GLM that uses its **canonical link**, and it is why one algorithm (IRLS / Newton) fits them all. The loss is convex, so the optimum is global.

4.3.4 Exposure and offsets

Counts often come from different-sized windows: one job was live for 3 days, another for 30. Model the *rate*: $\log \mu_i = \log(\text{exposure}_i) + \mathbf{w}^\top \mathbf{x}_i$. The term $\log(\text{exposure})$ is an **offset**: a feature whose coefficient is fixed at 1. Now e^{w_j} is a rate ratio.

4.3.5 Overdispersion and zero inflation

Real counts usually vary *more* than Poisson allows (variance > mean), because of unobserved differences between units. Symptoms: residual deviance much larger than its degrees of freedom; standard errors far too small.

- **Quasi-Poisson**: keep the Poisson mean model, scale the variance by $\phi > 1$ (fixes standard errors).
- **Negative binomial**: variance $\mu + \mu^2/k$; a Poisson whose rate itself varies (Gamma-distributed).
- **Zero-inflated** or **hurdle** models: when there are far more zeros than Poisson predicts (most jobs get zero applies from a given city), model “zero or not” and “how many, if any” separately.

4.4 Gamma regression: positive, skewed, multiplicative noise

The Gamma distribution lives on positive numbers and has variance proportional to μ^2 , i.e. a *constant coefficient of variation*: a 10 lakh salary varies by about ± 1.5 lakh and a 100 lakh salary by about ± 15 lakh. With a log link, $\log \mathbb{E}[y] = \mathbf{w}^\top \mathbf{x} + b$.

Why prefer it over “regress $\log y$ with OLS”?

- OLS on $\log y$ estimates $\mathbb{E}[\log y]$, and $e^{\mathbb{E}[\log y]}$ is the geometric mean, which is *below* the arithmetic mean $\mathbb{E}[y]$. You need a correction to predict average salary. Gamma-log models $\mathbb{E}[y]$ directly.
- It handles the heteroscedasticity through the variance function rather than through a transformation.

This is exactly the answer the candidate gave for violated homoscedasticity: “model the variance explicitly, keep a linear model on top”.

4.5 Tweedie: when many values are exactly zero

Revenue per recruiter this month, insurance claims, money spent by a user: most are exactly 0, the rest positive and skewed. The Tweedie family with power $1 < p < 2$ is a compound Poisson–Gamma: a Poisson number of Gamma-sized amounts, so it has a point mass at 0 and a continuous positive part. LightGBM and XGBoost offer `objective="tweedie"`.

4.6 Beyond the mean: other regression losses

GLMs change the distribution. Another way to vary regression is to change *what summary of y* you predict, or *how you treat large errors*.

4.6.1 Quantile regression

Instead of the mean, predict a quantile: the 90th percentile salary for a role, so that “90% of offers are below this”. The **pinball (quantile) loss** for quantile τ is

$$\ell_{\tau}(y, \hat{y}) = \begin{cases} \tau(y - \hat{y}) & y \geq \hat{y} \\ (1 - \tau)(\hat{y} - y) & y < \hat{y} \end{cases}.$$

Under-prediction costs τ per unit and over-prediction $1 - \tau$. At $\tau = 0.5$ it is half the absolute error, and the minimiser is the median.

Small example: why the pinball loss finds the 90th percentile

Values $1, 2, \dots, 10$, $\tau = 0.9$. If $\hat{y} = 9$: one value above (cost 0.9×1) and eight below (cost $0.1 \times (8 + 7 + \dots + 1) = 3.6$), total 4.5. If $\hat{y} = 10$: nothing above, cost $0.1 \times (9 + 8 + \dots + 1) = 4.5$. If $\hat{y} = 5$: above costs $0.9 \times (1 + 2 + 3 + 4 + 5) = 13.5$, below $0.1 \times 10 = 1$; total 14.5. The minimum is around the 90th percentile, where 90% of points are below.

Uses: **salary ranges** (10th–90th percentile band in a salary engine), delivery-time promises, prediction intervals from gradient boosting (`loss="quantile"`).

4.6.2 Robust regression: Huber and friends

Huber loss is quadratic for small errors ($|e| \leq \delta$) and linear for large ones: $\ell = \frac{1}{2}e^2$ if $|e| \leq \delta$, else $\delta(|e| - \frac{1}{2}\delta)$. Small errors are treated like OLS (efficient); large errors are treated like MAE (outliers cannot dominate). Other robust options: least absolute deviations (median regression), RANSAC (fit on random subsets, keep the model with most inliers), Theil–Sen.

4.6.3 Support Vector Regression (SVR)

The SVM idea applied to regression. Use the ϵ -**insensitive loss**: errors smaller than ϵ cost *nothing*; beyond that, cost grows linearly:

$$\ell_{\epsilon}(y, \hat{y}) = \max(0, |y - \hat{y}| - \epsilon).$$

Picture a tube of half-width ϵ around the prediction. Points inside the tube are ignored; points on or outside it become **support vectors** and alone determine the fit. Add an L2 penalty $\frac{1}{2}\|\mathbf{w}\|^2$ for flatness:

$$\min_{\mathbf{w}, b} \frac{1}{2}\|\mathbf{w}\|^2 + C \sum_i \max(0, |y_i - \mathbf{w}^T \mathbf{x}_i - b| - \epsilon).$$

Because the dual form involves only dot products, **kernels** apply (RBF SVR fits smooth curves). Hyperparameters: C (penalty for points outside the tube), ϵ (tube width, the tolerated error), and the kernel's γ . Chapter 9 develops SVMs fully.

Small example: the ϵ -insensitive loss

$\epsilon = 0.5$. Errors 0.3, -0.5, 1.2, -2.0 cost 0, 0, 0.7, 1.5.

4.6.4 Other regressors to be able to name

- **Kernel ridge regression:** ridge regression with the kernel trick; squared loss, so every point matters (unlike SVR's sparsity); closed-form solution $\alpha = (K + \lambda I)^{-1}y$.
- **Gaussian process regression:** a Bayesian kernel method that returns a mean *and* an uncertainty.
- **Isotonic regression:** fits the best monotone (non-decreasing) step function; used to calibrate classifier probabilities.
- **k-NN regression, regression trees, random forests, gradient boosting, neural networks:** all can regress.
- **Survival (time-to-event) models:** time-to-hire or time-to-churn with censoring (jobs still open), e.g. Cox proportional hazards. Ordinary regression on "days open" is biased because open jobs have not finished.

4.7 Choosing a regression model: a decision guide

Situation	Reach for
Continuous, roughly symmetric noise	OLS / Ridge
Counts (applies, clicks, calls)	Poisson (or negative binomial) with log link and exposure offset
Positive, right-skewed, spread $\propto$ level (salary, price)	Gamma GLM with log link, or OLS on $\log y$
Many exact zeros plus positive amounts (revenue)	Tweedie, or a two-part (hurdle) model
Need a range or a pessimistic estimate	Quantile regression (or quantile GBM)
Outliers in y	Huber / LAD / RANSAC
Non-linear, smallish data, smooth function	SVR or kernel ridge with RBF kernel, or Gaussian process
Big tabular data with interactions	Gradient boosting (with the matching objective: poisson, gamma, tweedie, quantile)
Time-to-event with censoring	Survival models

In practice: GLMs at InfoEdge

- **Applies per job per day** (Job Posting engine; 100 crore+ applies a year): Poisson/negative binomial with exposure, used to forecast which jobs will under-perform and need promotion via mailers or WhatsApp.
- **Salary engine** (Data Factory): Gamma-log or log-OLS for the mean, quantile regression for the band shown in Naukri TopTier's "salary intelligence" (10th/50th/90th percentile).
- **99acres:** number of enquiries per listing (Poisson); price per square foot (Gamma or log-OLS).
- **Revenue per recruiter** (many zero months): Tweedie.

4.8 Interview questions

Q4.1

★ **PYQ (reported by InfoEdge hire: count data)** You need to predict the number of applications a job posting will receive tomorrow. Why is ordinary linear regression a poor choice, and what would you use?

🕒 **In one line.** Counts are non-negative integers whose variance grows with their mean, and effects on them are naturally multiplicative; Poisson regression (a GLM with a log link), or negative binomial if overdispersed, models all three.

⚙️ **The full answer.** Problems with OLS on counts:

1. It can predict negative applies for unattractive jobs.
2. Its noise model assumes the same spread for a job expecting 1 apply and one expecting 500; in reality the spread grows like $\sqrt{\mu}$. OLS over-weights the noisy high-count jobs and its intervals are wrong.
3. Effects are multiplicative in reality ("remote jobs get 50% more applies", not "7 more applies" regardless of baseline). A linear model forces additivity.

Poisson regression: $y \sim \text{Poisson}(\mu)$, $\log \mu = \mathbf{w}^\top \mathbf{x} + b + \log(\text{hours live})$ (offset for exposure). Predictions are positive; e^{w_j} is a rate ratio; the variance equals the mean by assumption. **Check overdispersion:** if the Pearson χ^2 divided by degrees of freedom is well above 1 (say 3), switch to negative binomial or quasi-Poisson. If there are many structural zeros (expired jobs), use a zero-inflated model.

📊 **Worked example.** Fitted $\log \mu = 0.8 + 0.4 \text{ remote} + 0.3 \text{ top_company} - 0.07 \text{ days}$. A remote job at a top company on day 0: $\mu = e^{1.5} = 4.5$ applies/day; on day 10: $4.5 \times e^{-0.7} = 2.2$.

💡 **If the interviewer pushes.** "Could you just log-transform and use OLS?" $\log(0)$ is undefined; $\log(y + 1)$ is a hack that biases results for small counts. "Could gradient boosting do it?" Yes, with objective="poisson", which uses the same log link and Poisson deviance loss, capturing non-linearities.

⚠️ **Trap to avoid.** Saying "use classification because counts are discrete". Counts are ordered and unbounded; a regression with a count distribution is the right framing.

Q4.2

★ **PYQ (reported: GLM as the answer to heteroscedasticity)** What is a generalised linear model? Name its components and explain how it generalises linear regression.

🕒 **In one line.** A GLM has a linear predictor $\eta = \mathbf{w}^\top \mathbf{x} + b$, a link function g with $g(\mathbb{E}[y]) = \eta$, and a distribution from the exponential family that fixes how the variance depends on the mean; linear regression is the special case with Gaussian noise and identity link.

⚙️ **The full answer.** Linear regression ties together three choices: the mean is linear, the target is unbounded, the noise is constant and Gaussian. A GLM unties them. The link handles the range of the mean (log for positive, logit for (0, 1)). The distribution handles the noise shape

and the mean–variance relationship (Poisson: $\text{Var} = \mu$; Gamma: $\text{Var} \propto \mu^2$; Bernoulli: $\mu(1 - \mu)$). Fitting is by maximum likelihood, done with iteratively reweighted least squares: each Newton step is a weighted linear regression on a “working response”. Goodness of fit is measured by **deviance**, twice the log-likelihood gap between the model and a perfect (saturated) model; for the Gaussian it reduces to the sum of squared residuals.

 **Worked example.** Three GLMs on the same features: Gaussian-identity for “rating 1–5” treated as continuous, Bernoulli-logit for “applied or not”, Poisson-log for “number of profile views”.

 **If the interviewer pushes.** “Why must the distribution be in the exponential family?” Because then the likelihood has a common form $\exp\{(y\theta - b(\theta))/\phi + c(y, \phi)\}$, which gives a single fitting algorithm, convexity with the canonical link, and the neat gradient $(\mu - y)\mathbf{x}$.

Q4.3

★ **PYQ (reported: SVR)** Can an SVM be used for regression? Explain how.

 **In one line.** Yes: support vector regression fits a function that keeps as many points as possible inside a tube of half-width ε , ignores errors smaller than ε , penalises larger errors linearly with weight C , keeps the function flat with an L2 penalty, and uses kernels for non-linear fits.

 **The full answer.** Classification SVM: find the widest margin separating classes; only points on or inside the margin (support vectors) matter. SVR flips the picture: we want a function such that points lie *within* a band around it. Formally

$$\min \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_i (\xi_i + \xi_i^*) \quad \text{s.t.} \quad y_i - f(\mathbf{x}_i) \leq \varepsilon + \xi_i, \quad f(\mathbf{x}_i) - y_i \leq \varepsilon + \xi_i^*, \quad \xi, \xi^* \geq 0.$$

ξ and ξ^* measure how far a point sticks out above or below the tube. The dual involves only dot products $\mathbf{x}_i^\top \mathbf{x}_j$, so replacing them with a kernel $K(\mathbf{x}_i, \mathbf{x}_j)$ gives non-linear regression. The solution depends only on points on or outside the tube: **sparse** in the data, like SVM classification. **Hyperparameters:** ε (how much error is tolerated for free; larger ε means fewer support vectors and a smoother fit), C (how strongly violations are penalised), γ (RBF width). Scale features first.

 **Worked example.** $\varepsilon = 1$ lakh for a salary model: predictions within 1 lakh incur no loss, so the model concentrates on the large errors; points within the tube do not influence the fit.

 **If the interviewer pushes.** “SVR vs kernel ridge?” Same kernel trick; kernel ridge uses squared loss (every point matters, closed form, faster to fit for moderate n), SVR uses ε -insensitive loss (sparse, robust to small noise, linear in large errors so less outlier-sensitive). “When is SVR a good choice?” Small to medium data (up to tens of thousands of rows), smooth non-linear relations, a natural tolerance level.

 **Trap to avoid.** Saying “SVMs are only for classification”, or “points inside the tube are support vectors” (it is the reverse).

Q4.4

 **Think deeper** Is logistic regression a GLM? Is linear regression? What is a “link function” in each?

🎯 **In one line.** Both are GLMs: linear regression is Gaussian with the identity link, logistic regression is Bernoulli with the logit link $\log \frac{\mu}{1-\mu}$.

⚙️ **The full answer.** For a binary y , $\mu = P(y = 1)$ lives in $(0, 1)$. The logit maps $(0, 1)$ onto the real line, so the linear predictor can be any number. Its inverse is the sigmoid. The Bernoulli variance $\mu(1 - \mu)$ is the variance function. Other links for binary data exist: **probit** (inverse normal CDF, popular in economics) and **complementary log-log** (asymmetric, for rare events). Logit is the canonical link, giving the tidy gradient and interpretable odds ratios.

Q4.5

🧠 **Think deeper** Why does every canonical-link GLM have the gradient “(prediction – target) × input”?

🎯 **In one line.** Because in the exponential family the derivative of the log-partition function is the mean, so with the canonical link the score equation becomes $\sum (y_i - \mu_i) \mathbf{x}_i = 0$.

⚙️ **The full answer.** An exponential-family density can be written $p(y) = h(y) \exp(\theta y - A(\theta))$, where θ is the natural parameter and A the log-partition function, with $A'(\theta) = \mu$. The canonical link sets $\theta = \eta = \mathbf{w}^\top \mathbf{x}$. The log-likelihood of one point is $\theta y - A(\theta)$, whose derivative with respect to $\mathbf{w}$ is $(y - A'(\theta))\mathbf{x} = (y - \mu)\mathbf{x}$. Gaussian ($A = \theta^2/2$), Bernoulli ($A = \log(1 + e^\theta)$) and Poisson ($A = e^\theta$) all fit this.

💡 **If the interviewer pushes.** “Why does that matter?” One generic fitting routine (IRLS) and one gradient code path serve all of them; it is also why neural-network output layers pair sigmoid with BCE and softmax with cross-entropy.

Q4.6

🧠 **Think deeper** Salary is right-skewed and heteroscedastic. Compare: OLS on $\log(\text{salary})$, a Gamma GLM with log link, and gradient boosting with MSE.

🎯 **In one line.** Log-OLS and Gamma-log both model multiplicative effects and growing spread, but log-OLS predicts the geometric mean and needs a back-transform correction while Gamma-log predicts the mean directly; GBM with MSE fits non-linear effects but is dominated by the largest salaries unless you change its objective or target.

⚙️ **The full answer.**

- **Log-OLS:** fit $\mathbb{E}[\log y]$. Predicting $e^{\hat{\eta}}$ underestimates the mean by a factor $e^{\sigma^2/2}$ (Duan smearing fixes it). Coefficients are percent effects. Easy and robust.
- **Gamma-log:** $\log \mathbb{E}[y] = \eta$ and $\text{Var} \propto \mu^2$. Coefficients are percent effects on the mean; no back-transform; requires $y > 0$.
- **GBM with MSE on raw salary:** captures interactions (city × title) but squared error on raw scale makes the high earners dominate the loss. Better: GBM on $\log y$, or GBM with `objective="gamma"`.

Choose by the question: for an interpretable salary report, log-OLS or Gamma; for maximum accuracy, GBM with the gamma objective, validated with a scale-aware metric (MAPE or RMSLE).

Q4.7

What is overdispersion and how do you detect and handle it?

🎯 **In one line.** Overdispersion is count data varying more than the Poisson model's variance = mean allows; detect it with deviance or Pearson χ^2 per degree of freedom well above 1, and handle it with quasi-Poisson, negative binomial, or random effects.

⚙️ **The full answer.** Unmodelled differences (some recruiters write better job descriptions) make counts clump. Under Poisson the standard errors are then too small, so you would wrongly declare features significant. Negative binomial adds a Gamma-distributed multiplier to each rate, giving $\text{Var} = \mu + \mu^2/k$.

📊 **Worked example.** Mean daily applies 5, sample variance 23: ratio 4.6, strong overdispersion.

Q4.8

🧠 Think deeper You need a salary *range* for a job title, not a single number. How would you build it?

🎯 **In one line.** Use quantile regression (linear or gradient-boosted) to predict, for example, the 10th, 50th and 90th percentiles directly with the pinball loss; alternatively predict mean and variance with a distributional model.

⚙️ **The full answer.** A mean $\pm$ constant band is wrong under heteroscedasticity. Quantile regression lets each quantile depend on features differently (senior roles have wider bands). Train one model per quantile with loss ℓ_τ ; enforce non-crossing (sort the predictions or fit jointly). Evaluate with pinball loss and **coverage**: on validation data, about 80% of actual salaries should fall between the 10th and 90th predictions.

📊 **At InfoEdge scale.** This is what a “salary intelligence” card (Naukri TopTier: “you earn more than 68% of software engineers in Delhi”) needs: a conditional distribution, not a point.

Q4.9

Explain the pinball loss and prove that, for $\tau = 0.5$, its minimiser is the median.

🎯 **In one line.** It charges τ per unit of under-prediction and $1 - \tau$ per unit of over-prediction; at $\tau = 0.5$ both are 0.5, so it is half the absolute error, which is minimised by the median.

⚙️ **The full answer.** Minimise $\sum_i \ell_\tau(y_i, c)$ over a constant c . Moving c up by a tiny δ changes the loss by δ times $[(1 - \tau) \cdot \#\{y_i < c\} - \tau \cdot \#\{y_i > c\}]$. At the optimum this is zero: the fraction of points below c equals τ . So c is the τ -quantile; for $\tau = 0.5$, the median.

Q4.10

🧠 Think deeper What is Huber loss and when would you use it instead of MSE or MAE?

🎯 **In one line.** Huber is squared for small errors and linear for large ones; use it when the data

are mostly well-behaved but have some outliers, because it keeps OLS's efficiency on the bulk and MAE's robustness on the tail.

⚙️ **The full answer.** MSE's gradient is the error itself, so one huge error dominates the update. MAE's gradient is ± 1 everywhere: robust but not smooth at 0 and less efficient for Gaussian-like noise. Huber's gradient is the error, clipped at $\pm \delta$. δ sets where "outlier" starts; about 1.35σ gives 95% of OLS efficiency under Gaussian noise. Deep learning's "smooth L1" (in object detection) is the same idea; "gradient clipping" is its cousin.

Q4.11

💡 **Think deeper** Can a classification model be turned into a regression model, and vice versa?

🎯 **In one line.** Yes: regression can be framed as classification over bins (ordinal or binned targets), and a classifier's probability output is itself a regression of $\mathbb{E}[y]$ for a 0/1 target; many algorithms (trees, k -NN, SVM, neural nets) have both forms with only the loss or leaf rule changed.

⚙️ **The full answer.** Examples: predicting salary band (0–5, 5–10, 10–20 lakh) as multiclass when only bands are needed or when the distribution is multimodal; ordinal regression (cumulative logit) for star ratings 1–5; decision trees use mean in leaves for regression and majority vote for classification; SVM uses hinge or ϵ -insensitive loss. Binning loses ordering information unless you use an ordinal model.

Q4.12

What is deviance in a GLM, and how is it used?

🎯 **In one line.** Deviance is twice the difference in log-likelihood between a perfect (saturated) model and your model; it generalises the residual sum of squares, and differences in deviance compare nested models.

⚙️ **The full answer.** For Gaussian, deviance = SSE/σ^2 . For Poisson, $D = 2 \sum [y \log(y/\mu) - (y - \mu)]$. Null deviance (intercept only) minus residual deviance measures what the features explain; "pseudo- R^2 " = $1 - D_{\text{model}}/D_{\text{null}}$. Two nested models are compared by the deviance difference, roughly χ^2 with degrees of freedom equal to the number of extra parameters.

📊 **Worked example.** Null deviance 1,200, model deviance 800: pseudo- $R^2 = 0.33$.

Q4.13

💡 **Think deeper** How would you model "time to fill a vacancy" when 30% of jobs in your data are still open?

🎯 **In one line.** With a survival model, because open jobs are right-censored: we know they took at least their current age, and dropping them or treating their age as the outcome biases the estimate downward.

⚙️ **The full answer.** Ordinary regression on closed jobs only selects the fast-filling ones; using the current age of open jobs as if it were the fill time understates it. Survival analysis uses the full

information: Kaplan–Meier curves to describe, Cox proportional hazards ($h(t | x) = h_0(t)e^{\mathbf{w}^\top \mathbf{x}}$) to model, or accelerated failure time models. Gradient boosting has survival objectives too (survival:cox in XGBoost).

Q4.14

What is an offset in Poisson regression? Give an example.

🕒 **In one line.** A term added to the linear predictor with its coefficient fixed at 1, typically $\log(\text{exposure})$, so the model predicts a rate per unit of exposure.

⚙️ **The full answer.** $\log \mu_i = \log t_i + \mathbf{w}^\top \mathbf{x}_i$ means $\mu_i = t_i e^{\mathbf{w}^\top \mathbf{x}_i}$: expected count is exposure times rate. Example: listing enquiries on 99acres where listings were live 2 to 60 days: offset $\log(\text{days live})$. Without it, “old listings get more enquiries” masquerades as a feature effect.

Q4.15

🧠 **Think deeper** A colleague fits OLS to predict “probability of hire” and gets some predictions of 1.2. Fix the model without changing to a black box.

🕒 **In one line.** Use a binomial GLM (logistic regression) so predictions stay in $(0, 1)$; if they need a linear-probability reading, report average marginal effects from the logistic model.

⚙️ **The full answer.** The identity link has no range control; the logit link does. Average marginal effect $= \frac{1}{n} \sum w_j p_i (1 - p_i)$ gives “change in probability per unit of x_j ” averaged over the data, which is what the colleague wanted from OLS coefficients.

Q4.16

When would you use isotonic regression?

🕒 **In one line.** When you know the relationship is monotone but not its shape, most commonly to calibrate classifier scores into probabilities.

⚙️ **The full answer.** Isotonic regression finds the best non-decreasing step function (pool adjacent violators algorithm). To calibrate a gradient-boosted model, fit isotonic regression from its raw score to the 0/1 labels on a held-out set. It is flexible but can overfit with little data; Platt scaling (a 1-feature logistic regression) is the low-data alternative.

Q4.17

🧠 **Think deeper** Kernel ridge regression, SVR and Gaussian process regression all use kernels. How do they differ?

🕒 **In one line.** They share the kernel function but differ in loss and output: kernel ridge uses squared loss with a closed-form solution, SVR uses the ε -insensitive loss giving a sparse solution, and a GP is the Bayesian version of kernel ridge that also returns predictive uncertainty.

⚙️ **The full answer.** Kernel ridge: $\hat{f}(\mathbf{x}) = \sum_i \alpha_i K(\mathbf{x}_i, \mathbf{x})$ with $\boldsymbol{\alpha} = (K + \lambda I)^{-1} \mathbf{y}$; every training

point has non-zero α_i . SVR: most $\alpha_i = 0$ (points in the tube), so prediction is cheaper and small errors are ignored. GP: the posterior mean equals kernel ridge with $\lambda = \text{noise variance}$; the posterior variance tells you where the model is unsure (useful for active learning or pricing risk). All scale as $O(n^3)$ naively.

Q4.18**How does XGBoost or LightGBM do Poisson, Gamma or quantile regression?**

🎯 **In one line.** By swapping the loss: the trees fit gradients and Hessians of the chosen negative log-likelihood (Poisson, Gamma, Tweedie) or of the pinball loss, usually on the log scale, so the boosting machinery is unchanged.

⚙️ **The full answer.** For Poisson with log link and raw score F , the loss per row is $e^F - yF$; gradient $e^F - y = \mu - y$; Hessian $e^F = \mu$. These plug straight into the XGBoost split-gain formula (Chapter 8). Parameters: `objective="count:poisson"`, `"reg:gamma"`, `"reg:tweedie"`, `"reg:quantileerror"` with `quantile_alpha`.

CHAPTER 5

Gradient Descent and Optimisers

🎥 **Companion video for this chapter:** [StatQuest: Gradient Descent, Step-by-Step](#)

Almost every model in this book is trained by some form of gradient descent: logistic regression, neural networks, matrix factorisation, word embeddings, and (in function space) gradient boosting. InfoEdge has asked which is faster, SGD or batch gradient descent (OA), what the learning rate is and how to adjust it, and what to do if training gets stuck in a local optimum (interview).

5.1 Walking downhill in fog

You stand on a hillside in thick fog and want to reach the valley floor. You cannot see the valley, but you can feel the slope under your feet. Strategy: step in the steepest downhill direction, a small step, then feel again. That is gradient descent.

- The **landscape** is the loss $L(\mathbf{w})$: height at each setting of the parameters.
- The **slope** is the **gradient** $\nabla L(\mathbf{w})$: a vector of partial derivatives, pointing in the direction of steepest *ascent*. Its length says how steep.
- The **step** goes the other way: $\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla L(\mathbf{w})$.
- η (eta) is the **learning rate**: how big a step relative to the slope.

5.1.1 What a derivative is, in one paragraph

For a function of one number, $L(w)$, the derivative $L'(w)$ is the slope of the curve at w : how much L changes per tiny change in w . If $L'(w) > 0$, increasing w increases L , so we should decrease w ; if $L'(w) < 0$ we should increase it. “Move against the derivative” always goes downhill. With many parameters, the partial derivative $\partial L / \partial w_j$ is the slope along w_j alone, and the gradient collects all of them.

📊 Small example: three steps by hand

$L(w) = (w-3)^2$, so $L'(w) = 2(w-3)$. Start $w_0 = 0, \eta = 0.1$. Step 1: $L'(0) = -6, w_1 = 0 - 0.1 \cdot (-6) = 0.6$. Step 2: $L'(0.6) = -4.8, w_2 = 0.6 + 0.48 = 1.08$. Step 3: $L'(1.08) = -3.84, w_3 = 1.464$. Each step closes 20% of the remaining gap to 3 ($w_{t+1} - 3 = 0.8(w_t - 3)$). Steps shrink automatically as the slope flattens near the minimum.

5.2 The learning rate

- **Too small**: progress is painfully slow; training may stop (time budget) before converging.
- **Too large**: steps overshoot the valley floor and land higher on the other side; the loss oscillates or blows up (NaN).
- **Just right**: fast, steady decrease.

Small example: when it diverges

Same $L = (w-3)^2$. The update is $w_{t+1} - 3 = (1-2\eta)(w_t - 3)$. This shrinks the gap only if $|1-2\eta| < 1$, i.e. $0 < \eta < 1$. At $\eta = 0.5$ it jumps to 3 in one step. At $\eta = 0.9$ it oscillates ($1 - 1.8 = -0.8$) but converges. At $\eta = 1.1$, the factor is -1.2 : the gap grows by 20% each step, with alternating sign. In general, for a quadratic with curvature (second derivative) h , GD converges iff $\eta < 2/h$.

With several parameters, the steepest direction can have curvature 50 and the flattest 1; η must be below $2/50$ for stability, which makes progress along the flat direction 50 times slower. This ratio is the **condition number**, and it is why **feature scaling** matters: unscaled features (salary in rupees next to experience in years) create long, narrow valleys where GD zig-zags.

5.2.1 Schedules

- **Step decay**: divide η by 10 every few epochs.
- **Exponential decay**: $\eta_t = \eta_0 \gamma^t$.
- **Cosine annealing**: $\eta_t = \eta_{min} + \frac{1}{2}(\eta_{max} - \eta_{min})(1 + \cos(\pi t/T))$.
- **Warmup**: start tiny and ramp up over the first few hundred steps (standard for Transformers, where early gradients are unreliable and Adam's statistics are not yet warmed up).
- **Reduce on plateau**: cut η when validation loss stops improving.
- **LR range test**: increase η exponentially over one short run, plot loss, choose a value a bit below where loss is falling fastest.

5.3 Batch, stochastic and mini-batch

The loss is an average over n examples, so its gradient is the average of n per-example gradients.

- **Batch GD**: use all n examples per step. Exact gradient, smooth path, but one step costs a full pass over the data. With 10 million rows, that is very slow per step.
- **Stochastic GD (SGD)**: use one random example per step. Very cheap, very noisy steps.
- **Mini-batch GD**: use a batch of B examples (32–1024). The standard in practice: noise is reduced by averaging, and GPUs process a batch in parallel almost as fast as one example.

An **epoch** is one pass through the data. With $n = 50,000$ and $B = 250$, one epoch is 200 updates.

Intuition: why noisy steps can be good

A random sample of examples gives an unbiased but noisy estimate of the true gradient (on average correct). The noise jostles the parameters, which helps escape saddle points and shallow local minima and tends to settle in wider valleys, which generalise better. The price: SGD with a fixed learning rate never settles exactly; it jitters around the minimum in a cloud whose size is proportional to η . Decaying η shrinks the cloud.

5.4 Convexity, local minima and saddle points

- A function is **convex** if the straight line between any two points on its graph lies above the graph: a bowl. Any local minimum is global. Linear regression (MSE), logistic regression (log loss), SVM (hinge + L2) and Lasso are convex.

- Neural network losses are **non-convex**: many minima and many **saddle points** (flat points that go up in some directions and down in others, like a horse saddle).
- In high dimensions, saddle points are far more common than bad local minima: for a critical point to be a minimum, the curvature must be positive in *every* one of millions of directions. Empirically, most local minima of large networks have similar, good loss values. The real enemies are plateaus and saddles, where the gradient is tiny.

5.5 Momentum, RMSProp and Adam, built up step by step

5.5.1 Momentum

Plain GD in a long narrow valley zig-zags across the steep walls while creeping along the floor. **Momentum** keeps a running “velocity”:

$$\mathbf{v} \leftarrow \beta \mathbf{v} + \nabla L, \quad \mathbf{w} \leftarrow \mathbf{w} - \eta \mathbf{v}, \quad \beta \approx 0.9.$$

Across the valley the gradients alternate sign and cancel in $\mathbf{v}$; along the floor they agree and accumulate. Like a heavy ball rolling downhill. With a constant gradient g , $\mathbf{v}$ approaches $g/(1 - \beta) = 10g$: the effective step is up to 10 times larger along consistent directions.

Small example: momentum velocities

$\beta = 0.9$, constant gradient 1, $v_0 = 0$: $v_1 = 1, v_2 = 1.9, v_3 = 2.71, v_4 = 3.44, \dots \rightarrow 10$.

Nesterov momentum evaluates the gradient at the look-ahead point $\mathbf{w} - \eta\beta\mathbf{v}$: it corrects course before overshooting.

5.5.2 AdaGrad and RMSProp: a learning rate per parameter

Some parameters get large gradients, some small (rare features, rare words). RMSProp keeps a running average of *squared* gradients per parameter and divides by its square root:

$$s \leftarrow \rho s + (1 - \rho)g^2, \quad w \leftarrow w - \frac{\eta}{\sqrt{s} + \epsilon} g.$$

Parameters with consistently large gradients take smaller steps; rarely updated ones take larger steps. AdaGrad uses the *sum* of all past squared gradients instead of a moving average, so its effective learning rate only ever shrinks (good for sparse convex problems, stalls in deep learning).

5.5.3 Adam = momentum + RMSProp + bias correction

$$\begin{aligned} m &\leftarrow \beta_1 m + (1 - \beta_1)g && \text{(average gradient: momentum)} \\ v &\leftarrow \beta_2 v + (1 - \beta_2)g^2 && \text{(average squared gradient: scale)} \\ \hat{m} &= m/(1 - \beta_1^t), \quad \hat{v} = v/(1 - \beta_2^t) && \text{(bias correction)} \\ w &\leftarrow w - \eta \hat{m}/(\sqrt{\hat{v}} + \epsilon) \end{aligned}$$

Defaults: $\beta_1 = 0.9, \beta_2 = 0.999, \epsilon = 10^{-8}, \eta = 10^{-3}$. Because m and v start at 0, early averages are biased toward 0; dividing by $(1 - \beta^t)$ removes that bias.

Small example: the first Adam step

$g = 4$, $w = 1$, $\eta = 0.01$, $t = 1$. $m = 0.4$, $v = 0.016$. $\hat{m} = 0.4/0.1 = 4$, $\hat{v} = 0.016/0.001 = 16$. Step $= 0.01 \cdot 4/4 = 0.01$, so $w = 0.99$. The first step has size η whatever the gradient's scale: Adam is scale-invariant per parameter.

AdamW applies weight decay directly to the weights instead of through the gradient (Chapter 1); it is the default for Transformers.

5.6 Second-order methods

Newton's method uses curvature: $\mathbf{w} \leftarrow \mathbf{w} - H^{-1}\nabla L$, where H is the Hessian (matrix of second derivatives). It jumps straight to the minimum of a quadratic and converges in a handful of steps near the optimum. But H has d^2 entries and inverting costs d^3 : impossible for a million parameters. **L-BFGS** approximates H^{-1} from the last few gradients and is the default for logistic regression in scikit-learn. XGBoost uses a diagonal second-order step per leaf (Chapter 8).

5.7 Practical checklist when training is not going well

Symptom	Likely cause and fix
Loss explodes / NaN	learning rate too high, exploding gradients, $\log(0)$ or division by 0: lower η , gradient clipping, check data for inf/NaN, use numerically stable losses
Loss oscillates wildly	η too high or batch too small: lower η , larger batch, momentum
Loss decreases very slowly	η too low, unscaled features, saturated activations: raise η , scale features, better init, BatchNorm, Adam
Loss plateaus early at a poor value	saddle/plateau, dead ReLUs, too much regularisation, bug in labels: restart with different seed, warm restarts, cyclic LR, check data pipeline, overfit a tiny batch first
Train loss falls, val loss rises	overfitting: early stopping and regularisation

In practice: optimisation at InfoEdge scale

Click/apply models on billions of events are trained with mini-batch SGD/Adam on sparse features, often on Spark or GPU. The Transformers behind search and the fine-tuned LLMs in the GenAI slide (SFT, DPO, distillation) are trained with AdamW, warmup and cosine decay. A practical rule from these systems: when you double the batch size, scale the learning rate up (linearly for SGD, by about $\sqrt{2}$ for Adam) and add warmup.

5.8 Interview questions

Q5.1

★ **PYQ (OA: SGD vs batch gradient descent)** On a dataset with 10 million rows, which is faster, stochastic or batch gradient descent? Explain properly.

🎯 **In one line.** Per update, SGD (or mini-batch) is millions of times cheaper and in wall-clock time it reaches a good solution much faster; batch GD takes more exact steps but each step costs a full pass over the data.

⚙️ **The full answer.** Batch GD computes the exact gradient by touching all 10^7 rows before moving once. SGD moves after each row (or each mini-batch of, say, 256), so in the time batch GD makes 1 step, mini-batch SGD makes about 39,000 steps. Early in training, even a noisy direction is good enough to make huge progress. Near the optimum, the noise means SGD jitters rather than converging exactly; a decaying learning rate fixes that. Technically, to reach error ϵ on a strongly convex problem, batch GD needs $O(\log(1/\epsilon))$ iterations of cost $O(n)$ each, and SGD needs $O(1/\epsilon)$ iterations of cost $O(1)$ each: for large n and modest accuracy (which is all generalisation needs), SGD wins.

📊 **Worked example.** Logistic regression on 10M apply events: one epoch of mini-batch SGD with $B = 512$ gives about 19,500 updates and typically gets within 0.1% of the final log loss; batch GD would need hundreds of full passes.

💡 **If the interviewer pushes.** “Then why ever use batch GD?” Small datasets, when you need exact deterministic convergence, or with second-order methods (L-BFGS) that need accurate gradients. “Why mini-batch over pure SGD?” GPU parallelism and lower variance per step.

⚠️ **Trap to avoid.** Answering “batch GD, because it converges in fewer iterations”. Fewer iterations, but each is n times more expensive.

Q5.2

★ **PYQ (INT: learning rate)** What is the learning rate, how do you pick it, and how do you adjust it during training?

🎯 **In one line.** It is the step size multiplier on the gradient; pick it with an LR range test or a log-scale search, and adjust it with a schedule (warmup, then step/cosine decay or reduce-on-plateau).

⚙️ **The full answer.** Too small: slow; too large: oscillation or divergence. For a quadratic with curvature h , stability requires $\eta < 2/h$. **Picking:** try $10^{-1}, 10^{-2}, 10^{-3}, 10^{-4}$ for a few hundred steps and compare loss curves, or run a range test. Defaults: Adam 10^{-3} (from scratch), 2×10^{-5} to 5×10^{-5} for fine-tuning BERT; SGD with momentum 0.1 for ResNets with batch 256. **Adjusting:** large η early to make progress and escape poor regions, smaller later to settle. Warmup protects the start; decay sharpens the end. If validation loss plateaus, cut by 10.

💡 **If the interviewer pushes.** “What happens if you raise the batch size 8 times?” The gradient noise drops; you can raise η (linear scaling rule) with a warmup.

Q5.3

★ **PYQ (INT: stuck at a local optimum)** Your network's loss stops improving at a poor value. What could be happening and what do you do?

🕒 **In one line.** In deep nets it is more often a plateau, a saddle point, saturated or dead units, a learning rate problem or a data bug than a true bad local minimum; I would debug the pipeline first, then change the optimisation (learning rate schedule, momentum/Adam, restarts, initialisation, normalisation).

⚙️ **The full answer.**

1. **Bug check:** can the model overfit 32 examples to near-zero loss? If not, there is a bug (labels shuffled relative to inputs, wrong loss, frozen layers, wrong preprocessing).
2. **Learning rate:** too low means crawling on a plateau; too high means bouncing in a valley. Try an LR finder and a schedule; try cyclic or warm restarts to jump out.
3. **Saturation:** sigmoid/tanh units with huge inputs have near-zero gradients; ReLUs that always output 0 are dead. Fix with better initialisation (He/Xavier), BatchNorm, Leaky ReLU, a lower η .
4. **Optimiser:** momentum or Adam carry through flat regions; SGD noise (smaller batches) helps escape saddles.
5. **Capacity:** if training loss itself is high, the model may simply be too small.
6. **Restarts:** different random seeds; ensembles of them.

💡 **If the interviewer pushes.** "Are local minima a big problem in deep learning?" Mostly no, in large networks: most local minima are nearly as good as the global one; saddle points and plateaus are the practical issue.

Q5.4

Do three gradient-descent steps on $L(w) = (w - 3)^2$ from $w_0 = 0$ with $\eta = 0.1$. For which η does it converge?

🕒 **In one line.** $w = 0.6, 1.08, 1.464$; it converges for $0 < \eta < 1$.

⚙️ **The full answer.** Derivative $2(w - 3)$; the gap to the optimum multiplies by $(1 - 2\eta)$ each step; convergence needs $|1 - 2\eta| < 1$.

Q5.5

Why does feature scaling speed up gradient descent?

🕒 **In one line.** Unscaled features create a loss surface with very different curvature in different directions; the learning rate must suit the steepest direction, so progress along flat directions is slow and the path zig-zags; scaling makes the contours roughly circular.

⚙️ **The full answer.** For linear regression the curvature along w_j is proportional to the variance of feature j . Salary in rupees (variance 10^{11}) next to experience in years (variance 10) gives a condition number of 10^{10} . With standardisation both have variance 1 and GD heads straight to

the minimum.

 **If the interviewer pushes.** “Which models do not need scaling?” Trees and tree ensembles (they split on thresholds, invariant to monotone transforms). Naive Bayes mostly. Everything distance- or gradient-based needs it.

Q5.6

Explain momentum intuitively and mathematically.

 **In one line.** Momentum keeps an exponentially weighted average of past gradients and steps along it, so consistent directions accelerate and oscillating directions cancel, like a heavy ball rolling down a valley.

 **The full answer.** $\mathbf{v}_t = \beta \mathbf{v}_{t-1} + \nabla L$; the effective step for a constant gradient approaches $\eta g / (1 - \beta)$. In a ravine the across-ravine gradient flips sign every step and averages out, while the along-ravine component adds up.

 **Worked example.** See the velocity sequence $1, 1.9, 2.71, 3.44 \rightarrow 10$ for $\beta = 0.9$.

Q5.7

Explain Adam. Why does it need bias correction?

 **In one line.** Adam keeps moving averages of the gradient (momentum) and of its square (per-parameter scale) and steps by their ratio; because both averages start at zero they are biased toward zero early on, and dividing by $1 - \beta^t$ corrects this.

 **The full answer.** At $t = 1$, $m_1 = (1 - \beta_1)g = 0.1g$, ten times too small; $\hat{m}_1 = m_1 / (1 - 0.9) = g$. Without correction the first steps would be tiny for m and, more dangerously, v would be tiny too ($0.001g^2$), so $m / \sqrt{v}$ could be large and erratic.

 **If the interviewer pushes.** “Adam or SGD?” Adam converges faster with little tuning and is standard for Transformers and NLP; SGD with momentum often generalises slightly better for CNNs on vision. AdamW fixes Adam’s interaction with L2.

Q5.8

Compute one Adam step: $g = 4, w = 1, \eta = 0.01, \beta_1 = 0.9, \beta_2 = 0.999, t = 1$.

 **In one line.** $\hat{m} = 4, \hat{v} = 16$, step = 0.01, new $w = 0.99$.

 **The full answer.** $m = 0.1 \cdot 4 = 0.4$; $v = 0.001 \cdot 16 = 0.016$; $\hat{m} = 0.4 / 0.1 = 4$; $\hat{v} = 0.016 / 0.001 = 16$; $\sqrt{\hat{v}} = 4$; step = $0.01 \cdot 4 / 4$.

Q5.9

 **Think deeper** Why does SGD with a constant learning rate not converge exactly, and why might that be a feature rather than a bug?

🎯 **In one line.** Its gradient is a noisy estimate, so even at the minimum it keeps taking random steps of size proportional to η ; that noise also helps escape sharp minima and saddles and biases training toward flat minima that generalise better.

⚙️ **The full answer.** At the optimum the average gradient is zero but individual mini-batch gradients are not. The stationary cloud has variance roughly proportional to $\eta \cdot (\text{gradient variance})/B$. Sharp minima (high curvature) cannot hold a noisy SGD iterate, so it drifts to wide basins; wide basins are less sensitive to the shift between training and test data. To get exact convergence, decay η (Robbins–Monro: $\sum \eta_t = \infty, \sum \eta_t^2 < \infty$) or average the iterates.

Q5.10

🧠 **Think deeper** Large-batch training often generalises slightly worse. Why?

🎯 **In one line.** Larger batches mean less gradient noise and fewer updates per epoch, so training tends to settle into sharper minima and explores less; this can be largely recovered by scaling the learning rate, warmup and longer training.

⚙️ **The full answer.** Gradient noise scale is about η/B . Keeping η/B constant (linear scaling rule) restores the noise level. Methods like LARS/LAMB adapt per-layer learning rates for very large batches.

Q5.11

What is gradient clipping? Compute the clipped gradient for $g = (3, 4)$ with norm threshold 1.

🎯 **In one line.** If the gradient's length exceeds a threshold, rescale it to that length, keeping its direction; here $\|g\| = 5$ so the clipped gradient is $(0.6, 0.8)$.

⚙️ **The full answer.** Clipping prevents one exploding step from destroying training, standard for RNNs and Transformers. Clipping by norm keeps the direction; clipping by value (each component to $[-c, c]$) can change it.

Q5.12

🧠 **Think deeper** Is gradient descent guaranteed to find the global minimum of logistic regression? Of a neural network?

🎯 **In one line.** For logistic regression yes (with a suitable learning rate), because the loss is convex; for a neural network no guarantee, but in practice large networks reach minima that are nearly as good as any.

⚙️ **The full answer.** Convexity means every local minimum is global and GD with $\eta < 2/L$ (L = largest curvature) converges. Neural network losses are non-convex because of symmetries (swapping two hidden units gives the same function) and non-linearities. Over-parameterisation makes the landscape benign in practice.

💡 **If the interviewer pushes.** “What if the logistic data are separable?” There is no finite minimum; GD keeps growing $\|w\|$ (direction converges to the max-margin SVM solution, interestingly).

Q5.13

Why is Newton's method rarely used for deep learning, and where is second-order information used?

🕒 **In one line.** The Hessian of a network with d parameters has d^2 entries and inverting it costs d^3 , which is impossible for millions of parameters, and it can be indefinite (saddles); quasi-Newton L-BFGS is used for small convex problems and XGBoost uses per-leaf second-order steps.

⚙️ **The full answer.** With $d = 10^8$, the Hessian would need 10^{16} numbers. Newton also moves toward saddle points as readily as minima. Approximations (K-FAC, Shampoo) exist in research and large-scale training.

Q5.14

How many parameter updates happen in 20 epochs with 50,000 examples and batch size 250?

🕒 **In one line.** $50,000/250 = 200$ updates per epoch, so 4,000 updates.

Q5.15

💡 **Think deeper** Loss becomes NaN after a few epochs. Debug it out loud.

🕒 **In one line.** I would check, in order: input data for NaN/inf, numerically unstable operations (log of zero, division by zero, exp overflow), learning rate and exploding gradients, and mixed-precision overflow.

⚙️ **The full answer.**

1. Assert no NaN/inf in inputs and labels (`np.isfinite`). A single bad row can poison a batch.
2. Use stable loss functions: `BCEWithLogitsLoss` instead of sigmoid then log; log-softmax instead of softmax then log. Add ϵ in denominators.
3. Log the gradient norm per step: a spike before NaN means exploding gradients: lower η , clip gradients, add normalisation.
4. In fp16, loss scaling may overflow; use bf16 or dynamic loss scaling.
5. Reproduce on a small batch with anomaly detection (`torch.autograd.set_detect_anomaly(True)`).

Q5.16

What are the differences between AdaGrad, RMSProp and Adam?

🕒 **In one line.** All scale steps per parameter by past squared gradients; AdaGrad sums them forever (learning rate only decays), RMSProp uses a moving average (adapts to recent scale), Adam adds momentum and bias correction to RMSProp.

⚙️ **The full answer.** AdaGrad suits sparse convex problems (rare features get large steps). In deep nets its accumulated sum grows without bound, freezing learning. RMSProp's forgetting

fixes that. Adam's momentum smooths the direction.

Q5.17

💡 Think deeper Gradient boosting is called “gradient descent in function space”. What does that mean?

🎯 **In one line.** Instead of updating a parameter vector by minus the gradient, each boosting round adds a small tree fitted to the negative gradient of the loss with respect to the current predictions, so the predictions themselves take a gradient step.

⚙️ **The full answer.** Treat the vector of predictions $(F(x_1), \dots, F(x_n))$ as the thing to optimise. The gradient of $\sum \ell(y_i, F_i)$ with respect to F_i is $\partial \ell / \partial F_i$; for MSE, minus the gradient is the residual $y_i - F_i$. A tree fitted to these values generalises the step to new x . The learning rate (shrinkage) is η . Chapter 8 develops this.

Q5.18

What is early stopping, and why is it a form of regularisation?

🎯 **In one line.** Stop training when validation loss stops improving and keep the best checkpoint; because weights start small and grow during training, stopping early limits their size, similar to an L2 penalty.

⚙️ **The full answer.** For a quadratic loss with GD from zero, the number of steps t and learning rate η act like $1/\lambda$ in Ridge: directions with large curvature are fitted early, small-curvature (noisy) directions late. Use patience (for example 5 epochs without improvement) to avoid stopping on noise.

Q5.19

💡 Think deeper Why does warmup help when training Transformers with Adam?

🎯 **In one line.** At the start, Adam's second-moment estimates are based on very few gradients and are unreliable, and the network is in an unstable region (especially with post-LayerNorm); a small learning rate at first prevents a few huge, badly scaled steps from wrecking the model.

⚙️ **The full answer.** Early $\hat{v}$ can be tiny for some parameters, making $\hat{m}/\sqrt{\hat{v}}$ large; the attention softmax can saturate. Linear warmup over the first 1–10% of steps lets statistics accumulate. RAdam is a variant that corrects the variance of the adaptive term automatically.

Q5.20

💡 Think deeper A colleague says “we use the closed-form solution, so we do not need to worry about optimisation”. When can the closed form be worse than gradient descent?

🎯 **In one line.** When the matrix is ill-conditioned or singular (numerical error), when the

data do not fit in memory, when d is huge, and when early-stopped GD gives useful implicit regularisation.

⚙️ **The full answer.** $(X^T X)^{-1}$ amplifies rounding error by the condition number; QR or SVD solvers are safer. With $d = 10^6$ sparse text features, forming $X^T X$ is impossible. For underdetermined problems, GD from zero gives the minimum-norm solution, often desirable.

Part II

Classical Models

CHAPTER 6


Decision Trees

🎥 **Companion video for this chapter:** [StatQuest: Decision and Classification Trees, Clearly Explained](#)

Decision trees are the building block of random forests, gradient boosting, XGBoost and LightGBM, which together win most tabular problems, including most at InfoEdge. Glassdoor reports for InfoEdge include “pseudo-code a decision tree”; the experience notes include reducing overfitting in trees. Be able to compute impurity by hand and write the algorithm from memory.

6.1 A tree is a game of twenty questions

To guess whether a seeker will apply to a job, ask a sequence of yes/no questions:



- Each internal **node** tests one feature against a threshold (“experience ≤ 3.5 ”) or a category set.
- Each **branch** is an answer; each **leaf** holds a prediction: the majority class (or class proportions) for classification, the mean target for regression.
- A prediction walks from the root to one leaf. Depth D means at most D questions.

In feature space, each split cuts along one axis, so a tree partitions the space into axis-aligned rectangles (boxes), each with a constant prediction.

6.2 How to choose a question: impurity

At each node we want the question that makes the two children as *pure* as possible (each child mostly one class). We need a number for impurity.

6.2.1 Gini impurity

$$\text{Gini} = 1 - \sum_k p_k^2,$$

where p_k is the fraction of class k in the node. Interpretation: the probability that two items picked at random (with replacement) from the node have different classes. Pure node: 0. Two classes 50/50: $1 - 0.25 - 0.25 = 0.5$ (the maximum for two classes).

6.2.2 Entropy

$$H = - \sum_k p_k \log_2 p_k.$$

Interpretation: the average number of yes/no questions (bits) needed to identify the class of a random item. Pure: 0; 50/50: 1 bit; four equally likely classes: 2 bits.

Small example: impurities of one node

8 positives, 2 negatives: $p = (0.8, 0.2)$. Gini = $1 - 0.64 - 0.04 = 0.32$. Entropy = $-0.8 \log_2 0.8 - 0.2 \log_2 0.2 = 0.8(0.322) + 0.2(2.322) = 0.258 + 0.464 = 0.722$ bits.

6.2.3 Information gain

The quality of a split is how much it reduces impurity, weighting each child by its share of the rows:

$$\text{Gain} = I(\text{parent}) - \sum_c \frac{n_c}{n} I(c).$$

Small example: compare two splits

Parent: 6 apply, 6 do not. Entropy = 1. **Split A** ("preferred city?"): left (5 yes, 1 no), right (1 yes, 5 no). Each child: $p = (5/6, 1/6)$, entropy = 0.650. Weighted: 0.650. Gain = $1 - 0.650 = 0.350$. **Split B** ("remote?"): left 8 rows (6 yes, 2 no), right 4 rows (0 yes, 4 no). Left entropy = 0.811, right = 0. Weighted = $\frac{8}{12} 0.811 = 0.541$. Gain = 0.459. **Split B wins**, even though its left child is not pure, because the right child is perfectly pure and holds a third of the data. With Gini: parent 0.5; A: each child $1 - (25 + 1)/36 = 0.278$, gain 0.222; B: left $1 - (0.5625 + 0.0625) = 0.375$, weighted 0.25, gain 0.25. Same winner.

6.2.4 Regression trees: variance reduction

For a numeric target, impurity is the variance (or the sum of squared errors around the node mean). The leaf predicts the mean.

Small example: one regression split

Targets in a node: 2, 4, 6, 20. Mean 8; SSE = $36 + 16 + 4 + 144 = 200$. Split "{2, 4, 6} vs {20}": left mean 4, SSE = $4 + 0 + 4 = 8$; right SSE = 0. Reduction = 192. The outlier gets its own leaf.

6.2.5 How thresholds are searched

For a numeric feature, sort the node's values; candidate thresholds are midpoints between consecutive distinct values ($m - 1$ candidates for m distinct values). Evaluate the gain for every feature and every threshold; pick the best. With running sums, all thresholds of one sorted feature can be scored in one pass. Histogram methods (LightGBM, XGBoost hist) bucket values into about 256 bins first, making this much faster.

6.3 The algorithm (pseudo-code)

```
def build_tree(rows, depth):
    # stopping rules
    if depth == MAX_DEPTH or len(rows) < MIN_SAMPLES_SPLIT or is_pure(rows):
        return Leaf(prediction(rows)) # majority class / mean
    best = None
    for feature in features:
        for threshold in candidate_thresholds(rows, feature):
            left, right = split(rows, feature, threshold)
            if len(left) < MIN_SAMPLES_LEAF or len(right) < MIN_SAMPLES_LEAF:
                continue
            gain = impurity(rows) - (len(left)*impurity(left) + len(right)*impurity(right)
            ) / len(rows)
            if best is None or gain > best.gain:
                best = Split(feature, threshold, gain, left, right)
    if best is None or best.gain < MIN_GAIN:
        return Leaf(prediction(rows))
    return Node(best.feature, best.threshold,
                build_tree(best.left, depth + 1),
                build_tree(best.right, depth + 1))

def predict(node, x):
    while not node.is_leaf:
        node = node.left if x[node.feature] <= node.threshold else node.right
    return node.prediction
```

Complexity. Training: with pre-sorting, about $O(d n \log n)$ per level, so $O(d n \log n \cdot \text{depth})$ overall. Prediction: $O(\text{depth})$, independent of n : very fast.

6.3.1 Greedy, not optimal

The algorithm picks the best split *now*, never reconsidering. Finding the globally optimal tree is NP-hard. Greedy choice can miss splits that look useless alone but are powerful together (XOR: neither x_1 nor x_2 alone reduces impurity, but both together separate perfectly).

6.4 The classic algorithms

	ID3	C4.5	CART
Criterion	information gain (entropy)	gain ratio	Gini (classification), MSE (regression)
Splits	multiway on categoricals	multiway; handles numeric thresholds	binary only
Tasks	classification	classification	classification and regression
Missing values	no	fractional instances	surrogate splits
Pruning	no	error-based	cost-complexity

scikit-learn implements an optimised CART.

6.4.1 The bias toward many-valued features, and gain ratio

A feature like `customer_id` with a unique value per row splits the data into pure singletons: maximum information gain, zero generalisation. **Gain ratio** divides gain by the split information $-\sum \frac{n_c}{n} \log_2 \frac{n_c}{n}$, which is large for many small children.

6.5 Overfitting and how to stop it

A tree grown until every leaf is pure memorises the training data: training accuracy 100%, high variance (change a few rows and the whole structure changes, because an early split change cascades).

Pre-pruning (early stopping): `max_depth`, `min_samples_split`, `min_samples_leaf`, `max_leaf_nodes`, `min_impurity_decrease`, `max_features`.

Post-pruning: grow a full tree, then cut back. **Cost-complexity pruning (CART)** minimises

$$R_\alpha(T) = R(T) + \alpha|T|,$$

training error plus α times the number of leaves. Each α gives a nested subtree; choose α by cross-validation (scikit-learn: `ccp_alpha`, and `cost_complexity_pruning_path`).

Small example: is pruning worth it?

$\alpha = 0.02$. Subtree T_1 : error 0.10, 5 leaves: $R = 0.10 + 0.10 = 0.20$. Pruned T_2 : error 0.13, 2 leaves: $R = 0.13 + 0.04 = 0.17$. Prune.

Post-pruning sees the whole tree, so it avoids the “horizon effect” of pre-pruning (stopping before a weak split that enables strong later ones).

6.6 Strengths and weaknesses

Strengths: interpretable (a small tree is a flowchart a recruiter can read); handles numeric and categorical features; no scaling needed (only the order of values matters, so any monotone transform leaves the tree unchanged); captures interactions and non-linearity automatically; fast predictions; some implementations handle missing values.

Weaknesses: high variance and instability; axis-aligned boundaries approximate diagonal boundaries with staircases; piecewise-constant predictions, so **no extrapolation** beyond the training range; greedy; biased toward high-cardinality features; poorly calibrated probabilities from pure leaves.

6.7 Feature importance

Impurity-based (MDI): sum, over all nodes splitting on feature j , of the weighted impurity decrease; normalise to sum to 1. Fast, but biased toward high-cardinality and continuous features and computed on training data. **Permutation importance:** shuffle feature j in validation data and measure how much the metric drops. Model agnostic, uses held-out data, but splits credit between correlated features. **SHAP** values give per-prediction attributions with consistency guarantees (TreeSHAP is exact and fast for trees).

📦 In practice: trees at InfoEdge

Single trees are rarely the production model, but they are (1) the unit inside every GBM ranking and lead-scoring model; (2) a great **explanation tool**: fit a depth-3 tree to a black-box model's predictions to describe its logic to the 99acres sales team ("high intent = viewed ≥ 5 listings in one locality in 7 days and price band within 20%"); (3) quick **segment discovery** for the case-study round.

6.8 Interview questions

Q6.1

★ **PYQ (Glassdoor, InfoEdge: pseudo-code a decision tree)** Write the training algorithm for a classification tree and state its complexity.

🎯 **In one line.** Recursively, for each node, try every feature and every candidate threshold, pick the split with the largest impurity decrease, and recurse on the two children until a stopping rule fires; training costs about $O(d n \log n)$ per depth level and prediction costs $O(\text{depth})$.

⚙️ **The full answer.** I would write the pseudo-code above and narrate four design decisions:

1. **Impurity:** Gini (default, slightly faster, no logarithm) or entropy.
2. **Candidate thresholds:** midpoints between sorted distinct values; for categorical features either one-hot or, for binary classification, sort categories by positive rate and treat as ordinal (provably optimal for Gini).
3. **Stopping:** max depth, min samples per split/leaf, pure node, min gain.
4. **Leaf value:** majority class and class proportions (for `predict_proba`).

Complexity: sorting each feature costs $O(n \log n)$; scanning thresholds with running class counts is $O(n)$ per feature; at each depth level the nodes partition the n rows, so total per level is $O(d n \log n)$ (or $O(d n)$ with pre-sorting or histograms).

💡 **If the interviewer pushes.** "How would you make it a regression tree?" Replace Gini with variance (SSE) and the leaf value with the mean. "How do you handle a missing value at prediction time?" Send it down a learned default direction (XGBoost), use surrogate splits (CART), or impute.

Q6.2

★ **PYQ (INT: reducing overfitting in trees)** How do you stop a decision tree from overfitting?

🎯 **In one line.** Constrain growth (max depth, min samples per leaf, max leaf nodes, min impurity decrease), post-prune with cost-complexity pruning tuned by cross-validation, or replace the single tree by an ensemble (random forest or boosting).

⚙️ **The full answer.** A tree can always reach zero training error by giving every point its own leaf, which memorises noise. Each control limits how finely the space can be cut: `min_samples_leaf=50` means no prediction is based on fewer than 50 rows, so leaf estimates are stable. Cost-complexity pruning penalises leaves with $\alpha|T|$. Ensembles attack variance directly by averaging (bagging) or

by building in many small, shrunk steps (boosting).

 **Worked example.** Shiksha lead data: unlimited depth 100% train / 71% val; max_depth=8, min_samples_leaf=40: 83% / 80%; random forest: 88% / 85%.

 **If the interviewer pushes.** “Which is better, pre- or post-pruning?” Post-pruning is usually better because it sees later useful splits; pre-pruning is cheaper. In practice, tune min_samples_leaf and max_depth by CV.

Q6.3

★ **PYQ (OA: entropy and information gain)** A node has 8 positives and 2 negatives. Compute its entropy and Gini. Then compute the information gain of splitting a 6/6 parent into (5,1) and (1,5).

🎯 **In one line.** Entropy 0.722 bits, Gini 0.32; information gain $1 - 0.650 = 0.350$ bits.

⚙️ **The full answer.** As in the worked examples: $H(5/6, 1/6) = -\frac{5}{6} \log_2 \frac{5}{6} - \frac{1}{6} \log_2 \frac{1}{6} = 0.219 + 0.431 = 0.650$; both children have equal size and entropy, so the weighted child entropy is 0.650.

Q6.4

Gini or entropy: does it matter?

🎯 **In one line.** Rarely: they choose the same split in the vast majority of cases; Gini is slightly cheaper (no logarithm), entropy is slightly more sensitive to changes in small class probabilities.

⚙️ **The full answer.** Both are concave, zero for pure nodes and maximal for uniform ones; Gini is roughly a scaled second-order approximation of entropy. Empirical studies find them disagreeing on about 2% of splits. Mention **misclassification error** $1 - \max_k p_k$ as a third option that is a poor splitting criterion: it is not strictly concave, so it often sees zero gain for useful splits.

Q6.5

 **Think deeper** A regression tree is trained on flats priced 20 to 90 lakh. What does it predict for a mansion that is far larger than anything in training? Can a random forest or gradient boosting do better?

🎯 **In one line.** It predicts the mean of whichever leaf the mansion falls into, at most about 90 lakh: trees cannot extrapolate beyond the range of training targets, and forests and GBMs (sums or averages of such leaves) cannot either.

⚙️ **The full answer.** Each leaf predicts a training-data average, and every split beyond the largest training value sends the mansion to the rightmost leaf. Random forest predictions are averages of leaf means (bounded by the target range). Gradient boosting sums many trees' leaf values; in practice it also saturates. For trend extrapolation (prices rising over time, sales growing), detrend first and model residuals with the tree, use linear models or linear trees (LightGBM linear_tree), or add a linear component.

 **Worked example.** Training max price 90 lakh. A 5,000 sq ft villa: tree says 88 lakh (the mean of the largest-flat leaf). A linear model on log price extrapolates to about 3 crore.

 **If the interviewer pushes.** “Where does this bite in practice?” Forecasting with time trends: a GBM trained on 2023–2025 application counts cannot predict 2026 values above the historic maximum. Use differencing or a ratio target.

Q6.6

 **Think deeper** Why do decision trees not need feature scaling? Which preprocessing *can* change a tree?

 **In one line.** Splits depend only on the order of values, so any strictly increasing transform (scaling, log) leaves the tree unchanged; changes that alter order or grouping (one-hot vs ordinal encoding of categories, binning, combining features) do change it.

 **The full answer.** A threshold $x \leq t$ on x is identical to $\log x \leq \log t$. But ordinal-encoding “city” as 1, 2, 3, ... imposes an arbitrary order, so the tree can only split cities into contiguous blocks of that order. One-hot gives one city versus the rest per split. Creating a ratio feature (price per sq ft) can let one split do what needed many.

Q6.7

Why do trees prefer features with many unique values, and how is it fixed?

 **In one line.** More unique values offer more candidate splits, so by chance alone one of them looks good, and an ID-like feature can produce pure singleton leaves; gain ratio, minimum leaf sizes, permutation importance and dropping IDs fix it.

 **The full answer.** With m distinct values there are $m - 1$ thresholds to try; the maximum of many noisy gains is optimistic (multiple comparisons). ID3 with `customer_id` gives maximum information gain and zero generalisation. Fixes: C4.5 gain ratio, conditional inference trees (statistical tests), `min_samples_leaf`, and never feeding IDs.

Q6.8

 **Think deeper** Can a decision tree learn XOR? What does this tell you about greedy splitting?

 **In one line.** A depth-2 tree can represent XOR, but a greedy learner may not find it, because no single first split reduces impurity on its own; with real data’s noise it usually finds it, but in pure form the first split’s gain is zero.

 **The full answer.** Four points $(0,0) \rightarrow 0, (0,1) \rightarrow 1, (1,0) \rightarrow 1, (1,1) \rightarrow 0$: splitting on x_1 gives children $(0,1)$ and $(1,0)$, each 50/50: zero gain. Same for x_2 . A greedy algorithm with a “min gain > 0” rule stops at the root. Without that rule it splits anyway (ties broken arbitrarily) and the second level separates perfectly. Lesson: greedy, one-step look-ahead misses interactions whose components are individually useless; ensembles with random feature choices and deeper trees mitigate it.

Q6.9

How does a decision tree handle categorical features?

🎯 **In one line.** CART-style implementations need numeric input, so categories are one-hot or ordinal encoded; some libraries split categories natively by grouping them into two subsets, using the trick of sorting categories by target mean.

⚙️ **The full answer.** For binary classification or regression, sorting the k categories by mean target and treating them as ordered finds the best two-way partition among $2^{k-1} - 1$ possibilities in $O(k \log k)$. LightGBM and CatBoost do native categorical splits; recent scikit-learn also supports them in HistGradientBoosting. One-hot with many levels makes deep, unbalanced trees.

Q6.10

How does a tree compute `predict_proba`, and why are those probabilities poorly calibrated?

🎯 **In one line.** It returns the class proportions of the training rows in the leaf; deep trees have small, pure leaves that output 0 or 1, so probabilities are overconfident.

⚙️ **The full answer.** A leaf with 3 positives out of 3 says $p = 1.0$, but the true rate for such inputs might be 0.7. Remedies: larger `min_samples_leaf`, Laplace smoothing $(k + 1)/(n + 2)$, ensembling (averages are smoother), and post-hoc calibration (isotonic/Platt).

Q6.11

What is cost-complexity pruning? Work an example.

🎯 **In one line.** It chooses the subtree that minimises training error plus α times the number of leaves, with α tuned by cross-validation.

⚙️ **The full answer.** Increasing α from 0 gives a sequence of nested subtrees (weakest-link pruning): at each stage prune the internal node whose removal increases error least per leaf removed. Worked example in the theory section: 0.20 vs 0.17 means prune.

Q6.12

Explain impurity-based vs permutation feature importance. Which would you trust?

🎯 **In one line.** Impurity importance sums each feature's split gains on training data and is biased toward high-cardinality features; permutation importance measures the validation metric drop when a feature is shuffled and is usually more trustworthy, though it splits credit among correlated features.

⚙️ **The full answer.** A random ID column gets high MDI importance (many splits that fit noise) but near-zero permutation importance (shuffling it does not hurt validation). Correlated features: shuffling one leaves its twin intact, so both look unimportant. Fix by grouping correlated features or using SHAP with care.

Q6.13

 **Think deeper** Why are single decision trees called unstable, and why is that a good thing for ensembles?

🎯 **In one line.** A small change in the data can change an early split and therefore the whole tree below it; that high variance means different trees make different errors, which averaging cancels.

⚙️ **The full answer.** Bagging's variance reduction is $\rho\sigma^2 + \frac{1-\rho}{B}\sigma^2$: it needs models that individually have high variance σ^2 and are not too correlated (ρ small). Deep trees have low bias and high variance: ideal for bagging. Stable models like linear regression gain little from bagging.

Q6.14

What does a decision boundary of a tree look like in 2D? What kind of boundary is hard for it?

🎯 **In one line.** A set of axis-aligned rectangles; diagonal or curved boundaries need many small steps (a staircase), requiring a deep tree.

⚙️ **The full answer.** A boundary like $x_1 = x_2$ is approximated by a staircase; with limited depth the approximation is coarse. Fixes: add a feature $x_1 - x_2$ (one split then suffices), use oblique trees, or use ensembles (many staircases averaged look smoother).

Q6.15

How many leaves can a binary tree of depth 10 have? How many candidate thresholds for a feature with 100 distinct values?

🎯 **In one line.** At most $2^{10} = 1024$ leaves; 99 candidate thresholds.

Q6.16

 **Think deeper** You fit a depth-3 tree to explain a complex model's predictions to a business team. What is this called and what are its risks?

🎯 **In one line.** A global surrogate model; the risk is that it explains the surrogate, not the original, so you must report its fidelity (how often it agrees with the black box) and keep it shallow enough to read.

⚙️ **The full answer.** Fit the tree to the black-box predictions (not the true labels). If fidelity is 92%, say so; regions of disagreement deserve local explanations (SHAP). Useful for case-study rounds: "this is the logic the lead scorer learned".

Q6.17

Compute the Gini impurity of a node with class counts (2, 3, 5) and the entropy of one with counts (3, 3, 6).

🎯 **In one line.** Gini = $1 - (0.04 + 0.09 + 0.25) = 0.62$; entropy = 1.5 bits.

⚙️ **The full answer.** Entropy: $p = (0.25, 0.25, 0.5)$; $-2(0.25 \log_2 0.25) - 0.5 \log_2 0.5 = 2(0.5) + 0.5 = 1.5$.

Q6.18

How do trees handle missing values?

🎯 **In one line.** Options: impute first; CART surrogate splits (use a correlated feature when the split feature is missing); C4.5 sends the row down all branches with fractional weights; XGBoost/LightGBM learn a default direction for missing values at each split.

⚙️ **The full answer.** XGBoost tries sending all missing rows left, then right, and keeps the better gain; this also captures “missingness is informative” (salary blank often means fresher). scikit-learn’s `DecisionTreeClassifier` supports NaN in recent versions; HistGradientBoosting always has.

Q6.19

💡 **Think deeper** A tree’s top split is on “days since last login”. The product manager concludes recent login causes applying. How do you respond?

🎯 **In one line.** A split shows predictive association, not causation; recent login may simply mark already-active job seekers, so I would propose an experiment (re-engagement nudges with a holdout) to measure the causal effect.

⚙️ **The full answer.** Trees pick the most predictive feature, and the most predictive feature is often a proxy for intent. Causal questions need interventions or causal designs. This is a common case-study trap.

CHAPTER 7

Ensembles, Bagging and Random Forests

🎥 **Companion video for this chapter:** [StatQuest: Random Forests Part 1, Building, Using and Evaluating](#)

Random forests were asked about in InfoEdge interviews in detail (Glassdoor): intuition, bootstrap sampling, random feature selection, how each tree is trained, why a forest reduces variance compared with one tree, RF vs gradient boosting, and feature importance. The OA asked whether RF is bagging or boosting and the types of ensembles. The 2026–27 talk even jotted “Random forest . . . probabilistic”. This chapter answers all of that from first principles.

7.1 Wisdom of crowds

Ask one person to guess the number of sweets in a jar: they may be far off. Average the guesses of 500 people and the answer is usually remarkably close. Two conditions make this work: each guess must be *somewhat* informed (better than random), and the errors must be *different* from person to person (not all fooled the same way). Ensembles apply this to models.

7.2 Three families of ensembles

- **Bagging** (bootstrap aggregating): train many models independently, each on a random resample of the data, and average them (or vote). Reduces *variance*. Example: random forest.
- **Boosting**: train models one after another, each focusing on the mistakes of the ensemble so far, and add them up. Reduces *bias* (and variance somewhat). Examples: AdaBoost, gradient boosting, XGBoost, LightGBM, CatBoost.
- **Stacking**: train several different models, then train a *meta-model* that learns how to combine their predictions (using out-of-fold predictions to avoid leakage).

Also: **voting** (hard: majority class; soft: average probabilities) and **blending** (stacking with a single holdout set instead of out-of-fold predictions).

7.3 Why averaging reduces variance: the formula

Suppose B models each have prediction variance σ^2 and every pair has correlation ρ . The variance of their average is

$$\text{Var}\left(\frac{1}{B} \sum_b \hat{f}_b\right) = \rho\sigma^2 + \frac{1-\rho}{B}\sigma^2.$$

- If the models were independent ($\rho = 0$): σ^2/B . Ten models cut variance tenfold.
- If identical ($\rho = 1$): σ^2 . No gain at all.
- As $B \rightarrow \infty$, the variance falls to $\rho\sigma^2$: *correlation* sets the floor.

So an ensemble wants many, individually strong, *decorrelated* models. Averaging does not change the bias (the average of models with the same bias has that bias).

Small example: how much does averaging help?

$\sigma^2 = 4$, $B = 10$. Independent: 0.4. With $\rho = 0.3$: $1.2 + 0.7 \cdot 4/10 = 1.2 + 0.28 = 1.48$. With 1,000 models and $\rho = 0.3$: $1.2 + 0.0028 \approx 1.2$. Beyond a point, more trees barely help; lower correlation does.

7.4 Majority voting: the maths of the crowd

If B independent classifiers each have accuracy $p > 0.5$, majority vote accuracy rises with B (Condorcet's jury theorem).

Small example: three voters

Each 80% accurate, independent. The vote is right if at least 2 of 3 are right: $P = 0.8^3 + 3 \cdot 0.8^2 \cdot 0.2 = 0.512 + 0.384 = 0.896$. With five voters at 0.6: $P = \sum_{k=3}^5 \binom{5}{k} 0.6^k 0.4^{5-k} = 0.3456 + 0.2592 + 0.0778 = 0.683$. If each were only 40% accurate, three voters give $0.4^3 + 3 \cdot 0.4^2 \cdot 0.6 = 0.352$: voting makes bad models worse.

Real models are correlated, so gains are smaller, which is again why random forests work hard to decorrelate trees.

7.5 Bootstrap sampling

A **bootstrap sample** draws n rows *with replacement* from n rows. Some rows appear several times, some not at all. The probability that a particular row is never chosen is $(1 - 1/n)^n \rightarrow e^{-1} \approx 0.368$. So each bootstrap sample contains about **63.2%** of the distinct rows; the other 36.8% are **out-of-bag (OOB)** for that model.

Small example: OOB probability

$n = 1000$: $(0.999)^{1000} = 0.3677$. Expected unique rows in a bootstrap of 100: $100(1 - 0.99^{100}) = 63.4$.

7.6 Random forest = bagged trees + random feature subsets

1. For $b = 1, \dots, B$ (say 500):
 - (a) Draw a bootstrap sample of the rows.
 - (b) Grow a deep tree on it (often unpruned), but **at each split consider only a random subset of m features** (default $m = \sqrt{d}$ for classification, $d/3$ or all features for regression in various libraries).
2. Predict by majority vote (or averaging class probabilities) for classification, or averaging for regression.

Why the random feature subset? Suppose one feature (say "preferred city match") is very strong. Plain bagging would put it at the root of nearly every tree, so the trees look alike and are highly correlated, and the $\rho\sigma^2$ floor stays high. Forcing each split to choose from a random subset sometimes hides the dominant feature, so other features get to shape the tree. The trees become less correlated; their average has lower variance. Each individual tree gets slightly worse (a bit more bias), but the forest gets better.

Small example: the subset size

64 features, classification: $m = \sqrt{64} = 8$ features per split. If one feature dominates, the probability it is *not* available at a given split is $1 - 8/64 = 87.5\%$.

7.7 Out-of-bag error: free validation

Each row is OOB for about a third of the trees. Predict each row using only the trees that did not see it, and compare with its label. That OOB error is an almost unbiased estimate of test error, without a separate validation set. It is also how RF permutation importance was originally computed.

7.8 Feature importance in a forest

- **Mean decrease in impurity (MDI):** average over trees of the total impurity decrease from splits on each feature. Fast; biased toward high-cardinality/continuous features.
- **Permutation importance** (mean decrease in accuracy): shuffle a feature's values in OOB or validation data and see how much accuracy drops. Less biased; slower; correlated features share credit.
- **SHAP:** per-prediction attributions; TreeSHAP is exact for forests.

7.9 Variants and practical notes

- **Extra Trees** (extremely randomised trees): also pick thresholds at random (one per candidate feature) instead of the best one; usually no bootstrap. Even less correlation, faster training, slightly more bias.
- **Isolation Forest:** random trees with random splits used for anomaly detection: anomalies are isolated in fewer splits (shorter average path length).
- **Hyperparameters that matter:** `n_estimators` (more is never worse for accuracy, only slower; stops improving around 200–1000), `max_features` (the decorrelation knob), `min_samples_leaf`/`max_depth` (tree complexity), `class_weight` or balanced bootstrap for imbalance.
- **Parallel:** trees are independent, so training is embarrassingly parallel (`n_jobs=-1`).
- **Weaknesses:** large memory and slower prediction than one tree; no extrapolation; probabilities can be poorly calibrated (votes cluster away from 0 and 1); less accurate than well-tuned boosting on many tabular problems.

7.10 Random forest vs gradient boosting

	Random forest	Gradient boosting
Ensemble type	bagging (parallel, independent trees)	boosting (sequential, each tree fixes the previous)
Trees	deep, low bias, high variance	shallow (depth 3–8), high bias, low variance
Mainly reduces	variance	bias (and variance via shrinkage)
More trees	never overfits by itself (error plateaus)	can overfit; use early stopping
Tuning	forgiving, good defaults	sensitive (learning rate, depth, trees, regularisation)
Training	parallel across trees	sequential across trees (parallel within a tree)
Typical accuracy on tabular	very good	usually best
Noise / outliers in labels	robust	can chase noisy labels

7.11 Stacking done right

Level 0: several diverse models (logistic regression, RF, LightGBM, k -NN). Generate *out-of-fold* predictions for each training row (via k -fold CV) so every prediction is made by a model that did not see that row. Level 1: a simple model (often logistic regression) learns the best combination from these predictions. Without out-of-fold predictions, the meta-model learns to trust whichever base model overfit the most: a leakage bug.

In practice: ensembles at InfoEdge

- **Lead scoring / telecalling** (99acres): random forests are a robust first model on mixed user-behaviour features; their OOB score gives quick validation; permutation importance explains drivers to the sales team.
- **Fraud / fake-recruiter detection**: Isolation Forest flags unusual posting behaviour (hundreds of jobs in minutes, odd salary ranges) without labels.
- **Resdex ranking**: the “M10” ranking blocks on the Resdex slide are typically gradient-boosted learning-to-rank models; RF is the baseline they must beat.

7.12 Interview questions

Q7.1

★ **PYQ (OA: is Random Forest bagging or boosting?)** Is a random forest bagging or boosting? Justify.

🎯 **In one line.** Bagging: its trees are trained independently and in parallel on bootstrap samples and averaged; it adds random feature selection at each split to decorrelate the trees.

⚙️ **The full answer.** Boosting would train trees sequentially, each one on reweighted data or residuals from the ensemble so far. RF trees never see each other's errors. The “random” in random forest refers to two sources of randomness: bootstrap rows and random feature subsets

per split.

⚠ **Trap to avoid.** Calling RF “boosting because it combines weak learners”. RF’s trees are deep (strong but high-variance) learners.

Q7.2

★ **PYQ (OA: types of ensembles)** Name the types of ensemble methods, with an example and the error component each mainly attacks.

🎯 **In one line.** Bagging (random forest; reduces variance), boosting (AdaBoost, GBM, XGBoost, LightGBM; reduces bias), and stacking (meta-learner over diverse models; combines strengths), plus simple voting and blending.

⚙️ **The full answer.** Bagging averages high-variance models; boosting builds a strong model from weak, high-bias ones by sequential correction; stacking learns weights or interactions between models’ predictions. Mixture of experts is a related idea where a gating model chooses which expert to trust per input.

Q7.3

★ **PYQ (Glassdoor InfoEdge: Random Forest intuition)** Explain random forests end to end: bootstrap sampling, random feature selection, how each tree is trained and how predictions are made.

🎯 **In one line.** Build hundreds of deep trees, each on a bootstrap sample of the rows and each choosing every split from a random subset of features, then average their predictions (or vote); averaging many decorrelated, low-bias trees gives a low-variance model.

⚙️ **The full answer.** **Step 1, rows.** For each tree, draw n rows with replacement: about 63% unique rows; the rest are out-of-bag for that tree. **Step 2, features.** At every node, sample m of the d features (e.g. $\sqrt{d}$) and find the best split only among them. **Step 3, grow.** Usually deep (low bias); controlled by `min_samples_leaf`. No pruning by default. **Step 4, predict.** Classification: majority vote or mean of leaf class proportions (scikit-learn averages probabilities). Regression: mean of leaf means. **Why it works.** Deep trees have low bias but high variance; averaging cuts variance to about $\rho\sigma^2 + (1 - \rho)\sigma^2/B$; the two randomisations lower ρ .

📊 **Worked example.** $d = 64$, 500 trees, $m = 8$: each split sees 8 random features; each tree sees about 632 of 1,000 rows.

💬 **If the interviewer pushes.** “Is RF probabilistic?” Its `predict_proba` is the average of leaf class proportions across trees: a frequency, not a calibrated probability. It tends to be under-confident near 0 and 1 (averaging pulls toward the middle); calibrate with isotonic regression if probabilities are used for decisions.

Q7.4

★ **PYQ (Glassdoor InfoEdge)** Why does a random forest reduce variance and overfitting compared with a single tree?

🕒 **In one line.** Because averaging B models with variance σ^2 and pairwise correlation ρ gives variance $\rho\sigma^2 + (1 - \rho)\sigma^2/B$, and the forest's bootstrap and feature subsampling keep ρ low, while averaging leaves bias unchanged.

⚙️ **The full answer.** A single deep tree fits the noise of its particular sample; another sample gives a very different tree. Errors that come from sample-specific noise point in different directions for different trees and cancel in the average. Errors that come from a shared systematic blind spot (bias) do not cancel. Deep trees have little bias, so the forest's remaining error is small.

📊 **Worked example.** Single tree: 100% train, 78% test. 500-tree RF: 100% train, 87% test. Note RF training accuracy can still be 100%: it does not “fit less”, it generalises better.

💡 **If the interviewer pushes.** “Can a random forest overfit?” Adding trees cannot increase overfitting (the average converges), but the individual trees can overfit; with very noisy labels, limit `min_samples_leaf`. Also it can overfit through leakage, like any model.

Q7.5

★ **PYQ (Glassdoor InfoEdge)** Random forest vs gradient boosting: differences, and when would you choose each?

🕒 **In one line.** RF averages deep, independent trees to cut variance and is robust with little tuning; gradient boosting adds shallow trees sequentially, each fitting the current errors, to cut bias, and usually wins on accuracy when tuned, at the cost of sensitivity to hyperparameters and noise.

⚙️ **The full answer.** See the comparison table. Choose RF for a fast, strong baseline, small or noisy datasets, or when you need OOB validation and easy parallelism. Choose GBM (LightGBM/XGBoost) for maximum accuracy on large tabular data, custom losses (ranking, Poisson, quantile), and when you can afford tuning and early stopping.

💡 **If the interviewer pushes.** “Why does GBM overfit with more trees and RF not?” In RF each new tree is an independent sample; the average converges to a limit. In GBM each new tree fits the remaining residuals, which eventually are noise: it keeps reducing training loss by fitting that noise.

Q7.6

★ **PYQ (Glassdoor InfoEdge: feature importance)** How does a random forest compute feature importance, and what are its pitfalls?

🕒 **In one line.** By default (MDI) it sums the impurity decreases from all splits on each feature, averaged over trees; permutation importance (drop in OOB/validation score when a feature is shuffled) is the more reliable alternative.

⚙️ **The full answer.** MDI is computed on training data and favours features with many split points (continuous, high-cardinality), even pure noise. Permutation importance is computed on held-out data and reflects real predictive value, but correlated features hide each other. Use both; drop-column importance (retrain without the feature) is the gold standard but costly.

Q7.7

What is the probability that a given row is out-of-bag for one tree when $n = 1000$? Why does it approach 0.368?

🎯 **In one line.** $(1 - 1/1000)^{1000} \approx 0.368$, because $(1 - 1/n)^n \rightarrow e^{-1}$.

⚙️ **The full answer.** Each of the n draws misses the row with probability $1 - 1/n$; draws are independent. The limit is the definition of e^{-1} . Hence a bootstrap contains about 63.2% unique rows.

Q7.8

Compute the variance of the average of 10 models with variance 4 when (a) independent and (b) pairwise correlation 0.5.

🎯 **In one line.** (a) 0.4; (b) $0.5 \cdot 4 + 0.5 \cdot 4/10 = 2.2$.

⚙️ **The full answer.** Correlation dominates: even infinitely many models with $\rho = 0.5$ only halve the variance.

Q7.9

Three independent classifiers with 80% accuracy vote. What is the ensemble accuracy? What if each is 40% accurate?

🎯 **In one line.** 0.896 and 0.352: voting helps better-than-chance models and hurts worse-than-chance ones.

Q7.10

What is OOB error and how does it compare with cross-validation?

🎯 **In one line.** OOB error predicts each row using only the trees that did not train on it, giving a nearly free estimate of generalisation error, comparable to cross-validation with about 1/3 holdout per tree.

⚙️ **The full answer.** It is slightly pessimistic because each OOB prediction uses only about a third of the trees. It is valid only if rows are independent: with grouped or time data, OOB leaks just like a random split.

Q7.11

🧠 **Think deeper** Would bagging help a linear regression? A k -NN with $k = 1$? Why or why not?

🎯 **In one line.** Bagging helps unstable, high-variance learners (deep trees, 1-NN somewhat, neural nets) and barely helps stable ones like linear regression, whose predictions hardly change across bootstrap samples.

⚙️ **The full answer.** For OLS, the bagged average is very close to the full-data OLS fit: variance is already low and the models are highly correlated ($\rho \approx 1$). For 1-NN, bootstrapping changes which neighbour is nearest for some points, so it helps a bit, but 1-NN is surprisingly stable to bootstrap (each point is present with probability 63%). Subsampling features (random subspaces) helps k -NN more.

Q7.12

🧠 **Think deeper** Why does a random forest choose $\sqrt{d}$ features per split and not just 1, or all d ?

🕒 **In one line.** It trades individual tree strength against correlation: $m = d$ gives strong but highly correlated trees (plain bagging), $m = 1$ gives very decorrelated but weak, nearly random trees; $\sqrt{d}$ is an empirical sweet spot, and the right m should be tuned.

⚙️ **The full answer.** Breiman showed the forest's error bound depends on the ratio $\bar{\rho}/s^2$ (mean correlation over squared strength). Smaller m lowers both $\bar{\rho}$ and strength s . With many noisy features and few informative ones, a small m often misses the informative features at a split, so larger m (like $d/3$ or $0.5d$) is better.

Q7.13

🧠 **Think deeper** A random forest outputs probability 0.9. Is it calibrated? How would you check and fix it?

🕒 **In one line.** Not necessarily; RF probabilities are vote fractions that tend to be pulled away from 0 and 1; check with a reliability curve or Brier score on held-out data and fix with isotonic or Platt calibration.

⚙️ **The full answer.** For an RF to output 1.0, every tree must agree; feature subsampling makes some trees always disagree, so extremes are rare (under-confidence at the edges). Conversely, very deep trees with few rows per leaf give overconfident leaves. `CalibratedClassifierCV` fits the calibration map with cross-validation.

Q7.14

Explain stacking and how to avoid leakage in it.

🕒 **In one line.** Stacking trains a meta-model on the predictions of base models; to avoid leakage, the meta-model must be trained on out-of-fold predictions, never on predictions the base models made for rows they trained on.

⚙️ **The full answer.** If a base model overfits (100% training accuracy), its in-sample predictions are perfect, and the meta-model learns to trust it completely; on test data it fails. With k -fold, each row's base prediction comes from a model trained on the other folds. Test-time base predictions come from models retrained on all data (or the average of fold models).

Q7.15

Hard voting vs soft voting: compute both for five classifiers giving $P(\text{class } 1) = 0.9, 0.6, 0.4, 0.45, 0.3$.

🕒 **In one line.** Hard voting: two vote 1, three vote 0, so class 0; soft voting: average 0.53, so class 1.

⚙️ **The full answer.** Soft voting uses confidence: one very confident model can outweigh several lukewarm ones. It needs reasonably calibrated probabilities.

Q7.16

How does Isolation Forest detect anomalies?

🕒 **In one line.** It builds random trees with random features and random split points; anomalies, being few and different, get isolated in very few splits, so a short average path length means a high anomaly score.

⚙️ **The full answer.** A normal point in a dense region needs many random cuts to be separated from its neighbours. An isolated point (a recruiter posting 400 jobs in an hour) is cut off quickly. Score $s = 2^{-\mathbb{E}[h(x)]/c(n)}$; close to 1 means anomaly. No labels needed, linear time, works in high dimensions.

Q7.17

💡 **Think deeper** With 500 features of which only 2 are informative, how will a random forest behave and what would you do?

🕒 **In one line.** With $m = \sqrt{500} \approx 22$, most splits will not even see the two informative features, so many splits are on noise and accuracy suffers; increase m , do feature selection first, or use boosting, which picks the best feature at each split.

⚙️ **The full answer.** Probability a given informative feature is in a split's subset is $22/500 = 4.4\%$; both absent most of the time. Trees still find them sometimes, but noise splits add variance. Remedies: raise `max_features`, remove noise features (permutation or Boruta), or use L1 logistic regression / GBM.

Q7.18

Which hyperparameters of a random forest matter most and how would you tune them?

🕒 **In one line.** `max_features` (decorrelation), `min_samples_leaf`/`max_depth` (tree complexity), and enough `n_estimators`; tune the first two with CV or OOB, set trees high enough that OOB error has plateaued.

⚙️ **The full answer.** `n_estimators` is not a regularisation parameter: more trees only reduce noise in the average. `class_weight="balanced_subsample"` helps imbalance. `max_samples` (bootstrap size) can speed training and add diversity.

Q7.19

 **Think deeper** Is a random forest a “black box”? How would you explain one prediction to a recruiter?

🎯 **In one line.** It is hard to read globally but easy to explain locally: TreeSHAP gives each feature’s contribution to a specific prediction, and you can show the top contributing features in plain words.

⚙️ **The full answer.** Example output for a Resdex candidate: base score 0.10; +0.25 because skills match 9 of 10 required; +0.12 because notice period is 15 days; −0.05 because expected salary is above budget. This is the “Explainability: why this candidate, in plain language” idea from the PremiumX slide.

Q7.20

What is the difference between bagging and pasting? Between random forest and Extra Trees?

🎯 **In one line.** Bagging samples rows with replacement, pasting without; Extra Trees additionally choose split thresholds at random and usually use the whole dataset per tree, which further decorrelates trees and speeds training.

⚙️ **The full answer.** Extra Trees trade a little bias for less variance and are often as accurate as RF and faster, especially with many features.

CHAPTER 8

Boosting: AdaBoost, Gradient Boosting, XGBoost and LightGBM

🎥 **Companion video for this chapter:** [StatQuest: Gradient Boost Part 1, Regression Main Ideas](#)

Gradient-boosted trees are the workhorse of tabular machine learning: ranking in search, lead scoring, fraud, pricing. The syllabus names boosting, gradient boosting machines, XGBoost and LightGBM separately, so expect several questions. The aim of this chapter is that you can do one round of gradient boosting by hand, write the XGBoost gain formula, and explain what LightGBM changes.

8.1 The idea: learn from your mistakes, a little at a time

Imagine predicting flat prices. Your first guess for every flat is the average price. Look at the errors: big flats are under-predicted, small ones over-predicted. Fit a small model to those errors and add a fraction of it to your guess. Look at the new errors: maybe flats near metro stations are still under-predicted. Fit another small model to those errors, add a fraction. Repeat hundreds of times. Each step fixes some of what is left. That is boosting.

Key words:

- **Weak learner:** a simple model only a little better than guessing, usually a shallow tree (a “stump” has one split).
- **Additive model:** the final model is a sum $F_M(x) = F_0(x) + \nu \sum_{m=1}^M h_m(x)$.
- **Shrinkage / learning rate ν** (0.01–0.3): each new tree’s contribution is scaled down so no single tree dominates and later trees can correct earlier ones.

8.2 AdaBoost (adaptive boosting)

For binary labels $y \in \{-1, +1\}$:

1. Start with equal weights $w_i = 1/n$ on every example.
2. For $t = 1, \dots, T$:
 - (a) Fit a stump h_t minimising weighted error $\varepsilon_t = \sum_i w_i [h_t(x_i) \neq y_i]$.
 - (b) Give it a say $\alpha_t = \frac{1}{2} \ln \frac{1-\varepsilon_t}{\varepsilon_t}$ (accurate stumps get big say; $\varepsilon = 0.5$ gets 0).
 - (c) Reweight: $w_i \leftarrow w_i e^{-\alpha_t y_i h_t(x_i)}$ (misclassified points’ weights grow by e^{α_t} , correct ones shrink by $e^{-\alpha_t}$), then normalise to sum 1.
3. Final classifier: $\text{sign}(\sum_t \alpha_t h_t(x))$.

📖 Small example: one AdaBoost round

Five points, weights 0.2 each. A stump misclassifies one: $\varepsilon = 0.2$. $\alpha = \frac{1}{2} \ln(0.8/0.2) = \frac{1}{2} \ln 4 = 0.693$. Misclassified weight: $0.2e^{0.693} = 0.4$. Correct ones: $0.2e^{-0.693} = 0.1$ each. Sum = $0.4 + 0.4 = 0.8$; normalise: the misclassified point now has weight 0.5, the others 0.125. The next stump is forced

to care about the hard point: exactly half the total weight is on it, so the previous stump would now have weighted error 0.5 (useless), guaranteeing a different stump.

AdaBoost is equivalent to forward stagewise additive modelling with the **exponential loss** $e^{-yF(x)}$. That loss punishes misclassified points exponentially, so AdaBoost is sensitive to label noise and outliers.

8.3 Gradient boosting: gradient descent on predictions

8.3.1 The recipe

Choose any differentiable loss $\ell(y, F)$.

1. Initialise with the best constant: $F_0 = \arg \min_c \sum_i \ell(y_i, c)$ (the mean for MSE; the log-odds of the base rate for log loss).
2. For $m = 1, \dots, M$:
 - (a) Compute **pseudo-residuals**: $r_i = -\frac{\partial \ell(y_i, F)}{\partial F} \Big|_{F=F_{m-1}(x_i)}$. For MSE $(\frac{1}{2}(y - F)^2)$ this is exactly the residual $y_i - F_{m-1}(x_i)$.
 - (b) Fit a regression tree h_m to the pairs (x_i, r_i) .
 - (c) Set each leaf's value to the best constant step for the actual loss in that leaf.
 - (d) Update $F_m = F_{m-1} + \nu h_m$.

💡 Intuition: why "gradient"?

In ordinary gradient descent we move parameters by minus the gradient. Here, think of the n predictions $F(x_1), \dots, F(x_n)$ as the parameters. Minus the gradient of the total loss with respect to them is the vector of pseudo-residuals. A tree fitted to that vector is a step in (approximately) that direction that also tells us what to do for new x . So boosting is gradient descent in the space of functions.

📊 Small example: two rounds of gradient boosting by hand (MSE)

Targets $y = (3, 5, 10)$ for three flats (in tens of lakhs), $\nu = 0.1$. $F_0 = \text{mean} = 6$. Residuals $r = (-3, -1, 4)$. Suppose tree 1 (a stump) puts flats 1–2 in one leaf (mean residual -2) and flat 3 in another (residual 4). Update: $F_1 = (6 - 0.2, 6 - 0.2, 6 + 0.4) = (5.8, 5.8, 6.4)$. New residuals $(-2.8, -0.8, 3.6)$; MSE fell from $\frac{9+1+16}{3} = 8.67$ to $\frac{7.84+0.64+12.96}{3} = 7.15$. Hundreds of such small steps drive it down.

8.3.2 Gradient boosting for classification

With log loss and raw score F (log-odds), $p = \sigma(F)$, the pseudo-residual is $y - p$: "label minus predicted probability". $F_0 = \log \frac{y}{1-y}$. For 20% positives, $F_0 = \log(0.25) = -1.386$.

8.3.3 Regularising a boosted model

- Learning rate ν (smaller needs more trees but generalises better).
- Number of trees M chosen by **early stopping** on validation data.

- Tree size: max depth 3–8, or max leaves; min samples (or min Hessian) per leaf.
- **Stochastic gradient boosting**: each tree sees a random subsample of rows (e.g. 80%) and columns.
- L1/L2 penalties on leaf values (XGBoost).

8.4 XGBoost: regularised, second-order, engineered

8.4.1 The objective

XGBoost minimises loss plus a complexity penalty for each tree:

$$\text{Obj} = \sum_i \ell(y_i, \hat{y}_i) + \sum_m \Omega(h_m), \quad \Omega(h) = \gamma T + \frac{1}{2} \lambda \sum_{j=1}^T w_j^2,$$

T = number of leaves, w_j = leaf values. γ charges for each leaf; λ shrinks leaf values.

8.4.2 Second-order approximation

When adding tree h to current predictions, expand the loss with a Taylor series to second order. For each row, $g_i = \partial \ell / \partial \hat{y}_i$ (gradient) and $h_i = \partial^2 \ell / \partial \hat{y}_i^2$ (Hessian, curvature). For rows in leaf j , with sums $G_j = \sum g_i$ and $H_j = \sum h_i$:

$$\text{Obj}_j \approx G_j w_j + \frac{1}{2} (H_j + \lambda) w_j^2 + \gamma.$$

This is a parabola in w_j , minimised at

$$w_j^* = -\frac{G_j}{H_j + \lambda}, \quad \text{Obj}_j^* = -\frac{1}{2} \frac{G_j^2}{H_j + \lambda} + \gamma.$$

So the gain of splitting a node into left and right is

$$\text{Gain} = \frac{1}{2} \left[\frac{G_L^2}{H_L + \lambda} + \frac{G_R^2}{H_R + \lambda} - \frac{(G_L + G_R)^2}{H_L + H_R + \lambda} \right] - \gamma.$$

A split is made only if $\text{Gain} > 0$: γ is a minimum gain, which prunes useless splits.

Small example: leaf weight and gain

MSE loss $\frac{1}{2}(y - \hat{y})^2$: $g = \hat{y} - y = -\text{residual}$, $h = 1$. Leaf with residuals (2, 3, 4), $\lambda = 1$: $G = -9$, $H = 3$, $w^* = 9/4 = 2.25$ (the mean residual 3, shrunk by λ). Split gain: left $G_L = -6$, $H_L = 2$; right $G_R = 4$, $H_R = 3$; $\lambda = 1$, $\gamma = 0$: $\frac{1}{2}[36/3 + 16/4 - 4/6] = \frac{1}{2}[12 + 4 - 0.667] = 7.67$. For log loss with $p = 0.8$, $y = 0$: $g = p - y = 0.8$, $h = p(1 - p) = 0.16$.

Why second order? It gives the optimal leaf value in one Newton step (not just a direction), it works uniformly for any twice-differentiable loss (a custom objective only needs g and h), and the Hessian acts as a natural “weight” of each row (`min_child_weight` is a minimum sum of Hessians in a leaf).

8.4.3 Engineering features

Sparsity-aware split finding with a learned **default direction for missing values**; approximate and histogram split finding with weighted quantile sketches; column block structure for parallel split search within a tree; cache-aware access; out-of-core training; GPU support; column and row subsampling; monotonic and interaction constraints.

8.5 LightGBM: faster on big data

- **Histogram-based splits**: continuous features bucketed into at most 255 bins; split search costs $O(\#bins)$ per feature. A child's histogram is the parent's minus its sibling's ("histogram subtraction"), halving the work.
- **Leaf-wise (best-first) growth**: always split the leaf with the largest gain anywhere in the tree, instead of growing level by level. Lower loss for the same number of leaves; can grow deep, lopsided trees and overfit small data, so control it with `num_leaves`, `max_depth` and `min_data_in_leaf`.
- **GOSS** (gradient-based one-side sampling): keep all rows with large gradients (badly predicted), randomly sample a fraction of small-gradient rows and up-weight them by $(1 - a)/b$ to keep the gradient sums unbiased.
- **EFB** (exclusive feature bundling): bundle sparse features that are rarely non-zero together into one feature (one-hot columns), reducing the feature count.
- **Native categorical splits**.

CatBoost (for completeness): ordered target statistics for categorical features (encode a row using only rows "before" it in a random permutation, preventing target leakage) and ordered boosting; symmetric (oblivious) trees that are fast at inference.

8.6 Learning to rank with boosted trees

Search and recommendation (Resdex, job search, Jeevansathi) need *orderings*. **LambdaMART** is gradient boosting whose pseudo-gradients ("lambdas") come from swapping pairs of results and measuring the change in NDCG. LightGBM's `objective="lambdarank"` and XGBoost's `rank:ndcg` implement it. Training data are grouped by query.

In practice: boosting at InfoEdge

- **Resdex ranking blocks** ("M10" per segment on the slide): LambdaMART-style GBMs over query-candidate features (skill match, title similarity, experience fit, recency, past recruiter actions), then personalisation.
- **Telecalling lead ranking** (99acres): LightGBM classifier on user activity features; ranked list feeds the call queue.
- **Salary engine**: GBM with `objective="gamma"` or on log salary; quantile objectives for the band.
- Practical defaults: LightGBM, `learning_rate=0.05`, `num_leaves=31-127`, `min_data_in_leaf=50`, `feature_fraction=0.8`, `bagging_fraction=0.8`, early stopping with patience 100.

8.7 Interview questions

Q8.1

★ **PYQ (OA: types of ensemble, boosting)** What is boosting, and how is it different from bagging?

🎯 **In one line.** Boosting trains weak learners sequentially, each one concentrating on what the ensemble so far gets wrong, and adds them up with small weights; bagging trains strong learners independently on bootstrap samples and averages them.

⚙️ **The full answer.** Bagging reduces variance of low-bias models; boosting reduces bias of high-bias models (stumps or shallow trees), and its shrinkage and subsampling also control variance. Bagging is parallel; boosting is inherently sequential across rounds. Boosting can overfit with too many rounds; bagging plateaus.

Q8.2

★ **PYQ (Glassdoor InfoEdge: RF vs gradient boosting; INT: why GBM overfits with more trees)** Explain why adding trees to a gradient-boosted model can overfit while adding trees to a random forest does not.

🎯 **In one line.** Each boosting tree fits whatever residual remains; once the real signal has been fitted the residual is noise, and further trees keep fitting it; RF trees are independent draws whose average converges, so more trees only reduce the randomness of the average.

⚙️ **The full answer.** GBM's training loss decreases monotonically with M ; validation loss is U-shaped in M , so we use early stopping. Shrinkage ν makes the U flatter and wider. RF's validation error decreases and flattens as B grows; it never turns up from adding trees.

📊 **Worked example.** LightGBM with $\nu = 0.1$: validation AUC peaks at 420 trees (0.861) and falls to 0.853 by 2,000. With $\nu = 0.02$, the peak moves to about 2,100 trees at 0.864.

Q8.3

Walk through gradient boosting for regression with MSE, by hand, for targets 3, 5, 10.

🎯 **In one line.** Start with the mean 6, compute residuals $(-3, -1, 4)$, fit a small tree to them, add a fraction ν of its predictions, and repeat on the new residuals.

⚙️ **The full answer.** See the worked example: after one round with $\nu = 0.1$, predictions $(5.8, 5.8, 6.4)$, MSE $8.67 \rightarrow 7.15$.

💬 **If the interviewer pushes.** “What does each tree predict, exactly?” The negative gradient of the loss with respect to the current prediction: for MSE, the residuals; for log loss, $y - p$; for absolute loss, the sign of the residual (then leaf values are medians).

Q8.4

Why is it called *gradient* boosting?

🕒 **In one line.** Because each new tree is fitted to the negative gradient of the loss with respect to the current predictions, so adding it is a gradient-descent step in function space, which works for any differentiable loss.

⚙️ **The full answer.** For MSE the negative gradient happens to be the residual, which is why gradient boosting is often described as “fitting residuals”; for other losses it is not the residual. The learning rate is the step size; the tree approximates the gradient direction and generalises it to unseen x .

Q8.5

Compute AdaBoost’s α for weighted error 0.2, and the new weights for five equally weighted points with one mistake.

🕒 **In one line.** $\alpha = \frac{1}{2} \ln 4 = 0.693$; after normalising, the misclassified point has weight 0.5 and each correct point 0.125.

💡 **If the interviewer pushes.** “What if $\varepsilon = 0.5$?” $\alpha = 0$: the stump is ignored. “ $\varepsilon > 0.5$?” $\alpha < 0$: flip its predictions. “Which loss does AdaBoost minimise?” Exponential loss e^{-yF} .

Q8.6

Derive the XGBoost optimal leaf weight and split gain.

🕒 **In one line.** A second-order Taylor expansion makes each leaf’s objective $Gw + \frac{1}{2}(H + \lambda)w^2 + \gamma$, minimised at $w^* = -G/(H + \lambda)$ with value $-\frac{1}{2}G^2/(H + \lambda) + \gamma$; the split gain is the improvement from replacing one leaf by two.

⚙️ **The full answer.** Expand $\ell(y, \hat{y} + w) \approx \ell + gw + \frac{1}{2}hw^2$. Summing over rows in a leaf and adding $\frac{1}{2}\lambda w^2$ and γ gives the parabola. Setting the derivative $G + (H + \lambda)w = 0$ gives w^* . Plugging back gives the score. Gain = score(parent) – score(left) – score(right), which yields the boxed formula with the $-\gamma$ term.

📊 **Worked example.** Residuals (2, 3, 4), $\lambda = 1$: $w^* = 2.25$. The split example gives gain 7.67.

💡 **If the interviewer pushes.** “What does λ do?” Shrinks leaf values toward 0, more so for leaves with small Hessian sums (few rows): a data-dependent shrinkage. “And γ ?” Minimum loss reduction to make a split: post-pruning.

Q8.7

Why does XGBoost use the Hessian?

🕒 **In one line.** The second derivative gives the optimal step size per leaf (a Newton step) instead of only a direction, makes the method work for any twice-differentiable loss with one formula, and acts as a per-row weight.

⚙️ **The full answer.** For log loss, $h = p(1 - p)$: rows the model is already sure about have small Hessians and count less in deciding leaf values, which is like IRLS in logistic regression. For MSE, $h = 1$ and XGBoost reduces to ordinary gradient boosting with L2-shrunk leaf means.

Q8.8

★ **PYQ (INT and coding round: XGBoost and missing values)** How does XGBoost handle missing values?

🕒 **In one line.** At every split it learns a default direction: it tries sending all rows with a missing value left, then right, and keeps the direction with the higher gain; at prediction time missing values follow that learned direction.

⚙️ **The full answer.** This is “sparsity-aware split finding”. It only scans non-missing values for thresholds (fast on sparse data) and treats missingness as potentially informative (a blank current salary might indicate a fresher, which the default direction captures). It is learned per split, so the same feature’s missing values can go different ways in different parts of the tree.

Q8.9

Which XGBoost/LightGBM parameters control overfitting?

🕒 **In one line.** Learning rate with early stopping, tree complexity (max_depth, num_leaves, min_child_weight / min_data_in_leaf, gamma), regularisation (lambda, alpha), and randomness (subsample, colsample_bytree).

⚙️ **The full answer.** Order of tuning I use: fix $\nu = 0.05$ and early stopping; tune tree size (num_leaves/max_depth, min_data_in_leaf); then subsampling; then λ/α ; finally lower ν and retrain for the last bit.

Q8.10

What is LightGBM’s leaf-wise growth and why can it overfit?

🕒 **In one line.** It splits the single leaf with the largest gain anywhere in the tree rather than all leaves at one depth, reaching lower loss with fewer leaves but producing deep, unbalanced branches that can fit noise on small data.

⚙️ **The full answer.** Level-wise growth wastes splits on leaves with little gain; leaf-wise does not. But a few branches can become very deep. Constrain with num_leaves (keep it below $2^{\text{max_depth}}$), max_depth, and min_data_in_leaf.

Q8.11

Explain GOSS and EFB.

🕒 **In one line.** GOSS keeps all large-gradient rows and a random sample of small-gradient rows (up-weighted) to compute splits faster with little bias; EFB merges mutually exclusive sparse features into single features to cut the feature count.

⚙️ **The full answer.** Rows with small gradients are already well fitted and contribute little to gain. Keep top a fraction, sample b of the rest, multiply sampled rows' gradients by $(1 - a)/b$. Example: $a = 0.2, b = 0.1$: weight = 8. EFB uses the fact that one-hot columns are never non-zero together, so they can share one feature with offset value ranges.

Q8.12

💡 **Think deeper** Can you boost linear models? Can you boost a neural network?

🎯 **In one line.** Boosting linear regression with squared loss just converges to a linear model (the sum of linear models is linear), so it gains only regularisation-path behaviour, not new expressiveness; boosting non-linear weak learners like small trees or small nets is what adds power.

⚙️ **The full answer.** Componentwise linear boosting (one feature at a time, small steps) is equivalent to a form of forward stagewise regression and approximates the Lasso path: useful for feature selection. XGBoost has `booster="gblinear"`. Neural networks can be boosted (e.g. AdaNet, boosted residual networks), and ResNets can be viewed as an additive, boosting-like stack of corrections.

Q8.13

💡 **Think deeper** Why does boosting use shallow trees but random forests use deep ones?

🎯 **In one line.** Boosting reduces bias by adding many small corrections, so each learner should be weak and low-variance; bagging reduces variance by averaging, so each learner should be low-bias even if high-variance.

⚙️ **The full answer.** A depth- k tree captures interactions among up to k features; depth 4–8 is enough for most real interactions. Deep trees in boosting would fit noise in early rounds with nothing to average it away. Shallow trees in RF would all be similar and biased, and averaging cannot remove bias.

Q8.14

What are the pseudo-residual and initial value for log-loss gradient boosting? Compute F_0 for 20% positives.

🎯 **In one line.** Pseudo-residual $y - p$; $F_0 = \log \frac{0.2}{0.8} = -1.386$.

Q8.15

💡 **Think deeper** Your boosted model is excellent offline but its predicted probabilities seem too extreme. Why, and what do you do?

🎯 **In one line.** Many boosting rounds keep pushing scores of separable training points outward, especially with deep trees and little regularisation, so log-odds become overconfident; fix with

stronger regularisation and early stopping, then calibrate on held-out data.

⚙️ **The full answer.** Log loss keeps rewarding more confident correct predictions; with high capacity the model memorises. Check reliability curves; apply isotonic/Platt calibration; consider `min_child_weight` and λ .

Q8.16

How does CatBoost handle categorical features, and why is naive target encoding dangerous?

🎯 **In one line.** It uses ordered target statistics: each row's category is encoded with the average target of rows that come before it in a random permutation; naive target encoding uses the row's own label, leaking the target and overfitting rare categories.

⚙️ **The full answer.** A company that appears once with label 1 gets encoding 1.0: a perfect but fake signal. Out-of-fold or ordered encoding plus a prior ($\frac{\text{sum} + a \cdot \text{prior}}{\text{count} + a}$) solves it.

Q8.17

What are monotonic constraints and when would you use them?

🎯 **In one line.** They force predictions to only increase (or only decrease) with a feature, which encodes domain knowledge, stabilises the model and makes it defensible.

⚙️ **The full answer.** A salary model should not predict lower salary for more experience, all else equal; a credit-style lead score should not drop when a user views more listings in their price band. XGBoost/LightGBM: `monotone_constraints`. Often costs little accuracy and prevents embarrassing artefacts.

Q8.18

🧠 **Think deeper** Why might gradient boosting be hurt more than a random forest by mislabelled data?

🎯 **In one line.** Boosting concentrates effort on the examples it currently gets wrong, and mislabelled examples are permanently "wrong", so later trees keep chasing them; a random forest's averaging treats them as noise.

⚙️ **The full answer.** AdaBoost's exponential loss makes it extreme; gradient boosting with log loss is gentler, and with Huber or other robust losses even more so. Mitigations: robust loss, stronger regularisation, subsampling, early stopping, label cleaning.

Q8.19

How would you choose between XGBoost, LightGBM and CatBoost?

🎯 **In one line.** LightGBM for speed on large data, CatBoost when there are many high-cardinality categorical features and little time to engineer them, XGBoost for maturity and fine control;

accuracy differences are usually small after tuning.

⚙️ **The full answer.** In practice try LightGBM first, compare with CatBoost with default settings, pick by validation metric and latency.

Q8.20

What is LambdaMART and why is it used for search ranking?

🎯 **In one line.** It is gradient-boosted trees trained with “lambda” gradients derived from how much swapping two results would change a ranking metric like NDCG, so it optimises the ordering directly.

⚙️ **The full answer.** Pointwise models (predict relevance per item) do not care about order; pairwise (RankNet) cares about pairs; LambdaMART weights each pair by the NDCG change, focusing on the top of the list where users look. Data: (query, candidate, relevance grade) grouped by query.

📦 **At InfoEdge scale.** Resdex: for a recruiter query, rank 30 lakh profiles by fit; LambdaMART on retrieval features is the classic ranker.

Q8.21

Can gradient boosting be parallelised?

🎯 **In one line.** Not across rounds (each tree depends on the previous), but within each round the split search over features and histogram building is parallelised across threads, machines and GPUs.

Q8.22

💡 **Think deeper** A GBM ranks feature “user_id” as the most important. What is happening?

🎯 **In one line.** The model is memorising users: it is either leakage across the split (same users in train and validation) or overfitting to identity rather than behaviour; remove the ID, split by user, and use behavioural features.

⚙️ **The full answer.** High-cardinality IDs let trees isolate individual users; with a random split, the same users appear in validation, so it even seems to help. A group split by user will reveal the drop.

CHAPTER 9


Support Vector Machines and the Kernel Trick

► **Companion video for this chapter:** [StatQuest: Support Vector Machines Part 1, Main Ideas](#)

InfoEdge interviews reportedly asked “what is an SVM?” and expected, in order: hinge loss, kernels, margin and support vectors. The OA asked SVM statements and that hinge loss is a classification loss. The “kernel methods” line on the Tools slide shows the team uses them. This chapter builds SVMs slowly enough that you can derive the margin and explain the kernel trick to someone who has never seen it.

9.1 Which line is best?

Two classes of points can be separated by infinitely many lines. Which should we choose? Intuition: the line that stays as far as possible from both classes. New points are slightly different from training points; a line hugging one class would misclassify them at the first wobble. The distance from the line to the nearest points is the **margin**. An SVM picks the line with the **maximum margin**: the widest empty “street” between the classes.



Only the points on the edges of the street matter. They are the **support vectors**. Move or delete any other point (without crossing the street) and the solution does not change.

9.2 The geometry you need

A hyperplane is $\mathbf{w}^T \mathbf{x} + b = 0$. The vector $\mathbf{w}$ is perpendicular to it. The signed distance from a point $\mathbf{x}$ to the plane is

$$\frac{\mathbf{w}^T \mathbf{x} + b}{\|\mathbf{w}\|}.$$

📊 **Small example: distance to a plane**

$\mathbf{w} = (3, 4)$, $b = -5$, point $(1, 1)$: $\frac{3+4-5}{\sqrt{9+16}} = \frac{2}{5} = 0.4$.

Labels are $y \in \{-1, +1\}$. A point is correctly classified when $y(\mathbf{w}^T \mathbf{x} + b) > 0$.

► **Stuck on this page?** Watch: [StatQuest: Support Vector Machines Part 1, Main Ideas](#)

9.3 Hard margin: the optimisation problem

Scaling $\mathbf{w}$ and b by the same constant does not move the plane. Use that freedom: scale so the closest points satisfy $y(\mathbf{w}^\top \mathbf{x} + b) = 1$ exactly. Then their distance to the plane is $1/\|\mathbf{w}\|$ and the street width is $2/\|\mathbf{w}\|$. Maximising $2/\|\mathbf{w}\|$ is the same as minimising $\frac{1}{2}\|\mathbf{w}\|^2$ (squared for smoothness):

$$\min_{\mathbf{w}, b} \frac{1}{2}\|\mathbf{w}\|^2 \quad \text{subject to} \quad y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1 \text{ for all } i.$$

This is a convex quadratic program: a bowl-shaped objective with linear constraints, so there is one global solution.

Small example: the simplest SVM

Two points: $(1, 1)$ with $y = +1$ and $(-1, -1)$ with $y = -1$. By symmetry the boundary is perpendicular to the segment joining them and passes through the origin: $\mathbf{w} = (a, a)$, $b = 0$. Constraint for $(1, 1)$: $2a \geq 1$, smallest $\|\mathbf{w}\|$ at $a = 0.5$. So $\mathbf{w} = (0.5, 0.5)$, margin $2/\|\mathbf{w}\| = 2/0.707 = 2.83$, which is exactly the distance between the two points. Both points are support vectors.

9.4 Soft margin: allowing mistakes, and the C parameter

Real data overlap, and one outlier could force a tiny margin. Introduce a **slack** $\xi_i \geq 0$ per point that measures how much it violates the margin:

$$\min \frac{1}{2}\|\mathbf{w}\|^2 + C \sum_i \xi_i \quad \text{s.t.} \quad y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1 - \xi_i, \quad \xi_i \geq 0.$$

- $\xi_i = 0$: outside the street on the correct side. $0 < \xi_i \leq 1$: inside the street, still correct. $\xi_i > 1$: misclassified.
- **Large C**: violations are expensive; narrow margin; fits training data hard; low bias, high variance. $C \rightarrow \infty$ is the hard margin.
- **Small C**: wide margin; tolerates violations; more bias, less variance.

9.4.1 The same thing as a loss function: hinge loss

At the optimum, $\xi_i = \max(0, 1 - y_i f(\mathbf{x}_i))$. Substituting gives an unconstrained problem:

$$\min_{\mathbf{w}, b} \frac{1}{2}\|\mathbf{w}\|^2 + C \sum_i \max(0, 1 - y_i f(\mathbf{x}_i)).$$

The term $\max(0, 1 - yf)$ is the **hinge loss**. So a linear SVM is “hinge loss + L2 regularisation”, just as logistic regression is “log loss + L2”.

Small example: hinge loss values

$y = +1$: $f = 2$ gives 0 (confidently right, ignored); $f = 0.3$ gives 0.7 (right but inside the margin); $f = -2$ gives 3 (wrong). $y = -1$, $f = -1.5$: $yf = 1.5$, loss 0.

Hinge loss is exactly zero beyond the margin: confidently correct points exert no pull. That is why only support vectors matter.

9.5 The dual problem: where kernels come from

Using Lagrange multipliers $\alpha_i \geq 0$ (one per constraint), the solution has the form

$$\mathbf{w} = \sum_i \alpha_i y_i \mathbf{x}_i,$$

and the α 's solve the **dual** problem

$$\max_{\alpha} \sum_i \alpha_i - \frac{1}{2} \sum_{i,j} \alpha_i \alpha_j y_i y_j \mathbf{x}_i^{\top} \mathbf{x}_j \quad \text{s.t.} \quad 0 \leq \alpha_i \leq C, \quad \sum_i \alpha_i y_i = 0.$$

Three facts follow (from the KKT conditions):

- $\alpha_i = 0$ for points strictly outside the margin: they do not appear in $\mathbf{w}$.
- $0 < \alpha_i < C$: points exactly on the margin.
- $\alpha_i = C$: points inside the margin or misclassified.

And the crucial observation: **the data appear only through dot products $\mathbf{x}_i^{\top} \mathbf{x}_j$** , both in training and in prediction, $f(\mathbf{x}) = \sum_i \alpha_i y_i \mathbf{x}_i^{\top} \mathbf{x} + b$.

9.6 The kernel trick

9.6.1 The problem: curved boundaries

Points inside a circle (class +1) and outside it (class -1) cannot be separated by a line. But map each point $\mathbf{x} = (x_1, x_2)$ to a new 3-D feature vector $\phi(\mathbf{x}) = (x_1^2, \sqrt{2}x_1x_2, x_2^2)$. The circle $x_1^2 + x_2^2 = r^2$ becomes the *plane* $z_1 + z_3 = r^2$. Linear separation in the new space is a curved boundary in the old.

9.6.2 The trick

We could compute ϕ for every point and run a linear SVM. But ϕ can be huge (degree-3 polynomial features of 1,000 inputs: over 167 million features) or infinite. However, the dual needs only dot products, and for many ϕ there is a shortcut:

$$\phi(\mathbf{x})^{\top} \phi(\mathbf{z}) = x_1^2 z_1^2 + 2x_1 x_2 z_1 z_2 + x_2^2 z_2^2 = (x_1 z_1 + x_2 z_2)^2 = (\mathbf{x}^{\top} \mathbf{z})^2.$$

We never build ϕ ; we square the ordinary dot product. Any function $K(\mathbf{x}, \mathbf{z}) = \phi(\mathbf{x})^{\top} \phi(\mathbf{z})$ is a **kernel**. Replace every $\mathbf{x}_i^{\top} \mathbf{x}_j$ by $K(\mathbf{x}_i, \mathbf{x}_j)$.

Small example: a kernel by hand

$K(\mathbf{x}, \mathbf{z}) = (\mathbf{x}^{\top} \mathbf{z} + 1)^2$, $\mathbf{x} = (1, 2)$, $\mathbf{z} = (3, -1)$: $\mathbf{x}^{\top} \mathbf{z} = 3 - 2 = 1$, so $K = 4$. This equals a dot product in a 6-dimensional space of $(1, \sqrt{2}x_1, \sqrt{2}x_2, x_1^2, \sqrt{2}x_1x_2, x_2^2)$, computed with two multiplications and an addition.

9.6.3 Common kernels

- **Linear**: $K = \mathbf{x}^{\top} \mathbf{z}$ (no mapping). Best for high-dimensional sparse data like text.
- **Polynomial**: $K = (\gamma \mathbf{x}^{\top} \mathbf{z} + r)^d$: all interactions up to degree d .

- **RBF (Gaussian):** $K = \exp(-\gamma\|\mathbf{x} - \mathbf{z}\|^2)$: a similarity, 1 for identical points, decaying with distance. Corresponds to an infinite-dimensional feature space. The default non-linear choice.
- **Sigmoid:** $\tanh(\gamma\mathbf{x}^\top \mathbf{z} + r)$: not always a valid kernel.

9.6.4 What makes a valid kernel

K must equal a dot product in *some* space. Equivalently (Mercer's condition), for any set of points the **Gram matrix** $[K(\mathbf{x}_i, \mathbf{x}_j)]$ must be symmetric and positive semi-definite. Sums, positive multiples and products of valid kernels are valid, so kernels can be built from pieces (e.g. a text kernel plus a numeric kernel).

9.6.5 The RBF γ

γ sets how far each support vector's influence reaches. Large γ : tiny bumps around each point, wiggly boundary, overfitting (in the limit, memorisation). Small γ : smooth, nearly linear boundary, underfitting. Tune C and γ together on a log grid (e.g. $C \in \{0.1, 1, 10, 100\}$, $\gamma \in \{0.001, 0.01, 0.1, 1\}$).

9.7 Practicalities

- **Scale features:** the margin is a distance; an unscaled feature in rupees dominates one in years.
- **Cost:** kernel SVM training is between $O(n^2)$ and $O(n^3)$ and the kernel matrix needs $O(n^2)$ memory: fine for 10^4 , painful beyond about 10^5 rows. Linear SVMs (liblinear, SGD with hinge loss) scale to millions. Approximate kernels (random Fourier features, Nyström) make RBF-like models scalable.
- **Prediction cost:** proportional to the number of support vectors (for kernel SVMs).
- **Multiclass:** SVMs are binary; scikit-learn's SVC uses one-vs-one, LinearSVC one-vs-rest.
- **Probabilities:** SVM outputs signed distances; Platt scaling fits a sigmoid on cross-validated scores (probability=True, slower).
- **Imbalance:** class_weight scales C per class.
- **SVR:** regression with the ϵ -insensitive tube (Chapter 4).
- **One-class SVM:** novelty detection, a boundary around "normal" data.

In practice: SVMs and kernels at InfoEdge

Linear SVMs on TF-IDF features were the best text classifiers before deep learning and remain excellent fast baselines for classifying job titles or resumes into the taxonomy (500 categories, millions of rows: LinearSVC trains in seconds). The kernel idea, "compare items by a similarity instead of explicit features", reappears in attention (softmax of dot products) and in the Siamese networks on the Tools slide, which learn a similarity between a job and a CV.

9.8 Interview questions

Q9.1

★ **PYQ (INT: what is SVM; hinge loss, kernels, margin, support vectors)** Explain SVMs in the order the interviewer wants: margin, support vectors, hinge loss, kernels.

🕒 **In one line.** An SVM finds the separating hyperplane with the widest margin; the solution depends only on the points on or inside the margin (support vectors); with slack it is equivalent to minimising hinge loss plus an L2 penalty; and because it only needs dot products, kernels let it learn non-linear boundaries.

⚙️ **The full answer.** **Margin.** Among all separating hyperplanes choose the one farthest from the nearest points; scaled so the nearest satisfy $y(\mathbf{w}^\top \mathbf{x} + b) = 1$, the margin is $2/\|\mathbf{w}\|$, so we minimise $\frac{1}{2}\|\mathbf{w}\|^2$ subject to all points being on the correct side with functional margin ≥ 1 . **Support vectors.** In the dual, $\mathbf{w} = \sum \alpha_i y_i \mathbf{x}_i$ with $\alpha_i > 0$ only for points on or inside the margin. **Hinge loss.** Soft margin: $\min \frac{1}{2}\|\mathbf{w}\|^2 + C \sum \max(0, 1 - y_i f(\mathbf{x}_i))$. C trades margin width for training errors. **Kernels.** The dual and the prediction use only $\mathbf{x}_i^\top \mathbf{x}_j$; replacing with $K(\mathbf{x}_i, \mathbf{x}_j) = \phi(\mathbf{x}_i)^\top \phi(\mathbf{x}_j)$ fits a linear separator in a high (even infinite) dimensional space without computing ϕ .

📊 **Worked example.** RBF SVM on a “ring” dataset: linear SVM 52% accuracy; RBF with $C = 1$, $\gamma = 0.5$: 99%.

💡 **If the interviewer pushes.** “Why maximise the margin?” Larger margins mean the classifier is less sensitive to small perturbations of the data and, by statistical learning theory, generalisation bounds depend on margin rather than dimension, which is why SVMs work in very high dimensions.

Q9.2

★ **PYQ (OA: hinge loss)** Is hinge loss a classification or regression loss? Compute it for $(y, f) = (+1, 0.3), (+1, -2), (-1, -1.5)$.

🕒 **In one line.** Classification; the values are 0.7, 3 and 0.

⚙️ **The full answer.** $\max(0, 1 - yf)$: $1 - 0.3 = 0.7$; $1 + 2 = 3$; $yf = 1.5$, so 0. It is a convex upper bound on the 0–1 loss.

Q9.3

★ **PYQ (INT: kernels project to higher dimensions)** What does the kernel trick do, exactly, and why is it a “trick”?

🕒 **In one line.** It lets an algorithm that only uses dot products operate in a high-dimensional feature space by computing those dot products directly with a kernel function, without ever constructing the high-dimensional vectors.

⚙️ **The full answer.** “Trick” because the cost stays that of computing K in the original space. Example: $(\mathbf{x}^\top \mathbf{z})^2$ equals a dot product of degree-2 features. For RBF, the implicit space is infinite-dimensional, so explicit mapping is impossible, yet $K = e^{-\gamma\|\mathbf{x}-\mathbf{z}\|^2}$ is cheap. Any dot-product-only

algorithm can be kernelised: kernel ridge regression, kernel PCA, kernel k -means, Gaussian processes.

Q9.4

Derive why the SVM margin is $2/\|\mathbf{w}\|$.

🎯 **In one line.** After scaling so the closest points satisfy $|\mathbf{w}^\top \mathbf{x} + b| = 1$, their distance to the plane is $1/\|\mathbf{w}\|$, and the street spans both sides, so its width is $2/\|\mathbf{w}\|$.

⚙️ **The full answer.** Distance of $\mathbf{x}$ to the plane is $|\mathbf{w}^\top \mathbf{x} + b|/\|\mathbf{w}\|$ because $\mathbf{w}/\|\mathbf{w}\|$ is the unit normal. For $\mathbf{w} = (3, 4)$: margin $2/5 = 0.4$.

Q9.5

What does C control? What happens as $C \rightarrow \infty$ and as $C \rightarrow 0$?

🎯 **In one line.** C is the price of margin violations: large C gives a narrow margin that fits training data closely (low bias, high variance), small C a wide margin that tolerates errors (high bias, low variance); $C \rightarrow \infty$ is the hard margin, $C \rightarrow 0$ ignores the data.

⚙️ **The full answer.** In hinge-loss form, C is like $1/\lambda$. With outliers, a large C bends the boundary to accommodate them.

Q9.6

An RBF SVM gets 100% training accuracy and poor test accuracy. What do you change?

🎯 **In one line.** It is overfitting: decrease γ (smoother kernel) and/or decrease C (wider margin), tuned by cross-validation; also check scaling.

⚙️ **The full answer.** Large γ makes each support vector influence only its immediate neighbourhood, so the model memorises. The number of support vectors approaching n is a symptom.

Q9.7

Logistic regression vs linear SVM: which and when?

🎯 **In one line.** Both are linear classifiers with L2 regularisation differing only in loss; choose logistic regression for calibrated probabilities and interpretability, linear SVM when you want a max-margin boundary that ignores well-classified points, and either for text classification with similar accuracy.

⚙️ **The full answer.** Log loss is smooth and never zero; hinge is zero beyond the margin and has a kink. SVM gives sparse dual solutions; logistic gives probabilities. With kernels, SVM is the natural choice.

Q9.8

Why must features be scaled for SVMs?

🎯 **In one line.** Because the margin and the RBF kernel are based on Euclidean distances, so a feature with a large numeric range dominates the geometry regardless of its importance.

🔧 **Worked example.** Salary in rupees (range 10^6) and experience in years (range 30): $\|\mathbf{x} - \mathbf{z}\|^2$ is essentially only the salary difference.

Q9.9

🧠 **Think deeper** Why do SVMs work well in very high dimensions with few samples (e.g. text)?

🎯 **In one line.** Because their capacity is controlled by the margin, not by the number of features, and in high dimensions data are often nearly linearly separable with a good margin; the solution lives in the span of the support vectors, at most n of them.

🔧 **The full answer.** Generalisation bounds for margin classifiers scale with R^2/ρ^2 (data radius over margin squared), independent of dimension. TF-IDF vectors of documents with 100,000 dimensions are sparse and nearly separable; a linear SVM finds a robust boundary.

Q9.10

🧠 **Think deeper** Can you use the kernel trick with logistic regression, k -means or PCA?

🎯 **In one line.** Yes: any method that can be written purely in terms of dot products can be kernelised, giving kernel logistic regression, kernel k -means and kernel PCA; the catch is that without hinge loss's sparsity, all n points are needed at prediction time.

🔧 **The full answer.** Kernel logistic regression: $f(\mathbf{x}) = \sum_i \alpha_i K(\mathbf{x}_i, \mathbf{x})$ with log loss; every $\alpha_i \neq 0$, so prediction costs $O(n)$ and training $O(n^3)$. Kernel PCA: eigendecompose the centred kernel matrix to get non-linear components. Kernel k -means: distances in feature space via $K(\mathbf{x}, \mathbf{x}) - 2K(\mathbf{x}, \mathbf{c}) + \dots$; equivalent to spectral clustering in some forms.

Q9.11

🧠 **Think deeper** Is a neural network with a final linear layer “just a learned kernel”?

🎯 **In one line.** In a sense yes: the hidden layers learn a feature map ϕ , and the last layer is a linear model on $\phi(\mathbf{x})$, so the network learns its kernel $K(\mathbf{x}, \mathbf{z}) = \phi(\mathbf{x})^\top \phi(\mathbf{z})$ from data instead of fixing it in advance.

🔧 **The full answer.** An RBF SVM fixes the similarity; a network adapts it to the task. In the infinite-width limit, a randomly initialised network trained by gradient descent behaves like a kernel method with the “neural tangent kernel”. Siamese networks make this explicit: they learn an embedding where dot products measure job–CV similarity.

Q9.12

Compute the polynomial kernel $(\mathbf{x}^\top \mathbf{z} + 1)^2$ for $\mathbf{x} = (1, 2)$, $\mathbf{z} = (3, -1)$ and the RBF kernel with $\gamma = 0.5$ for $\mathbf{x} = (0, 0)$, $\mathbf{z} = (1, 1)$.

🕒 **In one line.** Polynomial: $(1 + 1)^2 = 4$; RBF: $e^{-0.5 \cdot 2} = e^{-1} = 0.368$.

Q9.13

What are the KKT conditions telling us about which points are support vectors?

🕒 **In one line.** Complementary slackness says $\alpha_i[y_i f(\mathbf{x}_i) - 1 + \xi_i] = 0$, so $\alpha_i > 0$ only when a point's constraint is tight (on the margin) or violated; points safely outside the margin have $\alpha_i = 0$.

⚙️ **The full answer.** Also $(C - \alpha_i)\xi_i = 0$: points with slack ($\xi_i > 0$) have $\alpha_i = C$. So: $\alpha = 0$ outside, $0 < \alpha < C$ on the margin, $\alpha = C$ inside or wrong.

Q9.14

What happens if you remove a non-support vector and retrain? A support vector?

🕒 **In one line.** Removing a non-support vector leaves the solution unchanged; removing a support vector can change it.

⚙️ **The full answer.** This is also why leave-one-out error of an SVM is bounded by (number of support vectors)/ n .

Q9.15

💡 **Think deeper** Your kernel SVM takes hours on 500,000 rows. What are your options?

🕒 **In one line.** Use a linear SVM if the data are high-dimensional, approximate the kernel with random Fourier features or Nyström and train a linear model on them, subsample, or switch to gradient boosting or a neural network.

⚙️ **The full answer.** Kernel SVM cost grows at least quadratically. Random Fourier features: $z(\mathbf{x}) = \sqrt{2/D} \cos(\Omega \mathbf{x} + \mathbf{b})$ with Ω drawn from a Gaussian approximates the RBF kernel with D features, so a linear SVM on z trains in linear time.

Q9.16

How do you get probabilities from an SVM, and why is it costly?

🕒 **In one line.** Platt scaling: fit a logistic regression mapping the SVM's decision values to probabilities, using cross-validated decision values so the calibration data were not used to train the SVM; the internal cross-validation makes training about five times slower.

Q9.17

 **Think deeper** When would you choose an SVM over gradient boosting today?

🎯 **In one line.** For small-to-medium data with many features and smooth or similarity-based structure (text with TF-IDF, bioinformatics, custom kernels on strings or graphs), for strong regularised baselines, and for one-class novelty detection; for large mixed tabular data, boosting usually wins.

⚙️ **The full answer.** Kernels let you encode domain similarity (string kernels for skill sequences). SVMs have few hyperparameters and convex training (reproducible). They lack native handling of categoricals and missing values and scale poorly.

Q9.18

How many support vectors would you expect for a well-separated vs a noisy dataset, and what does it indicate?

🎯 **In one line.** Few for well-separated data and many for noisy or overlapping data; a very high fraction suggests overlap, too small C or an overly complex kernel (overfitting with large γ).

Q9.19

What is a one-class SVM and where could InfoEdge use it?

🎯 **In one line.** It learns a boundary enclosing most of the “normal” data (in kernel space) and flags points outside as novel; useful for spotting unusual recruiter behaviour or anomalous job postings when labelled fraud is scarce.

⚙️ **The full answer.** Parameter ν upper-bounds the fraction of training points treated as outliers. Isolation Forest is often faster on large data.

CHAPTER 10

Three Classic Classifiers: k -NN, Naive Bayes and Fisher's LDA

► **Companion video for this chapter:** [StatQuest: Naive Bayes, Clearly Explained](#)

The InfoEdge hire was asked to build Naive Bayes up from Bayes' theorem "with intuition, not formulas", and to explain Fisher's linear discriminant ("push class means apart, shrink the variance within each class"). k -nearest neighbours is the simplest possible learner and a standard warm-up question. Each is a different philosophy: k -NN memorises, Naive Bayes models how data are generated, LDA finds the best projection.

10.1 k -nearest neighbours

10.1.1 The idea

To classify a new point, find the k training points closest to it and take a majority vote (or average their targets for regression). No training at all: the model *is* the data. This is called a **lazy** or **instance-based** learner.

Small example: 3-NN by hand

Training points (experience, skills matched) with labels "shortlisted" (S) or not (N): $A(1, 2)$ N, $B(2, 3)$ N, $C(5, 8)$ S, $D(6, 7)$ S, $E(4, 6)$ S. New candidate $Q = (4, 5)$. Euclidean distances: A : $\sqrt{9+9} = 4.24$; B : $\sqrt{4+4} = 2.83$; C : $\sqrt{1+9} = 3.16$; D : $\sqrt{4+4} = 2.83$; E : 1. Nearest three: E (S), B (N), D (S). Vote: S wins, 2 to 1.

10.1.2 Choosing k

- $k = 1$: the boundary follows every training point; zero training error; high variance (one mislabelled point creates an island).
- Large k : smooth boundary; high bias; with $k = n$ it always predicts the majority class.
- Choose k by cross-validation; odd k avoids ties in binary problems; distance-weighted voting ($1/d$) gives closer neighbours more say.

10.1.3 Distances and scaling

Euclidean $\sqrt{\sum (x_j - z_j)^2}$, Manhattan $\sum |x_j - z_j|$, cosine (for text and embeddings), Hamming (for categorical). **Scaling is essential**: a feature measured in rupees will dominate one measured in years.

10.1.4 The curse of dimensionality

In high dimensions almost all points are far from each other and distances become similar: the nearest neighbour is barely nearer than the farthest. Example: to capture 10% of uniformly spread data in 10

dimensions, a cube around the query must span $0.1^{1/10} = 79\%$ of each axis: “local” is no longer local. k -NN therefore needs dimensionality reduction or learned embeddings to work in high dimensions.

10.1.5 Cost

Training $O(1)$; prediction $O(nd)$ per query by brute force. Speed-ups: KD-trees and ball trees (good up to about 20 dimensions), and **approximate nearest neighbour** indexes (HNSW, IVF, product quantisation; libraries FAISS, ScaNN) for millions of embeddings. The PremiumX slide’s “vector index” is exactly this.

10.2 Naive Bayes, built from Bayes’ theorem

10.2.1 Bayes’ theorem in words

We want $P(\text{class} \mid \text{evidence})$. Bayes’ theorem turns it around:

$$P(C \mid x) = \frac{P(x \mid C) P(C)}{P(x)}.$$

- $P(C)$, the **prior**: how common the class is before looking at the evidence (2% of messages are spam).
- $P(x \mid C)$, the **likelihood**: how typical this evidence is for that class (“free offer” appears in 30% of spam and 1% of normal mail).
- $P(x)$, the **evidence**: the same for every class, so for choosing a class we can ignore it and compare $P(x \mid C)P(C)$ across classes.
- $P(C \mid x)$, the **posterior**: the updated belief.

10.2.2 The problem: too many combinations

The evidence is many features at once: thousands of words. Estimating $P(x_1, x_2, \dots, x_d \mid C)$ directly would need data for every combination of words, impossible.

10.2.3 The naive assumption

Assume the features are **conditionally independent given the class**: once you know it is spam, knowing it contains “free” tells you nothing extra about whether it contains “offer”. Then the joint likelihood factorises:

$$P(x_1, \dots, x_d \mid C) = \prod_{j=1}^d P(x_j \mid C), \quad \hat{C} = \arg \max_C P(C) \prod_j P(x_j \mid C).$$

Now each $P(x_j \mid C)$ is a simple count. The assumption is almost always false (“free” and “offer” co-occur), but the *ranking* of classes is often still right, which is all classification needs.

Small example: spam by hand

Prior: $P(S) = 0.3$, $P(H) = 0.7$. Word probabilities: $P(\text{free} \mid S) = 0.4$, $P(\text{free} \mid H) = 0.05$; $P(\text{meeting} \mid S) = 0.05$, $P(\text{meeting} \mid H) = 0.3$. Message: “free meeting”. Spam score: $0.3 \times 0.4 \times 0.05 = 0.006$. Ham score: $0.7 \times 0.05 \times 0.3 = 0.0105$. Normalise: $P(S \mid \text{msg}) = 0.006 / 0.0165 = 0.36$.

Classified as ham. “Meeting” outweighed “free”.

10.2.4 Variants

- **Multinomial NB**: features are word counts; $P(w | C) = \frac{\text{count}(w,C)+\alpha}{\text{total words in } C+\alpha|V|}$.
- **Bernoulli NB**: features are word presence/absence; also penalises the absence of words typical of a class.
- **Gaussian NB**: numeric features modelled as a normal distribution per class.
- **Complement NB**: estimates from the complement of each class; better with imbalanced text.

10.2.5 Laplace smoothing: the zero-frequency problem

If “lottery” never appeared in ham training mail, $P(\text{lottery} | H) = 0$, and one occurrence makes the whole ham product 0, whatever else the message says. Add α (usually 1) to every count:

Small example: smoothing

Spam class has 20 word tokens in total, vocabulary size 10, “offer” appears 3 times in spam. Unsmoothed: $3/20 = 0.15$. Laplace ($\alpha = 1$): $(3 + 1)/(20 + 10) = 4/30 = 0.133$. A word never seen in spam: $1/30 = 0.033$ instead of 0.

Smoothing is a MAP estimate with a Dirichlet (uniform-ish) prior (Chapter 12).

10.2.6 Log space

Multiplying hundreds of probabilities like 0.001 underflows to 0 in floating point. Use sums of logs: $\log P(C) + \sum_j \log P(x_j | C)$. The argmax is unchanged because log is increasing.

10.2.7 Why it works surprisingly well

Very few parameters (one per word per class), so low variance: excellent with little data and huge vocabularies; trains in one pass; predicts fast. Its probabilities are poorly calibrated (the independence assumption double-counts correlated evidence, so outputs are pushed toward 0 and 1), but its rankings are often good.

10.3 Fisher’s Linear Discriminant Analysis

10.3.1 The goal

Project data onto a single direction $\mathbf{w}$ (a line) such that, along that line, the two classes are as separated as possible. “Separated” means two things at once: the class **means** far apart, and each class **tightly bunched**.

Intuition: why means alone are not enough

Imagine two cigar-shaped clouds lying side by side, parallel. The direction joining their means may run along the cigars’ length, where each cloud is very spread out, so the projections overlap

heavily. A slightly rotated direction, across the cigars, separates them cleanly even though the means are a bit closer along it. Fisher's criterion finds that rotation.

10.3.2 The criterion

For a direction $\mathbf{w}$, the projected class means are $\tilde{m}_k = \mathbf{w}^\top \mathbf{m}_k$ and the projected within-class scatters are $\tilde{s}_k^2 = \sum_{i \in k} (\mathbf{w}^\top \mathbf{x}_i - \tilde{m}_k)^2$. Fisher maximises

$$J(\mathbf{w}) = \frac{(\tilde{m}_1 - \tilde{m}_2)^2}{\tilde{s}_1^2 + \tilde{s}_2^2} = \frac{\mathbf{w}^\top S_B \mathbf{w}}{\mathbf{w}^\top S_W \mathbf{w}},$$

with between-class scatter $S_B = (\mathbf{m}_1 - \mathbf{m}_2)(\mathbf{m}_1 - \mathbf{m}_2)^\top$ and within-class scatter $S_W = \sum_k \sum_{i \in k} (\mathbf{x}_i - \mathbf{m}_k)(\mathbf{x}_i - \mathbf{m}_k)^\top$. Setting the derivative to zero gives the solution

$$\mathbf{w} \propto S_W^{-1}(\mathbf{m}_1 - \mathbf{m}_2).$$

In words: start from the direction joining the means, then “undo” the shape of the clouds by multiplying by the inverse of the within-class scatter.

Small example: Fisher direction in 2-D

Class means $\mathbf{m}_1 = (2, 2)$, $\mathbf{m}_2 = (4, 3)$, so $\mathbf{m}_1 - \mathbf{m}_2 = (-2, -1)$. Suppose $S_W = \begin{pmatrix} 4 & 0 \\ 0 & 1 \end{pmatrix}$ (each class spread out along x_1). $S_W^{-1} = \text{diag}(0.25, 1)$. $\mathbf{w} \propto (-0.5, -1)$. Joining the means suggests weighting x_1 twice as much as x_2 ; Fisher weights x_2 twice as much as x_1 , because along x_1 the classes are noisy.

10.3.3 More than two classes, and LDA as a classifier

With K classes, LDA finds up to $K - 1$ directions: the top eigenvectors of $S_W^{-1} S_B$. It is a **supervised** dimensionality reduction, in contrast to PCA (unsupervised, maximises total variance, ignores labels). LDA is also a probabilistic classifier: assume each class is Gaussian with a *shared* covariance Σ ; then the Bayes decision boundary is linear, and its direction is exactly Fisher's $\Sigma^{-1}(\mathbf{m}_1 - \mathbf{m}_2)$. With class-specific covariances the boundary is quadratic: **QDA**.

10.3.4 Limits

Linear boundaries only; assumes roughly Gaussian classes with similar covariance for optimality; S_W must be invertible (regularise with shrinkage when d is large); at most $K - 1$ dimensions.

In practice: these three at InfoEdge

- **k -NN / vector search**: “similar jobs”, “similar candidates” and the dense half of PremiumX's hybrid retrieval are nearest-neighbour search over embeddings with an ANN index.
- **Naive Bayes**: a strong, instant baseline for classifying job titles, routing support tickets, or spam/abuse detection in Jeevansathi messages.
- **LDA**: a quick supervised projection to visualise how separable “placed” vs “not placed” students are, and a well-behaved linear classifier for small datasets.

10.4 Interview questions

Q10.1

★ **PYQ (INT: Naive Bayes from Bayes' theorem)** Build Naive Bayes from Bayes' theorem. What assumption does it make, why is it needed, and how is it used in practice?

🕒 **In one line.** Naive Bayes applies Bayes' theorem to get $P(\text{class} \mid \text{features}) \propto P(\text{class}) \prod_j P(x_j \mid \text{class})$, which requires assuming the features are independent given the class; that assumption turns an impossible joint estimate into simple per-feature counts.

⚙️ **The full answer.** Start: $P(C \mid \mathbf{x}) = P(\mathbf{x} \mid C)P(C)/P(\mathbf{x})$. The denominator is the same for all classes, so we compare numerators. The hard part is $P(\mathbf{x} \mid C)$ for a high-dimensional $\mathbf{x}$: with 10,000 binary word features there are 2^{10000} combinations, and we would need data for each. Conditional independence lets us write it as a product of 10,000 one-word probabilities, each estimated by counting. Note the independence is *conditional on the class*, not overall: “free” and “offer” are clearly correlated overall because both are spammy; the assumption is that within spam they occur independently. In practice: Multinomial NB on word counts with Laplace smoothing, computed in log space, as a fast baseline for text.

📊 **Worked example.** The “free meeting” example: spam score 0.006, ham 0.0105, posterior spam 0.36.

💡 **If the interviewer pushes.** “Why does it work when the assumption is false?” Classification needs only the correct argmax, not correct probabilities; errors from correlated features often affect all classes similarly. Also, its low variance beats more flexible models when data are scarce. “What about the probabilities?” Over-confident; calibrate if needed.

⚠️ **Trap to avoid.** Saying the features must be “independent” without “given the class”, or calling it a discriminative model (it is generative: it models $P(\mathbf{x} \mid C)P(C)$).

Q10.2

★ **PYQ (INT: Fisher's LDA)** Explain Fisher's linear discriminant: the formulation and the intuition.

🕒 **In one line.** Find a projection direction that maximises the squared distance between projected class means divided by the total within-class variance along that direction; the solution is $\mathbf{w} \propto S_W^{-1}(\mathbf{m}_1 - \mathbf{m}_2)$.

⚙️ **The full answer.** Projecting onto a line reduces each point to one number. A good line keeps the classes apart (numerator: between-class separation) and each class compact (denominator: within-class spread). Maximising the ratio $J(\mathbf{w}) = \frac{\mathbf{w}^T S_B \mathbf{w}}{\mathbf{w}^T S_W \mathbf{w}}$: since only the direction matters, set the gradient to zero, giving $S_B \mathbf{w} = \lambda S_W \mathbf{w}$, a generalised eigenproblem. Because $S_B \mathbf{w}$ always points along $\mathbf{m}_1 - \mathbf{m}_2$, the answer is $S_W^{-1}(\mathbf{m}_1 - \mathbf{m}_2)$. For a classifier, threshold the projection (e.g. at the midpoint of projected means, adjusted for priors).

📊 **Worked example.** See the 2-D example: means difference $(-2, -1)$ becomes Fisher direction $(-0.5, -1)$ after whitening by S_W^{-1} .

🗣️ **If the interviewer pushes.** “How is it different from PCA?” PCA ignores labels and keeps directions of maximum total variance, which may be useless for classification (the cigars example); LDA uses labels. “How many components?” At most $K - 1$. “When does it fail?” Non-linear class structure, multimodal classes, very unequal covariances (use QDA), $d > n$ (use shrinkage).

Q10.3

How does k -NN work, and how do you choose k ?

🕒 **In one line.** It predicts by majority vote (or average) of the k closest training points; small k gives low bias and high variance, large k the reverse, so choose k by cross-validation.

⚙️ **The full answer.** Worked example above: 3-NN classifies Q as S . With $k = 1$ the decision boundary is the Voronoi tessellation of training points.

🗣️ **If the interviewer pushes.** “Training and prediction cost?” Training free; prediction $O(nd)$ brute force; use KD-trees or ANN indexes.

Q10.4

🧠 **Think deeper** Why does k -NN suffer from the curse of dimensionality? Give a number.

🕒 **In one line.** In high dimensions volumes grow so fast that a neighbourhood holding a fixed fraction of data must span most of each axis, and all pairwise distances become nearly equal, so “nearest” neighbours are not similar.

⚙️ **The full answer.** To contain 1% of uniform data in $d = 100$ dimensions, a sub-cube needs side $0.01^{1/100} = 0.955$ of the full range. Also the ratio (farthest – nearest)/nearest distance tends to 0 as d grows for many distributions. Remedies: feature selection, PCA, learned embeddings (where the meaningful structure is low-dimensional), and metric learning.

Q10.5

Why must features be scaled for k -NN, and does it matter for Naive Bayes?

🕒 **In one line.** k -NN uses distances, so unscaled features with large ranges dominate; Naive Bayes models each feature separately per class, so scaling does not change its decisions (Gaussian NB just rescales its means and variances).

Q10.6

Compute the Laplace-smoothed probability of “offer” in spam if spam has 20 tokens, vocabulary size 10, and “offer” appears 3 times. Why is smoothing needed?

🕒 **In one line.** $(3 + 1)/(20 + 10) = 0.133$; without smoothing, a word never seen in a class would give probability 0 and wipe out the whole product.

Q10.7

Why do we compute Naive Bayes in log space?

🕒 **In one line.** Products of many small probabilities underflow to 0 in floating point; summing logs is numerically stable and preserves the argmax.

🧮 **Worked example.** 200 words each with probability 0.001: product 10^{-600} , below double precision's 10^{-308} ; log sum -1381.6 is fine.

Q10.8

🧠 Think deeper Naive Bayes vs logistic regression: they are both linear for text. What is the difference?

🕒 **In one line.** Multinomial Naive Bayes gives a linear decision rule in word counts, just like logistic regression, but NB estimates the weights generatively from class-conditional counts (fast, low variance, good with little data) while logistic regression fits them discriminatively to maximise conditional likelihood (better with more data, handles correlated features).

⚙️ **The full answer.** NB's log-odds = $\log \frac{P(S)}{P(H)} + \sum_j x_j \log \frac{P(w_j|S)}{P(w_j|H)}$: linear in counts x_j . Ng and Jordan showed NB reaches its (higher) asymptotic error faster; logistic regression reaches a lower error with enough data. Generative–discriminative pair.

Q10.9

🧠 Think deeper If two features are exact duplicates, what happens to Naive Bayes? To logistic regression?

🕒 **In one line.** Naive Bayes counts the same evidence twice, so its posterior becomes more extreme (over-confident); logistic regression splits the weight between the duplicates and its predictions are unchanged.

⚙️ **The full answer.** NB multiplies $P(x_1 | C)P(x_1 | C)$, squaring the likelihood ratio. Logistic regression's predictions depend on $w_1 + w_2$, which the fit learns correctly (with L2 it splits evenly). This is the clearest demonstration of what the naive assumption costs.

Q10.10

What is the difference between generative and discriminative classifiers? Classify NB, LDA, logistic regression, SVM, k -NN.

🕒 **In one line.** Generative models learn how data are produced, $P(\mathbf{x} | y)P(y)$ (Naive Bayes, LDA/QDA, GMM classifiers, GANs and VAEs as generators); discriminative models learn $P(y | \mathbf{x})$ or a boundary directly (logistic regression, SVM, trees, neural nets); k -NN is a non-parametric discriminative method.

⚙️ **The full answer.** Generative models can generate samples, handle missing features naturally, and do well with little data; discriminative models usually win on accuracy with enough data.

Q10.11

LDA vs PCA: when would PCA's first component be useless for classification?

🕒 **In one line.** When the direction of largest total variance is unrelated to the class difference, for example two parallel elongated clusters separated across their narrow direction.

⚙️ **The full answer.** PCA picks the long axis (where both clusters spread); projecting on it mixes the classes completely. LDA picks the narrow axis.

Q10.12

🧠 **Think deeper** Under what assumptions is LDA the Bayes-optimal classifier, and what changes with QDA?

🕒 **In one line.** If each class is Gaussian with the same covariance, the log-posterior ratio is linear in $\mathbf{x}$ and LDA is optimal; with different covariances the quadratic terms no longer cancel and the optimal boundary is quadratic (QDA).

⚙️ **The full answer.** $\log \frac{P(1|\mathbf{x})}{P(2|\mathbf{x})} = \mathbf{x}^\top \Sigma^{-1}(\boldsymbol{\mu}_1 - \boldsymbol{\mu}_2) + \text{const}$ when the $\mathbf{x}^\top \Sigma^{-1} \mathbf{x}$ terms cancel. QDA needs a covariance per class: $Kd(d+1)/2$ parameters, more variance; regularised DA interpolates.

Q10.13

How would you speed up nearest-neighbour search over 50 lakh job embeddings?

🕒 **In one line.** Use an approximate nearest-neighbour index such as HNSW or IVF with product quantisation (FAISS/ScaNN), trading a little recall for orders-of-magnitude speed, and pre-filter by hard constraints like location.

⚙️ **The full answer.** HNSW builds a navigable small-world graph; search hops greedily toward the query in $O(\log n)$. IVF clusters vectors and searches only the nearest clusters; PQ compresses vectors to bytes. Measure recall@k against brute force on a sample.

Q10.14

🧠 **Think deeper** Can k -NN be used for regression and for imputation? What are the risks?

🕒 **In one line.** Yes: average the neighbours' targets for regression, and fill a missing value with the neighbours' values for imputation; risks are scaling, irrelevant features dominating distances, high dimensions, and leakage if the imputer is fitted on test data.

📊 **Worked example.** Neighbours with value 10 at distance 1 and 20 at distance 2, weights $1/d$: $(10 + 10)/(1 + 0.5) = 13.3$.

Q10.15

Gaussian Naive Bayes: how would it classify a candidate with test score 70 if passers have mean 75, sd 5 and failers mean 60, sd 10, with equal priors?

🎯 **In one line.** Compare normal densities: passers $\phi(-1)/5 = 0.0484$, failers $\phi(1)/10 = 0.0242$; predict “pass” with posterior 0.667.

⚙️ **The full answer.** $\phi(1) = 0.242$. Passers: $z = (70 - 75)/5 = -1$, density $0.242/5 = 0.0484$. Failers: $z = 1$, $0.242/10 = 0.0242$. Ratio 2:1.

Q10.16

🧠 **Think deeper** Is k -NN with $k = 1$ ever a good model?

🎯 **In one line.** Yes, when classes are well separated, labels are clean and the distance is meaningful (e.g. near-duplicate detection, or embeddings trained for similarity); its asymptotic error is at most twice the Bayes error.

⚙️ **The full answer.** The Cover–Hart result: as $n \rightarrow \infty$, 1-NN error $\leq 2 \times$ Bayes error. With good embeddings, 1-NN retrieval is the basis of duplicate-resume detection and “find the canonical skill for this raw string” in taxonomy normalisation.

CHAPTER 11

Evaluation Metrics and the Confusion Matrix

🎥 **Companion video for this chapter:** [StatQuest: The Confusion Matrix](#)

At the 2026–27 pre-placement talk, the panel used a spam example to explain what they look for: “do you look at recall, not just accuracy; can you create the right dataframe; do you frame it properly as classification or regression”. The Round 3 task asked for a confusion matrix, precision, recall, F1 and accuracy. The OA asked what a false negative means. Metrics are where business sense meets ML, so interviewers use them to judge maturity.

11.1 The confusion matrix

For a binary classifier, every prediction falls into one of four boxes. Call the class we care about (spam, fraud, “will buy”) **positive**.

	Predicted positive	Predicted negative
Actually positive	TP (true positive)	FN (false negative)
Actually negative	FP (false positive)	TN (true negative)

Read the names as: *was the prediction right (True/False) + what did we predict (Positive/Negative)*. A **false negative** is a positive we called negative: a spam message that reached the inbox, a fraud we missed, a high-intent buyer we did not call. A **false positive** is a false alarm: a genuine mail sent to spam.

In statistics language: FP is a **Type I error** (rejecting a true null “this is normal”), FN is a **Type II error**.

11.2 The metrics, each answering one question

Running example: a spam filter on 1,000 messages, 100 actually spam. It flags 90 messages, of which 80 are spam. So TP = 80, FP = 10, FN = 20, TN = 890.

Metric	Formula	Question it answers	Example
Accuracy	$\frac{TP+TN}{\text{all}}$	What fraction of all predictions were right?	970/1000 = 0.97
Precision	$\frac{TP}{TP+FP}$	Of what I flagged, how much was really positive?	80/90 = 0.889
Recall (sensitivity, TPR)	$\frac{TP}{TP+FN}$	Of all real positives, how many did I catch?	80/100 = 0.80
Specificity (TNR)	$\frac{TN}{TN+FP}$	Of all real negatives, how many did I leave alone?	890/900 = 0.989
FPR	$\frac{FP}{FP+TN} = 1 - \text{spec.}$	How often do negatives raise a false alarm?	10/900 = 0.011
F1	$\frac{2PR}{P+R}$	Single balance of precision and recall	0.842
NPV	$\frac{TN}{TN+FN}$	Of what I cleared, how much was really negative?	890/910 = 0.978

11.2.1 Why accuracy misleads on imbalanced data

If only 1% of messages are spam, a model that says “not spam” to everything is 99% accurate and completely useless: recall 0. With rare positives, accuracy is dominated by the easy negatives. This is exactly the point the panel made.

11.2.2 Precision vs recall: which matters?

- **Recall matters** when missing a positive is costly: disease screening, fraud, safety, and (for the panel’s spam example) letting dangerous phishing through. Also lead generation: a missed high-intent buyer is lost revenue.
- **Precision matters** when false alarms are costly: moving a recruiter’s genuine mail to spam, auto-banning a user, sending a telecaller to a cold lead (caller time is limited).
- Usually both matter, and the threshold sets the balance.

11.2.3 F1 and F-beta

F1 is the **harmonic mean** of precision and recall. The harmonic mean is dragged down by the smaller value, so a model cannot score well by being excellent at one and terrible at the other.

Small example: harmonic vs arithmetic

$P = 0.9, R = 0.1$: arithmetic mean 0.5 (looks fine), $F1 = 2(0.09)/1.0 = 0.18$ (reveals the problem).

$F_\beta = (1 + \beta^2) \frac{PR}{\beta^2 P + R}$ weights recall β times as much as precision. F_2 for recall-heavy tasks, $F_{0.5}$ for precision-heavy.

Small example: F-beta

$P = 0.5, R = 0.8$: $F_2 = \frac{5 \times 0.4}{4 \times 0.5 + 0.8} = \frac{2}{2.8} = 0.714$. $F_{0.5} = \frac{1.25 \times 0.4}{0.125 + 0.8} = 0.541$.

11.2.4 Matthews correlation coefficient and Cohen’s kappa

$MCC = \frac{TP \cdot TN - FP \cdot FN}{\sqrt{(TP+FP)(TP+FN)(TN+FP)(TN+FN)}}$ uses all four cells, ranges from -1 to 1 , and stays informative under imbalance. Kappa compares accuracy with the accuracy expected by chance agreement.

11.3 Thresholds and threshold-free metrics

11.3.1 The threshold moves the trade-off

A probabilistic classifier outputs a score; we choose a cut-off t . Raising t : fewer positives predicted; precision usually up, recall down. Lowering t : the opposite. The model did not change, only the decision.

11.3.2 ROC curve and AUC

Plot TPR (recall) against FPR for every threshold. A random model is the diagonal; a perfect model goes straight up to (0, 1). The **area under the ROC curve (AUC)** has a beautiful meaning:

$$\text{AUC} = P(\text{a random positive gets a higher score than a random negative}).$$

Small example: AUC by counting pairs

Scores: positives 0.8, 0.4; negatives 0.5, 0.3, 0.1. Pairs: 6. Positive 0.8 beats all 3 negatives; positive 0.4 beats 0.3 and 0.1 but not 0.5. Correctly ordered pairs = 5; $\text{AUC} = 5/6 = 0.833$. Ties count half.

AUC is threshold-free and insensitive to class balance (both axes are rates within a class), and invariant to any monotone transform of scores (it only looks at order).

11.3.3 Precision–recall curve

Plot precision against recall over thresholds. With rare positives it is far more revealing than ROC: a model can have ROC AUC 0.95 while precision is terrible, because a small FPR on a huge negative class is still many false positives. The baseline of the PR curve is the positive rate (a random model's precision), not 0.5.

Small example: why ROC can look great while precision is poor

1,000,000 negatives, 1,000 positives. TPR 0.9 at FPR 0.01 looks excellent on ROC. But $\text{FP} = 10,000$ and $\text{TP} = 900$: $\text{precision} = 900/10,900 = 8\%$.

11.3.4 Precision from sensitivity, specificity and prevalence

Bayes' theorem links them: $\text{precision} = \frac{\text{sens} \cdot \pi}{\text{sens} \cdot \pi + (1 - \text{spec})(1 - \pi)}$.

Small example: a "95% accurate" fraud test

Sensitivity 0.9, specificity 0.95, prevalence 2%: $\frac{0.018}{0.018 + 0.049} = 0.269$. Only 27% of flags are real fraud.

11.4 Multiclass averaging

Compute per-class precision/recall/F1 by treating each class as positive in turn, then:

- **Macro average:** plain mean over classes. Every class counts equally, so rare classes matter.
- **Weighted average:** mean weighted by class frequency (support).
- **Micro average:** pool all TP, FP, FN, then compute. For single-label multiclass, micro F1 = accuracy.

11.5 Calibration metrics

Brier score = $\frac{1}{n} \sum (p_i - y_i)^2$ (lower is better). **Log loss.** **Reliability diagram:** bucket predictions (0–0.1, 0.1–0.2, ...) and plot mean prediction vs actual positive rate; a calibrated model lies on the diagonal. **Expected calibration error** averages the gaps.

11.6 Regression and ranking metrics (summary)

Regression: MAE, RMSE, MAPE, R^2 (Chapter 2); RMSLE for skewed targets; pinball loss for quantiles.
 Ranking: **Precision@k**, **Recall@k**, **MRR** (mean of $1/\text{rank}$ of the first relevant result), **MAP**, **NDCG@k**:
 $\text{DCG@k} = \sum_{i=1}^k \frac{2^{rel_i} - 1}{\log_2(i+1)}$, divided by the ideal DCG.

Small example: MRR

First relevant result at ranks 1, 3, 2 for three queries: $\text{MRR} = (1 + 1/3 + 1/2)/3 = 0.611$.

11.7 Offline metrics vs business metrics

Offline metrics are proxies. The real question is the business outcome: applies per job, recruiter response rate, calls converted per caller-hour, time to hire. A model that wins offline should be confirmed with an **online A/B test**; guardrail metrics (complaints, unsubscribes, latency) must not degrade.

In practice: picking the metric at InfoEdge

- **Spam/abuse**: recall at a fixed high precision (or the reverse), never accuracy.
- **Telecalling** (99acres): callers can make, say, 200 calls/day: precision@200 and conversions per call.
- **Resdex search**: NDCG@10 or recruiter actions (views, contacts) on the first page.
- **AI REX screening agent**: recall of genuinely fit candidates (do not drop good people) with a precision floor (recruiter time).

11.8 Interview questions

Q11.1

★ **PYQ (OA: what a false negative means)** In a fraud model where positive means fraud, what is a false negative, and which error type is it?

🎯 **In one line.** A fraudulent transaction that the model called legitimate, i.e. a missed fraud; it is a Type II error.

⚙️ **The full answer.** “False” because the prediction was wrong, “negative” because the prediction was negative. The null hypothesis is “this transaction is normal”; failing to reject it when it is false is Type II. Its rate is $1 - \text{recall}$.

Q11.2

★ **PYQ (Round 3 coding round)** Given $\text{TP} = 40$, $\text{FP} = 10$, $\text{FN} = 20$, $\text{TN} = 130$, compute accuracy, precision, recall, F1 and specificity, and say which you would report for a placement model.

🎯 **In one line.** Accuracy 0.85, precision 0.80, recall 0.667, F1 0.727, specificity 0.929; I would report the full confusion matrix plus precision, recall and F1 for the class the decision is about, and ROC AUC.

⚙️ **The full answer.** Total 200; accuracy 170/200. Precision 40/50. Recall 40/60. $F1 = 2(0.8)(0.667)/1.467 = 0.727$. Specificity 130/140. For a placement model, which error matters depends on use: if it is used to target students for extra coaching, recall of the “not placed” class matters most (do not miss students who need help), so I would treat “not placed” as the positive class.

💡 **If the interviewer pushes.** “MCC?” $(40 \cdot 130 - 10 \cdot 20) / \sqrt{50 \cdot 60 \cdot 140 \cdot 150} = 5000/7937 = 0.63$.

Q11.3

★ **PYQ (2026–27 talk: the spam example)** You are building a spam classifier. Walk me through framing, data and evaluation; why is accuracy not enough?

🎯 **In one line.** It is binary classification with rare positives, so I would build a labelled message dataframe with features from text and sender behaviour, split by time, and evaluate with precision, recall, PR AUC and the confusion matrix at a chosen threshold, because accuracy is dominated by the majority class.

⚙️ **The full answer.** **Framing.** Target: spam (1) vs ham (0). Not regression (there is no quantity to predict), unless we predict a “spamminess” score, which is still a classifier’s probability. **Dataframe.** One row per message: text, sender ID, sender account age, messages sent in the last hour, fraction of recipients who replied, links count, label. Avoid leakage: features must be computable at send time (“number of users who reported this message” is a post-hoc signal, fine only if the model runs later). **Split.** By time (spam evolves), and by sender (so the model is not just memorising known spammers). **Model.** TF-IDF + logistic regression baseline; then GBM on combined text and behaviour features. **Metrics.** If 2% is spam, an “all ham” model is 98% accurate. Report recall (spam caught) and precision (genuine mail not lost) at the operating threshold, and PR AUC. Choose the threshold with business costs: losing a recruiter’s interview invite to the spam folder is worse than letting one promotional spam through, so precision of the spam label must be high; for phishing, recall must be high. Possibly two thresholds: block above 0.98, spam-folder above 0.8.

💡 **If the interviewer pushes.** “Recall went up but users complain. Why?” Precision fell: more genuine mails are being hidden. Check FP examples by segment (new recruiters with templated messages look spammy).

Q11.4

★ **PYQ (INT: imbalanced data)** 1% of your data is positive. How do you train and evaluate?

🎯 **In one line.** Evaluate with precision, recall, F1, PR AUC (and ROC AUC) on a stratified, untouched test set; train with class weights or resampling (undersample, oversample, SMOTE) only on training folds, and tune the threshold.

⚙️ **The full answer.** Class weights multiply the loss for positives (scikit-learn `class_weight="balanced": weight  $n/(2n_c)$`). Undersampling negatives speeds training; over-sampling/SMOTE add positives. Resampling distorts probabilities (Chapter 3 correction). Never resample the test set. Often the simplest effective approach is a well-tuned GBM with `scale_pos_weight` and a tuned threshold. Also consider anomaly detection if positives are extremely rare and diverse.

Q11.5

Compute ROC AUC for positive scores 0.8, 0.4 and negative scores 0.5, 0.3, 0.1, and explain what AUC means.

🎯 **In one line.** $AUC = 5/6 = 0.833$; it is the probability that a randomly chosen positive is scored above a randomly chosen negative.

⚙️ **The full answer.** Counting pairs as shown. Equivalent to the normalised Mann–Whitney U statistic. 0.5 is random; below 0.5 means the scores are inverted.

Q11.6

💡 **Think deeper** A model has ROC AUC 0.95 but PR AUC 0.08 with 0.1% positives. Is it good?

🎯 **In one line.** It ranks positives above negatives very well on average, but because negatives are so numerous, even a tiny false-positive rate produces many false alarms, so precision at useful recall is poor; whether that is “good” depends on whether the use case can tolerate low precision.

⚙️ **The full answer.** PR AUC’s baseline is 0.001, so 0.08 is 80 times better than random, yet only about 8% of flags would be real at typical operating points. For human review queues with capacity limits, report precision@k. For a retrieval stage feeding a better second-stage model, high recall with modest precision is fine.

Q11.7

Why is F1 a harmonic mean rather than an arithmetic mean?

🎯 **In one line.** Because the harmonic mean is dominated by the smaller of precision and recall, so a model must be good at both; also F1 is naturally a harmonic mean because $1/F1$ is the average of $1/P$ and $1/R$, which treats precision and recall as rates over different denominators.

⚙️ **The full answer.** $P = 0.9, R = 0.1$: arithmetic 0.5, F1 0.18. Another view: $F1 = \frac{2TP}{2TP+FP+FN}$: it ignores TN entirely, which is appropriate when negatives are abundant and uninteresting.

Q11.8

💡 **Think deeper** What metric would you optimise for a telecalling team that can make only 100 calls a day out of 50,000 leads?

🎯 **In one line.** Precision@100 (or expected conversions in the top 100), because the operational decision is which 100 to call; overall AUC or recall over all leads matters much less.

⚙️ **The full answer.** A model can have a lower AUC but better ordering at the very top. Evaluate on a time-based holdout: of the top 100 scored each day, how many converted? Also check calibration so you can compute expected revenue per call and staff accordingly. Over time, run an A/B test against the current rule-based list.

Q11.9

When would you prefer specificity over sensitivity? Give a job-portal example.

🎯 **In one line.** When false positives are expensive: for example auto-hiding recruiter accounts flagged as fraudulent, where blocking a genuine paying recruiter is very costly, so we need high specificity (few false alarms).

Q11.10

Compute precision for a test with sensitivity 0.9, specificity 0.95, prevalence 2%.

🎯 **In one line.** About 0.269: $0.9 \times 0.02 / (0.018 + 0.05 \times 0.98) = 0.018 / 0.067$.

⚙️ **The full answer.** Base-rate neglect: with rare positives, even good tests mostly flag negatives.

Q11.11

Macro vs micro vs weighted F1: which would you report for classifying jobs into 30 functions where 5 functions have 70% of jobs?

🎯 **In one line.** Report macro F1 if every function matters equally (rare functions should not be ignored), weighted or micro F1 if overall volume matters; I would show both plus the per-class table.

⚙️ **The full answer.** Micro F1 equals accuracy for single-label problems and is dominated by the big classes. Macro reveals poor performance on small classes.

Q11.12

💡 **Think deeper** Your offline AUC improved from 0.80 to 0.83, but the A/B test shows no change in applies. Why might that be?

🎯 **In one line.** The offline metric may not reflect what users see (only the top few positions matter), the improvement may be in parts of the ranking nobody views, there may be a train-serve mismatch or feedback loop, or the test may be underpowered.

⚙️ **The full answer.** Check NDCG@5 instead of AUC; check that features at serving match training ("train-serve skew"); position bias in logged data (users click top results regardless) can mislead offline evaluation; compute the minimum detectable effect of the A/B test.

Q11.13

What does a reliability curve below the diagonal mean?

🕒 **In one line.** The model is over-confident: when it predicts, say, 0.8, the actual positive rate is lower, like 0.6.

⚙️ **The full answer.** Fix with Platt or isotonic calibration on held-out data; check after resampling or class weighting, which commonly cause it.

Q11.14

Compute NDCG@3 for relevances (3, 2, 0) in the returned order when the ideal order is (3, 2, 1).

🕒 **In one line.** $DCG = 7 + 3/1.585 + 0 = 8.893$; ideal $DCG = 7 + 1.893 + 1/2 = 9.393$; $NDCG = 8.893/9.393 = 0.947$.

⚙️ **The full answer.** Gains $2^{rel} - 1$: 7, 3, 0 vs 7, 3, 1; discounts $1/\log_2(i + 1)$: 1, 0.631, 0.5.

Q11.15

💡 **Think deeper** Can a model with lower accuracy be better? Give an example.

🕒 **In one line.** Yes, whenever errors have different costs or classes are imbalanced: a fraud model at 95% accuracy catching 80% of fraud beats one at 99% accuracy catching none.

Q11.16

How do you choose a classification threshold in practice?

🕒 **In one line.** From business costs (minimise expected cost), from a capacity constraint (top- k), from a required precision or recall, or by maximising F1 or Youden's J on validation data; never on the test set.

⚙️ **The full answer.** Expected cost minimisation: predict positive if $p \cdot C_{FN} > (1 - p)C_{FP}$, i.e. $p > \frac{C_{FP}}{C_{FP} + C_{FN}}$ for calibrated p . Youden's J = TPR - FPR picks the ROC point farthest above the diagonal.

Q11.17

What is the Brier score of predictions (0.9, 0.2, 0.6) for labels (1, 0, 0)?

🕒 **In one line.** $(0.01 + 0.04 + 0.36)/3 = 0.137$.

Q11.18

💡 **Think deeper** Why is accuracy fine for MNIST digit classification but not for fraud?

🎯 **In one line.** MNIST classes are balanced and every error costs the same; fraud is extremely imbalanced and missed fraud costs far more than a false alarm, so a single accuracy number hides what matters.

Q11.19

How would you evaluate a regression model predicting salaries shown to job seekers?

🎯 **In one line.** With a scale-aware error such as MAPE or RMSLE (a 1 lakh error matters more at 5 lakh than at 50 lakh), plus coverage of the predicted range, and segment-wise errors by city, experience and function to catch systematic bias.

⚙️ **The full answer.** Also check bias (mean signed error) per segment: a model consistently under-predicting salaries in tier-2 cities misleads users. For ranges, measure the fraction of true salaries inside the band and its average width.

CHAPTER 12

Maximum Likelihood (MLE) and Maximum a Posteriori (MAP)

🎥 **Companion video for this chapter:** [StatQuest: Maximum Likelihood, Clearly Explained](#)

MLE and MAP are named in the syllabus and quietly sit under almost everything else: least squares, log loss, cross-entropy, Naive Bayes counts, Laplace smoothing, Ridge and Lasso, EM for mixtures. If you can explain where a loss function *comes from*, interviewers see that you understand ML rather than having memorised it.

12.1 Probability vs likelihood: flipping the question

- **Probability:** the parameter is known; how likely is the data? “If the coin is fair ($p = 0.5$), what is the chance of 7 heads in 10 tosses?” Answer: $\binom{10}{7}0.5^{10} = 0.117$.
- **Likelihood:** the data are known; how plausible is each parameter value? “I saw 7 heads in 10 tosses. How well does $p = 0.5$ explain that? $p = 0.7$? $p = 0.9$?”

The formula is the same, $L(p) = P(\text{data} \mid p)$, but we treat it as a function of p with the data fixed.

📊 **Small example: compare candidate parameters**

7 heads in 10. $L(p) \propto p^7(1-p)^3$. $L(0.5) = 0.00098$; $L(0.7) = 0.00222$; $L(0.9) = 0.00048$. The data are best explained by $p = 0.7$. Note that the likelihood values do not sum to 1 over p : likelihood is not a probability distribution over parameters.

12.2 Maximum likelihood estimation: the recipe

MLE chooses the parameter that makes the observed data most probable: $\hat{\theta} = \arg \max_{\theta} L(\theta)$.

1. Write the probability of one observation given θ .
2. Assume observations are independent: the likelihood is the product over observations.
3. Take the log (products become sums; the maximiser is unchanged because log is increasing; numerically stable).
4. Differentiate with respect to θ , set to zero, solve. (Check it is a maximum.)

12.2.1 Coin (Bernoulli)

k heads in n tosses: $\ell(p) = k \log p + (n - k) \log(1 - p)$. $\ell'(p) = \frac{k}{p} - \frac{n-k}{1-p} = 0 \Rightarrow \hat{p} = \frac{k}{n}$. The MLE is the observed fraction. 12 heads in 20 tosses: $\hat{p} = 0.6$.

12.2.2 Gaussian

Data $x_1, \dots, x_n$ from $\mathcal{N}(\mu, \sigma^2)$: $\ell = -\frac{n}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum (x_i - \mu)^2$.

- $\partial\ell/\partial\mu = 0 \Rightarrow \hat{\mu} = \bar{x}$ (the sample mean).
- $\partial\ell/\partial\sigma^2 = 0 \Rightarrow \hat{\sigma}^2 = \frac{1}{n} \sum (x_i - \bar{x})^2$: divides by n , not $n - 1$, so it is *biased* (slightly too small); the unbiased sample variance divides by $n - 1$.

Small example: Gaussian MLE

Data 2, 4, 6, 8: $\hat{\mu} = 5$, $\hat{\sigma}^2 = (9 + 1 + 1 + 9)/4 = 5$ (unbiased version: $20/3 = 6.67$).

12.2.3 Poisson and exponential

Poisson counts: $\hat{\lambda} = \bar{x}$. Applications per hour 3, 5, 4, 0, 3: $\hat{\lambda} = 3$. Exponential waiting times with rate λ : $\hat{\lambda} = 1/\bar{x}$. Times 2, 3, 5 minutes: $\bar{x} = 3.33$, $\hat{\lambda} = 0.3/\text{min}$.

12.2.4 A case where calculus fails: Uniform(0, θ)

$L(\theta) = \theta^{-n}$ if $\theta \geq \max x_i$, else 0. It decreases in θ , so the MLE is the smallest allowed value: $\hat{\theta} = \max x_i$. Data 1.2, 3.4, 2.7: $\hat{\theta} = 3.4$ (biased low: the true θ is almost surely a bit bigger).

12.3 Why ML losses are negative log-likelihoods

This is the single most useful idea in the chapter.

- **Regression with Gaussian noise:** $y = f(\mathbf{x}) + \varepsilon$, $\varepsilon \sim \mathcal{N}(0, \sigma^2)$. $-\log L = \frac{1}{2\sigma^2} \sum (y_i - f(\mathbf{x}_i))^2 + \text{const}$. **Minimising squared error = MLE.**
- **Regression with Laplace noise:** $-\log L \propto \sum |y_i - f(\mathbf{x}_i)|$. **MAE = MLE** under Laplace noise.
- **Binary classification:** $y \sim \text{Bernoulli}(p)$. $-\log L = -\sum [y \log p + (1 - y) \log(1 - p)]$. **Log loss = MLE.**
- **Multiclass:** categorical distribution. **Cross-entropy = MLE.**
- **Counts:** Poisson. **Poisson deviance = MLE** (Chapter 4).

So choosing a loss is choosing a noise model. Minimising cross-entropy between the data distribution and the model is also the same as minimising the **KL divergence** from data to model, since the data's own entropy is a constant.

12.4 MAP: adding prior beliefs

12.4.1 Bayes' theorem for parameters

$$\underbrace{P(\theta | D)}_{\text{posterior}} = \frac{\overbrace{P(D | \theta)}^{\text{likelihood}} \overbrace{P(\theta)}^{\text{prior}}}{P(D)}.$$

The **prior** encodes what we believe before seeing data. The **maximum a posteriori** estimate is the posterior's peak:

$$\hat{\theta}_{MAP} = \arg \max_{\theta} [\log P(D | \theta) + \log P(\theta)].$$

It is MLE plus a term that pulls toward plausible values: a regulariser.

12.4.2 Coin with a Beta prior

A Beta(a, b) prior on p acts like $a - 1$ imaginary heads and $b - 1$ imaginary tails. With k heads in n :

$$\hat{p}_{MAP} = \frac{k + a - 1}{n + a + b - 2}, \quad \text{posterior mean} = \frac{k + a}{n + a + b}.$$

Small example: three heads in three tosses

MLE: $3/3 = 1$: "this coin never lands tails". Absurd after three tosses. With a Beta(2,2) prior (mild belief in fairness): MAP = $(3 + 1)/(3 + 2) = 0.8$; posterior mean = $5/7 = 0.714$. With more data the prior's influence fades.

Small example: click-through rate smoothing

A new job got 1 click in 2 views: MLE CTR = 50%. Ranking by raw CTR would put it on top of jobs with 300 clicks in 10,000 views (3%). Use a Beta prior matching the platform average 5%, e.g. Beta(2, 38): posterior mean = $(1 + 2)/(2 + 40) = 7.1\%$. Small samples are pulled toward the average; large samples keep their own rate.

12.4.3 Regularisation is MAP

- Gaussian prior $w_j \sim \mathcal{N}(0, \tau^2)$: $-\log P(\mathbf{w}) = \frac{1}{2\tau^2} \|\mathbf{w}\|^2 + c$. With Gaussian noise σ^2 , MAP minimises $\frac{1}{2\sigma^2} \text{SSE} + \frac{1}{2\tau^2} \|\mathbf{w}\|^2$, i.e. **Ridge** with $\lambda = \sigma^2/\tau^2$.
- Laplace prior $p(w_j) \propto e^{-|w_j|/b}$: **Lasso**. The Laplace density has a sharp peak at 0, which is why it produces exact zeros.
- Laplace smoothing in Naive Bayes: MAP (or posterior mean) with a Dirichlet prior.

Small example: Ridge's λ from priors

Noise variance $\sigma^2 = 4$, prior variance $\tau^2 = 2$: $\lambda = 4/2 = 2$.

12.4.4 Conjugate priors

When prior and posterior are in the same family, updating is just adding counts: Beta–Bernoulli/Binomial, Gamma–Poisson, Dirichlet–Multinomial, Gaussian–Gaussian (known variance).

Small example: Gamma–Poisson

Prior Gamma($\alpha = 2, \beta = 1$) on a rate; observe counts 3 and 5. Posterior Gamma($2 + 8, 1 + 2$) = Gamma(10, 3); mean $10/3 = 3.33$.

Small example: Gaussian–Gaussian

Prior $\mu \sim \mathcal{N}(0, 1)$, one observation $x = 4$ with noise variance 1: posterior mean $= \frac{0/1+4/1}{1/1+1/1} = 2$, a precision-weighted average of prior mean and data.

12.5 MLE vs MAP vs full Bayesian inference

- **MLE**: one best value, no prior. Can overfit with little data (the 3/3 coin).
- **MAP**: one best value, with prior. Equivalent to regularised MLE. Not invariant to reparameterisation (the peak of a density moves when you change variables).
- **Full Bayes**: keep the whole posterior distribution and average predictions over it: $P(y \mid x, D) = \int P(y \mid x, \theta)P(\theta \mid D) d\theta$. Gives uncertainty; usually expensive (MCMC, variational inference).

As $n \rightarrow \infty$, all three agree (the likelihood dominates the prior).

12.6 Properties of MLE worth quoting

Under regularity conditions MLE is **consistent** (converges to the truth), **asymptotically normal**, and **asymptotically efficient** (achieves the Cramér–Rao lower bound: variance $\approx 1/(nI(\theta))$ with Fisher information I). It is **invariant**: the MLE of $g(\theta)$ is $g(\hat{\theta})$ (MLE of σ is $\sqrt{\hat{\sigma}^2}$). It can be biased in small samples ($\hat{\sigma}^2$, Uniform). Bernoulli Fisher information is $\frac{1}{p(1-p)}$, giving standard error $\sqrt{p(1-p)/n}$.

12.7 When there is no closed form: EM

With hidden (latent) variables (which mixture component generated each point), the likelihood is a sum inside a log and has no closed form. The **Expectation–Maximisation** algorithm alternates: E-step: compute the probability of each hidden assignment given current parameters; M-step: re-estimate parameters as if those soft assignments were observed. Each iteration never decreases the likelihood. Chapter 15 works through EM for Gaussian mixtures.

In practice: MLE and MAP at InfoEdge

- **CTR/apply-rate smoothing** for new jobs and new recruiters (Beta priors) before ranking.
- **Salary engine**: shrink company-specific salary estimates toward the industry average when a company has few data points (hierarchical/empirical Bayes = MAP with a learned prior).
- “Bayesian inference” is on the Tools slide: A/B test analysis with Beta posteriors (“probability variant B is better”).

12.8 Interview questions

Q12.1

What is the difference between probability and likelihood?

🕒 **In one line.** Probability fixes the parameters and asks how likely different data are; likelihood fixes the observed data and asks how well different parameter values explain it; same formula, different variable.

⚙️ **The full answer.** Probability sums to 1 over outcomes; likelihood does not sum or integrate to 1 over parameters. Worked example: $L(0.7) > L(0.5) > L(0.9)$ for 7/10 heads.

Q12.2

★ **PYQ (OA style: cross-entropy as likelihood)** Show that minimising log loss is maximum likelihood, and that minimising squared error is maximum likelihood under Gaussian noise.

🕒 **In one line.** For Bernoulli labels the negative log-likelihood is exactly the log loss; for $y = f(x)$ + Gaussian noise, the negative log-likelihood is $\frac{1}{2\sigma^2} \sum (y - f(x))^2$ plus a constant.

⚙️ **The full answer.** Bernoulli: $P(y | p) = p^y(1 - p)^{1-y}$; $-\log = -[y \log p + (1 - y) \log(1 - p)]$; summed over independent examples. Gaussian: $P(y | x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-(y-f)^2/(2\sigma^2)}$; $-\log = \frac{(y-f)^2}{2\sigma^2} + \log(\sqrt{2\pi}\sigma)$. Since σ does not depend on f , minimising SSE maximises likelihood.

💡 **If the interviewer pushes.** “So what does choosing MAE assume?” Laplace noise; its MLE location is the median.

Q12.3

★ **PYQ (INT: Bayes in real life; Naive Bayes smoothing as MAP)** Explain Laplace smoothing as a Bayesian (MAP) estimate.

🕒 **In one line.** Adding 1 to every count is the posterior estimate under a uniform Dirichlet prior over word probabilities: it is as if we had seen every word once in every class before looking at the data.

⚙️ **The full answer.** Dirichlet($\alpha_1, \dots, \alpha_V$) prior with multinomial counts c_w gives a Dirichlet($c_w + \alpha_w$) posterior; its mean is $\frac{c_w + \alpha}{N + V\alpha}$. With $\alpha = 1$ that is Laplace smoothing (posterior mean; the MAP with $\alpha = 2$ gives the same formula with $\alpha - 1 = 1$). It prevents zero probabilities and shrinks rare-word estimates toward uniform.

Q12.4

Derive the MLE of p for k successes in n Bernoulli trials.

🕒 **In one line.** $\hat{p} = k/n$.

⚙️ **The full answer.** $\ell(p) = k \log p + (n - k) \log(1 - p)$; $\ell' = k/p - (n - k)/(1 - p) = 0 \Rightarrow k(1 - p) = (n - k)p \Rightarrow p = k/n$; $\ell'' < 0$.

Q12.5

Derive the MLE of μ and σ^2 for a Gaussian. Is $\hat{\sigma}^2$ unbiased?

🎯 **In one line.** $\hat{\mu} = \bar{x}$ and $\hat{\sigma}^2 = \frac{1}{n} \sum (x_i - \bar{x})^2$, which is biased low by the factor $(n - 1)/n$.

⚙️ **The full answer.** The bias arises because $\bar{x}$ is fitted to the same data, so deviations from $\bar{x}$ are smaller on average than deviations from the true μ . Dividing by $n - 1$ (Bessel's correction) fixes it. Worked: data 2, 4, 6, 8 give 5 and 5 (unbiased 6.67).

Q12.6

Three heads in three tosses. Give the MLE, the MAP with a Beta(2,2) prior and the posterior mean. Which would you use?

🎯 **In one line.** MLE 1.0, MAP 0.8, posterior mean 0.714; with three tosses I would use the posterior mean or MAP because the MLE overfits to tiny data.

Q12.7

🧠 **Think deeper** Show that L2 regularisation is MAP with a Gaussian prior and derive λ .

🎯 **In one line.** With Gaussian noise σ^2 and prior $w_j \sim \mathcal{N}(0, \tau^2)$, the negative log-posterior is $\frac{1}{2\sigma^2} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|^2 + \frac{1}{2\tau^2} \|\mathbf{w}\|^2$; multiplying by $2\sigma^2$ gives Ridge with $\lambda = \sigma^2/\tau^2$.

⚙️ **The full answer.** Intuition: a strong prior (small τ^2) or noisy data (large σ^2) means large λ : trust the prior more. $\sigma^2 = 4$, $\tau^2 = 2$ gives $\lambda = 2$.

💡 **If the interviewer pushes.** "And L1?" Laplace prior. "Is the Lasso solution the posterior mean?" No, it is the posterior mode; the posterior mean under a Laplace prior is not sparse.

Q12.8

🧠 **Think deeper** A new job has 1 click from 2 impressions. How would you estimate its CTR for ranking, and why not just use 50%?

🎯 **In one line.** Use a Beta prior fitted to the platform's CTR distribution and take the posterior mean $(k + \alpha)/(n + \alpha + \beta)$; raw 50% from 2 impressions is extremely noisy and would wrongly rank the job above proven performers.

⚙️ **The full answer.** With Beta(2, 38) (mean 5%): $(1 + 2)/(2 + 40) = 7.1\%$. Fit α, β by method of moments on historic jobs (empirical Bayes). For exploration, sample from the posterior (Thompson sampling) instead of using the mean.

Q12.9

What is the MLE of θ for Uniform(0, θ) data 2.1, 4.7, 3.9, 1.2? Why can't you use calculus?

🕒 **In one line.** $\hat{\theta} = 4.7$; the likelihood θ^{-n} is decreasing on the allowed region $\theta \geq \max x_i$ and zero below it, so the maximum is at the boundary, where the derivative is not zero.

Q12.10

What is the MLE of λ for Poisson counts 3, 5, 4, 0, 3, and for an exponential with waiting times 2, 3, 5?

🕒 **In one line.** Poisson: $\bar{x} = 3$; exponential: $1/\bar{x} = 0.3$ per minute.

⚙️ **The full answer.** Poisson: $\ell = \sum [x_i \log \lambda - \lambda] + c$; $\ell' = \sum x_i / \lambda - n = 0$. Exponential: $\ell = n \log \lambda - \lambda \sum x_i$; $\ell' = n / \lambda - \sum x_i = 0$.

Q12.11

Why do we maximise the log-likelihood instead of the likelihood?

🕒 **In one line.** Log turns the product over observations into a sum (easier to differentiate and to optimise with SGD), avoids numerical underflow, and does not move the maximiser because log is strictly increasing.

Q12.12

Explain MLE vs MAP vs full Bayesian inference in three lines.

🕒 **In one line.** MLE: the parameter that best explains the data; MAP: the parameter that best explains the data and the prior (regularised MLE); full Bayes: keep the whole posterior and average predictions over it, which also gives uncertainty.

Q12.13

🧠 **Think deeper** Why is MAP not invariant to reparameterisation while MLE is?

🕒 **In one line.** A likelihood is a function of the parameter, so its maximiser maps through any reparameterisation; a posterior is a density, and changing variables multiplies it by a Jacobian, which can move its peak.

⚙️ **The full answer.** If $\phi = g(\theta)$, $p_\phi(\phi) = p_\theta(g^{-1}(\phi)) |dg^{-1}/d\phi|$. The extra factor shifts the mode. Example: the mode of a posterior over σ differs from the square root of the mode over σ^2 .

Q12.14

What does minimising cross-entropy have to do with KL divergence?

🕒 **In one line.** Cross-entropy $H(p, q) = H(p) + KL(p \parallel q)$; since the data entropy $H(p)$ does not depend on the model, minimising cross-entropy minimises the KL divergence from the data distribution to the model.

Q12.15

What are conjugate priors? Give three pairs and a worked update.

🕒 **In one line.** A prior is conjugate to a likelihood if the posterior stays in the prior's family: Beta–Binomial, Gamma–Poisson, Dirichlet–Multinomial; e.g. Gamma(2,1) prior and counts 3, 5 give Gamma(10, 3), mean 3.33.

Q12.16

💡 **Think deeper** Logistic regression on perfectly separable data has no MLE. How does a prior fix this?

🕒 **In one line.** The likelihood keeps increasing as $\|\mathbf{w}\| \rightarrow \infty$, so no maximiser exists; a Gaussian prior adds $-\|\mathbf{w}\|^2/(2\tau^2)$ to the log-posterior, which eventually outweighs the shrinking gains in likelihood, so a finite MAP exists.

Q12.17

What is Fisher information and what does it tell you?

🕒 **In one line.** It measures how sharply the log-likelihood is curved around the true parameter, i.e. how much information each observation carries; the MLE's variance is approximately $1/(nI(\theta))$.

🧮 **Worked example.** Bernoulli: $I(p) = 1/(p(1-p))$; with $p = 0.1$, $n = 900$: $SE \approx \sqrt{0.09/900} = 0.01$.

Q12.18

💡 **Think deeper** In a salary engine, a small startup has 3 salary records for “Data Scientist”. How do you estimate its typical salary?

🕒 **In one line.** Use a hierarchical (empirical Bayes) estimate that shrinks the company's mean toward the industry/city mean in proportion to how few records it has: $\hat{\mu}_c = \frac{n_c \bar{x}_c / \sigma^2 + \mu_0 / \tau^2}{n_c / \sigma^2 + 1 / \tau^2}$.

⚙️ **The full answer.** This is the Gaussian–Gaussian posterior mean. With $n_c = 3$ the industry mean carries real weight; with $n_c = 300$ the company's own mean dominates. It is also how mixed-effects models and target encoding with priors work.

Q12.19

How is MLE used to train neural networks?

🎯 **In one line.** Training with cross-entropy or MSE is maximising the likelihood of the labels under the network's output distribution; weight decay adds a Gaussian prior (MAP); and the output layer (sigmoid, softmax, linear) defines which distribution is assumed.

CHAPTER 13

Missing Data Imputation, Outliers and Class Imbalance

🎥 **Companion video for this chapter:** [Kaggle Learn: Missing Values \(Intermediate ML\)](#)

“Missing data imputation” is a syllabus item, and the reported Round 3 checked for missing values and outliers before modelling. The candidate who did *not* blindly delete genuine extreme marks was praised for “presence of mind”. This chapter teaches the reasoning: *why* values are missing decides *how* to fill them.

13.1 Why missing values are a modelling problem

Most algorithms (linear models, SVMs, neural networks, k -NN) cannot take a blank. More importantly, *the fact that a value is missing can carry information*, and careless filling can bias the model or leak the target.

📊 Small example: a missing value that means something

On Naukri, “current salary” is blank for 40% of freshers and 5% of experienced candidates. If you fill blanks with the overall median (say 8 lakh), freshers suddenly look like mid-level professionals. If you drop rows with blanks, you drop most freshers. Either way the model learns a distorted picture.

13.2 Three missingness mechanisms (Rubin)

- **MCAR (missing completely at random):** missingness has nothing to do with any value. A lab technician randomly lost 3% of forms. Deleting such rows loses data but does not bias results.
- **MAR (missing at random):** missingness depends on *observed* variables only. Freshers (experience = 0, recorded) often leave salary blank; *among freshers*, blanks are random. Model-based imputation using experience works well.
- **MNAR (missing not at random):** missingness depends on the *missing value itself*. People with very high incomes decline to report income. No imputation from observed data can fully correct it; you need domain knowledge, a missingness indicator, or modelling the missingness.

You cannot prove MAR vs MNAR from the data alone; you argue it from how the data were collected.

13.3 Option 1: deletion

- **Listwise** (drop rows with any missing value). Fine for small MCAR fractions. Danger: with many columns, few rows survive.

Small example: how fast rows disappear

12 independent columns each missing 5% of the time: complete rows = $0.95^{12} = 54\%$. Almost half the data gone.

- **Column deletion:** drop a column that is mostly missing (say 85%) *and* whose missingness is uninformative. If missingness predicts the target, keep a missing indicator instead.

13.4 Option 2: simple (univariate) imputation

- **Mean:** fine for symmetric distributions. **Median:** robust for skewed ones (salary, price). **Mode / most frequent** or a new category "Unknown" for categoricals. **Constant** (0, -999 for tree models).
- Time series: **forward fill** (last observation carried forward), **interpolation**.

13.4.1 What mean imputation does to the data

Small example: shrinking variance

Column: 2, 4, 6, NaN, NaN. Observed mean 4, population variance $(4 + 0 + 4)/3 = 2.67$. After filling with 4: values 2, 4, 6, 4, 4; variance $(4 + 0 + 4 + 0 + 0)/5 = 1.6$. The variance shrank by 40%, and correlations with other features are weakened, because the filled points sit exactly at the mean and carry no relationship.

Consequences: distorted distributions (a spike at the mean), attenuated correlations, over-confident inference. Still often fine for prediction with tree models, especially with a missing indicator.

13.5 Option 3: multivariate (model-based) imputation

- **Group-wise:** fill salary with the median of the candidate's experience band and city. Simple, uses MAR structure.
- **k-NN imputation:** average the value from the k most similar rows (on observed features). Scale first.

Small example: KNN imputation

Three nearest neighbours have age 24, 27, 30: impute 27. Distance-weighted with neighbours 10 at $d = 1$ and 20 at $d = 2$, weights $1/d$: $(10 + 10)/1.5 = 13.3$.

- **Regression imputation:** predict the missing column from the others. Deterministic predictions sit exactly on the regression line, understating variance; **stochastic regression imputation** adds a random residual.
- **MICE / IterativeImputer:** cycle through columns, regressing each on all others, repeat until stable.
- **Multiple imputation:** create m completed datasets with randomness, fit the model on each, combine (Rubin's rules). This propagates the uncertainty about the filled values into standard errors, which single imputation ignores.
- **Matrix factorisation / low-rank completion:** fill a ratings-like matrix from latent factors (Chapter 14).

13.6 Option 4: let the model handle it

XGBoost, LightGBM, CatBoost and scikit-learn's HistGradientBoosting accept NaN and learn, at each split, which side missing values go (Chapter 8). This automatically exploits informative missingness.

13.7 The missing-indicator trick

Add a binary column `salary_missing`, then impute. The model can learn “blank salary” as a signal of its own. It is the standard way to protect against MNAR and informative MAR.

13.8 Leakage: fit imputers on training data only

The median, the KNN neighbours, the MICE regressions must be learned from training data and *applied* to validation/test. Put the imputer inside a Pipeline so cross-validation refits it per fold.

```
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
pipe = Pipeline([("impute", SimpleImputer(strategy="median", add_indicator=True)),
                 ("model", LogisticRegression(max_iter=1000))])
pipe.fit(X_train, y_train)           # median learned from training rows only
```

13.9 Outliers: detect, understand, then decide

13.9.1 Detection

- **Domain ranges** first: a percentage outside 0–100 is an error; 99.5 is a genius, not an error. The InfoEdge coding round tested exactly this with `describe()`.
- **IQR rule**: outliers beyond $Q_1 - 1.5 \text{ IQR}$ or $Q_3 + 1.5 \text{ IQR}$, $\text{IQR} = Q_3 - Q_1$.

Small example: IQR

Data 3, 5, 7, 8, 9, 11, 13, 40 with $Q_1 = 6$, $Q_3 = 12$: $\text{IQR} = 6$, fences -3 and 21 : 40 is flagged.

- **z-score**: $|z| > 3$ (assumes rough normality; the outlier itself inflates the standard deviation). Robust version with median and MAD.

Small example: z-score

Mean 50, sd 8, $x = 76$: $z = 26/8 = 3.25$: flagged.

- Multivariate: Mahalanobis distance, Isolation Forest, Local Outlier Factor, DBSCAN noise points.

13.9.2 What to do

1. **Error?** (age 250, price Rs. 1, negative experience) Fix or set to missing.

2. **Genuine but extreme?** (a 5 crore salary; a 99th-percentile score) Keep. Use robust methods: median/MAE/Huber loss, tree models (insensitive to feature outliers), log transforms, **win-sorising/capping** at the 1st/99th percentile, robust scaling (RobustScaler uses median and IQR).
3. **A different population?** (bulk-posting staffing agencies among normal recruiters) Model separately or add a segment feature.
4. **The thing you are looking for?** (fraud) Then outliers are the signal: anomaly detection.

13.10 Class imbalance

When one class is rare (fraud 0.5%, high-intent buyers 2%):

- **Evaluate correctly:** precision, recall, F1, PR AUC, precision@k; stratified splits (Chapter 11).
- **Class weights:** scale the loss of the rare class. scikit-learn "balanced": $w_c = \frac{n}{Kn_c}$.

Small example: balanced weight

900 negatives, 100 positives: positive weight $1000/(2 \cdot 100) = 5$, negative $1000/1800 = 0.556$.
XGBoost `scale_pos_weight` = n_{neg}/n_{pos} , e.g. $9500/500 = 19$.

- **Random undersampling** of the majority: fast, loses information.
- **Random oversampling** of the minority: duplicates; risk of overfitting those rows.
- **SMOTE:** create synthetic minority points by interpolating between a minority point and one of its k minority neighbours: $\mathbf{x}_{new} = \mathbf{x} + u(\mathbf{x}_{nn} - \mathbf{x})$, $u \sim U(0, 1)$.

Small example: SMOTE

$\mathbf{x} = (2, 4)$, neighbour $(6, 8)$, $u = 0.25$: $(2, 4) + 0.25(4, 4) = (3, 5)$.

Variants: Borderline-SMOTE (only near the boundary), ADASYN (more synthetic points where the minority is harder), SMOTE-NC for categorical features. Cautions: synthetic points can land in majority regions; apply *only* to training folds.

- **NearMiss, Tomek links:** informed undersampling (remove majority points close to the minority / boundary noise).
- **Threshold tuning:** often the single most effective step.
- **Anomaly detection** when positives are extremely rare and varied.

In practice: data quality at InfoEdge scale

Naukri profiles are user-entered: free-text skills, blank salaries, "0" as a placeholder, 50 spellings of the same company. The SilverSkAI taxonomy slide (50 lakh raw skills normalised to 70K; 80 lakh company strings to 7.2 lakh) is the industrial answer to a data-quality problem: standardise first, then impute. For 99acres, listing prices posted as "1" (to mean "call for price") are not outliers to cap; they are missing values in disguise.

13.11 Interview questions

Q13.1

★ **PYQ (Coding round / INT: handling missing values)** You find 8% NaNs in a numeric column. Walk me through your options and how you would choose.

🕒 **In one line.** First ask why they are missing (MCAR, MAR, MNAR) and whether missingness predicts the target; then choose between deletion, simple imputation (median for skewed data), model-based imputation (KNN, MICE, group medians), native NaN handling in boosted trees, and always consider adding a missing indicator, fitting everything on training data only.

⚙️ **The full answer.**

1. **Investigate:** `df[col].isna().groupby(df[target]).mean()` and against other columns. If the missing rate differs by target or segment, missingness is informative.
2. **If MCAR and few rows:** dropping is acceptable.
3. **If skewed:** median, not mean.
4. **If related to other features (MAR):** group-wise median or KNN/IterativeImputer.
5. **If possibly MNAR:** add an indicator; consider domain-specific fills.
6. **Model choice:** with LightGBM, leave NaN as is.
7. **Leakage:** put the imputer in a Pipeline.

📊 **Worked example.** In the placement data, “work experience months” is blank for students without work experience: that is not missing, it is 0, and the binary “work experience” column confirms it. Reading the data dictionary resolves many “missing” values.

⚠️ **Trap to avoid.** Filling with the mean computed on the whole dataset before splitting (leakage), or filling a categorical with the mean of its codes.

Q13.2

★ **PYQ (INT: mean vs median imputation)** When do you use mean vs median imputation? Show the effect on variance.

🕒 **In one line.** Median for skewed or outlier-prone variables, mean for roughly symmetric ones; both shrink the variance and weaken correlations because all filled values sit at one point.

⚙️ **The full answer.** Salary is right-skewed: a few crore-level salaries pull the mean far above the typical value; median is representative. Variance example: $2.67 \rightarrow 1.6$.

Q13.3

★ **PYQ (INT: imputing categorical features)** How do you impute a missing categorical column like “city”?

🕒 **In one line.** With the most frequent category if missingness is rare and random, or (usually better) with an explicit “Unknown” category so the model can learn whether missingness matters;

model-based imputation (from pincode, company location) if available.

⚙️ **The full answer.** Mode imputation inflates the most common city; “Unknown” preserves information. For tree models with native categorical support, NaN can be its own branch.

Q13.4

★ **PYQ (INT: XGBoost and missing values)** How do gradient-boosted trees handle missing values, and should you still impute?

🎯 **In one line.** They learn a default direction for missing values at every split, choosing the side that gives the higher gain, so imputation is usually unnecessary and can even remove the signal in missingness.

⚙️ **The full answer.** Imputing with the median merges “missing” with “typical”, which the tree can no longer distinguish. Exceptions: if test data will have missingness patterns not present in training (then imputation can be more robust), or if the same features feed other models.

Q13.5

Explain MCAR, MAR and MNAR with examples from a job portal.

🎯 **In one line.** MCAR: a server bug randomly dropped 2% of profile fields; MAR: freshers leave current salary blank (depends on observed experience); MNAR: very high earners hide salary (depends on the missing salary itself).

⚙️ **The full answer.** Under MCAR, deletion is unbiased. Under MAR, imputation from observed features is valid. Under MNAR, imputing from others is biased; an indicator, sensitivity analysis, or a model of the missingness is needed.

Q13.6

🧠 **Think deeper** Why does multiple imputation create several completed datasets instead of one?

🎯 **In one line.** Because a single filled-in value pretends we know the missing number exactly; drawing several plausible values and pooling the results carries the uncertainty about the missing data into the final estimates and standard errors.

⚙️ **The full answer.** With single imputation, confidence intervals are too narrow and p-values too small. Rubin’s rules combine: total variance = within-imputation variance + $(1 + 1/m) \times$ between-imputation variance. For pure prediction tasks, single imputation is usually fine.

Q13.7

★ **PYQ (INT: outlier detection)** Data 3, 5, 7, 8, 9, 11, 13, 40. Which points are outliers by the IQR rule, and what would you do about them?

🎯 **In one line.** With $Q_1 = 6$, $Q_3 = 12$, $IQR = 6$, fences are -3 and 21 , so 40 is an outlier; whether

to remove it depends on whether it is an error, a genuine extreme, or the phenomenon of interest.

⚙️ **The full answer.** I would look at the row: if it is “years of experience = 40” for a 60-year-old, keep it; if “age = 40” in a column of ages of 3–13-year-olds, it is likely an error. Rather than delete, prefer robust methods or capping.

💡 **If the interviewer pushes.** “Why not z-scores?” With small or skewed samples, the outlier inflates the sd; here $\bar{x} = 12$, $s \approx 11.7$, $z(40) = 2.4$: not flagged at 3 even though it is clearly unusual. IQR or median/MAD is more robust.

Q13.8

★ **PYQ (INT: dealing with outliers)** Which models and statistics are robust to outliers?

🎯 **In one line.** Medians and IQR (not means and standard deviations), MAE/Huber/quantile losses (not MSE), tree-based models with respect to feature outliers, and rank-based methods like Spearman correlation and AUC.

⚙️ **The full answer.** Trees split on order, so an extreme feature value just lands in the extreme leaf. But trees are not robust to target outliers under MSE (a leaf mean can be dragged). Linear models, k -NN, k -means, PCA and SVMs (without robust scaling) are sensitive.

Q13.9

★ **PYQ (reported Round 3: presence of mind)** `describe()` shows a maximum of 98.5 in an exam-percentage column, far above the 75th percentile of 72. Do you remove it?

🎯 **In one line.** No: it lies within the valid 0–100 range, so it is a genuine high score; removing it would delete real information, and I would only treat values outside the valid range as errors.

⚙️ **The full answer.** Statistical outlier rules flag “unusual” not “wrong”. Blindly applying box-plot or z-score filtering removes the best students, biasing the model. This is exactly what the reported interview tested.

Q13.10

★ **PYQ (INT: imbalanced data)** Fraud is 0.5% of transactions. Which techniques help, and in what order would you try them?

🎯 **In one line.** Correct metrics and stratified splits first, then class weights or `scale_pos_weight`, threshold tuning, and only then resampling (undersampling, SMOTE) or anomaly detection, always applied to training folds only.

⚙️ **The full answer.** Weighting is simple and does not distort data. Resampling helps some models but breaks calibration. A capable GBM with good features and a tuned threshold often beats elaborate resampling. Collect more positives if possible (active learning on the review queue).

Q13.11

★ **PYQ (INT: SMOTE mechanics)** How does SMOTE generate a sample? Compute one for $\mathbf{x} = (2, 4)$, neighbour $(6, 8)$, $u = 0.25$.

🕒 **In one line.** It picks a minority point and one of its k nearest minority neighbours and creates a new point at a random position on the segment between them: here $(3, 5)$.

⚙️ **The full answer.** $\mathbf{x} + u(\mathbf{x}_{nn} - \mathbf{x}) = (2, 4) + 0.25(4, 4)$. It enlarges minority regions instead of duplicating points.

💡 **If the interviewer pushes.** “Weaknesses?” Can create points in majority territory (overlap), amplifies noise, meaningless for categorical/one-hot features without SMOTE-NC, and must be done inside CV.

Q13.12

★ **PYQ (INT: SMOTE vs NearMiss vs ADASYN)** Describe each.

🕒 **In one line.** SMOTE interpolates new minority points uniformly; ADASYN interpolates more around minority points that are hard to learn (many majority neighbours); NearMiss undersamples the majority by keeping majority points closest to the minority.

Q13.13

💡 **Think deeper** Why is applying SMOTE before the train–test split a serious mistake?

🕒 **In one line.** Synthetic points are interpolations of real minority points, so near-copies of test points end up in training; the model looks far better than it is.

⚙️ **The full answer.** A test positive $\mathbf{x}$ and a synthetic point between $\mathbf{x}$ and its neighbour are almost identical. Correct: split first, then oversample only the training fold (imblearn’s Pipeline does this per CV fold).

Q13.14

Compute the `class_weight="balanced"` weights for 900 negatives and 100 positives, and `scale_pos_weight` for 9,500 negatives and 500 positives.

🕒 **In one line.** Positive weight 5, negative 0.556; `scale_pos_weight` = 19.

Q13.15

💡 **Think deeper** A column is 85% missing. Drop it?

🕒 **In one line.** Not automatically: if its missingness or its observed values predict the target, keep it with a missing indicator (or as-is for GBMs); drop it only if both are uninformative.

⚙️ **The full answer.** Example: “portfolio URL” is 85% blank, but having one strongly predicts

recruiter interest. Test by comparing target rates for missing vs present and by permutation importance.

Q13.16

What is winsorising, and when would you use it?

🎯 **In one line.** Capping values at chosen percentiles (e.g. 1st and 99th) instead of deleting them; use it for linear models or neural nets when extreme but genuine values would otherwise dominate the fit.

⚙️ **The full answer.** It keeps the row and its order information, reduces leverage. The caps must be computed on training data.

Q13.17

💡 **Think deeper** Freshers' salary is missing. Someone fills it with the global median and the model learns "freshers earn 8 lakh". What would you do instead?

🎯 **In one line.** Recognise the MAR (or structural) pattern and impute within segments (experience band, degree, city), add a missing indicator, or let a GBM use NaN directly; for a fresher, "current salary" may be genuinely not applicable, which a separate category captures.

Q13.18

Why must an imputer be fitted only on training data, and how do you enforce it?

🎯 **In one line.** Because statistics computed on validation/test rows leak information about them into training; enforce it with a scikit-learn Pipeline (or ColumnTransformer) so that fit is called only on training folds.

Q13.19

💡 **Think deeper** Missing values only occur in the test data for a feature that was always present in training. What happens and what do you do?

🎯 **In one line.** The model never learned what to do with missingness for that feature: a GBM will send NaNs down its default (arbitrary) direction and a pipeline imputer will fill with the training median; investigate the upstream pipeline break, and design training to include realistic missingness (randomly masking values) if it will recur.

Q13.20

Which scaler would you use with outliers, and why?

🎯 **In one line.** RobustScaler, which centres on the median and scales by the IQR, so a few extreme values do not compress all the other values.

 **Worked example.** Values 1–10 plus 1,000: StandardScaler squeezes 1–10 into a tiny range near -0.3 ; RobustScaler keeps them spread.

CHAPTER 14

Matrix Factorisations: SVD, PCA, NMF, Recommenders and t-SNE

🎥 **Companion video for this chapter:** [StatQuest: Principal Component Analysis \(PCA\), Step by Step](#)

The InfoEdge hire was asked how PCA uses SVD, what a scree plot is, and how t-SNE and its perplexity work. Matrix factorisation is also how classic recommenders (“jobs you may like”, Jeevansathi’s collaborative models) work. This chapter starts from “what is a matrix” and ends with recommender maths you can do on paper.

14.1 What “factorising a matrix” means

A matrix is a table of numbers. Factorising it means writing it as a product of simpler matrices, just as $12 = 3 \times 4$. Why? Because the factors reveal structure and compress data.

📊 Small example: a rank-1 table

Ratings of 3 users for 3 jobs: $A = \begin{pmatrix} 2 & 4 & 6 \\ 1 & 2 & 3 \\ 3 & 6 & 9 \end{pmatrix} = \begin{pmatrix} 2 \\ 1 \\ 3 \end{pmatrix} \begin{pmatrix} 1 & 2 & 3 \end{pmatrix}$. Nine numbers are explained by six: a “user enthusiasm” and a “job attractiveness”. Every row is a multiple of every other: the **rank** is 1 (only one independent pattern).

Real data are never exactly low-rank, but they are often *approximately* low-rank: a few patterns explain most of the variation. Factorisation finds those patterns.

14.2 A quick refresher: vectors, matrices, eigenvectors

- Multiplying a matrix by a vector transforms the vector (stretches, rotates, projects).
- An **eigenvector** $\mathbf{v}$ of a square matrix A is a direction that A only stretches: $A\mathbf{v} = \lambda\mathbf{v}$; the stretch factor λ is the **eigenvalue**.
- A symmetric matrix (like a covariance matrix) has real eigenvalues and perpendicular eigenvectors.
- The **covariance matrix** Σ of centred data X ($n \times d$) is $\frac{1}{n-1} X^\top X$: entry (j, k) is the covariance of features j and k ; the diagonal holds variances.

14.3 SVD: the master factorisation

Every real $m \times n$ matrix can be written

$$A = U\Sigma V^\top,$$

- V ($n \times n$): orthonormal columns $\mathbf{v}_i$, the **right singular vectors**: directions in the input (column) space.

📺 **Stuck on this page?** Watch: [StatQuest: Principal Component Analysis \(PCA\), Step by Step](#)

- Σ : diagonal with **singular values** $\sigma_1 \geq \sigma_2 \geq \dots \geq 0$: how much A stretches each direction.
- U ($m \times m$): orthonormal columns $\mathbf{u}_i$, the **left singular vectors**: where those directions land.

Geometric meaning: any linear map is a rotation ($V^\top$), then a stretch along the axes (Σ), then another rotation (U). Equivalently $A = \sum_i \sigma_i \mathbf{u}_i \mathbf{v}_i^\top$: a sum of rank-1 layers, most important first.

Connections: $A^\top A = V \Sigma^2 V^\top$, so the $\mathbf{v}_i$ are eigenvectors of $A^\top A$ and σ_i^2 its eigenvalues. Rank = number of non-zero singular values. For a symmetric positive semi-definite matrix, SVD and eigendecomposition coincide. Thin SVD of a 1000×50 full-rank matrix: U is 1000×50 , Σ 50×50 , V 50×50 .

14.3.1 Low-rank approximation (Eckart–Young)

Keep only the top k terms: $A_k = \sum_{i=1}^k \sigma_i \mathbf{u}_i \mathbf{v}_i^\top$. This is the *best* rank- k approximation of A in both Frobenius and spectral norms. The squared Frobenius error equals the sum of the discarded σ_i^2 .

Small example: compression

Singular values 5, 3, 1. Rank-1 error = $\sqrt{9+1} = 3.16$. Rank 2 keeps $(25+9)/35 = 97\%$ of the squared norm. A 100×200 image stored with rank 10 needs $10(100+200+1) = 3,010$ numbers instead of 20,000.

14.4 PCA is SVD of the centred data

14.4.1 The goal

Find a few new axes (**principal components**) that capture as much of the data's variation as possible, so we can describe each point with fewer numbers while losing little information.

14.4.2 Step by step

1. **Centre** each feature (subtract its mean). Usually also **standardise** (divide by its standard deviation), otherwise features with large units dominate.
2. Compute the covariance matrix $\Sigma = \frac{1}{n-1} X^\top X$.
3. Find its eigenvectors (the principal directions) and eigenvalues (the variance along each).
4. Sort by eigenvalue; keep the top k . Project: $Z = X V_k$ (each point's coordinates on the new axes).

Via SVD: $X = U \Sigma V^\top$ directly gives the directions V ; the variances are $\sigma_i^2/(n-1)$; the projected coordinates are $U \Sigma$. Libraries use SVD because it avoids forming $X^\top X$ (which squares the condition number and loses precision).

Small example: PCA on two features by hand

Covariance matrix $\begin{pmatrix} 2 & 1 \\ 1 & 2 \end{pmatrix}$ (two standardised-ish features with covariance 1). Eigenvalues: $\det \begin{pmatrix} 2-\lambda & 1 \\ 1 & 2-\lambda \end{pmatrix} = (2-\lambda)^2 - 1 = 0 \Rightarrow \lambda = 3, 1$. Eigenvectors: $(1, 1)/\sqrt{2}$ for 3 and $(1, -1)/\sqrt{2}$ for 1. PC1 is “the sum of the two features” and carries $3/4 = 75\%$ of the variance.

14.4.3 How many components? The scree plot

A **scree plot** shows eigenvalues (or explained-variance ratio $\lambda_i / \sum \lambda$) against component number. Look for the “elbow” after which the curve flattens (“scree” is the rubble at the foot of a cliff). Alternatives: keep enough components for 90–95% cumulative variance; Kaiser rule (eigenvalue > 1 on standardised data); or choose k by cross-validated downstream performance.

Small example: explained variance

Eigenvalues 6, 3, 1: PC1 explains 60%, PC1+PC2 90%.

14.4.4 Limitations of PCA

Linear only; maximises variance, which need not be what matters for a target (Chapter 10); sensitive to scaling and outliers; components are mixtures of all features and hard to interpret; assumes the mean and covariance summarise the data.

14.5 NMF: non-negative parts

Non-negative matrix factorisation: $A \approx WH$ with all entries of W and H non-negative. Because nothing can cancel, factors are additive “parts”. On a TF-IDF document–term matrix, each row of H is a *topic* (a non-negative combination of words) and each row of W says how much of each topic a document contains. Fitted by multiplicative updates or alternating least squares; not unique; needs non-negative input.

14.6 Matrix factorisation for recommendations

14.6.1 The setting

A huge user $\times$ item matrix R (seekers $\times$ jobs) with a few observed entries (ratings, applies, clicks); most are unknown, not zero. Goal: predict the unknown entries.

14.6.2 The model

Give each user u a vector $\mathbf{p}_u \in \mathbb{R}^k$ and each item i a vector $\mathbf{q}_i \in \mathbb{R}^k$ (latent factors, $k \approx 32$ –200). Predict

$$\hat{r}_{ui} = \mu + b_u + b_i + \mathbf{p}_u^\top \mathbf{q}_i,$$

with global mean μ , user bias b_u (this user rates/applies a lot) and item bias b_i (popular job). Factors are learned so that dot products reproduce the observed entries:

$$\min \sum_{(u,i) \text{ observed}} (r_{ui} - \hat{r}_{ui})^2 + \lambda(\|\mathbf{p}_u\|^2 + \|\mathbf{q}_i\|^2 + b_u^2 + b_i^2).$$

Latent dimensions often become interpretable (“remote-friendly”, “startup vs MNC”, “data roles”).

Small example: predicting a rating

$\mu = 3.5$, $b_u = 0.2$, $b_i = -0.3$, $\mathbf{p}_u = (1, 2)$, $\mathbf{q}_i = (3, 0.5)$: $\mathbf{p}_u^\top \mathbf{q}_i = 3 + 1 = 4$; $\hat{r} = 3.5 + 0.2 - 0.3 + 4 = 7.4$ (on an unbounded scale; clip if needed).

14.6.3 Why not plain SVD?

SVD needs a complete matrix; filling unknowns with 0 says “the user dislikes every unseen job”, which is false. Matrix factorisation only fits the observed entries.

14.6.4 Training: SGD and ALS

SGD: for each observed (u, i) , error $e = r - \hat{r}$; update $\mathbf{p}_u \leftarrow \mathbf{p}_u + \eta(e \mathbf{q}_i - \lambda \mathbf{p}_u)$ and $\mathbf{q}_i \leftarrow \mathbf{q}_i + \eta(e \mathbf{p}_u - \lambda \mathbf{q}_i)$.

Small example: one SGD step

$r = 4$, $\mathbf{p} = (1, 0)$, $\mathbf{q} = (2, 1)$, $\eta = 0.1$, $\lambda = 0$: $\hat{r} = 2$, $e = 2$. $\mathbf{p} \leftarrow (1, 0) + 0.2(2, 1) = (1.4, 0.2)$;
 $\mathbf{q} \leftarrow (2, 1) + 0.2(1, 0) = (2.2, 1)$ (using the old $\mathbf{p}$).

ALS (alternating least squares): fix all item vectors and solve each user vector exactly (a small ridge regression), then fix users and solve items; repeat. Each subproblem is convex and independent per user, so ALS parallelises well (Spark MLlib) and handles implicit feedback.

14.6.5 Implicit feedback and cold start

Clicks and applies are *implicit*: an absent click is not a dislike. Implicit ALS treats every entry with a confidence weight $(1 + \alpha \cdot \text{count})$ and a binary preference. **Cold start** (new job, new user): no interactions, no factors. Use content features (title, skills, location embeddings), hybrid models (factorisation machines, two-tower networks), popularity priors, and exploration.

14.7 t-SNE and UMAP: seeing clusters

PCA is linear and preserves large-scale variance. For *visualising* clusters in high-dimensional data (resume embeddings), use non-linear methods that preserve *neighbourhoods*.

14.7.1 t-SNE in steps

1. In the original space, turn distances into probabilities that j is i 's neighbour: a Gaussian around each point, $p_{j|i} \propto \exp(-\|\mathbf{x}_i - \mathbf{x}_j\|^2 / 2\sigma_i^2)$; symmetrise.
2. Each point's bandwidth σ_i is chosen so that its neighbourhood has a set **perplexity** (roughly the effective number of neighbours, typically 5–50): dense regions get small σ_i , sparse regions large.
3. In 2-D, define similarities with a heavy-tailed **Student-t** distribution: $q_{ij} \propto (1 + \|\mathbf{y}_i - \mathbf{y}_j\|^2)^{-1}$.
4. Move the 2-D points by gradient descent to minimise $KL(P \parallel Q)$.

Why the t-distribution: its heavy tails let moderately distant points sit far apart in 2-D, relieving the “crowding problem” and producing clear, separated clusters.

14.7.2 How to read (and not read) a t-SNE plot

- Clusters and local neighbourhoods are meaningful.
- Cluster *sizes* and the *distances between clusters* are not reliably meaningful.
- Results depend on perplexity and random seed; run several.

- It is for visualisation: no transform for new points; do not feed t-SNE coordinates to a classifier.
- **Perplexity**: low values focus on very local structure (many small clumps); high values capture more global structure but can merge clusters.

UMAP is a faster alternative that preserves more global structure and can transform new points.

In practice: factorisations at InfoEdge

- **Collaborative models** (Jeevansathi slide: “content models, activity-based models, collaborative models”): MF/ALS on profile-view and interest interactions; “users who showed interest in similar profiles”.
- **Job recommendations** (“Reco” in the Seeker AI engines): MF or two-tower models on apply data, with content features for cold start (most jobs live only weeks, so cold start is the norm).
- **LSA** (truncated SVD of TF-IDF) as a quick semantic search baseline; PCA to compress embeddings for storage in a vector index.
- t-SNE/UMAP of resume embeddings to sanity-check that the taxonomy’s clusters (Data Science, DevOps, Sales) separate.

14.8 Interview questions

Q14.1

★ **PYQ (INT: PCA via SVD and the scree plot)** Explain PCA and how it uses SVD. What is a scree plot?

🎯 **In one line.** PCA finds orthogonal directions of maximum variance in centred data; the right singular vectors of the centred data matrix are those directions, and the squared singular values divided by $n - 1$ are the variances; a scree plot shows those variances in decreasing order to choose how many components to keep.

⚙️ **The full answer.** **Objective.** Find unit vector $\mathbf{v}$ maximising the variance of projections $\text{Var}(X\mathbf{v}) = \mathbf{v}^\top \Sigma \mathbf{v}$. With a Lagrange multiplier for $\|\mathbf{v}\| = 1$: $\Sigma \mathbf{v} = \lambda \mathbf{v}$, so the best $\mathbf{v}$ is the top eigenvector and the variance captured is λ . The next component maximises variance subject to being orthogonal to the first, and so on. **SVD link.** $X = U \Sigma_s V^\top$ gives $X^\top X = V \Sigma_s^2 V^\top$: eigenvectors V , eigenvalues σ_i^2 . So compute SVD of the centred X and read off V (loadings), $\sigma_i^2/(n - 1)$ (variances) and $U \Sigma_s$ (scores). SVD is numerically better than forming $X^\top X$. **Scree plot.** Plot λ_i or explained-variance ratio; choose k at the elbow or by a cumulative threshold.

📊 **Worked example.** Eigenvalues 6, 3, 1: 60%, 30%, 10%; two components keep 90%. Centred data with $n = 101$ and $\sigma_1 = 20$: variance along PC1 = $400/100 = 4$.

💡 **If the interviewer pushes.** “Why centre?” Without centring, the first component points toward the mean rather than along the spread. “Why standardise?” So a feature in rupees does not dominate one in years. “Is PCA feature selection?” No, each component mixes all features.

Q14.2

★ PYQ (INT: t-SNE and perplexity) Why is t-SNE useful for seeing clusters, and what does perplexity control?

🕒 **In one line.** t-SNE places points in 2-D so that points that are neighbours in high dimensions remain neighbours, using a heavy-tailed similarity in 2-D that pulls clusters apart; perplexity sets the effective number of neighbours each point considers, trading local against more global structure.

⚙️ **The full answer.** Perplexity = $2^{H(P_i)}$, where H is the entropy of point i 's neighbour distribution; each point's Gaussian width is tuned so the perplexity matches the chosen value, so dense and sparse regions are treated fairly. Low perplexity (5): fragmented, very local; high (50–100): smoother, more global, small clusters may merge.

💡 **If the interviewer pushes.** “Can I trust distances between clusters?” No. “Why not use it as features?” It is non-parametric, stochastic, and has no out-of-sample mapping; use UMAP or PCA for features.

Q14.3

★ PYQ (INT: PCA vs t-SNE vs LDA) Compare them.

🕒 **In one line.** PCA is linear, unsupervised and preserves global variance; LDA is linear, supervised and maximises class separation (at most $K - 1$ dimensions); t-SNE is non-linear, unsupervised and preserves local neighbourhoods for visualisation.

⚙️ **The full answer.** Use PCA for compression, denoising and decorrelation before modelling; LDA for supervised reduction or a linear classifier; t-SNE/UMAP for looking at cluster structure.

Q14.4

What is the SVD and what do U , Σ and V mean geometrically?

🕒 **In one line.** $A = U\Sigma V^T$: V^T rotates the input onto special orthogonal directions, Σ stretches each by a singular value, and U rotates the result into the output space.

⚙️ **The full answer.** The unit circle maps to an ellipse whose semi-axes have lengths σ_i along u_i . Singular values are always non-negative; the number of non-zero ones is the rank.

Q14.5

★ PYQ (reported linear algebra: rank, null space) What is the rank of a matrix and its null space, and how does SVD reveal both?

🕒 **In one line.** Rank is the number of linearly independent columns (the dimension of the column space); the null space is the set of vectors the matrix sends to zero; in the SVD, rank = number of non-zero singular values, and the right singular vectors with zero singular value span the null space.

⚙️ **The full answer.** Rank-nullity: $\text{rank} + \dim(\text{null space}) = \text{number of columns}$. In regression, a non-trivial null space of X means some combination of features is identically zero (perfect multicollinearity), so coefficients are not unique: adding any null-space vector to $\mathbf{w}$ leaves predictions unchanged.

📊 **Worked example.** $\begin{pmatrix} 1 & 2 \\ 2 & 4 \end{pmatrix}$: rank 1; null space spanned by $(2, -1)$ because $1 \cdot 2 + 2 \cdot (-1) = 0$.

Q14.6

Explain the Eckart–Young theorem and compute the rank-1 approximation error for singular values 5, 3, 1.

🎯 **In one line.** Truncating the SVD to the top k singular values gives the best rank- k approximation; the Frobenius error is $\sqrt{3^2 + 1^2} = 3.16$.

Q14.7

Find the principal components of the covariance matrix $\begin{pmatrix} 2 & 1 \\ 1 & 2 \end{pmatrix}$.

🎯 **In one line.** Eigenvalues 3 and 1 with eigenvectors $(1, 1)/\sqrt{2}$ and $(1, -1)/\sqrt{2}$; PC1 explains 75% of the variance.

Q14.8

🧠 **Think deeper** Can PCA hurt a classifier?

🎯 **In one line.** Yes: PCA keeps high-variance directions, and the direction that separates the classes may have low variance and be discarded; also it can mix informative and noisy features.

⚙️ **The full answer.** The parallel-cigars example: the class difference lies along the narrow axis, PC2, which a “keep 95% variance” rule might drop. Choose k by validation performance, or use supervised reduction (LDA, PLS).

Q14.9

🧠 **Think deeper** Two features are perfectly correlated ($x_2 = 2x_1$). What does PCA give?

🎯 **In one line.** One component with all the variance and a second with eigenvalue exactly zero: the data lie on a line, so the covariance matrix has rank 1.

⚙️ **The full answer.** The zero eigenvector $(2, -1)/\sqrt{5}$ spans the null space of the centred data: a combination that is always 0.

Q14.10

How does matrix factorisation predict a seeker's interest in a job they have never seen?

🕒 **In one line.** It represents each seeker and job as vectors learned from observed interactions so that dot products (plus biases) reproduce known interests; the dot product of an unseen pair is the prediction.

⚙️ **The full answer.** Similar seekers get similar vectors because they interacted with similar jobs; so a seeker inherits the tastes of their “neighbours” in latent space. Worked: $3.5 + 0.2 - 0.3 + 4 = 7.4$.

Q14.11

Why not run ordinary SVD on a sparse ratings matrix?

🕒 **In one line.** SVD requires every entry, and filling unknown ratings with zeros treats “not seen” as “hated”, biasing all factors toward predicting zero; matrix factorisation fits only observed entries with regularisation.

Q14.12

ALS vs SGD for matrix factorisation: which and why?

🕒 **In one line.** ALS alternates exact least-squares solves for users and items, which is easily parallelised and handles implicit-feedback weighting efficiently; SGD is simpler, works well for explicit ratings, and handles extra terms easily but needs learning-rate tuning.

🧮 **Worked example.** ALS with $k = 1$: a user rated items with $q = 1$ and $q = 2$ as 3 and 4; with $\lambda = 1$: $p = \frac{1 \cdot 3 + 2 \cdot 4}{1 + 4 + 1} = 11/6 = 1.83$.

Q14.13

💡 **Think deeper** **How do you recommend a job posted five minutes ago, with zero interactions?**

🕒 **In one line.** With content-based signals: embed the job from its title, skills, location and salary (text encoders, taxonomy features), match it to seekers whose profiles or past applies are similar, and use exploration to collect early feedback; hybrid models map content features into the same latent space as collaborative factors.

⚙️ **The full answer.** Two-tower models learn a job tower from content features, so new jobs get vectors immediately. Also use popularity of similar jobs as a prior and bandit-style exploration so new jobs get some impressions.

Q14.14

What is NMF and how does it differ from SVD?

🎯 **In one line.** NMF factorises a non-negative matrix into non-negative factors, giving additive, parts-based, interpretable components (like topics); SVD factors can be negative, are orthogonal and optimal in squared error, but are harder to interpret.

Q14.15

How many parameters does a factorisation with 10,000 users, 5,000 jobs and $k = 32$ have (no biases)?

🎯 **In one line.** $(10,000 + 5,000) \times 32 = 480,000$.

Q14.16

🧠 **Think deeper** LoRA fine-tunes a large model by learning $\Delta W = BA$ with rank r . **How is this a matrix factorisation, and how many parameters does it save for a 4096×4096 weight with $r = 8$?**

🎯 **In one line.** It constrains the weight update to be low-rank, the product of a 4096×8 and an 8×4096 matrix: 65,536 parameters instead of 16,777,216, a 256-fold saving.

⚙️ **The full answer.** The hypothesis is that the change needed to adapt a pretrained model has low intrinsic rank. This is exactly a truncated factorisation of the update.

Q14.17

What is the condition number of a matrix and why does it matter in ML?

🎯 **In one line.** The ratio of largest to smallest singular value; a large condition number means the problem is ill-conditioned: small changes in data cause large changes in the solution, and gradient descent converges slowly.

📊 **Worked example.** Singular values 100, 10, 0.1: condition number 1,000. Ridge improves it to $(\sigma_1^2 + \lambda)/(\sigma_{\min}^2 + \lambda)$ for the normal equations.

Q14.18

🧠 **Think deeper** Is PCA the same as a linear autoencoder?

🎯 **In one line.** A linear autoencoder with a k -unit bottleneck trained with squared error learns the same subspace as the top k principal components, though its weights need not be orthogonal or ordered.

⚙️ **The full answer.** Non-linear autoencoders generalise PCA to curved manifolds (Chapter 22).

Q14.19

What is Latent Semantic Analysis?

🎯 **In one line.** Truncated SVD of a document–term (TF-IDF) matrix; documents and words are represented in a low-dimensional “concept” space where synonyms that co-occur with similar words end up close.

⚙️ **The full answer.** Use TruncatedSVD rather than PCA because it works on sparse matrices without centring (centring would make them dense).

CHAPTER 15

Density Estimation and Clustering: Histograms, KDE, K-Means, GMM and EM

🎥 **Companion video for this chapter:** [StatQuest: K-means clustering](#)

The syllabus line is “Density Estimation (KDE, K-Means etc)”. These methods answer: *where does the data live, and in what groups?* They are unsupervised, so there is no label to check against, which is why interviewers probe whether you know how to choose K , what the algorithms assume, and when they fail.

15.1 What is density estimation?

A **probability density** $p(x)$ describes how data are spread: high where values are common, low where rare; it integrates to 1. **Density estimation** means estimating $p(x)$ from samples. Uses: understanding a variable’s shape (bimodal salaries), anomaly detection (low-density points are unusual), generative modelling, and as a building block (Naive Bayes, Bayes classifiers).

- **Parametric:** assume a family (Gaussian) and estimate a few parameters (mean, variance). Efficient if the family is right, badly wrong otherwise.
- **Non-parametric:** let the data shape the estimate: histograms, KDE, k -NN density.
- **Mixtures** sit between: a weighted sum of simple components (GMM).

15.2 Histograms

Split the range into bins of width h ; the density in a bin is $\frac{\text{count}}{n \cdot h}$ (so areas sum to 1).

📊 **Small example: histogram density**

100 points, 20 fall in a bin of width 0.5: density = $20 / (100 \times 0.5) = 0.4$.

Weaknesses: blocky (discontinuous), depends on bin width and on where bins start, scales badly to many dimensions (bins multiply).

15.3 Kernel Density Estimation (KDE)

15.3.1 The idea

Instead of counting points in fixed bins, put a small smooth “bump” (a **kernel**, e.g. a tiny Gaussian) on top of every data point and add the bumps up:

$$\hat{p}(x) = \frac{1}{nh} \sum_{i=1}^n K\left(\frac{x - x_i}{h}\right).$$

🎥 **Stuck on this page?** Watch: [StatQuest: K-means clustering](#)

h is the **bandwidth** (width of each bump). K integrates to 1, so $\hat{p}$ does too.

Small example: KDE by hand

Box kernel $K(u) = \frac{1}{2}$ for $|u| \leq 1$, $h = 1$, data $\{1, 2, 4\}$. At $x = 2.5$: $|2.5 - 1|/1 = 1.5$ (out), 0.5 (in), 1.5 (out). $\hat{p} = \frac{1}{3 \cdot 1} \cdot \frac{1}{2} = 0.167$. Gaussian kernel, $h = 1$, data $\{0, 2\}$, at $x = 1$: both points are 1 away: $\hat{p} = \frac{1}{2}(\phi(1) + \phi(1)) = 0.242$.

15.3.2 Bandwidth is everything

Small h : spiky estimate, one bump per point, high variance (overfitting). Large h : over-smoothed, real modes merged, high bias. The kernel shape matters much less. Rules of thumb: **Silverman** $h = 1.06 \hat{\sigma} n^{-1/5}$ (assumes roughly Gaussian data; with $\hat{\sigma} = 2$, $n = 32$: $h = 1.06 \cdot 2 \cdot 0.5 = 1.06$); or choose h by cross-validated likelihood.

15.3.3 Limits

Prediction cost $O(n)$ per point; boundary bias (density leaks below 0 for salaries; use log scale or reflection); curse of dimensionality (needs exponentially more data as dimension grows).

15.4 K-Means clustering

15.4.1 The algorithm (Lloyd's)

1. Choose K and initial centroids.
2. **Assign** each point to its nearest centroid.
3. **Update** each centroid to the mean of its assigned points.
4. Repeat 2–3 until assignments stop changing.

It minimises the **inertia** (within-cluster sum of squares) $\sum_i \|x_i - \mu_{c(i)}\|^2$. Each step cannot increase it, so it converges, but only to a local minimum that depends on initialisation.

Small example: K-Means by hand

1-D points $\{1, 2, 6, 7\}$, $K = 2$, initial centroids 1 and 2. Assign: $1 \rightarrow c1$; $2 \rightarrow c2$; $6, 7 \rightarrow c2$. Update: $c1 = 1$, $c2 = 5$. Assign: $1 \rightarrow c1$; 2 is 1 from $c1$ and 3 from $c2 \rightarrow c1$; $6, 7 \rightarrow c2$. Update: $c1 = 1.5$, $c2 = 6.5$. Assign again: unchanged. Inertia = $0.25 + 0.25 + 0.25 + 0.25 = 1$.

15.4.2 Initialisation: k-means++

Pick the first centroid at random; pick each next centroid with probability proportional to the squared distance to the nearest existing centroid. Spreads centroids out, gives provably better starts. Also run several restarts (`n_init`) and keep the lowest inertia.

Small example: k-means++ probabilities

Remaining points at distances 1, 2, 3 from the first centroid: weights 1, 4, 9; probabilities $1/14$, $4/14$, $9/14$.

15.4.3 Choosing K

- **Elbow:** plot inertia vs K ; it always decreases; pick where improvement flattens.
- **Silhouette:** for each point, a = mean distance to its own cluster, b = mean distance to the nearest other cluster; $s = \frac{b-a}{\max(a,b)} \in [-1, 1]$. Average over points; higher is better. Example: $a = 2$, $b = 5$: $s = 0.6$.
- **Gap statistic:** compare inertia with that expected under a structureless reference distribution.
- **Business usefulness:** segments must be actionable.

15.4.4 Assumptions and failure modes

K-Means assumes roughly spherical, similar-sized, similar-density clusters (it draws straight boundaries halfway between centroids: a Voronoi partition). It fails on elongated, nested or very unequal clusters; it is sensitive to outliers (means) and to feature scaling (Euclidean distance); categorical data need K-modes. Cost per iteration $O(nKd)$.

15.5 Gaussian Mixture Models and EM

15.5.1 The model

Assume each point was generated by first picking a component k with probability π_k , then drawing from $\mathcal{N}(\mu_k, \Sigma_k)$. The density is

$$p(\mathbf{x}) = \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x} \mid \mu_k, \Sigma_k).$$

A GMM is both a **density estimator** (flexible shapes from Gaussian bumps) and a **soft clustering** method (each point gets a probability of belonging to each component).

15.5.2 EM, step by step

We do not know which component generated each point (a hidden variable), so maximum likelihood has no closed form. EM alternates:

- **E-step: responsibilities** $\gamma_{ik} = \frac{\pi_k \mathcal{N}(\mathbf{x}_i \mid \mu_k, \Sigma_k)}{\sum_j \pi_j \mathcal{N}(\mathbf{x}_i \mid \mu_j, \Sigma_j)}$: how much component k “owns” point i (Bayes’ theorem).
- **M-step:** re-estimate with responsibilities as soft counts: $N_k = \sum_i \gamma_{ik}$, $\pi_k = N_k/n$, $\mu_k = \frac{1}{N_k} \sum_i \gamma_{ik} \mathbf{x}_i$, $\Sigma_k = \frac{1}{N_k} \sum_i \gamma_{ik} (\mathbf{x}_i - \mu_k)(\mathbf{x}_i - \mu_k)^\top$.

Each iteration never decreases the likelihood; converges to a local optimum.

Small example: one E-step and one M-step

1-D, $\pi = (0.5, 0.5)$, $\mu = (0, 4)$, $\sigma = 1$. Point $x = 1$: $\mathcal{N}(1 \mid 0, 1) = 0.242$, $\mathcal{N}(1 \mid 4, 1) = \phi(3) = 0.0044$. $\gamma_1 = 0.242/0.2464 = 0.982$. M-step example: points 0, 2, 4 with responsibilities 1, 0.5, 0 for component 1: $N_1 = 1.5$, $\mu_1 = (0 + 1 + 0)/1.5 = 0.667$, $\pi_1 = 0.5$.

15.5.3 K-Means is a special case

With equal, spherical covariances $\sigma^2 I$ shrinking to 0, responsibilities become 0/1 (hard assignments to the nearest mean) and EM becomes K-Means. So K-Means is “hard-assignment EM for a GMM with identical spherical clusters”.

15.5.4 Practical issues

Singularity: a component can collapse onto one point (variance $\rightarrow 0$, likelihood $\rightarrow \infty$); fix with a variance floor (reg_covar) or priors. **Choosing K:** BIC or AIC (likelihood penalised by parameter count). **Covariance types:** full, diagonal, spherical, tied (fewer parameters). Parameters for $K = 3$, $d = 2$, full covariance: means 6, covariances $3 \times 3 = 9$, weights 2: total 17.

15.6 Other clustering methods worth knowing

- **DBSCAN:** density-based. A point is a *core point* if at least minPts points lie within radius ϵ . Clusters are connected core points plus their border points; everything else is noise. Finds arbitrary shapes and outliers; no K needed; struggles with varying densities. (HDBSCAN handles varying densities.)
- **Hierarchical (agglomerative):** start with every point alone; repeatedly merge the closest pair. Linkage: single (min distance; chaining), complete (max), average, **Ward** (merge that least increases within-cluster variance). Produces a dendrogram; cut at any level.
- **Mean shift:** move each point uphill on a KDE to a mode; points ending at the same mode form a cluster.
- **Spectral clustering:** cluster using eigenvectors of a similarity graph; handles non-convex shapes.

In practice: clustering and densities at InfoEdge

- **Talent Pulse** cohorts and seeker segmentation (freshers in tier-2 cities looking for remote IT jobs vs mid-career sales): K-Means or GMM on behaviour features, then named by the business.
- **Salary distributions** for a title are often bimodal (freshers and seniors): a GMM or KDE represents them; a single mean misleads.
- **Anomaly detection:** low KDE/GMM density for posting behaviour flags suspicious recruiters.
- **Taxonomy building:** clustering embeddings of 50 lakh raw skill strings to discover synonym groups (“ML”, “machine learning”, “M.L.”), the first step of SilverSkAI-style normalisation.

15.7 Interview questions

Q15.1

Explain K-Means, its objective, and why it always converges.

🕒 **In one line.** Alternate between assigning points to the nearest centroid and moving each centroid to the mean of its points, which minimises the within-cluster sum of squared distances; each step can only lower that objective and there are finitely many assignments, so it converges, though only to a local minimum.

⚙️ **The full answer.** Assignment step: for fixed centroids, nearest-centroid assignment minimises each point's term. Update step: for fixed assignments, the mean minimises the sum of squared distances. Monotone decrease plus finite partitions implies termination.

📊 **Worked example.** The $\{1, 2, 6, 7\}$ example ends with centroids 1.5 and 6.5, inertia 1.

Q15.2

How do you choose K in K-Means?

🕒 **In one line.** Use the elbow of the inertia curve, the average silhouette score, the gap statistic, or stability across runs, and finally whether the segments are actionable for the business.

⚙️ **The full answer.** Inertia always falls with K (with $K = n$ it is 0), so look for diminishing returns. Silhouette $s = (b - a) / \max(a, b)$ rewards tight, well-separated clusters.

Q15.3

🧠 **Think deeper** Is K-Means a special case of a Gaussian mixture model?

🕒 **In one line.** Yes, in the limit: a GMM with equal mixing weights and identical spherical covariances $\sigma^2 I$ with $\sigma \rightarrow 0$ makes responsibilities hard 0/1 assignments to the nearest mean, and EM then reduces exactly to K-Means.

⚙️ **The full answer.** As $\sigma \rightarrow 0$, $\gamma_{ik} \propto \exp(-\|x_i - \mu_k\|^2 / 2\sigma^2)$ is dominated by the nearest mean. The M-step mean update is then the plain cluster mean. GMM generalises K-Means with soft assignments, ellipsoidal shapes, different sizes and a likelihood for model selection.

Q15.4

Explain EM for Gaussian mixtures and compute a responsibility.

🕒 **In one line.** EM alternates computing each point's probability of belonging to each component given current parameters (E-step) and re-estimating weights, means and covariances using those probabilities as soft counts (M-step); for $x = 1$ between components at 0 and 4 with unit variance, the responsibility of the first is 0.982.

⚙️ **The full answer.** EM maximises a lower bound on the log-likelihood; the E-step makes the

bound tight, the M-step maximises it. Each iteration increases (or keeps) the likelihood.

Q15.5

 **Think deeper** What does a GMM do when one component collapses onto a single point, and how do you prevent it?

🎯 **In one line.** Its variance shrinks toward zero and the likelihood goes to infinity, a degenerate solution; prevent it with a minimum variance (regularisation added to the covariance diagonal), priors (MAP-EM), or restarting.

Q15.6

What is KDE, and how does the bandwidth affect it?

🎯 **In one line.** KDE estimates a density by placing a kernel bump on each data point and averaging them; small bandwidth gives a spiky, high-variance estimate, large bandwidth an over-smoothed, high-bias one.

📊 **Worked example.** Silverman: $\hat{\sigma} = 2$, $n = 32$: $h \approx 1.06$.

💡 **If the interviewer pushes.** “How do you pick it properly?” Maximise cross-validated log-likelihood of held-out points.

Q15.7

Compute a box-kernel KDE at $x = 2.5$ for data $\{1, 2, 4\}$, $h = 1$.

🎯 **In one line.** Only the point 2 is within distance 1, so $\hat{p} = \frac{1}{3} \cdot \frac{1}{2} = 0.167$.

Q15.8

 **Think deeper** Salaries for “Data Scientist” look bimodal. How would you model the distribution, and why not report the mean?

🎯 **In one line.** Use a two-component GMM (on log salary) or a KDE; the mean falls in the valley between the fresher and senior humps and describes almost nobody.

⚙️ **The full answer.** Better still, condition on experience: the bimodality is probably two populations. Report the distribution or quantiles by experience band.

Q15.9

When does K-Means fail? Give three examples and alternatives.

🎯 **In one line.** Non-spherical shapes (two moons: use DBSCAN/spectral), very different cluster sizes or densities (use GMM or HDBSCAN), and outliers or unscaled features (scale, use K-Medoids or robust methods).

Q15.10

Explain DBSCAN and count the core points for 1-D points 0, 0.5, 1, 5 with $\epsilon = 1$, minPts = 3 (including the point itself).

🕒 **In one line.** DBSCAN grows clusters from dense core points and labels sparse points as noise; here 0 has {0, 0.5, 1} (3: core), 0.5 has 3 (core), 1 has {0, 0.5, 1} (3: core), 5 has 1 (noise): three core points, one cluster and one noise point.

Q15.11

Single vs complete vs Ward linkage in hierarchical clustering.

🕒 **In one line.** Single linkage merges by the closest pair of points and tends to chain clusters together; complete linkage uses the farthest pair and gives compact clusters; Ward merges the pair that least increases within-cluster variance and behaves like K-Means.

Q15.12

Why must features be scaled before K-Means? Which clustering methods do not need it?

🕒 **In one line.** K-Means uses Euclidean distance, so large-range features dominate; methods on a precomputed meaningful distance (e.g. cosine on normalised embeddings) or density on a single scale avoid the issue, but almost all distance-based clustering needs scaling.

Q15.13

🧠 **Think deeper** How do you evaluate clustering when there are no labels?

🕒 **In one line.** With internal metrics (silhouette, Davies–Bouldin, Calinski–Harabasz), stability under resampling, and above all external usefulness: do segments differ on outcomes that matter (apply rate, churn) and can the business act on them?

⚙️ **The full answer.** If some labels exist, use adjusted Rand index or NMI. In a case study, propose an A/B test of segment-specific campaigns.

Q15.14

How would you use density estimation for anomaly detection?

🕒 **In one line.** Fit a density (KDE, GMM) on normal behaviour and flag points whose estimated density is below a threshold chosen for an acceptable alert rate.

⚙️ **The full answer.** For recruiters: features like jobs posted per hour, distinct cities, salary range width; a GMM on log-transformed features; alerts on the lowest 0.1% density. Isolation Forest is a non-density alternative.

Q15.15

What is the silhouette of a point with mean intra-cluster distance 2 and mean nearest-cluster distance 5?

🎯 **In one line.** $(5 - 2)/5 = 0.6$.

Q15.16

🧠 **Think deeper** Can you use K-Means for something other than clustering?

🎯 **In one line.** Yes: vector quantisation (compressing data or images to K colours), creating features (distance to each centroid, cluster ID), initialising GMMs, building inverted-file indexes for approximate nearest neighbour search (IVF in FAISS), and summarising large datasets.

Q15.17

How many free parameters does a GMM with $K = 3$, $d = 2$ and full covariances have?

🎯 **In one line.** 17: $3 \times 2 = 6$ means, $3 \times 3 = 9$ covariance entries (3 per symmetric 2×2), and $3 - 1 = 2$ mixing weights.

Q15.18

K-Means vs GMM: when would you choose each?

🎯 **In one line.** K-Means for speed and simple, round, similar-sized clusters with hard assignments; GMM when clusters are elliptical or of different sizes, when you need soft membership probabilities or a density, and when you want a likelihood-based way (BIC) to choose K .

Part III

Natural Language Processing

CHAPTER 16

Basic NLP: Preprocessing, Bag of Words and TF-IDF

🎥 **Companion video for this chapter:** [Krish Naik: TF-IDF Intuition \(NLP\)](#)

InfoEdge is a text company: resumes, job descriptions, search queries, chat messages, matrimonial bios. The OA has asked about bag of words, tokenisation, stemming and lemmatisation; Round 3 lists “basic text processing”; the PremiumX slide shows hybrid retrieval where the “sparse” half is keyword (TF-IDF/BM25) search. This chapter makes you fluent in classical NLP.

16.1 Why text needs special treatment

Models need numbers; text is a sequence of symbols of variable length, with a huge vocabulary, ambiguity (“Java” the language vs the island), spelling variants (“Bengaluru”, “Bangalore”), abbreviations (“ML”, “Sr.”), and meaning that depends on order (“not good”). Classical NLP turns text into fixed-length vectors through a pipeline:

raw text → clean → tokenise → normalise → vectorise → model.

16.2 Tokenisation

Splitting text into units (**tokens**).

- **Word-level:** split on spaces and punctuation. “Sr. Data Scientist, 5+ yrs” → [Sr, ., Data, Scientist, ,, 5, +, yrs]. Problems: “C++”, “Node.js”, “e-mail”, URLs, Hindi–English code mixing; unknown words at test time (**OOV**, out of vocabulary).
- **Character-level:** no OOV, but long sequences and little meaning per token.
- **Subword** (BPE, WordPiece, SentencePiece/Unigram): frequent words stay whole, rare words split into pieces (“unbelievable” → “un”, “##believ”, “##able”). The standard for modern models: small vocabulary, no OOV.

16.2.1 Byte-pair encoding (BPE) in one example

Start from characters; repeatedly merge the most frequent adjacent pair into a new symbol.

📖 Small example: BPE merges

Corpus: “low” ×5, “lower” ×2, “newest” ×6, “widest” ×3. Count adjacent pairs: “e s” occurs in newest (6) and widest (3) = 9; “s t” also 9; “l o” 7; “w e” 8. Merge “e s” → “es” (first of the ties). Next, “es t” = 9 → “est”. Then “l o” → “lo”. After a few merges, “est” and “low” are single tokens, and an unseen word like “lowest” is “low” + “est”.

16.3 Normalisation

- **Lower-casing** (careful: “US” vs “us”, “IT” department vs “it”).
- **Removing punctuation, numbers, HTML**; or keeping them when meaningful (“C++”, “5+ yrs”).
- **Stop words**: very common words (the, is, and) carry little topical meaning; removing them shrinks the vocabulary. Dangerous for sentiment (“not”) and phrase meaning (“to be or not to be”).
- **Stemming**: chop suffixes with rules (Porter, Snowball): “running”, “runs” → “run”; “studies” → “studi”. Fast, crude: *over-stemming* merges unrelated words (“university”, “universe” → “univers”); *under-stemming* misses related ones.
- **Lemmatisation**: map to the dictionary form (lemma) using vocabulary and part of speech: “better” → “good”, “was” → “be”, “meeting” → “meet” (verb) or “meeting” (noun). Accurate, slower, needs POS context.
- **Domain normalisation**: synonyms and spelling variants (“Bangalore” = “Bengaluru”; “ML” = “machine learning”): the SilverSkAI taxonomy problem.

16.4 Bag of Words (BoW)

Fix a **vocabulary** (all words seen in training, perhaps the top 50,000). Represent each document by a vector of word counts, ignoring order.

Small example: BoW vectors

Vocabulary [ml, python, sql]. “python ml python” → (1, 2, 0). “sql and ml” → (1, 0, 1) (“and” is not in the vocabulary).

Variants: binary presence (0/1); term frequency normalised by document length; **n-grams**: add pairs (bigrams) or triples of adjacent words to capture short phrases (“machine learning”, “not good”). A sentence of 7 tokens has 6 bigrams and 5 trigrams ($n - k + 1$ k -grams).

Limitations: loses word order (“dog bites man” = “man bites dog”), huge sparse vectors, no notion of similarity between words (“ML” and “machine learning” are unrelated columns), common words dominate counts.

Small example: sparsity

10,000 resumes, vocabulary 50,000, about 100 distinct words per resume: non-zero fraction = $100/50,000 = 0.2\%$. Stored as a sparse matrix.

16.5 TF-IDF: down-weighting words everybody uses

Raw counts make “experience”, “team”, “work” look important in every resume. TF-IDF multiplies a word’s frequency in the document by how *rare* it is across documents.

$$\text{tf-idf}(t, d) = \text{tf}(t, d) \times \text{idf}(t), \quad \text{idf}(t) = \log \frac{N}{\text{df}(t)},$$

with N = number of documents and $\text{df}(t)$ = number containing t .

- A word in every document: $\text{idf} = \log 1 = 0$; it contributes nothing.
- A word in one document of 10: $\text{idf} = \log_{10} 10 = 1$.

Small example: TF-IDF of one word

A resume has 100 words; “Kubernetes” appears 3 times: $\text{tf} = 0.03$. It appears in 100 of 10,000 resumes: $\text{idf} = \log_{10} 100 = 2$. $\text{tf-idf} = 0.06$. “Team” appears 5 times ($\text{tf} 0.05$) but in 8,000 resumes: $\text{idf} = \log_{10} 1.25 = 0.097$; $\text{tf-idf} = 0.005$. Kubernetes is 12 times more important for this resume despite appearing less often.

16.5.1 scikit-learn’s version

`TfidfVectorizer` uses raw counts for tf and a *smoothed* $\text{idf} = \ln \frac{1+N}{1+\text{df}} + 1$ (so no word gets 0 and no division by zero), then L2-normalises each document vector. With $N = 3$, $\text{df} = 1$: $\ln(4/2) + 1 = 1.693$; $\text{df} = 3$: $\ln 1 + 1 = 1$. Useful options: `ngram_range=(1,2)`, `min_df` (drop rare typos), `max_df` (drop near-stopwords), `sublinear_tf=True` ($1 + \log \text{tf}$, so 20 mentions are not 20 times one mention), `stop_words`, `max_features`.

16.5.2 Comparing documents: cosine similarity

$$\cos(\mathbf{a}, \mathbf{b}) = \frac{\mathbf{a}^\top \mathbf{b}}{\|\mathbf{a}\| \|\mathbf{b}\|}.$$

It measures the angle, ignoring length, so a long resume and a short job description about the same skills are still similar. Example: $(1, 1, 0)$ and $(1, 0, 1)$: $1/(\sqrt{2}\sqrt{2}) = 0.5$.

16.5.3 BM25: TF-IDF for search

Search engines (Elasticsearch/Lucene, behind keyword retrieval in Resdex) use **BM25**, which (1) saturates term frequency (the 10th occurrence adds little), with parameter k_1 , and (2) normalises by document length relative to the average, with parameter b :

$$\text{score}(q, d) = \sum_{t \in q} \text{idf}(t) \frac{\text{tf}(t, d)(k_1 + 1)}{\text{tf}(t, d) + k_1(1 - b + b \frac{|d|}{\text{avgdl}})}.$$

16.6 Classic text classifiers

On TF-IDF features: Multinomial Naive Bayes (fast baseline), logistic regression and linear SVM (strong baselines, often within a few points of deep models on topic classification). For 500 job categories and millions of titles, TF-IDF (word + character n-grams) + linear SVM is a formidable baseline.

16.7 Other basic NLP vocabulary

- **N-gram language model**: $P(w_t \mid w_{t-1})$ estimated from counts; MLE $\frac{c(w_{t-1}w_t)}{c(w_{t-1})}$; smoothing for unseen pairs.

- **Perplexity:** $2^{\text{cross-entropy}}$; “how many choices the model is confused between”. A model giving every token probability $1/4$ has perplexity 4.
- **POS tagging, NER** (tag “Infosys” as ORG, “Pune” as LOC, “Python” as SKILL in resumes), **dependency parsing**.
- **Edit distance** (Levenshtein): minimum insertions, deletions, substitutions; “kitten” $\rightarrow$ “sitting” = 3. Fuzzy matching of company names.
- **Jaccard similarity:** $|A \cap B| / |A \cup B|$ for skill sets; {python, sql, ml} vs {python, ml, spark, aws}: $2/5 = 0.4$.
- **Regular expressions:** extract emails, phones, years of experience (“(\d+)\.?s*(yrs | years)”).

In practice: classic NLP inside Naukri

- **Query intelligence** (Resdex slide): normalise “sr data scientist blr 5+ yrs” into title = Senior Data Scientist, location = Bengaluru, experience ≥ 5 , using dictionaries, regex, NER and the taxonomy.
- **Sparse retrieval:** BM25 over resumes returns candidates containing the exact skills; the dense (embedding) half catches synonyms.
- **Title/skill normalisation:** character n-gram TF-IDF + cosine nearest neighbour maps “Sr. S/W Engg” to “Senior Software Engineer”.

16.8 Interview questions

Q16.1

★ **PYQ (OA: bag of words)** What is the bag-of-words representation? Give the vector for “python ml python” with vocabulary [ml, python, sql].

🕒 **In one line.** A document is represented by the counts of each vocabulary word, ignoring order and grammar; here (1, 2, 0).

⚙️ **The full answer.** Each vocabulary word is one dimension; most entries are zero. It works surprisingly well for topic classification because topics are signalled by which words appear, not their order.

💬 **If the interviewer pushes.** “Limitations?” Order lost, sparsity, no semantic similarity, sensitivity to vocabulary choice, dominated by frequent words (hence TF-IDF).

Q16.2

★ **PYQ (OA: stemming vs lemmatisation vs tokenisation)** Define each, and say which maps “better” to “good”.

🕒 **In one line.** Tokenisation splits text into units; stemming chops affixes by rules to a crude stem; lemmatisation maps a word to its dictionary form using vocabulary and part of speech, and only lemmatisation maps “better” to “good”.

⚙️ **The full answer.** A stemmer sees no suffix to remove in “better”; a lemmatiser knows it is the

comparative of the adjective “good”. Stemming is fast and language-agnostic-ish; lemmatisation is accurate but needs a dictionary and POS tags (“meeting” as noun vs verb).

Q16.3

★ **PYQ (OA: lemmatisation needs context)** Why does lemmatising “meeting” depend on context?

🕒 **In one line.** Because its lemma depends on its part of speech: as a verb (“I am meeting them”) it is “meet”, as a noun (“the meeting was long”) it stays “meeting”.

Q16.4

Explain TF-IDF from scratch and compute it for a word appearing 3 times in a 100-word resume and in 100 of 10,000 resumes ($\log_{10}$).

🕒 **In one line.** TF-IDF weights a word by how often it appears in this document times how rare it is across documents; here $0.03 \times \log_{10}(100) = 0.06$.

⚙️ **The full answer.** tf rewards words central to this document; idf penalises words that appear everywhere (they do not distinguish documents). A word in every document gets idf 0. The product highlights words that are frequent here and rare elsewhere: the document’s distinctive vocabulary.

💡 **If the interviewer pushes.** “Why log in idf?” Without it, a word in 1 of a million documents would get weight 10^6 vs 1 for a word in half of them; the log compresses this to roughly “information content” ($\text{idf} \approx -\log P(\text{word in doc})$).

Q16.5

★ **PYQ (OA: limitations of bag of words)** Give the limitations of BoW/TF-IDF and what fixes each.

🕒 **In one line.** No word order (use n-grams or sequence models), no synonymy (use embeddings), sparsity and high dimension (use dimensionality reduction or embeddings), OOV words (subword tokenisation), and no context (contextual models like BERT).

Q16.6

★ **PYQ (Round 3: basic text processing)** In pandas, count the top 10 most frequent words in a column of job titles, after lower-casing and removing punctuation.

🕒 **In one line.** Lower-case and strip punctuation with vectorised string methods, split into lists, explode to one word per row, and use `value_counts`.

⚙️ **The full answer.**

```
words = (df["title"].str.lower()
         .str.replace(r"[^a-z0-9+\s]", " ", regex=True) # keep c++, c#
```

```
.str.split()
.explode()
words = words[~words.isin({"and", "the", "of", "for"})]
print(words.value_counts().head(10))
```

Alternative: CountVectorizer then sum columns. Mention keeping “+” and “#” so “c++” and “c#” survive.

Q16.7

Why is cosine similarity preferred to Euclidean distance for TF-IDF document vectors?

🕒 **In one line.** Cosine compares direction (which words, in what proportions) and ignores length, so a long and a short document on the same topic are similar; Euclidean distance is dominated by document length.

⚙️ **The full answer.** After L2-normalising vectors (as TfidfVectorizer does), Euclidean distance and cosine are monotonically related: $\|\mathbf{a} - \mathbf{b}\|^2 = 2 - 2 \cos$.

Q16.8

🧠 **Think deeper** Should you remove stop words? Give a case where it hurts.

🕒 **In one line.** It helps topic classification and search efficiency, but hurts tasks where function words carry meaning: sentiment (“not”), phrase or question matching (“jobs *in* Delhi” vs “jobs *near* Delhi”), and any model that uses word order (BERT should see full text).

Q16.9

What are n-grams and what do they buy you? How many bigrams and trigrams does a 7-token sentence have?

🕒 **In one line.** Contiguous sequences of n tokens that capture local order and short phrases (“machine learning”, “not good”); 6 bigrams and 5 trigrams.

⚙️ **The full answer.** The vocabulary grows fast with n , so prune with `min_df`. Character n-grams (3–5) are robust to typos and morphology.

Q16.10

Explain subword tokenisation (BPE/WordPiece) and why modern models use it.

🕒 **In one line.** It splits words into frequent pieces learned from the corpus, so common words are single tokens and rare or new words are built from pieces; this keeps the vocabulary small (about 30K–50K) with no out-of-vocabulary words.

⚙️ **The full answer.** BPE: start with characters, repeatedly merge the most frequent adjacent pair. WordPiece (BERT): merge the pair that most increases training-data likelihood; continuation pieces are marked “##”.

Q16.11

 **Think deeper** A user searches “ML engineer”; resumes say “machine learning engineer”. Why does TF-IDF search miss them, and how would you fix it?

🎯 **In one line.** TF-IDF matches exact tokens, so “ML” and “machine learning” share nothing; fix with query expansion via a synonym taxonomy, normalisation at index time, or dense embedding retrieval that maps both to nearby vectors, ideally combined in hybrid retrieval.

⚙️ **The full answer.** This is exactly why PremiumX uses hybrid retrieval: sparse for exact skills and rare terms (a precise certification code), dense for meaning. Merge result sets (e.g. reciprocal rank fusion) and re-rank.

Q16.12

What does BM25 improve over TF-IDF?

🎯 **In one line.** Term-frequency saturation (extra occurrences give diminishing returns) and document-length normalisation relative to the average length, both tunable (k_1, b), which makes it the standard keyword ranking function.

Q16.13

Compute $P(\text{NLP} \mid \text{love})$ from the corpus “I love NLP . I love ML .” and explain why smoothing is needed.

🎯 **In one line.** $c(\text{love NLP})/c(\text{love}) = 1/2$; without smoothing any unseen bigram gets probability 0, making whole sentences impossible.

📊 **Worked example.** Laplace: $c(w_1) = 8, c(w_1 w_2) = 0, |V| = 12$: $(0 + 1)/(8 + 12) = 0.05$.

Q16.14

What is perplexity? A model gives each of 5 tokens probability 1/4: what is the perplexity?

🎯 **In one line.** Perplexity is exp of the average negative log-likelihood per token, the effective number of equally likely choices; here 4.

Q16.15

 **Think deeper** Why do TF-IDF + linear models still compete with deep learning on some text tasks?

🎯 **In one line.** For topic-like tasks where the presence of keywords decides the label, a linear model on TF-IDF is near-optimal, fast, and data-efficient; deep models shine when meaning depends on context, order, or paraphrase, or when labelled data are plentiful.

⚙️ **The full answer.** Job-category classification from titles: keywords like “accountant”, “Java”

are decisive. Sentiment with negation, or matching a CV to a JD with different vocabulary, needs context.

Q16.16

How would you detect duplicate job postings?

🕒 **In one line.** Normalise text, represent postings with TF-IDF or MinHash shingles (or embeddings), and find pairs with cosine or Jaccard similarity above a threshold using locality-sensitive hashing to avoid comparing all pairs; add rules on recruiter, company, location and posting time.

Q16.17

Compute the Levenshtein distance between “kitten” and “sitting” and give a use.

🕒 **In one line.** 3 (substitute k→s, e→i, insert g); used for fuzzy matching of company names and correcting typos in queries.

Q16.18

Where must the TF-IDF vocabulary and idf be fitted, and why?

🕒 **In one line.** On the training data only (inside a Pipeline), because idf computed with test documents leaks information about the test set into training.

Q16.19

🧠 **Think deeper** **How would you extract “years of experience” and “skills” from 1 crore unstructured resumes?**

🕒 **In one line.** Combine regular expressions and rules for structured patterns (dates, “5+ years”), dictionary and taxonomy matching for skills with fuzzy matching, and a trained NER/sequence-labelling model (fine-tuned BERT) for the long tail, validated on a hand-labelled sample.

⚙️ **The full answer.** Compute experience from employment date ranges, handling overlaps and “present”. Precision/recall per field on a labelled sample; human review for low-confidence extractions.

Q16.20

What is the Jaccard similarity of {python, sql, ml} and {python, ml, spark, aws}, and when would you use it instead of cosine?

🕒 **In one line.** $2/5 = 0.4$; use Jaccard for sets (presence/absence), cosine for weighted vectors like TF-IDF.

CHAPTER 17

Embeddings: Word2Vec, GloVe, FastText and BERT

🎥 **Companion video for this chapter:** [StatQuest: Word Embedding and Word2Vec, Clearly Explained](#)

The syllabus says “embeddings-based approaches (Word2Vec, FastText, BERT)”. The IIT Roorkee OA had questions on word embeddings; the Tools slide lists BERT, Siamese networks and label attention networks; PremiumX retrieval uses a vector index. Embeddings are how InfoEdge represents skills, titles, resumes and jobs so that “similar” becomes “nearby”.

17.1 From one-hot to dense vectors

One-hot: a vocabulary of 50,000 words; each word is a 50,000-long vector with a single 1. Problems: huge, and every pair of words is equally different (dot product 0): “Python” is as far from “Django” as from “Accounting”.

Embedding: each word is a short dense vector (say 300 numbers) learned so that words used in similar contexts get similar vectors. Similarity is measured by cosine. Dimensions are not individually named, but directions often encode meaning.

💡 **Intuition: the distributional hypothesis**

“You shall know a word by the company it keeps” (Firth). “PyTorch” and “TensorFlow” appear near the same words (“model”, “training”, “GPU”, “deep learning”), so a method that predicts context words will give them similar vectors, even if they never appear together.

17.2 Word2Vec

17.2.1 Two training tasks

Slide a window over text. For a centre word and the words within $\pm m$ positions:

- **Skip-gram:** predict each context word from the centre word.
- **CBOW** (continuous bag of words): predict the centre word from the average of its context vectors.

📖 **Small example: skip-gram training pairs**

Sentence “senior data scientist python pune” (5 words), window 2. Centre “scientist”: pairs (scientist, senior), (scientist, data), (scientist, python), (scientist, pune). Over all centres: $2 + 3 + 4 + 3 + 2 = 14$ pairs.

17.2.2 The model

Each word has two vectors: an input (centre) vector $\mathbf{v}_w$ and an output (context) vector $\mathbf{u}_w$. Score of context c given centre w : $\mathbf{u}_c^\top \mathbf{v}_w$; probability by softmax over the whole vocabulary:

$$P(c | w) = \frac{\exp(\mathbf{u}_c^\top \mathbf{v}_w)}{\sum_{c' \in V} \exp(\mathbf{u}_{c'}^\top \mathbf{v}_w)}.$$

It is a neural network with one linear hidden layer (the embedding lookup, no activation) and a softmax output. The embedding table has $|V| \times d$ parameters: 10,000 words $\times$ 300 = 3 million.

17.2.3 Making it fast: negative sampling

The softmax denominator sums over the whole vocabulary for every pair: far too slow. **Negative sampling** turns it into binary classification: is (w, c) a real pair (label 1) or a random fake pair (label 0)? For each real pair, draw k (5–20) random “negative” words. Loss:

$$-\log \sigma(\mathbf{u}_c^\top \mathbf{v}_w) - \sum_{n=1}^k \log \sigma(-\mathbf{u}_n^\top \mathbf{v}_w).$$

Each pair now costs $k + 1$ dot products instead of $|V|$. Negatives are drawn with probability $\propto f^{3/4}$ (frequency to the 3/4), which gives rare words a bit more chance than raw frequency.

Small example: the 3/4 power

Two words with counts 16 and 1: $16^{0.75} = 8$, $1^{0.75} = 1$: probabilities 8/9 and 1/9 instead of 16/17 and 1/17.

Subsampling frequent words: each occurrence of word w is dropped with probability $1 - \sqrt{t/f(w)}$ ($t \approx 10^{-5}$), so “the” does not flood training. For $f = 0.01$: $\sqrt{0.001} = 0.032$, so 97% of occurrences are dropped.

17.2.4 Famous property: analogies

$\mathbf{v}_{king} - \mathbf{v}_{man} + \mathbf{v}_{woman} \approx \mathbf{v}_{queen}$. Relationships become directions. Toy 2-D: king (3, 4), man (2, 1), woman (2, 3): result (3, 6); the nearest word to it would be “queen”. In job data: $\mathbf{v}_{Java\ developer} - \mathbf{v}_{Java} + \mathbf{v}_{Python} \approx \mathbf{v}_{Python\ developer}$.

17.2.5 Skip-gram vs CBOW

Skip-gram: better for rare words and small data (each occurrence produces several training pairs); slower. CBOW: faster, smoother for frequent words (averaging context).

17.3 GloVe

Global Vectors start from a word–word co-occurrence matrix X (how often j appears near i in the whole corpus) and fit vectors so that $\mathbf{w}_i^\top \tilde{\mathbf{w}}_j + b_i + \tilde{b}_j \approx \log X_{ij}$, with a weighting that down-weights rare and caps very frequent pairs. It is a weighted matrix factorisation of log co-occurrences; Word2Vec does something similar implicitly (skip-gram with negative sampling factorises a shifted PMI matrix).

17.4 FastText: embeddings from character n-grams

Represent each word as the sum of vectors of its character n-grams (3–6 characters, with boundary markers) plus the word itself. “where” with $n = 3$: <wh, whe, her, ere, re> (5 trigrams) plus <where>.

- **OOV words get vectors:** a typo “pythn” or a new skill “langchainx” shares n-grams with known words.
- **Morphology:** “developer”, “development”, “developing” share n-grams.
- Great for noisy user-generated text (resumes, queries) and morphologically rich languages (Hindi).

FastText is also a fast text classifier: average word/n-gram embeddings, then a linear softmax layer.

17.5 Static vs contextual embeddings

Word2Vec, GloVe and FastText give each word *one* vector regardless of context: “Java” (language) and “Java” (island), “bank” (river vs money), “lead” (team lead vs sales lead) get one blended vector. **Contextual** models produce a different vector for each occurrence, computed from the whole sentence: ELMo (bidirectional LSTMs), then BERT (Transformers).

17.6 BERT

17.6.1 Architecture

A stack of Transformer **encoder** layers (Chapter 21). BERT-base: 12 layers, hidden size 768, 12 attention heads (64 dimensions each), about 110M parameters; BERT-large: 24 layers, 1024 hidden, 16 heads, 340M. WordPiece vocabulary of about 30,000; maximum 512 tokens. Input embedding of each token = token embedding + position embedding + segment embedding (sentence A or B). Special tokens: [CLS] at the start (its final vector summarises the sequence for classification), [SEP] between/after sentences, [MASK].

Small example: BERT sizes

Token embedding table: $30,522 \times 768 = 23.4\text{M}$ parameters. Per-head dimension: $768/12 = 64$.

17.6.2 Pre-training objectives

- **Masked language modelling (MLM):** choose 15% of tokens; of those, replace 80% with [MASK], 10% with a random token, keep 10% unchanged; predict the original token. Because the model sees both left and right context, it is **bidirectional**. The 80/10/10 split exists because [MASK] never appears at fine-tuning time; mixing in random and unchanged tokens forces the model to build good representations for every token, not just masked ones. In a 512-token input, about 77 tokens are selected and about 61 actually replaced by [MASK].
- **Next sentence prediction (NSP):** is sentence B the real next sentence? Later shown to be of little use (RoBERTa dropped it).

17.6.3 Fine-tuning

Add a small task head on top and train the whole network briefly (2–4 epochs, learning rate $2-5 \times 10^{-5}$): [CLS] → linear layer for classification (resume → job function); token vectors → per-token labels for NER (skill extraction); two segments for pair tasks (does this CV match this JD?).

17.6.4 The family

RoBERTa (more data, longer training, dynamic masking, no NSP); DistilBERT (40% smaller, 60% faster, about 97% of performance via distillation); ALBERT (parameter sharing); domain models (a “Naukri-BERT” further pre-trained on resumes); multilingual BERT/IndicBERT. GPT models are decoder-only and left-to-right (generation), BERT is encoder-only (understanding).

17.7 Sentence and document embeddings

- **Average of word vectors** (optionally TF-IDF weighted): simple, surprisingly strong baseline; loses order.
- **BERT [CLS] or mean-pooled tokens**: off-the-shelf they are poor for similarity search.
- **Sentence-BERT**: a Siamese (twin) BERT fine-tuned so that cosine similarity of the two pooled outputs matches semantic similarity (contrastive or triplet loss). The standard for semantic search and matching; this is the “Siamese neural networks” on the Tools slide.
- **Bi-encoder vs cross-encoder**: bi-encoders embed query and document separately (fast, index millions, used for retrieval); a cross-encoder feeds both together through BERT (slow, more accurate, used to re-rank the top 100).

In practice: embeddings at InfoEdge

- **Job–CV matching engine**: a two-tower (bi-encoder) model embeds jobs and resumes; ANN search retrieves candidates; a cross-encoder or GBM re-ranks.
- **PremiumX hybrid retrieval**: vector index (dense) + keyword index (sparse), merged, then LLM-based ranking features.
- **Taxonomy (SilverSkAI)**: FastText/BERT embeddings of 50 lakh raw skill strings, nearest-neighbour to 70K canonical skills.
- **Label attention networks** (Tools slide): for multi-label tagging (which of thousands of skills a JD needs), each label gets its own attention over the text’s token embeddings.

17.8 Interview questions

Q17.1

★ **PYQ (OA (IIT Roorkee): word embeddings)** What are word embeddings and why are they better than one-hot vectors?

🕒 **In one line.** Dense, low-dimensional vectors learned from context so that words used similarly are close; unlike one-hot vectors they are compact, encode similarity, and generalise from one

word to its neighbours.

⚙️ **The full answer.** One-hot: dimension $|V|$, all pairs orthogonal, nothing learned about “Python” helps “Django”. Embeddings: 100–300 dimensions, cosine similarity reflects meaning, and a model trained with “Python” examples transfers to related words. They are learned by predicting context (Word2Vec), from co-occurrence statistics (GloVe), or inside a task model.

Q17.2

★ **PYQ (OA (IIT Roorkee): CBOW vs skip-gram) Explain CBOW and skip-gram and when to use each.**

🎯 **In one line.** CBOW predicts the centre word from its averaged context; skip-gram predicts each context word from the centre word; skip-gram is better for rare words and smaller corpora, CBOW is faster and good for frequent words.

⚙️ **The full answer.** Skip-gram generates $2m$ training pairs per position, giving rare words many updates; CBOW makes one prediction per position from an average, which smooths over context noise.

Q17.3

★ **PYQ (INT: BERT basics) Explain BERT: architecture, pre-training and fine-tuning.**

🎯 **In one line.** BERT is a stack of Transformer encoders pre-trained on large text with masked language modelling (and originally next-sentence prediction), giving bidirectional contextual token vectors; you fine-tune it end to end with a small task head.

⚙️ **The full answer.** **Architecture.** WordPiece tokens with token, position and segment embeddings; 12 (base) or 24 (large) encoder layers of multi-head self-attention and feed-forward sublayers with residuals and LayerNorm. **MLM.** 15% of tokens selected; 80% → [MASK], 10% random, 10% unchanged; predict the originals. Because attention is not masked, every token sees left and right context: bidirectional. **Fine-tuning.** [CLS] vector into a classifier; token vectors for tagging; sentence pairs for matching. Small learning rate, few epochs.

💡 **If the interviewer pushes.** “Why not just train left-to-right?” Many understanding tasks benefit from both sides of context (“the bank of the river”); MLM allows it without the model trivially seeing the word it predicts. “BERT vs GPT?” Encoder vs decoder; understanding vs generation; MLM vs next-token.

Q17.4

★ **PYQ (INT: Word2Vec vs FastText vs BERT) Compare them.**

🎯 **In one line.** Word2Vec learns one static vector per word from context prediction; FastText does the same but builds vectors from character n-grams, so it handles misspellings and unseen words; BERT produces contextual vectors per token occurrence using a deep Transformer, so the same word gets different vectors in different sentences.

⚙️ **The full answer.** Cost and use: Word2Vec/FastText are tiny and fast (good features,

similarity of skills); BERT is heavy but captures context (“lead” in “team lead” vs “sales lead”). For OOV: Word2Vec none, FastText yes, BERT via subwords.

Q17.5

Why is the full softmax slow in Word2Vec and how does negative sampling fix it?

🕒 **In one line.** The softmax denominator sums over the entire vocabulary for every training pair; negative sampling replaces it with a binary task on the true pair and k random negatives, costing $k + 1$ dot products instead of $|V|$.

⚙️ **The full answer.** With $|V| = 10^6$ and $k = 5$, that is a 160,000-fold reduction per pair. Hierarchical softmax (a binary tree, $\log_2 |V|$ steps) is the other option.

Q17.6

Count skip-gram pairs for a 5-word sentence with window 2, and FastText trigrams for “where”.

🕒 **In one line.** 14 pairs ($2 + 3 + 4 + 3 + 2$); 5 trigrams: <wh, whe, her, ere, re>.

Q17.7

🧠 **Think deeper** Why does $\text{king} - \text{man} + \text{woman} \approx \text{queen}$ work?

🕒 **In one line.** Because the training objective makes vector differences encode consistent changes in context: words that differ only by gender co-occur with similar words except for gendered ones, so their difference vectors align, and adding the gender direction to “king” lands near “queen”.

⚙️ **The full answer.** Formally, skip-gram with negative sampling approximately factorises a PMI matrix; log co-occurrence ratios that are constant across pairs (king:queen :: man:woman) become constant vector offsets. Analogies are not perfect, and embeddings also encode social biases (“engineer” closer to “he”), which matters for fair candidate ranking.

Q17.8

🧠 **Think deeper** Static embeddings give “Java” one vector. When does that hurt a job portal and what do you do?

🕒 **In one line.** When a word has several senses or the sense depends on context: “Java” (language vs coffee/island), “lead” (team lead vs sales lead), “Spark” (Apache Spark vs “spark plug” in automotive jobs); contextual embeddings (BERT) or phrase-level tokens and taxonomy disambiguation fix it.

Q17.9

Why does BERT mask with 80/10/10 rather than always using [MASK]?

🕒 **In one line.** Because [MASK] never appears in real fine-tuning text; replacing some selected tokens with random words or leaving them unchanged forces the model to keep a good representation for every token and reduces the pre-train/fine-tune mismatch.

Q17.10

What are the dimensions of BERT-base, and how many parameters are in its token embedding table?

🕒 **In one line.** 12 layers, hidden 768, 12 heads of 64, about 110M parameters; token embeddings $30,522 \times 768 \approx 23.4\text{M}$.

Q17.11

💡 **Think deeper** You need to search 50 lakh resumes semantically for each recruiter query in under 100 ms. Bi-encoder or cross-encoder?

🕒 **In one line.** A bi-encoder for retrieval, because resume vectors are pre-computed and indexed, so a query needs one encoding plus an ANN lookup; then a cross-encoder (or GBM) to re-rank the top 100–200, where accuracy matters and the cost is affordable.

⚙️ **The full answer.** A cross-encoder needs a full BERT pass per (query, resume) pair: 50 lakh passes per query is impossible. This retrieve-then-rank cascade is exactly the Resdex “retrieval + ranking” design.

Q17.12

What is a Siamese network, and how would you train one for job–CV matching?

🕒 **In one line.** Two identical encoders with shared weights map a job and a CV into the same vector space; trained with a contrastive, triplet or in-batch softmax loss so that matched pairs (applied and shortlisted) are close and unmatched pairs are far.

⚙️ **The full answer.** Positives: recruiter shortlisted the CV for the job. Negatives: random CVs (easy), CVs shown but not shortlisted (hard negatives). Triplet loss $\max(0, d(a, p) - d(a, n) + \text{margin})$. In-batch negatives (every other job in the batch) make training efficient. Evaluate with recall@k.

Q17.13

How would you build a sentence embedding from Word2Vec vectors, and what are its weaknesses?

🕒 **In one line.** Average the word vectors (optionally weighted by TF-IDF, removing the top

principal component); it ignores word order and negation and is dominated by frequent words, but is a fast, decent baseline.

Q17.14

 **Think deeper** How does FastText give a vector to a misspelled skill like “kubernets”?

🕒 **In one line.** By summing the vectors of its character n-grams, most of which (“kub”, “ube”, “ber”, “ern”, “rne”) are shared with “kubernetes”, so the result lands near the correct word.

Q17.15

What are typical embedding dimensions, and how do you choose one?

🕒 **In one line.** 100–300 for Word2Vec/FastText, 384–1024 for sentence Transformers; choose by downstream validation and the cost of storage and search (a vector index of 50 lakh $\times$ 768 floats is about 15 GB).

Q17.16

How do you evaluate embeddings?

🕒 **In one line.** Intrinsically (word similarity correlation with human scores, analogy accuracy, nearest-neighbour inspection) and extrinsically (performance of downstream tasks such as retrieval recall@k or classification F1); extrinsic evaluation is what matters.

Q17.17

 **Think deeper** Would you fine-tune BERT or use it as a frozen feature extractor for classifying 2 lakh resumes into 40 job functions?

🕒 **In one line.** Fine-tune (or fine-tune a distilled version), because with 2 lakh labelled examples end-to-end training adapts representations to resume language and usually gives clearly better accuracy; use frozen features only with very little data, tight compute, or many tasks sharing one embedding.

⚙️ **The full answer.** Practical: DistilBERT or a domain-adapted model, max length 256–512 (or chunk), learning rate 3×10^{-5} , 2–3 epochs, class weights for rare functions, macro F1 for evaluation. Compare against TF-IDF + linear SVM baseline.

Q17.18

What did RoBERTa and DistilBERT change?

🕒 **In one line.** RoBERTa trained longer on more data with bigger batches, dynamic masking and no NSP, improving accuracy; DistilBERT used knowledge distillation to produce a model 40% smaller and 60% faster with about 97% of BERT’s performance.

Q17.19

What is a label attention network (on InfoEdge's Tools slide), and why would InfoEdge use one?

🎯 **In one line.** A multi-label text classifier where each label has its own attention vector over the token representations, so each label looks at the parts of the text relevant to it; useful for tagging job descriptions or resumes with many skills from a large taxonomy.

⚙️ **The full answer.** A single pooled vector must serve thousands of labels; per-label attention lets "Kubernetes" focus on the DevOps paragraph and "Negotiation" on the sales paragraph, and it gives interpretable highlights ("search highlights on the profile" in PremiumX).

Part IV

Deep Learning

CHAPTER 18

Neural Network Foundations: Backpropagation, Activations, Initialisation and Regularisation

🎥 **Companion video for this chapter:** [3Blue1Brown: But what is a neural network? \(Deep Learning, Chapter 1\)](#)

Deep learning was “the most drilled section” for the InfoEdge hire: vanishing and exploding gradients, weight decay, how a network learns (backprop), preventing overfitting, Xavier vs He initialisation. The OA asked about the universal approximation theorem, the number of hidden layers in a “deep” network, classification losses, ResNets for vanishing gradients, and RBMs/DBNs. This chapter covers each from scratch, with a backprop computation you can do on paper.

18.1 From a neuron to a network

18.1.1 One neuron

A neuron computes a weighted sum of its inputs plus a bias, then applies a non-linear **activation** function:

$$z = \mathbf{w}^\top \mathbf{x} + b, \quad a = g(z).$$

📊 Small example: one neuron

$\mathbf{x} = (1, 2)$, $\mathbf{w} = (0.5, -1)$, $b = 1$: $z = 0.5 - 2 + 1 = -0.5$. ReLU: $a = \max(0, -0.5) = 0$. Sigmoid: $a = \sigma(-0.5) = 0.378$.

Logistic regression is exactly one neuron with a sigmoid.

18.1.2 Layers

A **layer** is many neurons reading the same inputs: $\mathbf{a}^{(l)} = g(W^{(l)}\mathbf{a}^{(l-1)} + \mathbf{b}^{(l)})$. Stacking layers gives a **multilayer perceptron (MLP)**. The first layer reads raw features; **hidden layers** build intermediate features; the output layer produces the prediction. A dense layer from n_{in} to n_{out} has $(n_{in} + 1)n_{out}$ parameters.

📊 Small example: counting parameters

MLP 784–128–10 (MNIST): $(784 + 1) \cdot 128 + (128 + 1) \cdot 10 = 100,480 + 1,290 = 101,770$.

18.1.3 Why non-linear activations are essential

Without them, layer after layer of $W\mathbf{x} + \mathbf{b}$ collapses into one linear map: $W_2(W_1\mathbf{x} + \mathbf{b}_1) + \mathbf{b}_2 = W'\mathbf{x} + \mathbf{b}'$. A 100-layer linear network is just linear regression. Non-linearities let the network bend; e.g. XOR needs a hidden layer with a non-linearity.

🎥 **Stuck on this page?** Watch: [3Blue1Brown: But what is a neural network? \(Deep Learning, Chapter 1\)](#)

18.1.4 “Deep” and the universal approximation theorem

A network is usually called **deep** when it has more than one hidden layer (often many). The **universal approximation theorem**: a network with a single hidden layer and enough units, using a non-polynomial activation (sigmoid, ReLU), can approximate any continuous function on a bounded region to any accuracy. It says such a network *exists*; it does not say how many units are needed (possibly exponentially many), nor that gradient descent will find the weights, nor that it will generalise. Depth helps because deep networks can represent many functions with exponentially fewer units by reusing intermediate features (edges $\rightarrow$ shapes $\rightarrow$ objects).

18.2 Activation functions

Name	Formula	Derivative	Notes
Sigmoid	$\frac{1}{1+e^{-z}}$	$\sigma(1 - \sigma) \leq 0.25$	outputs in (0, 1); saturates; not zero-centred; use at binary outputs
Tanh	$\frac{e^z - e^{-z}}{e^z + e^{-z}}$	$1 - \tanh^2 \leq 1$	zero-centred; still saturates; RNN/LSTM default
ReLU	$\max(0, z)$	0 or 1	cheap, no saturation for $z > 0$; “dying ReLU” when $z < 0$ always
Leaky ReLU	$\max(\alpha z, z)$	α or 1	fixes dying ReLU; $\alpha = 0.01$
GELU / SiLU	$z\Phi(z) / z\sigma(z)$	smooth	Transformers (BERT, GPT use GELU)
Softmax	$\frac{e^{z_k}}{\sum_j e^{z_j}}$		multiclass output layer

Leaky ReLU at -2 with $\alpha = 0.01$: -0.02 .

18.3 Loss functions

Regression: MSE, MAE, Huber. Binary classification: sigmoid output + binary cross-entropy. Multiclass: softmax + categorical cross-entropy. Multi-label: one sigmoid per label + BCE. Margin: hinge. Ranking/similarity: contrastive, triplet. The output activation and loss are chosen together as a likelihood (Chapter 12). Classification losses: cross-entropy / log loss, hinge, exponential, focal. MSE, MAE, Huber are regression losses.

18.4 Training: forward pass, loss, backward pass

1. **Forward**: compute activations layer by layer and the loss.
2. **Backward (backpropagation)**: compute the gradient of the loss with respect to every weight using the chain rule, from the output back to the input.
3. **Update**: $W \leftarrow W - \eta \partial L / \partial W$ (or Adam).

18.4.1 The chain rule, in words

If L depends on a , which depends on z , which depends on w , then a tiny change in w changes z by $\frac{\partial z}{\partial w}$ times as much, which changes a by $\frac{\partial a}{\partial z}$ times that, which changes L by $\frac{\partial L}{\partial a}$ times that:

$$\frac{\partial L}{\partial w} = \frac{\partial L}{\partial a} \cdot \frac{\partial a}{\partial z} \cdot \frac{\partial z}{\partial w}.$$

Backprop applies this systematically, *reusing* the shared factors: the gradient flowing into layer l is computed once from the gradient of layer $l + 1$, so the cost of the backward pass is about the same as the forward pass (not one pass per weight).

18.4.2 Backprop for a dense layer

Let $\delta^{(l)} = \partial L / \partial \mathbf{z}^{(l)}$ (the “error” at layer l ’s pre-activations). Then

$$\frac{\partial L}{\partial \mathbf{W}^{(l)}} = \delta^{(l)} \mathbf{a}^{(l-1)\top}, \quad \frac{\partial L}{\partial \mathbf{b}^{(l)}} = \delta^{(l)}, \quad \delta^{(l-1)} = (\mathbf{W}^{(l)\top} \delta^{(l)}) \odot g'(\mathbf{z}^{(l-1)}).$$

At the output with softmax/sigmoid + cross-entropy, $\delta = \hat{\mathbf{y}} - \mathbf{y}$.

Small example: full backprop by hand

Network: input $x = 1$; hidden $h = \sigma(w_1 x)$ with $w_1 = 0.5$; output $\hat{y} = w_2 h$ with $w_2 = 2$; loss $L = \frac{1}{2}(\hat{y} - y)^2$ with $y = 1$. **Forward:** $z_1 = 0.5$, $h = \sigma(0.5) = 0.622$, $\hat{y} = 1.245$, $L = \frac{1}{2}(0.245)^2 = 0.030$. **Backward:** $\partial L / \partial \hat{y} = 0.245$. $\partial L / \partial w_2 = 0.245 \cdot h = 0.152$. $\partial L / \partial h = 0.245 \cdot w_2 = 0.490$. $\partial h / \partial z_1 = h(1 - h) = 0.235$. $\partial L / \partial z_1 = 0.115$. $\partial L / \partial w_1 = 0.115 \cdot x = 0.115$. **Update** ($\eta = 0.1$): $w_2 = 2 - 0.0152 = 1.985$; $w_1 = 0.5 - 0.0115 = 0.4885$. Both moves reduce $\hat{y}$ toward 1.

18.5 Vanishing and exploding gradients

18.5.1 Why they happen

The gradient at an early layer is a *product* of many factors, one per later layer: each factor involves a weight matrix and an activation derivative.

- If the factors are typically less than 1 in size, the product shrinks exponentially: **vanishing gradients**; early layers barely learn. Sigmoid’s derivative is at most 0.25: through 8 sigmoid layers the factor is at most $0.25^8 \approx 1.5 \times 10^{-5}$.
- If typically greater than 1, it grows exponentially: **exploding gradients**; updates become huge, loss jumps to NaN. A factor of 1.5 per step over 20 RNN time steps: $1.5^{20} \approx 3,325$.

RNNs are the extreme case (the same weight matrix multiplied at every time step, Chapter 20).

18.5.2 Fixes

Fix	Why it works
ReLU-family activations	derivative 1 on the active side: no shrinking from activations
Careful initialisation (Xavier, He)	keeps signal and gradient variance roughly constant across layers
Batch / Layer normalisation	keeps pre-activations in a well-scaled range, out of saturation
Residual (skip) connections	gradient flows through the identity path: $\partial(x + F(x))/\partial x = 1 + F'$
LSTM/GRU gates (RNNs)	additive cell-state path with forget gate near 1
Gradient clipping	caps exploding gradients (by norm)
Smaller learning rate, warmup	avoids large unstable steps

18.6 Weight initialisation: Xavier and He

18.6.1 Why not zeros or a constant?

If all weights in a layer start equal, every neuron computes the same output and receives the same gradient: they stay identical forever (**symmetry**). Random initialisation breaks symmetry. But the *scale* matters: too large and activations saturate or explode; too small and the signal dies after a few layers.

18.6.2 The variance argument

For $z = \sum_{i=1}^{n_{in}} w_i x_i$ with independent, zero-mean w_i and x_i : $\text{Var}(z) = n_{in} \text{Var}(w) \text{Var}(x)$. To keep $\text{Var}(z) \approx \text{Var}(x)$ we need $\text{Var}(w) = 1/n_{in}$.

- **Xavier/Glorot** (tanh, sigmoid): balance forward and backward passes: $\text{Var}(w) = \frac{2}{n_{in} + n_{out}}$. Uniform version: $U[-\sqrt{6/(n_{in} + n_{out})}, +\sqrt{6/(n_{in} + n_{out})}]$.
- **He/Kaiming** (ReLU): ReLU zeroes half the inputs, halving the variance, so double it: $\text{Var}(w) = \frac{2}{n_{in}}$.

Small example: initialisation numbers

He normal for $n_{in} = 512$: $\text{std} = \sqrt{2/512} = 0.0625$. Xavier uniform for $n_{in} = 100, n_{out} = 50$: limit $\sqrt{6/150} = 0.2$.

18.7 Residual connections (ResNet)

Degradation problem: before ResNet, simply adding more layers to a deep network made even the *training* error worse (not overfitting: an optimisation difficulty). A residual block outputs $\mathbf{x} + F(\mathbf{x})$: the layers learn a correction to the identity. If the best thing is to do nothing, set $F = 0$, which is easy. Gradients flow straight through the “+ $\mathbf{x}$ ” path, so very deep networks (152 layers and more) train well. Transformers use the same idea around every sublayer.

18.8 Normalisation layers

- **Batch Normalisation:** for each feature, normalise across the mini-batch to mean 0, variance 1, then scale and shift with learnable γ, β : $\hat{x} = (x - \mu_B)/\sqrt{\sigma_B^2 + \epsilon}$, $y = \gamma\hat{x} + \beta$. At test time use running averages of μ, σ^2 . Allows higher learning rates, reduces sensitivity to initialisation, mild regulariser. Depends on batch size; awkward for RNNs and tiny batches.

Small example: BatchNorm by hand

Batch (1, 2, 3): $\mu = 2, \sigma^2 = 2/3, \sigma = 0.816$. For $x = 3$: $\hat{x} = 1.225$; with $\gamma = 2, \beta = 1$: $y = 3.449$. A BatchNorm layer on 64 features has $2 \times 64 = 128$ trainable parameters.

- **Layer Normalisation:** normalise across features within each example: independent of batch size; standard in Transformers and RNNs.

18.9 Regularisation in deep learning

Covered conceptually in Chapter 1; specifics:

- **Dropout**: during training, zero each unit with probability p ; **inverted dropout** scales the survivors by $1/(1 - p)$ so the expected activation is unchanged and nothing needs to change at test time (dropout off). With $p = 0.2$, a kept activation of 5 becomes 6.25. Interpretation: trains an ensemble of thinned networks sharing weights; prevents co-adaptation.
- **Weight decay / L2, early stopping, data augmentation, label smoothing** (with $\epsilon = 0.1$ and 10 classes, the target for the true class becomes $0.9 + 0.01 = 0.91$ and 0.01 for each other), **batch norm**, smaller models, transfer learning.

18.10 Learning rate and getting unstuck

See Chapter 5. In deep nets: warmup, cosine or step decay, reduce-on-plateau; momentum or Adam to traverse plateaus; restarts; check for dead ReLUs (fraction of units always 0); always first verify the model can overfit one small batch.

18.11 RBMs and Deep Belief Networks (historical)

A **Restricted Boltzmann Machine** is a two-layer generative energy-based model (visible and hidden units, connections only between layers), trained with contrastive divergence. A **Deep Belief Network** stacks RBMs, trained greedily layer by layer, then fine-tuned; it was how deep networks were first trained successfully (2006) before ReLU, better initialisation and GPUs made direct training work. Today mostly of historical interest, but the OA asked about them.

In practice: neural networks at InfoEdge

The Tools slide lists ANNs, PyTorch and Keras/TensorFlow, and “TensorFlow optimisation”. Typical uses: two-tower recommendation and matching networks; deep click models combining embeddings of seeker, job and context; and fine-tuned Transformers. Practical lessons that interviewers like to hear: start with a GBM baseline on tabular data; use embeddings for high-cardinality IDs; monitor gradient norms; use mixed precision and early stopping; serve with quantised or distilled models for latency.

18.12 Interview questions

Q18.1

★ **PYQ (INT: how does a network learn / backprop)** Explain how a neural network learns, from the forward pass to the weight update.

🕒 **In one line.** Inputs flow forward through weighted sums and non-linearities to a prediction; a loss measures the error; backpropagation applies the chain rule from the output back to every weight to get gradients, reusing intermediate results; and an optimiser nudges each weight against its gradient.

⚙️ **The full answer.** **Forward.** Each layer computes $\mathbf{a}^{(l)} = g(W^{(l)}\mathbf{a}^{(l-1)} + \mathbf{b}^{(l)})$; store $\mathbf{z}$ and $\mathbf{a}$ for the backward pass. **Loss.** Cross-entropy or MSE. **Backward.** Start with $\delta^{(L)} = \partial L / \partial \mathbf{z}^{(L)}$ ($= \hat{\mathbf{y}} - \mathbf{y}$ for softmax + CE). For each layer going backward: weight gradient $\delta^{(l)} \mathbf{a}^{(l-1)\top}$, then pass the error back through the weights and the activation derivative. **Update.** $W \leftarrow W - \eta \nabla_W L$ on a mini-batch; repeat for many batches and epochs. Backprop is efficient because each layer's gradient reuses the next layer's: total cost about 2–3 forward passes, not one per weight.

📊 **Worked example.** The hand computation above: $\partial L / \partial w_2 = 0.152$, $\partial L / \partial w_1 = 0.115$, updated weights 1.985 and 0.4885.

💡 **If the interviewer pushes.** “What is autograd?” Frameworks record the computation graph during the forward pass and apply the chain rule automatically (reverse-mode automatic differentiation). “Why reverse mode?” One scalar loss and millions of parameters: reverse mode gets all gradients in one backward pass.

Q18.2

★ **PYQ (INT: vanishing and exploding gradients)** Why do vanishing and exploding gradients happen, and how do you fix each?

🕒 **In one line.** Because the gradient reaching early layers is a product of many per-layer factors (weights times activation derivatives); if those are mostly below 1 it shrinks exponentially, above 1 it blows up; fixes are ReLU-type activations, proper initialisation, normalisation, residual connections, gated RNNs, gradient clipping and smaller learning rates.

⚙️ **The full answer.** $\frac{\partial L}{\partial \mathbf{a}^{(1)}} = \prod_{l=2}^L (W^{(l)\top} \text{diag}(g'(\mathbf{z}^{(l)}))) \frac{\partial L}{\partial \mathbf{a}^{(L)}}$ (schematically). Sigmoid derivatives ≤ 0.25 make the product tiny. Large weights (or a recurrent matrix with largest singular value above 1) make it huge. **Symptoms.** Vanishing: early layers' weights barely change, training stalls; gradient norms per layer fall by orders of magnitude toward the input. Exploding: loss spikes or NaN; gradient norm spikes. **Fixes.** Vanishing: ReLU/GELU, He/Xavier init, BatchNorm/LayerNorm, residual connections, LSTM/GRU, shorter paths (attention). Exploding: gradient clipping, lower learning rate, proper init, normalisation, weight decay.

📊 **Worked example.** 8 sigmoid layers: at most $0.25^8 = 1.5 \times 10^{-5}$. RNN with factor 1.5 over 20 steps: about 3,325.

⚠️ **Trap to avoid.** Saying “ReLU completely solves vanishing gradients”: it solves the activation part, not the weight-product part, and introduces dead units.

Q18.3

★ **PYQ (OA and INT: Xavier vs He initialisation)** What are Xavier and He initialisation, and why does He use 2 in the numerator?

🕒 **In one line.** Both set random weight variance by layer size to keep activation and gradient variances stable across layers: Xavier uses $2/(n_{in} + n_{out})$ for tanh/sigmoid, He uses $2/n_{in}$ for ReLU, the 2 compensating for ReLU zeroing half of its inputs.

⚙️ **The full answer.** $\text{Var}(z) = n_{in} \text{Var}(w) \text{Var}(x)$. For a symmetric zero-mean input, ReLU passes half the mass, so $\mathbb{E}[\text{ReLU}(z)^2] = \frac{1}{2} \text{Var}(z)$; to keep the variance, double $\text{Var}(w)$: $2/n_{in}$. Xavier

averages n_{in} and n_{out} to balance forward and backward.

 **Worked example.** He normal $n_{in} = 512$: std 0.0625. Xavier uniform $n_{in} = 100$, $n_{out} = 50$: limit 0.2.

 **If the interviewer pushes.** “What happens with zero init?” Symmetry: all units in a layer stay identical. “Biases?” Usually zero (fine, symmetry is broken by weights).

Q18.4

★ **PYQ (OA: ResNets and vanishing gradients)** How do residual connections help train very deep networks?

🎯 **In one line.** Each block outputs $\mathbf{x} + F(\mathbf{x})$, so the gradient has a direct identity path backward ($1 + \partial F / \partial \mathbf{x}$), which does not vanish, and each block only needs to learn a small correction, which makes identity mappings easy and solves the degradation problem.

⚙️ **The full answer.** Without skips, a 56-layer plain network had higher training error than a 20-layer one (optimisation failure). With skips, depth helps. Unrolled, a ResNet is like an ensemble of many paths of different lengths.

Q18.5

★ **PYQ (OA: universal approximation theorem)** What does the universal approximation theorem say and not say?

🎯 **In one line.** A one-hidden-layer network with enough units and a suitable non-linear activation can approximate any continuous function on a compact set arbitrarily well; it does not say how many units are needed, that training will find the weights, or that the result will generalise.

⚙️ **The full answer.** Depth gives exponential efficiency for compositional functions; this is why we use deep, not just wide, networks.

Q18.6

★ **PYQ (OA: number of hidden layers in a deep network)** When is a network called “deep”?

🎯 **In one line.** Conventionally when it has more than one hidden layer, i.e. at least two hidden layers between input and output; modern networks have tens to hundreds.

Q18.7

★ **PYQ (OA: classification losses)** Which of these are classification losses: cross-entropy, hinge, MSE, Huber, focal, exponential?

🎯 **In one line.** Cross-entropy, hinge, focal and exponential; MSE and Huber are regression losses.

⚙️ **The full answer.** Focal loss $-(1-p)^{\gamma} \log p$ down-weights easy examples (object detection, imbalance). Exponential loss is AdaBoost's. Hinge for $y = +1$, score 0.3: 0.7.

Q18.8

★ **PYQ (OA: RBMs and DBNs)** What are RBMs and deep belief networks?

🎯 **In one line.** An RBM is a two-layer undirected generative model with visible and hidden units connected only across layers, trained by contrastive divergence; a DBN stacks RBMs trained greedily layer by layer, historically used to pre-train deep networks before modern techniques made direct training possible.

Q18.9

★ **PYQ (INT: dropout)** How does dropout work at train and test time, and why does it regularise?

🎯 **In one line.** During training each unit is zeroed with probability p and survivors are scaled by $1/(1-p)$ (inverted dropout); at test time dropout is off and no scaling is needed; it regularises by preventing co-adaptation and approximately averaging an ensemble of thinned networks.

📊 **Worked example.** $p = 0.2$, kept activation 5: $5/0.8 = 6.25$.

💡 **If the interviewer pushes.** “Dropout with BatchNorm?” Their interaction can shift variance between train and test; place dropout after BN or use it mainly in the classifier head. “Monte Carlo dropout?” Keep dropout on at test time and average many passes to estimate uncertainty.

Q18.10

★ **PYQ (INT: weight decay; learning rate adjustment)** How do weight decay and learning-rate schedules interact in deep learning training?

🎯 **In one line.** Weight decay shrinks weights each step by a factor proportional to the learning rate (for SGD), so schedules change effective regularisation; with Adam, use AdamW so decay is decoupled from adaptive scaling; typical recipes pair warmup and cosine decay with weight decay 0.01–0.1.

Q18.11

Why do we need non-linear activations? What happens to a 10-layer network without them?

🎯 **In one line.** Without non-linearities, compositions of linear layers are a single linear map, so a 10-layer network can represent nothing a linear model cannot (it cannot learn XOR).

Q18.12

Compare sigmoid, tanh and ReLU. What is a dying ReLU?

🎯 **In one line.** Sigmoid and tanh saturate (gradients near 0 for large $|z|$), sigmoid is also not zero-centred; ReLU does not saturate for positive inputs and is cheap, but a unit whose input is always negative outputs 0 with 0 gradient forever: a dead ReLU.

⚙️ **The full answer.** Causes of dying: large learning rate or large negative bias. Fixes: Leaky ReLU, ELU, GELU, He init, lower learning rate.

Q18.13

Explain Batch Normalisation. What happens at inference?

🎯 **In one line.** It normalises each feature over the mini-batch to zero mean and unit variance and then applies a learnable scale and shift; at inference it uses running averages of the training mean and variance, so predictions do not depend on the batch.

📊 **Worked example.** Batch (1, 2, 3), $\gamma = 2$, $\beta = 1$: output for 3 is 3.449.

💡 **If the interviewer pushes.** “Why does it help?” Keeps activations in a good range, smooths the loss landscape, allows larger learning rates. “BN vs LayerNorm?” LN normalises across features per example; used in Transformers and RNNs; independent of batch size.

Q18.14

🧠 **Think deeper** Why can't we initialise all weights to the same non-zero value?

🎯 **In one line.** Because every unit in a layer would compute the same output and receive the same gradient, so they would remain identical after every update and the layer would behave like a single unit.

Q18.15

🧠 **Think deeper** Is a neural network with a sigmoid output and cross-entropy loss just logistic regression?

🎯 **In one line.** Its last layer is exactly a logistic regression, but on learned features from the hidden layers rather than raw inputs; with no hidden layers it is logistic regression.

⚙️ **The full answer.** This is the cleanest way to explain what hidden layers do: learn a representation in which the classes become linearly separable.

Q18.16

Count the parameters of an MLP 784–128–10, and of a BatchNorm layer on 64 features.

🎯 **In one line.** 101,770 and 128 trainable (plus 128 running statistics that are not trained).

Q18.17

Gradient of binary cross-entropy with respect to the logit when $p = 0.8, y = 1$?

🎯 **In one line.** $p - y = -0.2$.

Q18.18

🧠 **Think deeper** Your training loss is flat from the first epoch. List what you check.

🎯 **In one line.** That the model can overfit a single small batch (if not, a bug), that labels match inputs, that the loss and output activation are consistent, that gradients are non-zero and not NaN, that the learning rate is in a sensible range, that inputs are normalised, and that layers are not frozen.

🔧 **The full answer.** Common culprits: applying softmax before CrossEntropyLoss (which expects logits), forgetting `optimizer.zero_grad()`, data loader returning the same batch, a dead network from bad init.

Q18.19

🧠 **Think deeper** Deep learning vs gradient-boosted trees for a tabular dataset with 5 lakh rows and 60 features. Which would you start with?

🎯 **In one line.** Gradient-boosted trees: on tabular data with heterogeneous features they usually match or beat neural networks with far less tuning; I would consider a neural network if there are high-cardinality IDs or text/image inputs that benefit from learned embeddings, or to ensemble with the GBM.

Q18.20

What is label smoothing? Compute the targets for $\epsilon = 0.1$ with 10 classes.

🎯 **In one line.** Replace the one-hot target with $(1 - \epsilon)$ on the true class plus ϵ/K spread over all classes: true class 0.91, others 0.01 each; it discourages over-confident logits and improves calibration.

Q18.21

🧠 **Think deeper** A 3-layer linear chain $h_1 = w_1x, h_2 = w_2h_1, \hat{y} = w_3h_2$ with all weights 0.5 and $x = 1$. What is $\partial\hat{y}/\partial w_1$, and what happens with 50 such layers?

🎯 **In one line.** $\partial\hat{y}/\partial w_1 = w_3w_2x = 0.25$; with 50 layers the corresponding product is $0.5^{49} \approx 1.8 \times 10^{-15}$: a vanishing gradient caused by weights alone.

CHAPTER 19

Convolutional Neural Networks

🎥 **Companion video for this chapter:** [3Blue1Brown: But what is a convolution?](#)

InfoEdge interviews reportedly asked CNN filters, stride and padding with the output-size formula, max pooling, and whether CNNs are rotation and translation invariant. 99acres has millions of property photos; Naukri parses resume PDFs and images; Jeevansathi has profile photos (and fake-photo detection). This chapter builds convolution from a 3×3 example to modern architectures.

19.1 Why not a dense network on images?

A 224×224 colour image has $224 \times 224 \times 3 = 150,528$ numbers. A dense layer to just 1,000 units needs 150 million weights. Worse, a dense layer treats pixel (10, 10) and pixel (11, 10) as unrelated, and a cat in the top-left corner is a completely different input from the same cat in the bottom-right. Images have two properties a model should exploit:

- **Locality:** nearby pixels are related; edges and textures are local patterns.
- **Translation:** the same pattern (an edge, a window, a face) can appear anywhere.

A convolution layer uses *small* filters (locality) and *slides the same filter* everywhere (weight sharing for translation).

19.2 The convolution operation, by hand

A **filter** (kernel) is a small grid of weights, e.g. 3×3 . Slide it over the image; at each position, multiply element-wise with the patch under it and sum: one output number. The grid of outputs is a **feature map**.

📐 Small example: a vertical-edge detector

Input (4×4) and filter:

$$X = \begin{pmatrix} 10 & 10 & 0 & 0 \\ 10 & 10 & 0 & 0 \\ 10 & 10 & 0 & 0 \\ 10 & 10 & 0 & 0 \end{pmatrix}, \quad K = \begin{pmatrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{pmatrix}.$$

Top-left position: patch columns 1–3: $(10 + 10 + 10) \cdot 1 + 0 + (0 + 0 + 0) \cdot (-1) = 30$. Top-right position: columns 2–4: $(10 + 10 + 10) - 0 = 30$. Output (2×2) is all 30: strong response where brightness drops left to right: a vertical edge. On a uniform patch the output would be 0.

(Technically deep-learning libraries compute cross-correlation, without flipping the kernel; since kernels are learned, it does not matter.)

With C_{in} input channels (RGB: 3), each filter has depth C_{in} ($3 \times 3 \times 3$) and sums over all channels. Using C_{out} filters gives C_{out} output channels, each a different detected pattern.

19.3 Output size: the formula

For input size n , filter f , padding p , stride s :

$$n_{out} = \left\lfloor \frac{n + 2p - f}{s} \right\rfloor + 1.$$

- **Padding** adds p rows/columns of zeros around the border, so edges are processed and size can be kept. “Same” padding for stride 1: $p = (f - 1)/2$ (1 for 3×3).
- **Stride** is the step between positions; stride 2 halves the size (downsampling).

Small example: output sizes

32×32 , $f = 5$, $p = 0$, $s = 1$: 28. AlexNet layer 1: 227 , $f = 11$, $s = 4$, $p = 0$: $(216/4) + 1 = 55$. 28, $f = 3$, same padding: $p = 1$, output 28. 7, $f = 3$, $s = 2$: 3. 8, $f = 3$, $s = 2$: $\lfloor 5/2 \rfloor + 1 = 3$ (floor matters). 64, $f = 5$, $s = 3$, $p = 2$: $\lfloor 63/3 \rfloor + 1 = 22$.

19.4 Parameters and computation

A conv layer has $(f \cdot f \cdot C_{in} + 1) \cdot C_{out}$ parameters (one bias per filter), *independent of image size*.

Small example: parameter counts

64 filters of 3×3 on RGB: $(27 + 1) \cdot 64 = 1,792$. 64 filters of 5×5 on 32 channels: $(800 + 1) \cdot 64 = 51,264$. A 1×1 conv from 256 to 64 channels: $(256 + 1) \cdot 64 = 16,448$. Compare a dense layer from a flattened $32 \times 32 \times 3$ image to 100 units: 307,200 weights.

Multiplications: output size $\times$ per-output cost: producing $28 \times 28 \times 64$ from 3 channels with 3×3 filters: $28 \cdot 28 \cdot 64 \cdot 27 = 1,354,752$.

19.5 Pooling

Max pooling (usually 2×2 , stride 2) keeps the largest value in each window: halves width and height, adds small-shift robustness, has *no parameters*. **Average pooling** takes the mean. **Global average pooling** averages each whole channel to one number ($7 \times 7 \times 512 \rightarrow 512$), replacing huge dense layers at the end of modern CNNs.

Small example: max vs average pooling

Window $\begin{pmatrix} 1 & 3 \\ 2 & 8 \end{pmatrix}$: max = 8, average = 3.5.

During backprop through max pooling, the gradient goes only to the position that held the maximum (others get 0).

19.6 Invariance vs equivariance

- **Equivariance**: shifting the input shifts the output by the same amount. Convolution is (approximately) **translation equivariant**: a cat moved 10 pixels right produces the same feature map moved 10 pixels right.

- **Invariance:** the output does not change. Pooling (and global pooling at the end) adds some **translation invariance**: the final “cat” score barely depends on where the cat is.
- CNNs are **not** inherently **rotation** or **scale** invariant: a filter tuned to vertical edges does not fire on horizontal ones. We get approximate rotation/scale robustness from **data augmentation** (random rotations, crops, zoom) or special architectures. Strides and pooling also break exact equivariance (small shifts can change outputs).

19.7 Receptive field and the feature hierarchy

The **receptive field** of a unit is the input region that can affect it. Stacking three 3×3 convolutions (stride 1) gives a 7×7 receptive field, with fewer parameters ($3 \cdot 9C^2 = 27C^2$) than one 7×7 layer ($49C^2$) and more non-linearities. Two 3×3 layers ($18C^2$) see as much as one 5×5 ($25C^2$): the VGG insight. Dilated convolutions (gaps between filter taps; 3×3 with dilation 2 has an effective size 5) enlarge the receptive field cheaply.

Early layers detect edges and colours; middle layers textures and parts (windows, door handles); deep layers whole objects (kitchen, bathroom).

19.8 Landmark architectures

- **LeNet-5** (1998): digits; conv–pool–conv–pool–dense.
- **AlexNet** (2012): won ImageNet by a large margin; ReLU, dropout, data augmentation, GPUs, overlapping max pooling.
- **VGG** (2014): only 3×3 convs, very deep (16–19 layers); about 138M parameters, most in the final dense layers.
- **GoogLeNet/Inception** (2014): parallel 1×1 , 3×3 , 5×5 branches; **1×1 convolutions** reduce channels cheaply (a “bottleneck”).
- **ResNet** (2015): residual connections, 50–152 layers.
- **MobileNet**: **depthwise separable** convolutions: a 3×3 filter per input channel, then a 1×1 to mix. For $64 \rightarrow 128$ channels: $9 \cdot 64 + 64 \cdot 128 = 8,768$ weights vs $9 \cdot 64 \cdot 128 = 73,728$.
- **EfficientNet**: scaling depth, width and resolution together.
- **Vision Transformers (ViT)**: split the image into 16×16 patches (224×224 gives 196 patches, 197 tokens with [CLS]) and apply a Transformer; weaker built-in locality bias, so they need more data or pre-training.

19.9 Transfer learning

Take a CNN pre-trained on ImageNet; its early and middle layers are general. With little data (2,000 labelled property photos), replace the final layer and train it (feature extraction), then optionally unfreeze the top blocks with a small learning rate (fine-tuning). Use the same input normalisation as pre-training. This is usually the single biggest win on small image datasets.

19.10 1-D CNNs

Convolutions also work on sequences: text (filters over word embeddings detect n-gram patterns, e.g. TextCNN) and time series (daily applications). They are fast and parallel, with a limited receptive field.

In practice: CNNs at InfoEdge

- **99acres:** room-type classification of listing photos, image quality scoring (blurry/dark photos ranked lower), duplicate-listing detection with CNN embeddings, watermark/contact-number detection in images.
- **Jeevansathi:** photo moderation and fake-profile detection (stock photos, celebrity images, GAN-generated faces).
- **Naukri:** OCR and layout models for resume parsing (CNN backbones in document AI).

19.11 Interview questions

Q19.1

★ **PYQ (INT: output size formula)** Input 32×32 , filter 5×5 , padding 0, stride 1: output size? AlexNet: input 227, filter 11, stride 4: output size? State the general formula and why the floor is there.

🕒 **In one line.** $\lfloor (n + 2p - f)/s \rfloor + 1$: 28 and 55; the floor appears because a filter that would run past the padded border is not applied.

⚙️ **The full answer.** Count positions: the filter's left edge can start at $0, s, 2s, \dots$ as long as it fits within $n + 2p$, i.e. $\text{start} \leq n + 2p - f$. Number of starts = $\lfloor (n + 2p - f)/s \rfloor + 1$.

📊 **Worked example.** $8 \times 8, f = 3, s = 2$: $\lfloor 5/2 \rfloor + 1 = 3$; the last column is skipped.

Q19.2

★ **PYQ (INT: same padding)** What padding keeps a 28×28 input at 28×28 with a 3×3 filter and stride 1? In general?

🕒 **In one line.** $p = 1$; in general $p = (f - 1)/2$ for odd f and stride 1.

Q19.3

★ **PYQ (INT: filters and parameters)** How many parameters and outputs does a conv layer with 64 filters of 3×3 on a $32 \times 32 \times 16$ input (same padding, stride 1) have?

🕒 **In one line.** $(3 \cdot 3 \cdot 16 + 1) \cdot 64 = 9,280$ parameters; output $32 \times 32 \times 64$.

⚙️ **The full answer.** Each filter spans all 16 input channels; parameters do not depend on the 32×32 spatial size. This weight sharing is why CNNs are parameter-efficient.

Q19.4

★ **PYQ (INT: max pooling)** What does max pooling do and why is it used? What happens to gradients through it?

🎯 **In one line.** It keeps the largest activation in each window, shrinking the feature map, adding small-shift robustness and reducing computation with no parameters; in backprop the gradient flows only to the maximum element of each window.

⚙️ **The full answer.** “Did this feature appear anywhere in this region?” Max keeps the strongest evidence. Average pooling is smoother; global average pooling summarises each channel for classification. Strided convolutions can replace pooling (learned downsampling).

Q19.5

★ **PYQ (INT: rotation and translation invariance)** Is a CNN translation invariant? Rotation invariant?

🎯 **In one line.** Convolution layers are translation *equivariant* (shifted input, shifted feature map); pooling and global pooling make the final output approximately translation *invariant*; CNNs are not rotation (or scale) invariant by design, so we use augmentation or special architectures for that.

⚙️ **The full answer.** A filter is applied identically at every position: equivariance. Rotate an image by 90° and horizontal edges become vertical: different filters respond. Max pooling’s window and strides only give invariance to small shifts; in fact aliasing can make outputs sensitive to 1-pixel shifts.

💡 **If the interviewer pushes.** “How would you make a property-photo classifier robust to rotated uploads?” Random rotation augmentation; or detect orientation (EXIF) and correct it.

Q19.6

Why do CNNs work better than dense networks on images?

🎯 **In one line.** Local connectivity and weight sharing exploit the locality and translation structure of images, giving far fewer parameters, better generalisation, and a hierarchy of features from edges to objects.

📊 **Worked example.** Dense from $32 \times 32 \times 3$ to 100 units: 307,200 weights; a conv layer with 64 filters of 3×3 : 1,792.

Q19.7

What does a 1×1 convolution do?

🎯 **In one line.** It mixes channels at each pixel independently, a per-pixel dense layer; used to reduce or expand the number of channels cheaply (bottlenecks) and add non-linearity.

 **Worked example.** $256 \rightarrow 64$ channels: 16,448 parameters. In a bottleneck block, $256 \rightarrow 64$ (1×1), $64 \rightarrow 64$ (3×3), $64 \rightarrow 256$ (1×1) is much cheaper than 3×3 at 256 channels.

Q19.8

What is the receptive field of three stacked 3×3 convolutions? Why stack small filters?

 **In one line.** 7×7 ; stacking small filters gives the same receptive field with fewer parameters and more non-linearities than one large filter.

 **The full answer.** Each extra 3×3 layer adds 2 to the receptive field: $3 \rightarrow 5 \rightarrow 7$. Parameters $27C^2$ vs $49C^2$.

Q19.9

Explain depthwise separable convolution and compute its saving for 3×3 , $64 \rightarrow 128$ channels.

 **In one line.** Split a convolution into a per-channel spatial filter (depthwise) and a 1×1 channel mixer (pointwise): $576 + 8,192 = 8,768$ weights vs 73,728, about 8.4 times fewer.

Q19.10

 **Think deeper** Why did VGG-16 have 138M parameters, and how did later networks shrink this?

 **In one line.** Most parameters sit in the first fully connected layer ($7 \times 7 \times 512 \rightarrow 4096$, about 103M); later networks replaced the dense head with global average pooling and used bottlenecks and depthwise convolutions.

Q19.11

 **Think deeper** You have 2,000 labelled 99acres photos across 6 room types. How do you build a classifier?

 **In one line.** Transfer learning: start from an ImageNet-pretrained ResNet/EfficientNet, replace the head with a 6-way classifier, train the head first, then fine-tune the top blocks with a small learning rate, with augmentation, class balancing and a held-out test set split by listing.

 **The full answer.** Split by listing (photos of the same flat are near-duplicates: leakage). Augmentations: flips, crops, colour jitter, small rotations. Evaluate macro F1 and the confusion matrix (kitchen vs dining confusions). Consider CLIP zero-shot as a baseline. Expect 90%+ quickly.

Q19.12

 **Think deeper** Can you use a CNN on tabular data? On text?

🕒 **In one line.** On text and time series, yes: 1-D convolutions over the sequence detect local patterns (n-grams, short-term motifs); on ordinary tabular data, no natural spatial order exists, so convolution's locality assumption is meaningless and trees or MLPs are better.

Q19.13

What does a transposed convolution do? Compute the output for input 4, kernel 4, stride 2, padding 1.

🕒 **In one line.** It upsamples (the gradient of a strided convolution), used in decoders and GAN generators; output = $(n - 1)s - 2p + f = 3 \cdot 2 - 2 + 4 = 8$.

Q19.14

How many multiplications to produce a $28 \times 28 \times 64$ output from 3 channels with 3×3 filters?

🕒 **In one line.** $28 \cdot 28 \cdot 64 \cdot 27 = 1,354,752$.

Q19.15

Global average pooling on a $7 \times 7 \times 512$ tensor: output shape, and why use it?

🕒 **In one line.** 512 numbers (one per channel); it removes huge dense layers, reduces overfitting, and accepts variable input sizes.

Q19.16

🧠 **Think deeper** What features does a CNN learn at different depths, and how could you check?

🕒 **In one line.** Edges and colour blobs early, textures and parts in the middle, object-level concepts deep; check by visualising filters, feature maps, maximally activating image patches, or saliency maps such as Grad-CAM.

Q19.17

CNN vs Vision Transformer: what is the trade-off?

🕒 **In one line.** CNNs have strong built-in biases (locality, translation equivariance) and do well with less data; ViTs have weaker biases and global attention from the first layer, so they need more data or pre-training but scale better and capture long-range relations.

Q19.18

How would you detect duplicate property listings that use the same photos with small edits?

🎯 **In one line.** Embed photos with a CNN (or perceptual hashes for exact/near copies), index embeddings with ANN, and flag listing pairs with several highly similar photos, combined with text and location similarity.

CHAPTER 20

Recurrent Networks: RNN, LSTM, GRU and Sequence Models

🎥 **Companion video for this chapter:** [StatQuest: Long Short-Term Memory \(LSTM\), Clearly Explained](#)

The InfoEdge hire was asked which deep learning models to use for time series; he covered RNNs and their drawbacks, then LSTMs and GRUs and their internal structure, and “the interviewers were satisfied”. The OA has asked where LSTMs are used (speech recognition). Be ready to draw an LSTM cell, write its equations, and explain exactly why it fixes the vanishing gradient.

20.1 Why sequences need memory

The meaning of “bank” depends on earlier words; tomorrow’s application count depends on the last few weeks; the next job a seeker clicks depends on what they clicked before. A feed-forward network takes a fixed-size input and has no memory. We want a model that reads a sequence one step at a time, keeps a running summary (a **hidden state**), and handles any length.

20.2 The vanilla RNN

At each time step t , with input $\mathbf{x}_t$ and previous hidden state $\mathbf{h}_{t-1}$:

$$\mathbf{h}_t = \tanh(W_{hh}\mathbf{h}_{t-1} + W_{xh}\mathbf{x}_t + \mathbf{b}), \quad \hat{\mathbf{y}}_t = W_{hy}\mathbf{h}_t + \mathbf{c}.$$

The *same* weights are used at every step (weight sharing across time). Unrolled over time, an RNN is a very deep network with tied weights.

📊 Small example: RNN by hand

Scalar RNN $h_t = \tanh(0.5h_{t-1} + x_t)$, $h_0 = 0$, $x_1 = 1$, $x_2 = -1$. $h_1 = \tanh(1) = 0.762$. $h_2 = \tanh(0.381 - 1) = \tanh(-0.619) = -0.551$.

Parameters (input size d , hidden size h): W_{xh} : hd , W_{hh} : h^2 , bias h : total $h(d + h + 1)$. For $d = 10$, $h = 20$: $20 \cdot 31 = 620$.

20.2.1 Shapes of sequence problems

- Many-to-one: sentiment of a review; “will this seeker apply” from their session.
- One-to-many: image captioning.
- Many-to-many (aligned): POS/NER tagging of resume tokens.
- Many-to-many (seq2seq): translation, summarising a JD.

🎥 **Stuck on this page?** Watch: [StatQuest: Long Short-Term Memory \(LSTM\), Clearly Explained](#)

20.3 Backpropagation through time (BPTT)

Unroll the RNN over the sequence, compute the loss (summed over steps), and backpropagate through the unrolled graph; gradients for the shared weights are summed over time steps. For long sequences, **truncated BPTT** backpropagates only through the last k steps (e.g. 50), carrying the hidden state forward but cutting the gradient: cheaper, but cannot learn dependencies longer than k through gradients.

20.4 The vanishing/exploding problem in RNNs

The gradient of the loss at step t with respect to $\mathbf{h}_{t-k}$ contains the product $\prod_j \frac{\partial \mathbf{h}_j}{\partial \mathbf{h}_{j-1}} = \prod_j \text{diag}(\tanh'(\cdot))W_{hh}$: the same matrix W_{hh} multiplied k times. If its largest singular value (times $\tanh' \leq 1$) is below 1, the product shrinks exponentially: the network cannot learn that a word 30 steps back matters. If above 1, it explodes.

Small example: how fast it vanishes

Factor 0.5 per step: after 10 steps $0.5^{10} = 0.001$. It first drops below 10^{-3} at step 10 ($0.5^{10} = 9.8 \times 10^{-4}$).
Factor 1.5: after 20 steps about 3,325.


Exploding is handled by gradient clipping. Vanishing needs an architectural fix: gates.

20.5 LSTM: a memory with gates

The LSTM adds a separate **cell state** $\mathbf{c}_t$, a conveyor belt of memory, and three **gates** (sigmoid outputs in $(0, 1)$ that act as soft switches).

$\mathbf{f}_t = \sigma(W_f[\mathbf{h}_{t-1}, \mathbf{x}_t] + \mathbf{b}_f)$	forget gate: how much old memory to keep
$\mathbf{i}_t = \sigma(W_i[\mathbf{h}_{t-1}, \mathbf{x}_t] + \mathbf{b}_i)$	input gate: how much new information to write
$\tilde{\mathbf{c}}_t = \tanh(W_c[\mathbf{h}_{t-1}, \mathbf{x}_t] + \mathbf{b}_c)$	candidate content
$\mathbf{c}_t = \mathbf{f}_t \odot \mathbf{c}_{t-1} + \mathbf{i}_t \odot \tilde{\mathbf{c}}_t$	update memory (additive!)
$\mathbf{o}_t = \sigma(W_o[\mathbf{h}_{t-1}, \mathbf{x}_t] + \mathbf{b}_o)$	output gate: how much memory to reveal
$\mathbf{h}_t = \mathbf{o}_t \odot \tanh(\mathbf{c}_t)$	hidden state / output

$[\mathbf{h}_{t-1}, \mathbf{x}_t]$ means the two vectors concatenated; $\odot$ is element-wise multiplication.



Small example: one LSTM step

$c_{t-1} = 2, f = 0.5, i = 0.8, \tilde{c} = 0.5, o = 0.5$. $c_t = 0.5 \cdot 2 + 0.8 \cdot 0.5 = 1.4$. $h_t = 0.5 \tanh(1.4) = 0.5 \cdot 0.885 = 0.443$.

20.5.1 Why the LSTM does not vanish

The cell update is *additive*: $\mathbf{c}_t = \mathbf{f}_t \odot \mathbf{c}_{t-1} + \dots$, so $\partial \mathbf{c}_t / \partial \mathbf{c}_{t-1} = \mathbf{f}_t$ (plus indirect terms), with no squashing non-linearity and no repeated weight matrix on this path. If the forget gate stays near 1, gradients flow back over many steps almost unchanged (the “constant error carousel”). The network *learns* when to keep and when to forget.

Small example: cell path vs vanilla

Forget gate 0.9 for 10 steps: $0.9^{10} = 0.35$. Vanilla RNN with factor 0.5: 0.001. Same 10 steps, 350 times more gradient.

Forget gate bias initialised to 1: so $f \approx \sigma(1) = 0.73$ at the start: the cell remembers by default until it learns otherwise.

20.5.2 Parameter count

Four blocks (f, i, $\tilde{c}$, o), each with weights for $[\mathbf{h}, \mathbf{x}]$ and a bias: $4h(d + h + 1)$. For $d = 10$, $h = 20$: $4 \cdot 20 \cdot 31 = 2,480$. Keras LSTM(128) on 64 features: $4 \cdot 128 \cdot (64 + 128 + 1) = 98,816$. (PyTorch uses two bias vectors, adding $4h$.)

20.6 GRU: a lighter LSTM

Two gates, no separate cell state:

$$\begin{aligned} \mathbf{z}_t &= \sigma(W_z[\mathbf{h}_{t-1}, \mathbf{x}_t]) && \text{update gate} \\ \mathbf{r}_t &= \sigma(W_r[\mathbf{h}_{t-1}, \mathbf{x}_t]) && \text{reset gate} \\ \tilde{\mathbf{h}}_t &= \tanh(W[\mathbf{r}_t \odot \mathbf{h}_{t-1}, \mathbf{x}_t]) \\ \mathbf{h}_t &= (1 - \mathbf{z}_t) \odot \mathbf{h}_{t-1} + \mathbf{z}_t \odot \tilde{\mathbf{h}}_t. \end{aligned}$$

The update gate interpolates between keeping the old state and writing the new one (merging LSTM’s forget and input gates); if $\mathbf{z}_t = 0$ the state is copied unchanged. Parameters $3h(d + h + 1)$: $3 \cdot 20 \cdot 31 = 1,860$. GRUs train faster with similar accuracy on many tasks; LSTMs can be slightly better on very long or complex dependencies.

20.7 Variants and training tricks

- **Stacked** (deep) RNNs: the hidden sequence of layer 1 is the input of layer 2.
- **Bidirectional**: one RNN reads left-to-right, another right-to-left; concatenate (hidden 128 per direction gives 256). Good for tagging and classification; *not* usable for forecasting or real-time generation, which cannot see the future.
- **Seq2seq** encoder–decoder: the encoder compresses the input into a context vector; the decoder generates the output. A single fixed vector is a bottleneck for long inputs; **attention** (Bahdanau) lets the decoder look at all encoder states at every step (20 input $\times$ 15 output tokens gives 300 attention weights), which led to Transformers.
- **Teacher forcing**: during training, feed the true previous token to the decoder rather than its own prediction (fast, stable; mismatch at inference = exposure bias).

- **Padding and masking, packed sequences** for variable lengths; **gradient clipping**; **dropout** between layers (and variational dropout across time).
- **CTC loss** (connectionist temporal classification) for speech recognition and handwriting, where input frames and output characters are not aligned.

20.8 Time series with deep learning

Framing: turn a series into supervised windows: use the last L values (lookback) to predict the next value (or next H). 1,000 daily values, lookback 30, next 1: $1000 - 30 = 970$ samples. **Before an LSTM, beat the baselines:** naive (tomorrow = today), seasonal naive (same day last week), moving averages, ARIMA/ETS, gradient boosting on lag features. They are often hard to beat on short, noisy series. **Good practice:** scale using training statistics only; split by time (walk-forward validation); add calendar features (day of week, holidays, month-end hiring cycles) and exogenous variables (job postings drive applications); predict differences or ratios for trending series. **Model options:** LSTM/GRU; 1-D CNN/TCN (dilated causal convolutions); Transformers for long horizons (Informer, PatchTST); DeepAR (probabilistic, many related series); N-BEATS; and GBMs on lags remain very strong.

20.9 Why Transformers replaced RNNs

RNNs process tokens sequentially (no parallelism over time), struggle with very long dependencies (information must survive many steps), and compress everything into a fixed-size state. Self-attention connects every pair of positions in one step and parallelises across the sequence (Chapter 21). RNNs remain useful for streaming and on-device settings with small memory.

In practice: sequence models at InfoEdge

- **Session-based recommendation:** predict the next job a seeker views from their click sequence (GRU4Rec historically; now Transformers like SASRec).
- **Forecasting:** daily applies per category, recruiter demand by city (Talent Pulse “expected hike”), call-centre volume for telecalling staffing.
- **Speech:** the GenAI slide’s voice agents (calling premium seekers, AI REX bot calling) use ASR and TTS; older ASR used LSTMs with CTC, today Transformer models.

20.10 Interview questions

Q20.1

★ **PYQ (INT: DL models for time series)** Which deep learning models would you use for time series, and what are their trade-offs?

🕒 **In one line.** RNN-family models (vanilla RNN, LSTM, GRU) process the series step by step with a memory; LSTM/GRU fix the vanilla RNN’s vanishing gradients with gates; 1-D CNNs/TCNs and Transformers are parallel alternatives; and I would always compare against simple baselines and gradient boosting on lag features.

⚙️ **The full answer.** **Vanilla RNN:** simple, but cannot learn long-range dependencies (vanishing gradients through repeated W_{hh} and $\tanh$). **LSTM:** cell state with additive updates and forget/input/output gates; remembers over hundreds of steps; more parameters ($4\times$ RNN). **GRU:** update and reset gates, no cell state; $3\times$ RNN parameters; trains faster; similar accuracy. **TCN / 1-D CNN:** dilated causal convolutions; parallel; fixed receptive field. **Transformers:** attention across the whole window; best for long horizons with lots of data; quadratic cost in window length. **Practical recipe:** windows with lookback L ; scale with training stats; walk-forward validation; calendar and exogenous features; compare MAE/MAPE with seasonal-naive and LightGBM on lags; ensemble.

📊 **Worked example.** Daily applications to Data Science jobs: seasonal naive (same weekday last week) MAPE 14%; LightGBM with lags 1–28 and calendar features 8%; LSTM 9%; ensemble 7.5%.

💡 **If the interviewer pushes.** “Why might an LSTM lose to LightGBM here?” Short, noisy series and few examples; GBMs exploit engineered lags and calendar effects very efficiently.

Q20.2

★ **PYQ (INT: LSTM internal structure)** Draw the LSTM cell, write its equations, and explain each gate.

🕒 **In one line.** The LSTM keeps a cell state updated additively as $\mathbf{c}_t = \mathbf{f}_t \odot \mathbf{c}_{t-1} + \mathbf{i}_t \odot \tilde{\mathbf{c}}_t$, where the forget gate decides what to keep, the input gate what to write, and the output gate what part of $\tanh(\mathbf{c}_t)$ becomes the hidden state.

⚙️ **The full answer.** Each gate is a sigmoid layer over $[\mathbf{h}_{t-1}, \mathbf{x}_t]$, producing numbers in $(0, 1)$ per memory slot. The candidate $\tilde{\mathbf{c}}_t$ is a tanh layer proposing new content. Intuition with a resume: on reading “Previously at TCS” the forget gate may clear “current employer” information; on “Currently at Infosys” the input gate writes it; the output gate exposes it when the model needs to predict “company”.

📊 **Worked example.** $c_t = 1.4$, $h_t = 0.443$ in the example above.

Q20.3

★ **PYQ (INT: how the LSTM avoids vanishing gradients)** Why do LSTMs mitigate vanishing gradients?

🕒 **In one line.** Because the cell state is updated additively with a learned forget gate, so the gradient along the cell path is multiplied by $\mathbf{f}_t$ (which can stay near 1) rather than by a repeated weight matrix and a squashing derivative.

⚙️ **The full answer.** $\partial \mathbf{c}_t / \partial \mathbf{c}_{t-1} \approx \text{diag}(\mathbf{f}_t)$. Over 10 steps with $f = 0.9$: 0.35, versus 0.001 for a vanilla RNN with factor 0.5. LSTMs do not eliminate the problem (forget gates below 1 still decay), but the network controls the decay. Exploding gradients can still occur via other paths: clip.

Q20.4

★ PYQ (INT: BPTT) What is backpropagation through time, and what is truncated BPTT?

🕒 In one line. BPTT is ordinary backpropagation applied to the RNN unrolled across time steps, summing gradients of the shared weights over steps; truncated BPTT limits the backward pass to the last k steps to save memory and time.

⚙️ The full answer. Memory grows with sequence length because activations at each step must be stored. Truncation trades long-range learning for cost.

Q20.5

★ PYQ (OA: LSTM applications) Where are LSTMs typically used?

🕒 In one line. Speech recognition (historically with CTC), language modelling and machine translation before Transformers, handwriting recognition, time-series forecasting, and sequence tagging.

Q20.6

★ PYQ (INT: GRU vs LSTM) Compare GRU and LSTM.

🕒 In one line. A GRU merges the forget and input gates into one update gate, adds a reset gate and drops the separate cell state, so it has about 25% fewer parameters and trains faster, with similar accuracy on many tasks; LSTMs can be better for very long, complex dependencies.

⚙️ The full answer. GRU: $h_t = (1 - z)h_{t-1} + z\tilde{h}_t$, a convex combination. If $z = 0$, the state is copied: the GRU's version of the constant error carousel.

Q20.7

★ PYQ (INT: sliding windows for time series) 1,000 daily values, lookback 30, predict the next day: how many training samples, and how do you split them?

🕒 In one line. 970 samples; split by time (e.g. last 20% as validation) or walk-forward, never randomly, and fit scalers on the training period only.

Q20.8

Compute h_2 for $h_t = \tanh(0.5h_{t-1} + x_t)$ with $h_0 = 0$, $x_1 = 1$, $x_2 = -1$.

🕒 In one line. $h_1 = 0.762$, $h_2 = \tanh(-0.619) = -0.551$.

Q20.9

How many parameters do an RNN, LSTM and GRU with input 10 and hidden 20 have (one bias per gate)?

🎯 **In one line.** 620, 2,480 and 1,860.

Q20.10

Why is tanh, not ReLU, the default activation inside RNNs and LSTMs?

🎯 **In one line.** tanh is bounded, keeping the recurrent state from growing without limit when the same weights are applied repeatedly; ReLU's unbounded output can make the state and gradients explode in recurrences (though ReLU RNNs with identity initialisation exist).

Q20.11

Why initialise the LSTM forget-gate bias to 1?

🎯 **In one line.** So that initially $f \approx 0.73$ and the cell remembers by default, letting gradients flow over long ranges early in training instead of the cell forgetting everything at $f \approx 0.5$.

Q20.12

🧠 **Think deeper** Can you use a bidirectional LSTM for stock-price or application-count forecasting?

🎯 **In one line.** Not for the forecast itself: the backward direction reads future values that are not available at prediction time, which is leakage; bidirectional models suit tasks where the whole sequence is known (tagging, classification).

Q20.13

What is teacher forcing, and what problem does it create?

🎯 **In one line.** Training the decoder with the true previous token as input instead of its own prediction; it speeds and stabilises training but creates exposure bias: at inference the model sees its own mistakes, which it never practised recovering from (mitigated by scheduled sampling).

Q20.14

Why did attention help seq2seq models?

🎯 **In one line.** A single fixed-length context vector cannot hold all of a long input; attention lets the decoder build a different weighted summary of all encoder states at each output step, focusing on the relevant input words.

Q20.15

 **Think deeper** Why do Transformers outperform RNNs on long text, and when would you still use an RNN?

🎯 **In one line.** Self-attention connects any two positions in one step (path length 1 vs n for RNNs) and processes all positions in parallel on GPUs; I would still use an RNN for streaming or on-device settings with strict memory limits, or short sequences where its simplicity suffices.

Q20.16

What is CTC loss used for?

🎯 **In one line.** For sequence problems where input frames and output labels are not aligned (speech, handwriting); it sums over all alignments using a blank symbol, so no frame-level labels are needed.

Q20.17

 **Think deeper** A seeker's click sequence today: J1, J5, J9. How would you predict the next job they will view?

🎯 **In one line.** Treat it as next-item prediction: embed job IDs (plus content features for new jobs), run a GRU or a self-attention model over the session, and score candidate jobs by dot product with the final state, trained with sampled softmax on historic sessions, evaluated by recall@20 and MRR on later sessions.

⚙️ **The full answer.** Add time gaps, filters (location), and blend with the seeker's long-term profile. This is what "Reco" in the Seeker AI engines needs.

Q20.18

How do you handle sequences of different lengths in a batch?

🎯 **In one line.** Pad them to a common length and use a mask so padded steps do not affect outputs or the loss, or use packed sequences (PyTorch) so the RNN skips padding; bucketing similar lengths reduces wasted computation.

CHAPTER 21

Transformers and Attention

🎥 **Companion video for this chapter:** [3Blue1Brown: Attention in transformers, step-by-step](#)

An IIT Kharagpur candidate reported that Round 2 “tested your ability to explain the architecture of ML and DL models” on pen and paper, and advised knowing Transformers in depth. Glassdoor InfoEdge reports list Transformers; the pre-placement talk mentioned Transformers in production. This chapter builds attention from a toy lookup, computes it by hand, and assembles the full architecture.

21.1 The idea of attention, in plain words

When you read “The recruiter rejected the candidate because *she* lacked Python”, to understand “she” you look back at “candidate”. Attention lets every word build its new representation as a *weighted average of all words’ information*, with weights saying how relevant each word is.

21.1.1 A soft dictionary lookup

A Python dict returns the value whose key exactly matches the query. Attention is a soft version:

- Each position emits a **query** $\mathbf{q}$ (“what am I looking for?”), a **key** $\mathbf{k}$ (“what do I contain?”) and a **value** $\mathbf{v}$ (“what will I hand over?”).
- The relevance of position j to position i is the dot product $\mathbf{q}_i^\top \mathbf{k}_j$.
- Softmax turns relevances into weights that sum to 1.
- The output for i is the weighted sum of values.

21.2 Scaled dot-product attention

Stack queries, keys and values as matrices Q, K, V (one row per position):

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^\top}{\sqrt{d_k}}\right)V.$$

$Q = XW_Q, K = XW_K, V = XW_V$ are learned linear projections of the input X .

📖 Small example: attention by hand

One query $\mathbf{q} = (1, 0)$; keys $\mathbf{k}_1 = (1, 0), \mathbf{k}_2 = (0, 1)$; values $\mathbf{v}_1 = (10, 0), \mathbf{v}_2 = (0, 10)$; $d_k = 2$. Scores: $\mathbf{q} \cdot \mathbf{k}_1 = 1, \mathbf{q} \cdot \mathbf{k}_2 = 0$; scaled: $1/\sqrt{2} = 0.707, 0$. Softmax: $e^{0.707} = 2.028, e^0 = 1$; weights 0.670 and 0.330. Output: $0.670(10, 0) + 0.330(0, 10) = (6.70, 3.30)$. The query mostly retrieves the first value.

21.2.1 Why divide by $\sqrt{d_k}$?

If $\mathbf{q}$ and $\mathbf{k}$ have d_k independent components with mean 0 and variance 1, their dot product has variance d_k (standard deviation $\sqrt{d_k}$). With $d_k = 64$, scores of size ± 8 to ± 16 are common; softmax of such scores is almost one-hot, so its gradient is nearly zero (saturation). Dividing by $\sqrt{d_k}$ restores unit variance.

21.2.2 Each row of attention weights is a probability distribution

Non-negative and summing to 1 (softmax). If two keys are identical (same score), they get equal weight: with scores 3 and 3 and scalar values 4 and 8, the output is 6.

21.3 Multi-head attention

One attention can focus on one kind of relationship. **Multi-head attention** runs h attentions in parallel, each with its own smaller projections ($d_k = d_{\text{model}}/h$), concatenates their outputs and projects back with W_O :

$$\text{MHA}(X) = \text{Concat}(\text{head}_1, \dots, \text{head}_h)W_O.$$

Different heads learn different patterns (syntax, coreference, position, “skill–seniority” links). With $d_{\text{model}} = 768$ and 12 heads, each head has 64 dimensions; total cost similar to one full-width head. Parameters: four $d \times d$ matrices plus biases: $4(d^2 + d)$. For $d = 512$: 1,050,624.

21.4 The Transformer block

1. **Multi-head self-attention** (each position attends to all positions in the same sequence).
2. **Add & LayerNorm**: residual connection then layer normalisation: $\mathbf{x} \leftarrow \text{LN}(\mathbf{x} + \text{MHA}(\mathbf{x}))$.
3. **Position-wise feed-forward network**: the same 2-layer MLP applied to each position independently, usually $d \rightarrow 4d \rightarrow d$ with ReLU/GELU: parameters $2d \cdot d_{ff} + d_{ff} + d$; for $d = 512$, $d_{ff} = 2048$: 2,099,712.
4. **Add & LayerNorm** again.

Attention mixes information *across* positions; the FFN processes each position’s information (and is where much of the “knowledge” is stored). A rule of thumb: about $12d^2$ parameters per layer ($4d^2$ attention + $8d^2$ FFN); 12 layers with $d = 768$: $12 \cdot 12 \cdot 768^2 \approx 85\text{M}$ (BERT-base’s 110M adds embeddings). Modern models put LayerNorm before each sublayer (**pre-LN**), which trains more stably.

21.5 Positional information

Self-attention is *permutation-equivariant*: shuffle the input tokens and the outputs are shuffled the same way; it has no notion of order. We add position information:

- **Sinusoidal** (original): $PE(pos, 2i) = \sin(pos/10000^{2i/d})$, $PE(pos, 2i+1) = \cos(\dots)$; deterministic, can extrapolate; relative offsets are linear functions. Position 1, dimension 0: $\sin(1) = 0.841$.
- **Learned absolute** embeddings (BERT, GPT-2): limited to the trained maximum length (512 for BERT).

- **Relative / RoPE** (rotary, used in LLaMA, Qwen) and **ALiBi**: encode relative distance directly in attention; better length generalisation.

21.6 Masks

- **Padding mask**: ignore pad tokens (set their scores to $-\infty$ before softmax so their weight is 0).
- **Causal (look-ahead) mask**: in decoders, position i may attend only to positions $\leq i$, so the model cannot peek at the token it must predict. For 4 tokens: $1 + 2 + 3 + 4 = 10$ allowed pairs.

21.7 Three families

- **Encoder-only** (BERT, RoBERTa): bidirectional; understanding tasks: classification, NER, retrieval embeddings.
- **Decoder-only** (GPT, LLaMA, Qwen): causal; trained to predict the next token; generation, chat, and (with scale) most tasks via prompting.
- **Encoder-decoder** (original Transformer, T5, BART): the encoder reads the input; the decoder generates with **cross-attention**, whose queries come from the decoder and keys/values from the encoder output. Translation, summarisation.

21.8 Cost and efficiency

Self-attention computes an $n \times n$ score matrix per head: $O(n^2d)$ time and $O(n^2)$ memory in sequence length n . Going from 1,024 to 4,096 tokens multiplies attention memory by 16. Sequence length 512 with 12 heads: $12 \cdot 512^2 = 3.1\text{M}$ scores per layer. Remedies: **FlashAttention** (exact attention computed in tiles in fast on-chip memory, never materialising the full matrix), sparse/local attention, linear-attention approximations, **KV cache** during generation (store keys and values of previous tokens so each new token costs $O(n)$ instead of recomputing everything), grouped-query attention, quantisation.

21.9 Training and generation

- **Training**: AdamW, warmup then decay, dropout, label smoothing (original paper), large batches, mixed precision.
- **Decoding**: greedy (always the top token), beam search (keep the top B partial sequences; good for translation), **temperature** sampling (divide logits by T : $T < 1$ sharper, $T > 1$ flatter; logits $(2, 0)$ at $T = 2$ give $\sigma(1) = 0.731$ for the first token), top- k and nucleus (top- p) sampling.

21.10 Transformers beyond text

Vision Transformers (patches as tokens), speech (Whisper), recommendation (sequences of user actions: SASRec, BERT4Rec), tabular (FT-Transformer), multimodal (CLIP aligns images and text).

📦 In practice: Transformers at InfoEdge

- **Search:** BERT-style bi-encoders for dense retrieval, cross-encoders for re-ranking (PremiumX “contextual match”).
- **Naukri foundational models** (GenAI slide): “Naukri Ls”, 7B/9B series fine-tuned for search and recommendation; distilled small models (0.6B–4B) for cheap serving.
- **JD generator**, free-text search, chat agents for job discovery: decoder-only LLMs.

21.11 Interview questions

Q21.1

★ **PYQ (INT: explain the Transformer architecture (pen and paper))** Explain the Transformer end to end.

🎯 **In one line.** Tokens become embeddings plus positional information; each layer applies multi-head self-attention (every token gathers information from all tokens with learned weights), then a position-wise feed-forward network, each wrapped in a residual connection and layer normalisation; encoders stack these layers, decoders add causal masking and cross-attention to the encoder; a final linear layer and softmax predict tokens.

⚙️ **The full answer.** I would draw: input tokens $\rightarrow$ embedding + position $\rightarrow$ [MHA $\rightarrow$ Add&Norm $\rightarrow$ FFN $\rightarrow$ Add&Norm] $\times N \rightarrow$ output. Then expand MHA: $Q = XW_Q$, $K = XW_K$, $V = XW_V$; per head $\text{softmax}(QK^T / \sqrt{d_k})V$; concat; W_O . Explain why each piece exists: attention for context mixing, FFN for per-token transformation, residuals for gradient flow, LayerNorm for stable scales, positions because attention is order-blind, masks for padding and causality.

📊 **Worked example.** Attention example: output (6.70, 3.30).

💡 **If the interviewer pushes.** “Complexity?” $O(n^2d)$ per layer for attention, $O(nd^2)$ for the FFN. “Why did it replace RNNs?” Parallelism and constant path length between any two tokens.

Q21.2

★ **PYQ (INT: compute attention)** For $q = (1, 0)$, keys $(1, 0)$ and $(0, 1)$, values $(10, 0)$ and $(0, 10)$, compute scaled dot-product attention.

🎯 **In one line.** Scores 1 and 0, scaled 0.707 and 0, weights 0.670 and 0.330, output (6.70, 3.30).

Q21.3

★ **PYQ (INT: why scale by $\sqrt{d_k}$)** Why divide by $\sqrt{d_k}$?

🎯 **In one line.** Because dot products of d_k -dimensional random vectors have variance d_k , so large d_k gives huge scores that saturate the softmax into near one-hot outputs with vanishing gradients; dividing by $\sqrt{d_k}$ keeps the variance at 1.

📊 **Worked example.** $d_k = 64$: standard deviation 8; after scaling, 1.

Q21.4

★ PYQ (INT: multi-head attention) Why use multiple heads?

🕒 **In one line.** Each head has its own projections and can attend to a different kind of relationship (syntax, coreference, position) in a different subspace; together they capture more than a single attention of the same total width, at about the same cost.

⚙️ **The full answer.** A single softmax produces one weighted average per position, which blends different relations; several heads keep them separate. Some heads can be pruned after training with little loss, showing redundancy.

Q21.5

★ PYQ (INT: encoder vs decoder) What is the difference between encoder and decoder blocks, and between BERT and GPT?

🕒 **In one line.** Encoder blocks use bidirectional self-attention; decoder blocks use causally masked self-attention (and, in encoder-decoder models, cross-attention to the encoder); BERT is encoder-only for understanding, GPT is decoder-only for generation.

Q21.6

★ PYQ (INT: positional encodings) Why do Transformers need positional encodings, and what are the options?

🕒 **In one line.** Self-attention treats its input as a set (permutation-equivariant), so order must be injected; options are sinusoidal, learned absolute, and relative schemes like RoPE and ALiBi.

⚙️ **The full answer.** Without positions, “recruiter rejected candidate” and “candidate rejected recruiter” look identical.

Q21.7

★ PYQ (INT: why Transformers beat RNNs) Give the advantages of self-attention over recurrence.

🕒 **In one line.** Parallel computation across positions (fast training on GPUs), a path length of 1 between any two positions (long-range dependencies), and flexible, content-based mixing; the cost is quadratic memory in sequence length.

Q21.8

How is the causal mask implemented? How many pairs are allowed for 4 tokens?

🕒 **In one line.** Add $-\infty$ to scores above the diagonal before softmax so future positions get zero weight; 10 pairs.

Q21.9

Count the parameters of one multi-head attention layer and one FFN with $d = 512$, $d_{ff} = 2048$.

🕒 **In one line.** MHA: $4(512^2 + 512) = 1,050,624$; FFN: $2 \cdot 512 \cdot 2048 + 2048 + 512 = 2,099,712$.

Q21.10

What is the KV cache?

🕒 **In one line.** During autoregressive generation, keys and values of past tokens do not change, so they are stored and reused; each new token computes only its own query, key and value and attends to the cache, making each step linear instead of quadratic in length.

⚙️ **The full answer.** Memory cost: layers $\times 2 \times$ sequence length $\times d$ per sequence: the main memory bottleneck in LLM serving; reduced by grouped-query attention and quantised caches.

Q21.11

🧠 **Think deeper** Is attention a kernel smoother?

🕒 **In one line.** Yes: the output is a weighted average of values with weights given by a normalised similarity $\exp(\mathbf{q}^\top \mathbf{k} / \sqrt{d})$, which is exactly a Nadaraya–Watson kernel regression with a learned, asymmetric exponential kernel.

⚙️ **The full answer.** Nadaraya–Watson: $\hat{f}(x) = \sum_i K(x, x_i) y_i / \sum_i K(x, x_i)$. Replace x by the query, x_i by keys, y_i by values. This links Chapter 9 (kernels) to Transformers and explains why attention is a similarity-based lookup.

Q21.12

🧠 **Think deeper** What does the feed-forward sublayer do, if attention already mixes information?

🕒 **In one line.** Attention only moves and averages information between positions (a linear combination of values); the FFN applies a learned non-linear transformation to each position's vector, providing most of the model's capacity, and acts somewhat like a key–value memory storing facts.

Q21.13

Why does attention cost $O(n^2)$, and how do long-context models cope?

🕒 **In one line.** Every position compares itself with every other, giving an $n \times n$ score matrix per head; long-context models use FlashAttention (exact but memory-efficient), sparse or sliding-window attention, linear approximations, and positional schemes that extrapolate (RoPE scaling).

📊 **Worked example.** 1,024 to 4,096 tokens: 16 times more score memory.

Q21.14

Explain temperature, top- k and top- p sampling.

🎯 **In one line.** Temperature divides logits by T (lower T more deterministic, higher more random); top- k samples from the k most likely tokens; top- p samples from the smallest set whose probability mass exceeds p .

🧮 **Worked example.** Logits (2, 0): $T = 1$ gives 0.881 for the first token; $T = 2$ gives 0.731.

Q21.15

What is cross-attention?

🎯 **In one line.** Attention where queries come from one sequence (the decoder) and keys and values from another (the encoder output), letting the generated text consult the input; also used to fuse modalities (text queries over image features).

Q21.16

🧠 **Think deeper** Why are residual connections and LayerNorm critical in Transformers? What changed with pre-LN?

🎯 **In one line.** Residuals give gradients a direct path through dozens of layers and let each sublayer learn a refinement; LayerNorm keeps activations well scaled; placing LayerNorm before each sublayer (pre-LN) keeps the residual stream un-normalised, which makes deep Transformers train stably with less warmup.

Q21.17

How would you use a Transformer for Resdex candidate ranking under a 100 ms budget?

🎯 **In one line.** Use a bi-encoder Transformer to embed recruiter queries and profiles for fast retrieval from a vector index, and apply a small distilled cross-encoder only to the top 100 candidates, combined with structured features in a GBM ranker; cache embeddings and quantise models.

Q21.18

What is a ViT's sequence length for a 224×224 image with 16×16 patches?

🎯 **In one line.** $14 \times 14 = 196$ patches, plus a [CLS] token: 197.

Q21.19

🧠 **Think deeper** Self-attention is permutation-equivariant. Could that ever be an advantage?

🎯 **In one line.** Yes, for inputs that are sets rather than sequences (a candidate's unordered skills, items in a basket, points in a cloud): omitting positional encodings gives a model that naturally ignores order (Set Transformer).

CHAPTER 22

Generative Models: GANs, Autoencoders and VAEs

🎥 **Companion video for this chapter:** [Computerphile: Generative Adversarial Networks \(GANs\)](#)

GANs are in the syllabus's architectures list. The InfoEdge hire had worked on GANs, so expect interviewers to push beyond the slogan "a forger and a detective" into the objective, the optimal discriminator, mode collapse and alternatives. Autoencoders and VAEs complete the picture of how neural networks learn to *generate*.

22.1 Discriminative vs generative, once more

A discriminative model learns $P(y | \mathbf{x})$: "is this profile photo real?". A **generative model** learns the data distribution $P(\mathbf{x})$ itself, so it can *create* new samples: a realistic face, a property photo, a job description. Explicit models give a density (GMMs, VAEs approximately, autoregressive LLMs); implicit models (GANs) only give a sampler.

22.2 GANs: the forger and the detective

- The **generator** G takes random noise $\mathbf{z}$ (e.g. 100 numbers from a Gaussian) and outputs a fake sample $G(\mathbf{z})$.
- The **discriminator** D takes a sample and outputs the probability it is real.
- They are trained against each other: D learns to tell real from fake; G learns to fool D . As D gets better, G is forced to produce more realistic samples, until (ideally) fakes are indistinguishable.

22.2.1 The minimax objective

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{data}} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_z} [\log(1 - D(G(\mathbf{z})))]$$

D maximises: high $D(\mathbf{x})$ on real data, low $D(G(\mathbf{z}))$ on fakes. This is exactly binary cross-entropy with real = 1, fake = 0. G minimises the second term: it wants $D(G(\mathbf{z}))$ high.

22.2.2 The optimal discriminator

For a fixed G whose samples have density p_g , at each point $\mathbf{x}$ the discriminator maximises $p_{data} \log D + p_g \log(1 - D)$. Setting the derivative to zero:

$$D^*(\mathbf{x}) = \frac{p_{data}(\mathbf{x})}{p_{data}(\mathbf{x}) + p_g(\mathbf{x})}$$

Small example: optimal discriminator

$p_{data}(\mathbf{x}) = 0.6, p_g(\mathbf{x}) = 0.2: D^* = 0.75$. At the global optimum $p_g = p_{data}: D^* = 0.5$ everywhere (the detective is reduced to guessing), and $V(D^*, G^*) = \log 0.5 + \log 0.5 = -2 \log 2 = -1.386$.

Plugging D^* in, G minimises $-\log 4 + 2 \cdot \text{JSD}(p_{data} \parallel p_g)$: the **Jensen–Shannon divergence** between real and generated distributions (maximum value $\log 2 = 0.693$ when they do not overlap). So in the ideal case, training a GAN minimises JSD.

22.2.3 Training loop

Alternate: (1) update D on a batch of real (label 1) and fake (label 0) samples; (2) update G through D (backprop through D into G ; D 's weights frozen in this step). Discriminator loss for one real with $D = 0.9$ and one fake with $D = 0.2$: $-\ln 0.9 - \ln 0.8 = 0.105 + 0.223 = 0.328$.

22.2.4 Non-saturating generator loss

Early in training fakes are obviously fake: $D(G(\mathbf{z})) \approx 0$, so $\log(1 - D(G(\mathbf{z})))$ is flat and gives G almost no gradient. In practice G instead *maximises* $\log D(G(\mathbf{z}))$ (minimises $-\log D(G(\mathbf{z}))$): same fixed point, strong gradients when G is losing. With $D(G(\mathbf{z})) = 0.1$: $\text{loss} = -\ln 0.1 = 2.303$.

22.3 Why GANs are hard to train

- **Non-convergence**: two players with opposing objectives can oscillate rather than settle (a game, not a minimisation).
- **Mode collapse**: G produces only a few kinds of samples (all faces look alike) that currently fool D , ignoring the diversity of the data.
- **Vanishing gradients**: a too-strong D (accuracy 100%) gives G no useful signal; when real and fake distributions do not overlap, JSD is constant ($\log 2$) and uninformative.
- **Sensitivity** to architecture and hyperparameters; no single loss value tracks quality (both losses can look flat while samples improve).

Remedies: DCGAN guidelines (strided convolutions instead of pooling, BatchNorm, ReLU in G with **tanh** output for pixels in $[-1, 1]$, LeakyReLU in D , Adam with $\eta = 2 \times 10^{-4}, \beta_1 = 0.5$); **one-sided label smoothing** (real targets 0.9); **WGAN** (below); **spectral normalisation** of D (controls its Lipschitz constant); gradient penalty; minibatch discrimination; two-timescale learning rates; progressive growing.

22.3.1 Wasserstein GAN

Replace JSD with the **Wasserstein (earth mover's) distance**: the minimum “work” to move probability mass from p_g to p_{data} . Unlike JSD, it decreases smoothly as distributions move closer even when they do not overlap, so G always gets useful gradients. The discriminator becomes a **critic** outputting an unbounded score (no sigmoid), constrained to be 1-Lipschitz (weight clipping in the original, gradient penalty in WGAN-GP); the critic is updated several times (e.g. 5) per generator step. The critic's loss correlates with sample quality.

22.4 Important variants

- **DCGAN**: convolutional GAN; generator starts from a 4×4 map and doubles resolution with transposed convolutions (kernel 4, stride 2, padding 1): $4 \rightarrow 8 \rightarrow 16 \rightarrow 32 \rightarrow 64$.
- **Conditional GAN**: feed a label or condition y to both G and D to control what is generated (“generate a 2BHK living-room photo”).
- **pix2pix**: paired image-to-image translation (sketch $\rightarrow$ photo) with an L1 loss plus adversarial loss.
- **CycleGAN**: unpaired translation (day $\leftrightarrow$ night) using a **cycle-consistency** loss: translating there and back should recover the original.
- **StyleGAN**: style-based generator, photorealistic faces; the source of many fake profile pictures.

22.5 Evaluating generative models

- **Inception Score (IS)**: high when each generated image is confidently classified (quality) and the classes are diverse (diversity). Ignores real data.
- **Fréchet Inception Distance (FID)**: fit Gaussians to Inception features of real and generated images and compute the Fréchet distance between them; lower is better; captures quality and diversity; standard metric.
- Precision/recall for generative models; human evaluation.

22.6 Autoencoders

An **autoencoder** compresses an input to a small code and reconstructs it: encoder $\mathbf{z} = f(\mathbf{x})$, decoder $\hat{\mathbf{x}} = g(\mathbf{z})$, trained to minimise reconstruction error. The bottleneck forces it to keep the important structure. A linear autoencoder with squared error learns the PCA subspace; non-linear ones learn curved manifolds. Variants: **denoising** (reconstruct clean input from a corrupted one), **sparse**, **contractive**. Uses: compression, feature learning, **anomaly detection** (normal data reconstruct well; anomalies have high reconstruction error), pre-training. A plain autoencoder is a poor *generator*: its code space has holes; random codes decode to garbage.

22.7 Variational autoencoders (VAEs)

- The encoder outputs a *distribution* over codes, $q(\mathbf{z} | \mathbf{x}) = \mathcal{N}(\boldsymbol{\mu}(\mathbf{x}), \text{diag } \sigma^2(\mathbf{x}))$.
- Sample $\mathbf{z} = \boldsymbol{\mu} + \sigma \odot \boldsymbol{\epsilon}$, $\boldsymbol{\epsilon} \sim \mathcal{N}(0, I)$: the **reparameterisation trick**, which moves the randomness into $\boldsymbol{\epsilon}$ so gradients flow through $\boldsymbol{\mu}$ and σ .
- Loss (negative ELBO) = reconstruction error + $KL(q(\mathbf{z} | \mathbf{x}) || \mathcal{N}(0, I))$. The KL term pulls all codes toward a standard normal, filling the latent space without holes, so sampling $\mathbf{z} \sim \mathcal{N}(0, I)$ and decoding gives plausible new data.

VAEs train stably and give a latent space with meaningful interpolation, but samples tend to be blurrier than GANs’.

22.8 GANs vs VAEs vs diffusion vs autoregressive

	Training	Samples	Notes
GAN	adversarial, unstable	sharp, fast (one pass)	mode collapse; no likelihood
VAE	stable (ELBO)	blurrier	meaningful latent space; approximate likelihood
Diffusion	stable (denoising)	best quality and diversity	slow sampling (many steps); powers modern image generators
Autoregressive	stable (next-token)	excellent for text	sequential sampling; exact likelihood; LLMs

Diffusion models learn to reverse a gradual noising process: add Gaussian noise over T steps until the image is pure noise; train a network to predict the noise at each step; generate by starting from noise and denoising step by step.

In practice: generative models at InfoEdge

- **Defensive:** detecting GAN/diffusion-generated profile photos on Jeevansathi and fake recruiter profiles (artefact detectors, reverse image search, behaviour signals). Fraud is adversarial too: as detectors improve, generators adapt.
- **Data augmentation:** synthetic tabular data (CTGAN) or images for rare classes; must be validated on real data.
- **Anomaly detection:** autoencoder reconstruction error on recruiter behaviour.
- **Text generation** (JD generator, content feed's "content manifestation agent"): autoregressive LLMs, not GANs.

22.9 Interview questions

Q22.1

Explain how a GAN works, including the objective.

🎯 **In one line.** A generator maps random noise to samples and a discriminator classifies samples as real or fake; they play a minimax game where the discriminator maximises binary cross-entropy on real vs fake and the generator tries to make the discriminator label its samples real; at the optimum the generator's distribution matches the data and the discriminator outputs 0.5.

⚙️ **The full answer.** Objective $\min_G \max_D \mathbb{E}[\log D(\mathbf{x})] + \mathbb{E}[\log(1 - D(G(\mathbf{z})))]$. Training alternates gradient steps. In practice G uses the non-saturating loss $-\log D(G(\mathbf{z}))$. With the optimal D , the generator minimises the Jensen–Shannon divergence.

💡 **If the interviewer pushes.** "What does G never see?" Real data directly: it learns only through D 's gradients.

Q22.2

Derive the optimal discriminator and compute it for $p_{data} = 0.6$, $p_g = 0.2$.

🕒 **In one line.** $D^* = p_{data} / (p_{data} + p_g) = 0.75$.

⚙️ **The full answer.** Maximise $f(D) = a \log D + b \log(1 - D)$; $f'(D) = a/D - b/(1 - D) = 0 \Rightarrow D = a/(a + b)$.

Q22.3

What is the value of the GAN objective at the global optimum, and what divergence is being minimised?

🕒 **In one line.** $V = -\log 4 = -1.386$, with $D^* = 0.5$ everywhere; the generator minimises the Jensen–Shannon divergence between data and generated distributions.

Q22.4

Why is the non-saturating generator loss used?

🕒 **In one line.** Because early on the discriminator easily rejects fakes, making $\log(1 - D(G(\mathbf{z})))$ flat with tiny gradients; maximising $\log D(G(\mathbf{z}))$ has the same fixed point but large gradients exactly when the generator is doing badly.

📊 **Worked example.** At $D(G(\mathbf{z})) = 0.1$: gradient of $\log(1 - D)$ w.r.t. D is $-1/0.9 = -1.11$; of $-\log D$ is -10 .

Q22.5

What is mode collapse and how do you mitigate it?

🕒 **In one line.** The generator produces only a narrow set of outputs that fool the current discriminator, ignoring other modes of the data; mitigations include WGAN/WGAN-GP, minibatch discrimination, unrolled GANs, spectral normalisation, feature matching and diversity-promoting losses.

⚙️ **The full answer.** The generator is rewarded for any sample that fools D , not for covering the data; D then learns to reject that mode, G jumps to another, and they cycle.

Q22.6

Why are GANs unstable, and what did WGAN change?

🕒 **In one line.** They solve a two-player game rather than a minimisation, and with non-overlapping distributions the JS divergence gives no useful gradient; WGAN uses the Wasserstein distance with a Lipschitz-constrained critic, which provides smooth, informative gradients and a loss that tracks quality.

Q22.7

 **Think deeper** If the discriminator reaches 100% accuracy early and stays there, what is wrong and what do you do?

🎯 **In one line.** The discriminator overpowers the generator, so the generator receives vanishing or uninformative gradients; weaken or regularise D (fewer updates, label smoothing, noise on inputs, spectral norm, dropout), use the non-saturating or Wasserstein loss, or strengthen G .

Q22.8

How are GANs evaluated?

🎯 **In one line.** With FID (distance between Gaussian fits of Inception features of real and generated images; lower is better), Inception Score (confident and diverse classifications), precision/recall for generative models, and human judgement.

Q22.9

Conditional GAN, pix2pix and CycleGAN: what does each do?

🎯 **In one line.** A conditional GAN conditions both networks on extra information to control outputs; pix2pix does paired image-to-image translation with adversarial plus L1 loss; CycleGAN does unpaired translation using cycle-consistency between two generators.

Q22.10

What is an autoencoder, and how can it detect anomalies?

🎯 **In one line.** A network that compresses an input to a low-dimensional code and reconstructs it; trained only on normal data, it reconstructs normal inputs well and anomalies poorly, so high reconstruction error flags anomalies.

⚙️ **The full answer.** Example: recruiter behaviour vectors; the top 0.1% reconstruction errors go to the trust-and-safety queue. Choose the threshold on validation data with known incidents.

Q22.11

 **Think deeper** Why can a VAE generate new data but a plain autoencoder usually cannot?

🎯 **In one line.** A VAE's KL term forces codes to fill a known prior (a standard normal), so random draws from the prior decode to plausible samples; a plain autoencoder's codes occupy an irregular region with gaps, so random codes decode to nonsense.

Q22.12

What is the reparameterisation trick?

🕒 **In one line.** Writing a random sample as a deterministic function of the parameters and independent noise, $z = \mu + \sigma \odot \epsilon$, so gradients can flow through μ and σ during backpropagation.

Q22.13

💡 **Think deeper** GAN vs VAE vs diffusion: which would you pick to generate realistic synthetic property photos for augmentation?

🕒 **In one line.** A diffusion model (or fine-tuned pre-trained one) for the best quality and diversity, accepting slower sampling; a GAN if real-time sampling matters and you can manage training instability; a VAE if you need a smooth latent space more than sharpness.

⚙️ **The full answer.** For augmentation also ask whether synthetic images help at all: validate on real held-out photos; synthetic data can teach artefacts.

Q22.14

Why does a DCGAN generator use tanh at the output?

🕒 **In one line.** Because images are scaled to $[-1, 1]$, and tanh produces outputs in exactly that range with zero-centred, non-saturating behaviour near the middle.

Q22.15

💡 **Think deeper** How would you detect AI-generated fake profile photos on a matrimonial platform?

🕒 **In one line.** Combine an image classifier trained on real vs generated faces (with continual updates as generators change), artefact checks (eyes, backgrounds, frequency patterns), reverse image search and duplicate detection, and behavioural signals (new account, many messages, copied bio text), routing high scores to human review.

⚙️ **The full answer.** A single detector is brittle (adversarial adaptation); metadata and behaviour are harder to fake. Measure precision at the review queue's capacity; monitor drift.

Q22.16

In one GAN training iteration, which network's weights change during the generator step?

🕒 **In one line.** Only the generator's; gradients flow through the discriminator to reach the generator, but the discriminator is not updated in that step.

CHAPTER 23

GenAI Basics: LLMs, Fine-tuning, RAG and Agents

🎥 **Companion video for this chapter:** [Andrej Karpathy: Intro to Large Language Models](#)

The 2026–27 talk said: “a lot more emphasis will be given to ML and DL; GenAI is an evolving field; *basic* knowledge of LLMs is expected; if you don’t know topics like quantisation it will be fine”. But the slides showed a great deal of GenAI work: chat and voice agents, fine-tuned LLMs (SFT, DPO, GRPO, distillation), Naukri foundational models, RAG and Graph RAG, speculative decoding, quantisation, and AI REX, an agentic hiring system. This chapter gives you a clear, correct basic layer, so you can talk about these products intelligently.

23.1 What a large language model is

An LLM is a decoder-only Transformer (Chapter 21) trained to predict the next token on trillions of tokens of text:

$$P(\text{text}) = \prod_t P(x_t \mid x_1, \dots, x_{t-1}).$$

Training minimises cross-entropy (MLE) of the next token. To do that well across all of the internet, the model must learn grammar, facts, styles, reasoning patterns and code. Generation samples one token at a time and appends it to the input.

23.1.1 The training stages

1. **Pre-training:** next-token prediction on huge unlabeled corpora. Produces a “base model” that continues text but does not follow instructions well.
2. **Supervised fine-tuning (SFT):** train on curated (instruction, good response) pairs. The model learns to follow instructions and a format.
3. **Preference alignment:** make outputs more helpful and safe using human (or AI) preferences between responses.
 - **RLHF:** train a reward model on preference pairs, then optimise the LLM with reinforcement learning (PPO) to maximise reward with a KL penalty to stay close to the SFT model.
 - **DPO** (direct preference optimisation): skip the reward model; directly increase the likelihood of preferred responses relative to rejected ones with a simple classification-like loss.
 - **GRPO:** an RL method that scores a group of sampled answers to the same prompt and uses their relative rewards (no separate value network); popular for reasoning with verifiable rewards (maths, code).

23.1.2 Tokens, context and cost

Text is split into subword tokens (Chapter 16); roughly 3–4 English characters per token. The **context window** is the maximum number of tokens the model attends to at once. Cost and latency grow with input and output tokens; attention cost grows quadratically with context (Chapter 21).

23.2 Ways to adapt an LLM to a task

Method	What it is	When
Prompting	instructions and examples in the input (zero-/few-shot); chain-of-thought	first try; no training
RAG	retrieve relevant documents and put them in the prompt	knowledge that changes or is private (job listings, policies)
Full fine-tuning	update all weights	lots of data and compute; deep behaviour change
PEFT / LoRA	train small low-rank adapters $\Delta W = BA$ (Chapter 14)	most fine-tuning today; cheap, swappable
Distillation	train a small “student” to imitate a large “teacher”	cheaper serving (Naukri’s 0.6B–4B distilled models)

23.3 Retrieval-augmented generation (RAG)

LLMs know only what was in their training data (up to a cut-off), can **hallucinate** (confidently produce false statements), and do not know private data. RAG fixes this by giving the model the facts at query time:

1. **Index:** split documents into **chunks** (by size or semantically), embed each chunk (Chapter 17), store in a vector index (plus a keyword index for hybrid search).
2. **Retrieve:** embed the user’s question; fetch the top- k chunks by similarity (and BM25); optionally re-rank with a cross-encoder.
3. **Generate:** put the chunks into the prompt with instructions (“answer only from the context; cite sources; say if unsure”).

Graph RAG (on the slide) builds a knowledge graph of entities and relations (skills, titles, companies: exactly the Naukri taxonomy) and retrieves connected subgraphs, which helps multi-hop questions (“which companies in Pune hire for skills adjacent to MLOps?”).

Evaluating RAG: retrieval recall@ k (is the answer in the retrieved chunks?), answer faithfulness (is every claim supported by the context?), answer relevance, and human or LLM-judge review on a fixed test set.

23.4 Agents

An **agent** is an LLM in a loop that can *act*: it decides which **tool** to call (search jobs, query a database, send a WhatsApp message, schedule a call), observes the result, and continues until the goal is met. Key ideas:

- **Planning and tool use** (function calling with structured arguments).
- **Memory:** short-term (conversation) and long-term (stored facts about a recruiter’s preferences).
- **Orchestration:** several specialised agents coordinated by an orchestrator.
- **Guardrails:** the talk’s note: “when an agent starts it is given a set of boundaries and can make decisions only within them; we define lots of guardrails”. Examples: allowed tools, budgets, approval steps for irreversible actions, PII handling, refusal rules.

23.4.1 AI REX, read as an ML system

The slide shows an orchestration agent with: a **mandate agent** (refines hiring criteria, generates the JD and screening questions), a **sourcing agent** (“leverages 10+ proprietary recommendation engines”), a **multi-channel outreach agent** (6 channels by candidate preference), a **screening agent** (evaluates fitment from candidate responses, including bot calling), and a **calibration agent** (monitors performance and “dynamically triggers more sourcing if more candidates are needed”). Underneath: LLMs for language tasks, classical rankers and recommenders for sourcing, channel-propensity models for outreach, and metrics/feedback loops for calibration.

23.5 Making LLMs cheaper and faster (basic awareness)

- **Quantisation**: store weights in 8 or 4 bits instead of 16; smaller and faster with small quality loss (quantisation-aware training reduces the loss).
- **Distillation** and **small language models**.
- **KV cache, batching, speculative decoding** (a small draft model proposes several tokens; the big model verifies them in one pass; exact same output distribution, faster), **disaggregated prefill/decode** (process the prompt and generate tokens on separately optimised hardware).

23.6 Speech: ASR and TTS

ASR (speech-to-text, e.g. Whisper: an encoder–decoder Transformer on spectrograms) and **TTS** (text-to-speech) power the voice agents on the slides (calling premium seekers, voice search). The note “converting TTS + ASR + LLM pipelines into end-to-end models” refers to speech-native models that avoid passing through text.

23.7 Risks and evaluation

Hallucination; prompt injection (instructions hidden in a resume or job post that hijack an agent); bias (an LLM screening candidates may disadvantage groups: must be audited); privacy (resumes contain PII); cost and latency. Evaluation: task metrics on a fixed test set, human review, LLM-as-judge with calibration against humans, online A/B tests.

In practice: GenAI products from the slides

Chat agents for job discovery and home-buyer guidance; talent matching; JD generator; free-text search; “Why in-house” models (cost, data privacy, control); Naukri foundational models for search and recommendation; content feed (insight generation agent → content manifestation agent → personalised recommendation); PremiumX “LLM-based ranking features” and “AI reasons for why a candidate is recommended”.

23.8 Interview questions

Q23.1

What is an LLM and how is it trained, in three stages?

🕒 **In one line.** A large decoder-only Transformer trained first to predict the next token on massive text (pre-training), then on instruction–response pairs (supervised fine-tuning), then aligned to human preferences (RLHF, DPO or similar).

⚙️ **The full answer.** Pre-training gives knowledge and language ability; SFT gives instruction following and format; preference tuning gives helpfulness, honesty and safety behaviours. All three optimise likelihood-based or reward-based objectives with gradient descent (AdamW).

Q23.2

What is RAG and why use it instead of fine-tuning to answer questions about current job listings?

🕒 **In one line.** RAG retrieves relevant, up-to-date documents and puts them in the prompt so the model answers from them; it suits fast-changing or private knowledge because updating the index is instant and answers can cite sources, while fine-tuning bakes knowledge into weights that go stale and is costly to update.

⚙️ **The full answer.** Listings change hourly; fine-tuning cannot keep up and cannot easily delete facts. Fine-tuning is better for teaching style, format or a skill (e.g. writing JDs in Naukri's style), often combined with RAG.

Q23.3

🧠 **Think deeper** A RAG chatbot for job seekers gives wrong answers. How do you debug it?

🕒 **In one line.** Separate retrieval from generation: check whether the correct chunk was retrieved (recall@k on a labelled question set); if not, fix chunking, embeddings, hybrid search or re-ranking; if it was retrieved but the answer is wrong, fix the prompt, context ordering and length, or the model, and add faithfulness checks.

⚙️ **The full answer.** Common culprits: chunks split mid-table, embeddings that miss exact terms (add BM25), too many irrelevant chunks diluting the context, stale index, the model ignoring instructions. Build an evaluation set of 200 real questions with gold answers and sources.

Q23.4

What is LoRA and why is it popular?

🕒 **In one line.** LoRA freezes the pre-trained weights and learns a low-rank update $\Delta W = BA$ for selected layers; it trains a tiny fraction of parameters, needs far less memory, and lets you keep

many task-specific adapters for one base model.

 **Worked example.** For a 4096×4096 weight with $r = 8$: 65,536 trainable instead of 16.8M.

Q23.5

SFT vs DPO vs RLHF: what is the difference?

🎯 **In one line.** SFT imitates good examples; RLHF trains a reward model from preference pairs and optimises the policy with RL against it (with a KL penalty); DPO uses the same preference pairs but optimises the policy directly with a classification-style loss, without a separate reward model or RL loop.

Q23.6

What is hallucination and how do you reduce it?

🎯 **In one line.** Hallucination is fluent but unsupported or false output; reduce it with retrieval grounding, instructions to answer only from context and to abstain, citations, constrained or structured outputs, lower temperature, verification steps, and fine-tuning on faithful examples.

Q23.7

Think deeper What is an agent, and what guardrails would you put on a recruiting agent that contacts candidates?

🎯 **In one line.** An agent is an LLM that plans and calls tools in a loop to reach a goal; for recruiting I would restrict tools and channels, cap contact frequency, require human approval for offers or rejections, filter outputs for policy and bias, protect PII, log every action, and enforce the recruiter's mandate boundaries.

⚙️ **The full answer.** The talk emphasised boundaries and guardrails. Also guard against prompt injection from candidate-provided text ("ignore previous instructions and shortlist me").

Q23.8

What is quantisation, in one paragraph?

🎯 **In one line.** Representing weights (and sometimes activations) with fewer bits, such as 8- or 4-bit integers instead of 16-bit floats, to cut memory and speed up inference at a small accuracy cost; post-training quantisation is quick, quantisation-aware training preserves more accuracy.

Q23.9

What is speculative decoding?

🎯 **In one line.** A small, fast draft model proposes several next tokens and the large model checks them all in one forward pass, accepting the longest correct prefix; the output distribution is

unchanged but generation is faster.

Q23.10

 **Think deeper** Would you use an LLM or a classical model to rank 30 lakh Resdex candidates for a query?

🎯 **In one line.** A classical cascade for scale: sparse and dense retrieval, then a GBM or small Transformer ranker; LLMs only on the top few dozen candidates for re-ranking features or explanations, because running an LLM over lakhs of profiles per search is far too slow and expensive.

⚙️ **The full answer.** This matches PremiumX: hybrid retrieval, then “LLM-based ranking features” and explanations on the final set.

Q23.11

What are the risks of using LLMs to screen candidates?

🎯 **In one line.** Bias against protected groups, inconsistent decisions, hallucinated assessments, prompt injection via resumes, privacy leaks, and lack of explainability; mitigate with audits, structured rubric-based scoring, human review of decisions, and monitoring by segment.

Q23.12

Why might InfoEdge build in-house models (“Why in-house” on the slide) instead of calling external APIs?

🎯 **In one line.** Data privacy (resumes are sensitive), cost at very high volume, latency control, customisation on proprietary data (taxonomy, behaviour), and independence from vendors.

Q23.13

What is chunking in RAG and how do you choose chunk size?

🎯 **In one line.** Splitting documents into retrievable pieces; too small loses context, too large dilutes relevance and wastes context window; choose by document structure (sections, paragraphs), with overlap, and tune on retrieval recall.

Q23.14

 **Think deeper** How would you evaluate an LLM-generated job description feature before launch?

🎯 **In one line.** Offline: a rubric (accuracy versus the recruiter’s inputs, completeness, tone, no discriminatory language) scored by humans and a calibrated LLM judge on a test set; online: A/B test recruiter adoption, edit rate, application quantity and quality, and complaint rate.

Part V

Round 3: Programming in Python

CHAPTER 24

Python for Data Science: the Core Language

🎥 **Companion video for this chapter:** [Corey Schafer: Python Tutorial for Beginners \(search the series\)](#)

The pre-placement talk said programming MCQs in the online test “will be quite easy”, giving the example of converting a string to an int and spotting the right output. Round 3 then puts you in a live notebook. This chapter teaches the language features that such questions test, assuming no prior Python. Every output shown here was checked by running the code (Python 3.11).

24.1 Values, types and variables

A **variable** is a name that refers to an object. Every object has a **type**:

Type	Example	Notes
int	42, -7	unlimited size
float	3.14, 1e-3	64-bit; <code>0.1 + 0.2 == 0.3</code> is False
bool	True, False	a subclass of int: <code>True + True == 2</code>
str	"Pune"	immutable sequence of characters
list	[1, "a", 2.5]	ordered, mutable
tuple	(1, 2), (1,)	ordered, immutable; (1) is just 1
dict	{"city": "Pune"}	key → value; insertion-ordered
set	{1, 2, 3}	unordered, unique elements
NoneType	None	“no value”

24.1.1 Converting between types (the OA’s favourite)

```
int("12") + 1          # 13
int(3.9), int(-3.9)    # (3, -3)  truncates toward zero
int("3.5")             # ValueError: invalid literal for int() with base 10: '3.5'
int(float("3.5"))      # 3
"3" + "4"             # '34'      string concatenation
"3" * 3               # '333'
str(10) + "5"         # '105'
int("101", 2)         # 5         binary
round(2.5), round(3.5) # (2, 4)  banker's rounding: ties go to the even number
round(2.675, 2)       # 2.67     2.675 is stored as 2.67499999...
```

24.1.2 Truthiness

In an if, these are false: 0, 0.0, "", [], {}, set(), None, False. Everything else is true, including "0" and [0]. So bool("0") is True.

24.2 Operators

```

7 // 2, -7 // 2      # (3, -4)   floor division rounds DOWN, not toward zero
7 % 3, -7 % 3       # (1, 2)    result takes the sign of the divisor
5 / 2, 4 / 2        # (2.5, 2.0) / always returns a float
2 ** 3 ** 2         # 512      ** is right-associative: 2 ** 9
divmod(17, 5)       # (3, 2)
1 == 1.0            # True     equal values
[] == [], [] is []  # (True, False) == compares values, `is` compares identity

```

24.3 Strings

Strings are immutable sequences: indexing and slicing work like lists, but you cannot assign to a character.

```

s = "DataScience"
s[0], s[-1]        # ('D', 'e')
s[2:6]             # 'taSc'    start inclusive, stop exclusive
s[-3:]             # 'nce'
s[:2]              # 'DtSine'  every second character
s[::-1]            # 'ecneicSataD' reversed
len(s)             # 11
"a,b,c".split(",") # ['a', 'b', '', 'c']
"x ".strip()       # 'x'
"-".join(["a", "b"]) # 'a-b'
"abc".find("z")     # -1      (index() would raise ValueError)
"12".isdigit(), "1.2".isdigit() # (True, False)
f"{3.14159:.2f}"    # '3.14'
s[0] = "d"          # TypeError: 'str' object does not support item assignment

```

24.4 Lists, tuples, dicts and sets

```

L = [3, 1, 2]
L.sort()           # sorts in place and RETURNS None
sorted(L, reverse=True) # [3, 2, 1] returns a new list
L.append(4); L.extend([5, 6]); L.pop(); L.insert(0, 9)
nums = [1, 2, 3, 4]
nums[1:-1], nums[::-1] # ([2, 3], [4, 3, 2, 1])
a, *rest = [1, 2, 3, 4] # a = 1, rest = [2, 3, 4]

d = {"a": 1}
d.get("b", 0)       # 0      (d["b"] would raise KeyError)
for k, v in d.items(): ...
{k: v for k, v in [{"a": 1}, {"b": 2}] if v > 1} # {'b': 2}

s, s2 = {1, 2, 3}, {2, 3, 4}
s & s2, s | s2, s - s2, s ^ s2 # ({2, 3}, {1, 2, 3, 4}, {1}, {1, 4})

```

Membership: x in list is $O(n)$; x in set and x in dict (keys) are $O(1)$ on average because they use hashing. Dict keys and set elements must be hashable (immutable): a tuple can be a key, a list cannot.

24.5 Mutability, references and copying

This is the source of most “what is the output” traps.

```
a = [1, 2, 3]; b = a; b.append(4); print(a)      # [1, 2, 3, 4] same object
a = [1, 2, 3]; b = a[:]; b.append(4); print(a)  # [1, 2, 3] a copy
m = [[0] * 2] * 2; m[0][0] = 5; print(m)      # [[5, 0], [5, 0]] both rows are ONE list
m = [[0] * 2 for _ in range(2)]                # correct: two separate rows

import copy
a = [[1, 2], [3]]
b = copy.copy(a)          # shallow: new outer list, same inner lists
c = copy.deepcopy(a)      # deep: everything copied
a[0].append(9)
print(b, c)               # [[1, 2, 9], [3]] [[1, 2], [3]]

t = ([1], 2); t[0].append(3); print(t)        # ([1, 3], 2) tuple fixed, its list is mutable
```

Function arguments are passed by object reference (“call by sharing”): a function can mutate a mutable argument, but rebinding the name inside does not affect the caller.

```
x = [1, 2, 3]
def mod(l): l += [4]      # in-place: caller sees it
def mod2(l): l = l + [5]  # new list bound to local name: caller does not
mod(x); mod2(x); print(x) # [1, 2, 3, 4]
```

24.5.1 The mutable default argument trap

```
def f(x, L=[]):
    L.append(x)
    return L
print(f(1), f(2))      # [1, 2] [1, 2] the default list is created ONCE
def f(x, L=None):      # fix
    if L is None: L = []
    L.append(x); return L
```

24.6 Control flow and comprehensions

```
[i * i for i in range(5) if i % 2 == 0]      # [0, 4, 16]
{w: len(w) for w in ["ml", "python"]}       # {'ml': 2, 'python': 6}
list(enumerate("ab", start=1))               # [(1, 'a'), (2, 'b')]
list(zip([1, 2, 3], "ab"))                   # [(1, 'a'), (2, 'b')] stops at shortest
any([0, 0, 1]), all([1, 1, 0]), all([])     # (True, False, True)
```

24.7 Functions: *args, **kwargs, lambda, scope

```
def k(*args, **kw): return args, kw
k(1, 2, a=3)          # ((1, 2), {'a': 3})
list(map(lambda v: v * 2, [1, 2, 3]))       # [2, 4, 6]
```

```
list(filter(None, [0, 1, "", 2]))           # [1, 2]
sorted(["b", "A", "c"])                    # ['A', 'b', 'c']  uppercase sorts first
sorted(["b", "A", "c"], key=str.lower)     # ['A', 'b', 'c']
sorted({"b": 2, "a": 3}.items(), key=lambda kv: kv[1], reverse=True) # [('a', 3), ('b', 2)]
```

Scope (LEGB): names are looked up in Local, Enclosing function, Global, Built-in order. Assigning to a global inside a function needs `global`. **Late binding of closures:**

```
funcs = [lambda: i for i in range(3)]; [f() for f in funcs]      # [2, 2, 2]
funcs = [lambda i=i: i for i in range(3)]; [f() for f in funcs] # [0, 1, 2]
```

24.8 Iterators and generators

A **generator** produces values lazily, one at a time, with `yield`; it remembers where it paused. It uses constant memory, ideal for streaming a 10 GB log file line by line.

```
def countdown(n):
    while n > 0:
        yield n
        n -= 1
list(countdown(3))      # [3, 2, 1]
g = (i for i in range(3)) # generator expression
sum(g), sum(g)          # (3, 0)  a generator is exhausted after one pass
```

`next(g)` on an exhausted generator raises `StopIteration`.

24.9 Decorators

A decorator is a function that takes a function and returns a new function that adds behaviour (logging, timing, caching, retries, auth checks).

```
import functools, time
def timer(fn):
    @functools.wraps(fn)                # keep fn's name and docstring
    def wrapper(*a, **k):
        t = time.perf_counter()
        result = fn(*a, **k)
        print(f"{fn.__name__} took {time.perf_counter() - t:.4f}s")
        return result
    return wrapper

@timer                                # same as: add = timer(add)
def add(a, b): return a + b
add(2, 3)                             # prints the timing line, returns 5
```

`@functools.lru_cache` is a built-in caching decorator.

24.10 Exceptions

```

try:
    value = int(row["salary"])
except (ValueError, TypeError):
    value = None
else:
    pass                # runs if no exception
finally:
    pass                # always runs

```

Catch specific exceptions, not bare `except:`.

24.11 Classes in one page

```

class Candidate:
    count = 0                # class attribute, shared
    def __init__(self, name, years):
        self.name, self.years = name, years # instance attributes
        Candidate.count += 1
    def is_senior(self):
        return self.years >= 5
    def __repr__(self):
        return f"Candidate({self.name!r}, {self.years})"

```

`__eq__` defines `==`; without it, `==` falls back to identity.

24.12 Performance notes

Vectorise with NumPy/pandas instead of Python loops (Chapter 25). Use sets/dicts for membership. Use generators for big streams. The **GIL** (global interpreter lock) lets only one thread execute Python bytecode at a time, so CPU-bound work is parallelised with multiprocessing (or done in NumPy, which releases the GIL), while threads suit I/O-bound work.

In practice: Python in the InfoEdge rounds

OA: short output-prediction MCQs on types, slicing, conversion, mutability. Round 3: data handling in pandas and an sklearn model (Chapters 26–27); clean, readable code with small functions and clear names counts.

24.13 Interview and OA questions

Q24.1

★ **PYQ (2026–27 talk: string to int)** What do `int("12") + 1`, `int("3.5")`, `int(3.9)` and `"3" + "4"` give?

🕒 **In one line.** 13; a `ValueError`; 3; and the string `'34'`.

⚙️ **The full answer.** `int()` parses a string only if it is a valid integer literal; "3.5" is not, so convert via float first: `int(float("3.5")) == 3`. `int(3.9)` truncates toward zero (not rounding). `+` on strings concatenates.

💡 **If the interviewer pushes.** "How do you safely convert a messy salary column?" `pd.to_numeric(df["salary"], errors="coerce")` turns unparsable values into NaN.

Q24.2

What is printed: `a = [1,2,3]; b = a; b.append(4); print(a)`? And with `b = a[:]`?

🎯 **In one line.** `[1, 2, 3, 4]` in the first case (both names refer to the same list) and `[1, 2, 3]` in the second (slicing makes a shallow copy).

⚙️ **The full answer.** Assignment never copies in Python; it binds another name to the same object. Check with `a is b`.

Q24.3

What does `m = [[0]*2]*2; m[0][0] = 5; print(m)` print, and how do you fix it?

🎯 **In one line.** `[[5, 0], [5, 0]]`, because the outer `*2` repeats a reference to the same inner list; build rows with a comprehension: `[[0]*2 for _ in range(2)]`.

Q24.4

Explain shallow vs deep copy with an example.

🎯 **In one line.** A shallow copy creates a new outer container that still references the same inner objects; a deep copy recursively copies everything.

📦 **Worked example.** After `a[0].append(9)`: shallow copy shows `[[1, 2, 9], [3]]`, deep copy `[[1, 2], [3]]`. In pandas, `df.copy()` is deep by default for data.

Q24.5

What is wrong with `def f(x, L=[])`? Show the output of `f(1)`, `f(2)`.

🎯 **In one line.** The default list is created once at definition time and shared across calls, so both calls append to the same list: `[1, 2] [1, 2]`; use `L=None` and create a new list inside.

Q24.6

Lists vs tuples vs sets vs dicts: when do you use each?

🎯 **In one line.** List for an ordered, changeable sequence; tuple for a fixed record or a dictionary key; set for uniqueness and fast membership; dict for key-to-value lookups.

⚙️ **The full answer.** Membership in a list is $O(n)$, in a set or dict $O(1)$ average. Tuples are hashable (if their contents are), so (city, title) can key a dict of average salaries.

Q24.7

Give the outputs: `7 // 2, -7 // 2, -7 % 3, 2 ** 3 ** 2, 5 / 2, round(2.5)`.

🕒 **In one line.** `3, -4, 2, 512, 2.5, 2`.

⚙️ **The full answer.** Floor division rounds toward $-\infty$; % has the divisor's sign so that $(a // b) * b + a \% b == a$; ** is right-associative; / is true division; round uses banker's rounding.

Q24.8

Slice `s = "DataScience"`: `s[2:6], s[-3:], s[::-1], s[::-2]`.

🕒 **In one line.** `'taSc', 'nce', 'ecneicSataD', 'DtSine'`.

Q24.9

What is the difference between `==` and `is`?

🕒 **In one line.** `==` compares values (calls `__eq__`); `is` checks whether two names refer to the same object; use `is` only for singletons like `None`.

📖 **Worked example.** `[] == []` is True; `[] is []` is False.

Q24.10

What is a generator, and why would you use one on a 10 GB log of job applications?

🕒 **In one line.** A function with `yield` that produces values lazily one at a time and keeps its state between calls; it lets you process the file line by line in constant memory.

📖 **Worked example.**

```
def parse(path):
    with open(path) as f:
        for line in f:
            uid, job, ts = line.rstrip("\n").split("\t")
            yield uid, job
applies_per_job = Counter(job for _, job in parse("applies.tsv"))
```

Q24.11

What is a decorator? Write one that times a function.

🕒 **In one line.** A function that wraps another function to add behaviour without changing its code; `@timer` above is equivalent to `add = timer(add)`.

⚙️ **The full answer.** Use `functools.wraps` so the wrapped function keeps its name. Real uses: caching (`lru_cache`), retries for flaky API calls, logging model latency.

Q24.12

💡 **Think deeper** Why does `[lambda: i for i in range(3)]` return functions that all give 2?

🎯 **In one line.** Closures capture variables, not values; all lambdas refer to the same `i`, which is 2 after the loop ends; bind the current value with a default argument `lambda i=i: i`.

Q24.13

What are `*args` and `**kwargs`?

🎯 **In one line.** `*args` collects extra positional arguments into a tuple, `**kwargs` extra keyword arguments into a dict; also used to unpack when calling, `f(*lst, **dict)`.

📊 **Worked example.** `k(1, 2, a=3)` returns `((1, 2), {'a': 3})`.

Q24.14

Count word frequencies in a list of job titles without pandas.

🎯 **In one line.** Use `collections.Counter` on the split words.

📊 **Worked example.**

```
from collections import Counter
titles = ["Data Scientist", "Senior Data Scientist", "Data Analyst"]
c = Counter(w.lower() for t in titles for w in t.split())
c.most_common(2)      # [('data', 3), ('scientist', 2)]
```

Q24.15

What does the GIL mean for data science code?

🎯 **In one line.** Only one thread runs Python bytecode at a time, so pure-Python CPU-bound loops do not speed up with threads; use vectorised NumPy/pandas (which run in C and release the GIL), multiprocessing, or libraries with their own parallelism (`n_jobs=-1`).

Q24.16

💡 **Think deeper** Why is `0.1 + 0.2 == 0.3` false, and how should you compare floats?

🎯 **In one line.** Because 0.1 and 0.2 have no exact binary representation, so their sum is 0.30000000000000004; compare with a tolerance using `math.isclose` or `np.isclose`.

Q24.17

What does `L.sort()` return? What about `sorted(L)`?

🎯 **In one line.** `L.sort()` sorts in place and returns `None`; `sorted(L)` returns a new sorted list and leaves `L` unchanged.

📖 **Worked example.** `print(L.sort())` prints `None` (a classic OA trap).

Q24.18

Sort a dict of skill counts by count, descending.

🎯 **In one line.** `sorted(d.items(), key=lambda kv: kv[1], reverse=True)`.

Q24.19

What happens with `d["missing"]` vs `d.get("missing")`?

🎯 **In one line.** The first raises `KeyError`; `get` returns `None` or a given default; for counting, `defaultdict(int)` or `Counter` avoid the issue.

Q24.20

🧠 **Think deeper** A function receives a list and does `l = l + [5]`; another does `l += [5]`. Which changes the caller's list and why?

🎯 **In one line.** Only `l += [5]`, because for lists `+=` mutates in place (`list.__iadd__`), while `l + [5]` builds a new list and rebinds the local name.

CHAPTER 25

NumPy: Arrays, Vectorisation and Broadcasting

🎥 **Companion video for this chapter:** [freeCodeCamp: NumPy Full Course \(search\)](#)

Round 3 Part 1 names NumPy explicitly. Under every pandas column and every scikit-learn model is a NumPy array, so understanding shapes, axes, vectorisation and broadcasting makes you fast and lets you debug “shape mismatch” errors in seconds. All outputs below were verified by running the code (NumPy 2.x).

25.1 Why NumPy

A Python list of a million numbers stores a million separate Python objects; a loop over them runs the interpreter a million times. A NumPy **array** stores raw numbers of one type in one contiguous block of memory, and operations run in compiled C loops. Typical speed-up: 10–100×. This is **vectorisation**: write the operation on the whole array, not on each element.

```
import numpy as np
a = np.array([1, 2, 3])
a * 2, a + a, a ** 2      # [2 4 6] [2 4 6] [1 4 9]   element-wise
a.dtype                  # int64
np.array([1, 2.5]).dtype  # float64   (upcast to a common type)
np.array([1, "a"]).dtype # <U21     everything became strings!
```

`[1, 2, 3] * 2` on a *list* gives `[1, 2, 3, 1, 2, 3]`; on an array it doubles each element.

25.2 Shape, axes and creation

```
M = np.arange(6).reshape(2, 3)  # [[0 1 2]
                                #  [3 4 5]]
M.shape, M.T.shape              # (2, 3) (3, 2)
np.zeros((2, 2)); np.ones(3); np.eye(2); np.linspace(0, 1, 5) # [0. .25 .5 .75 1.]
np.arange(10).reshape(2, -1).shape # (2, 5)   -1 means "infer"
```

Axis is the dimension along which an operation *collapses*: `axis=0` goes down the rows (one result per column), `axis=1` across the columns (one result per row).

```
M.sum()          # 15
M.sum(axis=0)     # [3 5 7]   column sums
M.sum(axis=1)     # [ 3 12]   row sums
M.mean(axis=0)    # [1.5 2.5 3.5]
```

25.3 Indexing

```
M[1, 2]          # 5         row 1, column 2
M[:, 1]          # [1 4]     a column
```

```
M[0]          # [0 1 2]    a row
M[M > 2]      # [3 4 5]    boolean mask
M[[0, 1], [1, 2]] # fancy indexing: elements (0,1) and (1,2)
s = np.array([1, 2, 3, 4]); s[s > 2] = 0    # [1 2 0 0]
np.where(v > 1, "hi", "lo")    # element-wise if/else
```

25.3.1 Views vs copies

Basic slicing returns a **view**: a window onto the same memory. Changing it changes the original. Fancy (list) and boolean indexing return copies.

```
b = M[0]; b[0] = 100; print(M[0])    # [100 1 2]  M changed!
c = M[0].copy(); c[0] = -1           # M unchanged
```

25.4 Broadcasting

When shapes differ, NumPy stretches dimensions of size 1 to match, without copying data. **Rule:** compare shapes from the right; two dimensions are compatible if they are equal or one of them is 1; missing leading dimensions count as 1.

Small example: broadcasting step by step

x has shape (2, 1): $\begin{pmatrix} 1 \\ 2 \end{pmatrix}$; y has shape (3,): (10, 20, 30). Align right: (2, 1) vs (, 3) $\rightarrow$ (2, 1) vs (1, 3).
Both stretch to (2, 3): $x + y = \begin{pmatrix} 11 & 21 & 31 \\ 12 & 22 & 32 \end{pmatrix}$. (2, 3) + (2, 2): last dims 3 vs 2, neither is 1: `ValueError: operands could not be broadcast together`.

The everyday use: standardising a feature matrix in one line.

```
X = np.array([[1., 2.], [3., 4.], [5., 6.]])    # shape (3, 2)
Z = (X - X.mean(axis=0)) / X.std(axis=0)        # (3,2) - (2,) broadcasts per column
# [[-1.22 -1.22] [0. 0.] [1.22 1.22]]
w = np.array([1, 2, 3])
w.reshape(-1, 1).shape, w[:, None].shape, w[None, :].shape # (3,1) (3,1) (1,3)
```

25.5 Useful functions

```
v = np.array([3, 1, 2])
np.argsort(v), np.sort(v), v.argmax()    # [1 2 0] [1 2 3] 0
np.unique([3, 1, 3, 2], return_counts=True) # (array([1, 2, 3]), array([1, 1, 2]))
np.percentile([1, 2, 3, 4, 5, 6, 7, 8], [25, 50, 75]) # [2.75 4.5 6.25]
np.corrcoef([1, 2, 3], [2, 4, 7])[0, 1]    # 0.9934
np.cumsum([1, 2, 3]), np.diff([1, 4, 9])  # [1 3 6] [3 5]
np.clip([-1, 0.5, 2], 0, 1)                # [0. 0.5 1.]
(p > 0.5).astype(int)                      # threshold probabilities into labels
np.concatenate([[1, 2], [3]]); np.vstack(...); np.hstack(...)
```

25.6 Linear algebra

```
A = np.array([[2, 1], [1, 2]])
A @ A           # [[5 4] [4 5]]  matrix product (not A * A, which is element-
               wise)
np.linalg.inv(A) # [[ 0.667 -0.333] [-0.333  0.667]]
np.linalg.eigh(A)[0] # [1. 3.]  eigenvalues of a symmetric matrix
np.dot([1, 2], [3, 4]) # 11
np.linalg.norm([3, 4]) # 5.0
w = np.linalg.solve(X.T @ X, X.T @ y) # normal equations, safer than inv
w, *_ = np.linalg.lstsq(X, y, rcond=None) # least squares via SVD (best)
```

25.7 Missing values and randomness

`np.nan == np.nan` is `False` (NaN is not equal to anything); use `np.isnan`. `np.mean([1, np.nan, 3])` is `nan`; `np.nanmean` gives 2.0. Use a seeded generator for reproducibility: `rng = np.random.default_rng(42)`; `rng.normal(size=5)`.

In practice: NumPy habits that impress in Round 3

Vectorise instead of looping; check `.shape` after every transformation; use `axis` deliberately; prefer `np.where` and boolean masks over Python `ifs` on arrays; set random seeds; never compare floats with `==`.

25.8 Interview and OA questions

Q25.1

Why is NumPy faster than Python lists for numeric work?

🕒 **In one line.** Arrays store one fixed type in contiguous memory and run operations in compiled, vectorised C loops (often using SIMD), avoiding per-element Python object overhead and interpreter loops.

🧮 **Worked example.** Summing 10^7 numbers: a Python loop takes about a second; `arr.sum()` a few milliseconds.

Q25.2

Explain broadcasting and predict the shape of $(2,1) + (3,)$ and whether $(2,3) + (2,2)$ works.

🕒 **In one line.** Shapes are compared from the right and dimensions of size 1 stretch to match; $(2,1) + (3,)$ gives $(2,3)$; $(2,3) + (2,2)$ fails because 3 and 2 are incompatible.

Q25.3

What do `axis=0` and `axis=1` mean in `M.sum(axis=...)` for a 2×3 matrix?

🎯 **In one line.** `axis=0` collapses the rows, giving one sum per column (`[3 5 7]`); `axis=1` collapses the columns, giving one sum per row (`[3 12]`).

Q25.4

💡 **Think deeper** After `b = M[0]`; `b[0] = 100`, has `M` changed? What about after `c = M[[0]]`; `c[0, 0] = -1`?

🎯 **In one line.** Yes in the first case, because basic slicing returns a view sharing memory; no in the second, because fancy indexing returns a copy.

⚙️ **The full answer.** Views save memory and time but cause surprising side effects; use `.copy()` when you intend to modify independently. Check with `np.shares_memory(a, b)`.

Q25.5

Standardise each column of a matrix `X` in one line.

🎯 **In one line.** `(X - X.mean(axis=0)) / X.std(axis=0)`, which broadcasts the per-column means and standard deviations over the rows.

💡 **If the interviewer pushes.** “In a train/test setting?” Compute mean and std on the training set only and apply them to test (or use `StandardScaler`).

Q25.6

What is the difference between `A * B`, `A @ B` and `np.dot(A, B)`?

🎯 **In one line.** `*` is element-wise multiplication (with broadcasting); `@` is matrix multiplication; `np.dot` is matrix multiplication for 2-D arrays and the dot product for 1-D arrays.

Q25.7

What is the dtype of `np.array([1, "a"])`, and why does it matter?

🎯 **In one line.** A Unicode string type (`<U21`): mixing types upcasts everything to a common type, so the numbers became strings and arithmetic will fail or misbehave.

⚙️ **The full answer.** The same thing happens in pandas when one bad value (“N/A” typed in a salary column) makes the column object; detect with `df.dtypes`.

Q25.8

Why is `np.nan == np.nan` false, and how do you count NaNs in an array?

🕒 **In one line.** By the IEEE floating-point standard NaN is unequal to everything, including itself; count with `np.isnan(a).sum()`.

Q25.9

Compute the sigmoid of an array and threshold it at 0.5, without loops.

🕒 **In one line.** `p = 1 / (1 + np.exp(-z)); labels = (p >= 0.5).astype(int)`.

Q25.10

Implement one gradient-descent step for linear regression in NumPy.

🕒 **In one line.** Compute predictions $X @ w$, the gradient $X.T @ (X @ w - y) * 2 / n$, and update `w -= lr * grad`.

📊 **Worked example.**

```
def gd_step(X, y, w, lr=0.1):
    n = len(y)
    grad = 2 / n * X.T @ (X @ w - y)
    return w - lr * grad
```

Q25.11

🧠 **Think deeper** Solve least squares in NumPy: why prefer `lstsq` or `solve` over `inv`?

🕒 **In one line.** Explicitly inverting $X^T X$ squares the condition number and amplifies rounding errors; `np.linalg.solve` solves the system directly and `lstsq` uses SVD, which is stable even for rank-deficient X .

Q25.12

What does `np.argsort` return and how would you get the top-3 scoring candidates?

🕒 **In one line.** The indices that would sort the array; top 3: `np.argsort(scores)[::-1][:3]` (or `np.argpartition` for speed on large arrays).

📊 **Worked example.** `np.argsort([3, 1, 2])` is `[1 2 0]`.

Q25.13

What does `reshape(-1, 1)` do, and why does scikit-learn often need it?

🕒 **In one line.** It turns a 1-D array of length n into a 2-D column of shape $(n, 1)$; scikit-learn expects features as a 2-D array (rows $\times$ features), even with one feature.

Q25.14**Compute the 25th, 50th and 75th percentiles of 1–8 with NumPy's default method.**

🎯 **In one line.** 2.75, 4.5, 6.25 (linear interpolation between order statistics).

Q25.15**How do you make random results reproducible?**

🎯 **In one line.** Use a seeded generator, `rng = np.random.default_rng(42)`, pass `random_state` to scikit-learn, and set seeds in deep-learning frameworks.

CHAPTER 26

Pandas and Visualisation: Data Handling for Round 3

🎥 **Companion video for this chapter:** [Corey Schafer: Pandas Tutorial \(Part 8\), Grouping and Aggregating](#)

Round 3 Part 1 is exactly this chapter: “Data Handling (Python, NumPy, Pandas): data reading, column naming, statistical parameters computation, group-by operations, basic visualisations, basic text processing.” The reported task used an MBA placement dataset. We use a small version of it throughout, and every output below was produced by running the code (pandas 3.0; older versions show text columns as object rather than str).

```
sl_no,gender,ssc_p,hsc_p,degree_p,degree_t,workex,mba_p,status,salary
1,M,67.0,91.0,58.0,Sci&Tech,No,58.8,Placed,270000
2,M,79.3,78.3,77.5,Sci&Tech,Yes,66.3,Placed,200000
3,M,65.0,68.0,64.0,Comm&Mgmt,No,57.8,Placed,250000
4,M,56.0,52.0,52.0,Sci&Tech,No,59.4,Not Placed,
5,M,85.8,73.6,73.3,Comm&Mgmt,No,55.5,Placed,425000
6,F,55.0,49.8,67.3,Sci&Tech,Yes,51.6,Not Placed,
7,F,46.0,49.2,79.0,Comm&Mgmt,No,53.3,Not Placed,
8,M,82.0,64.0,66.0,Sci&Tech,Yes,62.1,Placed,252000
```

26.1 The two objects: Series and DataFrame

A **Series** is a one-dimensional labelled array (one column). A **DataFrame** is a table: columns are Series sharing a row **index**. Each column has one dtype (int64, float64, str/object, bool, category, datetime64).

26.2 Reading and first look

```
import pandas as pd, numpy as np
df = pd.read_csv("placement.csv")      # also read_excel, read_json, read_parquet,
    read_sql
df.shape                               # (8, 10)
df.head(); df.tail(3); df.sample(5, random_state=0)
df.info()                              # dtypes, non-null counts, memory
df.dtypes
df.isna().sum()                        # salary: 3 missing
df.describe()                          # count, mean, std, min, quartiles, max of numeric columns
df.describe(include="all")
df["status"].value_counts()            # Placed 5, Not Placed 3
df["status"].value_counts(normalize=True) # 0.625, 0.375
df.nunique(); df["degree_t"].unique()  # ['Sci&Tech', 'Comm&Mgmt']
df.duplicated().sum()                  # 0
```

🎥 **Stuck on this page? Watch:** [Corey Schafer: Pandas Tutorial \(Part 8\), Grouping and Aggregating](#)

Useful `read_csv` arguments: `sep`, `header`, `names`, `usecols`, `dtype`, `parse_dates`, `na_values=["NA", "-", "?"]`, `nrows`, `chunksize` (stream a huge file), `encoding`.

26.3 Column naming

```
df = df.rename(columns={"ssc_p": "ssc_pct", "hsc_p": "hsc_pct"}) # selected columns
df.columns = (df.columns.str.strip() # clean all names at once
              .str.lower()
              .str.replace(" ", "_"))
df = df.add_prefix("f_") # or add_suffix
df.columns.tolist()
```

Good names are lower-case, no spaces, descriptive (`mba_pct`), so `df.mba_pct` and formulas work.

26.4 Selecting and filtering

```
df["mba_p"] # a Series
df[["sl_no", "mba_p"]] # a DataFrame
df.loc[0:2, "sl_no"] # label-based, END INCLUSIVE -> rows 0,1,2
df.iloc[0:2, 0:3] # position-based, end exclusive -> rows 0,1
df.loc[df["mba_p"] > 60, ["sl_no", "mba_p"]] # rows 2 and 8
df.query("ssc_pct > 80 and workex == 'Yes'") # sl_no 8
df[(df.gender == "F") & (df.workex == "Yes")] # & | ~ with parentheses, not and/or
df.sort_values("mba_p", ascending=False).head(2) # sl_no 2, 8
df.nlargest(2, "salary") # sl_no 5, 1
```

⚠ Common mistake

`loc` slices include the end label, `iloc` slices exclude it. Use `&`, `|`, `~` with parentheses for conditions; Python's `and/or` raise "truth value of a Series is ambiguous". Assign with `df.loc[mask, "col"] = value`, not chained indexing `df[mask]["col"] = value` (which may modify a copy).

26.5 Statistical parameters

```
df["salary"].mean(), df["salary"].median() # 279400.0 252000.0 (NaN skipped)
df["ssc_p"].std(), df["ssc_p"].var() # sample (n-1) by default: ddof=1
df["mba_p"].quantile([0.25, 0.5, 0.75])
df["mba_p"].skew(), df["mba_p"].kurt()
df["degree_t"].mode()
num = ["ssc_pct", "hsc_pct", "degree_p", "mba_p", "placed"]
df[num].corr()["placed"] # ssc 0.85, hsc 0.85, degree 0.09, mba 0.58
df[num].corr(method="spearman") # rank correlation, robust to outliers
```

Note that pandas' `std` uses $n - 1$ (sample) while NumPy's `np.std` uses n by default.

26.6 Creating and transforming columns

```

df["placed"] = (df["status"] == "Placed").astype(int)           # vectorised
df["gender_code"] = df["gender"].map({"M": 0, "F": 1})
df["mba_band"] = pd.cut(df["mba_p"], bins=[0, 55, 60, 100],
                        labels=["low", "mid", "high"])           # low 2, mid 4, high 2
df["mba_q"] = pd.qcut(df["mba_p"], q=4)                         # equal-count bins
df["school_gap"] = df["ssc_pct"] - df["hsc_pct"]                # prefer this...
df.apply(lambda r: r["ssc_pct"] - r["hsc_pct"], axis=1)         # ...over row-wise apply
                                                                (slow)
np.where(df["mba_p"] > 60, "high", "other")
pd.get_dummies(df, columns=["workex", "degree_t"], drop_first=True)

```

26.7 Group-by: split, apply, combine

💡 Intuition: what groupby does

Split the rows into groups by a key (“work experience: Yes / No”), apply a function to each group (mean placement rate), combine the results into a table. Every “by segment” question in Round 3 and Round 4 is a groupby.

```

df.groupby("workex")["placed"].mean()                          # No 0.600, Yes 0.667
df.groupby("degree_t").agg(n=("placed", "size"),
                           rate=("placed", "mean"),
                           avg_mba=("mba_p", "mean"))

#           n  rate  avg_mba
# Comm&Mgmt  3  0.67    55.53
# Sci&Tech   5  0.60    59.64
df.groupby(["gender", "workex"])["placed"].mean()             # MultiIndex result
df.groupby("degree_t")["mba_p"].transform("mean")             # same length as df: group mean per
row
df.groupby("degree_t").filter(lambda g: len(g) >= 4)          # keep only big groups
pd.crosstab(df["workex"], df["status"])                        # counts table
df.pivot_table(index="degree_t", columns="workex", values="placed", aggfunc="mean")

```

agg reduces each group to one row; transform returns a value per original row (useful for features like “salary relative to city average”); filter keeps or drops whole groups; apply is the general fallback.

26.8 Missing values

```

df.isna().sum(); df.isna().mean()                             # counts and fractions
df.dropna()                                                    # (5, ...) here: drops every non-placed row!
df.drop(columns=["salary"])                                     # what the leakage-aware candidate did
df["salary"].fillna(0)                                          # only when 0 is meaningful
df["x"].fillna(df.groupby("degree_t")["x"].transform("median")) # group-wise imputation
pd.to_numeric(pd.Series(["10", "x", "3.5"]), errors="coerce")  # [10.0, NaN, 3.5]

```

In this dataset dropna() silently deletes all non-placed students (their salary is missing): a disastrous bias. Always ask *why* values are missing.

26.9 Combining tables

```
pd.merge(left, right, on="id", how="left")    # inner / left / right / outer joins
# left ids [1,2,3], right ids [2,3,4] -> left join keeps 1,2,3; b is NaN for id 1
pd.concat([df1, df2])                        # stack rows (axis=0) or columns (axis=1)
```

Check row counts after a merge: a many-to-many key silently multiplies rows (validate="one_to_one" catches it).

26.10 Basic text processing with the .str accessor

```
s = pd.Series(["Sr. Data Scientist ", "data analyst", "ML Engineer"])
s.str.strip().str.lower()                  # ['sr. data scientist', 'data analyst', 'ml engineer']
s.str.contains("data", case=False)        # [True, True, False]
s.str.len()                               # [19, 12, 11]
s.str.split().str[0]                      # ['Sr.', 'data', 'ML']
s.str.replace(r"^[a-zA-Z ]", "", regex=True)
s.str.extract(r"(\d+)\s*(?:yrs|years)")    # numbers before "yrs"/"years"
words = s.str.lower().str.split().explode(); words.value_counts() # word frequencies
```

26.11 Dates

```
df["posted"] = pd.to_datetime(df["posted"], errors="coerce")
df["posted"].dt.dayofweek, df["posted"].dt.month
df.set_index("posted").resample("W")["applies"].sum()    # weekly totals
df["applies"].rolling(7).mean()                        # 7-day moving average
```

26.12 Basic visualisation

```
import matplotlib.pyplot as plt, seaborn as sns
df["mba_p"].hist(bins=10); plt.xlabel("MBA %"); plt.title("Distribution"); plt.show()
sns.boxplot(data=df, x="status", y="mba_p")             # distribution by group, outliers
sns.countplot(data=df, x="degree_t", hue="status")      # category counts
sns.scatterplot(data=df, x="ssc_p", y="mba_p", hue="status")
sns.heatmap(df[num].corr(), annot=True, cmap="coolwarm", vmin=-1, vmax=1)
df.groupby("workex")["placed"].mean().plot(kind="bar")
```

Which plot for which question:

Question	Plot
Distribution of one numeric variable	histogram, KDE, box plot
Numeric vs category	box/violin plot, bar of means
Two numerics	scatter (with hue for a third), hexbin for many points
Counts of categories	bar / count plot
Correlations among many numerics	heatmap
Trend over time	line plot (with rolling mean)

Always label axes and title; avoid pie charts with many slices; use log scale for skewed salaries.

26.13 Performance

Vectorise (`df["a"] * df["b"]`) instead of `apply`/loops; use `category` dtype for repeated strings; read only needed columns; process big files in chunks; avoid growing a `DataFrame` row by row in a loop (collect in a list, build once).

In practice: a Round 3 data-handling checklist

1. Load, shape, head, info, dtypes; fix types (`to_numeric`, `to_datetime`).
2. Clean column names.
3. Missing values per column, and *why*; duplicates.
4. `describe()`: check valid ranges (percentages within 0–100) before calling anything an outlier.
5. `value_counts` for every categorical; target balance.
6. Leakage check: any column that exists only because of the outcome (salary).
7. Group-bys of the target by key segments; correlation matrix; 2–3 clear plots.
8. Narrate findings in one line each.

26.14 Round 3 style questions

Q26.1

★ **PYQ (Round 3 Part 1)** Load the placement CSV and give a first summary: size, types, missing values, target balance and anything suspicious.

🕒 **In one line.** 8 rows and 10 columns; salary is the only column with missing values (3), exactly for the “Not Placed” rows, so it leaks the target and must be dropped; 62.5% are placed; percentages are all within 0–100, so extreme marks are genuine.

⚙️ **The full answer.**

```
df = pd.read_csv("placement.csv")
print(df.shape); df.info()
print(df.isna().sum())
print(df["status"].value_counts(normalize=True))
print(df.describe())
print(df.groupby("status")["salary"].apply(lambda s: s.isna().mean())) # 1.0 for Not
Placed
```

The last line proves the leakage: missingness of salary is 100% for one class and 0% for the other.

💬 **If the interviewer pushes.** “Would you drop rows with missing salary?” No: that deletes every negative example.

Q26.2

★ **PYQ (Round 3: column naming)** Clean messy column names like " SSC Percentage " and "Work Exp".

🕒 **In one line.** Vectorised string methods on `df.columns`: `strip`, `lower`, replace spaces with underscores, then rename specific ones with a mapping.

⚙️ **The full answer.** `df.columns = df.columns.str.strip().str.lower().str.replace(r"\s+", "_", regex=True)`; then `df.rename(columns={"ssc_percentage": "ssc_pct"})`.

Q26.3

★ **PYQ (Round 3: group-by)** What is the placement rate by work experience and by degree type? Show sizes too.

🕒 **In one line.** `df.groupby("workex")["placed"].mean()` gives No 0.60, Yes 0.67; by degree type: Comm&Mgmt 0.67 ($n = 3$), Sci&Tech 0.60 ($n = 5$), computed with named aggregation.

⚙️ **The full answer.** Always report group sizes next to rates: with $n = 3$, a rate of 0.67 is very uncertain. `agg(n=("placed", "size"), rate=("placed", "mean"))`.

Q26.4

★ **PYQ (Round 3: statistical parameters)** Compute mean, median, standard deviation and quartiles of MBA percentage, and the correlation of each score with placement.

🕒 **In one line.** `df["mba_p"].describe()` gives count, mean (58.1), std (4.74), min, quartiles and max; `df[num].corr()["placed"]` shows school scores (0.85) relate to placement more than degree percentage (0.09).

⚙️ **The full answer.** Correlation with a binary target is the point-biserial correlation. With 8 rows these numbers are illustrative only; say so. Use Spearman if outliers exist.

Q26.5

What is the difference between `loc` and `iloc`? What does `df.loc[0:2]` return vs `df.iloc[0:2]`?

🕒 **In one line.** `loc` selects by labels and includes the end of a slice; `iloc` selects by integer positions and excludes the end; with a default index, `loc[0:2]` returns 3 rows, `iloc[0:2]` returns 2.

Q26.6

agg vs transform vs apply in groupby: when do you use each?

🕒 **In one line.** `agg` to reduce each group to summary values (one row per group); `transform` to broadcast a group statistic back to every row (same length as the input, ideal for features); `apply`

for arbitrary functions when the others do not fit (slowest).

 **Worked example.** `df.groupby("degree_t")["mba_p"].transform("mean")` gives 59.64 for each Sci&Tech row and 55.53 for each Comm&Mgmt row; `df["mba_p"]` - that is "MBA score relative to peers".

Q26.7

 **Think deeper** Why is `df.dropna()` dangerous on the placement data?

 **In one line.** Because salary is missing for every non-placed student, so dropping rows with any NaN removes all negative examples and leaves a dataset where everyone is placed.

 **The full answer.** General lesson: inspect missingness by target and segment before dropping; drop the column instead (here it leaks anyway).

Q26.8

Make a cross-tabulation of work experience vs placement status, and a pivot table of placement rate by degree type and work experience.

 **In one line.** `pd.crosstab(df.workex, df.status)` gives counts (No: 2 not placed, 3 placed; Yes: 1, 2); `pivot_table(index="degree_t", columns="workex", values="placed", aggfunc="mean")` gives rates, with NaN where a combination has no rows (Comm&Mgmt with work experience).

Q26.9

 **PYQ (Round 3: basic text processing)** Given a column of job titles, normalise them, flag data roles, and count the most common words.

 **In one line.** Use the `.str` accessor: `strip` and lower-case, `str.contains("data|analyst|scientist", regex=True)` for the flag, then `split`, `explode` and `value_counts`.

 **Worked example.** `s.str.contains("data", case=False)` gives `[True, True, False]` for the three sample titles.

Q26.10

Convert a salary column containing values like "12,00,000", "NA" and "8.5 LPA" to numbers.

 **In one line.** Standardise text with `.str` (remove commas, convert "LPA" to rupees by multiplying by 10^5), then `pd.to_numeric(errors="coerce")`, and inspect what became NaN.

 **Worked example.**

```
s = df["salary"].astype(str).str.strip().str.upper()
lpa = s.str.contains("LPA")
num = pd.to_numeric(s.str.replace(r"^[0-9.]", "", regex=True), errors="coerce")
```

```
df["salary_rs"] = np.where(lpa, num * 1e5, num)
```

Q26.11

How do you merge a candidates table with an applications table, and what can go wrong?

🕒 **In one line.** `pd.merge(apps, cands, on="candidate_id", how="left");` check row counts and use `validate="many_to_one"` because duplicate keys on the right silently multiply rows, and check for unmatched keys (NaNs) after a left join.

Q26.12

Which plots would you draw to explore placement drivers, and what does each show?

🕒 **In one line.** A count plot of status (balance), box plots of each score by status (separation), a count plot of degree type and work experience with `hue = status`, a scatter of SSC vs MBA coloured by status, and a correlation heatmap.

Q26.13

🧠 **Think deeper** Why is `df.apply(func, axis=1)` slow and how do you replace it?

🕒 **In one line.** It calls a Python function once per row, losing vectorisation; replace it with column arithmetic, `np.where/np.select`, `map` with a dict, or `.str/.dt` methods.

📊 **Worked example.** `df["ssc_pct"] - df["hsc_pct"]` vs `apply(lambda r: r.ssc_pct - r.hsc_pct, axis=1)`: same result `[-24.0, 1.0, -3.0, ...]`, 100× faster on large data.

Q26.14

What is `SettingWithCopyWarning` (chained assignment) and how do you avoid it?

🕒 **In one line.** It warns that `df[mask]["col"] = v` may modify a temporary copy rather than `df`; use `df.loc[mask, "col"] = v`, or `.copy()` when you intentionally create a new frame (pandas 3 uses copy-on-write, making chained assignment never modify the original).

Q26.15

Compute a 7-day moving average of daily applications and weekly totals.

🕒 **In one line.** `df.set_index("date")["applies"].rolling(7).mean()` and `.resample("W").sum()` after parsing dates with `pd.to_datetime`.

Q26.16

 **Think deeper** A colleague bins MBA percentage with `pd.cut` and with `pd.qcut`. What is the difference?

🎯 **In one line.** `cut` uses fixed value boundaries (equal width or given edges), so bins can have very different counts; `qcut` uses quantiles, so bins have (roughly) equal counts but uneven widths.

📊 **Worked example.** With edges `[0, 55, 60, 100]`: counts low 2, mid 4, high 2.

CHAPTER 27

Scikit-learn and the Complete Round 3 Walkthrough

🎥 **Companion video for this chapter:** [Data School: Machine learning in Python with scikit-learn \(search\)](#)

Round 3 Part 2 is “ML Basics (Sklearn): train–test split, training an ML model, performance testing”. The 2026–27 talk described it as “an ML case study: cleaning data, applying an ML model to it, doing some testing of performance, how good the model is”. The reported task was the MBA placement dataset ending in logistic regression, a confusion matrix, precision, recall, F1 and accuracy. This chapter teaches the scikit-learn API from scratch and then walks through that round end to end with real outputs.

27.1 The scikit-learn API in one idea

Every object follows the same pattern:

- **Estimators** learn from data with `fit(X, y)`.
- **Predictors** then `predict(X)` (and classifiers often `predict_proba(X)`).
- **Transformers** fit on training data and transform any data (scalers, encoders, imputers, TF-IDF); `fit_transform` does both on the training set.
- Hyperparameters go in the constructor (`LogisticRegression(C=0.1)`); learned values end with an underscore (`coef_`, `mean_`).

`X` is a 2-D table (rows $\times$ features); `y` a 1-D target.

27.2 Splitting

```
from sklearn.model_selection import train_test_split
Xtr, Xte, ytr, yte = train_test_split(X, y, test_size=0.2, stratify=y, random_state=42)
```

`stratify=y` keeps the class ratio equal in both parts; `random_state` makes the split reproducible. For time data use `TimeSeriesSplit`; for grouped data `GroupKFold`.

27.3 Preprocessing with Pipeline and ColumnTransformer

Numeric and categorical columns need different treatment; every step must be fitted on training data only. A **Pipeline** chains steps so that `fit` fits them in order on training data and `predict` applies them to new data; a **ColumnTransformer** routes columns to different sub-pipelines.

```
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.impute import SimpleImputer
```

🎥 **Stuck on this page?** Watch: [Data School: Machine learning in Python with scikit-learn \(search\)](#)

```

num = X.select_dtypes("number").columns.tolist()
cat = X.columns.difference(num).tolist()
pre = ColumnTransformer([
    ("num", Pipeline([("imp", SimpleImputer(strategy="median")),
                      ("sc", StandardScaler())]), num),
    ("cat", Pipeline([("imp", SimpleImputer(strategy="most_frequent")),
                      ("oh", OneHotEncoder(handle_unknown="ignore"))]), cat)])

```

Why this matters: (1) no leakage, because cross-validation refits the imputer and scaler inside each fold; (2) the same object transforms production data identically; (3) `handle_unknown="ignore"` survives categories never seen in training.

Scalers: `StandardScaler` (mean 0, sd 1), `MinMaxScaler` ([0, 1]), `RobustScaler` (median/IQR, for outliers). Encoders: `OneHotEncoder` (nominal), `OrdinalEncoder` (ordered categories), target encoding (`TargetEncoder`, cross-fitted internally).

27.4 Models you should be able to call

```

from sklearn.linear_model import LogisticRegression, LinearRegression, Ridge, Lasso
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, HistGradientBoostingClassifier
from sklearn.svm import SVC, LinearSVC, SVR
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import MultinomialNB, GaussianNB
from sklearn.cluster import KMeans
from sklearn.decomposition import PCA

```

27.5 Performance testing

```

from sklearn.metrics import (confusion_matrix, classification_report, accuracy_score,
                             precision_score, recall_score, f1_score, roc_auc_score,
                             mean_absolute_error, mean_squared_error, r2_score)
from sklearn.model_selection import cross_val_score, StratifiedKFold, GridSearchCV

```

`confusion_matrix(y_true, y_pred)` returns rows = actual, columns = predicted, ordered by label: for 0/1 it is $\begin{pmatrix} TN & FP \\ FN & TP \end{pmatrix}$. Note: scikit-learn puts the negative class first, the opposite of many textbooks' layout.

27.6 The complete Round 3 walkthrough

To make the outputs concrete we use a *simulated* dataset of 215 students with the same columns as the public MBA campus-recruitment dataset the round reportedly used (gender, SSC/HSC/degree percentages, HSC stream, degree type, work experience, employability test, MBA specialisation and percentage, status, salary). The numbers below are what this code prints on the simulated data; on the real data they will differ, but every step and every judgement is the same.

27.6.1 Step 1: load and understand

```
df = pd.read_csv("placement.csv")
df.shape                                # (215, 13)
df.isna().sum()[lambda s: s > 0]       # salary: 105
df["status"].value_counts()             # Placed 110, Not Placed 105 (balanced)
df.groupby("status")["salary"].apply(lambda s: s.isna().mean())
# Not Placed 1.0, Placed 0.0  <- salary is missing exactly for non-placed students
```

Say aloud: “Salary exists only for placed students, so it is a consequence of the target. Keeping it, or even its missingness, would leak the answer. I will drop it.” **Also:** `sl_no` is an ID: drop it.

27.6.2 Step 2: sanity-check values and outliers

`df.describe()`: all percentages lie within 0–100. High values like 95 are genuine top performers, not errors; I will not remove them. `value_counts()` of each categorical: `degree_t = "Others"` is small (about 5%): keep, but note its estimates will be noisy.

27.6.3 Step 3: quick EDA

```
df.groupby("workex")["status"].apply(lambda s: (s == "Placed").mean())
df[num + ["placed"]].corr()["placed"]
sns.boxplot(data=df, x="status", y="ssc_p")
```

27.6.4 Step 4: split and build the pipeline

```
X = df.drop(columns=["sl_no", "status", "salary"])
y = (df["status"] == "Placed").astype(int)
num = ['ssc_p', 'hsc_p', 'degree_p', 'etest_p', 'mba_p']
cat = ['degree_t', 'gender', 'hsc_s', 'specialisation', 'workex']
Xtr, Xte, ytr, yte = train_test_split(X, y, test_size=0.2, stratify=y, random_state=42)
len(Xtr), len(Xte)                        # 172 43 (positive rate 0.512 in both)
clf = Pipeline([("pre", pre), ("lr", LogisticRegression(max_iter=1000))])
clf.fit(Xtr, ytr)
```

27.6.5 Step 5: evaluate

```
pred = clf.predict(Xte); prob = clf.predict_proba(Xte)[: , 1]
confusion_matrix(yte, pred)              # [[17  4]
                                          # [ 6 16]]  TN=17 FP=4 FN=6 TP=16
print(classification_report(yte, pred, digits=3))
#               precision recall f1-score support
#      0      0.739   0.810   0.773     21
#      1      0.800   0.727   0.762     22
#   accuracy                0.767     43
roc_auc_score(yte, prob)                  # 0.846
```

Interpretation aloud: “Accuracy 76.7%, sensible here because classes are balanced. Of students predicted placed, 80% were (precision); of students actually placed, we identified 72.7% (recall). AUC

0.85 means a random placed student outscores a random non-placed one 85% of the time. With only 43 test rows, these numbers have a margin of error of roughly $\pm 6-7$ points, so I will check cross-validation."

27.6.6 Step 6: cross-validate and compare

```
cv = StratifiedKFold(5, shuffle=True, random_state=42)
s = cross_val_score(clf, X, y, cv=cv, scoring="roc_auc")
s.round(3), s.mean(), s.std()           # [0.764 0.909 0.838 0.807 0.851] 0.834 0.048
GridSearchCV(clf, {"lr__C": [0.01, 0.1, 1, 10]}, cv=cv, scoring="roc_auc").fit(Xtr, ytr)
# best C = 0.01, CV AUC 0.843
rf = Pipeline([("pre", pre), ("rf", RandomForestClassifier(n_estimators=300,
                                                           min_samples_leaf=3, random_state=42))])
cross_val_score(rf, X, y, cv=cv, scoring="roc_auc").mean()  # 0.823
```

"Logistic regression (0.83–0.84) is as good as the random forest (0.82) on 215 rows; I prefer the simpler, interpretable model."

27.6.7 Step 7: explain the model

```
names = clf.named_steps["pre"].get_feature_names_out()
coefs = clf.named_steps["lr"].coef_[0]
sorted(zip(names, coefs), key=lambda t: -abs(t[1]))[:5]
# ssc_p 1.46, workex_Yes 0.62, workex_No -0.62, degree_t_Others 0.56, degree_p 0.45
```

"Scaled SSC percentage has the largest effect: one standard deviation higher multiplies the odds of placement by $e^{1.46} \approx 4.3$. Work experience raises the odds. The degree_t = Others coefficient rests on very few students, so I would not trust it."

27.6.8 Step 8: choose a threshold for the use case

```
for t in [0.3, 0.5, 0.7]:
    p = (prob >= t).astype(int)
    print(t, precision_score(yte, p), recall_score(yte, p))
# 0.3 0.667 0.818
# 0.5 0.800 0.727
# 0.7 0.824 0.636
```

"If the college uses this to find students who need extra support, the important class is 'Not Placed', and we want high recall for it; I would pick the threshold accordingly."

27.6.9 Step 9: demonstrate what leakage would have done

```
X_leaky = X.assign(salary_missing=df["salary"].isna().astype(int))
cross_val_score(leaky_pipeline, X_leaky, y, cv=cv, scoring="roc_auc").mean()  # 1.0
```

A perfect AUC on a noisy human outcome is a red flag, not a success.

27.7 Regression and other tasks in scikit-learn

```
reg = Pipeline([("pre", pre), ("m", Ridge(alpha=1.0))]).fit(Xtr, ytr_salary)
mean_absolute_error(yte_salary, reg.predict(Xte)); r2_score(...)
TransformedTargetRegressor(regressor=Ridge(), func=np.log1p, inverse_func=np.expm1) # log
target
```

Text classification:

```
Pipeline([("tfidf", TfidfVectorizer()), ("clf", LinearSVC())])
KMeans(n_clusters=4, n_init=10).fit_predict(Xs) # clustering
```

Saving: `joblib.dump(clf, "model.pkl")`.

In practice: what interviewers watch in Round 3

Leakage spotted first; no blind outlier deletion; a sound split (stratified); preprocessing inside a pipeline; a simple baseline before anything fancy; the right metrics read correctly; uncertainty acknowledged on small data; clear narration; clean code. Exactly the steps the successful candidate reported.

27.8 Round 3 style questions

Q27.1

★ **PYQ (Round 3 Part 2: train-test split, training, performance testing)** Walk me through building and evaluating a placement classifier in scikit-learn.

🕒 **In one line.** Drop leaking and ID columns, split stratified into train and test, build a Pipeline with a ColumnTransformer (impute + scale numerics, impute + one-hot categoricals) and logistic regression, fit on train, then report the confusion matrix, precision, recall, F1, accuracy and ROC AUC on test, confirmed by stratified cross-validation.

⚙️ **The full answer.** The nine steps above. The key judgement calls: leakage, outliers, metric choice, threshold, honest uncertainty.

Q27.2

★ **PYQ (Round 3: leakage)** Why do you drop salary before modelling placement? How would you prove the leak to the interviewer?

🕒 **In one line.** Because salary is only recorded after placement, so it is unavailable at prediction time and encodes the label; the proof is that its missingness is 100% for non-placed and 0% for placed students, and a model using it reaches AUC 1.0.

Q27.3

Why put the scaler and imputer inside a Pipeline instead of transforming the whole DataFrame first?

🎯 **In one line.** So that their statistics are learned only from training folds, avoiding leakage into validation and test, and so the identical transformation is applied at prediction time.

⚙️ **The full answer.** Scaling before splitting uses test means and variances; small with big data, but a real bias with 215 rows, and it is the kind of mistake interviewers look for.

Q27.4

Read this confusion matrix from scikit-learn: $\begin{bmatrix} 17 & 4 \\ 6 & 16 \end{bmatrix}$. Compute precision, recall and F1 for the positive class.

🎯 **In one line.** Rows are actual, columns predicted, negative class first: TN 17, FP 4, FN 6, TP 16; precision $16/20 = 0.80$, recall $16/22 = 0.727$, F1 = 0.762.

Q27.5

fit, transform, fit_transform, predict, predict_proba: explain each.

🎯 **In one line.** fit learns parameters from data; transform applies a fitted transformer; fit_transform does both (training data only); predict outputs labels or values; predict_proba outputs class probabilities.

⚠️ **Trap to avoid.** Calling fit_transform on the test set (re-fits on test: leakage and inconsistent features).

Q27.6

Why stratify=y and random_state?

🎯 **In one line.** stratify keeps class proportions equal across train and test (vital for imbalanced or small data); random_state makes results reproducible.

Q27.7

How do you tune hyperparameters properly?

🎯 **In one line.** With GridSearchCV or RandomizedSearchCV over the whole Pipeline on the training set with stratified CV and a metric that matches the goal, then evaluate the chosen model once on the test set.

📊 **Worked example.** {"lr__C": [0.01, 0.1, 1, 10]}: parameter names are step__param.

Q27.8

💡 **Think deeper** Your test accuracy is 0.77 on 43 rows. How confident are you?

🎯 **In one line.** Not very: the standard error is about $\sqrt{0.77 \cdot 0.23/43} \approx 0.064$, so the true accuracy could plausibly be anywhere from about 0.64 to 0.90; cross-validation (0.83 ± 0.05 AUC) gives a

steadier estimate.

Q27.9

How do you handle a categorical value in test that never appeared in training?

🕒 **In one line.** `OneHotEncoder(handle_unknown="ignore")` encodes it as all zeros; alternatives are grouping rare categories into “other” at training time or using encoders that map unknowns to a default.

Q27.10

Logistic regression vs random forest on this small dataset: which do you present and why?

🕒 **In one line.** Logistic regression: cross-validated AUC is equal or better (0.83 vs 0.82), it is interpretable (odds ratios), calibrated and less prone to overfit 215 rows.

Q27.11

How do you extract feature names and coefficients from a fitted pipeline?

🕒 **In one line.** `pipe.named_steps["pre"].get_feature_names_out()` with `pipe.named_steps["lr"].coef_[0]`; sort by absolute value; for tree models use `feature_importances_` or `permutation_importance`.

Q27.12

💡 **Think deeper** The interviewer asks you to predict salary for placed students. What changes?

🕒 **In one line.** It becomes a regression on the placed subset only (salary is not defined for others): use a log target or Gamma GLM because salary is skewed, evaluate with MAE/MAPE, and remember the model applies only to students who get placed (a two-stage model: placement probability, then salary given placement).

⚙️ **The full answer.** Expected salary = $P(\text{placed}) \times \mathbb{E}[\text{salary} \mid \text{placed}]$. That is the hurdle model from Chapter 4.

Q27.13

How do you save and reload a trained pipeline?

🕒 **In one line.** `joblib.dump(pipe, "model.joblib")` and `joblib.load`; save the whole Pipeline (preprocessing + model), record the library versions, and never unpickle untrusted files.

Q27.14

 **Think deeper** Accuracy, precision or recall: which would you optimise for this placement model?

🎯 **In one line.** It depends on the decision: for allocating extra coaching to at-risk students, recall of the “Not Placed” class; for a recruiter shortlist predicted to be placed, precision of “Placed”; since classes are balanced, accuracy is a fair overall summary, and AUC compares models independent of threshold.

Part VI

Round 4: Case Studies and the Final Conversation

CHAPTER 28

How to Answer a Data Science Case Study

🎥 **Companion video for this chapter:** [InfoEdge DS interview experience: Round 4 case studies \(at 9:07\)](#)

Round 4 (45–60 minutes, a senior data scientist): “A case study based on a real world data science problem to assess the candidate’s talent in applying machine learning theory to practice. The candidate shall be assessed on her/his ability to understand data related challenges, modelling and design choices as well as model assessment skills. There is generally no single best solution. The interviewer will rather assess the candidate on her/his thought process and on how she/he may fare in a team-working environment.” The 2026–27 talk added: “an AI case study: asked how you would approach the problem, what type of data you will need”, and that they want people “who actually want to understand business”.

This chapter gives you a repeatable structure. Chapter 29 applies it to InfoEdge’s own products.

28.1 The ten-step skeleton

Memorise it; say the steps aloud as you go, so the interviewer can follow and steer.

1. **Clarify the goal and the decision.** What business outcome? Who acts on the prediction, and how? (“Rank leads for 50 telecallers who make 200 calls each per day.”) If clarifying questions are not allowed, *state your assumptions* explicitly.
2. **Define success.** One primary business metric (conversions per call, applies per job, revenue), guardrails (complaints, unsubscribes, latency), and the offline proxy metric you will optimise (precision@k, NDCG, AUC, MAPE).
3. **Frame the ML problem.** Classification, regression, ranking, recommendation, clustering, forecasting, anomaly detection? What is one row (user, user–day, user–listing pair)? What is the label and over what time window? When is the prediction made?
4. **Data.** What data exist (logs, profiles, listings, CRM, call outcomes)? What is missing? How are labels obtained (natural outcomes, human labels, proxies)? Biases in how data were collected (only called leads have call outcomes: selection bias).
5. **Features.** Brainstorm by entity and by time window: user, item, user×item, context, recency/frequency/monetary, trends, ratios; text and image embeddings. Check leakage for every feature: “is this known at prediction time?”
6. **Baseline.** A rule or simple model the business already uses or would accept (most-recent-activity rule, popularity, logistic regression).
7. **Model.** Justify the choice by data type, size, latency, interpretability (GBM for tabular, two-tower for retrieval, Transformer for text). Handle imbalance and cold start.
8. **Offline evaluation.** Time-based split, the right metric at the operating point, segment analysis (new vs returning users, cities), calibration.
9. **Online evaluation and deployment.** A/B test with a clear hypothesis and sample size; batch vs real-time scoring; monitoring (data drift, prediction drift, metric decay); retraining cadence; feedback loops.

10. **Risks and iteration.** Fairness, privacy, gaming, cold start, exploration vs exploitation; next steps.

💡 Intuition: why structure wins

The interviewer is imagining you in their team next month. A structured answer shows you will not jump into modelling before knowing what decision the model serves, and that you will catch leakage and evaluation mistakes on your own.

28.2 Feature engineering: a generator of ideas

When the interviewer “keeps pushing until you run out of answers” (as reported for the 99acres question), use this grid to generate features systematically.

Dimension	Examples (99acres buyer intent)
Recency	days since last session, since last enquiry
Frequency	sessions in last 7/30 days, listings viewed, searches
Depth / intensity	time on listing pages, photos viewed, floor plans opened, EMI calculator used, contact number revealed
Consistency (concentration)	share of views within the dominant price band (the price-band concentration feature that impressed the interviewer), share in top locality, entropy of property types viewed (low entropy = focused buyer)
Trend	views this week vs last week; narrowing of price range over time
Context	device, time of day, city tier, festival season, interest rates
Profile	stated budget, loan pre-approval, family size, owner vs tenant
User × item	distance of viewed listings from workplace; match between budget and viewed prices
Social / market	inventory in the target locality, price trend of that locality
Text	search queries (“ready to move”, “near metro”), enquiry messages

For each feature, state the hypothesis (“a narrow price band means a decided buyer”) and the leakage check.

28.3 Metrics: from business to offline

Business decision	Offline metric	Online metric
Who to call with limited capacity	precision@k, expected conversions in top k	conversions per call-hour
Rank search results	NDCG@10, MRR, recall@k for retrieval	CTR, contacts/applies per search, time to first contact
Recommend jobs	recall@k, NDCG, coverage, novelty	applies per user, recruiter responses
Flag fraud/spam	recall at fixed precision, PR AUC	fraud caught, false-positive complaints
Estimate salary	MAPE/RMSLE, interval coverage	user trust, engagement with salary card
Forecast demand	MAPE, bias	staffing efficiency

28.4 A/B testing in one page

- **Hypothesis:** the new lead scorer increases conversions per call by at least 5%.
- **Randomisation unit:** caller or lead (avoid interference: two variants sharing one caller's queue contaminate each other).
- **Sample size:** depends on baseline rate, minimum detectable effect, significance α (0.05) and power (0.8). For a proportion p and absolute effect δ , roughly $n \approx 16p(1-p)/\delta^2$ per arm. Baseline 5%, detect +0.5 points: $16 \cdot 0.0475/0.000025 \approx 30,400$ per arm.
- **Guardrails, novelty effects** (run long enough), **multiple testing** (pre-register the primary metric), **peeking** (do not stop at the first significant day).

28.5 Product and business questions

InfoEdge is a marketplace company: Naukri (jobs), 99acres (real estate), Jeevansathi (matrimony), Shiksha (education), plus investments. Its revenue is mostly from the paying side of each marketplace (recruiters buying job postings and Resdex database access; property developers and agents buying listings and leads). Marketplaces have **network effects**: more recruiters attract more seekers and vice versa. Good product answers grow both sides, improve match quality (which raises willingness to pay), and use data products (Talent Pulse) to open new revenue. A useful skeleton for “grow revenue”: *more customers* (acquisition, new segments), *more revenue per customer* (upsell, pricing, premium tiers), *less churn* (better outcomes, success management), *new products* (data/insight products, AI tools).

28.6 Communication and team fit

Think aloud; summarise every few minutes; ask “shall I go deeper here or move on?”; accept pushback gracefully (“good point; that would bias the labels; I would fix it by...”); do not claim certainty you do not have; quantify where possible; end with a crisp summary and next steps.

28.7 Framework questions

Q28.1

★ **PYQ (Round 4 format: no clarifying questions allowed)** How do you structure a case answer when the interviewer says you may not ask clarifying questions?

🕒 **In one line.** State your assumptions explicitly and briefly, choose the most plausible interpretation, and keep going with the usual structure; mention how the answer would change under a different assumption.

⚙️ **The full answer.** The reported Naukri revenue question disallowed questions to test company knowledge and decisiveness. “I will assume most revenue comes from recruiters buying postings and database access, and that growth in seekers drives recruiter value. If instead the focus were seeker subscriptions, I would...”

Q28.2

What are the most common mistakes in ML case interviews?

🎯 **In one line.** Jumping to a model before defining the decision and metric, ignoring how labels are created, missing leakage, using random splits on temporal data, choosing the wrong metric (accuracy on imbalanced data), forgetting deployment and monitoring, and never connecting back to the business outcome.

Q28.3

🧠 **Think deeper** Your model will change who gets shown which jobs, which changes future training data. Why is that a problem and what do you do?

🎯 **In one line.** It creates a feedback loop: the model only observes outcomes for items it chose to show, so its training data become biased toward its own past choices and it can lock in errors or popularity bias; counter with exploration (randomised or bandit traffic), propensity-weighted training and logging of what was shown.

⚙️ **The full answer.** Position bias too: users click top results regardless. Use inverse propensity scoring or position-aware models; reserve a small random slice of traffic for unbiased evaluation.

Q28.4

How do you choose the label for “will this user buy a property”?

🎯 **In one line.** Define an observable event and a window tied to the decision: for example “a site visit booked or a transaction recorded within 30 days of the prediction date”; check that it is measured reliably, that it is available for all users (not only those called), and that the window matches the business cycle.

⚙️ **The full answer.** A purchase may happen offline and never be logged; proxies (site visit, developer-confirmed lead) may be better measured. Different windows give different models; choose by how long telecalling can usefully wait.

Q28.5

🧠 **Think deeper** Offline evaluation looks great but you can only A/B test with 5% of traffic. How do you de-risk?

🎯 **In one line.** Use a careful time-based offline backtest, shadow deployment (score live traffic without acting), a small ramped A/B test with guardrails and pre-computed power, interleaving for ranking (very sensitive with little traffic), and clear rollback criteria.

Q28.6

How would you explain a model’s output to a non-technical stakeholder, such as a sales head?

🎯 **In one line.** Translate into the decision and its impact (“calling the top 200 leads converts 3 times more than the current list”), show a few concrete examples with the top reasons per lead, and state limitations plainly.

Q28.7

Estimate the sample size for an A/B test with baseline conversion 5% and a minimum detectable absolute lift of 0.5 percentage points.

🎯 **In one line.** About $16p(1-p)/\delta^2 = 16 \times 0.0475/0.005^2 \approx 30,400$ users per arm (for 5% significance and 80% power).

Q28.8

🧠 **Think deeper** When would you *not* use machine learning for a business problem?

🎯 **In one line.** When a simple rule achieves nearly the same outcome, when there is not enough reliable labelled data, when the cost of errors or lack of explainability is unacceptable, or when the decision is rare and better made by people; ML must beat a baseline by enough to justify maintenance cost.

CHAPTER 29

Worked Case Studies from InfoEdge's Own Products

🎥 **Companion video for this chapter:** [InfoEdge DS interview experience: the 99acres and Naukri cases](#)

Each case below is written as a full answer you could give in 15–25 minutes, following the ten-step skeleton of Chapter 28. The first two were reported from real Round 4 interviews; the rest are built from systems InfoEdge described in its 2026–27 pre-placement talk (Resdex search, the job–CV matching engine, 99acres ML telecalling, Jeevansathi recommendations, the Data Factory salary engine, the SilverSkAI taxonomy, AI REX, the content feed). Practise them aloud with a friend who keeps asking “and then?”.

Q29.1

★ **PYQ (Round 4, reported)** Given a user's full browsing data on 99acres, design features (and a model) to predict whether they will buy a property.

🕒 **In one line.** Frame it as a time-windowed binary classification (“will this user transact or book a site visit within N days”), engineer intent features along recency, frequency, depth, consistency (price-band and locality concentration) and trend, train a gradient-boosted classifier on a time-based split, and evaluate by precision at the top of the ranked list that sales will act on.

⚙️ **The full answer.** **1. Goal and decision.** Prioritise high-intent users for telecalling, developer leads or personalised offers. The decision is a ranking, so the top of the list matters most. **2. Label.** “Bought” is often offline; use the best observable proxy: confirmed transaction (developer CRM), site visit booked, or a lead marked “converted” by an agent, within 30–60 days after the prediction date. Note selection bias: outcomes are mostly observed for users who were called or who enquired. **3. Row.** One user–date snapshot; features computed only from data before the snapshot date (no leakage). **4. Features** (I would generate them systematically):

- *Recency and frequency:* days since last visit; sessions, searches and listing views in the last 1/7/30 days; returning after a long gap (a life event).
- *Depth:* time on listing pages, photo and floor-plan views, EMI calculator use, shortlists, contact-number reveals, enquiries, repeated visits to the same listing.
- *Price-band concentration* (the feature that impressed the interviewer): the share of viewed listings within $\pm 15\%$ of the median viewed price; also the coefficient of variation of viewed prices. Someone who looks mostly at 30-lakh flats with a few 10- and 60-lakh outliers is a 30-lakh buyer; the outliers are noise. Narrowing price spread over time signals a decision forming.
- *Location focus:* localities of the last 10–15 searches; entropy of locality distribution (low = focused); distance between viewed listings and the user's city; a Noida searcher will not buy in Bangalore however similar the listing.
- *Property-type focus:* BHK and type entropy; ready-to-move vs under-construction preference.
- *Stage signals:* loan-eligibility checks, legal or registration guides read, “ready to move”

filters, visiting builder pages.

- *Profile*: stated budget vs viewed prices (mismatch means early stage), owner vs tenant, account age.
- *Market context*: inventory and price trend in the focus locality; festive season.

5. Baseline. Rule: “enquired in last 7 days”. **6. Model.** LightGBM on these features (handles non-linearity, missing values, mixed scales); class weights for the rare positives. **7. Offline evaluation.** Train on earlier months, test on later; metrics: precision@k for k equal to daily call capacity, PR AUC, calibration (to estimate expected conversions); segment by city and by new vs returning users. **8. Online.** A/B test against the rule-based list: conversions per call; guardrails: user complaints about calls, unsubscribes. **9. Risks.** Feedback loop (only called users get labels: reserve random calls), privacy (consent for calls), seasonality, drift after price changes.

 **If the interviewer pushes.** “What if two users have identical behaviour but one is a broker?” Add account-type and behaviour features (many cities, many enquiries per day) and possibly a separate broker classifier, because brokers are a different business line. “How would you handle someone who searches only at night on mobile?” Context features; but ask whether it causally relates to intent or is just a demographic proxy.

Q29.2

★ **PYQ (Round 4, reported)** You are the sales head of Naukri.com. Give three things you would do to increase revenue. No clarifying questions allowed.

 **In one line.** Assuming most revenue comes from recruiters paying for job postings and Resdex database access, and that recruiter value depends on finding good candidates quickly, I would (1) grow and activate the seeker side to raise match quality, (2) increase revenue per recruiter with outcome-focused premium products and smarter pricing, and (3) open new recruiter segments such as campus and SME hiring through partnerships, all measured by revenue per recruiter, renewal rate and time to hire.

 **The full answer.** **Framing.** Naukri is a two-sided marketplace with a flywheel: more and better seekers → more successful hires → recruiters pay and renew → more jobs → more seekers. Revenue = number of paying recruiters × revenue per recruiter × retention. **1. Strengthen the seeker side (match quality drives willingness to pay).** Reactivate dormant profiles with personalised job alerts on the channel each seeker prefers (the “seeker acquisition and reactivation engine” and multi-channel touchpoints: mailers, notifications, WhatsApp); improve profile completeness with AI resume help; this increases applies per job, the metric recruiters actually buy. **2. Revenue per recruiter.** Outcome-based and tiered products: PremiumX-style premium discovery with explainable shortlists; AI REX-style automated sourcing and screening sold as time-to-hire reduction; usage-based Resdex pricing for heavy users; upsell based on predicted hiring need (a model of which accounts are likely to expand). **3. New segments and partnerships.** Campus hiring: partner with colleges by giving them a mini placement portal on Naukri where companies post openings (the candidate’s reported idea), which brings fresher supply and campus-recruiting budgets; SME self-serve hiring with low-cost packages; data products (Talent Pulse) sold to enterprises and HR consultancies. **Measurement.** For each, an experiment with revenue per account, renewal rate, and a guardrail on seeker experience (spam complaints).

 **If the interviewer pushes.** “Which one first?” The one with the fastest, most measurable

impact: reactivation of seekers via channel-optimised alerts, because it lifts applies per job for all existing recruiters, which improves renewals.

Q29.3

 **Think deeper** 99acres ML-based telecalling (from the 2026–27 slide): design the system that turns users and projects into a ranked call list of high-intent leads.

🎯 **In one line.** Three models in sequence, as on the slide: infer each user's requirement (budget, location, configuration), match users to project bundles, and rank user-bundle pairs by predicted conversion (and value), then fill the callers' queues from the top, measuring conversions per call and learning from call outcomes.

⚙️ **The full answer.** **User requirement model.** From browsing and search history, predict the user's budget band, preferred localities and BHK (multi-output classification or simply robust statistics such as the mode of the price band and localities weighted by recency). Output a structured requirement with confidence. **User-project matching.** Retrieve projects that satisfy hard constraints (city, budget $\pm$ tolerance, BHK) and score soft fit (amenities, distance to work, developer reputation); a two-tower or GBM matcher trained on past enquiries and site visits. Bundle 2–3 projects per user so one call can pitch several. **User ranking.** Predict $P(\text{site visit or conversion} \mid \text{call})$ for each user-bundle with LightGBM; multiply by expected value (project commission) for revenue-aware ranking; calibrate so expected conversions can be planned. **Operations.** Capacity: C callers $\times k$ calls/day; assign by language and city; avoid re-calling too soon; respect consent and do-not-call lists. **Labels and bias.** Only called users have outcomes: reserve a small random-call slice for unbiased training and evaluation; use call dispositions (no answer, not interested, visit booked). **Evaluation.** Offline: precision@k on a later period. Online: A/B test conversions per call-hour against the current list; guardrails: complaint rate, user unsubscribes. **Iteration.** Learn the best time to call (uplift by hour), and which users the call actually influences (uplift modelling) rather than who would convert anyway.

💡 **If the interviewer pushes.** “Why uplift instead of conversion probability?” Calling a user who would buy anyway wastes a call; the business value is the *increase* in conversion caused by the call. Uplift models (two-model or class-transformation approaches) need randomised call/no-call data.

Q29.4

 **Think deeper** Resdex search: a recruiter types “sr data scientist blr 5+ yrs python”. Design the search system (the slide's query intelligence $\rightarrow$ retrieval + ranking $\rightarrow$ personalisation).

🎯 **In one line.** Parse and normalise the query into structured intent with the taxonomy, retrieve candidates with hybrid sparse and dense retrieval within hard filters, rank them with a learning-to-rank model per function segment, and personalise using the recruiter's past actions; evaluate with NDCG and recruiter actions per search.

⚙️ **The full answer.** **Query intelligence (QI).** Tokenise; map “sr” $\rightarrow$ Senior, “blr” $\rightarrow$ Bengaluru, “5+ yrs” $\rightarrow$ experience ≥ 5 , “python” $\rightarrow$ skill; identify department (Data Science) and premium/-targeting needs (the slide's M1–M4: taxonomy, department, premium, targeting models). Use

dictionaries, NER and the SilverSkAI taxonomy; expand synonyms (“ML”=“machine learning”). **Retrieval.** From 30 lakh+ daily-active candidate pool: sparse BM25 on skills/titles (precise, handles rare skills), dense bi-encoder embeddings (semantic: “data scientist” $\approx$ “ML engineer”), hard filters (location, experience); merge (reciprocal rank fusion); the slide shows separate retrieval models per segment (finance, core engineering, sales, IT) and per field (skill, title, skill+title), plus “company boost” and “catch-all” streams. **Ranking.** LambdaMART/GBM per segment (the “M10” rankers) on features: skill match coverage, title similarity, experience fit, notice period, current CTC vs budget, recency of profile update, activity (likely to respond), past recruiter engagement with similar profiles, profile completeness. Labels from recruiter actions (view < download/contact < shortlist), debiased for position. **Personalisation.** Recruiter-specific preferences (enterprise vs consultant, the slide’s “ER” and “Consultant” models): preferred colleges, companies, past shortlists. **Evaluation.** Offline NDCG@10 on held-out searches; online: contacts per search, time to first contact, recruiter retention. **Constraints.** 16 lakh searches/day: latency budget about 100–200 ms; precompute candidate embeddings; ANN index; cache popular queries. **Risks.** Fairness (avoid proxies for gender/age), freshness (stale profiles), gaming (keyword stuffing).

Q29.5

 **Think deeper** Design the job–CV matching engine that decides which applicants a recruiter sees first in the apply pool.

 **In one line.** Score every application for fit to the job with a model trained on recruiter feedback (shortlist vs reject), combining structured matching features with text-similarity from a fine-tuned cross-encoder, then rank the pool and explain the top reasons.

 **The full answer.** **Labels.** Recruiter actions on applications: shortlisted/contacted (positive), rejected (negative), ignored (weak negative; position-biased). **Features.** Must-have skills coverage (via taxonomy), experience range fit, location/relocation, salary expectation vs budget, education requirements, title trajectory, notice period, semantic similarity of JD and CV (bi-encoder cosine; cross-encoder score for the top part), application timing. **Model.** GBM ranker or a fine-tuned cross-encoder whose score is a GBM feature; per-function models if behaviour differs. **Evaluation.** NDCG@20 of shortlisted applicants; online: recruiter time to shortlist, shortlist rate of top-20. **Explainability.** SHAP or rule-based reasons (“matches 8/9 required skills; 6 yrs vs 5–8 required”), mirroring PremiumX’s “AI reasons” and “search highlights”. **Fairness.** Exclude protected attributes and audit proxies; monitor selection rates by group.

Q29.6

 **Think deeper** Jeevansathi: recommend profiles to a user (slide: preferences → pool generation → content, activity and collaborative models → ranked profiles → filters and personalisation).

 **In one line.** Generate a candidate pool that satisfies stated preferences, score each candidate with content (profile similarity), activity (likelihood to respond) and collaborative (people like you liked) models combined in a ranker that predicts *mutual* interest, then apply filters and diversity, and evaluate by mutual acceptances.

⚙️ **The full answer.** **Two-sided:** success requires both users to accept, so model $P(A \text{ likes } B) \times P(B \text{ likes } A)$ (reciprocal recommendation), otherwise popular profiles are flooded and everyone else ignored. **Pool:** hard preferences from the slide's example (community, cities), relaxed slightly if the pool is small. **Models:** content (age, education, profession, location similarity), activity (online recency, response rate), collaborative (matrix factorisation or two-tower on interest/accept interactions). **Ranking:** GBM over these scores and features, optimising mutual acceptance; diversity and fairness of exposure. **Metrics:** interests sent, acceptance rate, mutual matches, conversations started; guardrails: reports/blocks. **Trust:** fake-profile detection upstream (photo checks, behaviour).

Q29.7

💡 **Think deeper** Build the salary engine behind “Salary Intelligence” (Data Factory/-TopTier slides): given title, city, experience, company, skills, show the typical salary and a range.

🕒 **In one line.** Clean and validate salary data, normalise titles and companies with the taxonomy, model log salary (or a Gamma GLM / GBM) for the centre and quantile models for the 10th–90th band, shrink sparse company estimates toward their industry, and validate with MAPE and interval coverage by segment.

⚙️ **The full answer.** **Data:** self-reported CTCs (noisy, MNAR: high earners hide), offers, job-post salary ranges. Clean units (LPA vs monthly), remove impossible values, deduplicate. **Normalisation:** title and company to canonical IDs (SilverSkAI); experience bands. **Model:** GBM on log salary with features title, function, city tier, company (target encoding with shrinkage), skills (embedding or top-k binary), education, industry; quantile GBMs for the band; enforce monotonicity in experience. **Sparse cells:** hierarchical shrinkage (company → industry → function). **Evaluation:** MAPE/RMSLE on held-out recent reports; 80% interval should cover about 80% of truths; bias checks by city and gender. **Product:** “you earn more than 68% of software engineers in Delhi” is a percentile from the predicted conditional distribution; communicate uncertainty when data are thin.

💬 **If the interviewer pushes.** “Heteroscedasticity?” Expected: spread grows with level; log target or Gamma GLM handles it (Chapter 4).

Q29.8

💡 **Think deeper** SilverSkAI taxonomy: normalise 50 lakh raw skill strings (“ML”, “machine-learning”, “M.L.”, “Machine Larning”) into about 70K canonical skills.

🕒 **In one line.** Clean strings, generate candidate canonical matches with character n-gram and embedding nearest-neighbour search, decide matches with a trained classifier using string, semantic and co-occurrence features, route uncertain cases to human review, and keep the canonical set curated with clustering for discovering new skills.

⚙️ **The full answer.** **Steps:** (1) normalise case, punctuation, whitespace; expand abbreviations from a dictionary; (2) candidate generation: TF-IDF on character 3-grams + cosine, FastText/BERT embeddings + ANN, edit distance; (3) features per (raw, canonical) pair: edit distance, Jaccard of tokens, embedding cosine, how often they co-occur with the same other skills in resumes,

frequency; (4) classifier (GBM) trained on human-labelled pairs; (5) thresholds: auto-map high confidence, human-review medium, new-skill candidate otherwise; (6) cluster unmapped strings to propose new canonical skills ("LangChain" when it appears). **Evaluation:** precision/recall of mappings on a labelled sample, weighted by string frequency; downstream: search recall improvement. **Pitfalls:** ambiguous strings ("Java" island vs language; "Spark" Apache vs automotive) need context (the rest of the resume); hierarchy ("PyTorch" is-a "Deep learning framework").

Q29.9

 **Think deeper** Spam and fake recruiters (the panel's own example): design detection for job postings and recruiter messages.

 **In one line.** Treat it as rare-event classification combining content (text of posts and messages), behaviour (posting velocity, payment, response patterns) and network signals, with recall-focused but precision-guarded thresholds feeding automated blocks and a human review queue, evaluated with recall at fixed precision and refreshed often because adversaries adapt.

 **The full answer.** **Labels:** user reports, moderation decisions, confirmed frauds (asking candidates for money, fake companies). Noisy and delayed. **Features:** text (TF-IDF/BERT: "registration fee", "WhatsApp only", unrealistic salary for title); recruiter account age, verification status, company domain match, posts per hour, near-duplicate posts across accounts, apply-to-response ratio, candidate complaints, payment method. **Model:** GBM on tabular + text model scores; anomaly detection (Isolation Forest) for new patterns. **Metrics:** the panel's point: *look at recall, not accuracy*; also precision at the review queue capacity; time to detection. **Actions:** tiered: block (very high score, high precision), hold for review, monitor. **Adversarial:** rotate features, monitor drift weekly, red-team.

Q29.10

 **Think deeper** AI REX: how would you evaluate the screening agent that decides which candidates move forward after automated pre-screening?

 **In one line.** Measure it like a classifier against human judgement and hiring outcomes: recall of genuinely suitable candidates (not losing good people), precision of those forwarded (recruiter time), agreement with expert recruiters on a labelled sample, calibration, fairness across groups, and end-to-end time to hire in an A/B test.

 **The full answer.** **Offline:** sample conversations and answers; expert recruiters label fit/no-fit; compute precision/recall/F1, Cohen's kappa vs humans; error analysis by job type and language. **Online:** A/B (AI REX vs standard process) on time to shortlist, interview-to-offer ratio, recruiter satisfaction, candidate drop-off and complaints. **Guardrails:** no use of protected attributes; adverse-impact ratio monitoring; human approval for rejections in sensitive roles; audit logs; prompt-injection tests (candidate answers trying to manipulate the agent). **Calibration agent loop:** if forwarded candidates are too few, trigger more sourcing (as the slide notes).

Q29.11

 **Think deeper** Seeker reactivation: which dormant job seekers should receive which message on which channel (mail, notification, WhatsApp)?

🎯 **In one line.** Model, for each seeker and channel, the uplift in the probability of returning and applying caused by a message (not just the probability of returning), choose the best channel and content per seeker under frequency caps, and test with holdout groups.

⚙️ **The full answer.** Propensity models target people who would return anyway; uplift models target persuadable users. Features: last activity, past response by channel, profile freshness, market demand for their skills, time of day. Content: personalised job recommendations. Constraints: daily caps, unsubscribe risk (guardrail), WhatsApp cost. Evaluation: randomised holdout; metrics: reactivation rate, applies per reactivated user, unsubscribe rate.

Q29.12

 **Think deeper** Shiksha: recommend colleges to a student who has filled a short form (exam score, preferred city, course, budget).

🎯 **In one line.** Filter colleges by hard eligibility (cut-offs, course, budget), score fit with a model trained on past students' applications and admissions (content features plus collaborative signals from similar students), balance "reach, match, safety" options, and evaluate by applications and admissions.

⚙️ **The full answer.** Cold start for new students is the norm (one form): content-based and similar-student (KNN on profile) methods dominate; popularity priors; explain recommendations ("students with your rank were admitted here last year"). Fairness: avoid steering by socio-economic proxies.

Q29.13

★ **PYQ (Round 4 follow-up, reported)** After your feature list, the interviewer asks: "What are the assumptions of linear regression, and what do you do if homoscedasticity is violated?" Answer as you would in the room.

🎯 **In one line.** Linearity, independent errors, constant variance, normal errors, no perfect multicollinearity and exogeneity; under heteroscedasticity coefficients stay unbiased but standard errors are wrong, so I would transform the target (log or Box-Cox), use weighted least squares or robust standard errors, or model the variance directly with a GLM such as Gamma with a log link.

⚙️ **The full answer.** Full answers are in Chapters 2 and 4; in the room, mention both the transformation the interviewer expected and the GLM the candidate proposed, and say when you would pick each.

Q29.14

 **Think deeper** Content feed for job seekers (slide: insights generation agent → content manifestation agent → formats → personalised recommendation). How would you make sure the generated insights are true and the feed is useful?

🎯 **In one line.** Generate insights from verified aggregates in the Data Factory with statistical checks (sample size, significance, period-over-period), let the LLM only phrase pre-computed facts with citations to the numbers, review automatically for factual consistency, and rank content per user by predicted engagement and usefulness, evaluated by dwell time, saves and downstream applies.

⚙️ **The full answer.** Example insight from the slide: “AI jobs up 20% in Kolkata, flat in Pune and Chennai”. Risks: small-number noise (20% of 50 jobs), hallucinated numbers, clickbait optimisation. Guardrails: minimum counts, confidence intervals, numeric consistency checks between the generated text and source tables.

CHAPTER 30

The Final Conversation: Resume Drill, Company Knowledge and HR

🎥 **Companion video for this chapter:** [InfoEdge DS interview experience: HR round and “Why InfoEdge” \(at 12:26\)](#)

30.1 Everything on your resume will be drilled

The hire’s interviewers followed his projects (PCA, t-SNE) into depth. For every project and technique on your resume prepare:

1. A 60-second story: problem, why it mattered, data, approach, your specific contribution, result with a number, what you would do differently.
2. The maths of every method you name, at the level of this book.
3. The baseline you beat and how you evaluated (split, metric, significance).
4. One failure or surprise and what you learned.

30.2 Company knowledge

InfoEdge (India) Ltd: founded in the mid-1990s by Sanjeev Bikhchandani; Naukri.com (launched 1997, India’s largest job portal), Jeevansathi.com (matrimony), 99acres.com (real estate), Shiksha.com (education), and an investment arm with stakes in companies such as Zomato and PolicyBazaar. Listed on Indian stock exchanges in 2006. *Verify years and figures from InfoEdge’s annual report before quoting them.* From the 2026–27 talk: 100 crore+ applies a year through the job-posting engine; Resdex serves about 60K recruiters and 16 lakh searches a day; the Data Factory runs 400+ AI models; GenAI work includes in-house fine-tuned and distilled LLMs, voice agents and the AI REX agentic hiring product. Work model (as reported): mostly in-office.

30.3 “Why InfoEdge?”

The hire called this the most important HR question. A strong answer has three specific parts:

1. **Real problems at scale:** ML that touches crores of job seekers and lakhs of recruiters; search, ranking, matching and recommendation are core, not side projects.
2. **Technical depth:** cite something from the pre-placement talk (hybrid retrieval in PremiumX, the SilverSkAI taxonomy, AI REX’s agents, in-house LLMs) and why it excites you.
3. **Breadth and business:** several products (jobs, real estate, matrimony, education) let you learn many problem types; the company wants data scientists who understand business, which matches what you want to become.

30.4 Questions to ask them

How do teams decide between a classical ranker and an LLM-based approach given latency and cost? How is model impact measured (A/B culture)? How do new data scientists ramp up? What does a successful first six months look like?

30.5 Questions

Q30.1

Tell me about yourself (for a data scientist role at InfoEdge).

🎯 **In one line.** A 60–90 second story: who you are academically, the two or three ML experiences most relevant to search, ranking, NLP or tabular modelling with one concrete result each, and why that leads you to InfoEdge now.

⚙️ **The full answer.** Example structure: “I am an M.Tech AI student at IIT Jodhpur. In my project on X I built Y, which improved Z by N%. I enjoy problems where models meet users at scale, which is why the matching and search work you presented, especially hybrid retrieval in PremiumX, is exactly what I want to work on.”

Q30.2

Describe a project where your model did not work as expected. What did you do?

🎯 **In one line.** Pick a real case, explain the symptom (for example great offline, poor online), how you diagnosed it (leakage, distribution shift, wrong metric), what you changed, and what you now do by default because of it.

Q30.3

How do you keep up with ML, and what recent development do you find most useful for a company like InfoEdge?

🎯 **In one line.** Name concrete sources and one development with a clear application: for example retrieval-augmented generation with hybrid search for grounded job-discovery chat, or distillation of LLMs into small models to serve ranking features cheaply.

Part VII

The InfoEdge Hire's Checklist: EDA, Extra ML/DL, Time Series, Probability, Statistics and Linear Algebra

Part VII Guide: the Checklist from an InfoEdge Data Scientist

A candidate who secured an InfoEdge data scientist offer shared the checklist below of what to know for the interviews. This page maps every item to where it is taught. Items that the earlier chapters already covered fully point there; items that were missing or only touched on are taught in the new chapters 31–40, which follow the same pattern as the rest of the book: concepts from zero, many small examples, a “where this is used in ML” section, and interview questions with full answers at the end. Probability, statistics and linear algebra are also covered at OA depth in your separate Maths book; here they are taught again with an interview and ML-application focus.

Checklist item	Where
EDA	
Reason behind every step; data visualisations and how to read each graph	Ch. 31 (new), Ch. 26
Hands-on EDA in pandas/NumPy on real datasets (placement, HR, Titanic, Pima)	Ch. 41–44 + notebook
Univariate statistics: mean, median, mode, range, standard deviation, variance	Ch. 31 (new)
Missing values: MCAR/MAR/MNAR, dropping, mean/median/mode imputation pros and cons	Ch. 13, Ch. 33 (new)
ML models to predict missing values (KNN, MICE, MissForest, XGBoost imputer)	Ch. 33 (new), Ch. 13
Leaky variables	Ch. 1, Ch. 27, Ch. 33 (new)
Outliers: z-score, 68–95–99.7 rule, IQR formula, box plot interpretation	Ch. 33 (new), Ch. 13
Scaling: standardisation vs normalisation, min–max, robust	Ch. 34 (new)
Numerical to categorical: binning, clustering	Ch. 34 (new)
Categorical to numerical: ordinal, one-hot, target, frequency encoding; dummy variable trap	Ch. 34 (new)
Imbalanced data: under/oversampling, weighting, NearMiss, random deletion, SMOTE, Borderline-SMOTE, ADASYN	Ch. 13, Ch. 34 (new)
Bivariate: Pearson, Spearman (ranks), point-biserial, phi; correlation matrix	Ch. 32 (new)
Multivariate: multicollinearity, VIF in terms of R	Ch. 32 (new), Ch. 2
Skewed data: log, power, Box–Cox transforms	Ch. 34 (new), Ch. 31
Basic ML	
OLS, 5 assumptions (detect and fix), MLE for linear regression, loss	Ch. 2, Ch. 12
L1, L2, elastic net; Huber; homo-/heteroscedasticity	Ch. 1, Ch. 2, Ch. 4
GLM maths; Poisson and logistic as GLMs	Ch. 4, Ch. 3
SVM, hinge, kernels (polynomial, Gaussian/RBF), Mercer, SVR, margin, hyperplane	Ch. 9, Ch. 4
Naive Bayes (3 types), Laplace smoothing, independence assumption	Ch. 10
K-Means and variants (k-medoids, k-modes, k-prototypes, k-means++), soft vs hard clustering	Ch. 15, Ch. 35 (new)
GMM, EM, DBSCAN, Swiss roll	Ch. 15, Ch. 35 (new)
PCA (scree, explained variance, eigen intuition), SVD, Fisher LDA	Ch. 14, Ch. 10, Ch. 40 (new)
t-SNE, SNE, perplexity, Student-t; linear vs non-linear and deterministic vs non-deterministic DR	Ch. 14, Ch. 35 (new)
Decision trees, Gini, entropy, information gain, Extra Trees, RF, bagging, boosting, AdaBoost vs GB	Ch. 6–8
Bias vs variance with graphs; tackling under- and overfitting	Ch. 1
Deep Learning	
ANN, perceptron structure, activations, training, backprop	Ch. 18, Ch. 36 (new: perceptron algorithm)
GD variants (stochastic, batch, mini-batch), momentum, Adam; loss functions	Ch. 5, Ch. 18
Transfer learning, CNN layers, $\frac{n+2p-f}{s} + 1$	Ch. 19
RNN, GRU, LSTM	Ch. 20
Basics of CV, NLP, reinforcement learning, federated learning	Ch. 36 (new), Ch. 16–17
Time series: ANOVA, ARIMA, SARIMAX	Ch. 37 (new), Ch. 20
Statistics, Probability, Linear Algebra	
Descriptive vs inferential, population vs sample, CLT, p-value, hypothesis testing	Ch. 39 (new)
Distributions with real-life examples; PDF, PMF; expectation, variance, moments	Ch. 38 (new)
Skewness and kurtosis and their types	Ch. 31 (new), Ch. 38 (new)
Eigenvectors, eigenvalues, rank	Ch. 40 (new)

The checklist also says that what is asked in time series and applied deep learning “completely depends on the projects you put in your resume”. Use Chapter 30’s resume drill for those.

CHAPTER 31

EDA I: Univariate Analysis, Descriptive Statistics and Reading Graphs

🎥 **Companion video for this chapter:** [StatQuest: Histograms, Clearly Explained \(search\)](#)

The checklist's first line is "should know the reason behind all things". Exploratory data analysis (EDA) is not a ritual of plots: every statistic and every graph answers a specific question that changes a modelling decision. This chapter teaches the one-variable (univariate) toolkit from zero and, for every tool, says *why* you compute it and *what you would do differently* depending on the answer.

31.1 Why we do EDA at all

Before modelling, we need to know:

1. **Is the data what we think it is?** Units, ranges, types, duplicates, impossible values. (A percentage of 105 is an error; 98.5 is a topper.)
2. **What does each variable look like?** Centre, spread, shape, outliers. This decides scaling, transformations, robust losses.
3. **What is missing and why?** Decides imputation (Chapter 33).
4. **How do variables relate to the target and to each other?** Decides features, leakage checks, multicollinearity handling (Chapter 32).
5. **Is the target balanced?** Decides metrics and sampling (Chapters 11, 34).

Every finding should end in a sentence of the form "*because* I saw X, I will do Y".

31.2 Types of variables

Type	Meaning	Example	What you may compute
Nominal	categories with no order	city, job function	counts, mode
Ordinal	ordered categories, unequal gaps	seniority (junior < mid < senior), rating 1–5	median, percentiles, mode
Discrete numeric	countable values	number of applications	everything; often skewed
Continuous numeric	any value in a range	salary, MBA percentage	everything

Numeric scales are also called **interval** (differences meaningful, zero arbitrary: temperature in °C) or **ratio** (true zero: salary; "twice as much" makes sense).

31.3 Measures of central tendency: where is the middle?

Running example: salaries (lakh per year) of seven Data Analysts: 4, 5, 5, 6, 7, 9, 20.

📺 **Stuck on this page?** Watch: [StatQuest: Histograms, Clearly Explained \(search\)](#)

31.3.1 Mean

$\bar{x} = \frac{1}{n} \sum_i x_i = \frac{4+5+5+6+7+9+20}{7} = \frac{56}{7} = 8$. The balancing point: the sum of deviations $\sum(x_i - \bar{x})$ is always 0. It uses every value, so one extreme value (20) drags it: here 6 of 7 people earn less than the mean.

31.3.2 Median

Sort and take the middle value (average of the two middle values if n is even). Here the 4th of 7: **6**. Robust: change 20 to 200 and the median is still 6. It minimises the sum of absolute deviations (that is why MAE regression predicts medians).

31.3.3 Mode

The most frequent value: **5**. The only centre for nominal data ("most common city"). Can have several modes (**bimodal**: freshers and seniors).

31.3.4 Others

Weighted mean $\sum w_i x_i / \sum w_i$ (average CTR across jobs weighted by impressions). **Trimmed mean** (drop top and bottom 5%). **Geometric mean** $(\prod x_i)^{1/n}$ for growth rates (10% then 20% growth: $\sqrt{1.1 \times 1.2} = 1.149$, i.e. 14.9% per year, not 15%). **Harmonic mean** $n / \sum(1/x_i)$ for rates (F1 score is the harmonic mean of precision and recall).

☰ Remember: which centre to report

Symmetric, no outliers: mean. Skewed or outliers (salary, price, time): median. Categories: mode. Always report mean *and* median for money: their gap tells you about skew.

31.4 Measures of spread: how scattered?

31.4.1 Range

$\max - \min = 20 - 4 = 16$. Simple, but determined by the two most extreme points only.

31.4.2 Variance and standard deviation

Variance is the average squared distance from the mean. Deviations from 8: -4, -3, -3, -2, -1, 1, 12; squares 16, 9, 9, 4, 1, 1, 144; sum 184.

- **Population variance** (data are the whole population): $\sigma^2 = \frac{184}{7} = 26.29$, $\sigma = 5.13$.
- **Sample variance** (data are a sample used to estimate the population): $s^2 = \frac{184}{n-1} = \frac{184}{6} = 30.67$, $s = 5.54$.

Why $n - 1$ (Bessel's correction)? We measured deviations from the *sample* mean, which is by construction the point closest to the sample; deviations from it are systematically a bit smaller than deviations from the true population mean. Dividing by $n - 1$ instead of n corrects that under-estimate exactly (on average). Another view: once $\bar{x}$ is fixed, only $n - 1$ deviations are free (they must sum to 0): $n - 1$ **degrees of freedom**. pandas .std() uses $n - 1$; NumPy np.std uses n unless ddof=1.

The standard deviation is in the same units as the data (lakh), which is why we report it rather than the variance (lakh²).

31.4.3 Quartiles, IQR and MAD

Quartiles split sorted data into four parts. With NumPy's default (linear interpolation) on our data: $Q_1 = 5$, $Q_2 = 6$, $Q_3 = 8$; **IQR** = $Q_3 - Q_1 = 3$. The IQR is the spread of the middle 50%: robust to the 20. (Different software uses slightly different quartile conventions; in an exam, state the one you use.) **Median absolute deviation** (MAD): median of $|x_i - \text{median}| = \text{median of } 2, 1, 1, 0, 1, 3, 14 = 1$. Very robust.

31.4.4 Coefficient of variation

$CV = s/\bar{x}$: spread relative to the level, unit-free. Compare variability of fresher salaries ($\bar{x} = 4$, $s = 1$, CV 25%) with senior salaries ($\bar{x} = 40$, $s = 8$, CV 20%): seniors vary more in lakh but less relative to their level.

31.5 Shape: skewness and kurtosis


31.5.1 Moments

The k -th **central moment** is $\mu_k = \frac{1}{n} \sum (x_i - \bar{x})^k$. The 1st is 0, the 2nd is the variance, the 3rd measures asymmetry, the 4th measures tail heaviness. Standardising them gives skewness and kurtosis.

31.5.2 Skewness

$$\text{skewness} = \frac{\mu_3}{\sigma^3} = \frac{\frac{1}{n} \sum (x_i - \bar{x})^3}{\sigma^3}.$$

Cubing keeps the sign: big positive deviations dominate if the right tail is long.



- **Positive (right) skew**: long tail to the right; typically **mean** > **median** > **mode**. Salaries, house prices, time on site, applications per job, income. Our data: mean 8 > median 6 > mode 5; skewness = 1.70 (population formula).
- **Negative (left) skew**: long tail to the left; mean < median < mode. Marks in an easy exam (most score high, a few very low), age at retirement.
- **Zero skew**: symmetric (normal distribution, heights).

Rule of thumb: $|\text{skew}| < 0.5$ roughly symmetric, 0.5–1 moderate, > 1 highly skewed. **Pearson's second skewness coefficient** $3(\bar{x} - \text{median})/s = 3(8 - 6)/5.54 = 1.08$: a quick check without cubes.

31.5.3 Kurtosis

$$\text{kurtosis} = \frac{\mu_4}{\sigma^4}, \quad \text{excess kurtosis} = \frac{\mu_4}{\sigma^4} - 3.$$

The normal distribution has kurtosis 3 (excess 0). Kurtosis is mainly about **tails** (how often extreme values occur), not “peakedness”.

- **Mesokurtic**: excess ≈ 0 (normal).
- **Leptokurtic**: excess > 0 , heavy tails, more outliers than normal (stock returns, claim sizes, daily traffic spikes; Student- t , Laplace).
- **Platykurtic**: excess < 0 , light tails, fewer extremes (uniform distribution: excess -1.2).

Our salaries: excess kurtosis 1.38 (population formula): heavier tails than normal, driven by the 20.

⚠ Common mistake

Libraries differ: `scipy.stats.kurtosis` returns *excess* kurtosis by default; `pandas .kurt()` also returns excess kurtosis but with a small-sample correction (it gives 5.11 for our data, versus 1.38 for the plain population formula). Always say which definition you use.

31.6 Reading graphs: what each plot answers

31.6.1 Histogram


Bars over bins; height = count (or density). **Read**: centre, spread, number of peaks (modes), skew direction (where the tail is), gaps, impossible values, spikes at round numbers (“salary = 10 lakh” entered by many people: rounding; spike at 0: missing coded as 0). **Caution**: bin width changes the story; try several.

31.6.2 Density plot (KDE)

A smoothed histogram (Chapter 15). Good for comparing distributions of two groups (placed vs not placed MBA scores) on one plot.

31.6.3 Box plot

Five-number summary in one picture.



- The box spans Q_1 to Q_3 (the middle 50%); the line inside is the median.
- Whiskers extend to the most extreme data points within $Q_1 - 1.5 \text{ IQR}$ and $Q_3 + 1.5 \text{ IQR}$ (the **fences**); here the fences are 0.5 and 12.5, so the upper whisker stops at 9 (the largest value ≤ 12.5).

- Points beyond the fences are drawn individually as **outliers**: 20.
- **Reading skew**: median closer to Q_1 and a longer upper whisker/outliers $\Rightarrow$ right skew (our case).
- Side-by-side box plots compare groups (MBA % by placement status): non-overlapping boxes mean a strong difference.

Why 1.5? For normal data, the fences sit at about $\pm 2.7\sigma$, flagging roughly 0.7% of points: a sensible “unusual” level.

31.6.4 Violin plot

Box plot plus a mirrored KDE: shows bimodality that a box plot hides.

31.6.5 Bar chart and count plot

For categorical variables: counts or means per category. Sort bars; use horizontal bars for long labels. **Pie charts**: only for 2–3 parts of a whole; humans compare angles poorly.

31.6.6 Q–Q plot

Plots the quantiles of your data against those of a theoretical distribution (usually normal). Points on the straight line $\Rightarrow$ the distribution matches. **Read**: an upward curve at the right end (points above the line) $\Rightarrow$ heavy right tail; S-shape $\Rightarrow$ heavy (or light) tails on both sides. Used to check the normality of regression residuals.

31.6.7 ECDF

Empirical cumulative distribution: for each value, the fraction of data $\leq$ it. Reads percentiles directly (“80% of jobs get fewer than 30 applies”); no binning choices.

31.6.8 Line plot

Values over time: trend, seasonality, sudden level shifts (a tracking change), missing periods.

In practice: where univariate EDA changes ML decisions

- Right-skewed salary target $\rightarrow$ log-transform or use a Gamma/Tweedie objective; report MAPE not RMSE (Chapters 4, 34).
- Heavy tails (leptokurtic) in a feature $\rightarrow$ robust scaling, winsorising, tree models.
- Bimodal feature $\rightarrow$ maybe two populations: add a segment feature or model separately.
- Spike at 0 or at a round number $\rightarrow$ disguised missing values; recode.
- Near-constant column (99.9% one value) $\rightarrow$ drop (no information) unless the rare value is the signal.
- High-cardinality categorical (50,000 companies) $\rightarrow$ grouping, frequency or target encoding (Chapter 34).

31.7 Interview questions

Q31.1

Compute the mean, median, mode, range, population and sample variance and standard deviation of salaries 4, 5, 5, 6, 7, 9, 20 (lakh). Which centre would you report and why?

🎯 **In one line.** Mean 8, median 6, mode 5, range 16, population variance 26.29 (sd 5.13), sample variance 30.67 (sd 5.54); report the median because the distribution is right-skewed by one large salary, so the mean overstates what a typical analyst earns.

⚙️ **The full answer.** Deviations from 8 sum to 0; squared deviations sum to 184. Divide by 7 for population, 6 for sample. The mean exceeds the median, signalling right skew.

💡 **If the interviewer pushes.** “What if 20 were a data-entry error (should be 2.0)?” Mean becomes $(56 - 20 + 2)/7 = 5.43$, median 5: the error moved the mean by 2.6 lakh, the median by 1.

Q31.2

Why do we divide by $n - 1$ for the sample variance?

🎯 **In one line.** Because deviations are measured from the sample mean, which is fitted to the same data and so sits closer to the points than the true mean; dividing by n would underestimate the population variance on average, and $n - 1$ (the degrees of freedom left after estimating the mean) makes it unbiased.

⚙️ **The full answer.** $\mathbb{E}[\sum (x_i - \bar{x})^2] = (n - 1)\sigma^2$. Intuition with $n = 1$: one point has zero deviation from its own mean, which says nothing about spread; dividing by $n - 1 = 0$ correctly refuses to estimate.

Q31.3

When is the median better than the mean, and when is the mean better?

🎯 **In one line.** The median is better for skewed data or data with outliers (salaries, prices, durations) because it is robust; the mean is better for symmetric data and when totals matter (expected revenue = mean $\times$ count), and it is statistically more efficient when the data are roughly normal.

Q31.4

What are the types of skewness? Give the order of mean, median and mode for each, with real examples.

🎯 **In one line.** Positive (right) skew: long right tail, mean $>$ median $>$ mode (salaries, property prices, applications per job); negative (left) skew: long left tail, mean $<$ median $<$ mode (scores on an easy exam); zero skew: symmetric, all three equal (heights).

🗣️ **If the interviewer pushes.** “Is the order always that?” It is a strong tendency, not a theorem; multimodal or discrete distributions can break it.

Q31.5

What is kurtosis and what are its types? Is high kurtosis “a peaked distribution”?

🎯 **In one line.** Kurtosis is the standardised fourth moment and measures tail heaviness: mesokurtic (normal, excess 0), leptokurtic (heavy tails, excess > 0), platykurtic (light tails, excess < 0); it is primarily about tails and outliers, not peakedness.

⚙️ **The full answer.** A distribution can have a flat top and heavy tails, or a sharp peak and light tails. Practical reading: high kurtosis means extreme values occur more often than a normal model predicts, so normal-based intervals and z-score rules underestimate risk.

Q31.6

Interpret a box plot of salary where the median line is near the bottom of the box, the upper whisker is long and several points lie above it.

🎯 **In one line.** The distribution is right-skewed: the lower half of the middle 50% is compressed and the upper half stretched, and several unusually high salaries lie beyond $Q_3 + 1.5 \text{ IQR}$.

🗣️ **If the interviewer pushes.** “Would you delete those points?” Not without checking: high salaries for senior roles are genuine. Consider a log transform or robust methods.

Q31.7

How do you read a Q–Q plot, and where is it used in ML?

🎯 **In one line.** Each point compares a data quantile with the quantile expected under a reference distribution; a straight line means a good match, curvature at the ends means heavier or lighter tails, and a consistent bend means skew; it is used to check normality of regression residuals and to decide on transformations.

Q31.8

🧠 **Think deeper** Two datasets have the same mean, median, variance, skewness and kurtosis. Can they still look very different?

🎯 **In one line.** Yes: summary statistics compress a distribution into a few numbers, and very different shapes (bimodal vs unimodal, clustered vs spread) can share them; always plot.

⚙️ **The full answer.** Anscombe’s quartet and the “Datasaurus dozen” show datasets with identical means, variances and correlations but completely different scatter plots.

Q31.9

Histogram of “years of experience” shows a spike at exactly 0 and another at 30. What could they mean?

🕒 **In one line.** The spike at 0 is probably genuine freshers mixed with missing values coded as 0; the spike at 30 may be a default or capped value from the form; both must be investigated before modelling.

⚙️ **The full answer.** Cross-check: freshers should have no previous employer; a cap at 30 means the true values are censored.

Q31.10

What is the coefficient of variation and when would you use it?

🕒 **In one line.** $CV = s/\bar{x}$, the standard deviation relative to the mean; use it to compare variability across groups with very different scales or units, such as salary variability for freshers vs senior leaders.

Q31.11

What is the geometric mean and when is it the right average?

🕒 **In one line.** The n -th root of the product of values; it is the right average for multiplicative quantities like growth rates: growth of 10% then 20% averages to 14.9% per period, not 15%.

Q31.12

💡 **Think deeper** Why do we square deviations in the variance instead of taking absolute values?

🕒 **In one line.** Squaring gives a smooth function with clean algebra (variances of independent variables add; the mean minimises squared error), and connects to the normal distribution; absolute deviations are more robust but mathematically less convenient, and lead to the median and MAD instead.

Q31.13

Which plot would you use to show: (a) the distribution of MBA percentage, (b) MBA percentage by placement status, (c) the number of students per degree type, (d) applications per day over a year?

🕒 **In one line.** (a) histogram or KDE; (b) side-by-side box or violin plots; (c) sorted bar (count) plot; (d) line plot, ideally with a 7-day moving average.

Q31.14

 **Think deeper** A feature is 99.5% zeros. Drop it or keep it?

🎯 **In one line.** Keep it if the rare non-zero values carry signal for the target (for example “has a patent” predicting senior roles), which you check by comparing target rates; drop it if it is noise. Near-zero-variance filters should not be applied blindly.

Q31.15

What are moments of a distribution?

🎯 **In one line.** Expected values of powers: the first raw moment is the mean, the second central moment the variance, the third standardised moment the skewness, and the fourth standardised moment the kurtosis; together they describe location, spread, asymmetry and tail weight.

CHAPTER 32

EDA II: Bivariate and Multivariate Analysis, Correlations and Multicollinearity

🎥 **Companion video for this chapter:** [StatQuest: Pearson's Correlation, Clearly Explained](#)

Univariate analysis describes one column. Bivariate analysis asks how *two* variables move together (“does MBA percentage relate to placement?”), and multivariate analysis asks how many variables relate at once (“are experience, age and graduation year telling the same story?”). The checklist names Pearson, Spearman (with ranks), point-biserial and phi correlations, reading a correlation matrix, multicollinearity and the VIF formula in terms of R . Each is derived here with a worked example.

32.1 Which measure for which pair of variable types?

Variable 1	Variable 2	Measure
numeric	numeric	Pearson r (linear), Spearman ρ or Kendall τ (monotonic, robust)
numeric	binary	point-biserial correlation (a special case of Pearson)
numeric	multi-class categorical	ANOVA F -test, correlation ratio η^2 , box plots
binary	binary	phi coefficient ϕ ; odds ratio
categorical	categorical	chi-square test, Cramér's V
ordinal	ordinal / numeric	Spearman, Kendall

32.2 Covariance

Covariance measures whether two variables move together:

$$\text{Cov}(x, y) = \frac{1}{n-1} \sum_i (x_i - \bar{x})(y_i - \bar{y}).$$

If above-average x tends to come with above-average y , the products are mostly positive: positive covariance. Its size depends on units (rupees vs lakh), so it is hard to interpret alone; correlation fixes this.

32.3 Pearson correlation

$$r = \frac{\text{Cov}(x, y)}{s_x s_y} = \frac{\sum (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum (x_i - \bar{x})^2} \sqrt{\sum (y_i - \bar{y})^2}}.$$

It is the covariance of the standardised variables: unit-free, between -1 and $+1$.

Small example: Pearson by hand

Years of experience $x = 1, 2, 3, 4, 5$; rating $y = 2, 4, 5, 4, 5$. $\bar{x} = 3$, $\bar{y} = 4$. Deviations x : $-2, -1, 0, 1, 2$; y : $-2, 0, 1, 0, 1$. Products: $4, 0, 0, 0, 2$; sum = 6. $\sum(x - \bar{x})^2 = 10$, $\sum(y - \bar{y})^2 = 4 + 0 + 1 + 0 + 1 = 6$. $r = 6/\sqrt{10 \cdot 6} = 6/7.746 = 0.775$. Covariance = $6/4 = 1.5$.

Properties and cautions

- $r = \pm 1$: points exactly on a line; $r = 0$: no *linear* relationship (there may be a strong curved one: $y = x^2$ on symmetric x has $r = 0$).
- Unchanged by shifting or positive scaling of either variable (rupees or lakh: same r).
- Sensitive to outliers: one extreme point can create or destroy a correlation.
- r^2 is the fraction of variance of y explained by a linear fit on x (Chapter 2).
- Correlation is not causation: ice-cream sales and drownings correlate through temperature (a confounder).

32.4 Spearman rank correlation

Replace each variable's values by their **ranks**, then compute Pearson on the ranks. It measures **monotonic** association (does y tend to rise when x rises, in any shape?) and is robust to outliers and to non-linear but monotone relationships. Without ties, a shortcut:

$$\rho = 1 - \frac{6 \sum d_i^2}{n(n^2 - 1)}, \quad d_i = \text{rank}(x_i) - \text{rank}(y_i).$$

Small example: Spearman by hand (no ties)

Interview scores $a = 10, 20, 30, 40, 50$ (ranks 1–5); manager ratings $b = 1, 3, 2, 5, 4$ (ranks 1, 3, 2, 5, 4). $d = 0, -1, 1, -1, 1$; $\sum d^2 = 4$. $\rho = 1 - \frac{6 \cdot 4}{5 \cdot 24} = 1 - 0.2 = 0.8$.

Small example: ties

For $y = 2, 4, 5, 4, 5$ the two 4s share ranks 2 and 3: both get 2.5; the two 5s get 4.5. Ranks: 1, 2.5, 4.5, 2.5, 4.5. With ties, compute Pearson on these ranks (the shortcut is only approximate): $\rho = 0.738$ against $x = 1, \dots, 5$, slightly lower than Pearson's 0.775 here.

Kendall's τ : based on concordant vs discordant pairs; more robust with small samples and many ties.

32.5 Point-biserial correlation (numeric vs binary)

Code the binary variable as 0/1 and compute Pearson; equivalently

$$r_{pb} = \frac{\bar{x}_1 - \bar{x}_0}{s_x} \sqrt{p(1-p)},$$

where $\bar{x}_1, \bar{x}_0$ are the means of the numeric variable in the two groups, s_x its (population) standard deviation and p the fraction in group 1.

Small example: point-biserial

MBA % = 60, 65, 70, 75, 80, 85; placed = 0, 0, 1, 0, 1, 1. $\bar{x}_1 = 78.33$, $\bar{x}_0 = 66.67$, $s_x = 8.54$, $p = 0.5$.
 $r_{pb} = \frac{11.67}{8.54} \times 0.5 = 0.683$. Placed students have clearly higher MBA scores.

It is the natural measure of how a numeric feature relates to a binary target; a t -test between the two groups tests the same thing.

32.6 Phi coefficient (binary vs binary)

For a 2×2 table $\begin{pmatrix} a & b \\ c & d \end{pmatrix}$:

$$\phi = \frac{ad - bc}{\sqrt{(a+b)(c+d)(a+c)(b+d)}}.$$

Small example: phi

Rows: work experience yes/no; columns: placed yes/no. $a = 30$ (exp, placed), $b = 10$ (exp, not), $c = 20$ (no exp, placed), $d = 40$ (no exp, not). $\phi = \frac{1200 - 200}{\sqrt{40 \cdot 60 \cdot 50 \cdot 50}} = \frac{1000}{2449.5} = 0.408$. Also $\phi = \sqrt{\chi^2/n} = \sqrt{16.67/100} = 0.408$.

For larger tables, **Cramér's V** = $\sqrt{\chi^2/(n(k-1))}$ with $k = \min(\text{rows}, \text{cols})$, between 0 and 1.

32.7 Numeric vs multi-class categorical: ANOVA and η^2

Compare the means of a numeric variable across several groups (MBA % by degree type). The **correlation ratio** $\eta^2 = \frac{\text{between-group sum of squares}}{\text{total sum of squares}}$ is the fraction of variance explained by the group; the ANOVA F -test asks whether the differences are larger than chance (Chapter 39).

32.8 Reading a correlation matrix

A correlation matrix lists the correlation of every pair of numeric variables; usually shown as a heatmap.

	ssc	hsc	degree	mba	placed
ssc	1.00	0.51	0.54	0.39	0.61
hsc	0.51	1.00	0.43	0.35	0.49
degree	0.54	0.43	1.00	0.40	0.48
mba	0.39	0.35	0.40	1.00	0.08
placed	0.61	0.49	0.48	0.08	1.00

(Illustrative values.) How to read it, in order:

1. The diagonal is 1; the matrix is symmetric, so read one triangle.
2. **Target row/column** first: ssc has the strongest linear association with placement (0.61); mba almost none (0.08), surprisingly.
3. **Feature-feature** cells: high values ($|r| > 0.8$) warn of multicollinearity or redundancy; here moderate (0.4–0.54): the academic scores share a common “academic ability” factor.

4. **Blocks** of mutually correlated features suggest groups (PCA would find one component for them).
5. **Caveats:** only linear relationships (use Spearman for monotone), sensitive to outliers, hides non-linear and interaction effects, and with many variables some correlations are large by chance. Low correlation with the target does not make a feature useless (it may matter in interaction or non-linearly).

32.9 Simpson's paradox: when aggregate correlation lies

A trend can reverse when data are split by a third variable. Overall, more certifications might correlate with lower salary; within each experience band the correlation is positive, because juniors collect more certificates. Always check important relationships within key segments.

32.10 Multicollinearity

32.10.1 What and why it matters

Features are **multicollinear** when one can be (nearly) predicted from the others: age, years of experience and graduation year. Effects:

- Linear/logistic regression coefficients become unstable, with huge standard errors and sometimes the wrong sign; interpretation breaks.
- Predictions are usually still fine, as long as the collinearity pattern holds in new data.
- Tree models are unaffected in accuracy, but importance is split among the correlated features.

Pairwise correlations can miss it: $x_3 = x_1 + x_2$ may have only moderate pairwise correlations yet be perfectly collinear.

32.10.2 VIF, derived

Regress feature j on all the other features and record its R_j^2 (how well the others explain it). Then

$$\text{VIF}_j = \frac{1}{1 - R_j^2}, \quad \text{tolerance}_j = 1 - R_j^2.$$

The variance of the coefficient $\hat{w}_j$ is multiplied by VIF_j compared with uncorrelated features; its standard error by $\sqrt{\text{VIF}_j}$. **With only two predictors**, R_j^2 is just the squared correlation between them: $\text{VIF} = \frac{1}{1-r^2}$.

Small example: VIF values

Two predictors with $r = 0.9$: $\text{VIF} = 1/(1 - 0.81) = 5.26$; standard errors 2.3× larger. $R_j^2 = 0.9$ gives $\text{VIF} 10$; $R_j^2 = 0.99$ gives 100.

Rules of thumb: $\text{VIF} > 5$ worth a look, > 10 serious. **Condition number** of X (ratio of largest to smallest singular value, Chapter 14) > 30 also signals it.

32.10.3 Fixes

Drop one of a correlated group (keep the most interpretable); combine them (an index, or “experience – age”); PCA; Ridge regression; collect data where they vary independently; for prediction-only tasks, often leave it.

32.10.4 The dummy-variable trap is perfect multicollinearity

One-hot encoding a k -level category gives k columns that always sum to 1, the same as the intercept column: perfect collinearity, $X^T X$ singular (Chapter 34). Drop one dummy (`drop_first=True`) when the model has an intercept and no regularisation.

In practice: where these measures are used in ML

- **Feature screening:** rank features by $|r|$ or point-biserial correlation with the target as a first pass (then confirm with model-based importance).
- **Redundancy removal:** drop one of each pair with $|r| > 0.95$ before linear models; check VIF.
- **Categorical features:** Cramér’s V between candidate features (city and state are nearly the same information).
- **Monitoring:** a change in the correlation between a feature and the target after deployment signals concept drift.
- **Leakage detection:** a single feature with $|r| \approx 1$ with the target is suspicious (the salary-placement case).

32.11 Interview questions

Q32.1

Compute the Pearson correlation of $x = 1, 2, 3, 4, 5$ and $y = 2, 4, 5, 4, 5$, and interpret it.

🕒 **In one line.** $r = 6/\sqrt{10 \times 6} = 0.775$: a fairly strong positive linear association; about $0.775^2 = 60\%$ of the variance in y is explained by a straight line in x .

Q32.2

Pearson vs Spearman: what does each measure, and when do you prefer Spearman?

🕒 **In one line.** Pearson measures linear association of the raw values; Spearman is Pearson on ranks and measures monotonic association; prefer Spearman for ordinal data, non-linear but monotone relationships, and data with outliers or heavy skew.

📊 **Worked example.** $y = e^x$ on $x = 1..5$: Spearman exactly 1 (perfectly monotone), Pearson about 0.89 (not linear).

Q32.3

Write the Spearman formula and compute it for ranks (1, 2, 3, 4, 5) and (1, 3, 2, 5, 4).

🎯 **In one line.** $\rho = 1 - 6 \sum d^2 / (n(n^2 - 1)) = 1 - 6 \cdot 4 / 120 = 0.8$.

💡 **If the interviewer pushes.** “With ties?” Assign average ranks and compute Pearson on ranks.

Q32.4

How do you measure the association between a numeric feature and a binary target? Compute it for the MBA example.

🎯 **In one line.** With the point-biserial correlation, $r_{pb} = \frac{\bar{x}_1 - \bar{x}_0}{s_x} \sqrt{p(1-p)}$; for the example it is 0.683.

⚙️ **The full answer.** It equals Pearson with the binary variable coded 0/1. Significance via a two-sample t -test; a model-free alternative is the AUC of the feature alone.

Q32.5

What is the phi coefficient? Compute it for the 2×2 table $[[30, 10], [20, 40]]$.

🎯 **In one line.** The correlation between two binary variables, $\phi = (ad - bc) / \sqrt{(a+b)(c+d)(a+c)(b+d)} = 1000 / 2449.5 = 0.408$; it equals $\sqrt{\chi^2 / n}$.

Q32.6

How would you measure association between two categorical variables with many levels, like city and job function?

🎯 **In one line.** Use the chi-square test of independence for significance and Cramér’s $V = \sqrt{\chi^2 / (n(k-1))}$ for strength.

Q32.7

🧠 **Think deeper** $r = 0$ between two variables. Are they independent?

🎯 **In one line.** Not necessarily: $r = 0$ only rules out a linear relationship; a strong non-linear dependence (like $y = x^2$ for x symmetric around 0) can have zero correlation. For jointly normal variables, zero correlation does imply independence.

⚙️ **The full answer.** Use mutual information, Spearman (for monotone), or plots to detect non-linear dependence.

Q32.8

How do you read a correlation matrix in an interview?

🕒 **In one line.** Ignore the diagonal, read one triangle, look first at the target column for candidate predictors, then at feature pairs with high $|r|$ for redundancy and multicollinearity, look for blocks of related features, and remember it shows only linear, pairwise, possibly confounded relationships.

Q32.9

Define VIF, write it in terms of R , and compute it for two predictors with correlation 0.9.

🕒 **In one line.** $VIF_j = 1/(1 - R_j^2)$ where R_j^2 is from regressing feature j on the others; with two predictors $R_j^2 = r^2$, so $VIF = 1/(1 - 0.81) = 5.26$.

🔧 **The full answer.** It is the factor by which collinearity inflates the variance of the coefficient estimate.

Q32.10

💡 **Think deeper** Why can pairwise correlations miss multicollinearity?

🕒 **In one line.** Because collinearity can involve three or more variables jointly; for example $x_3 = x_1 + x_2$ with independent x_1, x_2 gives pairwise correlations of only about 0.71, yet x_3 is perfectly predictable from the other two (VIF infinite).

Q32.11

Does multicollinearity hurt prediction? Does it hurt random forests?

🕒 **In one line.** It mainly hurts coefficient interpretation and stability in linear models, not prediction accuracy (if the pattern persists); random forests predict fine but split importance among correlated features, making each look less important.

Q32.12

What is Simpson's paradox? Give a job-portal example.

🕒 **In one line.** A relationship that holds in every subgroup reverses (or vanishes) when the groups are combined, because of a confounder; for example, recommendations might have a lower apply rate overall but a higher rate within each experience band because they are shown more to juniors, who apply less.

Q32.13

 **Think deeper** A feature has correlation 0.02 with the target. Should you drop it?

🎯 **In one line.** Not on that alone: correlation captures only linear marginal effects, and the feature may matter non-linearly or through interactions (for example only for one city); check model-based importance, mutual information and segment-wise plots first.

Q32.14

Why is correlation not causation? How could you establish causation?

🎯 **In one line.** Because a correlation can come from a confounder, reverse causation, selection bias or chance; causation is established by randomised experiments (A/B tests) or, failing that, by careful causal designs (controlling for confounders, instrumental variables, difference-in-differences, regression discontinuity).

CHAPTER 33

EDA III: Outliers, Missing Values and Leaky Variables, in Depth


🎥 **Companion video for this chapter:** [StatQuest: The Normal Distribution, Clearly Explained \(search\)](#)

Chapter 13 introduced missing data, outliers and imbalance from a modelling angle. The checklist asks for more detail: the 68–95–99.7 rule behind the z-score, the IQR formula and box-plot interpretation, the pros and cons of mean, median and mode imputation, ML models that predict missing values (“XGB imputer”), and leaky variables. This chapter completes those, with numbers.

33.1 The normal curve and the 68–95–99.7 rule

Many measurements (heights, test scores, measurement errors, averages of many things) follow a bell-shaped **normal distribution** with mean μ and standard deviation σ (Chapter 38). The **empirical rule**:

- about **68%** of values lie within $\mu \pm 1\sigma$,
- about **95%** within $\mu \pm 2\sigma$ (exactly 95% within $\pm 1.96\sigma$),
- about **99.7%** within $\mu \pm 3\sigma$.



📊 Small example: using the rule

MBA percentage ~ normal with mean 60, sd 5. About 68% of students score 55–65; 95% score 50–70; 99.7% score 45–75. Of 1,000 students we expect about 683, 954 and 997 in those ranges; only about 3 outside 45–75, so a score of 80 ($z = 4$) is genuinely rare. Also: by symmetry, 16% score above 65 ($\frac{100-68}{2}$) and 2.5% above 70.

33.2 z-scores

$$z = \frac{x - \mu}{\sigma}$$

says “how many standard deviations from the mean”. It puts different variables on one scale and, for normal data, converts directly into rarity via the empirical rule.

Small example: z-score outlier check

Daily applications to a job: mean 50, sd 8. A day with 76: $z = 26/8 = 3.25$: beyond 3, flagged (for a normal variable, $P(|Z| > 3) = 0.27\%$).

Limitations: (1) assumes roughly normal data; for skewed data (salaries) many genuine values exceed 3; (2) **masking:** the outlier inflates the mean and sd it is judged by, so in small samples an outlier may not even reach $|z| = 3$ (the 40 in Chapter 13 had $z = 2.4$). **Modified z-score** (robust): $M_i = 0.6745 \frac{x_i - \text{median}}{\text{MAD}}$, flag $|M_i| > 3.5$. The constant 0.6745 makes MAD comparable to σ for normal data.

Small example: modified z

Salaries 4, 5, 5, 6, 7, 9, 20: median 6, MAD 1. For 20: $M = 0.6745 \times 14/1 = 9.4$: clearly flagged, while the ordinary $z = (20 - 8)/5.54 = 2.2$ is not.

33.3 IQR rule and box plots

$$\text{IQR} = Q_3 - Q_1, \quad \text{lower fence} = Q_1 - 1.5 \text{ IQR}, \quad \text{upper fence} = Q_3 + 1.5 \text{ IQR}.$$

Points outside the fences are outliers (“far out” beyond $3 \times \text{IQR}$). It uses quartiles, so it is robust and does not assume normality; it is what box-plot whiskers show (Chapter 31). For right-skewed data it flags many points on the high side only, which is correct behaviour (the high tail is long), but those points are usually genuine.

Small example: IQR

Data 3, 5, 7, 8, 9, 11, 13, 40 with $Q_1 = 6$, $Q_3 = 12$: IQR 6, fences -3 and 21 : 40 is an outlier.

33.4 Multivariate and model-based detection

A point can look normal on every single feature but be unusual in combination (age 22 with 15 years of experience). Methods: **Mahalanobis distance** $\sqrt{(\mathbf{x} - \boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1}(\mathbf{x} - \boldsymbol{\mu})}$ (a z-score that accounts for correlations; squared, it follows a χ_d^2 for normal data); **Isolation Forest** (Chapter 7); **Local Outlier Factor** (density relative to neighbours); **DBSCAN** noise points; autoencoder reconstruction error.

33.5 What to do with an outlier: a decision procedure

1. **Is it impossible?** (negative age, percentage 105, price Rs. 1) → data error: correct from source or set to missing.
2. **Is it genuine but extreme?** (a 2-crore salary for a CTO) → keep; reduce its influence: log transform, winsorise, robust loss (Huber, MAE), robust scaler, tree models.
3. **Is it from a different population?** (bulk staffing agencies among recruiters) → segment or add an indicator.
4. **Is it the signal?** (fraud) → anomaly detection; never remove.

Always decide using training data only and apply the same rule to test data; never drop test rows.

33.6 Missing values: the full picture

33.6.1 Diagnosing the mechanism

MCAR, MAR, MNAR (Chapter 13). Practical checks: compare the distributions of other variables for rows with and without the missing value (if they differ, it is not MCAR); compare target rates for missing vs present; ask how the data were collected. Little's MCAR test exists but is rarely decisive.

33.6.2 Dropping

- **Listwise deletion** (drop rows): unbiased only under MCAR; loses data fast with many columns ($0.95^{12} = 54\%$ complete rows).
- **Pairwise deletion** (use available pairs for each correlation): keeps more data but statistics can be inconsistent with each other.
- **Dropping a column**: when mostly missing *and* uninformative.
- Never drop rows whose missingness is tied to the target (salary missing for every non-placed student).

33.6.3 Mean, median and mode imputation: advantages and disadvantages

Method	Advantages	Disadvantages
Mean	simple; keeps the mean unchanged; fine for symmetric data, small missing fractions, MCAR	pulled by outliers; shrinks variance; weakens correlations; distorts distribution (spike at the mean); ignores relationships with other features; biased under MAR/MNAR
Median	robust to outliers and skew (salary, price); simple	same variance shrinkage and correlation weakening; spike at the median
Mode	works for categoricals and discrete values	over-represents the most frequent category; can be arbitrary when several modes; hides the signal in missingness (better: a separate "Unknown" category)
Constant (−999, 0)	preserves the information "was missing" for tree models	meaningless for linear models and distances unless paired with an indicator

Small example: variance shrinkage, simulated

500 salaries, 20% removed at random, true variance 29.2. After mean imputation the column's variance is 21.5: the filled values sit at the centre and pretend we are more certain than we are.

33.6.4 ML models that predict missing values

The idea: treat the column with gaps as a *target*, learn it from the other columns on rows where it is observed, then predict it where it is missing.

- **KNN imputer**: average of the k most similar rows (scale first).
- **Regression imputation**: linear model from the other columns (add noise for stochastic imputation).

- **MICE / IterativeImputer**: cycle through every incomplete column, modelling each from all the others, for several rounds.
- **MissForest**: IterativeImputer with random forests: handles non-linearity, interactions and mixed types.
- **“XGB imputer”**: the same idea with gradient boosting (XGBoost, LightGBM, scikit-learn’s HistGradientBoosting), whose ability to accept NaN in the *predictor* columns means rows with several gaps can still be used.

```

from sklearn.experimental import enable_iterative_imputer  # noqa: enables the class
from sklearn.impute import IterativeImputer, KNNImputer
from sklearn.ensemble import RandomForestRegressor, HistGradientBoostingRegressor

knn = KNNImputer(n_neighbors=5)
mice = IterativeImputer(random_state=0)  # Bayesian ridge per
column
miss_forest = IterativeImputer(estimator=RandomForestRegressor(n_estimators=50), max_iter
=5)

def gbm_impute(df, col):  # "XGB imputer" for one column
    obs = df[col].notna()
    feats = df.columns.drop(col)
    m = HistGradientBoostingRegressor().fit(df.loc[obs, feats], df.loc[obs, col])
    out = df[col].copy()
    out[~obs] = m.predict(df.loc[~obs, feats])
    return out

```

Small example: how much better? (simulated)

Salary depends on experience; 20% of salaries removed at random. RMSE of the imputed values: mean 5.81, median 5.79, KNN 1.69, MICE 1.64, MissForest 1.72, gradient-boosting imputer 1.68. Model-based imputation is more than 3 times more accurate because it uses the relationship with experience.

Cautions: fit the imputer on training data only (inside a Pipeline); model-based imputations are “too clean” (no noise), which can understate uncertainty; for inference prefer multiple imputation; add a missing indicator when missingness may be informative; for GBM models, native NaN handling is often as good and simpler.

33.7 Leaky variables

A **leaky variable** contains information about the target that would not be available at prediction time. Catalogue with detection tips:

Type	Example	How to catch it
Outcome recorded after the event	salary for placement; “offer letter generated” for hiring; “refund issued” for fraud	ask “when is this value created?”; missingness perfectly tied to the target
Proxy of the target	“number of interviews” when predicting “hired”; account status “closed” for churn	suspiciously high importance or AUC of one feature
Future information	features aggregated over the full period including after the prediction date	time-based validation drops sharply
Preprocessing on all data	scaler, imputer, TF-IDF, target encoding fitted before the split	always use Pipelines; fit transforms in folds
Duplicates / same entity ID or timestamp ordering	same candidate in train and test row ID correlated with label due to sorted export	group split by entity drop IDs; check importance

In practice: where this chapter shows up in ML

z-scores standardise features (StandardScaler) and power anomaly alerts on metrics dashboards (“applications today are 4 sd below normal: check the pipeline”); IQR rules drive winsorising; imputation choices change model accuracy directly; and most “too good to be true” offline results are leakage.

33.8 Interview questions

Q33.1

State the 68–95–99.7 rule and use it: MBA scores are normal with mean 60 and sd 5. What fraction score above 70? Between 55 and 65?

 **In one line.** About 68%, 95% and 99.7% of a normal variable lie within 1, 2 and 3 standard deviations of the mean; above 70 is $z > 2$, about 2.5%; between 55 and 65 is within $\pm 1\sigma$, about 68%.

 **The full answer.** Above $\mu + 2\sigma$: $(100 - 95)/2 = 2.5\%$ by symmetry.

Q33.2

How do you detect outliers with the z-score, and what are its weaknesses?

 **In one line.** Standardise $z = (x - \mu)/\sigma$ and flag $|z| > 3$ (or 2.5); it assumes approximate normality and suffers from masking because the outlier inflates μ and σ ; robust alternatives are the modified z-score with median and MAD, or the IQR rule.

 **Worked example.** Salaries with 20: ordinary $z = 2.2$ (missed), modified $z = 9.4$ (caught).

Q33.3

Write the IQR rule and explain why the multiplier is 1.5.

🎯 **In one line.** Outliers lie below $Q_1 - 1.5 \text{ IQR}$ or above $Q_3 + 1.5 \text{ IQR}$; with 1.5 the fences fall at about $\pm 2.7\sigma$ for normal data, flagging roughly 0.7% of points, a sensible threshold for “unusual” without assuming normality for the rule itself.

Q33.4

Would you remove outliers before training? Explain your decision process.

🎯 **In one line.** Only data errors; genuine extremes are information, so I reduce their influence (transforms, winsorising, robust losses, tree models) and treat target-relevant anomalies as signal; decisions are made on training data and applied consistently.

Q33.5

Give the advantages and disadvantages of mean, median and mode imputation.

🎯 **In one line.** All are fast and simple; mean preserves the average but is outlier-sensitive; median is robust for skewed data; mode works for categories; all shrink variance, weaken correlations, create artificial spikes, ignore other features and are biased unless data are MCAR.

Q33.6

🧠 **Think deeper** How can an ML model be used to impute missing values? Describe an “XGB imputer”.

🎯 **In one line.** Treat the incomplete column as a target, train a model (KNN, linear, random forest, gradient boosting) on rows where it is observed using the other columns as features, and predict it for rows where it is missing; with gradient boosting, the predictor columns may themselves contain NaNs, which the trees handle natively; iterating over all incomplete columns gives MICE/MissForest-style imputation.

📊 **Worked example.** Simulated: RMSE 5.8 for mean imputation vs about 1.7 for model-based.

💬 **If the interviewer pushes.** “Risks?” Leakage if fitted on test data; using the target as an imputation feature for test rows (not available at prediction time); understated uncertainty; extra complexity.

Q33.7

🧠 **Think deeper** Should you use the target variable when imputing features?

🎯 **In one line.** For pure prediction, no: at prediction time the target is unknown, so an imputer that relied on it cannot be applied consistently and leaks; for statistical inference with multiple imputation, including the outcome in the imputation model is actually recommended, because

the analysis model will condition on it.

Q33.8

What is a leaky variable? Give three examples and how you would detect each.

🕒 **In one line.** A feature carrying information about the target that would not exist at prediction time; e.g. salary for placement (missingness perfectly tied to the label), “number of interviews” for hiring (a post-outcome proxy with extreme importance), and full-period aggregates for forecasting (validation collapses with a time split).

Q33.9

What is the Mahalanobis distance and why is it better than per-feature z-scores for multivariate outliers?

🕒 **In one line.** It measures distance from the mean in units that account for the covariance between features, so it flags combinations that are unusual given the correlations (age 22 with 15 years of experience) even when each value alone is normal.

Q33.10

Listwise vs pairwise deletion?

🕒 **In one line.** Listwise drops any row with a missing value (consistent but wasteful); pairwise uses all available data for each pair of variables (keeps more data but different statistics use different subsets and a correlation matrix may not be valid).

CHAPTER 34

EDA IV: Scaling, Transforming Skew, Binning, Encoding and Imbalance

🎥 **Companion video for this chapter:** [scikit-learn User Guide: Preprocessing data](#)

This chapter turns the raw columns that EDA has understood into model-ready features. Each technique comes with the reason it exists, a small numeric example (computed with scikit-learn/SciPy), and which models need it.

34.1 Why scale at all?

Many algorithms compare features through distances, dot products or gradient steps. If one feature is salary in rupees (range 10^6) and another is years of experience (range 30), the salary dominates every distance and creates a stretched loss surface. Scaling puts features on comparable ranges.

Needs scaling	Does not need scaling
<i>k</i> -NN, K-Means, SVM (especially RBF), PCA, linear/logistic regression with regularisation, neural networks, anything trained by gradient descent	decision trees, random forests, gradient-boosted trees (split on order), Naive Bayes (per-feature models)

34.2 The scalers, compared on one column

Column: 10, 20, 30, 40, 1000 (one extreme value).

Scaler	Formula	Result	Behaviour
Standardisation (z-score)	$\frac{x-\mu}{\sigma}$	-0.54, -0.51, -0.49, -0.46, 2.00	mean 0, sd 1; outlier squashes the rest
Min-max normalisation	$\frac{x-\min}{\max-\min}$	0, 0.01, 0.02, 0.03, 1	range [0, 1]; very outlier-sensitive
Robust scaling	$\frac{x-\text{median}}{\text{IQR}}$	-1, -0.5, 0, 0.5, 48.5	normal values keep their spread; outlier stays extreme
Max-abs	$\frac{x}{\max x }$	0.01, 0.02, 0.03, 0.04, 1	keeps zeros (sparse data friendly)

34.2.1 Standardisation vs normalisation: the difference

- **Standardisation** centres to mean 0 and scales to standard deviation 1. No fixed range. Natural for algorithms that assume centred data or use variances (PCA, linear models with penalties, SVM, neural networks). Less distorted by outliers than min-max, but still affected.
- **Normalisation** usually means **min-max scaling** to [0, 1]. Fixed range; useful when a bounded input is needed (image pixels /255, some neural-network inputs, distance-based methods on features with known bounds). One outlier compresses all other values near 0.

- “Normalisation” is also used for **row-wise L2 normalisation**: scale each *row* to unit length, $(3, 4) \rightarrow (0.6, 0.8)$, used for TF-IDF vectors and embeddings before cosine similarity.

Robust scaling (median and IQR) is the choice when outliers exist and you do not want them to dictate the scale.

⚠ Common mistake

Fit every scaler on the training data only and apply it to test data (Pipeline). Scaling does not change the shape of a distribution: a skewed feature stays skewed after standardisation; that needs a transformation.

34.3 Curing skew: log, power, Box–Cox, Yeo–Johnson

34.3.1 Why

Right-skewed variables (salary, price, counts) make linear models chase the few large values, inflate the influence of outliers, produce heteroscedastic residuals, and make distances dominated by the tail. Transforming toward symmetry often helps linear models, k -NN and neural networks (not trees).

34.3.2 Log transform

$x' = \log x$ (or $\log(1 + x)$ when zeros exist). It turns multiplicative differences into additive ones: 1, 2, 4, 8, 16 becomes 0, 0.69, 1.39, 2.08, 2.77: equally spaced. Skewness $0.89 \rightarrow 0$. Requires $x > 0$ (or $x \geq 0$ with $\log(1+x)$).

34.3.3 Other power transforms

Square root (milder, for counts), reciprocal $1/x$ (strong), square (for left skew).

34.3.4 Box–Cox

A family indexed by λ , chosen by maximum likelihood to make the data as normal as possible:

$$x^{(\lambda)} = \begin{cases} \frac{x^\lambda - 1}{\lambda} & \lambda \neq 0 \\ \log x & \lambda = 0 \end{cases} \quad (x > 0).$$

$\lambda = 1$: no change (shift); $\lambda = 0.5$: square-root-like; $\lambda = 0$: log; $\lambda = -1$: reciprocal-like.

📊 Small example: Box–Cox

Fibonacci-like values 1, 2, 3, 5, 8, 13, 21, 34, 55 (skewness 1.25). SciPy picks $\lambda = 0.028$ (almost a log); skewness after ≈ -0.01 . With $\lambda = 0.5$: 1, 2, 4, 8, 16 become 0, 0.83, 2, 3.66, 6.

34.3.5 Yeo–Johnson

A Box–Cox variant that also accepts zero and negative values (profit/loss, change in salary). scikit-learn’s `PowerTransformer` uses it by default and standardises the result. **Quantile transformer**: maps values to a uniform or normal distribution by ranks; very strong, removes outliers’ influence, but non-linear distortion of distances.

34.3.6 Transforming the target

If you log-transform y for training, predictions must be back-transformed with $\exp$, and $\exp(\mathbb{E}[\log y])$ underestimates $\mathbb{E}[y]$ (Chapter 2): apply a smearing correction or use a Gamma/Tweedie GLM objective instead. `scikit-learn`: `TransformedTargetRegressor(func=np.log1p, inverse_func=np.expm1)`.

34.4 Numerical to categorical: binning

Why bin? To capture non-linear effects in linear models (age bands), to reduce the influence of noise and outliers, to make results explainable (“freshers”, “mid”, “senior”), or because the business thinks in bands (salary slabs).

- **Equal-width:** split the range into equal intervals (`pd.cut`). Simple; sensitive to outliers (most points in one bin).
- **Equal-frequency (quantile):** each bin has the same count (`pd.qcut`). Robust to skew; bin edges are data-dependent.
- **Domain-based:** experience 0, 1–3, 4–7, 8–12, 13+.
- **Clustering-based:** run 1-D K-Means on the values and use cluster IDs as bins (`KBinsDiscretizer(strategy="kmeans")`): edges fall in natural gaps.
- **Supervised / tree-based:** fit a shallow decision tree of the target on the feature and use its split points (optimal binning, used in credit scoring with weight of evidence).

Costs: loss of information within bins, arbitrary edges, more parameters. Tree models do not need binning.

Small example: three binnings of MBA %

Values 51.6, 53.3, 55.5, 57.8, 58.8, 59.4, 62.1, 66.3. Edges [0, 55, 60, 100] (domain): counts 2, 4, 2. Quartiles (`qcut(q=4)`): 2, 2, 2, 2. Equal-width into 3 over [51.6, 66.3] (width 4.9): [51.6, 56.5) has 3, [56.5, 61.4) has 3, [61.4, 66.3] has 2.

34.5 Categorical to numerical: encodings

Models need numbers. The encoding must not invent information that is not there (a fake order) and must not leak the target.

34.5.1 Label / ordinal encoding

Map categories to integers. Correct for **ordinal** variables where order matters: junior $\rightarrow$ 0, mid $\rightarrow$ 1, senior $\rightarrow$ 2. For **nominal** variables (city), integers create a fake order and fake distances (“Pune = 2 is between Delhi = 1 and Agra = 3”), which misleads linear models and k -NN; tree models tolerate it somewhat (they can split repeatedly), and it is memory-efficient for very high cardinality in GBMs.

34.5.2 One-hot encoding (OHE)

One binary column per category. City $\in$ {Delhi, Pune, Agra}: Pune $\rightarrow$ (0, 1, 0). No fake order. Drawbacks: many columns for high cardinality (50,000 companies), sparse; unseen categories at test time (`handle_unknown="ignore"`).

34.5.3 The dummy variable trap

With k one-hot columns and an intercept, the columns always sum to 1, which equals the intercept column: **perfect multicollinearity**. In OLS, $X^T X$ is singular and coefficients are not unique (adding c to every city coefficient and subtracting c from the intercept gives identical predictions). **Fix**: drop one dummy (drop_first=True); the dropped category becomes the **reference level** and each remaining coefficient is “the difference from the reference”. Regularised models (Ridge, Lasso, logistic regression in scikit-learn) and trees do not need the drop (the penalty picks a unique solution), and keeping all k is more interpretable for them.

Small example: the trap in numbers

Rows Delhi, Pune, Agra encoded (1, 0, 0), (0, 1, 0), (0, 0, 1) plus intercept 1: column sum $d + p + a = 1 = \text{intercept}$. Coefficients (intercept 5, Delhi 1, Pune 2, Agra 3) and (intercept 6, Delhi 0, Pune 1, Agra 2) predict exactly the same.

34.5.4 Frequency (count) encoding

Replace each category by its frequency (or count) in training data. Pune, Pune, Pune, Delhi, Delhi, Agra $\rightarrow$ 0.5, 0.5, 0.5, 0.333, 0.333, 0.167. One column whatever the cardinality; captures “popularity” (common companies vs rare); different categories with equal frequency collide.

34.5.5 Target (mean) encoding

Replace each category by the mean of the target for that category. Powerful for high-cardinality features (company, skill, locality), but dangerous:

- **Leakage**: computing a row’s encoding with its own label lets the model “see” the answer; rare categories become perfect predictors. Use **out-of-fold** encoding (each row’s value computed from the other folds) or ordered/leave-one-out schemes.
- **Noise for rare categories**: smooth toward the global mean:

$$\text{enc}(c) = \frac{n_c \bar{y}_c + m \bar{y}}{n_c + m}.$$

Small example: smoothed target encoding

Placed (1/0) by city: Pune 1, 1, 0; Delhi 0, 1; Agra 1. Global mean $4/6 = 0.667$, smoothing $m = 2$. Agra: raw 1.0 from one student; smoothed $\frac{1 \cdot 1 + 2 \cdot 0.667}{3} = 0.778$. Delhi: raw 0.5, smoothed 0.583. Pune: raw and smoothed 0.667. Rare categories are pulled toward the global mean.

It is the MAP/empirical-Bayes idea of Chapter 12. scikit-learn’s TargetEncoder cross-fits automatically; CatBoost does ordered target statistics.

34.5.6 Other encodings

Binary encoding (integer code written in binary bits: $\log_2 k$ columns), **hashing** (hash categories into a fixed number of columns; collisions but constant memory; used in large click models), **weight of evidence** $\ln \frac{P(\text{cat}|y=1)}{P(\text{cat}|y=0)}$ (credit scoring), **learned embeddings** (neural networks: each category gets a trainable vector; the standard for user/job IDs).

Situation	Encoding
Ordered categories	ordinal
Few nominal categories, linear model	one-hot (drop one if unregularised with intercept)
Many categories, tree model	native categorical (LightGBM/CatBoost), target or frequency encoding
Many categories, linear model	target encoding (out-of-fold, smoothed), hashing, or grouping rare levels into "Other"
Neural network with IDs	embeddings

34.6 Imbalanced datasets: the complete toolkit

Chapter 13 introduced the main ideas; here is the full list the checklist names.

- **Random undersampling** ("random deletion"): randomly delete majority-class rows until the ratio is acceptable. Fast; throws away information; good with huge data.
- **Random oversampling**: duplicate minority rows. Keeps all data; risk of overfitting the duplicated points.
- **Weighting the dataset (class weights / sample weights)**: multiply each row's loss by a weight, larger for the minority class ($w_c = \frac{n}{K n_c}$). No data changed; supported by most models (`class_weight="balanced"`, `scale_pos_weight`).
- **NearMiss** (informed undersampling): keep majority points that are *close* to minority points (version 1: smallest average distance to the 3 nearest minority points), so the model focuses on the boundary. Sensitive to noise.
- **Tomek links**: pairs of opposite-class points that are each other's nearest neighbours; removing the majority member cleans the boundary. **Edited nearest neighbours** removes majority points misclassified by their neighbours.
- **SMOTE**: synthetic minority points by interpolating between a minority point and one of its k minority neighbours: $\mathbf{x} + u(\mathbf{x}_{nn} - \mathbf{x})$.
- **Borderline-SMOTE**: first find minority points "in danger": those for which more than half (but not all) of their m nearest neighbours (of any class) are majority; generate synthetic points only from those. Points surrounded entirely by majority are treated as noise and skipped. Focuses effort where the boundary is.
- **ADASYN**: generate more synthetic points for minority points that have more majority neighbours (harder to learn), in proportion to that difficulty.
- **SMOTE-NC** (categorical features), **SMOTE + Tomek** / **SMOTE + ENN** (oversample then clean).
- **Threshold moving** and **proper metrics** (PR AUC, recall at precision): often the most effective and simplest step.
- **Algorithm-level**: focal loss, balanced random forest, EasyEnsemble (many undersampled subsets, ensembled).

Small example: Borderline-SMOTE danger set

$m = 5$. Minority point A has neighbours: 3 majority, 2 minority ($3/5 > \text{half}$, not all): *danger*, used for synthesis. Point B: 1 majority, 4 minority: *safe*, skipped. Point C: 5 majority: *noise*, skipped.

Always resample only the training fold (imblearn Pipeline); never the test set; recalibrate probabilities afterwards if they are used.

** In practice: where these choices matter in ML at InfoEdge**

Salary engine: log or Box–Cox target, robust scaling of experience, target-encoded company with smoothing. Resdex ranking: frequency-encoded skills, embeddings of titles. Fraud detection: class weights and threshold tuning rather than heavy SMOTE (fraud patterns change, synthetic points can mislead). Seeker segmentation: robust scaling before K-Means.

34.7 Interview questions

Q34.1

What is the difference between standardisation and normalisation? When do you use each?

🕒 **In one line.** Standardisation rescales to mean 0 and standard deviation 1, with no fixed range; normalisation (min–max) rescales to $[0, 1]$; use standardisation for most models (linear with penalties, SVM, PCA, neural networks), min–max when a bounded range is needed and there are no extreme outliers, and robust scaling when outliers exist.

📊 **Worked example.** Column 10, 20, 30, 40, 1000: min–max gives 0, 0.01, 0.02, 0.03, 1 (normal values crushed); robust gives $-1, -0.5, 0, 0.5, 48.5$ (normal values keep their spread).

Q34.2

Which models need feature scaling, and why don't tree models?

🕒 **In one line.** Distance-, dot-product- and gradient-based models (k-NN, K-Means, SVM, PCA, regularised linear models, neural networks) need it; trees split on thresholds of single features, and a monotone rescaling does not change which rows fall on each side, so they are unaffected.

Q34.3

How do you fix a right-skewed feature? Explain Box–Cox.

🕒 **In one line.** Apply a log (or $\log(1 + x)$), square root, or a fitted power transform; Box–Cox is the family $(x^\lambda - 1)/\lambda$ (log at $\lambda = 0$) with λ chosen by maximum likelihood to make the data as normal as possible; it needs positive data, and Yeo–Johnson extends it to zero and negative values.

📊 **Worked example.** Values 1–55 with skewness 1.25: Box–Cox chooses $\lambda \approx 0.03$, skewness ≈ 0 .

Q34.4

💡 **Think deeper** Does standardising a skewed feature make it normal?

🕒 **In one line.** No: standardisation is a linear shift and rescale, so it preserves the shape, including skewness and kurtosis; only non-linear transforms (log, Box–Cox, quantile) change

shape.

Q34.5

What is the dummy variable trap and how do you avoid it?

🕒 **In one line.** With one dummy per category plus an intercept, the dummies sum to the intercept column, creating perfect multicollinearity, so OLS coefficients are not unique; drop one category (the reference level) or omit the intercept; regularised models and trees do not need the drop.

Q34.6

One-hot vs ordinal vs target vs frequency encoding: when do you use each?

🕒 **In one line.** Ordinal for genuinely ordered categories; one-hot for low-cardinality nominal features (especially for linear models); target encoding for high-cardinality features when done out-of-fold and smoothed; frequency encoding when popularity itself is informative and you need a single compact column.

Q34.7

💡 **Think deeper** Why is naive target encoding dangerous, and how do you do it safely? Compute the smoothed value for a city with one student (placed) when the overall placement rate is 0.667 and $m = 2$.

🕒 **In one line.** Because each row's encoding includes its own label, rare categories become perfect fake predictors (leakage and overfitting); compute encodings out-of-fold and smooth toward the global mean: $(1 \cdot 1 + 2 \cdot 0.667)/(1 + 2) = 0.778$.

Q34.8

How do you convert a numeric feature into categories, and why would you?

🕒 **In one line.** Bin it: equal-width, equal-frequency (quantiles), domain-defined bands, clustering-based (1-D K-Means), or tree-based optimal bins; useful to capture non-linearity in linear models, reduce noise, and produce explainable segments, at the cost of within-bin information.

Q34.9

List the techniques for imbalanced data and when you would use each.

🕒 **In one line.** Class/sample weighting (first choice, no data change), random undersampling or "random deletion" (very large data), random oversampling, SMOTE and its variants (Borderline-SMOTE near the boundary, ADASYN for hard regions), NearMiss and Tomek links (informed undersampling/cleaning), ensembles of undersampled sets, threshold tuning and the right

metrics.

Q34.10

How does Borderline-SMOTE differ from SMOTE and ADASYN?

🎯 **In one line.** SMOTE interpolates from all minority points; Borderline-SMOTE only from minority points “in danger” (more than half but not all neighbours are majority), ignoring safe and noisy points; ADASYN creates more points for minority samples with more majority neighbours, in proportion to their difficulty.

Q34.11

🧠 **Think deeper** After class weighting your model’s predicted probabilities are too high. Why?

🎯 **In one line.** Weighting (like oversampling) changes the effective class prior seen in training, so the model learns an inflated base rate; the ranking is fine but probabilities must be recalibrated on unweighted validation data (or the intercept corrected) if used as probabilities.

Q34.12

What is row-wise L2 normalisation and where is it used?

🎯 **In one line.** Dividing each row vector by its length so it has norm 1, e.g. $(3, 4) \rightarrow (0.6, 0.8)$; used for TF-IDF and embedding vectors so that dot products equal cosine similarities.

CHAPTER 35

Clustering Variants and the Dimensionality-Reduction Family

🎥 **Companion video for this chapter:** [StatQuest: t-SNE, Clearly Explained](#)

Chapter 15 taught K-Means, GMM/EM and DBSCAN; Chapter 14 taught PCA, SVD and t-SNE. The checklist adds the K-Means family (k-medoids, k-modes, k-prototypes, k-means++), soft vs hard clustering, the Swiss roll dataset, SNE vs t-SNE, and the classifications “linear vs non-linear” and “deterministic vs non-deterministic” dimensionality reduction. This chapter fills those in.

35.1 The K-Means family

K-Means minimises squared Euclidean distance to cluster *means*. That choice has consequences: it needs numeric data, means are pulled by outliers, and a centroid may not be a real data point. The variants change the “centre” and the distance.

35.1.1 K-Means++ (initialisation)

Choose the first centre at random; choose each next centre with probability proportional to $D(x)^2$, the squared distance to the nearest centre already chosen (Chapter 15). Spreads initial centres apart; fewer bad local minima; expected cost within $O(\log K)$ of optimal.

35.1.2 K-Medoids (PAM)

Each cluster’s centre must be an actual data point (the **medoid**: the point with the smallest total distance to the others in its cluster). Works with *any* distance (Manhattan, cosine, edit distance between job titles) and is robust to outliers.

📊 Small example: mean vs medoid

Points 1, 2, 3, 100 in one cluster. Mean = 26.5, a value nobody has, dragged by 100. Medoid: total absolute distances from 2 is $1 + 0 + 1 + 98 = 100$ and from 3 is also 100 (tie; either is a medoid), far smaller than from 100 ($99 + 98 + 97 = 294$). The medoid stays among the typical points.

Algorithm (PAM): assign points to nearest medoid; for each medoid try swapping it with a non-medoid; keep swaps that lower total cost; repeat. Cost is higher than K-Means ($O(K(n - K)^2)$ per iteration); CLARA/CLARANS sample to scale.

35.1.3 K-Modes (categorical data)

For purely categorical records (city, function, education), use **simple matching dissimilarity**: the number of attributes that differ. The centre of a cluster is its **mode** per attribute.

Small example: k-modes

Cluster records: (Delhi, IT, Fresher), (Delhi, IT, Mid), (Pune, IT, Fresher). Mode per attribute: (Delhi, IT, Fresher). New record (Pune, Sales, Fresher) differs in 2 attributes from this mode.

35.1.4 K-Prototypes (mixed data)

Mixed numeric and categorical records: cost = squared Euclidean distance on numeric features + $\gamma \times$ number of categorical mismatches; γ balances the two parts (often around half the average numeric standard deviation). Centre = means for numeric, modes for categorical. Natural for customer data like (salary, experience, city, function).

35.1.5 Mini-batch K-Means

Update centroids from small random batches: much faster on millions of points, slightly worse objective.

35.2 Hard vs soft clustering

- **Hard:** each point belongs to exactly one cluster: K-Means, k-medoids, k-modes, DBSCAN (or noise), agglomerative.
- **Soft (fuzzy/probabilistic):** each point has a degree of membership in every cluster: GMM (responsibilities γ_{ik} , Chapter 15), **fuzzy c-means**.

Fuzzy c-means membership with fuzzifier $m > 1$: $u_{ik} = \left[\sum_j \left(\frac{d_{ik}}{d_{jk}} \right)^{2/(m-1)} \right]^{-1}$.


Small example: fuzzy memberships

A point at distance 1 from centre A and 3 from centre B, $m = 2$: $u_A = \frac{1}{1+(1/3)^2} = 0.9$, $u_B = 0.1$.

Soft clustering suits overlapping groups: a candidate who is 60% “data analyst” and 40% “business analyst” should be recommended both kinds of jobs.

35.3 The Swiss roll: why “non-linear” matters

Take a flat 2-D sheet of points coloured by position and roll it up like a Swiss-roll cake in 3-D. Two points on adjacent layers of the roll are close in straight-line (Euclidean) distance but far apart *along the sheet* (geodesic distance).



(Cross-section of a Swiss roll: the two marked points are close in a straight line but far along the curve.)

- **PCA** (linear) projects the roll flat, overlapping the layers: colours mix.
- **K-Means** with Euclidean distance cuts across layers.
- **Isomap, LLE, UMAP**, which respect local neighbourhoods and geodesic distances, “unroll” it into the original sheet.
- **DBSCAN / spectral clustering** follow the sheet’s connectivity.

It is the standard demonstration that data can lie on a low-dimensional curved **manifold** inside a high-dimensional space.

35.4 SNE vs t-SNE

SNE (stochastic neighbour embedding): convert distances to conditional probabilities $p_{j|i}$ with a Gaussian around each point (bandwidth from perplexity), do the same in low dimension with a Gaussian $q_{j|i}$, and minimise $\sum_i KL(P_i \| Q_i)$. Two problems: the asymmetric cost is hard to optimise, and the **crowding problem**: in 2-D there is not enough room to place all moderately distant neighbours at moderate distances, so clusters squash together in the middle.

t-SNE fixes both: (1) symmetric joint probabilities $p_{ij} = \frac{p_{j|i} + p_{i|j}}{2n}$ (simpler gradient); (2) in the low-dimensional space use a **Student-*t* distribution with one degree of freedom** (Cauchy), $q_{ij} \propto (1 + \|\mathbf{y}_i - \mathbf{y}_j\|^2)^{-1}$, whose heavy tails let moderately dissimilar points be placed far apart without much penalty, giving well-separated clusters.

Small example: why heavy tails help

At low-dimensional distance 3, a Gaussian similarity $e^{-9/2} \approx 0.011$; the Student-*t* similarity $1/(1 + 9) = 0.1$, nine times larger. To match a small high-dimensional similarity, t-SNE can push such pairs further apart.

Perplexity (5–50) $\approx$ effective number of neighbours per point; see Chapter 14 for reading t-SNE plots.

35.5 A taxonomy of dimensionality reduction

Method	Linear?	Deterministic?	Supervised?	New points?	Preserves / idea
PCA	yes	yes	no	yes	global variance
SVD / truncated SVD	yes	yes (exact)	no	yes	low-rank structure
LDA (Fisher)	yes	yes	yes	yes	class separation
Factor analysis, ICA	yes	ICA no	no	yes	latent factors / independent sources
Random projection	yes	no	no	yes	distances approximately (Johnson–Lindenstrauss)
Kernel PCA	no	yes	no	yes (approx.)	variance in kernel feature space
Classical MDS	no*	yes	no	no	pairwise distances
Isomap	no	yes	no	limited	geodesic distances on a neighbour graph
LLE	no	yes	no	limited	local linear reconstructions
SNE / t-SNE	no	no	no	no	local neighbourhoods (probabilities)
UMAP	no	no (seeded)	optional	yes	local + more global topology
Autoencoder	no	no	no	yes	reconstruction through a bottleneck

(*Classical MDS with Euclidean distances is equivalent to PCA.)

- **Linear:** new coordinates are linear combinations of the original features ($Z = XW$): projections onto flat subspaces. Fast, interpretable loadings, can transform new points; cannot unroll curved manifolds.
- **Non-linear:** can capture curved structure; usually costlier, harder to interpret, often no simple transform for new points.
- **Deterministic:** the same data always give the same output (PCA solves an eigenproblem exactly; Isomap and LLE given the neighbour graph).
- **Non-deterministic (stochastic):** depend on random initialisation or sampling (t-SNE gradient descent from random positions, UMAP, random projections, autoencoders, randomised SVD). Set seeds; run several times; do not over-interpret one run.

35.5.1 Brief descriptions

Kernel PCA: PCA on the kernel matrix (Chapter 9), finding non-linear components (concentric circles become separable). **Isomap:** build a k -NN graph, compute shortest-path (geodesic) distances, run classical MDS on them. **LLE:** express each point as a weighted combination of its neighbours, then find low-dimensional points preserving those weights. **UMAP:** builds a fuzzy neighbour graph and optimises a low-dimensional layout; faster than t-SNE, better global structure, supports transform. **Random projection:** multiply by a random Gaussian matrix; by the Johnson–Lindenstrauss lemma, $O(\log n / \epsilon^2)$ dimensions preserve all pairwise distances within a factor $1 \pm \epsilon$: a cheap first reduction for huge data.

** In practice: where these are used in ML**

- **k-prototypes** for segmenting recruiters by (spend, jobs posted, industry, city); **k-medoids** for picking representative real job descriptions per cluster.
- **Soft clustering** (GMM) for seekers who straddle two career tracks.
- **PCA** to compress 768-dimensional embeddings before an ANN index; **random projection** or **product quantisation** for very large indexes.
- **UMAP/t-SNE** to inspect whether resume embeddings separate by job function; **LDA** for a supervised 2-D view of classes.
- **Autoencoders** for non-linear compression and anomaly detection.

35.6 Interview questions

Q35.1

How do k-medoids, k-modes and k-prototypes differ from K-Means, and when would you use each?

🕒 In one line. K-Means uses means and squared Euclidean distance (numeric data); k-medoids uses actual data points as centres and any distance (robust to outliers, custom metrics); k-modes uses modes and mismatch counts (categorical data); k-prototypes combines both for mixed numeric and categorical data.

📊 Worked example. Cluster 1, 2, 3, 100: mean 26.5, medoid 2.

Q35.2

What is k-means++ and why does it help?

🕒 In one line. An initialisation that picks each new centre with probability proportional to its squared distance from existing centres, spreading centres out, which reduces bad local minima and comes with a guarantee of being within $O(\log K)$ of optimal in expectation.

Q35.3

Hard vs soft clustering: definitions and examples.

🕒 In one line. Hard clustering assigns each point to exactly one cluster (K-Means, DBSCAN, hierarchical); soft clustering gives each point a membership probability or degree for every cluster (GMM, fuzzy c-means).

📊 Worked example. Fuzzy c-means with distances 1 and 3, $m = 2$: memberships 0.9 and 0.1.

Q35.4

💡 Think deeper What is the Swiss roll dataset and what does it demonstrate?

🕒 **In one line.** A 2-D sheet rolled up in 3-D; points on neighbouring layers are close in Euclidean distance but far along the sheet, demonstrating that data can lie on a curved low-dimensional manifold that linear methods (PCA, Euclidean K-Means) cannot recover, while neighbourhood-based non-linear methods (Isomap, LLE, UMAP) can unroll it.

Q35.5

What is the difference between SNE and t-SNE? Why the Student- t ?

🕒 **In one line.** t-SNE symmetrises SNE's conditional probabilities and replaces the Gaussian in the low-dimensional space with a heavy-tailed Student- t (one degree of freedom), which solves the crowding problem by letting moderately distant points spread out, producing clearer clusters and an easier optimisation.

Q35.6

Classify these as linear or non-linear and deterministic or not: PCA, LDA, t-SNE, UMAP, kernel PCA, autoencoder, random projection.

🕒 **In one line.** PCA: linear, deterministic; LDA: linear, deterministic, supervised; t-SNE: non-linear, stochastic; UMAP: non-linear, stochastic; kernel PCA: non-linear, deterministic; autoencoder: non-linear, stochastic; random projection: linear, stochastic.

Q35.7

💡 **Think deeper** Why can't you use t-SNE outputs as features for a production classifier?

🕒 **In one line.** t-SNE has no mapping for new points (it optimises positions for the given data only), its output changes between runs, and it distorts global distances; use PCA, UMAP (which has a transform) or learned embeddings instead.

Q35.8

What is the Johnson–Lindenstrauss lemma and why is it useful?

🕒 **In one line.** Any n points in high dimension can be projected by a random linear map into $O(\log n / \epsilon^2)$ dimensions while preserving all pairwise distances within a factor $1 \pm \epsilon$; it justifies random projections as a fast, data-independent dimensionality reduction for huge datasets.

Q35.9

How would you cluster recruiters described by monthly spend, jobs posted, industry and city?

🕒 **In one line.** With k-prototypes (numeric: scaled spend and jobs; categorical: industry and city), choosing K by cost elbow and business usefulness, or by converting categories to embeddings or

target statistics and using K-Means or GMM.

Q35.10

Isomap vs LLE vs kernel PCA: one line each.

🎯 **In one line.** Isomap preserves geodesic distances computed on a neighbour graph; LLE preserves each point's linear reconstruction from its neighbours; kernel PCA performs PCA in a kernel-induced feature space.

CHAPTER 36

The Perceptron and the Four Application Areas: CV, NLP, RL and Federated Learning

► Companion video for this chapter: [3Blue1Brown: But what is a neural network?](#)

The checklist asks for the perceptron's structure and for "basics of computer vision, natural language processing, reinforcement learning and federated learning". Chapter 18 built multilayer networks; this chapter starts one step earlier with the original perceptron and its learning rule, then gives a self-contained basic layer for each application area, with the ML ideas each one uses.

36.1 The perceptron

36.1.1 Structure

Rosenblatt's perceptron (1958) is a single artificial neuron with a **step** activation:

$$\hat{y} = \begin{cases} 1 & \mathbf{w}^\top \mathbf{x} + b > 0 \\ 0 & \text{otherwise} \end{cases}$$

Inputs $x_1, \dots, x_d \rightarrow$ multiplied by weights $w_1, \dots, w_d \rightarrow$ summed with a bias $b \rightarrow$ step function $\rightarrow$ output 0/1. Geometrically it is a linear classifier: the boundary $\mathbf{w}^\top \mathbf{x} + b = 0$ is a hyperplane.

36.1.2 The perceptron learning rule

For each training example, predict; if wrong, nudge the weights toward the correct side:

$$\mathbf{w} \leftarrow \mathbf{w} + \eta (y - \hat{y}) \mathbf{x}, \quad b \leftarrow b + \eta (y - \hat{y}).$$

If correct, $y - \hat{y} = 0$ and nothing changes. If a positive was missed ($y = 1, \hat{y} = 0$), $\mathbf{w}$ moves toward $\mathbf{x}$; if a negative was wrongly fired, $\mathbf{w}$ moves away from $\mathbf{x}$.

📖 Small example: learning the AND gate by hand ($\eta = 1$, start $\mathbf{w} = (0, 0)$, $b = 0$)

Epoch 1: (0, 0): score 0 $\rightarrow$ predict 0, correct. (0, 1): correct. (1, 0): correct. (1, 1), $y = 1$: score 0 $\rightarrow$ predict 0, wrong: $\mathbf{w} = (1, 1)$, $b = 1$. Epoch 2: (0, 0), $y = 0$: score 1 $> 0 \rightarrow$ predict 1, wrong: $b = 0$. (0, 1): score 1, wrong: $\mathbf{w} = (1, 0)$, $b = -1$ After five epochs with mistakes, epoch 6 has none: $\mathbf{w} = (2, 1)$, $b = -2$. Check: (1, 1): $2 + 1 - 2 = 1 > 0 \rightarrow 1$; (1, 0): $0 \rightarrow 0$; (0, 1): $-1 \rightarrow 0$; (0, 0): $-2 \rightarrow 0$. AND learned.

36.1.3 Convergence theorem and its limit

Perceptron convergence theorem: if the data are linearly separable, the rule finds a separating hyperplane in a finite number of updates (at most $(R/\gamma)^2$, with R the largest input norm and γ the margin). If the data are *not* separable it never converges. **XOR** is not linearly separable, so a single perceptron cannot learn it; this result (Minsky and Papert, 1969) stalled neural network research until

multilayer networks trained with backpropagation, which need differentiable activations (sigmoid, ReLU) instead of the step function.

36.1.4 Perceptron vs logistic regression vs SVM

All three are linear classifiers. The perceptron updates only on mistakes and stops at *any* separating line; logistic regression minimises log loss and gives probabilities; the SVM finds the *maximum-margin* line (hinge loss). The perceptron's update is SGD on the loss $\max(0, -y \mathbf{w}^T \mathbf{x})$ with $y \in \{-1, 1\}$.

36.1.5 From perceptron to MLP

Replace the step with a smooth activation, stack layers, and train with backpropagation (Chapter 18): the multilayer perceptron (ANN). One hidden layer suffices to represent XOR.

36.2 Computer vision basics

- **Images as tensors:** height $\times$ width $\times$ channels (RGB), pixel values 0–255, usually scaled to $[0, 1]$ or standardised per channel.
- **Tasks:** *classification* (one label per image: kitchen); *object detection* (boxes + labels: where are the windows); *semantic segmentation* (a label per pixel); *instance segmentation* (separate each object); *OCR* (text in images: resume PDFs, ID cards); *image similarity/retrieval* (duplicate listing photos); *face verification*.
- **Models:** CNNs (Chapter 19), Vision Transformers; detection with YOLO or Faster R-CNN; segmentation with U-Net; CLIP for image–text matching.
- **Detection metrics:** **IoU** (intersection over union) of predicted and true boxes; a detection counts as correct if $\text{IoU} \geq 0.5$; **non-maximum suppression** removes overlapping duplicate boxes; **mAP** averages precision over recall levels and classes.

Small example: IoU

Predicted box (0, 0)–(4, 4) (area 16), true box (2, 2)–(6, 6) (area 16). Intersection (2, 2)–(4, 4), area 4. Union $16 + 16 - 4 = 28$. $\text{IoU} = 4/28 = 0.14$: not a correct detection at the 0.5 threshold.

Augmentation, transfer learning from ImageNet, and class imbalance handling are the practical essentials.

36.3 NLP basics (a map of Chapters 16, 17 and 21)

Text $\rightarrow$ tokens (words or subwords) $\rightarrow$ vectors (BoW/TF-IDF, static embeddings, contextual embeddings) $\rightarrow$ model (linear, RNN, Transformer). Core tasks: text classification, NER/sequence labelling, semantic similarity and search, question answering, summarisation, translation, generation. Evaluation: accuracy/F1 for classification, exact match/F1 for QA, BLEU/ROUGE for generation (with human review), recall@k and NDCG for retrieval.

36.4 Reinforcement learning basics

36.4.1 The setting

An **agent** interacts with an **environment** over time. At each step it observes a **state** s , takes an **action** a , receives a **reward** r and moves to a new state. The goal is a **policy** $\pi(a | s)$ that maximises the expected **return** $G = r_1 + \gamma r_2 + \gamma^2 r_3 + \dots$, where the **discount** $\gamma \in [0, 1)$ makes near rewards worth more. Unlike supervised learning, there are no correct labels, only rewards, often delayed. This is formalised as a **Markov decision process** (states, actions, transition probabilities, rewards, γ): the next state depends only on the current state and action (the Markov property, Chapter 40).

Small example: discounted return

Rewards 1, 0, 2 with $\gamma = 0.9$: $G = 1 + 0.9 \cdot 0 + 0.81 \cdot 2 = 2.62$.

36.4.2 Value functions and Q-learning

$Q(s, a)$ = expected return from taking action a in state s and acting well afterwards. **Q-learning** learns it from experience:

$$Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma \max_{a'} Q(s', a') - Q(s, a)].$$

The bracket is the **temporal-difference error**: how much better or worse things went than expected.

Small example: one Q-learning update

$Q(s, a) = 0$, reward $r = 1$, $\gamma = 0.9$, best next value $\max_{a'} Q(s', a') = 2$, learning rate $\alpha = 0.5$: target = $1 + 0.9 \cdot 2 = 2.8$; $Q(s, a) \leftarrow 0 + 0.5(2.8 - 0) = 1.4$.

Deep Q-networks approximate Q with a neural network. **Policy gradient** methods (REINFORCE, PPO) adjust the policy directly to increase the probability of actions that led to high return; PPO and GRPO are used to fine-tune LLMs from human or verifiable rewards (Chapter 23).

36.4.3 Exploration vs exploitation and bandits

Always choosing the current best action (exploitation) may miss better ones; trying others (exploration) costs reward now. **ϵ -greedy**: explore at random with probability ϵ . **Multi-armed bandits** are one-step RL: choose which of several options (e.g. which job-alert template) to show, observe a reward (click), update; **Thompson sampling** samples from Beta posteriors of each option's success rate (Chapter 12); **contextual bandits** condition on user features. Bandits are the most common RL in industry: recommendation exploration, notification timing, ad and creative selection.

36.5 Federated learning basics

36.5.1 The problem

Data live on many devices or with many organisations (phones, hospitals, companies) and cannot be centralised because of privacy, regulation or bandwidth. **Federated learning** trains a shared model without moving the raw data.

36.5.2 FedAvg

1. The server sends the current model to a sample of clients.
2. Each client trains locally on its own data for a few epochs.
3. Clients send back model updates (weights or gradients), not data.
4. The server averages them, weighted by each client's number of examples: $\mathbf{w} = \sum_k \frac{n_k}{n} \mathbf{w}_k$.
5. Repeat.

Small example: weighted averaging

Client A has 100 examples and local weight 0.2; client B has 300 examples and local weight 0.6.
 New global weight = $\frac{100 \cdot 0.2 + 300 \cdot 0.6}{400} = 0.5$.

Challenges: non-IID data (each phone's data differ), communication cost, stragglers, privacy leakage through updates (mitigated by secure aggregation and **differential privacy**: adding calibrated noise), poisoning by malicious clients. **Uses:** next-word prediction on phone keyboards; cross-hospital models; for a job platform, on-device personalisation of a mobile app without uploading private activity.

In practice: these four areas at InfoEdge

CV: 99acres photo quality and room classification, duplicate-listing detection, Jeevansathi photo moderation. NLP: resume parsing, taxonomy, search, chat agents. RL: bandits for notification channel and timing (seeker reactivation), exploration for new jobs in recommendation, RL fine-tuning of LLMs (GRPO on the GenAI slide). Federated learning: on-device models in the mobile apps where privacy matters.

36.6 Interview questions

Q36.1

Describe the perceptron's structure and learning rule. Will it learn XOR?

 **In one line.** A weighted sum of inputs plus bias passed through a step function, trained by $\mathbf{w} \leftarrow \mathbf{w} + \eta(y - \hat{y})\mathbf{x}$ on mistakes; it converges whenever the data are linearly separable (AND, OR) but never for XOR, which is not linearly separable.

 **Worked example.** AND from zeros with $\eta = 1$ converges to $\mathbf{w} = (2, 1)$, $b = -2$.

Q36.2

Perceptron vs logistic regression vs linear SVM.

 **In one line.** All learn linear boundaries; the perceptron stops at any separating hyperplane using mistake-driven updates and has no probabilities; logistic regression minimises log loss and outputs calibrated probabilities; the SVM maximises the margin with hinge loss and is the most robust boundary.

Q36.3

 **Think deeper** Why did the step activation have to be replaced for multilayer networks?

🕒 **In one line.** Its derivative is zero everywhere (and undefined at 0), so gradients cannot flow backward through it; backpropagation needs differentiable activations like sigmoid, tanh or ReLU.

Q36.4

What are classification, detection and segmentation in computer vision? Compute IoU for boxes (0, 0)–(4, 4) and (2, 2)–(6, 6).

🕒 **In one line.** Classification gives one label per image, detection gives boxes with labels, segmentation labels every pixel; $\text{IoU} = 4/28 = 0.14$.

Q36.5

What is reinforcement learning and how does it differ from supervised learning?

🕒 **In one line.** An agent learns a policy by acting in an environment and receiving rewards, maximising discounted return; there are no labelled correct actions, feedback is delayed and depends on the agent's own choices, so exploration is needed.

Q36.6

Write the Q-learning update and compute it for $Q = 0$, $r = 1$, $\gamma = 0.9$, $\max Q' = 2$, $\alpha = 0.5$.

🕒 **In one line.** $Q \leftarrow Q + \alpha[r + \gamma \max Q' - Q] = 0 + 0.5(1 + 1.8) = 1.4$.

Q36.7

 **Think deeper** Where would a job platform use reinforcement learning or bandits?

🕒 **In one line.** Choosing which notification, channel and time to send each seeker (contextual bandits with click/apply rewards), allocating exposure to new jobs in recommendations (exploration), sequential recommendations optimising long-term engagement, and fine-tuning LLM agents with reward signals.

⚙️ **The full answer.** Full RL is rarely needed; bandits capture most value with simpler, safer learning. Always keep a holdout to measure true uplift.

Q36.8

What is federated learning, and what is FedAvg?

🎯 **In one line.** Training a shared model across many clients without collecting their raw data; FedAvg has clients train locally and the server average their weights, weighted by data size.

📊 **Worked example.** Clients with 100 and 300 examples and weights 0.2 and 0.6: global weight 0.5.

💡 **If the interviewer pushes.** “Is it fully private?” Not by itself: updates can leak information; add secure aggregation and differential privacy.

Q36.9

Explain the exploration–exploitation trade-off and ϵ -greedy.

🎯 **In one line.** Exploiting the current best option maximises short-term reward but may miss better options; exploring gathers information at a cost; ϵ -greedy explores a random option with probability ϵ and exploits otherwise.

Q36.10

What is a Markov decision process?

🎯 **In one line.** A formal RL setting of states, actions, transition probabilities, rewards and a discount factor, where the next state and reward depend only on the current state and action (the Markov property).

CHAPTER 37

Time Series: Components, Stationarity, ARIMA, SARIMAX and ANOVA

🎥 **Companion video for this chapter:** [ritvikmath: Time Series Talk, ARIMA Model \(search\)](#)

The checklist lists “time series (ANOVA, ARIMA, SARIMAX)” and warns that what is asked “completely depends on the projects you put in your resume”. Chapter 20 covered deep-learning forecasters. This chapter teaches the classical statistical toolkit from zero, which interviewers expect you to compare against any LSTM you mention, plus one-way ANOVA.

37.1 What makes time series different

Observations are ordered in time and **depend on the past**: today’s applications are related to yesterday’s. Two consequences: (1) random train/test splits leak the future (always split by time); (2) the “independent errors” assumption of ordinary regression fails, so special models are needed.

37.2 Components

A series is usually viewed as a combination of:

- **Trend**: long-run increase or decrease (job postings growing year on year).
- **Seasonality**: a pattern repeating with a fixed period (weekly: fewer applications on Sundays; yearly: hiring peaks after appraisals and in campus season).
- **Cycle**: rises and falls without a fixed period (economic cycles, hiring freezes).
- **Noise**: what remains.

Additive: $y_t = T_t + S_t + R_t$ (seasonal swing constant in size). **Multiplicative**: $y_t = T_t \times S_t \times R_t$ (seasonal swing grows with the level; take logs to make it additive). `statsmodels.tsa.seasonal_decompose` or `STL` separates them.

37.3 Stationarity

A series is (weakly) **stationary** if its mean, variance and autocorrelation structure do not change over time. Most classical models assume it. Trends and seasonality make a series non-stationary.

- **Check**: plot it (moving mean/variance); the **Augmented Dickey–Fuller (ADF)** test (null hypothesis: non-stationary, has a unit root; small p-value $\Rightarrow$ stationary); KPSS (null: stationary).
- **Fix**: **differencing** $\nabla y_t = y_t - y_{t-1}$ removes a linear trend; seasonal differencing $y_t - y_{t-s}$ removes seasonality of period s ; log transform stabilises a variance that grows with the level.

Small example: differencing

Monthly postings 100, 110, 125, 135, 150. First differences: 10, 15, 10, 15 (no trend left, mean 12.5). Second differences: 5, -5, 5 (over-differencing introduces oscillation, so stop at the first difference here).

37.4 Autocorrelation: ACF and PACF

The **autocorrelation at lag k** , ρ_k , is the correlation between y_t and y_{t-k} .

Small example: lag-1 autocorrelation

Series 2, 4, 6, 8, 10 (mean 6): deviations -4, -2, 0, 2, 4. Lag-1 products: $(-4)(-2) + (-2)(0) + (0)(2) + (2)(4) = 16$; total squares 40; $r_1 = 0.4$. A trending series shows slowly decaying positive ACF.

PACF at lag k : the correlation between y_t and y_{t-k} after removing the effect of the lags in between. Reading them chooses model orders:

Pattern	ACF	PACF
AR(p)	decays gradually	cuts off after lag p
MA(q)	cuts off after lag q	decays gradually
ARMA	both decay	both decay
Seasonal	spikes at lags $s, 2s, \dots$	spikes at s

37.5 The ARIMA family, built up

37.5.1 AR(p): autoregression

Today is a weighted sum of the last p values plus noise: $y_t = c + \phi_1 y_{t-1} + \dots + \phi_p y_{t-p} + \varepsilon_t$. A regression on its own past.

Small example: AR(1) forecast

$y_t = 4 + 0.6y_{t-1} + \varepsilon_t$, last value 15. One step ahead: $4 + 0.6 \cdot 15 = 13$. Two steps: $4 + 0.6 \cdot 13 = 11.8$. Long run: the forecasts approach the mean $c/(1 - \phi) = 4/0.4 = 10$. Stationarity needs $|\phi| < 1$.

37.5.2 MA(q): moving average (of past errors)

$y_t = \mu + \varepsilon_t + \theta_1 \varepsilon_{t-1} + \dots + \theta_q \varepsilon_{t-q}$: today depends on recent *shocks* (surprises). A one-day outage shock affects the next few days only. (Not the same as a rolling-mean smoother.)

37.5.3 ARMA(p, q) and ARIMA(p, d, q)

ARMA combines both for stationary series. **ARIMA** adds **I**ntegration: difference the series d times first, then fit ARMA(p, q). ARIMA(0,1,0) is a random walk (forecast = last value: the “naive” baseline). ARIMA(0,1,1) is equivalent to simple exponential smoothing.

37.5.4 Seasonal ARIMA: SARIMA(p, d, q)(P, D, Q) $_s$

Adds seasonal AR, differencing and MA terms at lags that are multiples of the season length s (7 for daily data with weekly pattern, 12 for monthly data with yearly pattern). Example: SARIMA(1,1,1)(1,1,1) $_7$ for daily applications.

37.5.5 SARIMAX: adding exogenous variables

The X adds external regressors x_t : $y_t = \beta^\top x_t + \text{SARIMA errors}$. For applications: number of new job postings, holidays, marketing campaigns, campus-season indicator. Caution: at forecast time you must know (or forecast) future values of the exogenous variables.

```
from statsmodels.tsa.statespace.sarimax import SARIMAX
m = SARIMAX(y_train, exog=X_train, order=(1, 1, 1),
            seasonal_order=(1, 1, 1, 7)).fit(dispatch=False)
forecast = m.forecast(steps=14, exog=X_future)
```

37.5.6 Choosing orders

Make stationary (choose d, D via ADF and plots), read ACF/PACF for p, q, P, Q , compare candidates by AIC/BIC (lower is better: fit penalised by complexity), check residuals are white noise (Ljung–Box test, residual ACF), or use auto-ARIMA. Validate with walk-forward (rolling-origin) evaluation.

37.6 Exponential smoothing

Simple exponential smoothing: level $\ell_t = \alpha y_t + (1 - \alpha)\ell_{t-1}$; forecast = current level. Recent observations weigh more.

Small example: SES

$\alpha = 0.5$, initial level 10, observations 12, 11, 13: levels 11, 11, 12. Forecast 12.

Holt adds a trend; **Holt–Winters** adds seasonality; the ETS framework unifies them. Prophet (trend + seasonality + holidays) is a popular practical tool.

37.7 Evaluating forecasts

Walk-forward validation; metrics MAE, RMSE, MAPE (fails near zero), sMAPE, MASE (scaled by the naive forecast's error: < 1 beats naive). Always compare with the naive and seasonal-naive baselines; report prediction intervals, not only point forecasts.

37.8 ANOVA (analysis of variance)

37.8.1 The question

Do several groups have the same mean? (Do average applications per job differ between Monday, Wednesday and Saturday postings? Do three email templates give different mean click rates?) Running many pairwise t -tests inflates false positives; ANOVA tests all at once.

37.8.2 The idea

Split total variation into **between-group** variation (how far group means are from the grand mean) and **within-group** variation (spread inside groups). If the group means are really equal, the between-group variation should be about as large as what noise alone produces.

$$F = \frac{MSB}{MSW} = \frac{SSB/(k-1)}{SSW/(N-k)}.$$

Small example: one-way ANOVA by hand

Applications per job for three posting days: A = 5, 6, 7 (mean 6), B = 8, 9, 10 (mean 9), C = 4, 5, 6 (mean 5). Grand mean $60/9 = 6.67$. $SSB = 3[(6 - 6.67)^2 + (9 - 6.67)^2 + (5 - 6.67)^2] = 3(0.444 + 5.444 + 2.778) = 26$. $SSW = (1 + 0 + 1) \times 3 = 6$. $k = 3$, $N = 9$: $MSB = 26/2 = 13$, $MSW = 6/6 = 1$, $F = 13$, p-value = 0.0066. Reject equal means: posting day matters. Post-hoc tests (Tukey HSD) then say which pairs differ. Effect size $\eta^2 = 26/32 = 0.81$.

Assumptions: independent observations, roughly normal residuals, equal variances across groups (Levene's test; use Welch's ANOVA if unequal); non-parametric alternative: Kruskal–Wallis. Two-way ANOVA tests two factors and their interaction. In time series, ANOVA-style comparisons check whether seasonal means (by weekday or month) differ, which justifies seasonal terms.

In practice: time series at InfoEdge

Forecasting daily applications and job postings by category and city (staffing of sales and support teams, inventory of jobs for campaigns); Talent Pulse “expected hike” trends; anomaly alerts when today's applications fall far below the seasonal forecast (pipeline breakages); A/B tests where ANOVA compares several variants. A strong practical recipe: SARIMAX or ETS baselines per series, LightGBM on lag/calendar features across many series, and an LSTM only if it beats both on walk-forward MASE.

37.9 Interview questions

Q37.1

What does ARIMA(p, d, q) mean? Explain each part.

 **In one line.** p autoregressive terms (dependence on the last p values), d rounds of differencing to make the series stationary, and q moving-average terms (dependence on the last q forecast errors).

 **Worked example.** AR(1) $y_t = 4 + 0.6y_{t-1}$ from 15 forecasts 13, 11.8, ... toward 10.

Q37.2

What is stationarity, why does it matter and how do you achieve it?

 **In one line.** Constant mean, variance and autocorrelation over time; ARMA-type models and many statistics assume it, otherwise relationships learned on the past do not hold in the future; achieve it with differencing (trend), seasonal differencing (seasonality) and log transforms

(growing variance), verified with ADF/KPSS tests and plots.

Q37.3

How do ACF and PACF help choose p and q ?

🎯 **In one line.** An $AR(p)$ process has a PACF that cuts off after lag p and a decaying ACF; an $MA(q)$ process has an ACF that cuts off after lag q and a decaying PACF; seasonal spikes at lag s suggest seasonal terms.

Q37.4

ARIMA vs SARIMA vs SARIMAX?

🎯 **In one line.** SARIMA adds seasonal AR, differencing and MA terms at multiples of the season length; SARIMAX further adds exogenous regressors (holidays, campaigns, job postings), which must be known or forecast for the future.

Q37.5

🧠 **Think deeper** You built an LSTM forecaster for your project. How would you defend it against “why not ARIMA?”

🎯 **In one line.** Show walk-forward results against naive, seasonal-naive, SARIMAX/ETS and a GBM on lags with a scale-free metric (MASE); argue for the LSTM only where it wins, typically with many related series, long histories, non-linear effects or rich exogenous inputs, and acknowledge ARIMA’s interpretability and small-data strength.

Q37.6

What is one-way ANOVA? Compute F for groups (5, 6, 7), (8, 9, 10), (4, 5, 6).

🎯 **In one line.** A test of whether several group means are equal, comparing between-group to within-group variance; $SSB = 26$, $SSW = 6$, $F = (26/2)/(6/6) = 13$, $p \approx 0.007$, so the means differ.

💡 **If the interviewer pushes.** “Why not three t -tests?” Each at 5% gives about a 14% chance of at least one false positive across three tests; ANOVA controls the overall error rate, then post-hoc tests with corrections compare pairs.

Q37.7

Additive vs multiplicative seasonality?

🎯 **In one line.** Additive: the seasonal swing has a constant size regardless of level; multiplicative: the swing grows proportionally with the level, which a log transform turns into additive.

Q37.8

How should you validate a time-series model?

🎯 **In one line.** With walk-forward (rolling or expanding window) validation that always trains on the past and tests on the future, comparing against naive baselines, with MAE/MASE, and checking that residuals are white noise.

Q37.9

Compute simple exponential smoothing with $\alpha = 0.5$, initial level 10 and observations 12, 11, 13.

🎯 **In one line.** Levels 11, 11, 12; the next forecast is 12.

Q37.10

What do AIC and BIC do in model selection?

🎯 **In one line.** They score models by goodness of fit (log-likelihood) penalised by the number of parameters (BIC penalises more for large samples); lower is better, preventing overfitting with too many AR/MA terms.

CHAPTER 38

Probability from Zero: Events, Random Variables, Moments and Distributions

🎥 **Companion video for this chapter:** [3Blue1Brown: Bayes theorem, the geometry of changing beliefs](#)

Probability is the language in which ML losses, Naive Bayes, MLE/MAP, GMMs, dropout, sampling and A/B tests are written. The checklist asks for PDFs and PMFs, expectation, variance, moments, and the common distributions with real-life examples. Your Maths book drills these for the OA; this chapter teaches them again from zero with many examples and, for each idea, a short “where in ML” note.

38.1 Experiments, outcomes and events

An **experiment** is anything with an uncertain outcome (roll a die, show a job to a seeker). The **sample space** S is the set of all outcomes; an **event** is a subset of outcomes.

📊 Small example: counting outcomes

One die: $S = \{1, \dots, 6\}$. Event “even” = $\{2, 4, 6\}$, probability $3/6 = 0.5$. Two dice: 36 equally likely outcomes; “sum is 7” has 6 outcomes: $1/6$.

For equally likely outcomes, $P(A) = \frac{|A|}{|S|}$. In general probabilities follow three **axioms**: $P(A) \geq 0$; $P(S) = 1$; for mutually exclusive events $P(A \cup B) = P(A) + P(B)$. Everything else follows:

- Complement: $P(\text{not } A) = 1 - P(A)$. Probability of at least one click in 5 independent impressions with CTR 0.1: $1 - 0.9^5 = 0.41$.
- Union: $P(A \cup B) = P(A) + P(B) - P(A \cap B)$. If 60% of seekers know Python, 40% SQL and 25% both: $0.6 + 0.4 - 0.25 = 0.75$ know at least one.

38.2 Conditional probability, independence and Bayes

38.2.1 Conditional probability

$P(A | B) = \frac{P(A \cap B)}{P(B)}$: the probability of A once we know B happened; we restrict attention to the world where B is true.

📊 Small example: conditional

Of 1,000 applicants, 200 have work experience; 120 of those are shortlisted; overall 300 are shortlisted. $P(\text{short} | \text{exp}) = 120/200 = 0.6$; $P(\text{short}) = 0.3$. Experience doubles the chance.

Multiplication rule: $P(A \cap B) = P(A | B)P(B)$.

38.2.2 Independence

A and B are independent if $P(A \cap B) = P(A)P(B)$, i.e. $P(A | B) = P(A)$: knowing B tells nothing about A . Coin flips are independent; “knows Python” and “applies to data jobs” are not. **Conditional independence** $P(A \cap B | C) = P(A | C)P(B | C)$ is the Naive Bayes assumption (Chapter 10). Independent and mutually exclusive are different: exclusive events with positive probability are always *dependent* (if one happens, the other cannot).

38.2.3 Law of total probability

If $B_1, \dots, B_k$ split the sample space: $P(A) = \sum_i P(A | B_i)P(B_i)$.

Small example: total probability

40% of applications come from referrals with 30% shortlist rate, 60% from the portal with 10%: $P(\text{short}) = 0.4 \cdot 0.3 + 0.6 \cdot 0.1 = 0.18$.

38.2.4 Bayes' theorem

$$P(B | A) = \frac{P(A | B)P(B)}{P(A)}.$$

Small example: the rare-condition trap

1% of recruiter accounts are fraudulent. A detector flags 99% of frauds and 5% of genuine accounts. Given a flag, $P(\text{fraud} | \text{flag}) = \frac{0.99 \cdot 0.01}{0.99 \cdot 0.01 + 0.05 \cdot 0.99} = \frac{0.0099}{0.0594} = 0.167$. Only 1 in 6 flags is fraud: the base rate dominates. Continuing the referral example: $P(\text{referral} | \text{short}) = 0.12/0.18 = 0.667$.

In ML: Naive Bayes, MAP estimation, calibration and precision under imbalance (Chapter 11), Bayesian A/B testing, posterior updating of CTRs.

38.3 Random variables

A **random variable** assigns a number to each outcome: X = number of applications a job gets tomorrow; T = time until the next click.

- **Discrete:** countable values. Described by a **probability mass function (PMF)** $p(x) = P(X = x)$, with $\sum_x p(x) = 1$.
- **Continuous:** any value in an interval. Described by a **probability density function (PDF)** $f(x)$: probabilities are *areas* under it, $P(a \leq X \leq b) = \int_a^b f(x) dx$, and $\int f = 1$. $P(X = x) = 0$ for any exact value. A density can exceed 1 (Uniform(0, 0.5) has density 2).
- **Cumulative distribution function (CDF)** $F(x) = P(X \leq x)$, for both types; it rises from 0 to 1. Percentiles and medians are read from it.

Small example: PMF vs PDF

Die: $p(x) = 1/6$ for $x = 1..6$. Uniform(0, 10) waiting time: $f(x) = 0.1$ on $[0, 10]$; $P(2 < X < 5) = 3 \times 0.1 = 0.3$.

38.4 Expectation, variance and moments

38.4.1 Expectation (mean)

The long-run average: $\mathbb{E}[X] = \sum_x x p(x)$ (discrete) or $\int x f(x) dx$ (continuous).

Small example: expectations

Die: $\mathbb{E}[X] = (1 + \dots + 6)/6 = 3.5$. A telecall costs Rs. 20 and converts with probability 0.05 to a Rs. 1,000 commission: expected profit = $0.05 \cdot 1000 - 20 = \text{Rs. } 30$ per call.

Linearity: $\mathbb{E}[aX + bY + c] = a\mathbb{E}[X] + b\mathbb{E}[Y] + c$ always, even for dependent X, Y . Expected number of applies from 1,000 impressions each with probability p_i is $\sum p_i$.

38.4.2 Variance and standard deviation

$$\text{Var}(X) = \mathbb{E}[(X - \mu)^2] = \mathbb{E}[X^2] - (\mathbb{E}[X])^2.$$

Small example: variance of a die

$$\mathbb{E}[X^2] = (1 + 4 + 9 + 16 + 25 + 36)/6 = 91/6; \text{Var} = 91/6 - 3.5^2 = 15.17 - 12.25 = 2.917.$$

Rules: $\text{Var}(aX + b) = a^2 \text{Var}(X)$ (shifts do not change spread); $\text{Var}(X + Y) = \text{Var}(X) + \text{Var}(Y) + 2\text{Cov}(X, Y)$, which is just the sum when independent. Averaging n independent copies divides the variance by n : the root of the CLT and of bagging (Chapter 7).

38.4.3 Moments, skewness, kurtosis

The k -th **moment** is $\mathbb{E}[X^k]$; the k -th **central moment** is $\mathbb{E}[(X - \mu)^k]$. Mean = 1st moment; variance = 2nd central; **skewness** = $\mathbb{E}[(X - \mu)^3]/\sigma^3$ (asymmetry: positive, negative, zero); **kurtosis** = $\mathbb{E}[(X - \mu)^4]/\sigma^4$ (tails: mesokurtic 3, leptokurtic > 3 , platykurtic < 3 ; Chapter 31). The **moment generating function** $M(t) = \mathbb{E}[e^{tX}]$ packs all moments: its k -th derivative at 0 is $\mathbb{E}[X^k]$.

38.5 The common distributions

38.5.1 Bernoulli(p)

One yes/no trial: $P(1) = p, P(0) = 1 - p$. Mean p , variance $p(1 - p)$. *Real life*: did this seeker click this job. *ML*: label model of logistic regression (log loss is its negative log-likelihood); dropout masks.

38.5.2 Binomial(n, p)

Number of successes in n independent Bernoulli trials: $P(X = k) = \binom{n}{k} p^k (1 - p)^{n-k}$. Mean np , variance $np(1 - p)$.

Small example: binomial

A recruiter contacts 10 candidates, each responds with probability 0.3.
 $P(\text{exactly } 2) = \binom{10}{2} 0.3^2 0.7^8 = 45 \times 0.09 \times 0.0576 = 0.233$. Mean 3, variance 2.1.

$$P(8 \text{ or more heads in 10 fair tosses}) = (45 + 10 + 1)/1024 = 0.055.$$

ML: number of conversions in an A/B arm; binomial GLM for proportions.

38.5.3 Geometric(p)

Number of trials until the first success: $P(X = k) = (1 - p)^{k-1}p$, mean $1/p$; **memoryless**.

Small example: geometric

Each cold call converts with $p = 0.2$: $P(\text{first conversion on call 3}) = 0.8^2 \times 0.2 = 0.128$; expected calls until first conversion = 5.

38.5.4 Poisson(λ)

Count of events in a fixed interval when they occur independently at a constant average rate: $P(X = k) = e^{-\lambda} \lambda^k / k!$. Mean = variance = λ .

Small example: Poisson

A job receives on average $\lambda = 3$ applications per hour. $P(0) = e^{-3} = 0.050$; $P(2) = 0.224$; $P(\text{at least 2}) = 1 - P(0) - P(1) = 0.80$.

Real life: calls to a support centre per minute, website errors per day. *ML*: Poisson regression for counts (Chapter 4). Binomial with large n and small p is approximately Poisson with $\lambda = np$.

38.5.5 Uniform(a, b)

Every value in $[a, b]$ equally likely: density $1/(b - a)$, mean $(a + b)/2$, variance $(b - a)^2/12$. *ML*: random initialisation (Xavier uniform), random search of hyperparameters, random numbers for sampling.

38.5.6 Normal (Gaussian) $\mathcal{N}(\mu, \sigma^2)$

$f(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-(x-\mu)^2/(2\sigma^2)}$: the bell curve; symmetric; the 68–95–99.7 rule (Chapter 33). Arises when many small independent effects add up (CLT, Chapter 39).

Small example: normal

MBA scores $\mathcal{N}(60, 5^2)$: $P(X > 70) = P(Z > 2) = 0.023$; $P(55 < X < 62.5) = P(-1 < Z < 0.5) = 0.533$.

Real life: heights, measurement errors, averages. *ML*: squared-error loss (Gaussian noise MLE), Gaussian priors (Ridge), GMMs, Gaussian Naive Bayes, weight initialisation, VAE latent space, the noise in diffusion models.

38.5.7 Exponential(λ)

Waiting time until the next event in a Poisson process: $f(t) = \lambda e^{-\lambda t}$, $t \geq 0$; mean $1/\lambda$; $P(T > t) = e^{-\lambda t}$; **memoryless** ($P(T > s + t \mid T > s) = P(T > t)$).

Small example: exponential

Clicks arrive at 0.5 per minute. $P(\text{no click for 4 minutes}) = e^{-2} = 0.135$ (the same as Poisson $P(0)$ with $\lambda = 0.5 \times 4 = 2$). Mean wait 2 minutes.

Real life: time between arrivals, lifetimes of components. *ML:* survival/time-to-event modelling, queueing for serving systems.

38.5.8 Gamma(k, θ)

Sum of k independent exponential waiting times (for integer k); positive and right-skewed; mean $k\theta$, variance $k\theta^2$.

Small example: gamma

Time until the 3rd application when applications arrive every 2 hours on average: Gamma($k = 3, \theta = 2$): mean 6 hours, variance 12.

Real life: claim sizes, rainfall, time to complete multi-step tasks, salaries (roughly). *ML:* Gamma GLM (Chapter 4), conjugate prior for a Poisson rate (Chapter 12).

38.5.9 Beta(α, β)

A distribution over probabilities on $[0, 1]$; mean $\alpha/(\alpha + \beta)$.

Small example: beta

Beta(2, 8): mean 0.2, variance 0.0145: a belief that a CTR is around 20% with some uncertainty. After 3 clicks in 10 views, posterior Beta(5, 15), mean 0.25.

ML: conjugate prior for Bernoulli/binomial (CTR smoothing), Thompson sampling.

38.5.10 Others to recognise

Negative binomial (overdispersed counts), **log-normal** (log X normal: incomes, prices, session lengths), **chi-square** (sum of squared standard normals: variance tests, χ^2 test), **Student- t** (heavier tails than normal: small-sample means, t-SNE), **categorical/multinomial** (softmax outputs, word counts), **Dirichlet** (prior over probability vectors: topic models).

Distribution	Type	Mean	Variance	Typical example
Bernoulli(p)	discrete	p	$p(1-p)$	click / no click
Binomial(n, p)	discrete	np	$np(1-p)$	responses out of 10 contacts
Geometric(p)	discrete	$1/p$	$(1-p)/p^2$	calls until first conversion
Poisson(λ)	discrete	λ	λ	applications per hour
Uniform(a, b)	continuous	$\frac{a+b}{2}$	$\frac{(b-a)^2}{12}$	random initialisation
Normal(μ, σ^2)	continuous	μ	σ^2	test scores, errors
Exponential(λ)	continuous	$1/\lambda$	$1/\lambda^2$	time between clicks
Gamma(k, θ)	continuous	$k\theta$	$k\theta^2$	time to k -th event, claim sizes
Beta(α, β)	continuous on $[0, 1]$	$\frac{\alpha}{\alpha+\beta}$	$\frac{\alpha\beta}{(\alpha+\beta)^2(\alpha+\beta+1)}$	uncertainty about a CTR

** In practice: probability inside ML, in one list**

Loss functions are negative log-likelihoods (Bernoulli $\rightarrow$ log loss, Gaussian $\rightarrow$ MSE, Poisson $\rightarrow$ Poisson deviance); classifiers output probability distributions (softmax); generative models (GMM, VAE, LLMs) model $P(x)$; Bayes' theorem drives Naive Bayes and MAP; expectation underlies expected-value decisions (who to call); variance underlies bias–variance and confidence; sampling (bootstrap, dropout, SGD mini-batches, Thompson sampling) is probabilistic.

38.6 Interview questions

Q38.1

What is the difference between a PMF and a PDF? Can a PDF value exceed 1?

🕒 In one line. A PMF gives probabilities of exact values of a discrete variable; a PDF gives density for a continuous variable, whose areas (not heights) are probabilities, so a PDF can exceed 1 (Uniform(0, 0.5) has density 2) while still integrating to 1.

Q38.2

1% of accounts are fraudulent; a detector catches 99% of fraud and flags 5% of genuine accounts. What is the probability an account is fraudulent given it was flagged?

🕒 In one line. $0.0099 / (0.0099 + 0.0495) = 0.167$.

⚙️ The full answer. Bayes' theorem; the low base rate means most flags are false alarms. Precision depends on prevalence, sensitivity and specificity together.

Q38.3

What is the difference between independent and mutually exclusive events?

🕒 In one line. Independent events do not affect each other's probability ($P(A \cap B) = P(A)P(B)$); mutually exclusive events cannot happen together ($P(A \cap B) = 0$); two exclusive events with positive probabilities are therefore dependent.

Q38.4

Define expectation and variance; compute them for a fair die.

🕒 In one line. Expectation is the probability-weighted average $\sum xp(x)$; variance is $\mathbb{E}[X^2] - \mathbb{E}[X]^2$; for a die, 3.5 and $91/6 - 12.25 = 2.917$.

Q38.5

A recruiter contacts 10 candidates; each responds with probability 0.3 independently. Probability exactly 2 respond? Expected number?

🕒 **In one line.** $\binom{10}{2} 0.3^2 0.7^8 = 0.233$; expected $np = 3$.

Q38.6

A job gets 3 applications per hour on average. Probability of none in an hour? At least 2?

🕒 **In one line.** Poisson: $e^{-3} = 0.050$; $1 - e^{-3}(1 + 3) = 0.80$.

Q38.7

Clicks arrive at 0.5 per minute. Probability of waiting more than 4 minutes for the next click? What is “memoryless”?

🕒 **In one line.** $e^{-0.5 \times 4} = 0.135$; memoryless means that having already waited s minutes does not change the distribution of the remaining wait.

Q38.8

Give a real-life example and an ML use for each: Bernoulli, binomial, Poisson, normal, exponential, gamma, beta.

🕒 **In one line.** Bernoulli: a click (logistic regression’s label model); binomial: conversions in an A/B arm (tests); Poisson: applications per hour (Poisson regression); normal: measurement errors (MSE, Gaussian priors, GMMs); exponential: time between clicks (survival models); gamma: claim sizes or salaries (Gamma GLM, Poisson-rate prior); beta: uncertainty about a CTR (smoothing, Thompson sampling).

Q38.9

State linearity of expectation and use it: 1,000 impressions with click probabilities p_i . Expected clicks?

🕒 **In one line.** $\mathbb{E}[\sum X_i] = \sum \mathbb{E}[X_i] = \sum p_i$, regardless of dependence between impressions.

Q38.10

🧠 **Think deeper** Why is the variance of an average of n independent variables σ^2/n , and where does ML rely on this?

🕒 **In one line.** $\text{Var}(\frac{1}{n} \sum X_i) = \frac{1}{n^2} \sum \text{Var}(X_i) = \frac{\sigma^2}{n}$ for independent variables; it underlies bagging and random forests, mini-batch gradient noise, the standard error of metrics, and A/B test sample sizes.

Q38.11

What are moments, and how do skewness and kurtosis relate to them?

🎯 **In one line.** Moments are expectations of powers; the mean is the first moment, variance the second central moment, skewness the third standardised moment (asymmetry) and kurtosis the fourth (tail heaviness).

Q38.12

🧠 **Think deeper** When is a binomial well approximated by a Poisson, and by a normal?

🎯 **In one line.** By a Poisson with $\lambda = np$ when n is large and p small (rare events); by a normal with mean np and variance $np(1 - p)$ when np and $n(1 - p)$ are both reasonably large (say ≥ 10).

Q38.13

A Beta(2, 8) prior on a CTR; you observe 3 clicks in 10 views. Posterior and its mean?

🎯 **In one line.** $\text{Beta}(2 + 3, 8 + 7) = \text{Beta}(5, 15)$, mean 0.25.

Q38.14

Compute $P(X > 70)$ and $P(55 < X < 62.5)$ for $X \sim \mathcal{N}(60, 25)$.

🎯 **In one line.** $P(Z > 2) = 0.023$; $P(-1 < Z < 0.5) = 0.533$.

CHAPTER 39

Statistics for Data Scientists: Sampling, CLT, Confidence Intervals and Hypothesis Testing

🎥 **Companion video for this chapter:** [StatQuest: p-values, what they are and how to interpret them](#)

The checklist recommends the Rice University statistics courses and lists descriptive vs inferential statistics, population vs sample, the central limit theorem, p-values and hypothesis testing. The reported Round 1 question was exactly “what is a p-value, what is alpha, how do you compare them”. This chapter teaches inference from zero with worked numbers, and then shows where it drives ML decisions (model comparison, A/B tests, drift).

39.1 Descriptive vs inferential statistics

- **Descriptive statistics** summarise the data you have: mean, median, standard deviation, histograms (Chapter 31). “The 200 applicants to this job had a median experience of 3 years.”
- **Inferential statistics** use a sample to draw conclusions about a larger population, with quantified uncertainty: estimates, confidence intervals, hypothesis tests. “Seekers who receive the new alert apply 12% more (95% CI: 5% to 19%).”

39.2 Population, sample, parameter, statistic

- **Population:** everyone or everything you care about (all Naukri seekers in India). Its numerical characteristics are **parameters** (μ, σ, p): fixed but usually unknown.
- **Sample:** the subset you observe (10,000 surveyed seekers). Its numbers are **statistics** ($\bar{x}, s, \hat{p}$): known but random (a different sample gives different values).

Sampling methods: simple random; **stratified** (sample within groups such as cities, then combine: guarantees coverage of small groups, as in stratified CV); cluster (sample whole groups such as colleges); systematic (every k -th record). **Sampling bias:** the sample differs systematically from the population (survey only active users; train only on users who were called by telecallers). No sample size fixes bias.

39.3 Sampling distributions and the standard error

If we repeated the sampling many times, the statistic (say $\bar{x}$) would vary. Its distribution is the **sampling distribution**; its standard deviation is the **standard error** (SE). For a mean of n independent observations:

$$SE(\bar{x}) = \frac{\sigma}{\sqrt{n}} \quad (\text{estimated by } s/\sqrt{n}).$$

For a proportion: $SE(\hat{p}) = \sqrt{p(1-p)/n}$. Four times the data halves the error.

Small example: standard error

Test accuracy 0.80 on 400 examples: $SE = \sqrt{0.8 \cdot 0.2/400} = 0.02$. On 43 examples (Chapter 27): $SE = 0.061$: the estimate could easily be off by 10 points.

39.4 The central limit theorem

Statement: for independent observations from *any* distribution with finite mean μ and variance σ^2 , the sample mean is approximately normal for large n :

$$\bar{X} \approx \mathcal{N}\left(\mu, \frac{\sigma^2}{n}\right),$$

however skewed the original distribution is. “Large” is often $n \geq 30$; more for very skewed data.

Small example: CLT with dice

One die is uniform on 1–6 (mean 3.5, variance 2.917). The average of 30 rolls is approximately normal with mean 3.5 and $SE \sqrt{2.917/30} = 0.31$. So about 95% of such averages fall within 3.5 ± 0.61 , i.e. 2.89 to 4.11, even though a single roll is anything but normal.

Why it matters: it justifies normal-based confidence intervals and z/t-tests for means and proportions (A/B tests on clicks), explains why averages are stable, and why sums of many small effects look Gaussian. Related: the **law of large numbers** (the sample mean converges to μ); the CLT adds *how fast* and *in what shape*.

39.5 Confidence intervals

A 95% confidence interval is a range computed from the sample by a procedure that captures the true parameter in 95% of repeated samples:

$$\bar{x} \pm z^* \frac{s}{\sqrt{n}} \quad (z^* = 1.96 \text{ for } 95\%), \quad \text{or } \bar{x} \pm t_{n-1}^* \frac{s}{\sqrt{n}} \text{ for small } n.$$

Small example: CI for a mean

Time-to-hire for 36 jobs: $\bar{x} = 12$ days, $s = 3$. $SE = 0.5$. 95% CI = $12 \pm 1.96 \times 0.5 = (11.02, 12.98)$ (with $t_{35}^* = 2.03$: 12 ± 1.02).

Correct interpretation: the method has 95% coverage; it is *not* “95% probability that μ is in this particular interval” (in frequentist terms μ is fixed; a Bayesian credible interval does have that reading). Wider intervals for higher confidence or smaller samples.

39.6 Hypothesis testing, step by step

1. **Null hypothesis** H_0 : the default, “no effect” claim (the new ranking does not change apply rate; the coin is fair).
2. **Alternative** H_1 : what we want evidence for (apply rate changed: two-sided; increased: one-sided).

3. **Significance level α** : the false-positive rate we accept, chosen before seeing data (0.05).
4. **Test statistic**: how far the data are from H_0 in standard-error units (z , t , χ^2 , F).
5. **p-value**: the probability, *assuming H_0 is true*, of getting a result at least as extreme as the one observed.
6. **Decision**: if $p \leq \alpha$, reject H_0 ("statistically significant"); otherwise fail to reject (not the same as proving H_0).

Small example: a z-test

Average applications per job are claimed to be 10 (with known $\sigma = 2$). A sample of 64 jobs averages 10.5. $z = \frac{10.5-10}{2/\sqrt{64}} = \frac{0.5}{0.25} = 2$. Two-sided p-value = $P(|Z| \geq 2) = 0.046 < 0.05$: reject H_0 at the 5% level (but not at 1%).

39.6.1 What a p-value is not

Not the probability that H_0 is true; not the probability the result is due to chance; not a measure of effect size or importance. A tiny, useless effect can be highly significant with huge n ; a large, important effect can be non-significant with small n . Always report an effect size and a confidence interval.

39.6.2 Errors and power

	H_0 true	H_0 false
Reject H_0	Type I error (false positive), prob. α	correct (power = $1 - \beta$)
Fail to reject	correct	Type II error (false negative), prob. β

Power rises with sample size, effect size and α , and falls with noise. Plan sample size before an experiment (Chapter 28: about $16\sigma^2/\delta^2$ per arm for 80% power at $\alpha = 0.05$).

39.7 The common tests

- **One-sample z/t-test**: is a mean equal to a value? (t when σ is estimated, $n - 1$ degrees of freedom.)
- **Two-sample t-test**: do two independent groups have equal means? (Welch's version for unequal variances.)

Small example: two-sample t

Interview scores, group A (12, 14, 11, 13, 15) mean 13; group B (10, 11, 9, 12, 10) mean 10.4. $t = 2.98$ with 8 df, $p = 0.018$: significant difference.

- **Paired t-test** (the syllabus's "paired means test"): the same units measured twice (before/after a feature, model A vs B on the same folds); test whether the mean *difference* is 0.

Small example: paired t

Five recruiters' time-to-shortlist before and after a tool; differences 2, 1, 3, 0, 4 hours. Mean 2, sd 1.58, SE 0.707, $t = 2.83$ (4 df), $p = 0.047$. Pairing removes between-recruiter variation, making the test more powerful.

- **Two-proportion z-test:** do two conversion rates differ? (A/B tests on clicks.)

Small example: an A/B test

Control: 200 applies out of 4,000 (5.0%); variant: 260 of 4,000 (6.5%). Pooled $\hat{p} = 460/8000 = 0.0575$; $SE = \sqrt{0.0575 \cdot 0.9425 \cdot (2/4000)} = 0.0052$; $z = 0.015/0.0052 = 2.88$; $p = 0.004$. The lift is significant; 95% CI for the difference $\approx 1.5 \pm 1.0$ points.

- **Chi-square test of independence:** are two categorical variables related? (Work experience vs placement: $\chi^2 = 16.7$, $p < 0.001$; Chapter 32.)
- **ANOVA:** equal means across several groups (Chapter 37).
- **Non-parametric:** Mann–Whitney U (two groups, ranks), Wilcoxon signed-rank (paired), Kruskal–Wallis (several groups), Kolmogorov–Smirnov (are two distributions the same: drift detection).

39.8 Multiple testing

Run 20 independent tests at $\alpha = 0.05$ with no real effects: $P(\text{at least one false positive}) = 1 - 0.95^{20} = 0.64$. Corrections: **Bonferroni** (use α/m : $0.05/10 = 0.005$ for 10 tests), Holm, and false discovery rate control (Benjamini–Hochberg). Relevant when checking many metrics, segments or features.

In practice: where statistics drives ML decisions

- **Model comparison:** a paired test (or bootstrap CI) on per-fold or per-example scores before claiming model B beats model A.
- **A/B tests:** two-proportion tests, sample-size planning, guardrails, multiple-testing control; this is how every InfoEdge model launch is judged.
- **Uncertainty in metrics:** standard errors of accuracy/AUC on small test sets; bootstrap confidence intervals.
- **Drift detection:** KS or chi-square tests comparing feature distributions between training and live data.
- **Feature significance** in linear/logistic regression (coefficient t-tests), with heteroscedasticity-robust SEs (Chapter 2).

39.9 Interview questions

Q39.1

★ **PYQ (reported Round 1: p-value vs alpha)** What is a p-value, what is alpha, and how do you compare them?

🎯 **In one line.** The p-value is the probability, assuming the null hypothesis is true, of seeing data at least as extreme as observed; alpha is the false-positive rate we decide in advance to tolerate; if $p \leq \alpha$ we reject the null, otherwise we fail to reject it.

⚙️ **The full answer.** Alpha is a property of the decision rule (fixed before the data); the p-value is a property of the observed data. Example: $z = 2$, $p = 0.046$: significant at $\alpha = 0.05$, not at 0.01.

🗣️ **If the interviewer pushes.** “Is $p = 0.03$ the probability that the null is true?” No. “What if $p = 0.06$?” Not significant at 5%; report the effect size and CI, and consider power.

Q39.2

Explain descriptive vs inferential statistics, and population vs sample, with examples.

🕒 **In one line.** Descriptive statistics summarise observed data; inferential statistics use a sample to make uncertain statements about a population; a population is the full group of interest (all seekers), a sample is the observed subset (10,000 surveyed seekers), whose statistics estimate population parameters.

Q39.3

State the central limit theorem and say why it matters for data science.

🕒 **In one line.** For independent observations with finite variance, the sample mean is approximately normal with mean μ and variance σ^2/n for large n , whatever the original distribution; it justifies normal-based confidence intervals and tests for means and proportions (A/B tests), and explains why averages are stable.

🧮 **Worked example.** Average of 30 die rolls: mean 3.5, SE 0.31.

Q39.4

What is a standard error and how does it differ from a standard deviation?

🕒 **In one line.** The standard deviation measures the spread of individual observations; the standard error measures the spread of an estimate (like a sample mean) across repeated samples, $\sigma/\sqrt{n}$, and shrinks as n grows.

Q39.5

Compute a 95% CI for mean time-to-hire with $n = 36$, $\bar{x} = 12$, $s = 3$, and interpret it correctly.

🕒 **In one line.** $12 \pm 1.96 \times 0.5 = (11.02, 12.98)$ days; the procedure captures the true mean in 95% of repeated samples (not a 95% probability statement about this fixed interval).

Q39.6

Type I vs Type II error, and power. How do you increase power?

🕒 **In one line.** Type I: rejecting a true null (false positive, rate α); Type II: failing to reject a false null (false negative, rate β); power = $1 - \beta$; increase it with more data, larger true effects, less noise (paired designs, variance reduction), or a higher α .

Q39.7

★ **PYQ (syllabus: paired means test)** When do you use a paired t-test? Compute it for differences 2, 1, 3, 0, 4.

🎯 **In one line.** When the same units are measured twice (before/after, model A vs B on the same data); test whether the mean difference is zero: mean 2, SE 0.707, $t = 2.83$ with 4 df, $p = 0.047$.

⚙️ **The full answer.** Pairing removes unit-to-unit variation, so the test is more powerful than an independent two-sample test on the same numbers.

Q39.8

Analyse an A/B test: control 200/4000 applies, variant 260/4000. Is the lift significant?

🎯 **In one line.** Two-proportion z-test: pooled rate 0.0575, SE 0.0052, $z = 2.88$, $p = 0.004$: yes, the 1.5-point lift (30% relative) is significant at 5%.

💡 **If the interviewer pushes.** “What else would you check?” Sample ratio mismatch, guardrail metrics, novelty effects, pre-registered metric, segment heterogeneity, practical significance.

Q39.9

🧠 **Think deeper** With 50 lakh users in an A/B test, almost everything is significant. How do you decide what matters?

🎯 **In one line.** Focus on effect sizes and confidence intervals relative to a pre-defined minimum practically important effect, cost and risk; statistical significance with huge n only says the effect is not exactly zero.

Q39.10

What is the multiple testing problem and how do you correct for it?

🎯 **In one line.** Testing many hypotheses inflates the chance of at least one false positive ($1 - 0.95^{20} = 0.64$ for 20 tests); correct with Bonferroni (α/m), Holm, or Benjamini–Hochberg FDR control.

Q39.11

🧠 **Think deeper** How would you test whether model B is really better than model A?

🎯 **In one line.** Evaluate both on the same folds or examples and use a paired test (paired t-test on fold scores, McNemar’s test on per-example correctness, or a bootstrap CI of the metric difference), then confirm online with an A/B test.

Q39.12

Which test would you use: (a) apply rate of two email templates, (b) mean salary in three cities, (c) is city related to job function, (d) has the distribution of experience changed since training?

Ⓒ **In one line.** (a) two-proportion z-test (or chi-square); (b) one-way ANOVA (Kruskal–Wallis if skewed); (c) chi-square test of independence; (d) Kolmogorov–Smirnov test (or PSI) for drift.

Q39.13

What is sampling bias? Give an ML example.

Ⓒ **In one line.** A systematic difference between the sample and the population; e.g. training a lead-conversion model only on users that telecallers chose to call, so the model never learns about the users they skipped.

CHAPTER 40

Linear Algebra for ML: Vectors, Matrices, Rank, Eigenvectors and Markov Chains

🎥 **Companion video for this chapter:** [3Blue1Brown: Eigenvectors and eigenvalues \(Essence of linear algebra\)](#)

The InfoEdge hire called linear algebra “the hardest part”: rank, eigenvalues and eigenvectors, null space (the one he could not answer) and a Markov chain question. The checklist repeats eigenvectors, eigenvalues and rank. This chapter teaches them from zero, with many small computations, and shows exactly where each appears in ML.

40.1 Vectors

A **vector** is an ordered list of numbers: a point or an arrow. A candidate described by (experience, skills matched, expected salary) = (5, 8, 12) is a vector in 3-D; a word embedding is a vector in 300-D.

- **Addition and scaling:** $(1, 2) + (3, 1) = (4, 3)$; $2(1, 2) = (2, 4)$.
- **Dot product:** $\mathbf{a} \cdot \mathbf{b} = \sum a_i b_i$. $(3, 4) \cdot (4, 3) = 12 + 12 = 24$.
- **Norm (length):** $\|\mathbf{a}\| = \sqrt{\mathbf{a} \cdot \mathbf{a}}$; $\|(3, 4)\| = 5$. L1 norm $\sum |a_i| = 7$.
- **Angle:** $\cos \theta = \frac{\mathbf{a} \cdot \mathbf{b}}{\|\mathbf{a}\| \|\mathbf{b}\|} = \frac{24}{25} = 0.96$: **cosine similarity** (search, embeddings). Dot product 0 means perpendicular (**orthogonal**): no overlap.
- **Projection of $\mathbf{a}$ onto $\mathbf{b}$:** $\frac{\mathbf{a} \cdot \mathbf{b}}{\mathbf{b} \cdot \mathbf{b}} \mathbf{b} = \frac{24}{25}(4, 3) = (3.84, 2.88)$: the shadow of $\mathbf{a}$ on $\mathbf{b}$'s line (the core of least squares and PCA).

40.2 Matrices as transformations

A matrix is a table of numbers, but more usefully a *function* that transforms vectors: $A\mathbf{x}$ is a linear combination of A 's columns weighted by $\mathbf{x}$.

📊 **Small example: what a matrix does**

$A = \begin{pmatrix} 2 & 0 \\ 0 & 0.5 \end{pmatrix}$ stretches the x -axis by 2 and squashes the y -axis by half: $(1, 1) \mapsto (2, 0.5)$.
 $R = \begin{pmatrix} 0 & -1 \\ 1 & 0 \end{pmatrix}$ rotates by 90° : $(1, 0) \mapsto (0, 1)$.

Matrix multiplication AB means “apply B , then A ”; entry (i, j) is row i of A dot column j of B ; sizes $(m \times n)(n \times p) \rightarrow m \times p$; not commutative. In a neural network, a dense layer is $W\mathbf{x} + \mathbf{b}$; a batch is XW .

Transpose $A^\top$ swaps rows and columns. **Identity** I leaves vectors unchanged. **Inverse** A^{-1} undoes A : $AA^{-1} = I$. For 2×2 :

$$\begin{pmatrix} a & b \\ c & d \end{pmatrix}^{-1} = \frac{1}{ad - bc} \begin{pmatrix} d & -b \\ -c & a \end{pmatrix}.$$

📺 **Stuck on this page?** Watch: [3Blue1Brown: Eigenvectors and eigenvalues \(Essence of linear algebra\)](#)

Small example: inverse and determinant

$$\begin{pmatrix} 3 & 1 \\ 2 & 4 \end{pmatrix}: \text{determinant } 12 - 2 = 10; \text{inverse } \frac{1}{10} \begin{pmatrix} 4 & -1 \\ -2 & 3 \end{pmatrix} = \begin{pmatrix} 0.4 & -0.1 \\ -0.2 & 0.3 \end{pmatrix}.$$

The **determinant** is the factor by which the transformation scales areas (volumes); determinant 0 means space is flattened onto a line or plane: information is lost and no inverse exists.

40.3 Systems of equations

$A\mathbf{x} = \mathbf{b}$. $2x + y = 5, x + 3y = 10$: $\mathbf{x} = A^{-1}\mathbf{b} = (1, 3)$. Linear regression's normal equations $X^T X \mathbf{w} = X^T \mathbf{y}$ are such a system (Chapter 2).

40.4 Linear independence, span, rank and null space

40.4.1 Linear independence

Vectors are **linearly independent** if none is a combination of the others. $(1, 2)$ and $(2, 4)$ are dependent ($2 \times$ the first); $(1, 0)$ and $(0, 1)$ are independent. The **span** of vectors is everything reachable by combining them.

40.4.2 Rank

The **rank** of a matrix is the number of linearly independent columns (equivalently rows): the dimension of its output space (column space).

Small example: rank by inspection

$$B = \begin{pmatrix} 1 & 2 & 3 \\ 2 & 4 & 6 \\ 1 & 0 & 1 \end{pmatrix}: \text{row 2} = 2 \times \text{row 1, so at most 2 independent rows; rows 1 and 3 are independent.}$$

Rank 2, determinant 0 (not invertible). In general, rank via row reduction (count non-zero rows of the echelon form) or the number of non-zero singular values (Chapter 14).

40.4.3 Null space

The **null space** of A is the set of vectors $\mathbf{x}$ with $A\mathbf{x} = \mathbf{0}$: the directions the transformation destroys.

Small example: null space

For B above, solve $B\mathbf{x} = \mathbf{0}$: from row 3, $x_1 + x_3 = 0$; from row 1, $x_1 + 2x_2 + 3x_3 = 0$. Setting $x_3 = -1$: $x_1 = 1$, then $1 + 2x_2 - 3 = 0$, $x_2 = 1$. Null space = all multiples of $(1, 1, -1)$. Check: $(1 + 2 - 3, 2 + 4 - 6, 1 + 0 - 1) = (0, 0, 0)$.

Rank-nullity theorem: rank + dimension of null space = number of columns ($2 + 1 = 3$). In ML: if the feature matrix X has a non-trivial null space (perfect multicollinearity, the dummy variable trap), adding any null-space vector to the weights leaves all predictions unchanged, so the coefficients are not unique (Chapters 2, 34). LoRA assumes the weight *update* has low rank (Chapter 23).

40.5 Eigenvalues and eigenvectors

40.5.1 Meaning

Most vectors change direction when a matrix acts on them. An **eigenvector** is a special direction that is only stretched (or flipped): $A\mathbf{v} = \lambda\mathbf{v}$; the stretch factor λ is its **eigenvalue**.

40.5.2 How to compute (2×2)

$A\mathbf{v} = \lambda\mathbf{v} \iff (A - \lambda I)\mathbf{v} = 0$ has a non-zero solution $\iff \det(A - \lambda I) = 0$ (the **characteristic equation**). Solve for λ , then for each λ find $\mathbf{v}$ in the null space of $A - \lambda I$.

Small example: eigenvalues and eigenvectors

$A = \begin{pmatrix} 4 & 1 \\ 2 & 3 \end{pmatrix}$. $\det \begin{pmatrix} 4-\lambda & 1 \\ 2 & 3-\lambda \end{pmatrix} = (4-\lambda)(3-\lambda) - 2 = \lambda^2 - 7\lambda + 10 = 0$, so $\lambda = 5$ or 2 . $\lambda = 5$: $\begin{pmatrix} -1 & 1 \\ 2 & -2 \end{pmatrix} \mathbf{v} = 0 \Rightarrow \mathbf{v} = (1, 1)$. Check: $A(1, 1) = (5, 5) = 5(1, 1)$. $\lambda = 2$: $\begin{pmatrix} 2 & 1 \\ 2 & 1 \end{pmatrix} \mathbf{v} = 0 \Rightarrow \mathbf{v} = (1, -2)$. Check: $A(1, -2) = (2, -4) = 2(1, -2)$.

Shortcuts: sum of eigenvalues = trace (sum of the diagonal: $4 + 3 = 7$); product = determinant ($12 - 2 = 10$). Diagonal or triangular matrices have their diagonal entries as eigenvalues (diag(2, 0.5): 2 and 0.5). A matrix is invertible iff no eigenvalue is 0.

40.5.3 Diagonalisation and symmetric matrices

If A has n independent eigenvectors, $A = PDP^{-1}$ (P : eigenvectors as columns; D : eigenvalues on the diagonal). Then $A^k = PD^kP^{-1}$: powers are easy, and the largest $|\lambda|$ dominates long-run behaviour. **Symmetric matrices** (covariance matrices, $X^T X$, kernel matrices) are especially nice: real eigenvalues and *orthogonal* eigenvectors, $A = Q\Lambda Q^T$ (the spectral theorem). If all eigenvalues are positive the matrix is **positive definite** ($\mathbf{x}^T A \mathbf{x} > 0$ for all $\mathbf{x} \neq 0$): a bowl-shaped quadratic, a valid covariance, a convex loss Hessian.

Small example: the PCA matrix

Covariance $\begin{pmatrix} 2 & 1 \\ 1 & 2 \end{pmatrix}$: eigenvalues 3 and 1 (trace 4, determinant 3), eigenvectors $(1, 1)/\sqrt{2}$ and $(1, -1)/\sqrt{2}$, orthogonal. PC1 explains 75% of variance (Chapter 14).

40.6 Markov chains: eigenvectors in action

A **Markov chain** moves between states with probabilities that depend only on the current state (no memory of the path). It is described by a **transition matrix** P where $P_{ij} = P(\text{next} = j \mid \text{now} = i)$; each row sums to 1. Think of it as a finite state machine with probabilities on its arrows (as the InfoEdge candidate did).

Small example: an active/dormant seeker chain

States: Active (A), Dormant (D). Each week an active seeker stays active with 0.9 and goes dormant with 0.1; a dormant seeker reactivates with 0.5 and stays dormant with 0.5: $P = \begin{pmatrix} 0.9 & 0.1 \\ 0.5 & 0.5 \end{pmatrix}$. Start active, $\pi_0 = (1, 0)$. After one week $\pi_1 = \pi_0 P = (0.9, 0.1)$; two weeks $(0.86, 0.14)$; three weeks $(0.844, 0.156)$. **Steady state:** $\pi = \pi P$ with $\pi_A + \pi_D = 1$. From $\pi_D = 0.1\pi_A + 0.5\pi_D$: $0.5\pi_D = 0.1\pi_A$, $\pi_A = 5\pi_D$, so $\pi = (5/6, 1/6) = (0.833, 0.167)$. In the long run 83% of seekers are active, whatever the start. The steady state is the eigenvector of P^T with eigenvalue 1.

Two-step probabilities are P^2 (Chapman–Kolmogorov). Chains with every state reachable and no periodic cycling converge to a unique steady state. **In ML:** PageRank (steady state of a random surfer; power iteration $\pi \leftarrow \pi P$), n-gram language models (next word depends on the last word), MDPs in reinforcement learning (Chapter 36), customer journey and churn modelling, MCMC sampling for Bayesian inference, hidden Markov models.

40.7 SVD in one paragraph (recap)

Every matrix factors as $A = U\Sigma V^T$: rotate, stretch along axes by the singular values, rotate (Chapter 14). Singular values are the square roots of the eigenvalues of $A^T A$; their count of non-zeros is the rank; truncating gives the best low-rank approximation (recommenders, LSA, compression, PCA).

In practice: where linear algebra lives in ML

Concept	ML use
Dot product, cosine	similarity search, attention scores, linear models
Matrix multiplication	every neural network layer, batched computation on GPUs
Inverse, solving systems	normal equations, Gaussian processes, Newton's method
Rank, null space	multicollinearity, dummy trap, low-rank models (LoRA, matrix factorisation)
Eigen-decomposition	PCA, spectral clustering, LDA (generalised eigenproblem), stability of RNNs (largest eigenvalue of W_{hh})
Positive definiteness	valid covariance/kernel matrices (Mercer), convexity of losses
SVD	PCA, LSA, recommenders, pseudo-inverse, compression
Markov chains	PageRank, language models, RL, MCMC

40.8 Interview questions

Q40.1

★ **PYQ (reported interview: rank)** What is the rank of a matrix? Find the rank of $\begin{pmatrix} 1 & 2 & 3 \\ 2 & 4 & 6 \\ 1 & 0 & 1 \end{pmatrix}$.

🕒 **In one line.** The number of linearly independent rows or columns (the dimension of the column space); here 2, because row 2 is twice row 1 while rows 1 and 3 are independent.

💡 **If the interviewer pushes.** “Why does rank matter in ML?” A feature matrix with rank below its number of columns has redundant features (perfect multicollinearity), making $X^T X$ singular.

Q40.2

★ **PYQ (reported interview: null space)** What is the null space of a matrix? Find it for the matrix above.

🕒 **In one line.** The set of all vectors the matrix maps to zero; here all multiples of $(1, 1, -1)$; by rank-nullity its dimension is $3 - 2 = 1$.

⚙️ **The full answer.** In regression, any null-space direction can be added to the weight vector without changing predictions, so coefficients are not identifiable; regularisation or dropping features fixes it.

Q40.3

★ **PYQ (reported interview: eigenvalues and eigenvectors)** Define eigenvalues and eigenvectors and compute them for $\begin{pmatrix} 4 & 1 \\ 2 & 3 \end{pmatrix}$.

🕒 **In one line.** Eigenvectors are directions a matrix only stretches, $A\mathbf{v} = \lambda\mathbf{v}$, with eigenvalue λ the stretch factor; here $\lambda = 5$ with $\mathbf{v} = (1, 1)$ and $\lambda = 2$ with $\mathbf{v} = (1, -2)$.

⚙️ **The full answer.** Characteristic equation $\lambda^2 - 7\lambda + 10 = 0$. Checks: trace $7 = 5 + 2$, determinant $10 = 5 \times 2$.

Q40.4

★ **PYQ (reported interview: Markov chain)** Seekers move weekly between Active and Dormant with $P = \begin{bmatrix} 0.9 & 0.1 \\ 0.5 & 0.5 \end{bmatrix}$. Where is an active seeker after two weeks, and what is the long-run split?

🕒 **In one line.** After two weeks $(0.86, 0.14)$; the steady state solves $\pi = \pi P$: $(5/6, 1/6) \approx (0.833, 0.167)$, the eigenvector of $P^\top$ for eigenvalue 1.

⚙️ **The full answer.** Model it as a finite state machine with probabilistic transitions (Markov property: the next state depends only on the current one). Multiply the state vector by P each step.

💡 **If the interviewer pushes.** “How would you estimate P from data?” Count observed transitions from each state and divide by the row totals (the MLE).

Q40.5

How are eigenvectors used in PCA?

🕒 **In one line.** The principal components are the eigenvectors of the data’s covariance matrix, ordered by eigenvalue, and each eigenvalue is the variance along its component; because the covariance is symmetric, the components are orthogonal.

Q40.6

What does a zero determinant mean geometrically and for ML?

🕒 **In one line.** The transformation flattens space into a lower dimension, so it is not invertible (some information is lost); in ML it signals singular $X^T X$ (redundant features) or a degenerate covariance (a GMM component collapsing).

Q40.7

What are the trace and determinant in terms of eigenvalues?

🕒 **In one line.** The trace equals the sum of the eigenvalues and the determinant equals their product.

Q40.8

🧠 **Think deeper** Why do symmetric matrices matter so much in ML?

🕒 **In one line.** Covariance, $X^T X$, kernel and Hessian matrices are symmetric, and symmetric matrices have real eigenvalues and orthogonal eigenvectors, so they can be decomposed as $Q\Lambda Q^T$ (PCA, whitening, convexity checks via positive eigenvalues).

Q40.9

🧠 **Think deeper** How does the largest eigenvalue of an RNN's recurrent matrix relate to vanishing and exploding gradients?

🕒 **In one line.** Backpropagation through time multiplies by W_{hh} (times activation derivatives) once per step; if the largest eigenvalue (more precisely singular value) is below 1 the product shrinks exponentially (vanishing), above 1 it grows (exploding).

Q40.10

Compute the cosine similarity of (3, 4) and (4, 3) and the projection of the first onto the second.

🕒 **In one line.** Cosine $24/25 = 0.96$; projection $(24/25)(4, 3) = (3.84, 2.88)$.

Q40.11

What is a positive definite matrix and why does it matter?

🕒 **In one line.** A symmetric matrix with $x^T A x > 0$ for every non-zero x (all eigenvalues positive); valid covariance matrices are positive semi-definite, Mercer kernels must produce positive semi-definite Gram matrices, and a positive definite Hessian means a strict local minimum (convex bowl).

Q40.12

Solve $2x + y = 5$, $x + 3y = 10$ using a matrix inverse.

⦿ **In one line.** Determinant 5; $\mathbf{x} = \frac{1}{5} \begin{pmatrix} 3 & -1 \\ -1 & 2 \end{pmatrix} \begin{pmatrix} 5 \\ 10 \end{pmatrix} = (1, 3)$.

Part VIII

EDA Lab: Hands-on Exploratory Data Analysis on Real Datasets (with a Colab Notebook)

CHAPTER 41

EDA Lab I: The EDA Playbook in Code

🎥 **Companion video for this chapter:** [Krish Naik: Exploratory Data Analysis with pandas \(search\)](#)

EDA carries heavy weight in the InfoEdge interviews: Round 3 is literally “clean the data, analyse it, model it, test it”, and the successful candidate’s first moves (spotting salary leakage, checking ranges before calling anything an outlier) were EDA judgements. This part of the book is hands-on. It comes with a **Colab-ready notebook** that runs every analysis shown here on real data:

📄 **Notebook:** InfoEdge_Interview_Companion/eda_lab/EDA_Lab.ipynb in your repository (branch `claudio/admiring-edison-vd2vbx`). Open it in Google Colab (File → Open notebook → GitHub, or upload the file) and run all cells. The real datasets (HR job change, Titanic, Pima diabetes) are read directly from your public GitHub repository. For the MBA placement case, upload Kaggle’s `Placement_Data_Full_Class.csv` to the Colab session; without it the notebook builds a simulated replica with the same columns.

Chapter 41 gives the method and a cookbook of pandas/NumPy recipes, each run on the real HR job-change dataset (19,158 data-science candidates; target = 1 if they are looking for a job change: very close to Naukri data). Chapters 42–44 are complete case studies.

41.1 The 12-step EDA workflow

Say the step, run the code, read the output, state the decision. Every step ends in “because I see X, I will do Y”.

Step	Question it answers
1 Frame	What is one row? What is the target? When is a prediction made?
2 Load and look	Shape, head, types, memory; does it look like the description?
3 Types and parsing	Numbers stored as text? Codes like “>20”, “never”? Dates?
4 Data quality	Duplicates, impossible values, malformed labels, constant columns, disguised missing values (0, -1, 999)
5 Missing values	How much, where, together with what, and does missingness relate to the target?
6 Target	Balance; base rate; metric implications
7 Univariate numeric	Centre, spread, skew, outliers (genuine or error?)
8 Univariate categorical	Levels, rare levels, cardinality
9 Bivariate with target	Which features separate the classes, linearly or not?
10 Feature–feature	Correlations, redundancy, multicollinearity
11 Multivariate	Interactions and segments (does an effect depend on another variable?)
12 Leakage and summary	Anything “too good”? Write the findings and the modelling plan

41.2 The reusable toolkit

Define these once at the top of any interview notebook (they are cell 3 of `EDA_Lab.ipynb`).

```
def overview(df):
```

🎥 **Stuck on this page?** Watch: [Krish Naik: Exploratory Data Analysis with pandas \(search\)](#)

```

# One row per column: type, missing %, number of unique values, an example value.
return pd.DataFrame({
    "dtype": df.dtypes.astype(str),
    "missing": df.isna().sum(),
    "missing_%": (df.isna().mean() * 100).round(1),
    "n_unique": df.nunique(),
    "example": df.apply(lambda s: s.dropna().iloc[0] if s.notna().any() else np.nan)})

def num_summary(df):
    # describe() + skew, kurtosis, IQR-outlier count and zero count for numeric columns.
    num = df.select_dtypes("number")
    q1, q3 = num.quantile(0.25), num.quantile(0.75); iqr = q3 - q1
    res = num.describe().T
    res["skew"], res["kurt"] = num.skew(), num.kurt()
    res["n_iqr_outliers"] = ((num < q1 - 1.5 * iqr) | (num > q3 + 1.5 * iqr)).sum()
    res["n_zeros"] = (num == 0).sum()
    return res.round(2)

def rate_by(df, col, target, bins=None, q=None):
    # Target rate and count per category, or per bin of a numeric column.
    key = df[col]
    if bins is not None: key = pd.cut(key, bins)
    elif q is not None: key = pd.qcut(key, q, duplicates="drop")
    key = key.astype(object).where(key.notna(), "MISSING")
    return df.groupby(key, observed=True)[target].agg(rate="mean", n="size").round(3)

def missing_vs_target(df, target):
    # Is a column's missingness related to the target? (informative missingness / leakage)
    rows = []
    for c in df.columns.drop(target):
        m = df[c].isna()
        if 0 < m.sum() < len(df):
            rows.append((c, m.mean() * 100, df.loc[m, target].mean(), df.loc[~m, target].mean()))
    return pd.DataFrame(rows, columns=["column", "missing_%", "rate_if_missing", "rate_if_present"]).round(3)

def point_biserial(df, num_cols, target):
    # numeric feature vs 0/1 target
    rows = [(c, *stats.pointbiserialr(df[[c, target]].dropna()[target], df[[c, target]].dropna()[c])) for c in num_cols]
    return pd.DataFrame(rows, columns=["feature", "r_pb", "p_value"]).round(4)

def cramers_v(x, y):
    # categorical vs categorical, 0..1
    t = pd.crosstab(x, y); chi2 = stats.chi2_contingency(t, correction=False)[0]
    return np.sqrt(chi2 / (t.values.sum() * (min(t.shape) - 1)))

def vif(df_num):
    # VIF = 1 / (1 - R^2_j), via numpy least squares
    X = df_num.dropna().values; out = {}
    for j, c in enumerate(df_num.columns):
        y = X[:, j]; A = np.column_stack([np.ones(len(X)), np.delete(X, j, axis=1)])
        coef, *_ = np.linalg.lstsq(A, y, rcond=None)
        r2 = 1 - ((y - A @ coef) ** 2).sum() / ((y - y.mean()) ** 2).sum()
        out[c] = 1 / (1 - r2)
    return pd.Series(out, name="VIF").round(2)

```


41.3 The cookbook: 25 recipes, run on real data

All outputs below come from running the code on `aug_train.csv` (HR job change). Load: `df = pd.read_csv("Data_gathering/dataset/aug_train.csv")`.

41.3.1 Steps 2–4: structure, types, quality

R1. First look.

```
print(df.shape); print(df.dtypes.value_counts()); print(df.memory_usage(deep=True).sum() /
    1e6, "MB")
```

Output (real run):

```
(19158, 14)    10 text columns, 2 int64, 2 float64    12.39 MB
```

R5. Duplicates (exact, and ignoring the ID column, which hides duplicates):

```
df.duplicated().sum(), df.drop(columns="enrollee_id").duplicated().sum(), df["enrollee_id"]
    ].is_unique
```

Output (real run):

```
(0, 49, True)    # 49 people identical in every field except their ID
```

Decision: 49 of 19,158 is negligible; check whether they are the same person registered twice (keep one) or coincidence (keep all). In a group split they must stay together.

R6. Cardinality and constant columns: `df.nunique().sort_values()` → `relevent_experience` 2, `target` 2, `enrolled_university` 3, `gender` 3, ..., `city` 123, `enrollee_id` 19158. ID-like columns (unique per row) are dropped; constant columns carry no information.

R8. Parsing numbers stored as text.

```
exp = df["experience"].replace({">20": "21", "<1": "0"})
df["exp_num"] = pd.to_numeric(exp, errors="coerce")    # unparsable -> NaN, never crash
df["exp_num"].isna().sum()                            # 65 = exactly the originally
    missing values
```

R7. Value counts with percentages, including missing.

```
vc = df["education_level"].value_counts(dropna=False)
pd.concat([vc, (vc / len(df) * 100).round(1)], axis=1, keys=["n", "%"])
```

Output (real run):

	n	%
Graduate	11598	60.5
Masters	4361	22.8
High School	2017	10.5
NaN	460	2.4
Phd	414	2.2
Primary School	308	1.6

R9. Grouping rare levels (cities with <0.5% of rows):

```
freq = df["city"].value_counts(normalize=True)
rare = freq[freq < 0.005].index
df["city_grp"] = df["city"].where(~df["city"].isin(rare), "Other")    # 89 rare cities -> "
    Other"; 35 levels left
```

41.3.2 Step 5: missing values

R2. Missing table, sorted.

```
df.isna().agg(["sum", "mean"]).T.sort_values("sum", ascending=False).head()
```

Output (real run):

	sum	mean
company_type	6140	0.320
company_size	5938	0.310
gender	4508	0.235
major_discipline	2813	0.147
education_level	460	0.024

R3. How many fields are missing per row? `df.isna().sum(axis=1).value_counts().sort_index()` → {0: 8955, 1: 3718, 2: 3654, 3: 1953, 4: 628, 5: 176, 6: 62, 7: 12}. Only 47% of rows are complete: dropping rows with any NaN would discard more than half the data.

R4. Do fields go missing together?

```
m = df[["company_size", "company_type", "gender"]].isna()
m.corr() # correlation of the missingness indicators
pd.crosstab(m.company_size, m.company_type)
```

Output (real run):

	company_size	company_type	gender
company_size	1.00	0.84	0.06
company_type	0.84	1.00	0.07
company_type	False	True	
company_size			
False	12440	780	
True	578	5360	

Reading: company size and type are usually missing *together* (5,360 rows), suggesting one cause: the candidate has no current company.

R23. Missingness vs target and a missing indicator.

```
df["company_size_missing"] = df["company_size"].isna().astype(int)
df.groupby("company_size_missing")["target"].mean() # {0: 0.179, 1: 0.406}
```

Blank company size more than doubles the job-change rate: informative missingness. Keep the indicator; never fill with the mode.

41.3.3 Steps 6–8: target and univariate

R10. Percentiles that reveal tails.

```
df["training_hours"].describe(percentiles=[.01, .05, .5, .95, .99])
```

Output (real run):

```
| mean 65.4  std 60.1  min 1  1% 3  5% 7  50% 47  95% 188  99% 302  max 336
```

Mean (65) far above median (47): right skew. **R11. Skew of every numeric column:** `df.select_dtypes("number").skew()` → city_development_index -1.00, training_hours 1.82, exp_num 0.41. **R25. Does a log fix it?** `np.log1p(df.training_hours).skew()` → -0.35 (from 1.82).

R12. IQR outlier flags.

```
q1, q3 = df.training_hours.quantile([.25, .75]); iqr = q3 - q1
flag = (df.training_hours < q1 - 1.5 * iqr) | (df.training_hours > q3 + 1.5 * iqr)
q3 + 1.5 * iqr, flag.sum(), flag.mean() * 100          # (185.5, 984, 5.1 %)
```

R13. z-score flags with NumPy: $z = (x - x.mean()) / x.std()$; $(np.abs(z) > 3).sum() \rightarrow 450$. Reading: on a right-skewed variable both rules flag hundreds of *genuine* heavy learners. They are not errors: keep them; log-transform or cap for linear models.

41.3.4 Step 9: features vs target

R14. Rate with counts per category.

```
df.groupby("relevent_experience")["target"].agg(rate="mean", n="size")
```

Output (real run):

	rate	n
Has relevent experience	0.215	13792
No relevent experience	0.338	5366

R15. Crosstab normalised by row:

```
pd.crosstab(df["enrolled_university"], df["target"], normalize="index")
```

→ job-change share: full-time course 0.381, part-time 0.252, no enrolment 0.211.

R16. Numeric feature in quantile bins vs target.

```
df.groupby(pd.qcut(df["training_hours"], 4))["target"].mean()
```

Output (real run):

(0.999, 23.0]	0.250
(23.0, 47.0]	0.258
(47.0, 88.0]	0.253
(88.0, 336.0]	0.237

Flat: training hours carry almost no signal (as point-biserial $r = -0.022$ also says).

R18. Group means and a Welch t-test.

```
a = df.loc[df.target == 1, "city_development_index"]; b = df.loc[df.target == 0, "city_development_index"]
a.mean(), b.mean(), stats.ttest_ind(a, b, equal_var=False).pvalue      # 0.756, 0.853, ~0
```

R19. Chi-square for a categorical feature:

```
stats.chi2_contingency(pd.crosstab(df["company_type"].fillna("Unknown"), df["target"]))[1]
# 4e-204
```

With 19,000 rows almost anything is “significant”; use Cramér’s V for strength.

R20. Pearson vs Spearman with the target.

```
num = df[["city_development_index", "training_hours", "exp_num", "target"]]
num.corr()["target"], num.corr(method="spearman")["target"]
```

Output (real run):

```
| pearson : CDI -0.342  training_hours -0.022  exp_num -0.177
```

```
| spearman: CDI -0.279 training_hours -0.014 exp_num -0.184
```

41.3.5 Step 11: interactions and segments

R17. Two-factor pivot.

```
df.pivot_table(index="education_level", columns="relevent_experience", values="target",
                aggfunc="mean")
```

Output (real run):

	Has relevent experience	No relevent experience
Graduate	0.23	0.43
High School	0.12	0.24
Masters	0.19	0.32
Phd	0.12	0.17
Primary School	0.15	0.13

Graduates without relevant experience are the most mobile group (43%): a candidate segment for targeted job alerts.

41.3.6 Feature-engineering recipes that come out of EDA

R21. Relative-to-group feature with transform:

```
df["exp_vs_city_avg"] = df["exp_num"] - df.groupby("city")["exp_num"].transform("mean")
```

R22. Frequency encoding:

```
df["city_freq"] = df["city"].map(df["city"].value_counts(normalize=True)) # city_103
-> 0.2273
```

R24. Multi-condition labels with np.select:

```
df["career_stage"] = np.select([df.exp_num <= 2, df.exp_num <= 10], ["junior", "mid"],
                                default="senior")
df.groupby("career_stage")["target"].mean() # junior 0.384, mid 0.284, senior 0.171
```

41.4 Plot recipes (Matplotlib/Seaborn)

```
fig, axes = plt.subplots(1, 4, figsize=(15, 3)) # grid of histograms
for ax, c in zip(axes, num_cols): sns.histplot(df[c], kde=True, ax=ax); ax.set_title(c)
sns.boxplot(data=df, x="target", y="city_development_index") # numeric by class
sns.kdeplot(data=df, x="city_development_index", hue="target", common_norm=False) #
overlap
rate_by(df, "company_type", "target")["rate"].plot(kind="barh") # rate by category
sns.heatmap(df[num_cols].corr(), annot=True, fmt=".2f", cmap="coolwarm", vmin=-1, vmax=1)
sns.pointplot(data=df, x="career_stage", y="target", hue="relevent_experience", errorbar=
None) # interaction
sns.scatterplot(data=df, x="x1", y="x2", hue="target", alpha=0.5) # two features at
once
```

Rules: common_norm=False when comparing distributions of unequal-sized classes; always show *n* beside rates; sort bar charts; log axes for skewed money.

41.5 How to narrate EDA in the interview

1. Start with the frame: “One row is a candidate; target is job change; 25% positive, so I will use PR AUC/F1 and stratified splits.”
2. Data quality before statistics: types, codes, duplicates, malformed labels.
3. Missingness with the target, not just counts.
4. Rank features by separation (point-biserial, Cramér’s V , binned rates) and say which you would expect from domain sense.
5. Name one interaction and one non-linearity.
6. Close with a written plan: drops (leaks, IDs), encodings, transforms, imputation, model choice, metric.

41.6 Interview questions

Q41.1

★ **PYQ (Round 3 Part 1)** You get a new CSV in a Colab notebook. What are your first ten lines of code, and why each?

🕒 **In one line.** Load, then shape/head/info, the overview table (types, missing %, unique counts), duplicates, target balance, describe with skew and zero counts, value counts of categoricals, missingness vs target, and a quick leakage check of any column that perfectly separates the target.

⚙️ **The full answer.** Each line answers a question that changes a decision (Section 1). Leakage and data quality first, statistics second, plots third, model last.

📊 **Worked example.**

```
df = pd.read_csv(path); print(df.shape); df.head()
overview(df)
df.duplicated().sum(), df.drop(columns=id_col).duplicated().sum()
df[target].value_counts(normalize=True)
num_summary(df)
for c in df.select_dtypes(exclude="number"): print(df[c].value_counts(dropna=False).head(10))
missing_vs_target(df, target)
point_biserial(df, num_cols, target)
```

Q41.2

🧠 **Think deeper** `isna().sum()` shows zero missing values. Are you done with missing values?

🕒 **In one line.** No: missing values are often disguised as 0, -1, 999, empty strings, “NA” text or a default date; check zero counts, impossible values, spikes at round numbers and value counts of text columns.

 **Worked example.** The Pima diabetes data (Chapter 44) has no NaN but 374 zero Insulin and 35 zero blood-pressure readings: impossible values that are really missing.

Q41.3

How do you find out whether missing values are informative?

🎯 **In one line.** Compare the target rate (or target mean) for rows where the field is missing vs present, and check what else is missing with it; a big difference means missingness carries signal and should be encoded (indicator or “Unknown”), not imputed away.

 **Worked example.** HR data: job-change rate 0.406 when company size is blank vs 0.179 when present; company size and type are missing together in 5,360 rows.

Q41.4

Why do you print counts next to rates in every group-by?

🎯 **In one line.** Because a rate from a tiny group is unreliable; 100% from 2 people is noise, while 40% from 6,000 is solid; counts tell the listener how much to trust each number and whether to group rare levels.

Q41.5

 **Think deeper** Training hours shows 984 IQR outliers and 450 z-score outliers. Do you remove them?

🎯 **In one line.** No: the variable is right-skewed, so both rules flag many genuine heavy learners; they are valid values, so I keep them and, for linear models, log-transform or cap; tree models are unaffected.

⚙️ **The full answer.** Outlier rules assume a roughly symmetric distribution; on skewed data they flag the long tail by design. Remove only impossible values.

Q41.6

How do you turn EDA findings into a modelling plan? Give the plan for the HR data.

🎯 **In one line.** Translate each finding into an action: drop the ID; parse experience and last job; fix the 10/49 label; keep company missingness as a category; frequency- or target-encode city; keep CDI (strongest, non-linear, so a tree model); stratified split; class weights; evaluate with PR AUC and F1.

Q41.7

What is the difference between `groupby().agg()`, `transform()` and `pivot_table()` in EDA?

🎯 **In one line.** `agg` gives one summary row per group (rates by city); `transform` returns a group

statistic aligned to every original row (features relative to the group); `pivot_table` lays out a statistic across two grouping variables (interactions).

Q41.8

 **Think deeper** Your target rate by a feature looks almost perfect (0% vs 100%). What do you suspect and do?

 **In one line.** Leakage: the feature is probably created by or after the outcome (salary for placement, “interview scheduled” for hiring); confirm with domain questions and timing, then drop it from the predictive model (it may still be useful for a different question).

CHAPTER 42

EDA Lab II: The MBA Placement Dataset, End to End

🎥 **Companion video for this chapter:** [InfoEdge DS interview experience: the Round 3 placement task \(at 2:39\)](#)

This is the dataset reported in InfoEdge Round 3. We walk through the full EDA exactly as you would narrate it in the Colab round, with code, real printed output and plots. Notebook section: *Case 1* of `EDA_Lab.ipynb`.

⚠ Common mistake: about the numbers in this chapter

The real file is Kaggle's "Campus Recruitment" `Placement_Data_Full_Class.csv` (215 students). It could not be downloaded inside this build environment, so the outputs below come from a **simulated replica with the same 15 columns and similar structure** (generated by `eda_lab/make_placement_replica.py`, seed 42). Put the real CSV next to the notebook and every cell runs unchanged; your numbers will differ in detail, but the method, the checks and the leakage finding are the same.

42.1 The data dictionary

Column	Type	Meaning
sl_no	ID	serial number (drop for modelling)
gender	nominal	M / F
ssc_p, ssc_b	numeric, nominal	10th-grade percentage and board (Central / Others)
hsc_p, hsc_b, hsc_s	numeric, nominal, nominal	12th-grade percentage, board, stream (Commerce / Science / Arts)
degree_p, degree_t	numeric, nominal	undergraduate percentage and field (Comm&Mgmt / Sci&Tech / Others)
workex	binary	prior work experience
etest_p	numeric	employability test percentage
specialisation, mba_p	nominal, numeric	MBA specialisation (Mkt&Fin / Mkt&HR) and MBA percentage
status	target	Placed / Not Placed
salary	numeric	offered package, <i>recorded only after placement</i>

Frame: one row is a student; the decision is made *before* placement season (e.g. whom to coach), so only information available then may be used.

42.2 Step 1: load and look

```
pl = pd.read_csv("Placement_Data_Full_Class.csv")      # or make_placement_replica()
print(pl.shape)
```

🎥 **Stuck on this page?** Watch: [InfoEdge DS interview experience: the Round 3 placement task \(at 2:39\)](#)


```
overview(pl)
```

Output (real run):

```
(215, 15)
```

	dtype	missing	missing_%	n_unique	example
sl_no	int64	0	0.0	215	1
gender	str	0	0.0	2	F
ssc_p	float64	0	0.0	207	52.67
ssc_b	str	0	0.0	2	Others
hsc_p	float64	0	0.0	212	56.04
hsc_b	str	0	0.0	2	Others
hsc_s	str	0	0.0	3	Science
degree_p	float64	0	0.0	203	57.82
degree_t	str	0	0.0	3	Sci&Tech
workex	str	0	0.0	2	No
etest_p	float64	0	0.0	185	98.0
specialisation	str	0	0.0	2	Mkt&HR
mba_p	float64	0	0.0	202	72.4
status	str	0	0.0	2	Not Placed
salary	float64	69	32.1	102	263000.0

Say: “215 students, 15 columns, 5 numeric scores, 7 categoricals, an ID, the target, and salary. Only salary has missing values: 69 of them, 32%. That number is suspicious: it is close to the share of non-placed students, so I will check whether salary is missing *because* the student was not placed.”

42.3 Step 2: data quality

```
print("duplicate rows:", pl.duplicated().sum(), "| duplicate ids:", pl["sl_no"].duplicated
().sum())
pct = ["ssc_p", "hsc_p", "degree_p", "etest_p", "mba_p"]
print("all percentages inside [0, 100]:", ((pl[pct] >= 0) & (pl[pct] <= 100)).all().all())
```

Output (real run):

```
duplicate rows: 0 | duplicate ids: 0
all percentages inside [0, 100]: True
```

Every percentage lies in its valid range. Therefore a 96.9 in hsc_p is a topper, not an error: **no outlier deletion** (the “presence of mind” test from the reported interview).

42.4 Step 3: the target

```
pl["status"].value_counts(), pl["status"].value_counts(normalize=True)
```

Output (real run):

```
Placed 146 (0.679)    Not Placed 69 (0.321)
```

Mildly imbalanced (68/32). Accuracy is usable, but report precision, recall and F1 for the class the decision is about, and use stratified splits.

42.5 Step 4: salary, leakage and “if the package is given, we know the outcome”

42.5.1 The check

```
pl.groupby("status")["salary"].apply(lambda s: s.isna().mean())
rule = np.where(pl["salary"].notna(), "Placed", "Not Placed")
print("Accuracy of 'salary present => Placed':", (rule == pl["status"]).mean())
pd.crosstab(pl["salary"].notna().rename("salary_present"), pl["status"])
```

Output (real run):

```
status
Not Placed    1.0
Placed        0.0
Name: fraction_salary_missing
Accuracy of 'salary present => Placed': 1.0
status      Not Placed  Placed
salary_present
False                69      0
True                 0     146
```

42.5.2 What it means, said properly


- Salary is missing for **100%** of non-placed students and **0%** of placed ones. So the single rule “if a package is recorded, the student was placed” is **100% accurate**. Any model given salary (or even “salary is missing”) will find this and score near-perfect AUC.
- This is **target leakage**: salary is a *consequence* of being placed, recorded after the outcome. When the model is used (before placements), no student has a salary yet, so the feature does not exist and the near-perfect offline score would collapse.
- The missingness is not MCAR or MAR; it is structural (“not applicable”): there is no salary to impute for a non-placed student. Imputing it (with the median) would be meaningless and would still leak through the imputed pattern.
- **Decision**: drop salary (and anything derived from it) from the placement model. This is the first thing the successful candidate said in Round 3.
- Salary is still valuable for a *different* question: “what package will a placed student get?”. That is a regression on placed students only (or a two-part model: $P(\text{placed}) \times \mathbb{E}[\text{salary} \mid \text{placed}]$, Chapter 4).

42.5.3 Salary EDA among placed students

```
placed = pl[pl["status"] == "Placed"].copy(); s = placed["salary"]
print(s.describe().round(0))
print("skew:", round(s.skew(), 2), "| skew of log(salary):", round(np.log(s).skew(), 2))
q1, q3 = s.quantile([.25, .75]); upper = q3 + 1.5 * (q3 - q1)
print(f"IQR upper fence = {upper:.0f}; packages above it: {(s > upper).sum()}")
placed.groupby("workex")["salary"].median(), placed.groupby("gender")["salary"].median()
```

Output (real run):

```
count 146   mean 315623   std 98410   min 200000   25% 253000   50% 305000   75% 339000   max 784000
skew: 2.39 | skew of log(salary): 0.98
IQR upper fence = 468,000; packages above it: 6
workex  No 286500   Yes 311000           gender  F 275000   M 312000
```




Reading: salary is strongly right-skewed (mean 3.16 lakh > median 3.05 lakh; skew 2.4) with a floor at 2 lakh and six packages above the IQR fence, the highest 7.84 lakh. These are plausible genuine offers, so keep them; for a salary model use $\log(\text{salary})$ (skew drops to about 1) or a Gamma/quantile objective, and report MAPE. Work experience is associated with higher packages. A gender gap in medians appears; with 146 rows and no controls it is a *question to investigate* (does it persist after controlling for specialisation and scores?), not a conclusion.

42.6 Step 5: univariate analysis of the scores

```
num_summary(pl[pct])
```

Output (real run):

	count	mean	std	min	25%	50%	75%	max	skew	kurt	n_iqr_outliers	n_zeros
ssc_p	215.0	66.92	10.79	40.0	59.58	67.57	73.73	89.40	-0.11	-0.32	0	0
hsc_p	215.0	66.21	11.24	37.0	57.91	67.43	73.88	96.87	-0.23	-0.07	0	0
degree_p	215.0	66.52	7.64	50.0	60.95	67.03	71.72	83.63	-0.07	-0.55	0	0
etest_p	215.0	72.97	12.38	50.0	63.74	73.18	80.80	98.00	0.08	-0.53	0	0
mba_p	215.0	62.28	5.42	51.2	58.31	62.15	66.11	76.00	0.07	-0.59	0	0




Reading: all five are roughly symmetric ($|\text{skew}| < 0.25$) and light-tailed (negative excess kurtosis); no IQR outliers; means $\approx$ medians. No transformation needed; standardise for linear models. degree_p and mba_p have smaller spread (sd 7.6 and 5.4): differences of a few points matter more there.

42.7 Step 6: which scores separate placed from not placed?

```
pl["placed"] = (pl["status"] == "Placed").astype(int)
pl.groupby("status")[pct].mean().round(1).T
point_biserial(pl, pct, "placed")
```

Output (real run):

status	Not Placed	Placed	feature	r_pb	p_value
ssc_p	57.8	71.2	ssc_p	0.583	0.000
hsc_p	58.5	69.9	hsc_p	0.475	0.000
degree_p	61.5	68.9	degree_p	0.456	0.000
etest_p	71.2	73.8	mba_p	-0.109	0.113
mba_p	63.1	61.9	etest_p	0.100	0.146



```
for c in pct:
    a, b = pl.loc[pl.placed == 1, c], pl.loc[pl.placed == 0, c]
    t, p = stats.ttest_ind(a, b, equal_var=False)      # Welch's t-test
```

Output (real run):

```
ssc_p    placed mean  71.2 vs  57.8  Welch t = 10.39, p = 8.3e-19
hsc_p    placed mean  69.9 vs  58.5  Welch t =  7.69, p = 3.8e-12
degree_p placed mean  68.9 vs  61.5  Welch t =  7.54, p = 5.8e-12
etest_p  placed mean  73.8 vs  71.2  Welch t =  1.46, p = 0.15
mba_p    placed mean  61.9 vs  63.1  Welch t = -1.61, p = 0.11
```


Reading: the box plots barely overlap for 10th, 12th and degree percentages; placed students score about 13 points higher in school. The employability test and MBA percentage do not differ significantly. **Interview insight:** recruiters appear to screen on the long-run academic record; the MBA score itself is not what separates the groups. (Ask: is that because MBA grading is compressed? Its sd is only 5.4.)

42.8 Step 7: thresholds a placement cell can act on

```
rate_by(pl, "ssc_p", "placed", bins=[0, 55, 60, 65, 70, 75, 100])
```

Output (real run):

```
rate  n
(0, 55]  0.161  31
(55, 60]  0.346  26
(60, 65]  0.625  32
(65, 70]  0.864  44
(70, 75]  0.806  31
(75, 100] 0.961  51
```




Reading: placement climbs steeply from 16% (10th-grade score ≤ 55) to 96% (above 75), with the biggest jump around 60–65. Below 60, about two-thirds are not placed: a clear, explainable at-risk segment for coaching. The small dip at 70–75 (31 students) is within noise. The curve is S-shaped, which is exactly what logistic regression models.

42.9 Step 8: categorical features

```
cats = ["gender", "ssc_b", "hsc_b", "hsc_s", "degree_t", "workex", "specialisation"]
for c in cats:
    t = pd.crosstab(pl[c], pl["placed"]); chi2, p, *_ = stats.chi2_contingency(t)
    print(c, round(cramers_v(pl[c], pl["placed"]), 3), round(p, 4), rate_by(pl, c, "placed")["rate"].to_dict())
```

Output (real run):

feature	cramers_v	chi2_p	placement rate by level
workex	0.171	0.019	No 0.627 Yes 0.800
specialisation	0.164	0.024	Mkt&Fin 0.750 Mkt&HR 0.596
hsc_s	0.060	0.678	Arts 0.800 (n=10) Commerce 0.681 Science 0.663
degree_t	0.048	0.784	Comm&Mgmt 0.696 Others 0.647 (n=17) Sci&Tech 0.650
gender	0.027	0.810	F 0.662 M 0.688
ssc_b	0.027	0.798	Central 0.667 Others 0.692
hsc_b	0.000	1.000	Central 0.679 Others 0.679



```
pd.crosstab(pl["workex"], pl["specialisation"], values=pl["placed"], aggfunc="mean").round(2)
```


Output (real run):

specialisation	Mkt&Fin	Mkt&HR
workex		
No	0.70	0.55
Yes	0.85	0.73

Reading: work experience (+17 points) and Mkt&Fin specialisation (+15 points) are the only categorical features with a significant association. The effects add up: Mkt&HR students without work experience are placed at 55%, Mkt&Fin with experience at 85%. Boards, stream and degree type show no meaningful difference. hsc_s = Arts (80%) rests on only 10 students: do not read anything into it.

42.10 Step 9: correlations and multicollinearity

```
sns.heatmap(pl[pct + ["placed"]].corr(), annot=True, fmt=".2f", cmap="coolwarm", vmin=-1,
            vmax=1)
vif(pl[pct])
```




Output (real run):

```
| ssc_p 2.28   hsc_p 1.70   degree_p 1.55   etest_p 1.02   mba_p 1.02
```

Reading: the three academic scores correlate with each other (0.37–0.64: a shared “academic ability” factor) and with placement; employability test and MBA percentage are unrelated to everything. All VIFs are below 2.3, so multicollinearity is not a problem for a logistic regression; coefficients will be stable.

42.11 Step 10: multivariate view and a sanity-check model

```
sns.scatterplot(data=pl, x="ssc_p", y="degree_p", hue="status", style="workex")
X = pd.get_dummies(pl[pct + ["workex", "specialisation", "gender"]], drop_first=True).
    astype(float)
auc = cross_val_score(make_pipeline(StandardScaler(), LogisticRegression(max_iter=1000)),
                      X, pl["placed"], cv=StratifiedKFold(5, shuffle=True, random_state=0),
                      scoring="roc_auc")
```

**Output (real run):**

| CV ROC AUC using EDA-selected features (no salary): 0.916 +/- 0.024

Reading: placed students cluster in the upper-right of the school-vs-degree plane, with a fuzzy diagonal boundary: a linear model fits. Without salary, a 5-fold cross-validated logistic regression reaches AUC 0.92, honestly; with salary it would reach 1.0, dishonestly.

42.12 Step 11: the summary you say at the end

1. 215 students, no duplicates, all percentages valid, so no outlier removal.
2. **Salary is recorded only for placed students: a perfect but leaky predictor. Drop it for placement; use it only for a separate salary model on placed students.**
3. 68% placed: mildly imbalanced; report precision/recall/F1 and AUC, stratify splits.
4. 10th-grade, 12th-grade and degree percentages are the strongest drivers (point-biserial 0.46–0.58); employability test and MBA percentage add little.
5. Work experience and Mkt&Fin specialisation raise placement by about 15 points each, additively; boards and stream do not matter.
6. Placement rises steeply with 10th-grade score: students below 60 form a clear at-risk group.
7. No multicollinearity concerns; a logistic regression on these features gives CV AUC ≈ 0.92 .
8. Salary (for the separate question) is right-skewed: log-transform; higher with work experience; a gender gap worth investigating with controls.

42.13 Interview questions

Q42.1

★ **PYQ (Round 3, reported)** In the placement data, if the package (salary) is given, can you tell whether the student was placed? What do you do with the column?

🕒 **In one line.** Yes, perfectly: salary is filled only for placed students, so “salary present” predicts

placement with 100% accuracy; because it is recorded after the outcome and is unavailable at prediction time, it is target leakage and must be dropped from the placement model.

⚙️ **The full answer.** Show it with two lines: the fraction of missing salary by status (1.0 vs 0.0) and the crosstab of `salary.notna()` vs `status` (no off-diagonal cells). Then explain timing: the model is meant to be used before placement. Mention the legitimate use: a separate salary regression on placed students, or a two-part expected-salary model.

💡 **If the interviewer pushes.** “What if salary were only partly missing for placed students?” Then it is still leaky (its presence is mostly determined by the outcome) and also messy; the answer does not change. “How would a model reveal the leak if you missed it?” Near-perfect AUC and a dominant importance for salary or its missing indicator.

Q42.2

Would you impute the missing salaries? Why or why not?

🎯 **In one line.** No: for non-placed students salary is not unknown, it is not applicable; there is nothing to estimate, and any imputed value or indicator would still encode the outcome.

Q42.3

Which scores matter for placement, and how did you establish it?

🎯 **In one line.** 10th, 12th and degree percentages (point-biserial 0.58, 0.48, 0.46; Welch t-tests with $p < 10^{-11}$; non-overlapping box plots), while employability test and MBA percentage show no significant difference.

Q42.4

💡 **Think deeper** MBA percentage does not differ between placed and non-placed students. Does that mean the MBA does not matter?

🎯 **In one line.** No: it means variation in *MBA percentage among MBA students* does not separate outcomes in this sample, perhaps because MBA grades are compressed (sd 5.4) or recruiters screen on earlier records; everyone in the data has the MBA, so its overall value cannot be measured here.

Q42.5

How would you present the 10th-grade threshold finding to a placement cell?

🎯 **In one line.** As a simple, explainable at-risk rule with uncertainty: “students with 10th-grade scores below 60 are placed only 16–35% of the time (57 students), versus over 80% above 65; prioritise them for coaching and mock interviews”, and propose measuring whether coaching changes their rate.

Q42.6

 **Think deeper** The `hsc_s = Arts` group has an 80% placement rate, the highest. Would you add “Arts stream” as an important finding?

🎯 **In one line.** No: it is based on 10 students, and the chi-square test for stream is far from significant ($p = 0.68$); a difference of one or two students would change it; report it with its count and treat it as noise.

Q42.7

How would you check whether the gender gap in salary is real?

🎯 **In one line.** Compare medians with confidence intervals (bootstrap), then control for specialisation, work experience and scores with a regression of log salary, check sample sizes, and say clearly that an observational difference is not proof of discrimination or of its absence.

Q42.8

After this EDA, what model and preprocessing would you choose?

🎯 **In one line.** Drop `sl_no` and `salary`; one-hot encode the categoricals (drop one level for an unregularised model); standardise the scores; stratified split; logistic regression as the interpretable baseline (CV AUC ≈ 0.92), compared with a small random forest; report the confusion matrix, precision, recall, F1 and AUC.

CHAPTER 43

EDA Lab III: Will a Data Scientist Change Jobs? (Real HR Data)

🎥 **Companion video for this chapter:** [pandas User Guide: Working with missing data](#)

This chapter uses a **real** dataset from your repository, `Data_gathering/dataset/aug_train.csv` (Kaggle “HR Analytics: Job Change of Data Scientists”): 19,158 candidates who took a company’s data-science training, with profile, education, current company, experience and training hours; the target is 1 if the candidate is looking for a new job. It is a near-perfect stand-in for Naukri data, and it is messy in exactly the ways real job-portal data are. Notebook section: *Case 2*. All outputs are from real runs.

43.1 Step 1: load, overview, target

```
hr = load_csv("Data_gathering/dataset/aug_train.csv"); print(hr.shape)
overview(hr)
hr["target"].value_counts(normalize=True)
```

Output (real run):

```
(19158, 14)
```

	dtype	missing	missing%	n_unique	example
enrollee_id	int64	0	0.0	19158	8949
city	str	0	0.0	123	city_103
city_development_index	float64	0	0.0	93	0.92
gender	str	4508	23.5	3	Male
relevent_experience	str	0	0.0	2	Has relevent experience
enrolled_university	str	386	2.0	3	no_enrollment
education_level	str	460	2.4	5	Graduate
major_discipline	str	2813	14.7	6	STEM
experience	str	65	0.3	22	>20
company_size	str	5938	31.0	8	50-99
company_type	str	6140	32.0	6	Pvt Ltd
last_new_job	str	423	2.2	6	1
training_hours	int64	0	0.0	241	36
target	float64	0	0.0	2	1.0

```
target: 0.0 -> 0.751, 1.0 -> 0.249
```

Say: “One row per candidate; 25% are looking to change: imbalanced, so PR AUC and F1, stratified splits, maybe class weights. Two numeric columns only; experience and last-new-job are numbers stored as text; city has 123 levels; company fields are a third missing.”

43.2 Step 2: types and data quality

```
for c in ["experience", "last_new_job", "company_size"]:
    print(c, hr[c].value_counts(dropna=False).to_dict())
```

Output (real run):

```
experience: {'>20': 3286, '5': 1430, '4': 1403, ..., '1': 549, '<1': 522, ..., '20': 148, nan: 65}
last_new_job: {'1': 8040, '>4': 3290, '2': 2900, 'never': 2452, '4': 1029, '3': 1024, nan: 423}
company_size: {nan: 5938, '50-99': 3083, '100-500': 2571, '10000+': 2019, '10/49': 1471,
```

```
'1000-4999': 1328, '<10': 1308, '500-999': 877, '5000-9999': 563}
```

Three issues an interviewer wants you to spot: (1) experience is ordinal text with codes >20 and <1; (2) last_new_job has >4 and never; (3) company_size has a malformed label 10/49 (every other band uses a hyphen, so it should be 10-49). Also notice >20 is the single most common experience value: a capped variable (we cannot distinguish 21 from 35 years).

```
hr["experience_num"] = hr["experience"].replace({">20": "21", "<1": "0"}).astype(float)
hr["last_new_job_num"] = hr["last_new_job"].replace({">4": "5", "never": "0"}).astype(float)
hr["company_size"] = hr["company_size"].replace({"10/49": "10-49"})
size_order = ["<10", "10-49", "50-99", "100-500", "500-999", "1000-4999", "5000-9999", "10000+"]
hr["company_size_ord"] = hr["company_size"].map({s: i for i, s in enumerate(size_order)})
print("duplicates ignoring id:", hr.drop(columns="enrollee_id").duplicated().sum()) # 49
```

company_size is ordinal, so an **ordinal map** keeps the order (0–7) instead of one-hot. The 49 duplicates (identical except ID) are 0.26% of rows: note and keep, or deduplicate if they are re-registrations.

43.3 Step 3: missing values, with the target

```
missing_vs_target(hr[["gender", "enrolled_university", "education_level", "major_discipline", "experience", "company_size", "company_type", "last_new_job", "target"]], "target")
```

Output (real run):

column	missing_%	target_rate_if_missing	target_rate_if_present
company_type	32.0	0.388	0.184
company_size	31.0	0.406	0.179
gender	23.5	0.308	0.231
major_discipline	14.7	0.195	0.259
education_level	2.4	0.226	0.250
last_new_job	2.2	0.364	0.247
enrolled_university	2.0	0.319	0.248
experience	0.3	0.354	0.249



Reading:

- When company size or type is blank, the job-change rate is about 40% versus about 18% when present: more than double. Combined with the fact that they are missing *together* (Chapter 41,

R4), the likely meaning is “not currently employed”, a strong reason to look for a job. This is **informative missingness** (not MCAR).

- Missing gender also carries a higher rate (31% vs 23%): people who skip optional fields may be less engaged or different.
- Missing major discipline goes the other way (20% vs 26%): probably candidates without a degree.




Decisions: never drop these rows (53% of rows have something missing); encode the missingness of company size, company type and gender as an explicit "Unknown" level (or a missing indicator); for GBMs leave NaN; do not impute company size with the mode (that would erase the 40% signal).

43.4 Step 4: numeric features

```
num_summary(hr[["city_development_index", "training_hours", "experience_num", "
last_new_job_num"]])
```

Output (real run):

	count	mean	std	min	25%	50%	75%	max	skew	kurt	n_iqr_outliers
n_zeros											
city_development_index	19158.0	0.83	0.12	0.45	0.74	0.9	0.92	0.95	-1.00	-0.54	17
0											
training_hours	19158.0	65.37	60.06	1.00	23.00	47.0	88.00	336.00	1.82	3.84	984
0											
experience_num	19093.0	10.10	6.78	0.00	4.00	9.0	16.00	21.00	0.41	-1.17	0
522											
last_new_job_num	18735.0	2.00	1.68	0.00	1.00	1.0	3.00	5.00	0.80	-0.76	0
2452											



```
rate_by(hr, "city_development_index", "target", q=5)
point_biserial(hr, ["city_development_index", "training_hours", "experience_num", "
last_new_job_num"], "target")
```

Output (real run):

	rate	n	feature	r_pb	p_value
(0.447, 0.691]	0.551	3869	city_development_index	-0.342	0.000
(0.691, 0.878]	0.206	3827	experience_num	-0.177	0.000
(0.878, 0.92]	0.178	8925	last_new_job_num	-0.083	0.000
(0.92, 0.949]	0.104	2537	training_hours	-0.022	0.003

Reading:

- **City development index (CDI)** is left-skewed and *clumped* (only 93 distinct values: it is a city-level number shared by everyone in a city). It is the strongest numeric signal and strongly **non-linear**: 55% of candidates in the least developed cities want to move, against 10–21%

elsewhere; the drop happens almost entirely at the bottom. A linear model will under-fit this; use bins, a spline, or a tree model.

- **training_hours** is right-skewed (skew 1.8, 984 IQR “outliers” that are genuine heavy learners) and almost unrelated to the target ($r = -0.02$; flat across quartiles in R16). The $p = 0.003$ is “significant” only because n is huge: statistically detectable, practically useless.
- Experience: more experience, less job change ($r = -0.18$); 522 candidates with <1 year.


43.5 Step 5: categorical features

```
for c in ["relevent_experience", "enrolled_university", "education_level", "company_type",
         "last_new_job"]:
    print(rate_by(hr, c, "target").sort_values("rate", ascending=False))
```

Output (real run):

```
relevent_experience    No 0.338 (5366) | Has 0.215 (13792)
enrolled_university   Full time 0.381 (3757) | MISSING 0.319 (386) | Part time 0.252 (1198) | none 0.211 (13817)
education_level       Graduate 0.280 | MISSING 0.226 | Masters 0.214 | High School 0.195 | Phd 0.140 | Primary
0.133
company_type          MISSING 0.388 (6140) | Other 0.240 | Early Stage Startup 0.235 | Public Sector 0.220
| NGO 0.186 | Pvt Ltd 0.181 (9817) | Funded Startup 0.140 (1001)
last_new_job          MISSING 0.364 | never 0.301 | 1 0.264 | 2 0.241 | 3 0.226 | 4 0.222 | >4 0.182
```

```
[(c, cramers_v(hr[c].fillna("MISSING"), hr["target"])) for c in [...]]
```



Output (real run):

```
city 0.396 | company_size 0.246 | company_type 0.224 | enrolled_university 0.156 | relevent_experience 0.128
education_level 0.093 | last_new_job 0.092 | gender 0.078 | major_discipline 0.058
```

Reading: city is the strongest categorical feature; company size and type are next, *mostly because of their MISSING level*; full-time students (38%) and people without relevant experience (34%) are the most likely to move; employees of funded startups the least (14%). last_new_job shows a monotone pattern: the longer since the last job change, the less likely to change now.

43.6 Step 6: the high-cardinality city

```
cs = hr.groupby("city").agg(n=("target", "size"), rate=("target", "mean"), cdi=(
    "city_development_index", "first"))
print("cities:", len(cs), "| cities with < 20 candidates:", (cs["n"] < 20).sum())
```

```
print("CDI constant within each city?", (hr.groupby("city")["city_development_index"].
    nunique() == 1).all())
print("corr(city rate, city CDI) =", round(cs["rate"].corr(cs["cdi"]), 3))
prior, m = hr["target"].mean(), 20
cs["smoothed_rate"] = (cs["n"] * cs["rate"] + m * prior) / (cs["n"] + m)
```

Output (real run):

```
cities: 123 | cities with < 20 candidates: 44
CDI constant within each city? True
corr(city rate, city CDI) = -0.705
   n  rate  cdi  smoothed_rate
city_103 4355 0.213 0.920         0.213
city_21  2702 0.591 0.624         0.589
city_16  1533 0.117 0.910         0.118
city_114 1336 0.100 0.926         0.102
```


Reading: 44 of 123 cities have fewer than 20 candidates, so their raw rates are noise; one-hot would add 122 sparse columns. CDI is a *property of the city* and explains much of the city effect (correlation -0.71 between a city's rate and its CDI). One city alone, city_21 (CDI 0.624, 2,702 candidates, 59% movers), drives a large share of the low-CDI signal: a business question in itself. **Encoding plan:** keep CDI; add city frequency; add out-of-fold **smoothed target encoding** (shown with $m = 20$); group rare cities.

43.7 Step 7: experience and an interaction

```
hr.groupby(pd.cut(hr["experience_num"], [-1, 2, 5, 10, 15, 21]))["target"].agg(["mean", "size"])
sns.pointplot(data=hr, x=pd.cut(hr["experience_num"], [-1, 2, 5, 10, 15, 21]).astype(str),
    y="target",
    hue="relevent_experience", errorbar=None)
```

Output (real run):

```
(-1, 2]  0.384  2198      (2, 5]  0.322  4187      (5, 10]  0.252  5011
(10, 15] 0.191  2829      (15, 21] 0.156  4868
```



Reading: job-change propensity falls steadily with experience (38% for 0–2 years to 16% for 15+). At every experience level, candidates without relevant experience are more likely to move: they are trying to switch into data science.

43.8 Step 8: findings and the modelling plan

1. 19,158 candidates, 25% movers: imbalanced; use stratified CV, PR AUC/F1, class weights; threshold by recruiter capacity.
2. Fix types: parse experience and last_new_job (note the cap at 20+), repair the 10/49 label, ordinal-encode company size; drop enrollee_id.
3. **Missing company size/type is the second-strongest signal** (40% vs 18%): keep as “Unknown” or let a GBM handle NaN; do not mode-impute; do not drop rows.
4. **City development index** is the strongest driver and non-linear (55% in the least developed cities): tree models or binning.
5. City has 123 levels: CDI + frequency + smoothed out-of-fold target encoding.
6. Juniors, people without relevant experience and full-time students move more; training hours are irrelevant (statistically significant only because of n).
7. Model: LightGBM/CatBoost (handles NaN and categoricals natively) vs logistic regression with the engineered features; validate with stratified 5-fold CV.

43.9 Interview questions

Q43.1

 **Think deeper** A third of candidates have no company size or company type. How do you handle it, and why not impute the mode?

 **In one line.** I first check missingness against the target: blank company fields come with a 40% job-change rate versus 18%, and they are missing together, which suggests “currently not employed”; so I keep missingness as its own category or indicator (or let a GBM use NaN), because mode imputation would label these people as working at the most common company size and destroy the strongest signal.

Q43.2

experience contains “>20” and “<1”. How do you convert it, and what information is lost?

 **In one line.** Map “<1” to 0 and “>20” to 21 (or keep an ordinal scale) and convert with `pd.to_numeric(errors="coerce")`; the cap means we cannot distinguish 21 from 35 years, so the top value is censored, which is fine for a tree but should be noted for any linear interpretation.

Q43.3

training_hours has a p-value of 0.003 against the target. Is it a useful feature?

 **In one line.** Hardly: the correlation is -0.02 and the job-change rate is flat across quartiles (0.250, 0.258, 0.253, 0.237); with 19,000 rows tiny effects become “significant”, so judge by effect size, not p-value.

Q43.4

How would you encode city (123 levels)?

🎯 **In one line.** Not one-hot: keep the city development index (a city-level summary), add frequency encoding, and add an out-of-fold smoothed target encoding with shrinkage toward the global rate for small cities; or use CatBoost's ordered target statistics.

📊 **Worked example.** Smoothed rate = $(n\bar{y}_c + m\bar{y})/(n + m)$ with $m = 20$: a city with 3 candidates all moving gets $(3 + 20 \times 0.249)/23 = 0.35$, not 1.0.

Q43.5

🧠 **Think deeper** The job-change rate is 55% in the lowest CDI quintile and 10–21% elsewhere. What does that suggest about the model and the business?

🎯 **In one line.** The relation is strongly non-linear (a threshold effect), so a linear term in CDI under-fits: use bins, splines or trees; for the business, candidates in less developed cities are far more mobile, a segment to target with relocation-friendly or remote job recommendations, though one large city drives much of it.

Q43.6

What malformed value did you find and how did you detect it?

🎯 **In one line.** `company_size = "10/49"`, while every other band uses a hyphen; it shows up immediately in `value_counts()`, which is why I print value counts of every categorical before analysis.

CHAPTER 44

EDA Lab IV: Titanic (Interactions and Text Features) and Pima Diabetes (Disguised Missing Values)

🎥 **Companion video for this chapter:** [Kaggle: Titanic, Machine Learning from Disaster \(search\)](#)

Two more **real** datasets from your repository, each chosen for one lesson interviewers love: Titanic for interactions, feature extraction from text and group-wise imputation; Pima diabetes for missing values that `isna()` cannot see. Notebook sections: *Case 3* and *Case 4*.

44.1 Titanic: who survived?

44.1.1 Overview

```
ti = load_csv("understanding_your_data_eda/train.csv"); overview(ti)
```

Output (real run):

	dtype	missing	missing_%	n_unique	example
PassengerId	int64	0	0.0	891	1
Survived	int64	0	0.0	2	0
Pclass	int64	0	0.0	3	3
Name	str	0	0.0	891	Braund, Mr. Owen Harris
Sex	str	0	0.0	2	male
Age	float64	177	19.9	88	22.0
SibSp	int64	0	0.0	7	1
Parch	int64	0	0.0	7	0
Ticket	str	0	0.0	681	A/5 21171
Fare	float64	0	0.0	248	7.25
Cabin	str	687	77.1	147	C85
Embarked	str	2	0.2	3	S


Note that `Pclass` is stored as an integer but is an *ordinal category*; `Name` and `Ticket` are text that may hide features; `Cabin` is 77% missing.

44.1.2 Sex and class, and their interaction

```
print("survival rate:", ti["Survived"].mean().round(3))
rate_by(ti, "Sex", "Survived"); rate_by(ti, "Pclass", "Survived")
pd.crosstab(ti["Pclass"], ti["Sex"], values=ti["Survived"], aggfunc="mean").round(3)
```

Output (real run):

survival rate: 0.384			Pclass: 1 0.630 (216) 2 0.473 (184) 3 0.242 (491)		
Sex: female 0.742 (314) male 0.189 (577)					
Sex	female	male			
Pclass					
1	0.968	0.369			
2	0.921	0.157			
3	0.500	0.135			



Reading: sex is the dominant factor (74% vs 19%), class the second. But the effect of class **depends on** sex: for women, moving from 2nd to 3rd class halves survival (92% to 50%); for men the big drop is from 1st to 2nd (37% to 16%). That is an **interaction**. A decision tree finds it automatically; a logistic regression needs an explicit Sex×Pclass feature.

44.1.3 Extracting a feature from text: the title

```
ti["Title"] = ti["Name"].str.extract(r",\s*([^\.,]+)\.")[0] # "Braund, Mr. Owen"
-> "Mr"
ti["Title"] = ti["Title"].replace({"Mlle": "Miss", "Ms": "Miss", "Mme": "Mrs"})
ti.loc[~ti["Title"].isin(["Mr", "Mrs", "Miss", "Master"]), "Title"] = "Rare"
ti.groupby("Title").agg(n=("Survived", "size"), survival=("Survived", "mean"), median_age=
   =("Age", "median"))
```

Output (real run):

	n	survival	median_age
Master	40	0.57	3.5
Miss	185	0.70	21.0
Mr	517	0.16	30.0
Mrs	126	0.79	35.0
Rare	23	0.35	48.5

Reading: the title combines sex, age and status in one feature: “Master” means a boy (median age 3.5), and boys survived at 57% while men (“Mr”) survived at 16%. The regular expression reads “after a comma and spaces, capture everything up to the first full stop”.

44.1.4 Group-wise imputation of Age

```
ti["Age_filled"] = ti["Age"].fillna(ti.groupby("Title")["Age"].transform("median"))
print("global median age:", ti["Age"].median(), "| ages imputed:", ti["Age"].isna().sum())
# 28.0, 177
```


Filling all 177 missing ages with the global median 28 would turn missing boys into adults and erase the “children first” pattern; filling by title gives boys 3.5 and married women 35. (In a modelling pipeline compute the medians on training folds only.)

44.1.5 Fare, family size and cabin

```
print("Fare skew:", ti["Fare"].skew().round(2), "| log1p skew:", np.log1p(ti["Fare"]).skew(
    ).round(2),
      "| zero fares:", (ti["Fare"] == 0).sum())
ti["FamilySize"] = ti["SibSp"] + ti["Parch"] + 1
rate_by(ti, "FamilySize", "Survived", bins=[0, 1, 4, 11])
ti.groupby(ti["Cabin"].isna())["Survived"].mean()
```

Output (real run):

```
Fare skew: 4.79 | log1p skew: 0.39 | zero fares: 15
(0, 1] 0.304 537 (1, 4] 0.579 292 (4, 11] 0.161 62
survival if Cabin missing vs present: {False: 0.667, True: 0.3}
```



Reading:

- Fare is extremely right-skewed (4.8; max 512): log it for linear models. Fifteen zero fares are suspicious (crew? complimentary tickets?): investigate, do not silently keep them as real prices.
- Family size is non-monotonic:** alone 30%, 2–4 people 58%, 5 or more 16%. A linear term would miss it; bins or trees capture it. This is an engineered feature (`SibSp + Parch + 1`) that is more meaningful than its parts.
- Missing Cabin goes with 30% survival versus 67% when present: cabin numbers were recorded mostly for first-class passengers, so missingness encodes class (informative, MNAR-like). Use “has cabin” as a feature rather than imputing cabins.
- The age KDE shows the survival advantage concentrated among young children.

44.1.6 Titanic findings

Sex $\gg$ class with a strong interaction; title is a powerful compact feature; impute age by title; log fare and check zero fares; family size is non-monotonic; cabin missingness is informative; Embarked has 2 missing values (fill with the mode, harmless). A depth-limited tree or gradient boosting will capture the interactions; logistic regression needs engineered interaction terms.

44.2 Pima Indians diabetes: missing values in disguise

44.2.1 The dataset says it has no missing values

```
pi = load_csv("Module_1_Foundation_of_ML_AI/dataset/diabetes.csv")
print(pi.shape, "| isna().sum() total:", pi.isna().sum().sum())
num_summary(pi)
```

Output (real run):

```
(768, 9) | isna().sum() total: 0
```

	count	mean	std	min	25%	50%	75%	max	skew	kurt	n_iqr_outliers	n_zeros
Pregnancies	768.0	3.85	3.37	0.00	1.00	3.00	6.00	17.0	0.90	0.16	4	111
Glucose	768.0	120.89	31.97	0.00	99.00	117.00	140.25	199.0	0.17	0.64	5	5
BloodPressure	768.0	69.11	19.36	0.00	62.00	72.00	80.00	122.0	-1.84	5.18	45	35
SkinThickness	768.0	20.54	15.95	0.00	0.00	23.00	32.00	99.0	0.11	-0.52	1	227
Insulin	768.0	79.80	115.24	0.00	0.00	30.50	127.25	846.0	2.27	7.21	34	374
BMI	768.0	31.99	7.88	0.00	27.30	32.00	36.60	67.1	-0.43	3.29	19	11
DiabetesPedigreeFunction	768.0	0.47	0.33	0.08	0.24	0.37	0.63	2.42	1.92	5.59	29	0
Age	768.0	33.24	11.76	21.00	24.00	29.00	41.00	81.0	1.13	0.64	9	0
Outcome	768.0	0.35	0.48	0.00	0.00	0.00	1.00	1.0	0.64	-1.60	0	500

Reading the clues: min = 0 for Glucose, BloodPressure, SkinThickness, Insulin and BMI. A living person cannot have a blood pressure, BMI or glucose of 0. The n_zeros column counts them: 374 Insulin, 227 skin thickness, 35 blood pressure, 11 BMI, 5 glucose. BloodPressure's skew of -1.84 and kurtosis 5.2 are artefacts of the spike at 0. These are **missing values coded as 0**. (Zero pregnancies is valid.)




44.2.2 Recode and re-examine

```
impossible_zero = ["Glucose", "BloodPressure", "SkinThickness", "Insulin", "BMI"]
pi2 = pi2.copy(); pi2[impossible_zero] = pi2[impossible_zero].replace(0, np.nan)
(pi2.isna().mean() * 100).round(1)
pi2.groupby("Outcome").median().round(1).T
point_biserial(pi2, pi2.columns.drop("Outcome"), "Outcome")
```

Output (real run):

```
missing %: Glucose 0.7 | BloodPressure 4.6 | SkinThickness 29.6 | Insulin 48.7 | BMI 1.4
medians by outcome after fixing:
Outcome      0      1
Glucose      107.0  140.0
0.171
BloodPressure  70.0  74.5
SkinThickness  27.0  32.0
Insulin       102.5  169.5
BMI           30.1  34.3
Age           27.0  36.0
```

point-biserial with Outcome:
 Glucose 0.495 | BMI 0.314 | Insulin 0.303 | SkinThickness 0.260
 Age 0.238 | Pregnancies 0.222 | Pedigree 0.174 | BloodPressure

**Reading:**

- Insulin is missing for almost half the patients (48.7%) and skin thickness for 30%: these columns need careful imputation or a missing indicator.
- **The trap:** with the zeros left in, the median insulin of diabetic patients is **0** (more than half were simply not measured), suggesting the absurd conclusion that diabetics have no insulin.

After recoding, diabetic patients have *higher* median insulin (169.5 vs 102.5), as expected with insulin resistance.

- Glucose is the dominant predictor ($r = 0.50$; medians 140 vs 107); then BMI and insulin (≈ 0.3), age and pregnancies.
- The blood-pressure histogram's spike at 0 is the visual signature of disguised missingness; always look at histograms, not just `isna()`.

Imputation plan: median or KNN/iterative imputation computed on training folds only, *without* using the outcome; add missing indicators for Insulin and SkinThickness; compare model performance with and without those two columns.

44.3 Interview questions

Q44.1

 **Think deeper** The Pima dataset has no NaN values. Why is that misleading, and how did you detect the problem?

 **In one line.** Missing measurements were stored as 0; I detected them because `describe()` shows minimum 0 for glucose, blood pressure, skin thickness, insulin and BMI (physically impossible), the zero counts are large, and histograms show spikes at 0; recoding them as NaN reveals 48.7% missing insulin.

 **The full answer.** The consequence of ignoring it: means are biased downward, distributions distorted, and the diabetic patients' median insulin appears to be 0.

Q44.2

Explain the sex-class interaction in Titanic and how different models handle it.

 **In one line.** Class matters differently for each sex (women: 97%, 92%, 50% across classes; men: 37%, 16%, 14%); trees and boosting learn this through successive splits, while logistic regression, being additive in log-odds, needs an explicit `Sex×Pclass` interaction feature.

Q44.3

How would you impute Titanic's missing ages, and why not the global median?

 **In one line.** By the median age of the passenger's title (Master 3.5, Miss 21, Mr 30, Mrs 35), computed on training data; the global median (28) would turn missing boys into adults and weaken the strong child-survival pattern.

Q44.4

Cabin is 77% missing. Drop it?

 **In one line.** Drop the raw cabin string, but keep "has cabin" (survival 67% vs 30%), because

missingness encodes passenger class and wealth; the deck letter can also be extracted where present.

Q44.5

What does a non-monotonic relationship like family size vs survival imply for modelling?

🎯 **In one line.** A single linear coefficient cannot represent a rise-then-fall pattern (alone 30%, 2–4 58%, 5+ 16%); use bins, polynomial or spline terms, or tree-based models.

Q44.6

Write the pandas code to extract the title from a name like “Braund, Mr. Owen Harris”.

🎯 **In one line.** `df["Name"].str.extract(r",\s*([^\.,]+)\.")[0]`: after the comma and optional spaces, capture characters that are not a full stop, up to the first full stop.

Q44.7

🧠 **Think deeper** Across the four datasets, what are the recurring EDA lessons?

🎯 **In one line.** Check how a value was recorded before trusting it (salary after placement, cabins for first class, zeros for unmeasured insulin); relate missingness to the target; look for non-linear and interaction effects with binned rates and pivots; extract features from text; and never call a value an outlier without a domain range.

Part IX

SQL for Data Scientists: From Zero to Interview Level (with a Colab Notebook)

CHAPTER 45

SQL Level 1: Tables, Rows and Reading Data from One Table

🎥 **Companion video for this chapter:** [Programming with Mosh: SQL Tutorial, Full Database Course for Beginners](#)

Why this part exists, and how much SQL InfoEdge asks

The InfoEdge candidate whose experience this book follows said Rounds 1 and 2 are “drilling” rounds with three or four questions from each of six topics, and **SQL is one of the six** (with Python, machine learning, deep learning, statistics and probability). He did not remember the exact SQL questions, so this part does not invent “PYQs”. Instead, questions marked 🟢 **Most asked** are the standard data-science SQL questions that appear in almost every analytics and data-science interview (Nth highest salary, joins and row counts, WHERE vs HAVING, ranking functions, duplicates, NULL traps, streaks, funnels, retention). They come first in each chapter’s question set.

The part is organised in four levels. Each level assumes only the previous ones.

Level	Chapter	What you can do at the end
1	45	Read any single table: choose columns, filter rows, sort, handle NULL, create new columns with CASE, strings and dates
2	46	Summarise with GROUP BY and HAVING; combine tables with every kind of JOIN; stack results with UNION
3	47	Nest queries (subqueries, CTEs, recursive CTEs); use window functions (ranking, LAG/LEAD, running totals); change data safely
4	48	Solve the classic interview patterns (Nth highest, top-N per group, duplicates, streaks, funnels, A/B, cohorts, median, pivot); build leak-free ML feature tables; explain indexes, query plans and schema design

💻 **Every query in this part was actually run.** The database is in `InfoEdge_Interview_Companion/sql_lab/schema.sql`; the queries are in `sql_lab/queries.sql`; the output printed under each query is the real SQLite output, generated by `sql_lab/run_queries.py`. The Colab notebook `sql_lab/SQL_Lab.ipynb` creates the database inside Python (no installation: SQLite ships with Python) and runs every query, plus practice exercises with hidden solutions. Type each query yourself before reading its output.

45.1 What a database is, from zero

45.1.1 Tables, rows and columns

Think of an Excel workbook. Each sheet holds one kind of thing: one sheet for candidates, one for jobs, one for applications. A **relational database** is the same idea made strict and fast:

- a **table** is one sheet; it holds one kind of thing (candidates);

- a **row** (also *record*) is one thing (one candidate, Aarav);
- a **column** (also *field* or *attribute*) is one property that every row has (city, years of experience);
- every column has a **data type** that the database enforces: INTEGER (whole numbers), DECIMAL(5,1) (numbers with up to 5 digits, 1 after the point), VARCHAR(20) (text up to 20 characters), DATE ('2026-03-01'), BOOLEAN. Excel would let you type "five" into an experience cell; a database refuses.

SQL (Structured Query Language, said "sequel" or "S-Q-L") is the language for asking a database questions. It is *declarative*: you describe *what* rows you want, not *how* to loop over them; the database's *query optimiser* decides how. That is why one line of SQL can replace twenty lines of Python loops, and why the same query can run on ten rows or ten billion.

45.1.2 NULL: the value that means "unknown"

A cell can be empty. SQL marks an empty cell with NULL, meaning *unknown* or *not applicable*. NULL is **not** zero and **not** the empty string ". A fresher has no current salary: NULL. A candidate who skipped the city field: NULL. Because NULL means "unknown", any comparison with NULL is also unknown (Section 45.5); this single fact causes more wrong interview answers than anything else in SQL.


45.1.3 Keys: how rows are identified and linked

- A **primary key** (PK) is a column (or set of columns) whose value is unique for every row and never NULL. `cand_id` identifies a candidate even if two candidates are both called Aarav.
- A **foreign key** (FK) is a column that stores another table's primary key, which creates a link. `applications.cand_id` says *which* candidate applied. The database can refuse an application for candidate 999 if no candidate 999 exists (*referential integrity*).
- A **unique key** is unique like a PK but may allow NULL and a table may have several (e.g. email).
- A **composite key** uses several columns together: in a skills table, (`cand_id`, `skill`) is unique but neither column alone is.

Why split data into several linked tables instead of one giant sheet? If the company name were typed into every job and application row, renaming "PayWave" would mean editing hundreds of rows and missing one would make the data contradict itself. Storing each fact once and linking by keys is called **normalisation** (Chapter 48).

45.2 The practice database: MiniNaukri

All queries in Part IX use one small job-portal database, deliberately close to what Naukri stores. It is small enough that you can check every answer by eye, and it contains the awkward things real data has: NULLs, a duplicate row, a company with no jobs, candidates who never applied, ties in salary.



An arrow points from the table holding the foreign key to the table it refers to. “many:1” means many jobs can belong to one company, but each job belongs to exactly one company.

Table	Rows	What one row is, and the awkward bits planted in it
companies	6	A hiring company. GreenGrid (Delhi) has posted no jobs.
candidates	10	A job seeker. Sana’s city is NULL; Kabir is a fresher with NULL salary; Sana and Priya never applied; referred_by points to another candidate.
jobs	10	A job posting. Job 210 (Hyderabad) received no applications.
applications	23	One candidate applying to one job, with the current status. Row 22 is an accidental duplicate of row 1.
logins	18	One login event. Candidate 109 logged in twice on 4 March.
impressions	16	A job card shown to a candidate during an A/B test of a new ranking (variant B), and whether it was clicked.
employees	11	An internal employee, for the classic salary questions. Zoya’s salary is NULL; Kavya and Manav tie at 45.

45.2.1 Running SQL from Python (the way you will in a notebook round)

```

import sqlite3, pandas as pd
con = sqlite3.connect(":memory:") # a database that lives in RAM
con.executescript(open("sql_lab/schema.sql").read()) # create and fill the tables
df = pd.read_sql("SELECT * FROM candidates WHERE exp_yrs >= 5", con)
print(df)

```

`pd.read_sql` returns a `DataFrame`, so you can move between SQL and pandas freely. In industry the connection is to MySQL, PostgreSQL, Redshift, BigQuery or Hive rather than SQLite, but the SQL in this part is standard and runs on all of them (dialect differences are listed in Section 45.10.1).

45.3 SELECT and FROM: choosing columns

The smallest query names a table (`FROM`) and the columns to show (`SELECT`). `*` means “all columns”.

```
SELECT * FROM companies;
```

Output (real run):

	company_id	name	city	industry	size
1		DataNest	Bengaluru	IT	large
2		PayWave	Mumbai	Fintech	mid
3		MediCore	Pune	Healthcare	small

```

4      | ShopKart | Bengaluru | E-commerce | large
5      | EduSpark | Noida     | Edtech     | mid
6      | GreenGrid | Delhi    | Energy     | small
(6 rows)

```

```

SELECT name, city, exp_yrs
FROM candidates;

```

Output (real run):

```

name | city      | exp_yrs
-----+-----+-----
Aarav | Bengaluru | 2
Diya  | Mumbai   | 5
Kabir  | Delhi    | 0
Meera  | Bengaluru | 7
Rohan  | Pune     | 3
Sana   | NULL     | 1
Vikram | Noida    | 10
Ananya | Mumbai   | 4
Ishaan | Bengaluru | 2
Priya  | Delhi    | 6
(10 rows)

```

Read the output like a small spreadsheet: one line per row, columns separated by |. Note NULL in Sana's city: the value is unknown.

⚠ Common mistake: SELECT * in production

SELECT * is fine for exploring. In production code and in interviews, list the columns you need: it transfers less data, documents intent, and does not break silently when someone adds a column. On columnar warehouses (BigQuery, Redshift) you also pay per column read.

45.4 WHERE: choosing rows

WHERE keeps only the rows for which a condition is true. Comparison operators: =, <> (or !=, "not equal"), <, <=, >, >=. Text values go in single quotes.

```

SELECT name, city, exp_yrs, ctc_lpa
FROM candidates
WHERE city = 'Bengaluru';

```

Output (real run):

```

name | city      | exp_yrs | ctc_lpa
-----+-----+-----+-----
Aarav | Bengaluru | 2       | 8
Meera  | Bengaluru | 7       | 30
Ishaan | Bengaluru | 2       | 9.5
(3 rows)

```

45.4.1 AND, OR and the precedence trap

AND needs both conditions true; OR needs at least one. **AND is evaluated before OR**, exactly as multiplication comes before addition. Suppose the recruiter asks for "candidates from Bengaluru or Mumbai with at least 5 years". The obvious query is wrong:

```

SELECT name, city, exp_yrs
FROM candidates

```

```
WHERE city = 'Bengaluru' OR city = 'Mumbai' AND exp_yrs >= 5;
```

Output (real run):

```
name | city | exp_yrs
-----+-----+-----
Aarav | Bengaluru | 2
Diya | Mumbai | 5
Meera | Bengaluru | 7
Ishaan | Bengaluru | 2
(4 rows)
```

SQL read it as `city = 'Bengaluru' OR (city = 'Mumbai' AND exp_yrs >= 5)`, so every Bengaluru candidate came back, even with 2 years (Aarav, Ishaan). Brackets fix it:

```
SELECT name, city, exp_yrs
FROM candidates
WHERE (city = 'Bengaluru' OR city = 'Mumbai') AND exp_yrs >= 5;
```

Output (real run):

```
name | city | exp_yrs
-----+-----+-----
Diya | Mumbai | 5
Meera | Bengaluru | 7
(2 rows)
```

Remember

Whenever a WHERE mixes AND and OR, add brackets, even if you think precedence already does what you want. Interviewers plant this bug in “find the error” questions.

45.4.2 IN, BETWEEN and LIKE

- `x IN ('a', 'b', 'c')` is a short way to write `x = 'a' OR x = 'b' OR x = 'c'`.
- `x BETWEEN 3 AND 10` means `x >= 3 AND x <= 10`: **both ends included**.
- `x LIKE 'Data%'` matches text patterns: `%` = any number of characters (including none), `_` = exactly one character. `'%Engineer'` = ends with Engineer; `'%data%'` = contains data; `'___'` = exactly three characters.

```
SELECT name, city, exp_yrs
FROM candidates
WHERE city IN ('Delhi', 'Pune', 'Noida')
AND exp_yrs BETWEEN 3 AND 10;
```

Output (real run):

```
name | city | exp_yrs
-----+-----+-----
Rohan | Pune | 3
Vikram | Noida | 10
Priya | Delhi | 6
(3 rows)
```

```
SELECT job_id, title, city
FROM jobs
WHERE title LIKE 'Data%';
```

Output (real run):

```

job_id | title          | city
-----+-----+-----
201    | Data Scientist | Bengaluru
203    | Data Analyst   | Mumbai
204    | Data Scientist | Mumbai
205    | Data Analyst   | Pune
206    | Data Scientist | Bengaluru
207    | Data Analyst   | Bengaluru
209    | Data Scientist | Noida
(7 rows)

```

⚠ Common mistake: BETWEEN with timestamps, and LIKE's case sensitivity

- If `applied_at` is a *timestamp*, `applied_at BETWEEN '2026-03-01' AND '2026-03-31'` stops at 31 March 00:00:00 and silently drops everything on the afternoon of the 31st. Use a half-open range: `applied_at >= '2026-03-01' AND applied_at < '2026-04-01'`.
- LIKE is case-insensitive for English letters in SQLite and (by default collation) MySQL, but **case-sensitive in PostgreSQL** (ILIKE is the insensitive version there). The portable way: `LOWER(title) LIKE '%data%'`.

45.5 NULL and three-valued logic

Because NULL means “unknown”, SQL logic has three values: TRUE, FALSE and UNKNOWN. Is an unknown city equal to Delhi? Unknown. Is an unknown city *not* equal to Delhi? Also unknown. Is NULL equal to NULL? Two unknowns might or might not be the same, so: unknown. **WHERE keeps a row only when the condition is TRUE**; UNKNOWN rows are dropped, exactly like FALSE rows.

p	q	$p \text{ AND } q$	$p \text{ OR } q$
TRUE	UNKNOWN	UNKNOWN	TRUE
FALSE	UNKNOWN	FALSE	UNKNOWN
UNKNOWN	UNKNOWN	UNKNOWN	UNKNOWN

Expression	Result
<code>NULL = NULL</code>	UNKNOWN
<code>NULL <> 5</code>	UNKNOWN
<code>NULL + 5</code>	NULL
<code>NOT UNKNOWN</code>	UNKNOWN
<code>NULL IS NULL</code>	TRUE

Rule of thumb: FALSE AND anything is FALSE; TRUE OR anything is TRUE; otherwise UNKNOWN spreads.

So = NULL never finds anything. You must use `IS NULL` / `IS NOT NULL`:

```
SELECT name, city FROM candidates WHERE city = NULL;
```

Output (real run):

```

name | city
-----+-----
(0 rows)

```

```
SELECT name, city FROM candidates WHERE city IS NULL;
```

Output (real run):

```

name | city
-----+-----
Sana | NULL
(1 row)

```

The more dangerous version hides inside ordinary filters. There are 10 candidates and 3 live in Bengaluru. How many are “not in Bengaluru”?

```
SELECT COUNT(*) AS not_bengaluru
FROM candidates
WHERE city <> 'Bengaluru';
```

Output (real run):

```
not_bengaluru
-----
6
(1 row)
```

Six, not seven: Sana’s NULL city makes `city <> 'Bengaluru'` UNKNOWN, so she is dropped. If the business question is “everyone except Bengaluru”, write `WHERE city <> 'Bengaluru' OR city IS NULL`, or `WHERE COALESCE(city, '') <> 'Bengaluru'`.

45.5.1 COALESCE: replacing NULL

`COALESCE(a, b, c)` returns the first argument that is not NULL. It is SQL’s `fillna`.

```
SELECT name,
       COALESCE(city, 'Unknown') AS city,
       COALESCE(ctc_lpa, 0)      AS ctc_lpa_filled
FROM candidates
WHERE city IS NULL OR ctc_lpa IS NULL;
```

Output (real run):

```
name | city      | ctc_lpa_filled
-----+-----+-----
Kabir | Delhi     | 0
Sana  | Unknown   | 5.5
(2 rows)
```

In practice: is “fill with 0” right?

Filling Kabir’s salary with 0 is correct for “total payroll” but wrong for “average salary” (it pulls the average down) and wrong as an ML feature unless you also add a flag `ctc_missing` (Chapter 48 builds exactly that). Missingness here means “fresher”, which is information, the same lesson as the EDA chapters’ informative missingness.

45.6 ORDER BY, LIMIT and OFFSET: sorting and taking the top

`ORDER BY col` sorts ascending (ASC, the default); DESC sorts descending. You can sort by several columns: `ORDER BY title, city` sorts by title and breaks ties by city. `LIMIT n` keeps the first n rows; `OFFSET k` skips k rows first (used for pages: page 2 of size 3 is `LIMIT 3 OFFSET 3`).

```
SELECT name, ctc_lpa
FROM candidates
ORDER BY ctc_lpa DESC
LIMIT 3;
```

Output (real run):

```
name | ctc_lpa
-----+-----
```

```
Vikram | 42
Meera  | 30
Priya  | 24
(3 rows)
```

```
SELECT name, ctc_lpa
FROM candidates
ORDER BY ctc_lpa DESC
LIMIT 3 OFFSET 3;
```

Output (real run):

```
name | ctc_lpa
-----+-----
Diya  | 18
Ananya | 15
Rohan | 12
(3 rows)
```

Where does NULL go when sorting? It depends on the database:

```
SELECT name, ctc_lpa
FROM candidates
ORDER BY ctc_lpa ASC
LIMIT 3;
```

Output (real run):

```
name | ctc_lpa
-----+-----
Kabir | NULL
Sana  | 5.5
Aarav | 8
(3 rows)
```

SQLite and MySQL treat NULL as the smallest value (first in ASC, last in DESC). PostgreSQL and Oracle treat it as the largest (last in ASC). Standard SQL lets you say exactly what you want: `ORDER BY ctc_lpa ASC NULLS LAST`.

⚠ Common mistake: rows have no order unless you ask

A table is a *set* of rows: it has no built-in order. Without `ORDER BY`, `LIMIT 3` returns *some* 3 rows, which may change from run to run. “Top 3” always needs `ORDER BY`, and a tie-breaker (`ORDER BY ctc_lpa DESC, cand_id`) if ties are possible.

45.7 DISTINCT: unique values

`DISTINCT` removes duplicate result rows. With several columns it removes duplicate *combinations*.

```
SELECT DISTINCT title FROM jobs;
```

Output (real run):

```
title
-----
Data Scientist
ML Engineer
Data Analyst
(3 rows)
```

```
SELECT DISTINCT title, city FROM jobs ORDER BY title, city;
```

Output (real run):

title	city
Data Analyst	Bengaluru
Data Analyst	Mumbai
Data Analyst	Pune
Data Scientist	Bengaluru
Data Scientist	Mumbai
Data Scientist	Noida
ML Engineer	Bengaluru
ML Engineer	Hyderabad
ML Engineer	Noida

(9 rows)

Three distinct titles, but nine distinct (title, city) pairs. `DISTINCT` applies to the whole row, never to one column only.

45.8 Computed columns, aliases and integer division

`SELECT` can compute new columns with arithmetic (+ - * /) and name them with `AS` (an *alias*).

```
SELECT name,
       ctc_lpa,
       ctc_lpa * 100000 / 12 AS monthly_ctc_rs,
       ctc_lpa * 1.30      AS expected_ctc_lpa
FROM candidates
WHERE ctc_lpa IS NOT NULL
ORDER BY expected_ctc_lpa DESC
LIMIT 4;
```

Output (real run):

name	ctc_lpa	monthly_ctc_rs	expected_ctc_lpa
Vikram	42	350000	54.6
Meera	30	250000	39
Priya	24	200000	31.2
Diya	18	150000	23.4

(4 rows)

Meera earns 30 LPA, i.e. $30 \times 100000 / 12 = 250000$ rupees a month; a 30% hike would be 39 LPA. Now a trap that costs candidates points:

```
SELECT 7 / 2      AS int_div,
       7 * 1.0 / 2 AS real_div,
       ROUND(2.0 / 3, 3) AS rounded;
```

Output (real run):

int_div	real_div	rounded
3	3.5	0.667

(1 row)

In SQLite, PostgreSQL and SQL Server, **integer divided by integer is integer division** ($7/2 = 3$). A conversion rate computed as offers / applications with two integer counts returns 0 for every row with fewer offers than applications. Multiply by 1.0 (or `CAST(x AS DECIMAL)`) first. MySQL is the exception: its `/` always returns a decimal.

45.9 CASE WHEN: if-else inside a query

CASE checks conditions top to bottom and returns the value of the first one that is true; ELSE catches everything else (without ELSE you get NULL). It is how you bin a numeric column, exactly like `pd.cut`.

```
SELECT name, exp_yrs,
       CASE
         WHEN exp_yrs = 0 THEN 'Fresher'
         WHEN exp_yrs <= 3 THEN 'Junior'
         WHEN exp_yrs <= 6 THEN 'Mid'
         ELSE 'Senior'
       END AS band
FROM candidates
ORDER BY exp_yrs;
```

Output (real run):

name	exp_yrs	band
Kabir	0	Fresher
Sana	1	Junior
Aarav	2	Junior
Ishaan	2	Junior
Rohan	3	Junior
Ananya	4	Mid
Diya	5	Mid
Priya	6	Mid
Meera	7	Senior
Vikram	10	Senior

(10 rows)

Because the conditions are checked in order, the second branch `exp_yrs <= 3` does not need `exp_yrs > 0`: zero was already caught.

45.10 String and date functions

```
SELECT name,
       UPPER(name) AS upper_name,
       LENGTH(name) AS n_chars,
       SUBSTR(name, 1, 3) AS first3,
       name || ' (' || COALESCE(city, '?') || ')' AS label
FROM candidates
LIMIT 4;
```

Output (real run):

name	upper_name	n_chars	first3	label
Aarav	AARAV	5	Aar	Aarav (Bengaluru)
Diya	DIYA	4	Diy	Diya (Mumbai)
Kabir	KABIR	5	Kab	Kabir (Delhi)
Meera	MEERA	5	Mee	Meera (Bengaluru)

(4 rows)

`||` joins strings (`CONCAT(a, b)` in MySQL). `SUBSTR(s, start, length)` counts from 1, not 0.

```
SELECT job_id, posted_date,
       strftime('%Y-%m', posted_date) AS month,
       CAST(julianday('2026-03-15') - julianday(posted_date) AS INTEGER) AS days_open,
```

```

    date(posted_date, '+30 days') AS expires_on
FROM jobs
ORDER BY posted_date
LIMIT 4;

```

Output (real run):

```

job_id | posted_date | month | days_open | expires_on
-----+-----+-----+-----+-----
201    | 2026-01-15 | 2026-01 | 59         | 2026-02-14
202    | 2026-01-20 | 2026-01 | 54         | 2026-02-19
203    | 2026-01-25 | 2026-01 | 49         | 2026-02-24
204    | 2026-02-05 | 2026-02 | 38         | 2026-03-07
(4 rows)

```

`strftime('%Y-%m', d)` extracts year and month (the grouping key for monthly reports); `julianday` turns a date into a day number so two dates can be subtracted; `date(d, '+30 days')` adds days.


45.10.1 Dialect cheat-sheet

The core of SQL is standard. Dates, string joining and a few limits differ. Say in the interview which dialect you are writing; the interviewer will not mind the small differences.

Task	SQLite (this book)	MySQL 8	PostgreSQL
Year-month	<code>strftime('%Y-%m', d)</code>	<code>DATE_FORMAT(d, '%Y-%m')</code>	<code>TO_CHAR(d, 'YYYY-MM')</code> or <code>DATE_TRUNC('month', d)</code>
Days between	<code>julianday(a) - julianday(b)</code>	<code>DATEDIFF(a, b)</code>	<code>a - b (dates)</code>
Add 7 days	<code>date(d, '+7 days')</code>	<code>DATE_ADD(d, INTERVAL 7 DAY)</code>	<code>d + INTERVAL '7 days'</code>
Join strings	<code>a b</code>	<code>CONCAT(a, b)</code>	<code>a b</code> or <code>CONCAT</code>
First <i>n</i> rows	<code>LIMIT n</code>	<code>LIMIT n</code>	<code>LIMIT n (SQL Server: TOP n)</code>
7/2	3	3.5000	3
NULL in ASC sort	first	first	last
Case-insensitive LIKE	default	default	<code>ILIKE</code>

45.11 The logical order of a query

We *write* a query in the order SELECT, FROM, WHERE, GROUP BY, HAVING, ORDER BY, LIMIT. The database *evaluates* it (logically) in a different order. Knowing this order answers half of all “why does this error?” questions.



Consequences you can now explain from first principles:

- An alias defined in SELECT (step 5) cannot be used in WHERE (step 2), because it does not exist yet. It *can* be used in ORDER BY (step 7). (SQLite and MySQL are lenient and sometimes accept it; PostgreSQL and SQL Server reject it. Do not rely on it.)
- Aggregates like `COUNT(*)` cannot appear in WHERE (rows have not been grouped yet); filters on aggregates go in HAVING (step 4).

- Window functions are computed in step 5, so they cannot be used in WHERE either; wrap the query and filter outside (Chapter 47).
- ORDER BY can sort by a column you did not select (except after DISTINCT), and LIMIT is applied last.

Small example: walk one query through the eight steps

```
SELECT city, COUNT(*) AS n FROM candidates WHERE exp_yrs >= 2 GROUP BY city HAVING
COUNT(*) >= 2 ORDER BY n DESC;
```

1. FROM: all 10 candidates.
2. WHERE exp_yrs >= 2: Kabir (0) and Sana (1) removed; 8 left.
3. GROUP BY city: Bengaluru {Aarav, Meera, Ishaan}, Mumbai {Diya, Ananya}, Delhi {Priya}, Pune {Rohan}, Noida {Vikram}.
4. HAVING COUNT(*) >= 2: Bengaluru (3) and Mumbai (2) survive.
5. SELECT: (Bengaluru, 3), (Mumbai, 2).
6. DISTINCT: not used, nothing happens.
7. ORDER BY n DESC: Bengaluru first, then Mumbai. (8. no LIMIT.)

45.12 Interview questions: Level 1

Q45.1

✓ **Most asked** What is the order in which a SQL query is executed, and why can't I use a SELECT alias in WHERE?

🕒 **In one line.** Logically: FROM/JOIN, WHERE, GROUP BY, HAVING, SELECT, DISTINCT, ORDER BY, LIMIT; aliases are created in SELECT, which runs after WHERE, so WHERE cannot see them.

⚙️ **The full answer.** The written order (SELECT first) is for humans; the logical order is how the result is defined. FROM builds the set of rows (joining tables), WHERE removes rows, GROUP BY collapses the survivors into groups, HAVING removes groups, SELECT computes the output columns (including aliases and window functions), DISTINCT removes duplicate output rows, ORDER BY sorts, LIMIT cuts. Each step only sees what earlier steps produced. “Logical” matters: the optimiser may physically push filters earlier or use an index, but the *result* must be as if this order were followed.

📊 **Worked example.** SELECT ctc_lpa*1.3 AS expected FROM candidates WHERE expected > 20 fails in PostgreSQL (“column expected does not exist”). Fix by repeating the expression (WHERE ctc_lpa*1.3 > 20) or by wrapping: SELECT * FROM (SELECT ctc_lpa*1.3 AS expected FROM candidates) t WHERE expected > 20. In ORDER BY the alias works: ORDER BY expected DESC.

💬 **If the interviewer pushes.** Then why does my alias work in WHERE on MySQL? MySQL and SQLite add a non-standard convenience; portable SQL must not rely on it. Where do window functions fit? Step 5, which is why WHERE ROW_NUMBER() OVER(...) = 1 is illegal.

⚠️ **Trap to avoid.** Saying “SQL executes top to bottom as written”.

Q45.2

✓ **Most asked** Why does `WHERE city = NULL` return nothing, and how many rows does `WHERE city <> 'Bengaluru'` return in MiniNaukri?

🕒 **In one line.** Any comparison with NULL is UNKNOWN, and WHERE keeps only TRUE rows; so `= NULL` matches nothing, and `<> 'Bengaluru'` returns 6, not 7, because Sana's NULL city is silently dropped.

⚙️ **The full answer.** NULL means “value unknown”. Asking “is unknown equal to unknown?” cannot be answered, so SQL returns UNKNOWN (third truth value). WHERE treats UNKNOWN as “not true” and discards the row. `IS NULL` is a different operator that asks “is this value missing?”, which has a definite answer.

📊 **Worked example.** 10 candidates, 3 in Bengaluru, 1 with NULL city: `<> 'Bengaluru'` gives $10 - 3 - 1 = 6$ (query `L1_not_equal_null` above). To include the unknown city: `WHERE city <> 'Bengaluru' OR city IS NULL` gives 7.

💬 **If the interviewer pushes.** What does `COUNT(city)` return? 9: aggregate functions skip NULLs, only `COUNT(*)` counts rows. What is `NULL + 5`? NULL; arithmetic with unknown is unknown. Is there a NULL-safe equality? Yes: `IS NOT DISTINCT FROM` (standard, PostgreSQL), `<=>` (MySQL), `IS` (SQLite).

⚠️ **Trap to avoid.** Using `= NULL`, or forgetting that `NOT IN` with a NULL in the list returns nothing (Chapter 47).

Q45.3

✓ **Most asked** Primary key vs unique key vs foreign key: what is the difference?

🕒 **In one line.** A primary key uniquely identifies each row and cannot be NULL (one per table); a unique key enforces uniqueness but a table can have several and they may allow NULL; a foreign key references another table's key and enforces that the referenced row exists.

⚙️ **The full answer.** PK: the identity of the row; the database usually builds an index on it automatically, and other tables point at it. Unique key: a business rule (“no two candidates share an email”); several are allowed. FK: a link; it prevents orphans (an application for a job that does not exist) and can cascade deletes (`ON DELETE CASCADE`).

📊 **Worked example.** `candidates.cand_id` is the PK; `applications.cand_id` is an FK to it; `candidates.referred_by` is an FK to the *same* table (a self-reference), which is how referral chains are stored. In a skills table, `(cand_id, skill)` is a composite PK.

💬 **If the interviewer pushes.** *Surrogate vs natural key?* A surrogate key is a meaningless id (`cand_id = 101`); a natural key is real data (PAN number). Surrogates never change, natural keys sometimes do (people change phone numbers), so surrogates are preferred as PKs with a unique constraint on the natural key.

⚠️ **Trap to avoid.** Saying a unique key “is the same as a primary key”.

Q45.4

 Most asked DELETE vs TRUNCATE vs DROP?

 **In one line.** DELETE removes chosen rows (with WHERE) and can be rolled back; TRUNCATE removes all rows quickly and keeps the table; DROP removes the table itself, structure included.

 **The full answer.** DELETE is a row-by-row data operation (DML): it fires triggers, logs each row, honours WHERE. TRUNCATE is a structural operation (DDL in most databases): it deallocates the data pages, so it is much faster on big tables, resets auto-increment counters, has no WHERE, and in MySQL and Oracle cannot be rolled back (PostgreSQL can roll it back inside a transaction). DROP deletes data, indexes, constraints and the definition.

 **Worked example.** Removing the duplicate application: `DELETE FROM applications WHERE app_id = 22`. Clearing a staging table before each nightly load: `TRUNCATE TABLE staging_apps`. Retiring an old table: `DROP TABLE old_apps`.

 **If the interviewer pushes.** Which is fastest? $\text{DROP} \approx \text{TRUNCATE} \gg \text{DELETE}$ on large tables. Can TRUNCATE work if other tables reference this one by FK? Usually not without disabling or cascading the constraint.

 **Trap to avoid.** “TRUNCATE is just DELETE without WHERE” (it is implemented very differently).

Q45.5

Write a query: candidates from Bengaluru or Mumbai with at least 5 years of experience. Then explain what goes wrong without brackets.

 **In one line.** `WHERE (city = 'Bengaluru' OR city = 'Mumbai') AND exp_yrs >= 5`, or `WHERE city IN ('Bengaluru', 'Mumbai') AND exp_yrs >= 5`; without brackets AND binds first and every Bengaluru candidate is returned.

 **The full answer.** AND has higher precedence than OR, as $\times$ over $+$. `A OR B AND C` means `A OR (B AND C)`.

 **Worked example.** The unbracketed query returned Aarav, Diya, Meera, Ishaan; the correct one returns Diya and Meera (outputs `L1_and_or` and `L1_and_or_fixed`).

 **If the interviewer pushes.** Why prefer IN? Shorter, no precedence risk, and easy to extend. What about NOT? NOT binds tighter than AND; `NOT a AND b` is `(NOT a) AND b`.

 **Trap to avoid.** Testing on data where the bug happens not to show (if no Bengaluru candidate had fewer than 5 years you would never notice).

Q45.6

CHAR vs VARCHAR, and which numeric type should store salary or a rate?

 **In one line.** CHAR(n) is fixed length (padded with spaces), VARCHAR(n) variable length; money and rates should be DECIMAL/NUMERIC (exact), not FLOAT (approximate).

⚙️ **The full answer.** CHAR suits values that always have the same length (country code 'IN', variant 'A'); VARCHAR saves space for names and titles. FLOAT stores binary approximations: $0.1 + 0.2 = 0.30000000000000004$, so summing millions of salaries in FLOAT drifts and equality tests fail. DECIMAL(p,s) stores exact decimal digits. Use INTEGER for counts and ids, DATE/TIMESTAMP for time (never text, which sorts wrongly unless in ISO format), BOOLEAN for flags.

📊 **Worked example.** ctc_lpa DECIMAL(5,1) holds up to 9999.9 lakh with one decimal; variant CHAR(1).

💬 **If the interviewer pushes.** Why does ISO format ('2026-03-01') matter if dates are text? Because then alphabetical order equals time order; '01/03/2026' sorts wrongly.

⚠️ **Trap to avoid.** Storing dates as '1/3/26' strings.

Q45.7

Find all job titles that contain the word “data” in any case, and say how your answer changes across databases.

🕒 **In one line.** WHERE LOWER(title) LIKE '%data%' is portable; PostgreSQL also offers ILIKE.

⚙️ **The full answer.** % matches any run of characters. LIKE's case sensitivity depends on the database and its collation: SQLite and MySQL (default) ignore case, PostgreSQL does not. Lower-casing both sides removes the ambiguity.

📊 **Worked example.** In MiniNaukri, 7 of 10 jobs match (all Data Scientist and Data Analyst jobs, output L1_like). A leading % also prevents the database from using an ordinary index on title (Chapter 48): fine for 10 rows, slow for 50 million.

💬 **If the interviewer pushes.** How would you search job descriptions at Naukri scale? With a full-text index (MySQL FULLTEXT, PostgreSQL tsvector) or a search engine such as Elasticsearch, not LIKE.

⚠️ **Trap to avoid.** Forgetting that _ is a wildcard: LIKE 'ML_%' matches “MLE” too; escape it when you mean a literal underscore.

Q45.8

🧠 **Think deeper** SQL is called “declarative”. What does that mean, and why does it matter for a data scientist?

🕒 **In one line.** You state *what* result you want and the database chooses *how* to compute it, so the same query can be run by very different physical plans, and performance depends on the optimiser, indexes and data layout rather than on how you ordered your loops.

⚙️ **The full answer.** In pandas you write a sequence of operations and they run in that order. In SQL you describe the result set: “rows from these tables matching this condition, grouped this way”. The query optimiser estimates the cost of alternative plans (which table to read first, whether to use an index, nested-loop vs hash join) using table statistics and picks the cheapest. Consequences: (1) two queries that mean the same thing usually run equally fast; (2) you speed

queries up by giving the optimiser options (indexes, sargable filters, fewer columns) rather than by reordering clauses; (3) you inspect the chosen plan with EXPLAIN (Chapter 48).

Worked example. `SELECT * FROM applications WHERE cand_id = 105;` without an index the plan is SCAN applications (read all rows); after `CREATE INDEX` it is SEARCH ... USING INDEX (jump to the matching rows). The query text did not change.

If the interviewer pushes. *So is the logical order irrelevant?* No: it defines the meaning (what is legal, what each clause sees); the physical order is the optimiser's business. *When would you still push work into SQL instead of pandas?* When data is large: filtering and aggregating in the database moves megabytes instead of gigabytes over the network.

Trap to avoid. Claiming that putting the most selective condition first in WHERE makes the query faster (optimisers reorder predicates).

Q45.9

Think deeper A recruiter-analytics dashboard shows the average CTC of Mumbai candidates as 16.5 LPA. A colleague's Python script says 11 LPA for the same candidates. Both read the same table. What could explain it?

In one line. Most likely NULL handling: SQL's AVG ignores NULLs, while the script filled missing salaries with 0 (or the reverse); other suspects are integer division, duplicates, or different filters on NULL cities.

The full answer. AVG computes sum of non-NULL values / number of non-NULL values. `df['ctc'].fillna(0).mean()` divides by all rows. If 1 of 3 Mumbai candidates had NULL salary with the others at 18 and 15, SQL gives $33/2 = 16.5$; the script gives $33/3 = 11$. Neither is "wrong"; they answer different questions (average among those who reported vs average with unknown treated as zero, which is rarely sensible).

Worked example. The employees table shows the same effect: `AVG(salary) = 53` (10 known salaries), while `SUM(salary)/COUNT(*) = 48.18` (11 rows, Zoya unknown) (output L2_agg_basic in Chapter 46).

If the interviewer pushes. *How do you settle it?* Decide the definition with the business, write it down ("average over candidates with a reported CTC"), and report the coverage alongside ("2 of 3 candidates reported"). *Other causes?* A join that duplicated rows (fan-out, Chapter 46), or `city <> ...` filters that silently dropped NULL cities.

Trap to avoid. Picking one number without explaining the definitional difference.

Q45.10

What is the difference between COUNT(*), COUNT(col) and COUNT(DISTINCT col)? Give the values for MiniNaukri.

In one line. COUNT(*) counts rows; COUNT(col) counts rows where col is not NULL; COUNT(DISTINCT col) counts different non-NULL values.

The full answer. Aggregates skip NULLs, except COUNT(*), which does not look at any column. DISTINCT inside the function deduplicates before counting.

 **Worked example.** employees: `COUNT(*) = 11`, `COUNT(salary) = 10` (Zoya NULL). applications: `COUNT(*) = 23` applications, `COUNT(DISTINCT cand_id) = 8` candidates, `COUNT(DISTINCT job_id) = 9` jobs (output `L2_count_distinct` in Chapter 46).

 **If the interviewer pushes.** *Is `COUNT(1)` faster than `COUNT(*)`?* No; every modern optimiser treats them identically. *Count distinct pairs?* `COUNT(DISTINCT cand_id || '-' || job_id)` or count rows of a `SELECT DISTINCT` subquery.

 **Trap to avoid.** Using `COUNT(col)` to count rows when `col` has NULLs.

CHAPTER 46

SQL Level 2: Aggregation, GROUP BY and Joins

🎥 **Companion video for this chapter:** [Alex The Analyst: SQL Joins, GROUP BY and HAVING \(beginner series\)](#)

Level 1 read one table row by row. Real questions are about *groups* (“offer rate per job”) and about *several tables* (“which company made each offer”). This chapter adds the two tools for that: aggregation and joins. Both have exact rules for how many rows come out; most interview traps are about getting that row count wrong.

46.1 Aggregate functions: many rows in, one number out

An aggregate function collapses a column of many values into one: COUNT, SUM, AVG, MIN, MAX (most databases also have STDDEV, VARIANCE; PostgreSQL has PERCENTILE_CONT). The rule to remember: **every aggregate ignores NULL, except COUNT(*)**.

```
SELECT COUNT(*)      AS n_rows,
       COUNT(salary) AS n_salary,
       SUM(salary)    AS total,
       AVG(salary)    AS avg_ignores_null,
       SUM(salary) * 1.0 / COUNT(*) AS avg_null_as_zero,
       MIN(salary)    AS min_sal,
       MAX(salary)    AS max_sal
FROM employees;
```

Output (real run):

n_rows	n_salary	total	avg_ignores_null	avg_null_as_zero	min_sal	max_sal
11	10	530	53	48.18181818	25	90

(1 row)

Read it carefully. There are 11 employees but only 10 known salaries. $AVG = 530/10 = 53$ (Zoya ignored). If you wanted unknown salaries to count as zero (rarely sensible), you would get $530/11 = 48.18$. Two analysts who “computed the average salary” can disagree for this reason alone.

```
SELECT COUNT(*)      AS applications,
       COUNT(DISTINCT cand_id) AS candidates_who_applied,
       COUNT(DISTINCT job_id)  AS jobs_with_applications
FROM applications;
```

Output (real run):

applications	candidates_who_applied	jobs_with_applications
23	8	9

(1 row)

23 application rows, but only 8 different candidates and 9 different jobs. “How many candidates applied?” is 8, not 23: the difference between COUNT(*) and COUNT(DISTINCT ...) is the difference between events and entities.

46.2 GROUP BY: one result row per group

GROUP BY title splits the rows into piles, one pile per distinct title, then computes the aggregates *inside each pile*. It is pandas' `df.groupby('title').agg(...)`: split, apply, combine.

```
SELECT title,
       COUNT(*)      AS n_jobs,
       AVG(max_ctc_lpa) AS avg_max_ctc,
       MAX(max_ctc_lpa) AS best_max_ctc
FROM jobs
GROUP BY title
ORDER BY avg_max_ctc DESC;
```

Output (real run):

title	n_jobs	avg_max_ctc	best_max_ctc
ML Engineer	3	34	45
Data Scientist	4	25.75	40
Data Analyst	3	10	12

(3 rows)

Small example: do it by hand

Jobs by title: ML Engineer max CTCs {35, 45, 22} $\Rightarrow$ count 3, average $102/3 = 34$, max 45. Data Scientist {20, 25, 40, 18} $\Rightarrow$ 4, $103/4 = 25.75$, 40. Data Analyst {12, 8, 10} $\Rightarrow$ 3, 10, 12. That is exactly the output.

Grouping by two columns makes one group per *combination*:

```
SELECT city, title, COUNT(*) AS n_jobs
FROM jobs
GROUP BY city, title
ORDER BY city, title;
```

Output (real run):

city	title	n_jobs
Bengaluru	Data Analyst	1
Bengaluru	Data Scientist	2
Bengaluru	ML Engineer	1
Hyderabad	ML Engineer	1
Mumbai	Data Analyst	1
Mumbai	Data Scientist	1
Noida	Data Scientist	1
Noida	ML Engineer	1
Pune	Data Analyst	1

(9 rows)

Common mistake: the golden rule of GROUP BY

Every column in SELECT must either be in GROUP BY or be inside an aggregate. `SELECT dept, name, MAX(salary) FROM employees GROUP BY dept` is ambiguous: a department has several names, which one should be shown? PostgreSQL and SQL Server raise an error. SQLite and old MySQL silently return *some* name (SQLite happens to return the name on the max row, MySQL an arbitrary one). Never rely on it; to get “the employee with the top salary per department” use a window function (Chapter 47).

46.3 HAVING: filtering groups

WHERE filters rows *before* grouping; HAVING filters groups *after* aggregation (recall the logical order: WHERE is step 2, HAVING step 4). So a condition on COUNT(*) must be in HAVING.

```
SELECT cand_id, COUNT(*) AS n_apps
FROM applications
GROUP BY cand_id
HAVING COUNT(*) >= 3
ORDER BY n_apps DESC, cand_id;
```

Output (real run):

```
cand_id | n_apps
-----+-----
101     | 4
102     | 3
103     | 3
104     | 3
105     | 3
109     | 3
(6 rows)
```

Both can appear together: first WHERE keeps only rejected applications, then GROUP BY counts them per candidate, then HAVING keeps candidates with at least one.

```
SELECT cand_id, COUNT(*) AS n_rejections
FROM applications
WHERE status = 'rejected'
GROUP BY cand_id
HAVING COUNT(*) >= 1
ORDER BY cand_id;
```

Output (real run):

```
cand_id | n_rejections
-----+-----
102     | 1
103     | 1
105     | 1
107     | 1
109     | 1
(5 rows)
```

💡 Intuition

If a condition can be checked by looking at one row (status = 'rejected'), put it in WHERE: it shrinks the data before the expensive grouping. Only conditions that need the whole group ("at least 3 applications") belong in HAVING.

46.4 Conditional aggregation: the pattern behind every rate

To count only some rows inside each group, put a CASE inside the aggregate: SUM(CASE WHEN status = 'offer' THEN 1 ELSE 0 END) counts offers. The average of a 0/1 column is the proportion of ones, i.e. a rate. This one idea gives offer rates, click-through rates, conversion rates, churn rates.

```
SELECT job_id,
       COUNT(*) AS applicants,
       SUM(CASE WHEN status = 'offer' THEN 1 ELSE 0 END) AS offers,
```

```

SUM(CASE WHEN status = 'rejected' THEN 1 ELSE 0 END) AS rejections,
ROUND(AVG(CASE WHEN status = 'offer' THEN 1.0 ELSE 0 END), 2) AS offer_rate
FROM applications
GROUP BY job_id
ORDER BY job_id;

```

Output (real run):

job_id	applicants	offers	rejections	offer_rate
201	5	0	0	0
202	1	1	0	1
203	3	0	2	0
204	3	1	0	0.33
205	2	0	1	0
206	2	0	1	0
207	3	1	1	0.33
208	2	1	0	0.5
209	2	0	0	0

(9 rows)

```

SELECT status, COUNT(*) AS n
FROM applications
GROUP BY status
ORDER BY n DESC;

```

Output (real run):

status	n
rejected	5
interview	5
applied	5
shortlisted	4
offer	4

(5 rows)

Job 207 had 3 applications and 1 offer: offer rate $1/3 = 0.33$. Note the 1.0 inside the CASE: with integer 1 the AVG would still work in SQLite, but writing 1.0 protects you from integer division on every database. PostgreSQL also offers `COUNT(*) FILTER (WHERE status = 'offer')`.

46.5 Joins: combining tables

46.5.1 The idea, from zero

The jobs table stores `company_id = 2`, not the name “PayWave”. To show the name, SQL must *match* each job row with the company row that has the same id. A join does exactly this: for every pair of rows (one from each table) it checks the ON condition and keeps the pairs for which it is true.


Small example: a join by hand

Left table L: (job 201, company 1), (job 203, company 2), (job 999, company 7). Right table R: (company 1, DataNest), (company 2, PayWave), (company 6, GreenGrid). Condition: `L.company = R.company`.

- **INNER JOIN**: only matching pairs: (201, DataNest), (203, PayWave). Job 999 and GreenGrid disappear.
- **LEFT JOIN**: all inner pairs, plus every left row with no match, padded with NULL: adds

(999, NULL).

- **RIGHT JOIN:** all inner pairs plus unmatched right rows: adds (NULL, GreenGrid).
- **FULL OUTER JOIN:** inner pairs plus unmatched rows from both sides: 4 rows.
- **CROSS JOIN:** every pair, no condition: $3 \times 3 = 9$ rows.



The circles are a memory aid only. A join is *not* a set intersection of rows: if one left row matches three right rows, it appears three times. The real rule is “every matching pair”.

46.5.2 INNER JOIN

```
SELECT j.job_id, j.title, c.name AS company, c.industry
FROM jobs AS j
JOIN companies AS c ON c.company_id = j.company_id
ORDER BY j.job_id
LIMIT 5;
```

Output (real run):

job_id	title	company	industry
201	Data Scientist	DataNest	IT
202	ML Engineer	DataNest	IT
203	Data Analyst	PayWave	Fintech
204	Data Scientist	PayWave	Fintech
205	Data Analyst	MediCore	Healthcare

(5 rows)

`j` and `c` are *table aliases*, short names so that `j.company_id` says which table's column you mean (both tables have a `company_id` and a `city`, so you must qualify them). Joins chain: to describe each offer we need application → candidate, application → job → company.

```
SELECT a.app_id, ca.name AS candidate, j.title, co.name AS company, a.status
FROM applications a
JOIN candidates ca ON ca.cand_id = a.cand_id
JOIN jobs j ON j.job_id = a.job_id
JOIN companies co ON co.company_id = j.company_id
WHERE a.status = 'offer'
ORDER BY a.app_id;
```

Output (real run):

app_id	candidate	title	company	status
2	Aarav	Data Analyst	ShopKart	offer
5	Diya	Data Scientist	PayWave	offer
10	Meera	ML Engineer	DataNest	offer
15	Vikram	ML Engineer	EduSpark	offer

(4 rows)

46.5.3 LEFT JOIN, and counting things that may be zero

“How many jobs has each company posted?” must include GreenGrid with 0. INNER JOIN would drop it, because it has no matching job. LEFT JOIN keeps every company:

```
SELECT c.name AS company, COUNT(j.job_id) AS n_jobs
FROM companies c
LEFT JOIN jobs j ON j.company_id = c.company_id
GROUP BY c.company_id, c.name
ORDER BY n_jobs DESC, company;
```

Output (real run):

```
company | n_jobs
-----+-----
DataNest | 2
EduSpark | 2
MediCore | 2
PayWave  | 2
ShopKart | 2
GreenGrid | 0
(6 rows)
```

Why COUNT(j.job_id) and not COUNT(*)? For GreenGrid the LEFT JOIN produced one row whose job columns are all NULL. COUNT(*) counts that row:

```
SELECT c.name AS company, COUNT(*) AS wrong_n_jobs
FROM companies c
LEFT JOIN jobs j ON j.company_id = c.company_id
GROUP BY c.company_id, c.name
ORDER BY wrong_n_jobs DESC, company;
```

Output (real run):

```
company | wrong_n_jobs
-----+-----
DataNest | 2
EduSpark | 2
MediCore | 2
PayWave  | 2
ShopKart | 2
GreenGrid | 1
(6 rows)
```

GreenGrid shows 1 job it does not have. After a LEFT JOIN, always count a column from the right table.

46.5.4 Anti-join: rows with no match

“Candidates who never applied”: LEFT JOIN, then keep rows where the right side is NULL. (Chapter 47 shows the equivalent NOT EXISTS.)

```
SELECT ca.cand_id, ca.name
FROM candidates ca
LEFT JOIN applications a ON a.cand_id = ca.cand_id
WHERE a.app_id IS NULL;
```

Output (real run):

```
cand_id | name
-----+-----
106     | Sana
```

```
110 | Priya
(2 rows)
```

46.5.5 The WHERE-after-LEFT-JOIN trap

We want every candidate with their offer job, or NULL if they have no offer. Putting the status filter in WHERE quietly turns the LEFT JOIN into an INNER JOIN:

```
SELECT ca.name, a.job_id, a.status
FROM candidates ca
LEFT JOIN applications a ON a.cand_id = ca.cand_id
WHERE a.status = 'offer';
```

Output (real run):

```
name | job_id | status
-----+-----+-----
Aarav | 207    | offer
Diya  | 204    | offer
Meera  | 202    | offer
Vikram | 208    | offer
(4 rows)
```

Only 4 candidates came back. The LEFT JOIN did produce rows for Kabir, Rohan, ..., but with a.status NULL, and WHERE a.status = 'offer' is UNKNOWN for them, so WHERE removed them. Put conditions on the right table *inside* ON: they then decide which right rows match, not which result rows survive.

```
SELECT ca.name, a.job_id AS offer_job
FROM candidates ca
LEFT JOIN applications a
      ON a.cand_id = ca.cand_id AND a.status = 'offer'
ORDER BY ca.cand_id;
```

Output (real run):

```
name | offer_job
-----+-----
Aarav | 207
Diya  | 204
Kabir | NULL
Meera  | 202
Rohan  | NULL
Sana  | NULL
Vikram | 208
Ananya | NULL
Ishaan | NULL
Priya  | NULL
(10 rows)
```

46.5.6 FULL OUTER JOIN

Supply and demand by city: jobs per city vs candidates per city, keeping cities that appear on only one side.

```
WITH job_cities AS (
  SELECT city, COUNT(*) AS n_jobs FROM jobs GROUP BY city
),
cand_cities AS (
  SELECT city, COUNT(*) AS n_cands FROM candidates WHERE city IS NOT NULL GROUP BY city
```

```

)
SELECT COALESCE(j.city, c.city) AS city, j.n_jobs, c.n_cands
FROM job_cities j
FULL OUTER JOIN cand_cities c ON c.city = j.city
ORDER BY city;

```

Output (real run):

city	n_jobs	n_cands
Bengaluru	4	3
Delhi	NULL	2
Hyderabad	1	NULL
Mumbai	2	2
Noida	2	1
Pune	1	1

(6 rows)

Delhi has 2 candidates and no jobs (unmet supply); Hyderabad has a job and no candidates (unmet demand): both are business insights that an inner join would hide. `COALESCE(j.city, c.city)` picks whichever side has the city. MySQL has no `FULL JOIN`; emulate it with `LEFT JOIN ... UNION ... RIGHT JOIN`.

46.5.7 SELF JOIN: a table joined to itself

`employees.manager_id` points to another row of `employees`. Join the table to itself under two aliases, one playing “employee” and one playing “manager”:

```

SELECT e.name AS employee, e.salary, m.name AS manager, m.salary AS manager_salary
FROM employees e
LEFT JOIN employees m ON m.emp_id = e.manager_id
ORDER BY e.emp_id;

```

Output (real run):

employee	salary	manager	manager_salary
Neha	90	NULL	NULL
Arjun	60	Neha	90
Kavya	45	Arjun	60
Manav	45	Arjun	60
Tara	30	Arjun	60
Rahul	70	Neha	90
Sneha	50	Rahul	70
Yash	75	Rahul	70
Pooja	40	Neha	90
Dev	25	Pooja	40
Zoya	NULL	Pooja	40

(11 rows)

`LEFT JOIN` keeps Neha, who has no manager. The same pattern finds referral pairs among candidates:

```

SELECT c.name AS candidate, r.name AS referred_by
FROM candidates c
JOIN candidates r ON r.cand_id = c.referred_by
ORDER BY c.cand_id;

```

Output (real run):

candidate	referred_by
Diya	Aarav
Meera	Diya


```
Rohan      | Aarav
Vikram     | Meera
Ishaan     | Rohan
(5 rows)
```

46.5.8 CROSS JOIN and non-equi joins

A CROSS JOIN pairs every row with every row. It is useful to build a complete grid (every variant × every day), so that a later LEFT JOIN can show zeros for missing combinations:

```
SELECT v.variant, d.shown_date
FROM (SELECT DISTINCT variant FROM impressions) v
CROSS JOIN (SELECT DISTINCT shown_date FROM impressions) d
ORDER BY v.variant, d.shown_date;
```

Output (real run):

```
variant | shown_date
-----+-----
A       | 2026-03-01
A       | 2026-03-02
A       | 2026-03-03
B       | 2026-03-01
B       | 2026-03-02
B       | 2026-03-03
(6 rows)
```

The ON condition does not have to be equality. A *non-equi join* matches candidates to the jobs they are eligible for (enough experience, same city):

```
SELECT ca.name, j.job_id, j.title, j.min_exp
FROM candidates ca
JOIN jobs j
  ON ca.exp_yrs >= j.min_exp
 AND ca.city = j.city
ORDER BY ca.name, j.job_id;
```

Output (real run):

```
name | job_id | title           | min_exp
-----+-----+-----+-----
Aarav | 201    | Data Scientist | 2
Aarav | 207    | Data Analyst   | 1
Ananya | 203    | Data Analyst   | 1
Ananya | 204    | Data Scientist | 3
Diya   | 203    | Data Analyst   | 1
Diya   | 204    | Data Scientist | 3
Ishaan | 201    | Data Scientist | 2
Ishaan | 207    | Data Analyst   | 1
Meera  | 201    | Data Scientist | 2
Meera  | 202    | ML Engineer    | 4
Meera  | 206    | Data Scientist | 5
Meera  | 207    | Data Analyst   | 1
Rohan  | 205    | Data Analyst   | 0
Vikram | 208    | ML Engineer    | 6
Vikram | 209    | Data Scientist | 2
(15 rows)
```

This is how a simple rule-based job recommender is written in SQL.

46.5.9 How many rows does a join return? Fan-out

- 1:1 key match: at most one row per left row.
- 1:many (one job, many applications): each job row is repeated once per application.
- many:many (joining on a non-unique column on both sides): the row count can explode ($m \times n$ per matching key).

Repetition is correct for row-level output, but **dangerous for sums**. Total advertised max-CTC across jobs:

```
SELECT SUM(max_ctc_lpa) AS true_sum, COUNT(*) AS n_jobs FROM jobs;
```

Output (real run):

```
true_sum | n_jobs
-----+-----
235      | 10
(1 row)
```

The same SUM after joining applications:

```
SELECT SUM(j.max_ctc_lpa) AS sum_after_join, COUNT(*) AS rows_after_join
FROM jobs j
JOIN applications a ON a.job_id = j.job_id;
```

Output (real run):

```
sum_after_join | rows_after_join
-----+-----
498            | 23
(1 row)
```

498 instead of 235: job 201 (max CTC 20) has 5 applications, so its 20 was added 5 times. Jobs with no application (210) vanished. This is *fan-out*. Fix: aggregate each table to the grain you need *before* joining (Chapter 48 builds an ML feature table this way).

Remember: joins

Row count after a join = number of matching pairs (+ unmatched rows for outer joins). Before writing a join, ask “what is one row of each table, and is the join key unique on each side?” Count rows before and after; if a SUM suddenly grows, suspect fan-out.

46.6 Set operations: UNION, UNION ALL, INTERSECT, EXCEPT

Joins put tables *side by side*; set operations stack results *on top of each other*. Both queries must return the same number of columns with compatible types.

- UNION: stack and remove duplicates (needs a sort or hash, slower).
- UNION ALL: stack and keep everything (faster; use it unless you need deduplication).
- INTERSECT: rows in both. EXCEPT (Oracle: MINUS): rows in the first but not the second.

```
SELECT city FROM companies
UNION
SELECT city FROM candidates WHERE city IS NOT NULL
ORDER BY city;
```

Output (real run):

```
city
-----
Bengaluru
Delhi
Mumbai
Noida
Pune
(5 rows)
```

```
SELECT 'company' AS source, city FROM companies WHERE city = 'Bengaluru'
UNION ALL
SELECT 'candidate', city FROM candidates WHERE city = 'Bengaluru';
```

Output (real run):

```
source | city
-----+-----
company | Bengaluru
company | Bengaluru
candidate | Bengaluru
candidate | Bengaluru
candidate | Bengaluru
(5 rows)
```

```
SELECT city FROM candidates
EXCEPT
SELECT city FROM jobs;
```

Output (real run):

```
city
-----
NULL
Delhi
(2 rows)
```

Candidate cities with no jobs: Delhi, and NULL. Unlike `=`, set operations treat two NULLs as the same value, so the NULL city appears once.

46.7 Interview questions: Level 2

Q46.1

✓ **Most asked** Table A has values 1, 1, 2, NULL and table B has 1, 1, 3, NULL (one column x). How many rows do INNER, LEFT, RIGHT, FULL and CROSS joins on `A.x = B.x` return?

🕒 **In one line.** INNER 4, LEFT 6, RIGHT 6, FULL 8, CROSS 16.

⚙️ **The full answer.** A join returns every matching pair. The two 1s in A each match the two 1s in B: $2 \times 2 = 4$ pairs. 2 and 3 match nothing. NULL never equals NULL, so the NULLs match nothing. LEFT adds A's unmatched rows (2 and NULL): $4 + 2 = 6$. RIGHT adds B's unmatched (3 and NULL): 6. FULL adds both: $4 + 2 + 2 = 8$. CROSS: $4 \times 4 = 16$.

📊 **Worked example.** Verified by running it:

```
WITH a(x) AS (VALUES (1), (1), (2), (NULL)),
```

```

        b(x) AS (VALUES (1), (1), (3), (NULL))
SELECT 'inner' AS join_type, COUNT(*) AS n_rows FROM a INNER JOIN b ON a.x = b.x
UNION ALL
SELECT 'left', COUNT(*) FROM a LEFT JOIN b ON a.x = b.x
UNION ALL
SELECT 'right', COUNT(*) FROM a RIGHT JOIN b ON a.x = b.x
UNION ALL
SELECT 'full', COUNT(*) FROM a FULL JOIN b ON a.x = b.x
UNION ALL
SELECT 'cross', COUNT(*) FROM a CROSS JOIN b;

```

Output (real run):

```

join_type | n_rows
-----+-----
inner     | 4
left      | 6
right     | 6
full      | 8
cross     | 16
(5 rows)

```

 **If the interviewer pushes.** What if A had three 1s? INNER = $3 \times 2 = 6$. How do you make NULLs match? ON A.x IS NOT DISTINCT FROM B.x (or COALESCE(A.x, -1) = COALESCE(B.x, -1) if -1 cannot occur).

 **Trap to avoid.** Answering “INNER = 1 row because only the value 1 is common” (thinking in sets instead of pairs), or matching the NULLs.

Q46.2

 **Most asked** What is the difference between WHERE and HAVING? Can a query have both?

 **In one line.** WHERE filters individual rows before grouping; HAVING filters groups after aggregation; yes, a query can and often does have both.

 **The full answer.** Logical order: FROM → WHERE → GROUP BY → HAVING → SELECT. Aggregates do not exist at WHERE time, so WHERE COUNT(*) > 2 is an error. A row-level condition placed in HAVING works (on grouped columns) but wastes effort grouping rows that will be discarded.

 **Worked example.** L2_where_vs_having: WHERE keeps the 5 rejected applications, GROUP BY counts per candidate, HAVING keeps groups with at least one. “Candidates with at least 3 applications” needs HAVING COUNT(*) >= 3: 101 (4 rows, including the duplicate), 102, 103, 104, 105, 109.

 **If the interviewer pushes.** Can HAVING be used without GROUP BY? Yes; the whole table is then one group (SELECT COUNT(*) FROM jobs HAVING COUNT(*) > 5). Where do you filter on a window function? In an outer query, since windows are computed in SELECT, after HAVING.

 **Trap to avoid.** Saying “HAVING is just WHERE for GROUP BY queries” without the before/after distinction.

Q46.3

✓ **Most asked** Find candidates who never applied to any job, in three different ways.

🕒 **In one line.** LEFT JOIN ... WHERE right.key IS NULL; NOT EXISTS (correlated subquery); NOT IN (subquery), the last only if the subquery cannot return NULL.

⚙️ **The full answer.** All three express an anti-join. LEFT JOIN/IS NULL keeps unmatched left rows. NOT EXISTS checks, for each candidate, whether any application row exists. NOT IN compares against a list; if that list contains NULL, $x \text{ NOT IN } (\dots, \text{NULL})$ is never TRUE, and the query returns nothing.

📄 **Worked example.**

```
-- 1
SELECT c.cand_id, c.name FROM candidates c
LEFT JOIN applications a ON a.cand_id = c.cand_id WHERE a.app_id IS NULL;
-- 2
SELECT c.cand_id, c.name FROM candidates c
WHERE NOT EXISTS (SELECT 1 FROM applications a WHERE a.cand_id = c.cand_id);
-- 3 (safe here because applications.cand_id is NOT NULL)
SELECT cand_id, name FROM candidates
WHERE cand_id NOT IN (SELECT cand_id FROM applications);
```

All return Sana (106) and Priya (110). Third option with EXCEPT: SELECT cand_id FROM candidates EXCEPT SELECT cand_id FROM applications.

💡 **If the interviewer pushes.** Which is fastest? Modern optimisers turn NOT EXISTS and LEFT JOIN/IS NULL into the same anti-join plan; NOT EXISTS is the safest default. Chapter 47 shows NOT IN returning 0 rows on employees.manager_id, which contains a NULL.

⚠️ **Trap to avoid.** NOT IN against a nullable column.

Q46.4

✓ **Most asked** UNION vs UNION ALL: which is faster and when must you use each?

🕒 **In one line.** UNION ALL simply concatenates and is faster; UNION also removes duplicate rows, which costs a sort or hash; use UNION only when duplicates are wrong.

⚙️ **The full answer.** Deduplication must compare all rows, $O(n \log n)$ sort or a hash table. When the two parts cannot overlap (e.g. January data UNION February data), UNION ALL gives the same result for free. When stacking event logs from two sources where the same event might appear twice, UNION (or explicit dedup) is needed.

📄 **Worked example.** L2_union returns 5 distinct cities; L2_union_all returns all 5 Bengaluru rows (2 companies + 3 candidates).

💡 **If the interviewer pushes.** Do column names come from the first or second query? The first. Can you ORDER BY inside each part? Only the final result can be ordered (wrap a part in a subquery if you need LIMIT inside it).

⚠️ **Trap to avoid.** Using UNION by habit on large tables, or mixing column orders between the two parts (types may line up silently but the meaning is swapped).

Q46.5

Count jobs per company including companies with zero jobs. Why does `COUNT(*)` give the wrong answer?

🕒 **In one line.** `companies LEFT JOIN jobs ... GROUP BY company ... COUNT(j.job_id);` `COUNT(*)` counts the NULL-padded row a `LEFT JOIN` creates for a company with no jobs, so it shows 1 instead of 0.

⚙️ **The full answer.** A `LEFT JOIN` keeps an unmatched left row by padding right columns with `NULL`. That padded row is a real row of the result. `COUNT(j.job_id)` skips the `NULL`, giving 0.

📊 **Worked example.** Outputs `L2_left_join` (GreenGrid 0) and `L2_left_count_star` (GreenGrid 1).

💡 **If the interviewer pushes.** And `SUM` of a right-table column for GreenGrid? `NULL`, not 0; wrap with `COALESCE(SUM(...), 0)`.

⚠️ **Trap to avoid.** Using `INNER JOIN`, which drops GreenGrid entirely.

Q46.6

Your `LEFT JOIN` returns fewer rows than the left table has. What happened?

🕒 **In one line.** A `WHERE` condition on a right-table column filtered out the `NULL`-padded rows, effectively converting the `LEFT JOIN` into an `INNER JOIN`; move the condition into `ON`.

⚙️ **The full answer.** For unmatched rows the right columns are `NULL`, so `WHERE right.col = 'x'` is `UNKNOWN` and the row is dropped. Conditions in `ON` restrict which right rows can match, while still keeping every left row.

📊 **Worked example.** `L2_left_where_trap` returns 4 rows; `L2_left_on_filter` returns all 10 candidates with offer job or `NULL`.

💡 **If the interviewer pushes.** *Is it ever correct to filter the right table in `WHERE`?* Yes, `WHERE right.key IS NULL` for an anti-join, or when you genuinely want inner-join semantics (then write `INNER JOIN`, for clarity). *What about conditions on the left table?* Those belong in `WHERE`; in `ON` of a `LEFT JOIN` they would not remove left rows, which surprises people.

⚠️ **Trap to avoid.** Diagnosing “missing rows” as a data problem.

Q46.7

After joining jobs to applications, the total of `max_ctc_lpa` doubled. Explain and fix.

🕒 **In one line.** Fan-out: each job row is repeated once per matching application, so its value is summed several times; aggregate to the job grain first, or sum over the jobs table alone.

⚙️ **The full answer.** Joining a one-side table (jobs) to a many-side table (applications) produces one row per application. Row-level questions (“list applications with job info”) are fine; job-level sums are not.

📊 **Worked example.** True sum 235 over 10 jobs vs 498 after the join over 23 rows (outputs

L2_fanout_truth, L2_fanout); job 201 alone contributes $5 \times 20 = 100$ instead of 20.

 **If the interviewer pushes.** How do you detect fan-out in a pipeline? Assert row counts and key uniqueness before and after each join ($\text{COUNT}(*) = \text{COUNT}(\text{DISTINCT job_id})$). Is $\text{SUM}(\text{DISTINCT max_ctc_lpa})$ a fix? No: two different jobs with the same value would be counted once.

 **Trap to avoid.** “Fixing” it with DISTINCT.

Q46.8

Write a query for each company’s offer rate (offers / applications), keeping companies with no applications.

 **In one line.** LEFT JOIN companies → jobs → applications, then $1.0 * \text{SUM}(\text{CASE WHEN status='offer' THEN 1 ELSE 0 END}) / \text{NULLIF}(\text{COUNT}(\text{app_id}), 0)$.

 **The full answer.** Conditional aggregation counts offers; 1.0 avoids integer division; $\text{NULLIF}(x, 0)$ turns a zero denominator into NULL, so the rate is NULL (undefined) instead of a division-by-zero error. Both joins are LEFT so that GreenGrid survives. The duplicate application 22 is excluded inside the ON.

```
SELECT co.name AS company,
       COUNT(a.app_id) AS applications,
       SUM(CASE WHEN a.status = 'offer' THEN 1 ELSE 0 END) AS offers,
       ROUND(1.0 * SUM(CASE WHEN a.status = 'offer' THEN 1 ELSE 0 END)
             / NULLIF(COUNT(a.app_id), 0), 2) AS offer_rate
FROM companies co
LEFT JOIN jobs j      ON j.company_id = co.company_id
LEFT JOIN applications a ON a.job_id    = j.job_id AND a.app_id <> 22
GROUP BY co.company_id, co.name
ORDER BY offer_rate DESC, company;
```

Output (real run):

company	applications	offers	offer_rate
EduSpark	4	1	0.25
DataNest	5	1	0.2
ShopKart	5	1	0.2
PayWave	6	1	0.17
MediCore	2	0	0
GreenGrid	0	0	NULL

(6 rows)

 **If the interviewer pushes.** Is EduSpark (1 of 4) really better than PayWave (1 of 6)? With these counts, no evidence; a Beta-smoothed rate or a confidence interval (Chapter 38–39) is needed before ranking companies. Why put $\text{app_id} \neq 22$ in ON, not WHERE? In WHERE it would drop GreenGrid’s NULL row.

 **Trap to avoid.** Dividing two integer counts.

Q46.9

List each candidate’s applications to jobs outside their own city. What happens to Sana?

🕒 **In one line.** Join applications to candidates and jobs and keep `ca.city <> j.city`; Sana would be dropped by the NULL comparison anyway, but she has no applications.

```
SELECT ca.name, ca.city AS home_city, j.city AS job_city, j.title
FROM applications a
JOIN candidates ca ON ca.cand_id = a.cand_id
JOIN jobs j      ON j.job_id   = a.job_id
WHERE ca.city <> j.city
      AND a.app_id <> 22
ORDER BY ca.name, j.job_id;
```

Output (real run):

name	home_city	job_city	title
Aarav	Bengaluru	Noida	Data Scientist
Diya	Mumbai	Bengaluru	Data Scientist
Ishaan	Bengaluru	Noida	Data Scientist
Kabir	Delhi	Mumbai	Data Analyst
Kabir	Delhi	Pune	Data Analyst
Kabir	Delhi	Bengaluru	Data Analyst
Meera	Bengaluru	Noida	ML Engineer
Rohan	Pune	Bengaluru	Data Scientist
Rohan	Pune	Mumbai	Data Scientist
Vikram	Noida	Bengaluru	Data Scientist

(10 rows)

⚙️ **The full answer.** Ten out-of-city applications: a relocation-willingness signal that a recommender can use. A candidate with NULL city would vanish from this list because `NULL <> 'Mumbai'` is UNKNOWN; if such candidates existed you would have to decide whether “unknown city” counts as outside.

💡 **If the interviewer pushes.** How would you turn this into a feature? Per candidate: share of applications outside home city, via conditional aggregation.

⚠️ **Trap to avoid.** Not stating how NULL cities are treated.

Q46.10

💡 **Think deeper** The average of the department averages is 58.1, but the company-wide average salary is 53. Which is “the average salary”?

🕒 **In one line.** 53 is the average per employee; 58.1 is the average per department, which gives a one-person department (Leadership, 90) the same weight as a four-person one; averages of averages are only equal to the overall average when groups have equal sizes.

```
SELECT (SELECT AVG(salary) FROM employees) AS overall_avg,
       (SELECT AVG(dept_avg)
        FROM (SELECT AVG(salary) AS dept_avg FROM employees GROUP BY dept)) AS
       avg_of_dept_avgs;
```

Output (real run):

overall_avg	avg_of_dept_avgs
53	58.125

(1 row)

⚙️ **The full answer.** The overall mean is $\sum_g n_g \bar{x}_g / \sum_g n_g$, a size-weighted mean of group means. The plain average of group means gives each group weight 1. Here department sizes (known salaries) are 4, 3, 1, 2 with means 45, 65, 90, 32.5: weighted $(180 + 195 + 90 + 65)/10 = 53$; unweighted $(45 + 65 + 90 + 32.5)/4 = 58.125$.

💡 **If the interviewer pushes.** *Where does this bite in analytics?* Averaging daily CTRs instead of total clicks / total impressions; averaging city-level conversion rates for a national rate; computing “average order value” per store. Always ask what the unit of analysis is. *Related paradox?* Simpson’s paradox: subgroup trends reversing in the aggregate (Chapter 48, A/B by day).

⚠️ **Trap to avoid.** Averaging pre-aggregated rates in dashboards.

Q46.11

💡 **Think deeper** How does a database physically execute a join? When is each method chosen?

🕒 **In one line.** Three main algorithms: nested-loop (good when one side is tiny or the inner side has an index), hash join (build a hash table on the smaller side, probe with the larger; best for large equality joins), and sort-merge join (sort both on the key and walk them together; good when inputs are already sorted or for range-like joins).

⚙️ **The full answer.** Nested loop: for each outer row, find matches in the inner table, $O(n \cdot m)$ without an index, $O(n \log m)$ with one. Hash join: $O(n + m)$ time but needs memory for the hash table and only works for equality conditions. Sort-merge: $O(n \log n + m \log m)$, cheap if already sorted (clustered index), handles inequality joins. The optimiser picks using table statistics (row counts, distinct values); EXPLAIN shows the choice.

📊 **Worked example.** Joining 23 applications to 10 jobs on job_id, the PK: a nested loop with an index lookup per application is trivial. Joining a billion impression rows to 10 million jobs in Spark/Hive: a hash join, or a broadcast join if the jobs table fits in memory on every worker.

💡 **If the interviewer pushes.** *What is a broadcast join?* In distributed engines (Spark), the small table is copied to every machine so the big table need not be shuffled across the network. *Why can stale statistics slow a query?* The optimiser may think a table is tiny and choose a nested loop over millions of rows.

⚠️ **Trap to avoid.** Assuming the written join order is the execution order.

CHAPTER 47

SQL Level 3: Subqueries, CTEs, Window Functions and Changing Data

🎥 **Companion video for this chapter:** [techTFQ: SQL Window Functions explained \(RANK, DENSE_RANK, ROW_NUMBER, LAG, LEAD\)](#)

Level 2 answered “one number per group”. Level 3 answers questions that need *a query inside a query* (“candidates earning more than the average”) and questions that need *each row compared with its neighbours or its group* (“rank employees within department”, “days since the previous application”). These are the tools behind almost every medium and hard SQL interview problem.

47.1 Subqueries: a query inside a query

A subquery is a SELECT in brackets used where a value, a list or a table is expected.

47.1.1 Scalar subquery: returns one value

“Candidates earning above the average CTC.” First compute the average, then compare each row with it:

```
SELECT name, ctc_lpa
FROM candidates
WHERE ctc_lpa > (SELECT AVG(ctc_lpa) FROM candidates)
ORDER BY ctc_lpa DESC;
```

Output (real run):

```
name | ctc_lpa
-----+-----
Vikram | 42
Meera | 30
Priya | 24
(3 rows)
```

The inner query `SELECT AVG(ctc_lpa) FROM candidates` returns one number ($164/9 = 18.2$; Kabir’s NULL is ignored), and the outer query compares each candidate with it. You cannot write `WHERE ctc_lpa > AVG(ctc_lpa)` directly: an aggregate cannot appear in WHERE.

47.1.2 IN subquery: returns a list

```
SELECT name
FROM candidates
WHERE cand_id IN (SELECT cand_id FROM applications WHERE status = 'offer')
ORDER BY name;
```

Output (real run):

```
name
-----
Aarav
```

```
Diya
Meera
Vikram
(4 rows)
```

47.1.3 The NOT IN NULL trap

“Employees who manage nobody”: their `emp_id` does not appear in the `manager_id` column. The natural query returns **nothing**:

```
SELECT name FROM employees
WHERE emp_id NOT IN (SELECT manager_id FROM employees);
```

Output (real run):

```
name
----
(0 rows)
```

Why? Neha’s `manager_id` is NULL, so the list is (NULL, 1, 2, 2, 2, 1, 6, 6, 1, 9, 9). `x NOT IN (a, b, NULL)` means `x <> a AND x <> b AND x <> NULL`. The last part is UNKNOWN for every `x`, and TRUE AND UNKNOWN is UNKNOWN, so no row is ever TRUE. One NULL in the list kills the whole query, silently. `NOT EXISTS` does not have this problem, because it only asks whether a matching row exists:

```
SELECT e.name FROM employees e
WHERE NOT EXISTS (SELECT 1 FROM employees x WHERE x.manager_id = e.emp_id)
ORDER BY e.emp_id;
```

Output (real run):

```
name
-----
Kavya
Manav
Tara
Sneha
Yash
Dev
Zoya
(7 rows)
```

47.1.4 Correlated subquery: re-evaluated for each outer row

A correlated subquery refers to a column of the outer query, so conceptually it runs once per outer row. “Employees earning more than their own department’s average”:

```
SELECT e.name, e.dept, e.salary
FROM employees e
WHERE e.salary > (SELECT AVG(x.salary) FROM employees x WHERE x.dept = e.dept)
ORDER BY e.dept, e.salary DESC;
```

Output (real run):

```
name | dept       | salary
-----+-----+-----
Arjun | DataScience | 60
Yash  | Engineering | 75
Rahul | Engineering | 70
Pooja | Sales       | 40
(4 rows)
```

For Arjun (DataScience) the inner query computes the DataScience average (45); $60 > 45$, keep. For Sneha (Engineering, 50) it computes the Engineering average ($195/3 = 65$); $50 < 65$, drop. Correlated subqueries are easy to read but can be slow on big tables (one inner evaluation per row, unless the optimiser rewrites it as a join); the window-function version below avoids that.

47.1.5 Derived tables: a subquery in FROM

A subquery in FROM acts as a temporary table, which lets you group by a computed column:

```
SELECT band, COUNT(*) AS n, AVG(ctc_lpa) AS avg_ctc
FROM (
  SELECT ctc_lpa,
         CASE WHEN exp_yrs <= 3 THEN '0-3' WHEN exp_yrs <= 6 THEN '4-6' ELSE '7+' END AS
         band
  FROM candidates
) t
GROUP BY band
ORDER BY band;
```

Output (real run):

```
band | n | avg_ctc
-----+-----
0-3  | 5 | 8.75
4-6  | 3 | 19
7+   | 2 | 36
(3 rows)
```

47.2 CTEs: naming the steps of a query

A **Common Table Expression** (WITH name AS (...)) is a derived table with a name, written before the main query. Several CTEs can follow each other and each can use the previous ones, so a complicated query reads like the steps of a pandas notebook.

```
WITH apps_per_job AS (
  SELECT job_id, COUNT(*) AS n_apps
  FROM applications
  GROUP BY job_id
),
job_info AS (
  SELECT j.job_id, j.title, c.name AS company, COALESCE(p.n_apps, 0) AS n_apps
  FROM jobs j
  JOIN companies c ON c.company_id = j.company_id
  LEFT JOIN apps_per_job p ON p.job_id = j.job_id
)
SELECT * FROM job_info
WHERE n_apps >= (SELECT AVG(n_apps) FROM job_info)
ORDER BY n_apps DESC, job_id;
```

Output (real run):

```
job_id | title           | company | n_apps
-----+-----
201    | Data Scientist | DataNest | 5
203    | Data Analyst  | PayWave | 3
204    | Data Scientist | PayWave | 3
207    | Data Analyst  | ShopKart | 3
```

(4 rows)

Read top-down: count applications per job; attach title and company, with 0 for jobs that have none (job 210); keep jobs at or above the average of 2.3 applications per job (23/10).

💡 Intuition: subquery, CTE, view or temp table?

- **CTE:** readability; exists only for this one query; can be referenced several times; can be recursive.
- **Subquery:** fine for one small step; nesting three levels deep becomes unreadable.
- **View** (CREATE VIEW): a saved query that other queries and people reuse; it stores no data (a *materialised view* does, and must be refreshed).
- **Temporary table:** materialises an intermediate result for the session; useful when it is reused by many later queries or needs an index.

47.2.1 Recursive CTEs: hierarchies and date spines

A recursive CTE has two parts joined by UNION ALL: an *anchor* (the starting rows) and a *recursive part* that refers to the CTE itself and produces the next level from the previous one. It stops when the recursive part returns no new rows. Org chart from the top:

```
WITH RECURSIVE chain(emp_id, name, manager_id, level, path) AS (
  SELECT emp_id, name, manager_id, 0, name
  FROM employees WHERE manager_id IS NULL
  UNION ALL
  SELECT e.emp_id, e.name, e.manager_id, c.level + 1, c.path || ' > ' || e.name
  FROM employees e
  JOIN chain c ON e.manager_id = c.emp_id
)
SELECT level, name, path FROM chain ORDER BY path;
```

Output (real run):

level	name	path
0	Neha	Neha
1	Arjun	Neha > Arjun
2	Kavya	Neha > Arjun > Kavya
2	Manav	Neha > Arjun > Manav
2	Tara	Neha > Arjun > Tara
1	Pooja	Neha > Pooja
2	Dev	Neha > Pooja > Dev
2	Zoya	Neha > Pooja > Zoya
1	Rahul	Neha > Rahul
2	Sneha	Neha > Rahul > Sneha
2	Yash	Neha > Rahul > Yash

(11 rows)

Step by step: anchor = Neha (level 0). Iteration 1: employees whose manager is in the previous result: Arjun, Rahul, Pooja (level 1). Iteration 2: their reports (level 2). Iteration 3 finds nobody: stop.

The second use is a **date spine**: a table of every date, so that days with no activity show 0 instead of disappearing.

```
WITH RECURSIVE days(d) AS (
  SELECT '2026-03-01'
  UNION ALL
  SELECT date(d, '+1 day') FROM days WHERE d < '2026-03-07'
```

```
)
SELECT days.d AS day, COUNT(DISTINCT l.cand_id) AS dau
FROM days
LEFT JOIN logins l ON l.login_date = days.d
GROUP BY days.d
ORDER BY days.d;
```

Output (real run):

```
day          | dau
-----+-----
2026-03-01 | 3
2026-03-02 | 2
2026-03-03 | 4
2026-03-04 | 3
2026-03-05 | 3
2026-03-06 | 2
2026-03-07 | 0
(7 rows)
```

7 March had no logins. Without the spine, a GROUP BY on logins alone would simply omit that day, and a chart of DAU (daily active users) would draw a straight line across the gap. Note COUNT(DISTINCT l.cand_id): candidate 109 logged in twice on 4 March but is one active user.

47.3 Window functions: aggregates that keep every row

47.3.1 The idea

GROUP BY collapses each group into one row:

```
SELECT dept, AVG(salary) AS dept_avg
FROM employees
GROUP BY dept
ORDER BY dept;
```

Output (real run):

```
dept          | dept_avg
-----+-----
DataScience | 45
Engineering | 65
Leadership | 90
Sales | 32.5
(4 rows)
```

Often you want the group's number *next to every row* ("each employee's salary and their department's average"). A window function does that: it computes over a set of rows related to the current row (its *window*) and returns one value per row, without collapsing anything.

function(...) OVER (PARTITION BY *groups* ORDER BY *order within group frame*)

- PARTITION BY: splits rows into groups (like GROUP BY, but nothing is collapsed). Omit it and the whole result is one partition.
- ORDER BY: order inside each partition; needed for ranks, LAG/LEAD and running totals.
- *frame* (ROWS BETWEEN ...): which rows around the current row are included in a running aggregate.

```

SELECT name, dept, salary,
       DENSE_RANK() OVER (PARTITION BY dept ORDER BY salary DESC) AS rank_in_dept,
       AVG(salary) OVER (PARTITION BY dept) AS dept_avg,
       salary - AVG(salary) OVER (PARTITION BY dept) AS diff_from_avg
FROM employees
WHERE salary IS NOT NULL
ORDER BY dept, salary DESC, name;

```

Output (real run):

name	dept	salary	rank_in_dept	dept_avg	diff_from_avg
Arjun	DataScience	60	1	45	15
Kavya	DataScience	45	2	45	0
Manav	DataScience	45	2	45	0
Tara	DataScience	30	3	45	-15
Yash	Engineering	75	1	65	10
Rahul	Engineering	70	2	65	5
Sneha	Engineering	50	3	65	-15
Neha	Leadership	90	1	90	0
Pooja	Sales	40	1	32.5	7.5
Dev	Sales	25	2	32.5	-7.5

(10 rows)

Every employee keeps their row, and gets their department's average and their difference from it: Yash earns 10 above the Engineering average of 65. Pandas equivalent: `df.groupby('dept')['salary'].transform('mean')`.

47.3.2 Ranking: ROW_NUMBER, RANK, DENSE_RANK

The three differ only in how they handle ties (here Kavya and Manav both earn 45):

```

SELECT name, dept, salary,
       ROW_NUMBER() OVER (ORDER BY salary DESC) AS row_num,
       RANK() OVER (ORDER BY salary DESC) AS rnk,
       DENSE_RANK() OVER (ORDER BY salary DESC) AS dense_rnk
FROM employees
WHERE salary IS NOT NULL
ORDER BY salary DESC, name;

```

Output (real run):

name	dept	salary	row_num	rnk	dense_rnk
Neha	Leadership	90	1	1	1
Yash	Engineering	75	2	2	2
Rahul	Engineering	70	3	3	3
Arjun	DataScience	60	4	4	4
Sneha	Engineering	50	5	5	5
Kavya	DataScience	45	6	6	6
Manav	DataScience	45	7	6	6
Pooja	Sales	40	8	8	7
Tara	DataScience	30	9	9	8
Dev	Sales	25	10	10	9

(10 rows)

Function	Ties	Sequence for 50, 45, 45, 40
ROW_NUMBER()	Always different numbers; tied rows ordered arbitrarily (add a tie-breaker!)	5, 6, 7, 8
RANK()	Ties share a rank; the next rank <i>skips</i> (Olympic medals: two golds, no silver)	5, 6, 6, 8
DENSE_RANK()	Ties share a rank; no gaps	5, 6, 6, 7

Which to use: “the 3rd highest *salary value*” $\Rightarrow$ DENSE_RANK; “exactly one row per group” (deduplication, latest record) $\Rightarrow$ ROW_NUMBER with a deterministic ORDER BY; competition-style ranking $\Rightarrow$ RANK.

47.3.3 Window functions cannot go in WHERE

Window functions are computed in the SELECT step, after WHERE (Section 45.11). WHERE DENSE_RANK() OVER (...) <= 3 fails with “*misuse of window function DENSE_RANK()*” in SQLite (similar errors elsewhere). Compute the rank in a subquery or CTE, then filter outside:

```
SELECT name, salary FROM (
  SELECT name, salary, DENSE_RANK() OVER (ORDER BY salary DESC) AS dr
  FROM employees
) t
WHERE dr <= 3
ORDER BY salary DESC;
```

Output (real run):

```
name | salary
-----+-----
Neha | 90
Yash | 75
Rahul | 70
(3 rows)
```

(Snowflake, BigQuery and Teradata offer QUALIFY dr <= 3 as a shortcut.)

47.3.4 LAG and LEAD: looking at the previous and next row

LAG(col) returns the value of col from the previous row of the partition (NULL for the first row); LEAD(col) from the next row. They answer “time since previous event” questions:

```
SELECT cand_id, applied_date,
  LAG(applied_date) OVER (PARTITION BY cand_id ORDER BY applied_date, app_id) AS
prev_app,
  LEAD(applied_date) OVER (PARTITION BY cand_id ORDER BY applied_date, app_id) AS
next_app,
  CAST(julianday(applied_date)
    - julianday(LAG(applied_date) OVER (PARTITION BY cand_id ORDER BY applied_date,
    app_id))
  AS INTEGER) AS days_since_prev
FROM applications
WHERE cand_id IN (101, 102)
ORDER BY cand_id, applied_date, app_id;
```

Output (real run):

cand_id	applied_date	prev_app	next_app	days_since_prev
101	2026-01-16	NULL	2026-01-16	NULL
101	2026-01-16	2026-01-16	2026-02-19	0
101	2026-02-19	2026-01-16	2026-03-05	34
101	2026-03-05	2026-02-19	NULL	14
102	2026-01-18	NULL	2026-01-26	NULL
102	2026-01-26	2026-01-18	2026-02-06	8
102	2026-02-06	2026-01-26	NULL	11

(7 rows)

Aarav applied twice on 16 January (the duplicate row), then waited 34 days. `LAG(col, 2)` looks two rows back; `LAG(col, 1, 0)` gives 0 instead of NULL for the first row. Without window functions you would need a self-join that finds, for each row, the latest earlier row:

```
SELECT cur.cand_id, cur.login_date,
       MAX(prev.login_date) AS prev_login
FROM (SELECT DISTINCT cand_id, login_date FROM logins) cur
LEFT JOIN (SELECT DISTINCT cand_id, login_date FROM logins) prev
  ON prev.cand_id = cur.cand_id AND prev.login_date < cur.login_date
WHERE cur.cand_id = 101
GROUP BY cur.cand_id, cur.login_date
ORDER BY cur.login_date;
```

Output (real run):

cand_id	login_date	prev_login
101	2026-03-01	NULL
101	2026-03-02	2026-03-01
101	2026-03-03	2026-03-02
101	2026-03-05	2026-03-03
101	2026-03-06	2026-03-05

(5 rows)

Same answer, but it is quadratic in the number of rows per candidate: one reason window functions exist.

47.3.5 Running totals, moving averages and frames

With `ORDER BY` inside `OVER`, `SUM` becomes a running total. A frame `ROWS BETWEEN 2 PRECEDING AND CURRENT ROW` restricts it to the current row and the two before it: a 3-row moving average.

```
WITH daily AS (
  SELECT applied_date AS day, COUNT(*) AS n
  FROM applications
  WHERE applied_date >= '2026-02-01' AND applied_date < '2026-03-01'
  GROUP BY applied_date
)
SELECT day, n,
       SUM(n) OVER (ORDER BY day) AS running_total,
       ROUND(AVG(n) OVER (ORDER BY day ROWS BETWEEN 2 PRECEDING AND CURRENT ROW), 2) AS
       moving_avg_3
FROM daily
ORDER BY day;
```

Output (real run):

day	n	running_total	moving_avg_3
2026-02-02	1	1	1

```

2026-02-06 | 1 | 2          | 1
2026-02-07 | 1 | 3          | 1
2026-02-11 | 2 | 5          | 1.33
2026-02-13 | 1 | 6          | 1.33
2026-02-16 | 1 | 7          | 1.33
2026-02-19 | 2 | 9          | 1.33
2026-02-21 | 1 | 10         | 1.33
2026-02-22 | 1 | 11         | 1.33
(9 rows)

```

Check 11 February: running total $1 + 1 + 1 + 2 = 5$; moving average of the last three days' counts $(1 + 1 + 2)/3 = 1.33$.

⚠ Common mistake: the default frame, and LAST_VALUE

With ORDER BY and no explicit frame, the default frame is RANGE BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW: from the first row up to the current row *and its ties*. Hence LAST_VALUE returns the current row (or its last tie), not the partition's last row:

```

SELECT name, salary,
       LAST_VALUE(name) OVER (ORDER BY salary) AS last_default,
       LAST_VALUE(name) OVER (ORDER BY salary
                             ROWS BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED FOLLOWING) AS last_full
FROM employees
WHERE dept = 'DataScience'
ORDER BY salary;

```

Output (real run):

```

name | salary | last_default | last_full
-----+-----+-----+-----
Tara  | 30     | Tara        | Arjun
Kavya | 45     | Manav       | Arjun
Manav | 45     | Manav       | Arjun
Arjun | 60     | Arjun       | Arjun
(4 rows)

```

For Kavya, last_default is Manav (her tie), not Arjun. Write the full frame (ROWS BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED FOLLOWING) or use FIRST_VALUE with the opposite order. ROWS counts physical rows; RANGE groups rows with equal ORDER BY values; with ties in a running sum the two give different answers.

47.3.6 Share of total, buckets and percentiles

A window with an empty OVER () spans the whole result, so SUM(COUNT(*)) OVER () is the grand total of the grouped counts, which gives a percentage in one pass:

```

SELECT title, COUNT(*) AS n_jobs,
       ROUND(100.0 * COUNT(*) / SUM(COUNT(*)) OVER (), 1) AS pct_of_jobs
FROM jobs
GROUP BY title
ORDER BY n_jobs DESC;

```

Output (real run):

```

title          | n_jobs | pct_of_jobs
-----+-----+-----
Data Scientist | 4      | 40
Data Analyst   | 3      | 30
ML Engineer    | 3      | 30
(3 rows)

```

NTILE(4) splits ordered rows into 4 nearly equal buckets (quartiles); PERCENT_RANK = $(\text{rank} - 1) / (n - 1)$; FIRST_VALUE picks the first row of the window:

```
SELECT name, ctc_lpa,
       NTILE(4) OVER (ORDER BY ctc_lpa)           AS quartile,
       FIRST_VALUE(name) OVER (ORDER BY ctc_lpa DESC) AS top_earner,
       ROUND(PERCENT_RANK() OVER (ORDER BY ctc_lpa), 3) AS pct_rank
FROM candidates
WHERE ctc_lpa IS NOT NULL
ORDER BY ctc_lpa;
```

Output (real run):

name	ctc_lpa	quartile	top_earner	pct_rank
Sana	5.5	1	Vikram	0
Aarav	8	1	Vikram	0.125
Ishaan	9.5	1	Vikram	0.25
Rohan	12	2	Vikram	0.375
Ananya	15	2	Vikram	0.5
Diya	18	3	Vikram	0.625
Priya	24	3	Vikram	0.75
Meera	30	4	Vikram	0.875
Vikram	42	4	Vikram	1

(9 rows)

With 9 rows, NTILE(4) makes buckets of sizes 3, 2, 2, 2 (the first buckets get the extra rows).

47.4 Changing data: INSERT, UPDATE, DELETE and DDL

DML (data manipulation) changes rows; DDL (data definition) changes structure.

```
INSERT INTO applications (app_id, cand_id, job_id, applied_date, status)
VALUES (24, 110, 210, '2026-03-11', 'applied');
UPDATE applications SET status = 'shortlisted' WHERE app_id = 24;
DELETE FROM applications WHERE app_id = 22;
SELECT app_id, cand_id, job_id, status FROM applications WHERE app_id IN (1, 22, 24);
```

Output (real run):

app_id	cand_id	job_id	status
1	101	201	interview
24	110	210	shortlisted

(2 rows)

Row 24 was inserted then updated to shortlisted; the duplicate row 22 was deleted. **Always run the WHERE as a SELECT first:** an UPDATE or DELETE without WHERE changes every row.

```
CREATE TABLE skills (
  cand_id INTEGER NOT NULL REFERENCES candidates(cand_id),
  skill VARCHAR(20) NOT NULL,
  level INTEGER CHECK (level BETWEEN 1 AND 5),
  PRIMARY KEY (cand_id, skill)
);
INSERT INTO skills VALUES (101, 'python', 4), (101, 'sql', 3), (102, 'sql', 5);
SELECT * FROM skills;
```

Output (real run):

cand_id	skill	level
101	python	4
101	sql	3
102	sql	5

```

-----+-----+-----
101    | python | 4
101    | sql   | 3
102    | sql   | 5
(3 rows)

```

NOT NULL, CHECK, PRIMARY KEY and REFERENCES make the database reject bad data at the door: inserting level 7 or a second (101, python) row would fail. Constraints are cheaper than cleaning data afterwards.

47.4.1 Transactions and ACID

A **transaction** groups statements so that they succeed or fail together (BEGIN; ... COMMIT; or ROLLBACK;). Moving a candidate's application to "offer" and inserting an offer-letter row must both happen or neither. Databases guarantee **ACID**:

- **Atomicity**: all statements of a transaction or none.
- **Consistency**: a transaction takes the database from one valid state (all constraints satisfied) to another.
- **Isolation**: concurrent transactions do not see each other's half-finished work (levels: read uncommitted, read committed, repeatable read, serialisable).
- **Durability**: once committed, the change survives a crash (written to a log on disk).

47.5 Interview questions: Level 3

Q47.1

✓ **Most asked** Explain ROW_NUMBER, RANK and DENSE_RANK with salaries 90, 75, 70, 60, 50, 45, 45, 40.

🎯 **In one line.** ROW_NUMBER gives 1–8 with no ties; RANK gives 1, 2, 3, 4, 5, 6, 6, 8 (skips 7); DENSE_RANK gives 1, 2, 3, 4, 5, 6, 6, 7 (no gap).

⚙️ **The full answer.** All three number rows in the ORDER BY of the window. ROW_NUMBER ignores ties (and so is non-deterministic between tied rows unless you add a tie-breaker column). RANK = 1 + number of rows strictly better, so ties share a rank and the next rank jumps by the size of the tie. DENSE_RANK = 1 + number of distinct better values.

📊 **Worked example.** Output L3_rownum_rank: Kavya and Manav get row numbers 6 and 7, rank 6 and 6, dense rank 6 and 6; Pooja after them gets rank 8 and dense rank 7.

💡 **If the interviewer pushes.** Which for "3rd highest salary"? DENSE_RANK (values, not rows). Which for deduplication? ROW_NUMBER with ORDER BY app_id so that exactly one row gets 1. Which for "top 3 per department including ties"? DENSE_RANK ≤ 3 or RANK ≤ 3, depending on whether the business wants three values or three positions.

⚠️ **Trap to avoid.** Using ROW_NUMBER for Nth highest salary (ties give the wrong answer).

Q47.2

✓ **Most asked** NOT IN vs NOT EXISTS: why can NOT IN return zero rows?

🎯 **In one line.** If the subquery returns any NULL, $x \text{ NOT IN (list)}$ evaluates to UNKNOWN for every x , so nothing passes WHERE; NOT EXISTS only checks for matching rows and is NULL-safe.

⚙️ **The full answer.** $x \text{ NOT IN (a, b, NULL)}$ expands to $x \neq a \text{ AND } x \neq b \text{ AND } x \neq \text{NULL}$; the last term is UNKNOWN, so the conjunction is at best UNKNOWN. NOT EXISTS evaluates EXISTS (SELECT 1 ... WHERE inner.k = outer.k), which is TRUE or FALSE, never UNKNOWN.

📊 **Worked example.** “Employees who manage nobody”: NOT IN on manager_id returns 0 rows (Neha’s manager_id is NULL); NOT EXISTS returns the 7 correct employees (outputs L3_not_in_null, L3_not_exists). Fix for NOT IN: add WHERE manager_id IS NOT NULL inside the subquery.

💡 **If the interviewer pushes.** Does IN have the same problem? No: $x \text{ IN (a, NULL)}$ is TRUE when $x = a$; the NULL only turns “no match” from FALSE into UNKNOWN, which WHERE drops anyway. Performance? Most optimisers plan NOT EXISTS as an anti-join; NOT IN with a nullable column often cannot be.

⚠️ **Trap to avoid.** Believing the two are interchangeable.

Q47.3

✓ **Most asked** Find each employee’s salary next to their department’s average, and list those above it. Do it with a window and with a correlated subquery.

🎯 **In one line.** AVG(salary) OVER (PARTITION BY dept) in a CTE, then filter salary > dept_avg; or WHERE salary > (SELECT AVG(salary) FROM employees x WHERE x.dept = e.dept).

⚙️ **The full answer.** The window computes each department’s average once and attaches it to every row; the correlated subquery (logically) recomputes it per row.

📊 **Worked example.**

```
WITH t AS (
  SELECT name, dept, salary, AVG(salary) OVER (PARTITION BY dept) AS dept_avg
  FROM employees
)
SELECT name, dept, salary, dept_avg FROM t WHERE salary > dept_avg;
```

Both return Arjun (60 > 45), Yash (75 > 65), Rahul (70 > 65), Pooja (40 > 32.5) (output L3_correlated). Neha (90 = 90) is not above her one-person department’s average.

💡 **If the interviewer pushes.** What about Zoya (NULL salary)? $\text{NULL} > x$ is UNKNOWN, so she is never listed; AVG ignored her when computing Sales’ 32.5.

⚠️ **Trap to avoid.** WHERE salary > AVG(salary): aggregates are illegal in WHERE.

Q47.4

✓ **Most asked** Compute a running total and a 7-day moving average of daily applications. What changes if some days have no applications?

🕒 **In one line.** `SUM(n) OVER (ORDER BY day)` and `AVG(n) OVER (ORDER BY day ROWS BETWEEN 6 PRECEDING AND CURRENT ROW)` on a daily-count CTE; if days are missing, ROWS counts the last 7 rows, not the last 7 days, so first left-join to a date spine (filling 0) or use a RANGE frame on dates where supported.

⚙️ **The full answer.** A “7-row” window over a table that skips empty days stretches across more than 7 calendar days and averages only active days, overstating the average. A recursive CTE (or a calendar table) supplies every day.

📅 **Worked example.** `L3_running` uses a 3-row window over February: the dates jump (2nd, 6th, 7th, 11th, ...), so the “3-day” average on 11 February actually covers 6–11 February. With a date spine those missing days would contribute zeros.

💬 **If the interviewer pushes.** What does PostgreSQL offer? `RANGE BETWEEN INTERVAL '6 days' PRECEDING AND CURRENT ROW` on a date column. Running total that resets each month? Add `PARTITION BY strftime('%Y-%m', day)`.

⚠️ **Trap to avoid.** Calling a row-based window a time-based one.

Q47.5

What is a correlated subquery, and how would you rewrite one for speed?

🕒 **In one line.** A subquery that references the outer query’s columns and so is logically evaluated once per outer row; rewrite as a join to a pre-aggregated derived table/CTE or as a window function.

⚙️ **The full answer.** For n outer rows the naive cost is n inner evaluations. Pre-aggregating (one pass to compute all department averages, then one join) is $O(n)$. Many optimisers do this rewrite automatically (“subquery unnesting”), but not always, especially with complex correlations.

📅 **Worked example.**

```
SELECT e.name, e.salary, d.dept_avg
FROM employees e
JOIN (SELECT dept, AVG(salary) AS dept_avg FROM employees GROUP BY dept) d
ON d.dept = e.dept
WHERE e.salary > d.dept_avg;
```

💬 **If the interviewer pushes.** When is a correlated EXISTS efficient? When the inner table is indexed on the correlation column: each probe is an index lookup that stops at the first match.

⚠️ **Trap to avoid.** Assuming every subquery is correlated (a scalar subquery with no outer reference runs once).

Q47.6

CTE vs subquery vs temporary table vs view: when do you use which?

🎯 **In one line.** CTEs for readable multi-step logic in one query (and for recursion); subqueries for small inline steps; temporary tables to materialise and reuse an intermediate result within a session; views to share a saved query definition; materialised views to cache an expensive result.

⚙️ **The full answer.** CTEs are usually inlined by the optimiser (PostgreSQL 12+, most engines), so they cost nothing; some engines materialise them, which can help or hurt. A temp table can be indexed and its statistics help the optimiser; it lives until the session ends. A view stores only the query text; every use reruns it. A materialised view stores the result and must be refreshed, a common pattern for dashboards (“daily applications per company”).

📊 **Worked example.** `L3_cte` names two steps, `apps_per_job` and `job_info`, and references `job_info` twice (once for the average), which a subquery would have had to repeat.

💡 **If the interviewer pushes.** *Can a CTE be referenced by a later query?* No, only by the statement it belongs to.

⚠️ **Trap to avoid.** Saying CTEs are always faster than subqueries.

Q47.7

What are ACID properties? Give a job-portal example of each.

🎯 **In one line.** Atomicity (all or nothing), Consistency (constraints hold before and after), Isolation (concurrent transactions do not see each other’s partial work), Durability (committed data survives crashes).

⚙️ **The full answer.** Atomicity: making an offer updates `applications.status` and inserts an offers row; if the second fails, the first is rolled back. Consistency: a CHECK that status is one of five values, an FK that the job exists. Isolation: two recruiters closing the last opening at the same moment must not both succeed (serialisable isolation or row locks). Durability: after COMMIT, a power cut does not lose the offer (write-ahead log).

💡 **If the interviewer pushes.** *What anomalies do weak isolation levels allow?* Dirty reads (seeing uncommitted data), non-repeatable reads (a row changes between two reads), phantom reads (new rows appear in a repeated range query). *Do analytics warehouses need full ACID?* Less so; they favour throughput and often use eventual consistency, while the OLTP database behind the product needs it fully.

⚠️ **Trap to avoid.** Listing the four words without an example.

Q47.8

💡 Think deeper Why does `LAST_VALUE(name) OVER (ORDER BY salary)` not return the highest-paid name?

🎯 **In one line.** Because with ORDER BY and no frame the default frame ends at the current row (and its ties), so “last value” means “the current row”, not the partition’s last row.

⚙️ **The full answer.** The default frame is RANGE BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW. This default exists so that `SUM(x) OVER (ORDER BY t)` is a running total. `FIRST_VALUE` looks at the frame's first row, which is the partition's first row, so it behaves as people expect; `LAST_VALUE` does not.

📊 **Worked example.** `L3_frame_trap`: for Tara the default returns Tara; for Kavya it returns Manav (her tie, because RANGE includes peers); with `ROWS BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED FOLLOWING` every row gets Arjun.

💡 **If the interviewer pushes.** *Simpler fix?* `FIRST_VALUE(name) OVER (ORDER BY salary DESC)`. When do ROWS and RANGE differ in a running sum? With ties in the ORDER BY: RANGE adds all tied rows at once, ROWS adds them one by one.

⚠️ **Trap to avoid.** Not knowing frames exist.

Q47.9

💡 **Think deeper** Rewrite “days since the previous login” without window functions. What does that teach about why window functions exist?

🎯 **In one line.** Self-join each login to all earlier logins of the same candidate and take the MAX earlier date; it works, but it compares every pair of rows per candidate ($O(k^2)$), whereas LAG reads the sorted rows once.

```
SELECT cur.cand_id, cur.login_date,
       MAX(prev.login_date) AS prev_login
FROM (SELECT DISTINCT cand_id, login_date FROM logins) cur
LEFT JOIN (SELECT DISTINCT cand_id, login_date FROM logins) prev
  ON prev.cand_id = cur.cand_id AND prev.login_date < cur.login_date
WHERE cur.cand_id = 101
GROUP BY cur.cand_id, cur.login_date
ORDER BY cur.login_date;
```

Output (real run):

cand_id	login_date	prev_login
101	2026-03-01	NULL
101	2026-03-02	2026-03-01
101	2026-03-03	2026-03-02
101	2026-03-05	2026-03-03
101	2026-03-06	2026-03-05

(5 rows)

⚙️ **The full answer.** The self-join produces, for a candidate with k logins, up to $k(k-1)/2$ pairs before grouping. Window functions sort each partition once and then stream through it, $O(k \log k)$. They also express “previous row” directly, so the query is shorter and less error-prone.

💡 **If the interviewer pushes.** *Older MySQL (5.7) has no window functions; what did people do?* Self-joins, correlated subqueries, or user variables. *How would pandas do it?* `df.sort_values(['cand_id', 'login_date']).groupby('cand_id')['login_date'].shift()`.

⚠️ **Trap to avoid.** Forgetting to deduplicate logins first (two logins on the same day would produce a 0-day gap).

Q47.10

 **Think deeper** A data scientist wrote `UPDATE applications SET status = 'rejected' WHERE job_id IN (SELECT job_id FROM jobs WHERE posted_date < '2026-02-01')` on production at 6 pm. What safeguards should have been in place?

 **In one line.** Run the WHERE as a SELECT and check the row count; wrap the UPDATE in a transaction and verify before COMMIT; have a backup or point-in-time recovery; and ideally not have write access to production at all.

 **The full answer.** An UPDATE is irreversible after commit without a backup. Analysts should work on replicas or warehouse copies. In a transaction (BEGIN; UPDATE ...; SELECT COUNT(*) ...; COMMIT/ROLLBACK) you can inspect the effect first. Auditing (who changed what) and soft deletes (a deleted_at column) make mistakes recoverable.

 **Worked example.** The subquery selects jobs 201, 202, 203 (posted in January). The SELECT version of the WHERE shows 9 application rows: 201's five (including the duplicate 22), 202's one, 203's three. Only after confirming that "9" is what the business intended should the UPDATE run.

 **If the interviewer pushes.** *What is a soft delete?* Marking a row as deleted instead of removing it, so history and joins remain intact.

 **Trap to avoid.** Treating SQL on production like SQL in a notebook.

CHAPTER 48

SQL Level 4: Interview Patterns, SQL for Machine Learning, Performance and Design

🎥 **Companion video for this chapter:** [techTFQ: Solving real SQL interview queries step by step](#)

This chapter is a catalogue of the patterns that most SQL interview problems are built from. For each: the problem in words, the idea in one sentence, the query, the real output, and how to explain it aloud. Then the data-scientist-specific part: building leak-free ML feature tables in SQL. Then the engineering questions every interviewer asks at the end: indexes, query plans, normalisation, OLTP vs OLAP.

☰ **Remember: how to answer any SQL problem in an interview**

1. **Restate the grain:** “one output row per *what?*” (per candidate, per day, per department).
2. **Ask about edge cases:** ties, NULLs, duplicates, days with no data, what to return if nothing qualifies.
3. **Build in steps** with CTEs; say each step aloud.
4. **Check with a tiny example:** walk two or three rows through your query.
5. **Mention performance** only after correctness (indexes on join/filter keys).

48.1 Nth highest value

Salaries (known): 90, 75, 70, 60, 50, 45, 45, 40, 30, 25. Third highest distinct value = 70.

Method 1: DENSE_RANK (best answer). Handles ties correctly and generalises to “per department”.

```
SELECT DISTINCT salary
FROM (SELECT salary, DENSE_RANK() OVER (ORDER BY salary DESC) AS dr
      FROM employees WHERE salary IS NOT NULL) t
WHERE dr = 3;
```

Output (real run):

```
salary
-----
70
(1 row)
```

Method 2: DISTINCT + LIMIT/OFFSET. Skip $N - 1$ distinct values. Returns no row (not NULL) if fewer than N values exist.

```
SELECT DISTINCT salary
FROM employees
WHERE salary IS NOT NULL
ORDER BY salary DESC
LIMIT 1 OFFSET 2;
```

Output (real run):

```
salary
-----
70
(1 row)
```

Method 3: correlated count (works on any SQL). A salary is the 3rd highest if exactly 2 distinct salaries are greater.

```
SELECT DISTINCT e.salary
FROM employees e
WHERE 2 = (SELECT COUNT(DISTINCT x.salary) FROM employees x WHERE x.salary > e.salary);
```

Output (real run):

```
salary
-----
70
(1 row)
```

Second highest with MAX only. The largest value below the maximum:

```
SELECT MAX(salary) AS second_highest
FROM employees
WHERE salary < (SELECT MAX(salary) FROM employees);
```

Output (real run):

```
second_highest
-----
75
(1 row)
```

⚠ Common mistake

ORDER BY salary DESC LIMIT 1 OFFSET 2 *without* DISTINCT returns the 3rd *row*, which is wrong when the top salaries tie. ROW_NUMBER has the same flaw. Also say what happens when there are fewer than N distinct salaries (method 1 returns no row; wrap in SELECT (subquery) AS nth to return NULL instead, which is what LeetCode-style graders expect).

48.2 Top N per group

“Top 2 salaries in each department.” Rank inside each partition, then filter outside:

```
WITH ranked AS (
  SELECT name, dept, salary,
         DENSE_RANK() OVER (PARTITION BY dept ORDER BY salary DESC) AS dr
  FROM employees
  WHERE salary IS NOT NULL
)
SELECT dept, name, salary
FROM ranked
WHERE dr <= 2
ORDER BY dept, salary DESC, name;
```

Output (real run):

dept	name	salary
DataScience	Arjun	60
DataScience	Kavya	45
DataScience	Manav	45
Engineering	Yash	75
Engineering	Rahul	70
Leadership	Neha	90
Sales	Pooja	40
Sales	Dev	25

(8 rows)

DataScience returns three people because Kavya and Manav tie for the second-highest salary. If the business wants exactly two people, use ROW_NUMBER with a tie-breaker (ORDER BY salary DESC, hire_date) and say so. The same pattern gives “latest application per candidate”, “most recent salary per employee”, “top 3 jobs by applications in each city”.

48.3 Duplicates: find, inspect, delete

Decide first what “duplicate” means: same *business key* (candidate, job, date), not the same id.

```
SELECT cand_id, job_id, applied_date, COUNT(*) AS copies
FROM applications
GROUP BY cand_id, job_id, applied_date
HAVING COUNT(*) > 1;
```

Output (real run):

cand_id	job_id	applied_date	copies
101	201	2026-01-16	2

(1 row)

ROW_NUMBER labels the copies; every row with rn > 1 is a duplicate to remove:

```
WITH r AS (
  SELECT app_id, cand_id, job_id,
         ROW_NUMBER() OVER (PARTITION BY cand_id, job_id, applied_date ORDER BY app_id) AS
         rn
  FROM applications
)
SELECT app_id, cand_id, job_id, rn FROM r WHERE cand_id = 101 ORDER BY app_id;
```

Output (real run):

app_id	cand_id	job_id	rn
1	101	201	1
2	101	207	1
3	101	209	1
22	101	201	2

(4 rows)

Delete everything except the smallest id in each duplicate group:

```
DELETE FROM applications
WHERE app_id NOT IN (
  SELECT MIN(app_id) FROM applications
  GROUP BY cand_id, job_id, applied_date
);
```

```
SELECT COUNT(*) AS rows_left FROM applications;
```

Output (real run):

```
rows_left
-----
22
(1 row)
```

23 rows became 22. In analytics you often do not delete at all: you deduplicate in a CTE (WHERE rn = 1) at the start of every query, keeping the raw table untouched.

48.4 Self-join classics: employee vs manager

```
SELECT e.name AS employee, e.salary, m.name AS manager, m.salary AS manager_salary
FROM employees e
JOIN employees m ON m.emp_id = e.manager_id
WHERE e.salary > m.salary;
```

Output (real run):

```
employee | salary | manager | manager_salary
-----+-----+-----+-----
Yash      | 75     | Rahul   | 70
(1 row)
```

Yash (75) earns more than his manager Rahul (70). Variants: “employees hired before their manager” (compare hire dates), “managers with at least 3 reports” (GROUP BY manager_id HAVING COUNT(*) >= 3: Neha and Arjun, 3 reports each).

48.5 Consecutive days: the gaps-and-islands trick

“Candidates who logged in on at least 3 consecutive days.” The trick: number each candidate’s distinct login dates 1, 2, 3, ... and subtract that number of days from the date. Consecutive dates move up by one day while the row number also moves up by one, so the difference stays constant; a gap breaks the run and changes the difference. Each constant value is one *island* (streak).

```
WITH d AS (SELECT DISTINCT cand_id, login_date FROM logins WHERE cand_id = 101)
SELECT login_date,
       ROW_NUMBER() OVER (ORDER BY login_date) AS rn,
       date(login_date, '-' || ROW_NUMBER() OVER (ORDER BY login_date) || ' days') AS
       grp_date
FROM d ORDER BY login_date;
```

Output (real run):

```
login_date | rn | grp_date
-----+-----+-----
2026-03-01 | 1  | 2026-02-28
2026-03-02 | 2  | 2026-02-28
2026-03-03 | 3  | 2026-02-28
2026-03-05 | 4  | 2026-03-01
2026-03-06 | 5  | 2026-03-01
(5 rows)
```

Aarav’s 1–3 March all map to 28 February (one island of length 3); 5–6 March map to 1 March (a new island of length 2). Group by that value and count:

```

WITH d AS (SELECT DISTINCT cand_id, login_date FROM logins),
g AS (
  SELECT cand_id, login_date,
         julianday(login_date)
         - ROW_NUMBER() OVER (PARTITION BY cand_id ORDER BY login_date) AS grp
  FROM d
)
SELECT cand_id, MIN(login_date) AS streak_start, MAX(login_date) AS streak_end,
       COUNT(*) AS streak_len
FROM g
GROUP BY cand_id, grp
HAVING COUNT(*) >= 3
ORDER BY streak_len DESC, cand_id;

```

Output (real run):

cand_id	streak_start	streak_end	streak_len
104	2026-03-02	2026-03-05	4
101	2026-03-01	2026-03-03	3
107	2026-03-04	2026-03-06	3

(3 rows)

Note the first CTE: SELECT DISTINCT removes candidate 109's double login on 4 March. Without it the row numbers would advance twice on one date and break the islands.

48.6 Daily, rolling and cohort metrics

48.6.1 Daily active users and rolling active users

Daily active users with every day shown is the date-spine query from Chapter 47 (L3_recursive_dates). "Active in the last 3 days" for each day uses a range join:

```

WITH d AS (SELECT DISTINCT cand_id, login_date FROM logins),
days AS (SELECT DISTINCT login_date AS day FROM logins)
SELECT days.day,
       COUNT(DISTINCT d.cand_id) AS users_last_3_days
FROM days
JOIN d ON d.login_date BETWEEN date(days.day, '-2 days') AND days.day
GROUP BY days.day
ORDER BY days.day;

```

Output (real run):

day	users_last_3_days
2026-03-01	3
2026-03-02	4
2026-03-03	5
2026-03-04	5
2026-03-05	5
2026-03-06	4

(6 rows)

On 3 March, users active on 1–3 March: 101, 102, 104, 105, 109 = 5.

48.6.2 Cohort retention

Group candidates by signup month (their *cohort*) and ask what fraction were active in the first week of March:

```
WITH cohort AS (
  SELECT cand_id, strftime('%Y-%m', signup_date) AS signup_month FROM candidates
),
active AS (
  SELECT DISTINCT cand_id FROM logins
  WHERE login_date BETWEEN '2026-03-01' AND '2026-03-07'
)
SELECT c.signup_month,
       COUNT(*) AS signed_up,
       COUNT(a.cand_id) AS active_first_week_march,
       ROUND(1.0 * COUNT(a.cand_id) / COUNT(*), 2) AS retention
FROM cohort c
LEFT JOIN active a ON a.cand_id = c.cand_id
GROUP BY c.signup_month
ORDER BY c.signup_month;
```

Output (real run):

signup_month	signed_up	active_first_week_march	retention
2026-01	4	3	0.75
2026-02	4	2	0.5
2026-03	2	1	0.5

(3 rows)

January signups (101–104): 101, 102 and 104 logged in, retention $3/4 = 0.75$. Real retention tables put cohorts on rows and “months since signup” on columns; the logic is the same, with `months_since = month(login) - month(signup)` as the column key.

48.6.3 Month-over-month growth

```
WITH monthly AS (
  SELECT strftime('%Y-%m', applied_date) AS month, COUNT(*) AS n_apps
  FROM applications
  GROUP BY 1
)
SELECT month, n_apps,
       LAG(n_apps) OVER (ORDER BY month) AS prev_month,
       ROUND(100.0 * (n_apps - LAG(n_apps) OVER (ORDER BY month))
            / LAG(n_apps) OVER (ORDER BY month), 1) AS mom_growth_pct
FROM monthly
ORDER BY month;
```

Output (real run):

month	n_apps	prev_month	mom_growth_pct
2026-01	6	NULL	NULL
2026-02	11	6	83.3
2026-03	6	11	-45.5

(3 rows)

February: $(11 - 6)/6 = +83.3\%$; March: $(6 - 11)/11 = -45.5\%$. (March is only partial data: say so before anyone panics. Compare like with like, e.g. month-to-date vs the same days of last month.)

48.7 Funnels and A/B tests

48.7.1 Hiring funnel

Applications move applied → shortlisted → interview → offer. The table stores only the current status, so map each status to the furthest stage reached and count how many reached at least each stage. (Simplification, said aloud: we treat rejected as “stopped at applied” because we do not know at which stage the rejection happened. A real events table with one row per stage change would remove this assumption.)

```
WITH stage AS (
  SELECT app_id,
         CASE status WHEN 'applied' THEN 1 WHEN 'rejected' THEN 1
                   WHEN 'shortlisted' THEN 2 WHEN 'interview' THEN 3
                   WHEN 'offer' THEN 4 END AS reached
  FROM applications
  WHERE app_id <> 22
)
SELECT SUM(reached >= 1) AS applied,
       SUM(reached >= 2) AS shortlisted,
       SUM(reached >= 3) AS interviewed,
       SUM(reached >= 4) AS offered,
       ROUND(1.0 * SUM(reached >= 2) / SUM(reached >= 1), 3) AS apply_to_short,
       ROUND(1.0 * SUM(reached >= 4) / SUM(reached >= 3), 3) AS interview_to_offer
FROM stage;
```

Output (real run):

applied	shortlisted	interviewed	offered	apply_to_short	interview_to_offer
22	12	8	4	0.545	0.5

(1 row)

In SQLite a comparison is 1 or 0, so SUM(reached >= 2) counts rows; elsewhere write SUM(CASE WHEN reached >= 2 THEN 1 ELSE 0 END). 22 applications (duplicate removed), 12 reached shortlisting (54.5%), 8 interviews, 4 offers: half of interviews convert.

48.7.2 A/B test click-through rate

```
SELECT variant,
       COUNT(*) AS impressions,
       SUM(clicked) AS clicks,
       ROUND(AVG(clicked * 1.0), 3) AS ctr
FROM impressions
GROUP BY variant;
```

Output (real run):

variant	impressions	clicks	ctr
A	8	2	0.25
B	8	5	0.625

(2 rows)

Variant B (new ranking) has CTR 0.625 vs 0.25. Is that real? A two-proportion z-test (Chapter 39): pooled $p = 7/16 = 0.4375$, standard error $\sqrt{p(1-p)(1/8 + 1/8)} = 0.248$, $z = (0.625 - 0.25)/0.248 = 1.51$,

two-sided p-value 0.13: **not significant** with 8 impressions per arm. SQL computes the counts; the decision needs statistics. Also check the split by day:

```
SELECT shown_date, variant, COUNT(*) AS imps, SUM(clicked) AS clicks,
       ROUND(1.0 * SUM(clicked) / COUNT(*), 2) AS ctr
FROM impressions
GROUP BY shown_date, variant
ORDER BY shown_date, variant;
```

Output (real run):

shown_date	variant	imps	clicks	ctr
2026-03-01	A	2	1	0.5
2026-03-01	B	2	1	0.5
2026-03-02	A	3	1	0.33
2026-03-02	B	2	2	1
2026-03-03	A	3	0	0
2026-03-03	B	4	2	0.5

(6 rows)

B wins or ties every day, which is reassuring; when a variant wins in every segment but loses overall (or vice versa) you are looking at Simpson's paradox, caused by unequal segment sizes between arms.

48.8 Other patterns worth knowing

48.8.1 Median without a MEDIAN function

Number the sorted values; the median is the middle row (odd n) or the average of the two middle rows (even n). With integer division, $(n+1)/2$ and $(n+2)/2$ are the same row when n is odd and the two middle rows when n is even.

```
WITH s AS (
  SELECT ctc_lpa,
         ROW_NUMBER() OVER (ORDER BY ctc_lpa) AS rn,
         COUNT(*)    OVER ()           AS n
  FROM candidates
  WHERE ctc_lpa IS NOT NULL
)
SELECT AVG(ctc_lpa) AS median_ctc
FROM s
WHERE rn IN ((n + 1) / 2, (n + 2) / 2);
```

Output (real run):

median_ctc
15

(1 row)

9 known CTCs; the 5th smallest is 15 (Ananya). PostgreSQL: PERCENTILE_CONT(0.5) WITHIN GROUP (ORDER BY ctc_lpa).

48.8.2 Pivot: rows to columns

```
SELECT co.name AS company,
       SUM(CASE WHEN j.title = 'Data Scientist' THEN 1 ELSE 0 END) AS data_scientist,
```

```

SUM(CASE WHEN j.title = 'Data Analyst' THEN 1 ELSE 0 END) AS data_analyst,
SUM(CASE WHEN j.title = 'ML Engineer' THEN 1 ELSE 0 END) AS ml_engineer
FROM companies co
LEFT JOIN jobs j ON j.company_id = co.company_id
GROUP BY co.company_id, co.name
ORDER BY co.company_id;

```

Output (real run):

company	data_scientist	data_analyst	ml_engineer
DataNest	1	0	1
PayWave	1	1	0
MediCore	0	1	1
ShopKart	1	1	0
EduSpark	1	0	1
GreenGrid	0	0	0

(6 rows)

Conditional aggregation is the portable pivot (pandas: `pd.crosstab(jobs.company, jobs.title)`).

48.8.3 First event per entity

```

SELECT cand_id, job_id, applied_date
FROM (
    SELECT a.*, ROW_NUMBER() OVER (PARTITION BY cand_id ORDER BY applied_date, app_id) AS rn
    FROM applications a
) t
WHERE rn = 1
ORDER BY cand_id;

```

Output (real run):

cand_id	job_id	applied_date
101	201	2026-01-16
102	201	2026-01-18
103	203	2026-01-27
104	202	2026-01-21
105	201	2026-02-02
107	206	2026-02-16
108	203	2026-02-21
109	201	2026-03-03

(8 rows)

48.8.4 Time between two events

```

SELECT j.title,
    ROUND(AVG(julianday(a.applied_date) - julianday(j.posted_date)), 2) AS
    avg_days_to_apply,
    COUNT(*) AS n_apps
FROM applications a
JOIN jobs j ON j.job_id = a.job_id
WHERE a.app_id <> 22
GROUP BY j.title
ORDER BY avg_days_to_apply;

```

Output (real run):

title	avg_days_to_apply	n_apps
-------	-------------------	--------

```

-----+-----+-----
ML Engineer | 1          | 3
Data Analyst | 5.88       | 8
Data Scientist | 8.73      | 11
(3 rows)

```

Candidates apply to ML Engineer postings within a day on average, but take almost 9 days for Data Scientist postings, which could mean DS postings are less visible or attract more deliberate applicants: a question for the product team, not a conclusion.

48.8.5 “Applied to all” (relational division)

“Candidates who applied to every Data Analyst job”: count each candidate’s distinct DA jobs and compare with the total number of DA jobs.

```

SELECT a.cand_id
FROM applications a
JOIN jobs j ON j.job_id = a.job_id
WHERE j.title = 'Data Analyst'
GROUP BY a.cand_id
HAVING COUNT(DISTINCT a.job_id) = (SELECT COUNT(*) FROM jobs WHERE title = 'Data Analyst');

```

Output (real run):

```

cand_id
-----
103
(1 row)

```

Kabir (103) applied to all three (203, 205, 207).

48.8.6 Non-equi join as a data-quality check

Candidates whose current CTC already exceeds the job’s maximum (likely to reject an offer, or a recommendation the ranker should not have made):

```

SELECT ca.name, ca.ctc_lpa, j.job_id, j.max_ctc_lpa
FROM applications a
JOIN candidates ca ON ca.cand_id = a.cand_id
JOIN jobs j ON j.job_id = a.job_id
WHERE ca.ctc_lpa > j.max_ctc_lpa
ORDER BY ca.name;

```

Output (real run):

```

name | ctc_lpa | job_id | max_ctc_lpa
-----+-----+-----+-----
Ananya | 15      | 203    | 12
Diya   | 18      | 203    | 12
Rohan  | 12      | 205    | 8
Vikram | 42      | 206    | 40
(4 rows)

```

48.9 SQL for machine learning: feature tables without leakage

In industry, the training table for a model is usually built in SQL: one row per entity (candidate), one column per feature, one label column. Two mistakes ruin such tables: fan-out and leakage.

48.9.1 Fan-out with two one-to-many tables

Joining candidates to applications *and* logins in one go multiplies rows: candidate 101 has 4 application rows and 5 login rows, so the join creates $4 \times 5 = 20$ rows and every count is inflated:

```
SELECT ca.cand_id,
       SUM(CASE WHEN a.status = 'rejected' THEN 1 ELSE 0 END) AS n_rejected_wrong,
       COUNT(l.login_date) AS n_logins_wrong
FROM candidates ca
LEFT JOIN applications a ON a.cand_id = ca.cand_id
LEFT JOIN logins l      ON l.cand_id = ca.cand_id
WHERE ca.cand_id IN (101, 107, 109)
GROUP BY ca.cand_id
ORDER BY ca.cand_id;
```

Output (real run):

cand_id	n_rejected_wrong	n_logins_wrong
101	0	20
107	3	6
109	3	9

(3 rows)

Correct: aggregate each source to one row per candidate in its own CTE, then LEFT JOIN the summaries:

```
WITH apps AS (
    SELECT cand_id,
           COUNT(DISTINCT job_id) AS n_jobs_applied,
           SUM(CASE WHEN status = 'rejected' THEN 1 ELSE 0 END) AS n_rejected,
           MAX(CASE WHEN status = 'offer' THEN 1 ELSE 0 END) AS got_offer
    FROM (SELECT DISTINCT cand_id, job_id, applied_date, status FROM applications) a
    GROUP BY cand_id
),
act AS (
    SELECT cand_id, COUNT(DISTINCT login_date) AS n_login_days
    FROM logins
    GROUP BY cand_id
)
SELECT ca.cand_id,
       ca.exp_yrs,
       COALESCE(ca.ctc_lpa, 0) AS ctc_lpa,
       CASE WHEN ca.ctc_lpa IS NULL THEN 1 ELSE 0 END AS ctc_missing,
       COALESCE(ap.n_jobs_applied, 0) AS n_jobs_applied,
       COALESCE(ap.n_rejected, 0) AS n_rejected,
       COALESCE(ac.n_login_days, 0) AS n_login_days,
       COALESCE(ap.got_offer, 0) AS got_offer
FROM candidates ca
LEFT JOIN apps ap ON ap.cand_id = ca.cand_id
LEFT JOIN act ac ON ac.cand_id = ca.cand_id
ORDER BY ca.cand_id;
```

Output (real run):

cand_id	exp_yrs	ctc_lpa	ctc_missing	n_jobs_applied	n_rejected	n_login_days	got_offer
101	2	8	0	3	0	5	1
102	5	18	0	3	1	2	1
103	0	0	1	3	1	0	0
104	7	30	0	3	0	4	1
105	3	12	0	3	1	1	0
106	1	5.5	0	0	0	0	0
107	10	42	0	2	1	3	1
108	4	15	0	2	0	0	0
109	2	9.5	0	3	1	2	0
110	6	24	0	0	0	0	0

(10 rows)

Every candidate appears exactly once, including those with no applications or logins (filled with 0), and ctc_missing keeps the information that the fresher had no salary.

48.9.2 Leakage: point-in-time correctness

The feature table above uses *all* history, including events after the moment a prediction would have been made. To train “will this candidate get an offer?”, features must use only data before a cutoff date, and the label must come from after it:

```
WITH params AS (SELECT '2026-02-15' AS cutoff),
feat AS (
    SELECT a.cand_id,
           COUNT(*) AS apps_before,
           SUM(CASE WHEN a.status = 'rejected' THEN 1 ELSE 0 END) AS rejections_before
    FROM applications a, params p
    WHERE a.applied_date < p.cutoff
    GROUP BY a.cand_id
),
label AS (
    SELECT a.cand_id,
           MAX(CASE WHEN a.status = 'offer' THEN 1 ELSE 0 END) AS offer_after
    FROM applications a, params p
    WHERE a.applied_date >= p.cutoff
    GROUP BY a.cand_id
)
SELECT c.cand_id,
       COALESCE(f.apps_before, 0) AS apps_before,
       COALESCE(f.rejections_before, 0) AS rejections_before,
       COALESCE(l.offer_after, 0) AS label_offer_after
FROM candidates c
LEFT JOIN feat f ON f.cand_id = c.cand_id
LEFT JOIN label l ON l.cand_id = c.cand_id
WHERE c.signup_date < (SELECT cutoff FROM params)
ORDER BY c.cand_id;
```

Output (real run):

cand_id	apps_before	rejections_before	label_offer_after
101	2	0	1
102	3	1	0
103	2	1	0
104	2	0	0

```

105      | 3          | 1          | 0
106      | 0          | 0          | 0
(6 rows)

```

Only candidates who had signed up before the cutoff are included (an entity must exist at prediction time). Aarav had 2 applications before 15 February and got an offer after it (label 1). In production, repeat this for many cutoffs (one snapshot per week) to build a larger training set, and validate on later cutoffs than you train on (time-based split, Chapter 37).

⚠ Common mistake: features that leak

got_offer computed over all time, “number of interviews” (which only happens after a shortlist), “days since last login” computed at query time rather than at the cutoff, and global target encodings (offer rate per company over all data) all leak future information. The model looks perfect offline and fails in production: the same lesson as salary in the placement data (Chapter 42).

48.10 Performance: indexes and query plans

48.10.1 What an index is

Without an index, finding candidate 105’s applications means reading every row (a *full table scan*), like finding a name in an unsorted pile of forms. An **index** is a separate, sorted structure (usually a B-tree: a balanced tree whose leaves hold the key values in order with pointers to the rows), like the alphabetical index at the back of a book. A lookup takes $O(\log n)$ steps instead of $O(n)$: about 30 steps for a billion rows.

EXPLAIN QUERY PLAN

```
SELECT * FROM applications WHERE cand_id = 105;
```

Output (real run):

```
QUERY PLAN: SCAN applications
```

```
CREATE INDEX idx_app_cand ON applications(cand_id);
```

EXPLAIN QUERY PLAN

```
SELECT * FROM applications WHERE cand_id = 105;
```

Output (real run):

```
QUERY PLAN: SEARCH applications USING INDEX idx_app_cand (cand_id=?)
```

SCAN = read everything; SEARCH ... USING INDEX = jump straight to the matching rows.

48.10.2 Sargable filters

An index on applied_date is sorted by the raw date. Wrapping the column in a function hides that order from the optimiser, so it scans:

```
CREATE INDEX idx_app_date ON applications(applied_date);
```

EXPLAIN QUERY PLAN

```
SELECT * FROM applications WHERE strftime('%Y', applied_date) = '2026';
```

Output (real run):

```
QUERY PLAN: SCAN applications
```

Rewrite as a range on the raw column (“Search ARGument-able”):

```
CREATE INDEX idx_app_date ON applications(applied_date);
EXPLAIN QUERY PLAN
SELECT * FROM applications WHERE applied_date >= '2026-01-01' AND applied_date < '
2027-01-01';
```

Output (real run):

```
| QUERY PLAN: SEARCH applications USING INDEX idx_app_date (applied_date>? AND applied_date<?)
```

Other non-sargable forms: WHERE salary * 12 > 600, LIKE '%data' (leading wildcard), WHERE LOWER(email) = ... (unless there is an index on LOWER(email)), implicit type casts (comparing a text column to a number).

48.10.3 Rules of thumb

- Index columns used in WHERE, JOIN ON, ORDER BY and GROUP BY of frequent queries; foreign keys almost always.
- A **composite index** on (cand_id, applied_date) serves filters on cand_id alone or on both, but not on applied_date alone (*leftmost-prefix rule*), like a phone book sorted by surname then first name.
- A **covering index** contains every column a query needs, so the table itself is never read.
- Indexes cost: every INSERT/UPDATE must update them, and they use disk. Low-cardinality columns (gender) rarely benefit.
- Read the plan (EXPLAIN / EXPLAIN ANALYZE) instead of guessing; select only needed columns; filter early; avoid SELECT DISTINCT as a fix for fan-out; in warehouses, filter on partition columns (e.g. date) so whole partitions are skipped.

48.11 Schema design: normalisation, OLTP vs OLAP, star schema

48.11.1 Normalisation in three steps

Start from one wide sheet: (app_id, candidate_name, candidate_city, skills, job_title, company_name, company_city, status).

- **1NF** (atomic values): skills = 'python, sql' holds a list in one cell; move skills to their own table, one row per (candidate, skill).
- **2NF** (no partial dependency): in a table keyed by (cand_id, job_id), candidate_city depends on cand_id alone; move it to candidates.
- **3NF** (no transitive dependency): company_city depends on company, which depends on job; move it to companies.

The result is exactly MiniNaukri's design: each fact stored once, so updates cannot create contradictions (*update anomalies*).

48.11.2 OLTP vs OLAP

	OLTP (the product database)	OLAP (the analytics warehouse)
Workload	Many small reads and writes (apply to a job)	Few large scans and aggregations (applications per city per month)
Design	Normalised (3NF), row-oriented	Denormalised star/snowflake schema, column-oriented
Examples	MySQL, PostgreSQL	Redshift, BigQuery, Snowflake, Hive, ClickHouse
Priorities	ACID, low latency	Throughput, compression, scanning only needed columns

A **star schema** has a central *fact* table (one row per event: an application, with numeric measures and foreign keys) surrounded by *dimension* tables (candidate, job, company, date) that describe it. Analysts join the fact to a few dimensions and aggregate. Data scientists usually query the warehouse, never the production OLTP database.

48.12 SQL ↔ pandas translation

SQL	pandas
SELECT a, b FROM t WHERE x > 5 ORDER BY a DESC LIMIT 3	t.loc[t.x > 5, ['a', 'b']] t.sort_values('a', ascending=False).head(3) or t.nlargest(3, 'a')
SELECT DISTINCT a GROUP BY g with COUNT(*), AVG(x) HAVING COUNT(*) >= 3 LEFT JOIN u ON t.k = u.k UNION ALL / UNION CASE WHEN COALESCE(x, 0) AVG(x) OVER (PARTITION BY g) ROW_NUMBER() OVER (PARTITION BY g ORDER BY d) DENSE_RANK() OVER (ORDER BY x DESC) LAG(x) OVER (PARTITION BY g ORDER BY d) SUM(x) OVER (ORDER BY d) Pivot with SUM(CASE ...)	t['a'].drop_duplicates() / t['a'].unique() t.groupby('g').agg(n=('x', 'size'), avg=('x', 'mean')) .query('n >= 3') after the groupby t.merge(u, on='k', how='left') pd.concat([t, u]) / .drop_duplicates() np.where, np.select, pd.cut t.x.fillna(0) t.groupby('g').x.transform('mean') t.sort_values('d').groupby('g').cumcount() + 1 t.x.rank(method='dense', ascending=False) t.sort_values('d').groupby('g').x.shift() t.sort_values('d').x.cumsum() pd.pivot_table / pd.crosstab

One semantic difference to mention: pandas' groupby drops NaN keys by default (dropna=False keeps them), whereas SQL's GROUP BY keeps NULL as its own group.

48.13 Interview questions: Level 4

Q48.1

✔ **Most asked** Find the second (and the Nth) highest salary. What if there are ties or fewer than N salaries?

🕒 **In one line.** DENSE_RANK() OVER (ORDER BY salary DESC) and keep rank N; ties are handled because DENSE_RANK ranks distinct values; if fewer than N distinct values exist, return NULL by wrapping the query as a scalar subquery.

⚙️ **The full answer.** Ranking distinct values is exactly what “Nth highest salary” means. `MAX(salary) WHERE salary < (SELECT MAX(salary) ...)` works for $N = 2$ only. `LIMIT 1 OFFSET N-1` needs `DISTINCT`. The correlated version (count distinct greater values = $N-1$) works on any SQL, at $O(n^2)$ cost.

📊 **Worked example.** Here: second highest = 75 (`L4_second_max`); third = 70 by all three methods (`L4_nth_dense`, `L4_nth_limit`, `L4_nth_correlated`).

```
SELECT (SELECT DISTINCT salary FROM
        (SELECT salary, DENSE_RANK() OVER (ORDER BY salary DESC) AS dr
         FROM employees WHERE salary IS NOT NULL) t
       WHERE dr = 20) AS nth_salary;  -- returns one row with NULL
```

💬 **If the interviewer pushes.** Per department? Add `PARTITION BY dept`. Without window functions or `LIMIT`? The correlated count. What about the `NULL` salary? Exclude it explicitly; in a `DESC` sort SQLite puts `NULL` last, PostgreSQL puts `NULL` first, and `RANK` would then give `NULL` rank 1 in PostgreSQL.

⚠️ **Trap to avoid.** `ROW_NUMBER` or `LIMIT` without `DISTINCT`.

Q48.2

✅ **Most asked** Find and delete duplicate rows, keeping one copy.

🎯 **In one line.** Define the business key, find duplicates with `GROUP BY key HAVING COUNT(*) > 1`, label copies with `ROW_NUMBER() OVER (PARTITION BY key ORDER BY id)`, and delete rows with `rn > 1` (or keep `MIN(id)` per key).

⚙️ **The full answer.** Duplicates almost never share the surrogate id; they share the business key (candidate, job, date). `ROW_NUMBER`'s `ORDER BY` decides which copy is kept (the first-inserted, the latest-updated, the most complete).

📊 **Worked example.** `L4_dupes_find` finds (101, 201, 2026-01-16) twice; `L4_dupes_delete` keeps `MIN(app_id)` and leaves 22 rows. In PostgreSQL:

```
DELETE FROM applications WHERE app_id IN (
  SELECT app_id FROM (
    SELECT app_id, ROW_NUMBER() OVER (PARTITION BY cand_id, job_id, applied_date
                                     ORDER BY app_id) AS rn
    FROM applications) t
  WHERE rn > 1);
```

💬 **If the interviewer pushes.** How do you prevent it? A `UNIQUE` constraint on the business key, or idempotent inserts (`INSERT ... ON CONFLICT DO NOTHING`). Table has no id at all? Use the database's row id (`ctid`, `rowid`) or rebuild with `CREATE TABLE ... AS SELECT DISTINCT`.

⚠️ **Trap to avoid.** Deleting all copies (`WHERE key IN (duplicates)`), including the one to keep.

Q48.3

✅ **Most asked** Find users who logged in on at least 3 consecutive days.

🕒 **In one line.** Deduplicate (user, date), subtract ROW_NUMBER (in days) from the date to get a constant per run, group by user and that constant, keep groups with count ≥ 3 .

⚙️ **The full answer.** Within a run of consecutive dates, both the date and the row number increase by 1 per row, so their difference is constant; any gap increases the date by more than the row number and starts a new group.

📅 **Worked example.** L4_streak_steps shows Aarav's dates mapping to 28 Feb, 28 Feb, 28 Feb, 1 Mar, 1 Mar; L4_streaks returns 104 (4 days), 101 (3) and 107 (3).

💡 **If the interviewer pushes.** *Alternative?* LAG: flag rows where $\text{date} - \text{LAG}(\text{date}) > 1$ as a new streak, then a running SUM of the flags gives the streak id. *Longest streak per user?* MAX(streak_len) per user over the grouped result. *Consecutive weeks?* Use week numbers instead of dates.

⚠️ **Trap to avoid.** Forgetting to remove same-day duplicates.

Q48.4

✅ **Most asked** Compute 7-day retention: of the users who signed up on each day, what fraction were active 7 days later?

🕒 **In one line.** Cohort table (user, signup_date) LEFT JOIN activity on user = user AND activity_date = signup_date + 7, group by signup date, COUNT(activity.user) / COUNT(*).

⚙️ **The full answer.** The LEFT JOIN keeps users who were not retained (they count in the denominator); the join condition defines "retained". Variants: "active on day 7" (exact) or "active any time in days 1–7" (BETWEEN).

📅 **Worked example.** The monthly-cohort version on MiniNaukri, L4_cohort: January 0.75, February 0.5, March 0.5.

```
SELECT c.signup_date, COUNT(*) AS cohort,
       COUNT(DISTINCT l.cand_id) AS retained_d7,
       ROUND(1.0 * COUNT(DISTINCT l.cand_id) / COUNT(*), 2) AS d7_retention
FROM candidates c
LEFT JOIN logins l ON l.cand_id = c.cand_id
                  AND l.login_date = date(c.signup_date, '+7 days')
GROUP BY c.signup_date;
```

On MiniNaukri this returns 0 for every signup day, because the login log only covers 1–6 March and no candidate's day 7 falls in it: a reminder to check that the activity data actually covers each cohort's window.

💡 **If the interviewer pushes.** *Why might the latest cohorts look bad?* They have not had 7 days yet (right-censoring); exclude cohorts younger than the window. *Why COUNT(DISTINCT) in the numerator?* A user can log in twice on day 7.

⚠️ **Trap to avoid.** Putting the activity condition in WHERE (drops non-retained users, giving 100% retention).

Q48.5

✓ **Most asked** Top 3 jobs by number of applications in each city.

🕒 **In one line.** Count applications per job (LEFT JOIN so zero-application jobs count), rank with DENSE_RANK() OVER (PARTITION BY city ORDER BY n DESC), keep rank ≤ 3.

📊 **Worked example.**

```
WITH n AS (
  SELECT j.job_id, j.city, COUNT(a.app_id) AS n_apps
  FROM jobs j LEFT JOIN applications a ON a.job_id = j.job_id
  GROUP BY j.job_id, j.city
), r AS (
  SELECT *, DENSE_RANK() OVER (PARTITION BY city ORDER BY n_apps DESC) AS dr FROM n
)
SELECT city, job_id, n_apps FROM r WHERE dr <= 3 ORDER BY city, n_apps DESC;
```

⚙️ **The full answer.** Grain: one row per job first, then rank within city. Bengaluru has 4 jobs (201: 5 including the duplicate, 207: 3, 206: 2, 202: 1), so 202 is cut.

💬 **If the interviewer pushes.** Ties? State whether you want DENSE_RANK (all ties) or ROW_NUMBER (exactly 3).

⚠️ **Trap to avoid.** Ranking before aggregating (ranking individual application rows).

Q48.6

Build a hiring funnel from an application_events table (one row per status change). How is that better than a status column?

🕒 **In one line.** Count distinct applications that reached each stage (COUNT(DISTINCT CASE WHEN stage = 'interview' THEN app_id END)), divide consecutive stages; an events table records the path, so “rejected after interview” is distinguishable from “rejected at screening”.

⚙️ **The full answer.** A single status column overwrites history; MiniNaukri’s funnel had to assume rejected applications stopped at stage 1 (L4_funnel). With events you also get time between stages (median days to shortlist) and where in the funnel rejections happen.

📊 **Worked example.** With the current-status approximation: 22 applied, 12 shortlisted (54.5%), 8 interviews, 4 offers (50% of interviews).

💬 **If the interviewer pushes.** Ordered funnel (must pass stage 2 after stage 1)? Compare per-application timestamps of each stage with MIN() per stage and require increasing times.

⚠️ **Trap to avoid.** Counting events instead of distinct applications (a stage can be logged twice).

Q48.7

Your query computing applications per day is slow on a 2-billion-row table. How do you investigate?

🕒 **In one line.** Read the plan with EXPLAIN (ANALYZE); check that the date filter is sargable

and hits a partition or index; select only needed columns; pre-aggregate (materialised daily table); check statistics and join strategy.

⚙️ **The full answer.** Typical culprits: WHERE DATE(applied_at) = ... (function on the column defeats indexes and partition pruning), SELECT * on a columnar warehouse (reads every column), joining before filtering, fan-out joins, COUNT(DISTINCT) on huge cardinalities (approximate with APPROX_COUNT_DISTINCT/HyperLogLog when exactness is not needed).

📊 **Worked example.** L4_explain_nonsarg shows SCAN for strftime('%Y', applied_date) = '2026'; the half-open range in L4_explain_sarg uses the index.

💬 **If the interviewer pushes.** Composite index for “applications of candidate X in March”? (cand_id, applied_date): equality column first, range column second. Why not index everything? Writes slow down and storage grows.

⚠️ **Trap to avoid.** Guessing optimisations without looking at the plan.

Q48.8

What is normalisation? When would you deliberately denormalise?

🎯 **In one line.** Organising tables so each fact is stored once (1NF atomic values, 2NF no partial dependency on part of a key, 3NF no transitive dependency), which prevents update anomalies; denormalise in analytics warehouses and read-heavy paths, where fewer joins matter more than update cost.

⚙️ **The full answer.** A normalised OLTP schema makes writes safe and small. Analytical queries then need many joins; a warehouse star schema (wide fact + dimensions) or a flat feature table trades redundancy for speed and simplicity. Updates there come from batch pipelines, not users, so anomalies are controlled.

📊 **Worked example.** The single wide application sheet with skills = 'python, sql' and company city repeated per row becomes MiniNaukri's 4 tables plus a skills table.

💬 **If the interviewer pushes.** What is BCNF? A stricter 3NF: every determinant is a candidate key. Star vs snowflake? Snowflake normalises the dimensions further (company → industry table).

⚠️ **Trap to avoid.** Claiming normalised is always better.

Q48.9

🧠 **Think deeper** You are asked for “average applications per active candidate per day”. Write down three different numbers this could mean, and the SQL for the one you would pick.

🎯 **In one line.** (1) total applications / total candidate-days with activity; (2) the mean over days of (applications that day / active candidates that day); (3) the mean over candidates of their own daily average; they differ because of weighting, and the choice must be agreed with the stakeholder.

⚙️ **The full answer.** (1) weights busy days more; (2) gives each day equal weight (an average of ratios); (3) gives each candidate equal weight, so a candidate active one day counts as much as

one active every day. “Active” also needs a definition: logged in, or applied? Each choice is a different metric.

 **Worked example.** Definition (2), with active = logged in:

```
WITH a AS (SELECT applied_date AS d, COUNT(*) AS apps FROM applications GROUP BY 1),
    u AS (SELECT login_date AS d, COUNT(DISTINCT cand_id) AS users FROM logins GROUP BY 1)
SELECT AVG(1.0 * COALESCE(a.apps, 0) / u.users) AS avg_apps_per_active_user
FROM u LEFT JOIN a ON a.d = u.d;
```

 **If the interviewer pushes.** Which would you report? The one matching the decision: capacity planning wants (1); a product health dashboard usually wants (2) with the numerator and denominator shown. Why show both parts? So that a change in the ratio can be traced to a change in applications or in users.

 **Trap to avoid.** Writing SQL before defining the metric.

Q48.10

 **Think deeper** You build a training table in SQL for “will this candidate get an offer in the next 30 days?” and the offline AUC is 0.99. What do you check first?

 **In one line.** Leakage: every feature must be computed only from data before the prediction cutoff, the label only from after it, and the population must be the candidates who existed (and were eligible) at the cutoff.

 **The full answer.** Features computed over all history (“number of interviews”, “has offer”, “days since last login” computed today) contain the answer. Global aggregates (company offer rate over all time) leak too. A random train/test split of snapshots also leaks, because the same candidate’s later snapshots inform earlier ones; split by time.

 **Worked example.** Q_point_in_time: features from applications strictly before 15 February, label from on or after it, only candidates who signed up before the cutoff. Aarav: 2 earlier applications, label 1 (offer on 19 February).

 **If the interviewer pushes.** How do feature stores help? They store timestamped feature values and do “as-of” joins automatically. Sanity checks? Feature importance dominated by one suspicious feature; performance collapsing on a later time-split; features whose values change for the past when recomputed.

 **Trap to avoid.** Celebrating the AUC.

Q48.11

 **Think deeper** Your A/B test query shows variant B with CTR 0.625 vs A with 0.25. The product manager wants to ship B today. What do you say?

 **In one line.** The counts are tiny (8 impressions per arm): a two-proportion z-test gives $z = 1.51$, $p = 0.13$, so the difference is not significant; we also need to check randomisation balance, the unit of analysis (users, not impressions) and novelty effects before shipping.

⚙️ **The full answer.** SQL gives counts; inference needs the sampling variability. Impressions from the same user are correlated, so the effective sample size is the number of users and the variance should be computed per user (or by the delta method). A sample-ratio-mismatch check (are arms the expected size?) catches broken assignment. Novelty: users click new things at first.

📊 **Worked example.** Pooled $p = 7/16$; $SE = \sqrt{0.4375 \times 0.5625 \times (1/8 + 1/8)} = 0.248$; $z = 0.375/0.248 = 1.51$; two-sided $p = 0.13$.

💬 **If the interviewer pushes.** *How many impressions would you need?* For 80% power to detect $0.25 \rightarrow 0.35$ at $\alpha = 0.05$: about $n \approx (1.96 + 0.84)^2 [0.25 \cdot 0.75 + 0.35 \cdot 0.65] / 0.1^2 \approx 325$ per arm. *Daily results?* L4_ctr_daily: B is never worse on any day, encouraging but not proof.

⚠️ **Trap to avoid.** Treating a SQL ratio as a conclusion.

Q48.12

💡 **Think deeper** Why is `WHERE YEAR(applied_at) = 2026` slower than `WHERE applied_at >= '2026-01-01' AND applied_at < '2027-01-01'` when there is an index on `applied_at`?

🎯 **In one line.** The index stores raw `applied_at` values in sorted order; the optimiser can binary-search a range on the raw column, but it cannot search for values of a function of the column without computing the function for every row, so it scans.

⚙️ **The full answer.** A B-tree supports lookups on its key's order. `YEAR(x) = 2026` is equivalent to a range, but the optimiser does not invert arbitrary functions. Expression indexes (`CREATE INDEX ... ON t (YEAR(applied_at))`) or partitioning on the year are alternatives.

📊 **Worked example.** SQLite plans: `SCAN applications` for the strftime filter vs `SEARCH applications USING INDEX idx_app_date (applied_date>? AND applied_date<?)` for the range (`L4_explain_nonsarg, L4_explain_sarg`).

💬 **If the interviewer pushes.** *Same issue elsewhere?* `WHERE amount + 0 = 5`, `WHERE CAST(id AS TEXT) = '5'`, `LIKE '%x'`, and comparing a VARCHAR column with an integer literal (implicit cast on the column).

⚠️ **Trap to avoid.** Answering “functions are slow”; the real reason is losing the index order.

Last-Week Checklist

Can you do these on paper, aloud, in under three minutes each?

- Draw the error-vs-complexity and error-vs-iterations curves; explain bias–variance with the formula and one number example.
- Derive OLS (one feature) and the normal equations; list linear-regression assumptions and the fix for each violation, including a GLM for heteroscedasticity.
- Derive the logistic-regression gradient $(p - y)x$; interpret a coefficient as an odds ratio.
- Explain GLMs (link + distribution); Poisson regression for counts; SVR's ϵ -tube; quantile regression's pinball loss.
- Do one gradient-descent, momentum and Adam step by hand.
- Compute entropy, Gini and information gain; write decision-tree pseudo-code; explain pruning.
- Explain random forests (bootstrap, feature subsampling, variance formula, OOB) and RF vs GBM.
- Do one gradient-boosting round by hand; write XGBoost's leaf weight and gain; explain LightGBM's histogram, leaf-wise growth, GOSS, EFB.
- Derive the SVM margin; explain hinge loss, support vectors, the dual and the kernel trick with a worked kernel.
- Build Naive Bayes from Bayes' theorem with smoothing; derive Fisher's LDA direction.
- Compute precision, recall, F1, specificity, AUC (pair counting) and choose metrics for spam, fraud and telecalling.
- Derive MLE for Bernoulli and Gaussian; explain MAP as regularisation; Beta smoothing of CTR.
- Explain MCAR/MAR/MNAR and imputation choices; SMOTE; why not to drop genuine outliers.
- Explain PCA via SVD, the scree plot, t-SNE perplexity; matrix factorisation for recommendations.
- Run K-Means and one EM step by hand; KDE bandwidth.
- Compute TF-IDF; compare BoW, Word2Vec, FastText and BERT; explain negative sampling and MLM.
- Backpropagate through a two-layer network by hand; explain vanishing/exploding gradients and the fixes; Xavier vs He; dropout; BatchNorm.
- CNN output size and parameter counts; invariance vs equivariance.
- Draw an LSTM cell with equations; explain why it does not vanish; GRU differences; time-series modelling plan.
- Compute scaled dot-product attention by hand; explain every part of a Transformer block.
- GAN objective and optimal discriminator; mode collapse; VAE vs autoencoder.
- Explain RAG, LoRA, SFT/DPO and agents with guardrails at a basic level.
- In a notebook: load, clean names, missing values, describe, value counts, groupby, plots, leakage check, pipeline, split, fit, confusion matrix, report, CV.
- Compute mean, median, mode, variance (population and sample) and read skewness and kurtosis; read a box plot, histogram and Q–Q plot.

- Compute Pearson, Spearman (ranks), point-biserial and phi correlations; read a correlation matrix; $VIF = 1/(1 - R^2)$.
- Apply the 68–95–99.7 rule, z-scores and the IQR rule; choose an imputation (including model-based) and spot leaky variables.
- Compare standardisation, min–max and robust scaling; Box–Cox; one-hot vs ordinal vs target vs frequency encoding; the dummy variable trap; Borderline-SMOTE and ADASYN.
- Explain k-medoids, k-modes, k-prototypes, soft vs hard clustering, the Swiss roll, SNE vs t-SNE, linear vs non-linear and deterministic vs stochastic DR.
- Run the perceptron rule by hand; IoU; a Q-learning update; FedAvg.
- Explain stationarity, ACF/PACF, ARIMA/SARIMAX; compute a one-way ANOVA.
- Use Bayes' theorem, PMF/PDF, expectation, variance and the common distributions with examples; explain CLT, CI, p-value vs alpha, paired t-test, A/B test.
- Compute rank, null space, eigenvalues and a Markov chain steady state.
- SQL: explain the logical order of a query; NULL three-valued logic ($=$ NULL, $<>$ dropping NULLs, NOT IN with NULL); COUNT(*) vs COUNT(col).
- SQL: predict row counts of INNER/LEFT/FULL/CROSS joins with duplicate keys; fix the LEFT JOIN + WHERE trap and fan-out; WHERE vs HAVING; conditional aggregation for rates.
- SQL: ROW_NUMBER vs RANK vs DENSE_RANK; Nth highest and top-N per group; LAG/LEAD; running totals and frames; CTEs and recursive CTEs.
- SQL: write duplicates, consecutive-days streaks, retention, funnel, MoM growth, median and pivot queries from memory; build a point-in-time feature table; explain indexes, sargability, normalisation, OLTP vs OLAP, ACID.
- Walk through two case studies with the ten-step skeleton; answer “why InfoEdge” specifically.

Book statistics

This edition contains **717** interview questions with full answers, of which **110** are reported InfoEdge questions (PYQ), **17** are “Most asked” standard SQL questions, and **184** are “Think deeper” questions (the thought-provoking or unconventional kind you asked to be at least 15% of the total).